

Realtime User Analytics with ClickHouse

Amir Vaza

Overview

- Motivation
- Appsee's pre-ClickHouse Data Model
- Requirements
- ClickHouse
- Future Work



Motivation

Why would we replace something that works?

Motivation

The motivation to replace our existing reliable, battle tested model comes from the product growth, allowing richer analytics features such as:

- Realtime dynamic slicing and dicing the data
- Realtime Funnels and Cohorts

Pre-ClickHouse Data Model

Appsee's Pre-ClickHouse Data Model



Aggregation is performed upon incoming data processing to allow an interactive dashboard experience.



Pre-aggregated segmentation is based on a set of fixed dimensions:

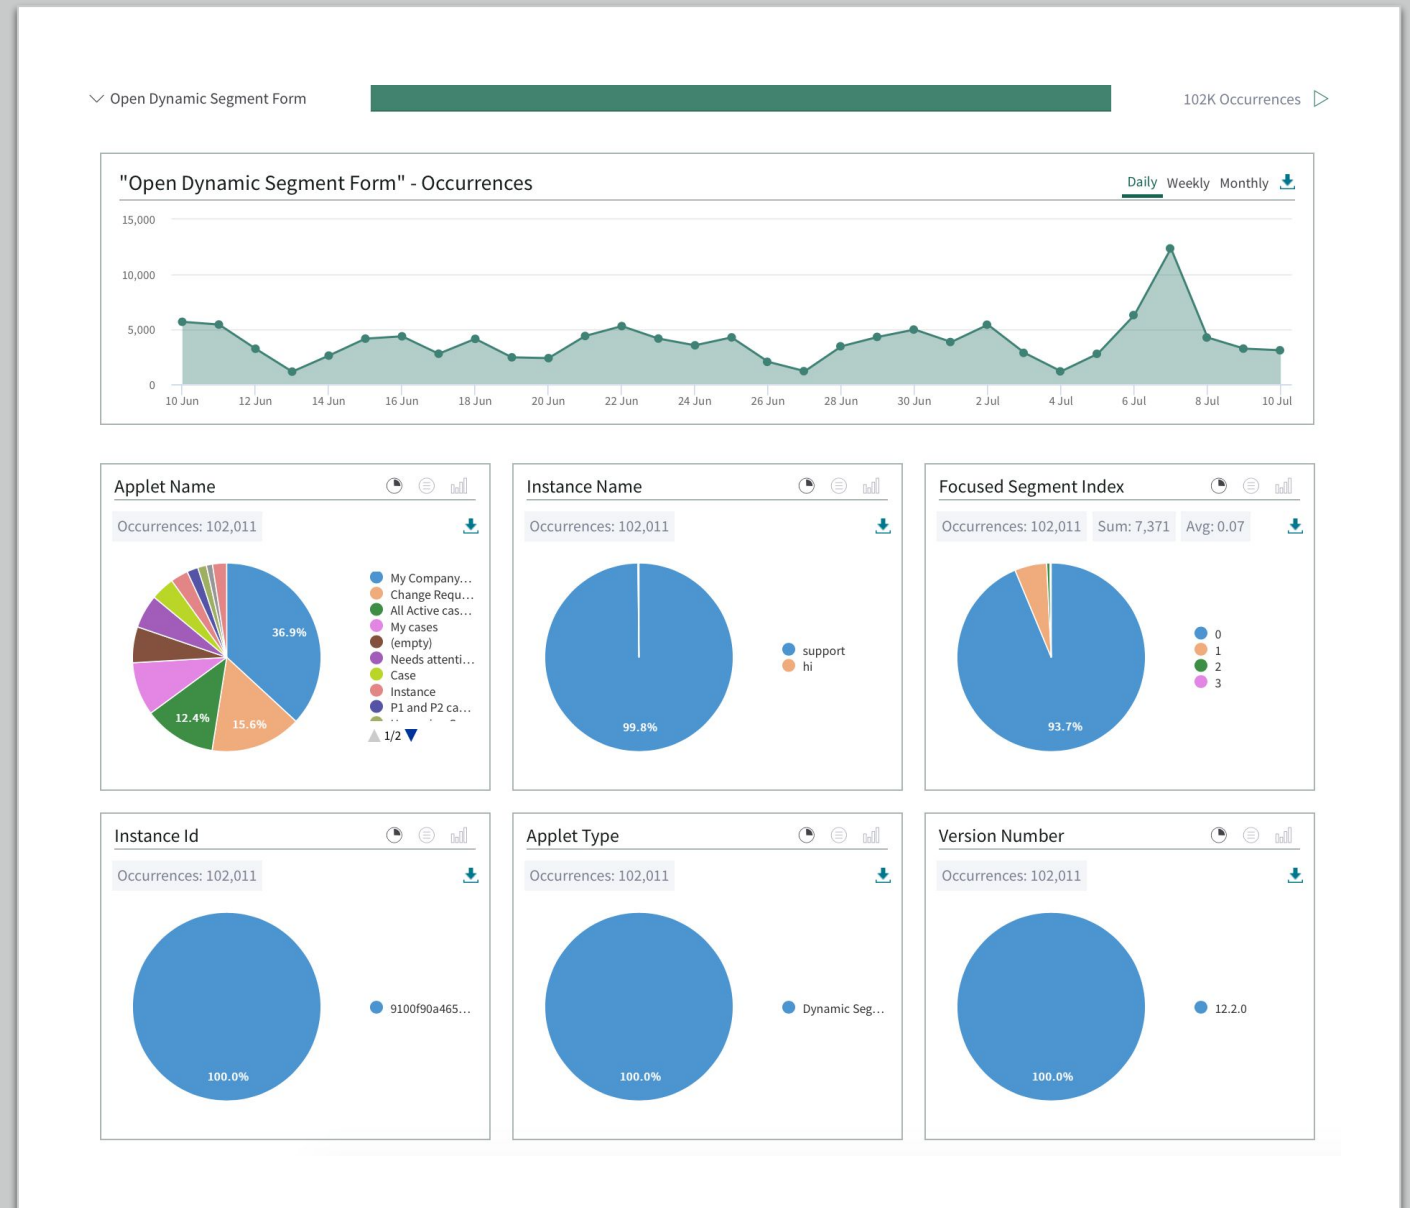
Application
App Version
Platform



Appsee holds ~30 different pre-aggregated reports, each segmented by the above dimensions + any other dimension relevant to the specific needs.

Custom Data Points

- Data points sent by customers (SN app teams / SN customers)
- Each data point is in the form {EventName, Properties<K,V>}
- Appsee provides information about the occurrences count of each event and a GROUP BY breakdown of each property value's occurrences



Pre Aggregated Data Model - Example

- AppSegments Table

Segment Id	App Id	Platform	App Version
1	Now Support		
2	Now Support	Android	1.0
...

- Event Occurrences Table

Segment Id	Day	Event Name	Count
1	01/07/2022	Search Started	53
1	01/07/2022	View Article	8
...


- Custom Event Properties Occurrences Table

Segment Id	Day	Event Name	Property Name	Property Value	Count
1	01/07/2022	Search Started	Search Term	Reset RSA Token	8
1	01/07/2022	Search Started	Search Term	Slow Load	44
1	01/07/2022	Search Started	Click Position	3	20
...

Processed Message

```
{
  Appld: Now Support,
  Platform: iOS,
  AppVersion: 1.0,
  EventName: 'Search Started',
  {
    'Search Term': 'Reset RSA Token',
    'Click Position' : 3,
    ...
  }
  ...
}
```


Appsee's High Level Pipeline



```
Metadata Upload Call
{
  Session Data Points
  User Data Points
  ...
}
```

```
Upload Result
{
  Accepted (Y / N)
  Skip upon next error
  ...
}
```



(MongoDB) Update the session entry with all collected MD

(MongoDB) Update the user entry

(MongoDB) Calculate all analytics metrics and update state (Funnels, Cohorts, Retention...)

(MongoDB) Update all user-facing reports

(Elasticsearch) Perform Elasticsearch Indexing

Requirements




Non Functional Requirements

- **Scalability**

- Horizontally scalable (CAP: prefer availability over consistency)
- Native sharding + replication support
- Write optimized
- Prefer products with simple topologies, No Single Point Of Failure (SPOF)

- **Maturity & Licensing**

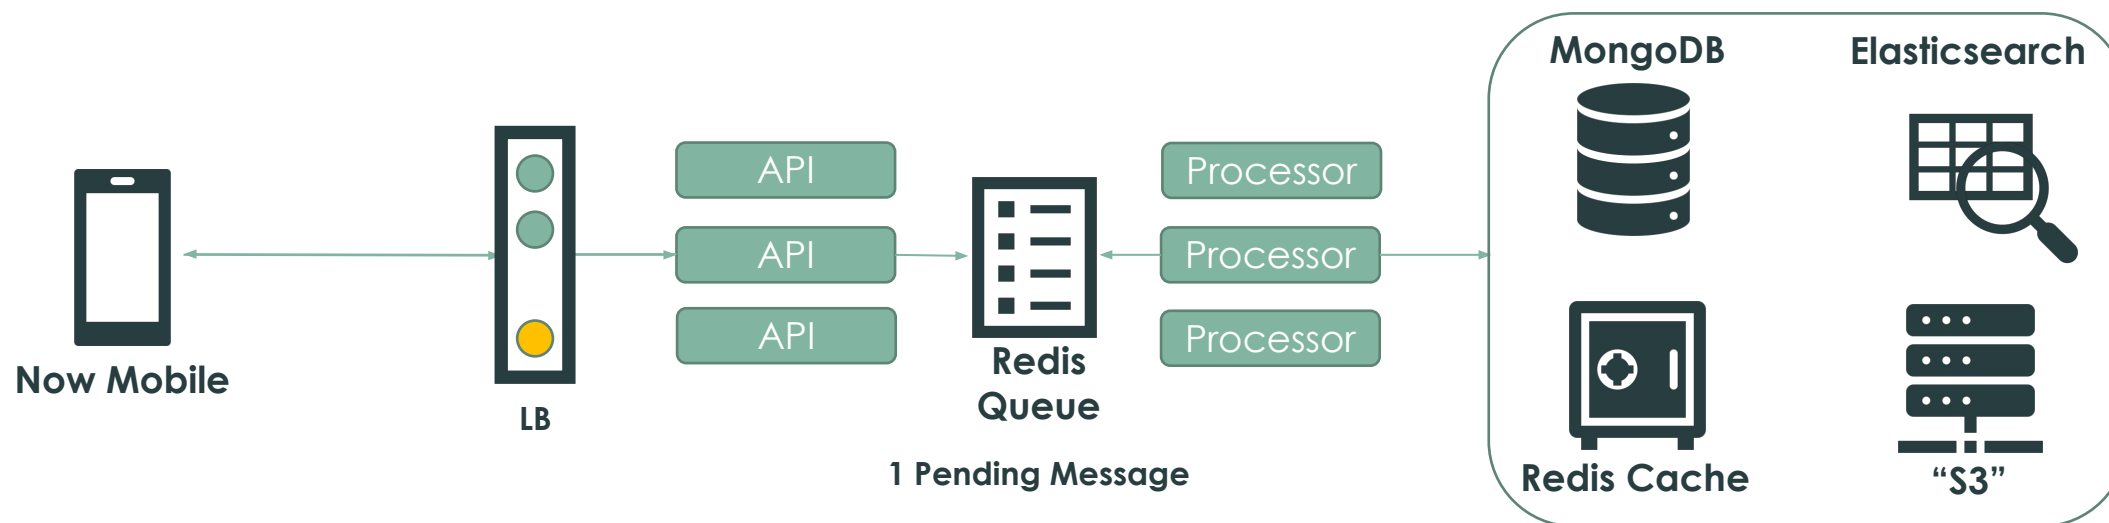
- Rich documentation + Community
 - Proper drivers (Python)
 - Fully open source, commercial support availability
- 

Functional Requirements

- Support all of Appsee's Existing dashboard aggregations:
 - Sub-second query time during insertion(1k data points/s per shard)
 - Same HW
- **Dynamic Segmentation**
 - **Allow segmenting by any session property**
 - **Allow segmenting by any event property value / combination of multiple properties**
 - **Same session "Join" - segment by multiple events (w/ prop values)**
- **Bonus: Data points \rightleftharpoons Users Join Aggregations**
(Users are mutable)




Appsee's High Level Pipeline




(MongoDB) Update the session entry with all collected MD

(MongoDB) Update the user entry



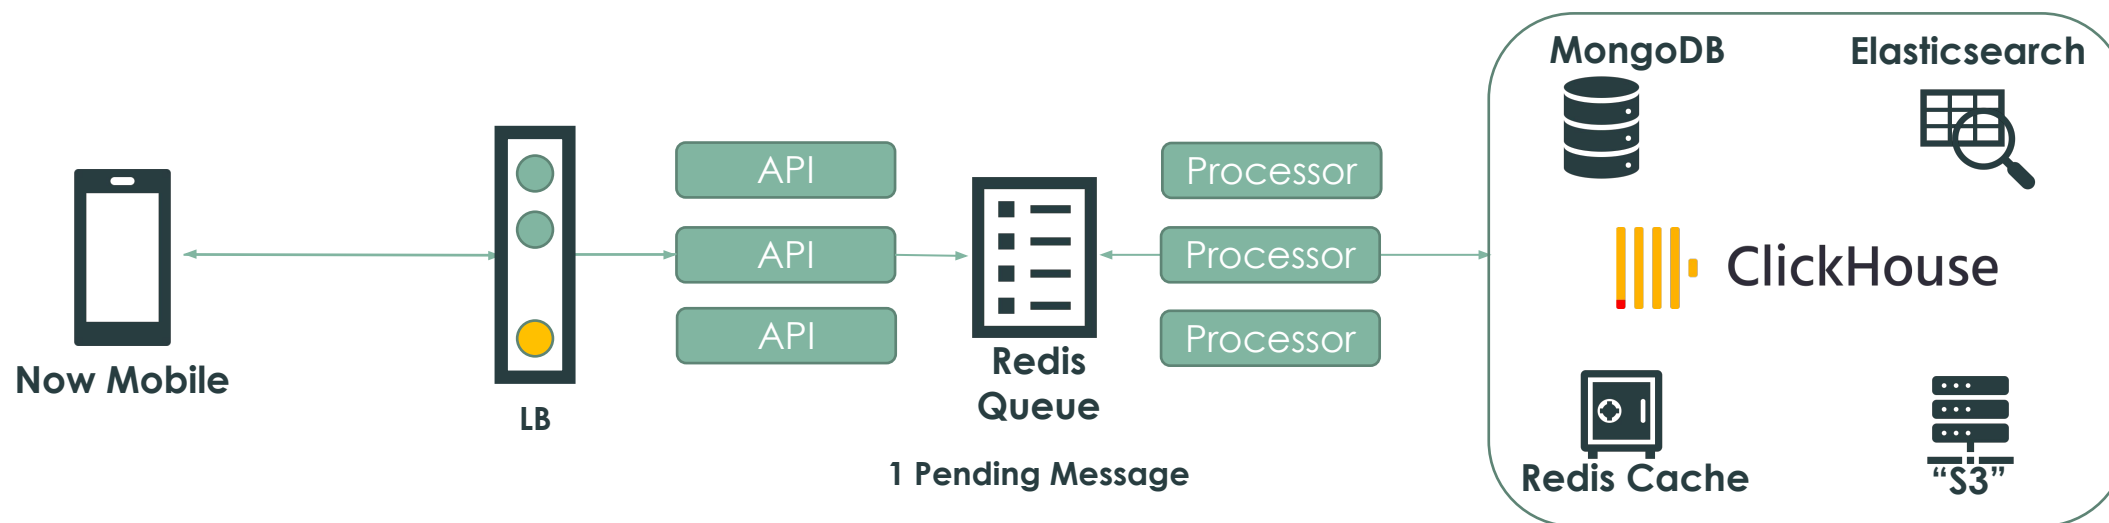
(MongoDB) Calculate all analytics metrics and update state (Funnels, Cohorts, Retention...)



(MongoDB) Update all user-facing reports

(Elasticsearch) Perform Elasticsearch Indexing

Appsee's High Level Pipeline



(MongoDB) Update the session entry with all collected MD

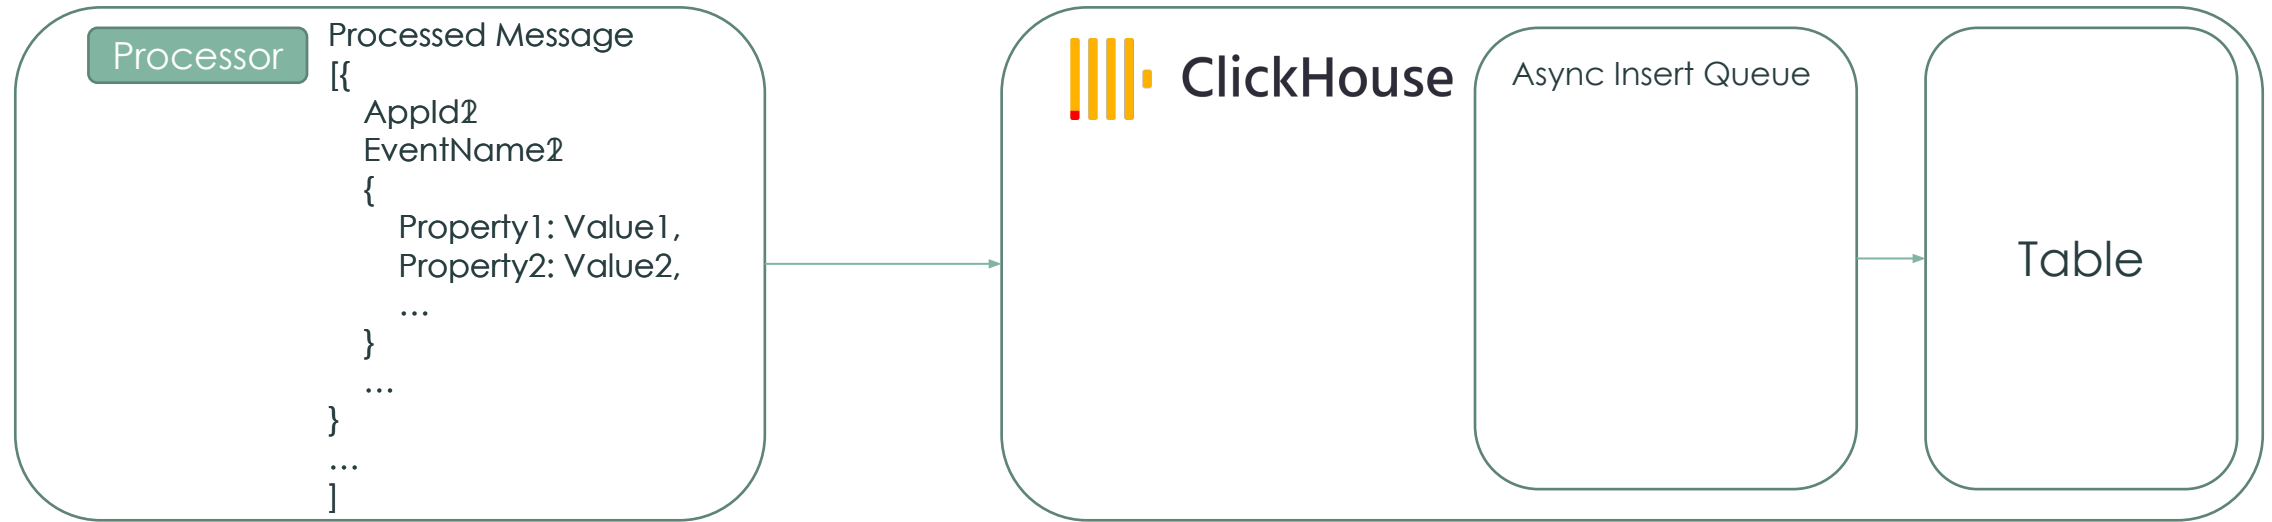
(MongoDB) Update the user entry

Break the MD file into data points

(ClickHouse) Insert all data points

(Elasticsearch) Perform Elasticsearch Indexing

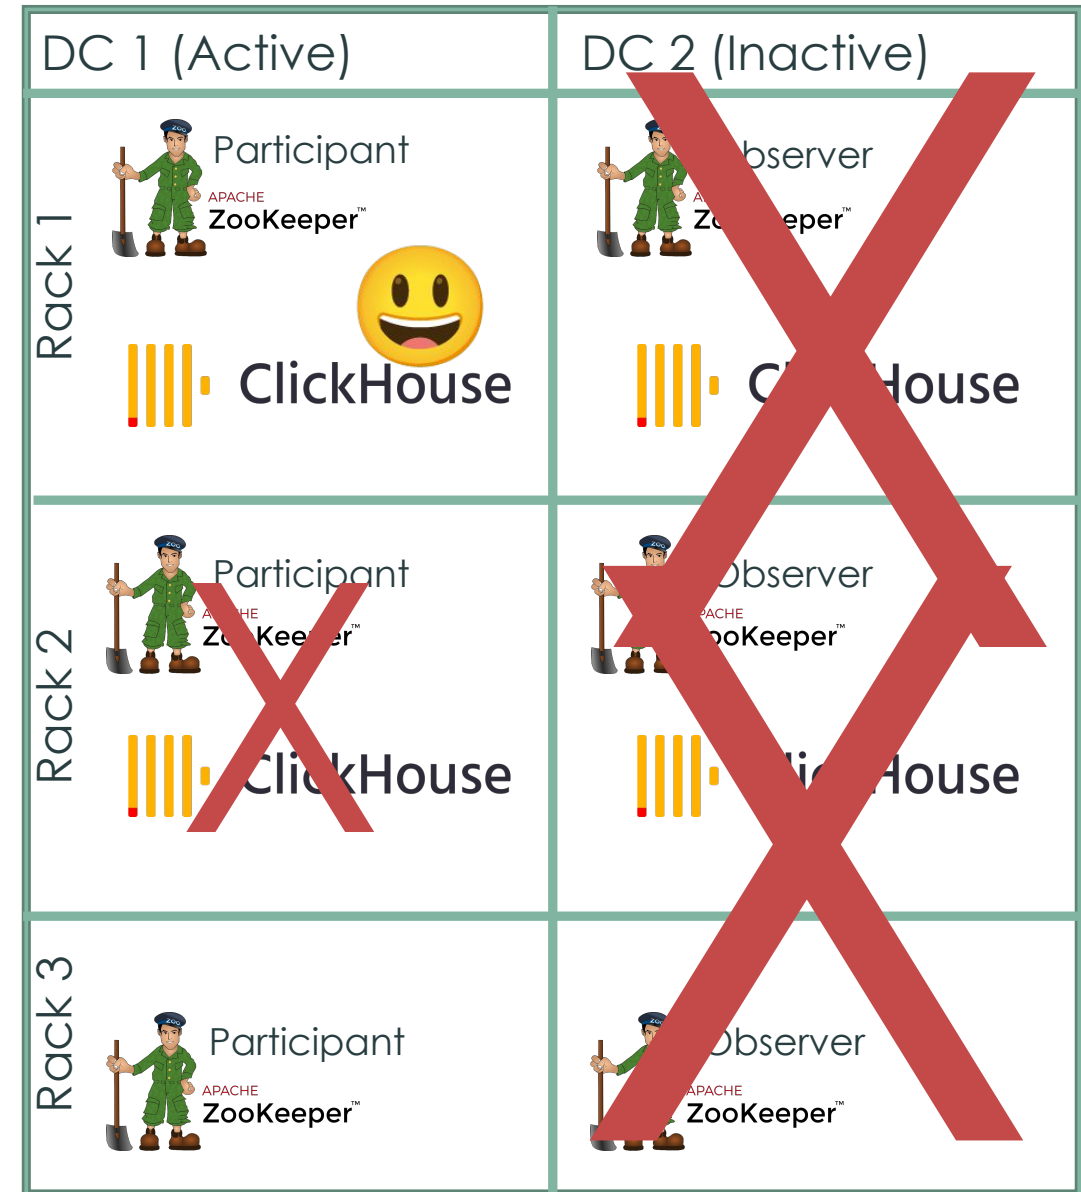
Insertion to ClickHouse



ClickHouse – Deployment

- 2 DCs (Active / Inactive)
- 3 Racks each DC
- A Zookeeper in each Rack
- Each CH shard has 2 replicas per DC (4 in total)
- 3 participant Zookeepers in the active DC
- 3 observer Zookeepers in the inactive DC

* Allowing DC + Rack failure




ClickHouse Benchmarks - Overview

Setup

- 2 Servers, each with 16 CPUs and 32GB RAM
- 430M Sessions (per model) over a single app == 2.1B Data points
- *Note:* Ingesting 4.5M data points took 124 seconds (~37.5K data points / second) and was bound on the loading machine and not the CH cluster
- Converted every Appsee session into a set of separate data points, each one representing an event occurrence
- Tested 6 different models with a set of 7 queries

ClickHouse – Models Benchmark



	Query	Result	Rows Scanned **	Row Per Property	Nested	Map Type	Mapping	Mapping & Ordinals	Map Type & Ordinals
1	COUNT 'StartScreen' events	964,553,500	970,536,535	2.29	0.76	0.77	0.79	0.13	0.12
2	COUNT 'Open Tab' events WITH property 'Instance Name' = 'ABC'	260,970,000	271,063,867	2.63	2.12	1.37	0.5	0.18	0.45
3	COUNT event occurrences GROUPED BY event name	-	2,154,496,000	8.46	2.27	2.26	2.26	0.59	0.57
4	COUNT UNIQUE sessions WITH event 'Open Tab' *	237,008,022	271,063,867	2.18	1.07	1.07	0.96	0.71	0.81
5	COUNT UNIQUE sessions WITH event 'Open Tab' AND property 'Instance Name' = 'ABC' *	233,070,147	271,063,867	3.17	2.3	1.71	0.73	0.35	0.68
6	COUNT UNIQUE sessions WITH event 'Open Tab' AND property 'Instance Name' = 'ABC' AND property 'Tab Name' = 'SafePass' *	224,091,596	271,063,867	14.23	2.67	2.04	1.01	0.46	0.87
7	COUNT UNIQUE sessions WITH 2 events and 1 property: (event 'Open Tab' AND property 'Instance Name' = 'ABC') AND (event 'User Login' AND property 'Version Number' = '11.0.0') *	2,083,958	357,072,779	5.82	4.83	4.01	2.85	2.45	2.78
	Storage(GB)			32	14.02	13.26	15.02	15.66	12.8

* Estimation ± 1% miss

** Irrelevant for model "Row Per Property" which has a different PK


ClickHouse – Chosen Data Model = Map Type

App Id	Event Time	Session Id	User Id	Platform	Country	...	Event Name	Map<Event Property Name, Event Property Value>
--------	------------	------------	---------	----------	---------	-----	------------	---

- Allows an arbitrary count of properties
- Sharded CH cluster using primary key: {AppId, startOfDay(EventTime), EventName, UserId}
- Ready for Ordinals migration in the future (The process of transforming String values into Integers)
- Easy to use

Future ClickHouse Work

1. Idempotent insertion (Exactly once insert)



Disable ClickHouse Async Insert and rely on its dedup feature based on bulk hash

2. Insert User entities to allow JOINS
3. Realtime Cohorts & Funnels
4. Move to K8S

Realtime User Analytics with ClickHouse

Amir Vaza @ Linked in
Architect & Lead at ServiceNow