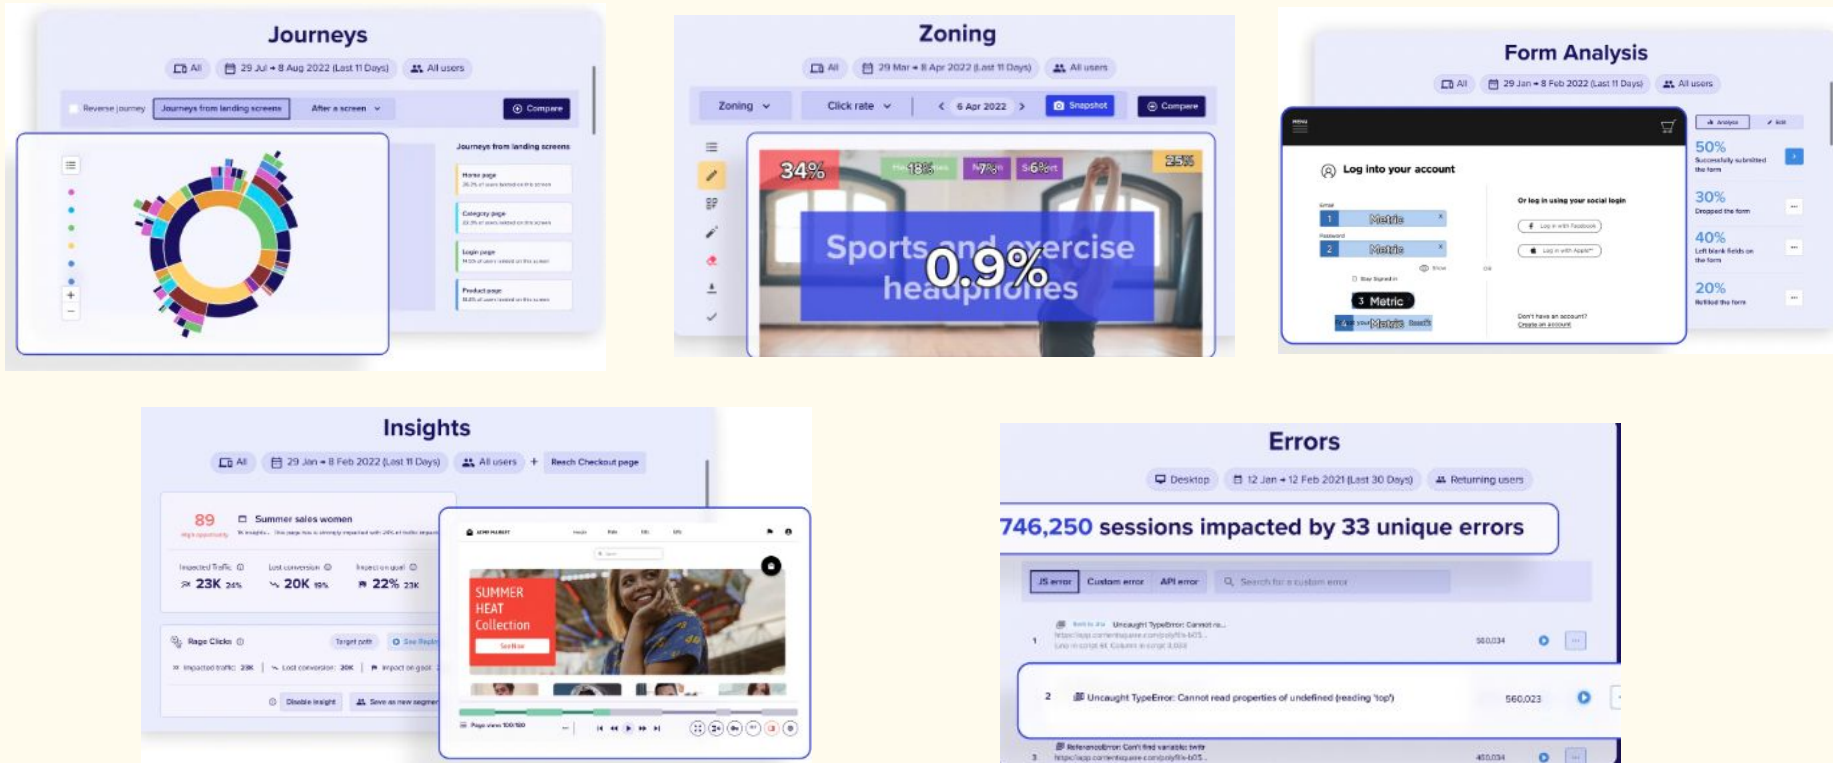




Clickhouse journey in Contentsquare

Doron Hoffman
Chief Architect @ Contentsquare

Contentsquare experience analytics platform



The Contentsquare community



1000+ clients



1M+ SMB customers



1B+ pages analyzed a day



1500+ employees




17 offices

Clickhouse clusters in Contentsquare



Up to 9 shards per cluster



2 replicates per shard



16 clusters on 4 regions



100+ instances with CH



2 cloud providers



2 regions per provider

Analytics cluster – ELK to Clickhouse



CH replace ELK 4 years ago



11x cheaper (with 6x more data)



10x faster on p99

Analytics cluster

- CollapsingMergeTree engine
- RAID of 5 st1 EBS disk per node
- 13 month of retention
- use cases require on the fly computation
- Clickhouse powerful functions enabled use cases like Population analysis on sequential behavior

more details in Contentsquare engineering blog
[Behavioral Time Series Segmentation in ClickHouse](#)



Real time events reconciliation cluster

- AggregationMergeTree engine
- RAID of 4 gp3 EBS disk per node
- short retention
- Enables analysis of recent events which occurred before session ended
- more details in Contentsquare engineering blog [Real-Time events reconciliation with ClickHouse](#)



Observability cluster (internal)

- S3 table engine
- analysis of internal audit events
- Contentsquare data team contributed Azure blob storage support
- cross cloud support

BI cluster (Internal)

- replaced Redshift
- cross cloud support
- reduced cost
- increased analysis capabilities vs. prior solution

Internal tooling and contribution

- schema management – chschema
public site – <https://chschema.org/>
planned to be opened sourced
- User management integrated with vault
- PII deletion tooling
- scala client – dynamically builds and execute queries based on AST with optimizer
- Chproxy – Contentsquare ownership on open source
public site – <https://www.chproxy.org/>
- contributing PRs to Clickhouse open source

Contentsquare engineering blog – more content

<https://engineering.contentsquare.com/>

- [How we scale out our ClickHouse cluster](#)



- [Full-Text Search with Quickwit and ClickHouse in a Cluster-to-Cluster Context](#)





Thank you

