

Behind the Scenes:

Observability with ClickHouse

Boris Tane

baseline.io
@boristane

25th May 2023



@boristane
Software Engineer



@boristane
Software Engineer

"observability without action is just storage"

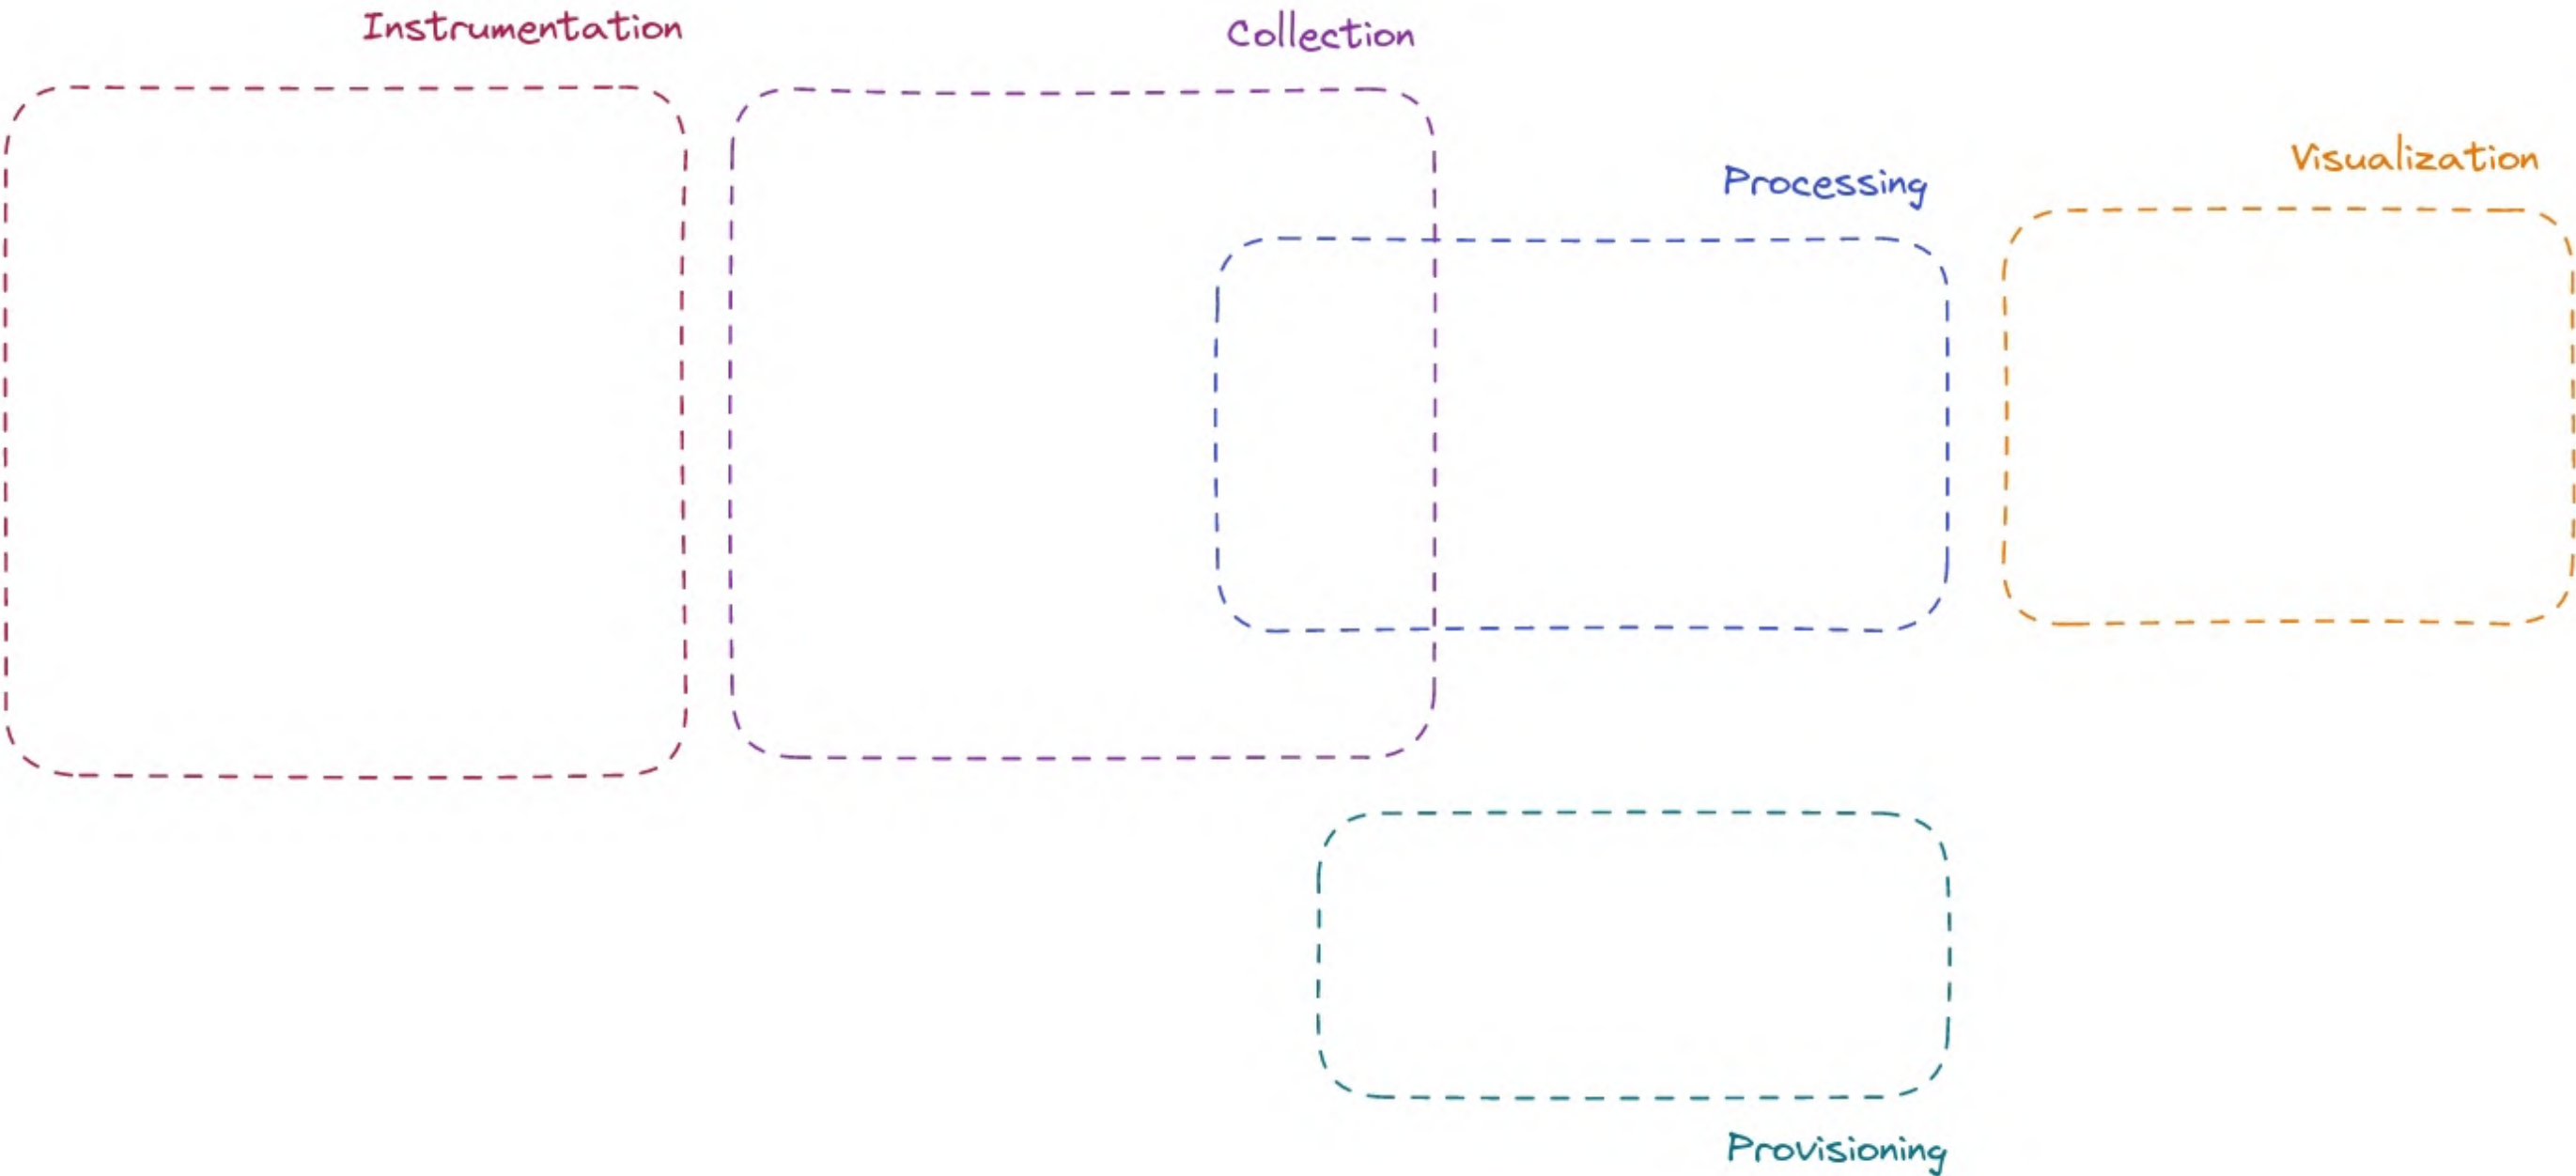
What is observability?




<https://youtu.be/cJSarYwejzs>

understand how your software works
by looking at its external outputs

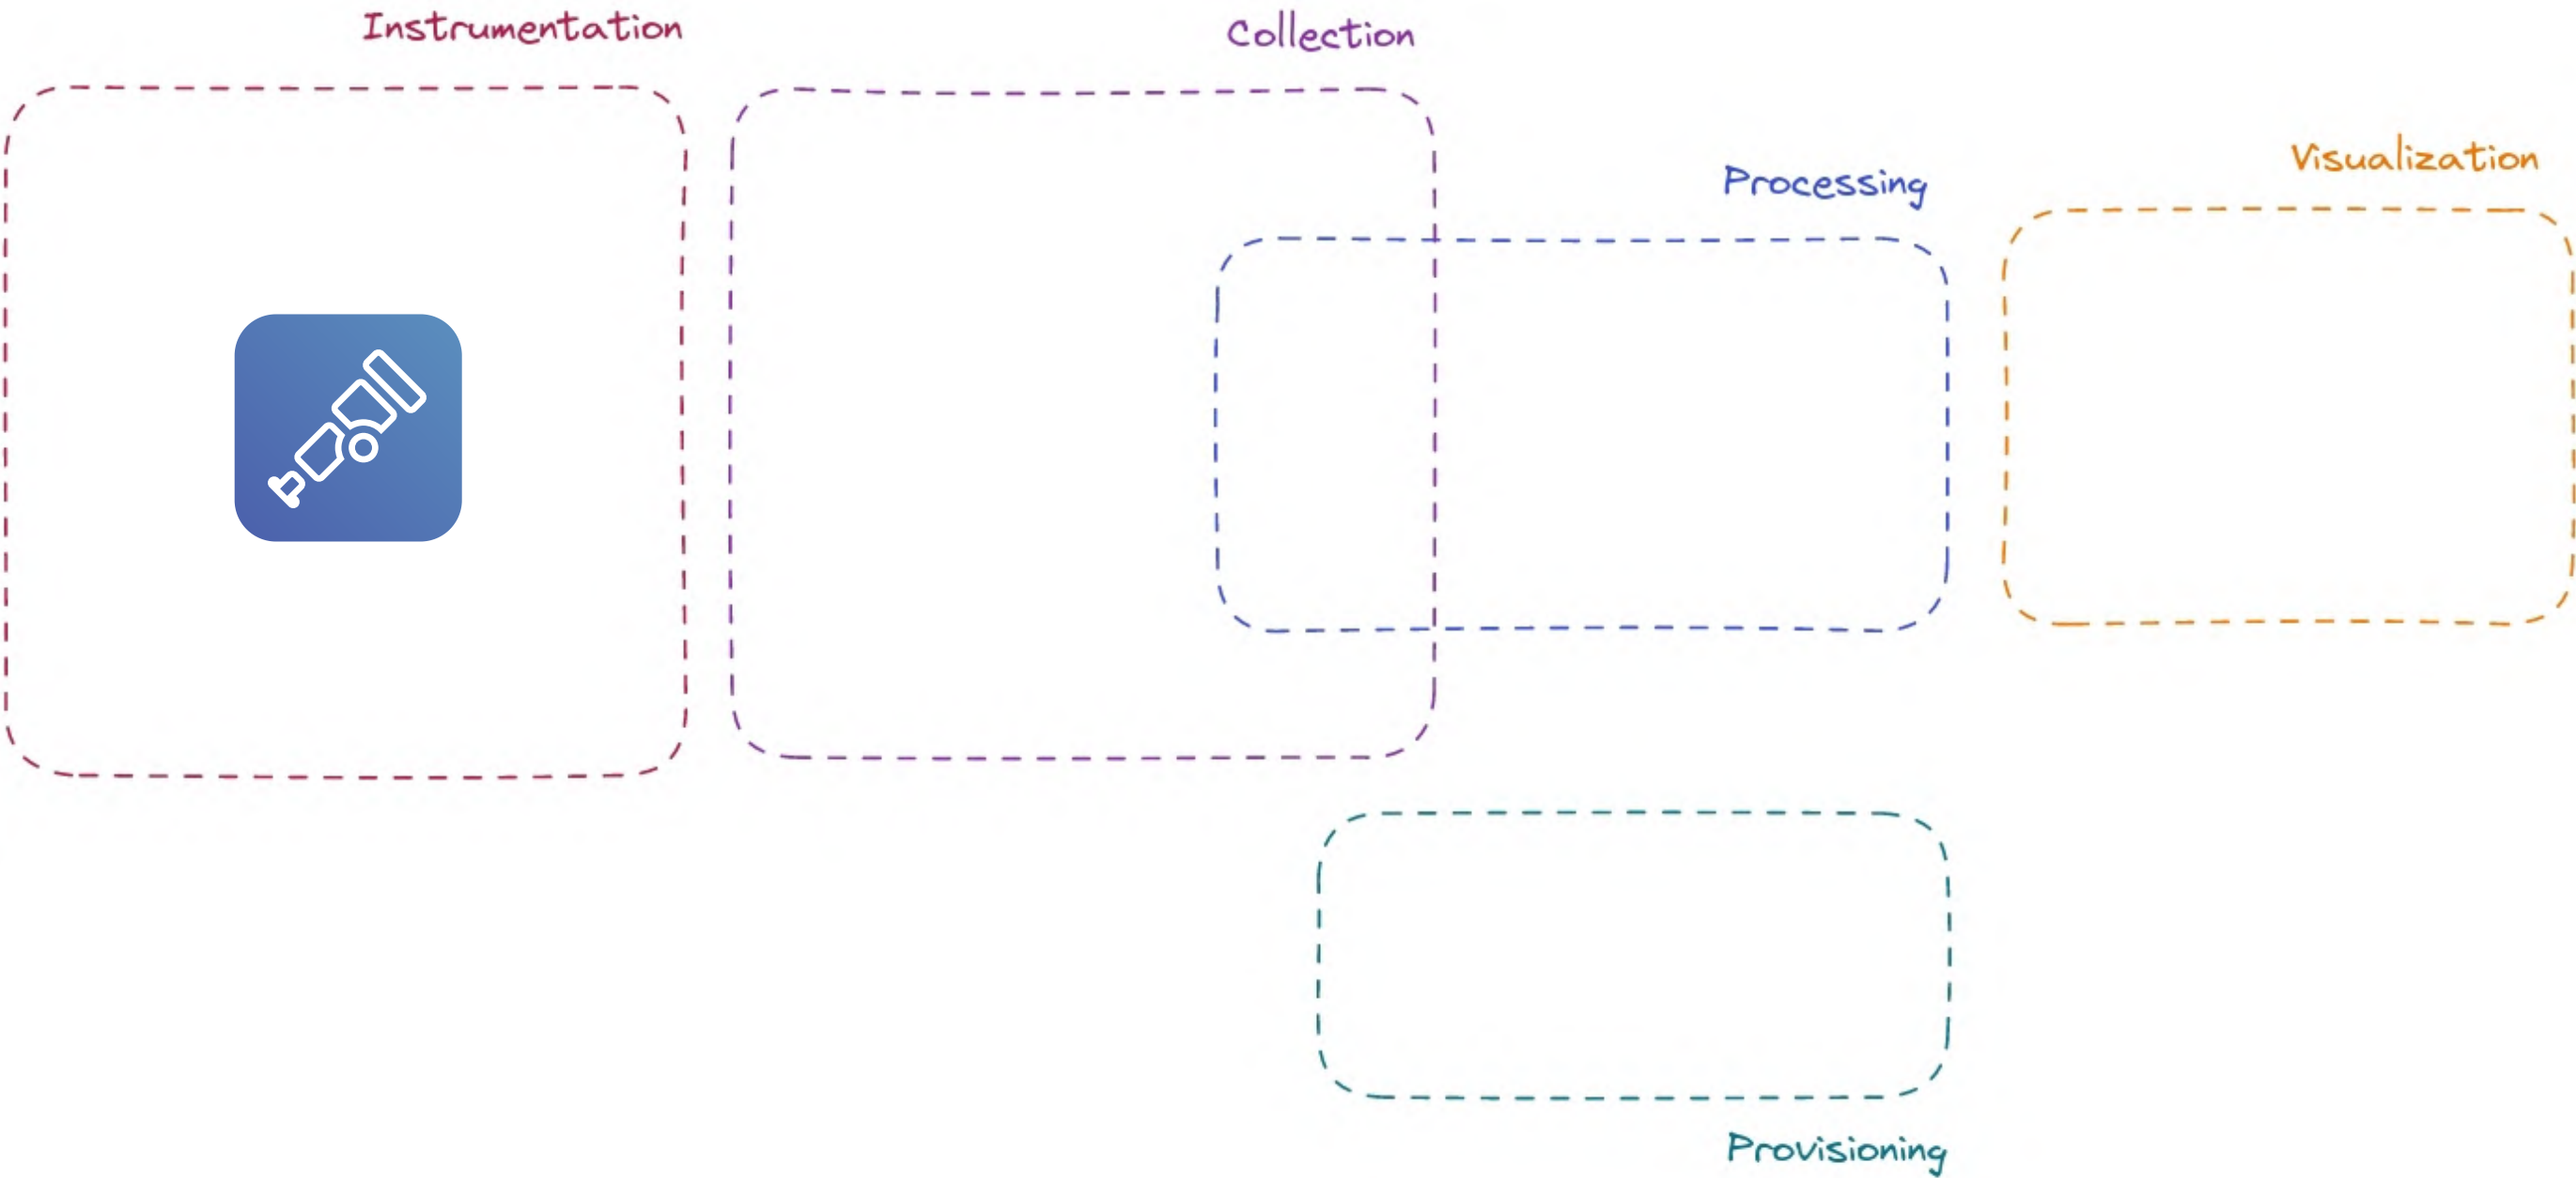
How do we get there?



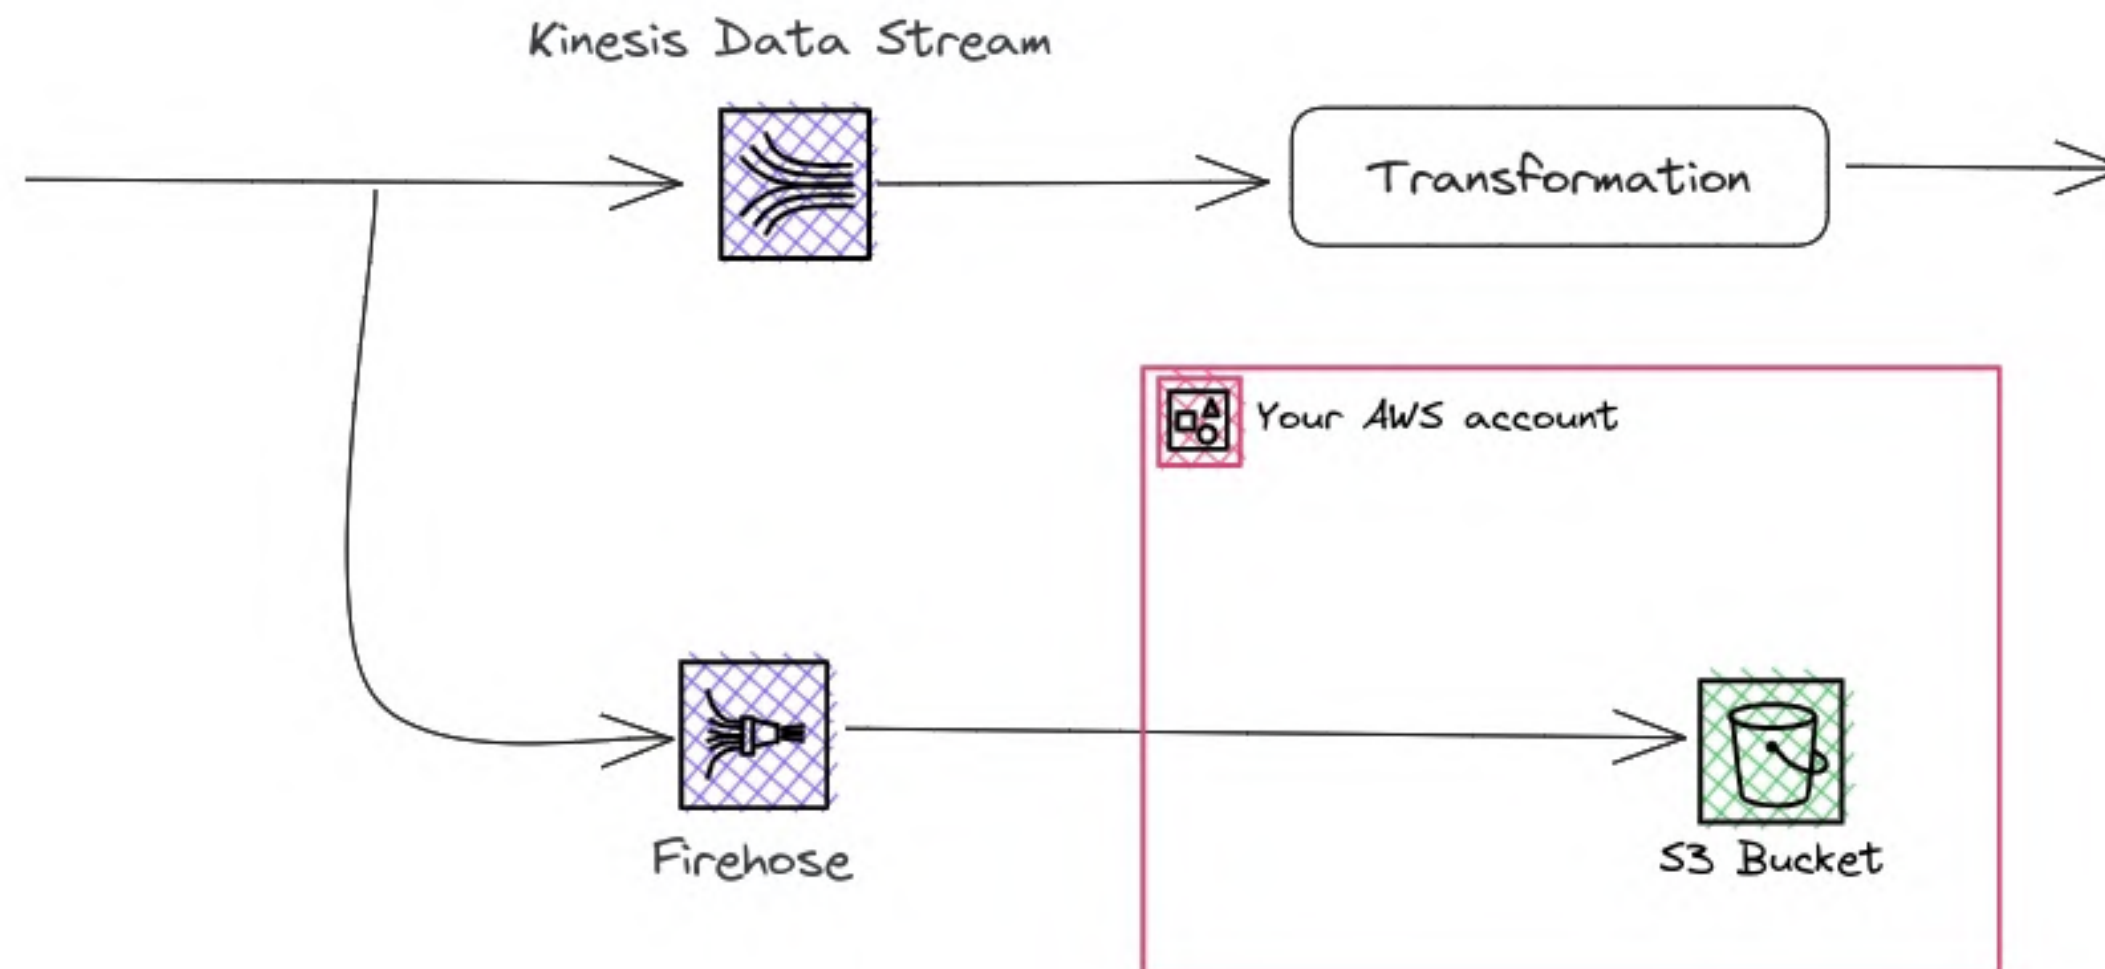
Instrumentation



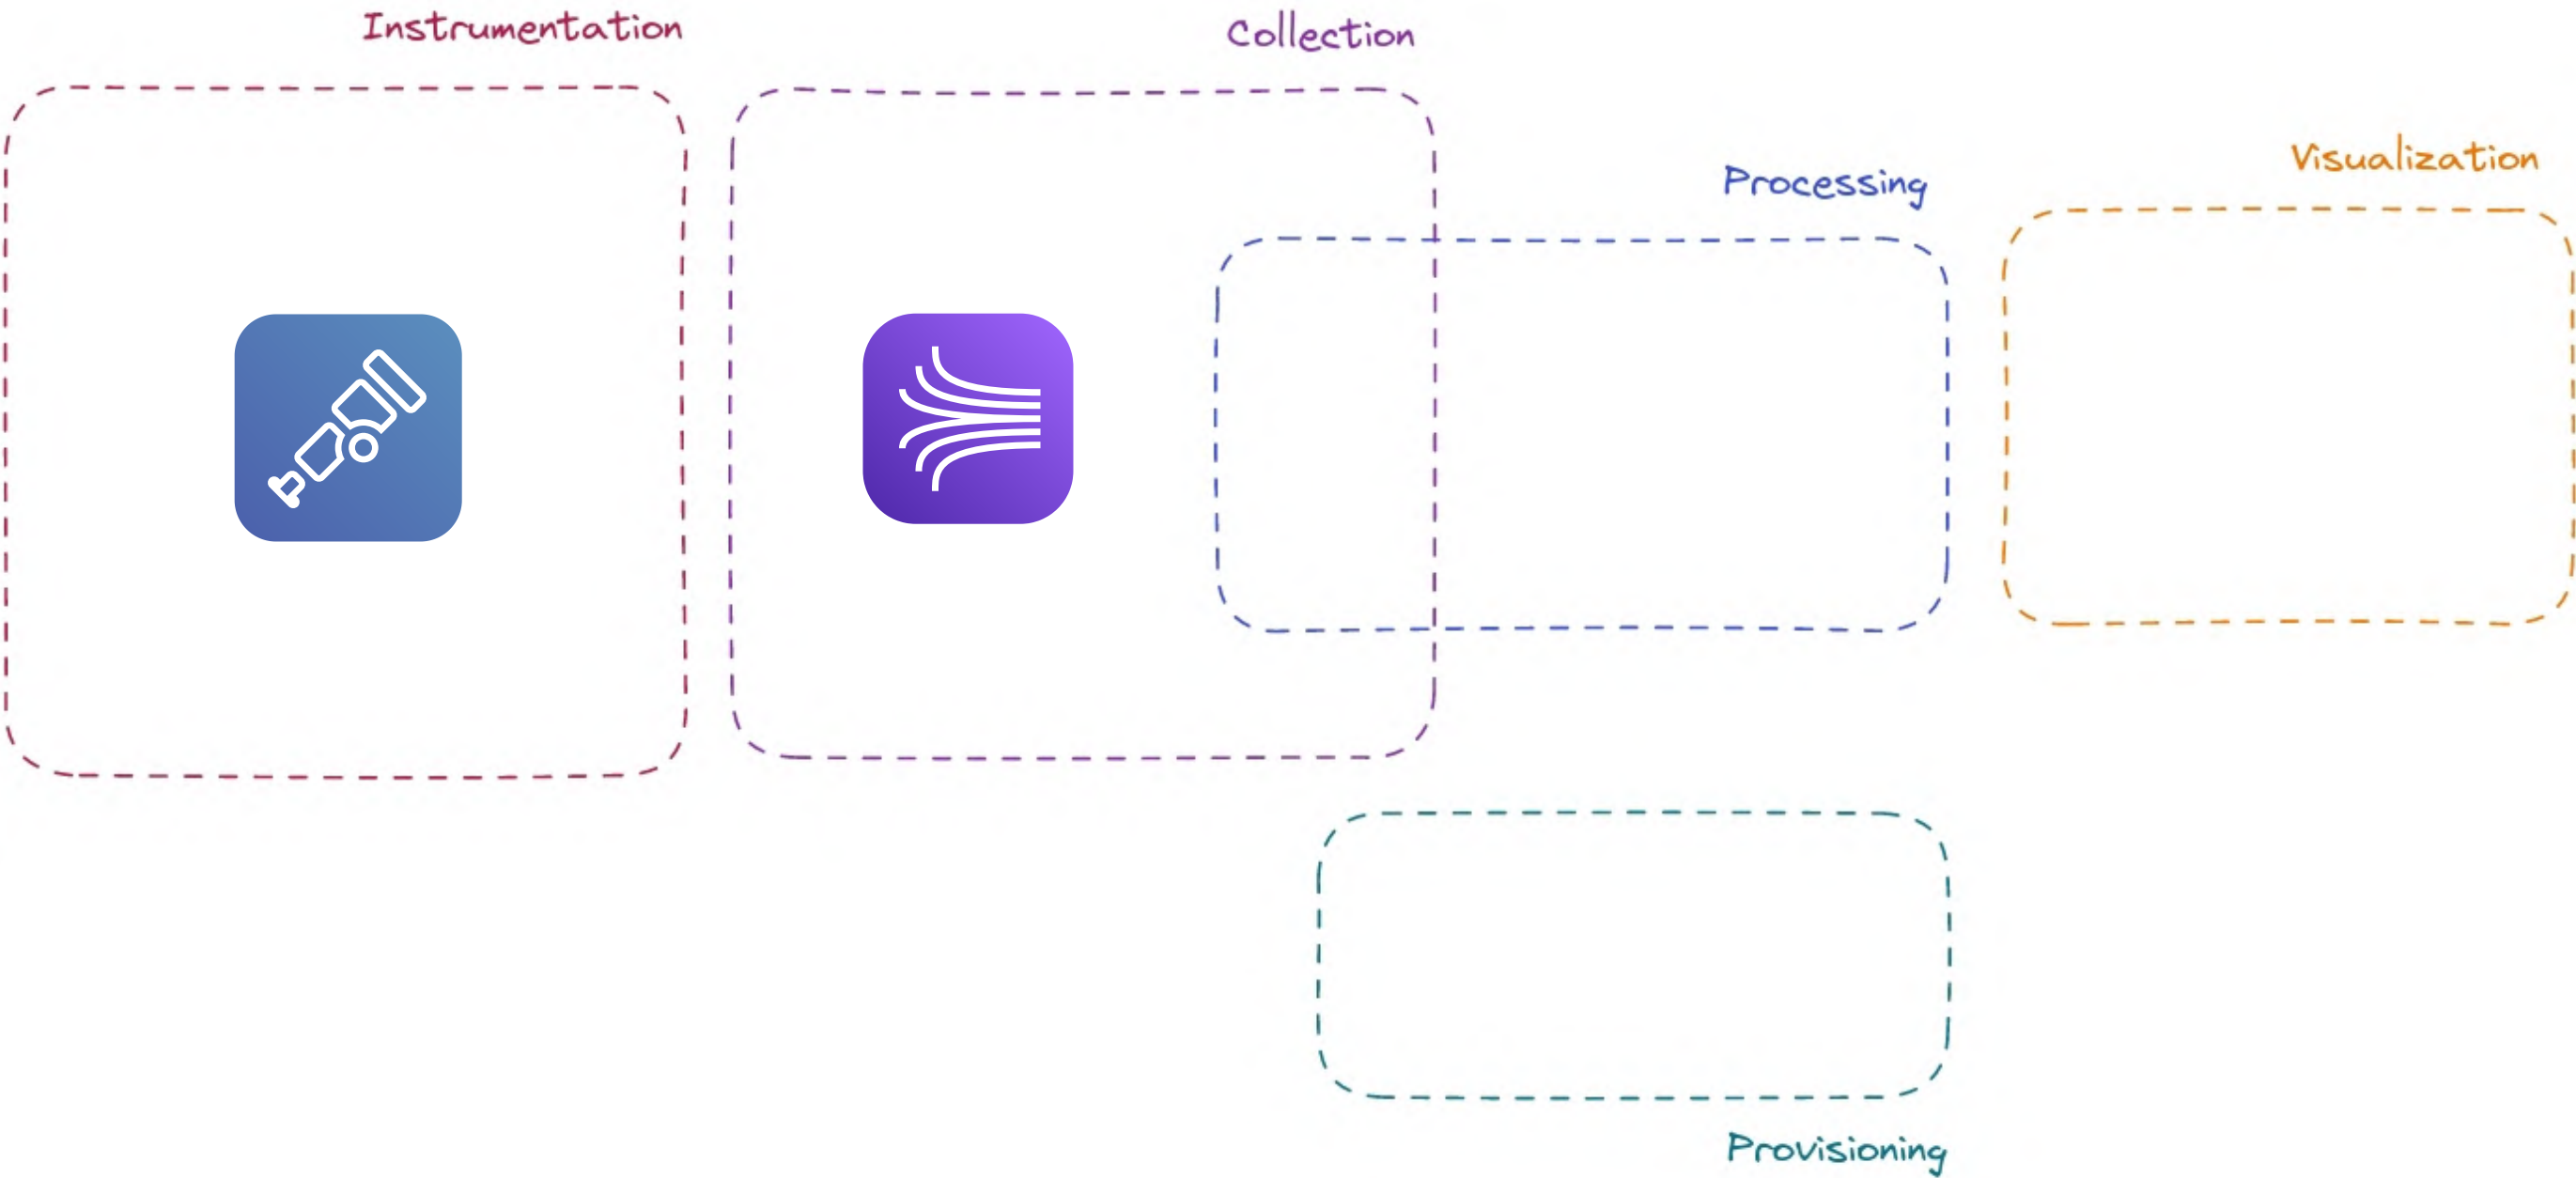
How do we get there?



Collection



How do we get there?




Provisioning

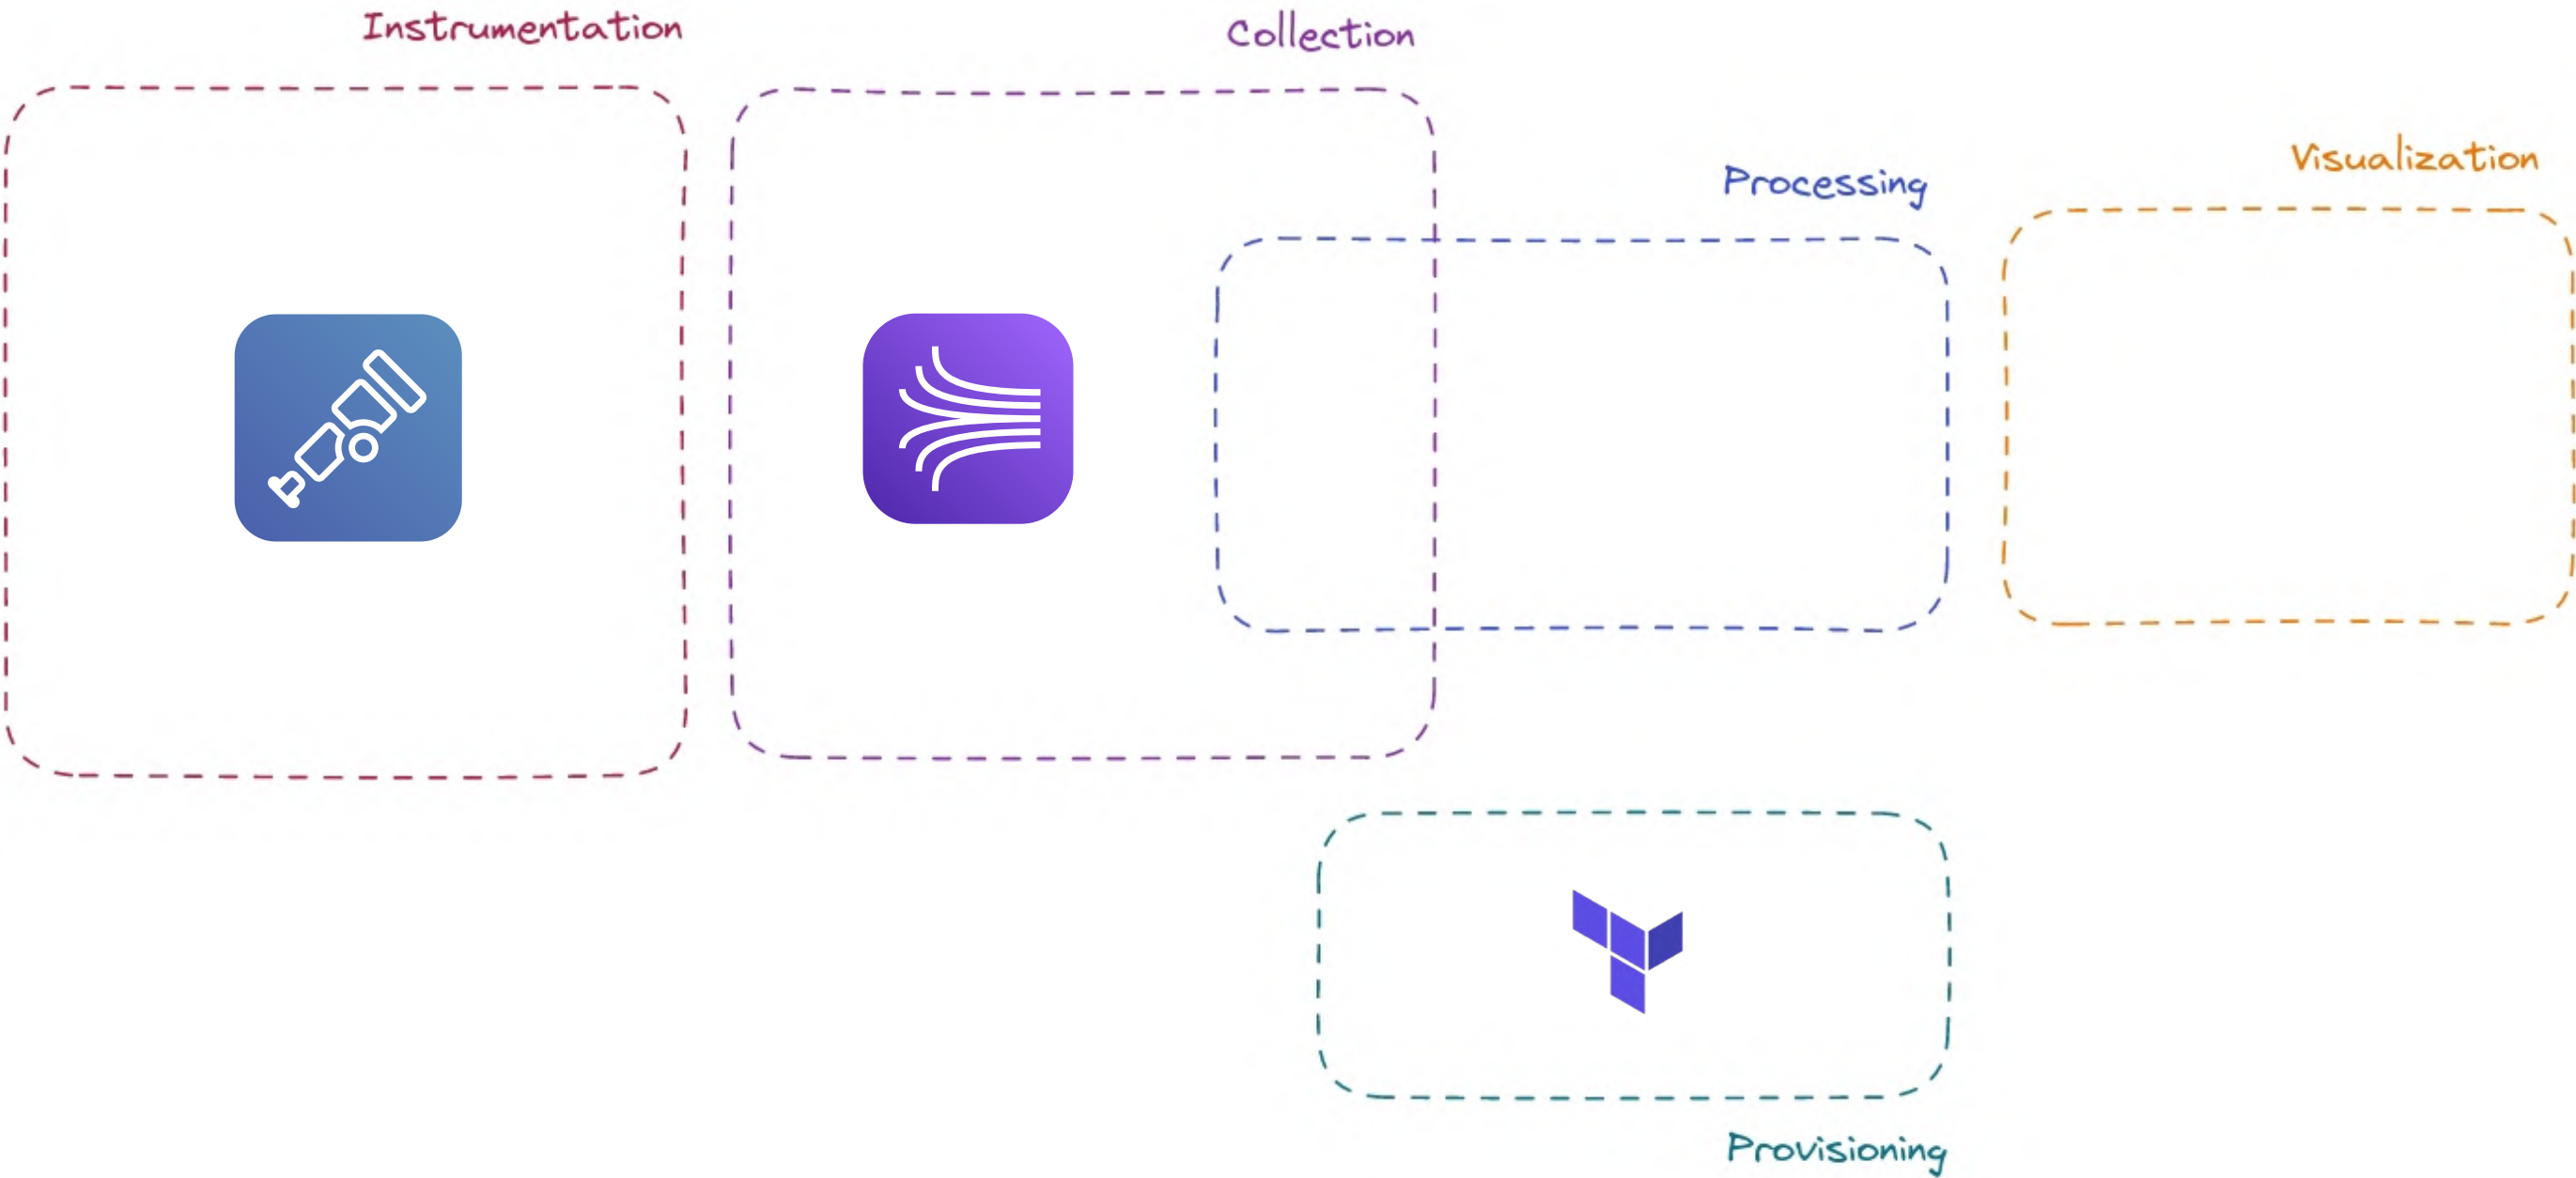
```
avg-invoice-error-cost:
  type: query
  properties:
    parameters:
      datasets:
        - lambda-logs
      calculations:
        - AVG(extra.price)
      filters:
        - message = Error sending the invoice to a workspace
```

Provisioning

```
avg-invoice-error-cost:
  type: query
  properties:
    parameters:
      datasets:
        - lambda-logs
      calculations:
        - AVG(extra.price)
      filters:
        - message = Error sending the invoice to a workspace
```



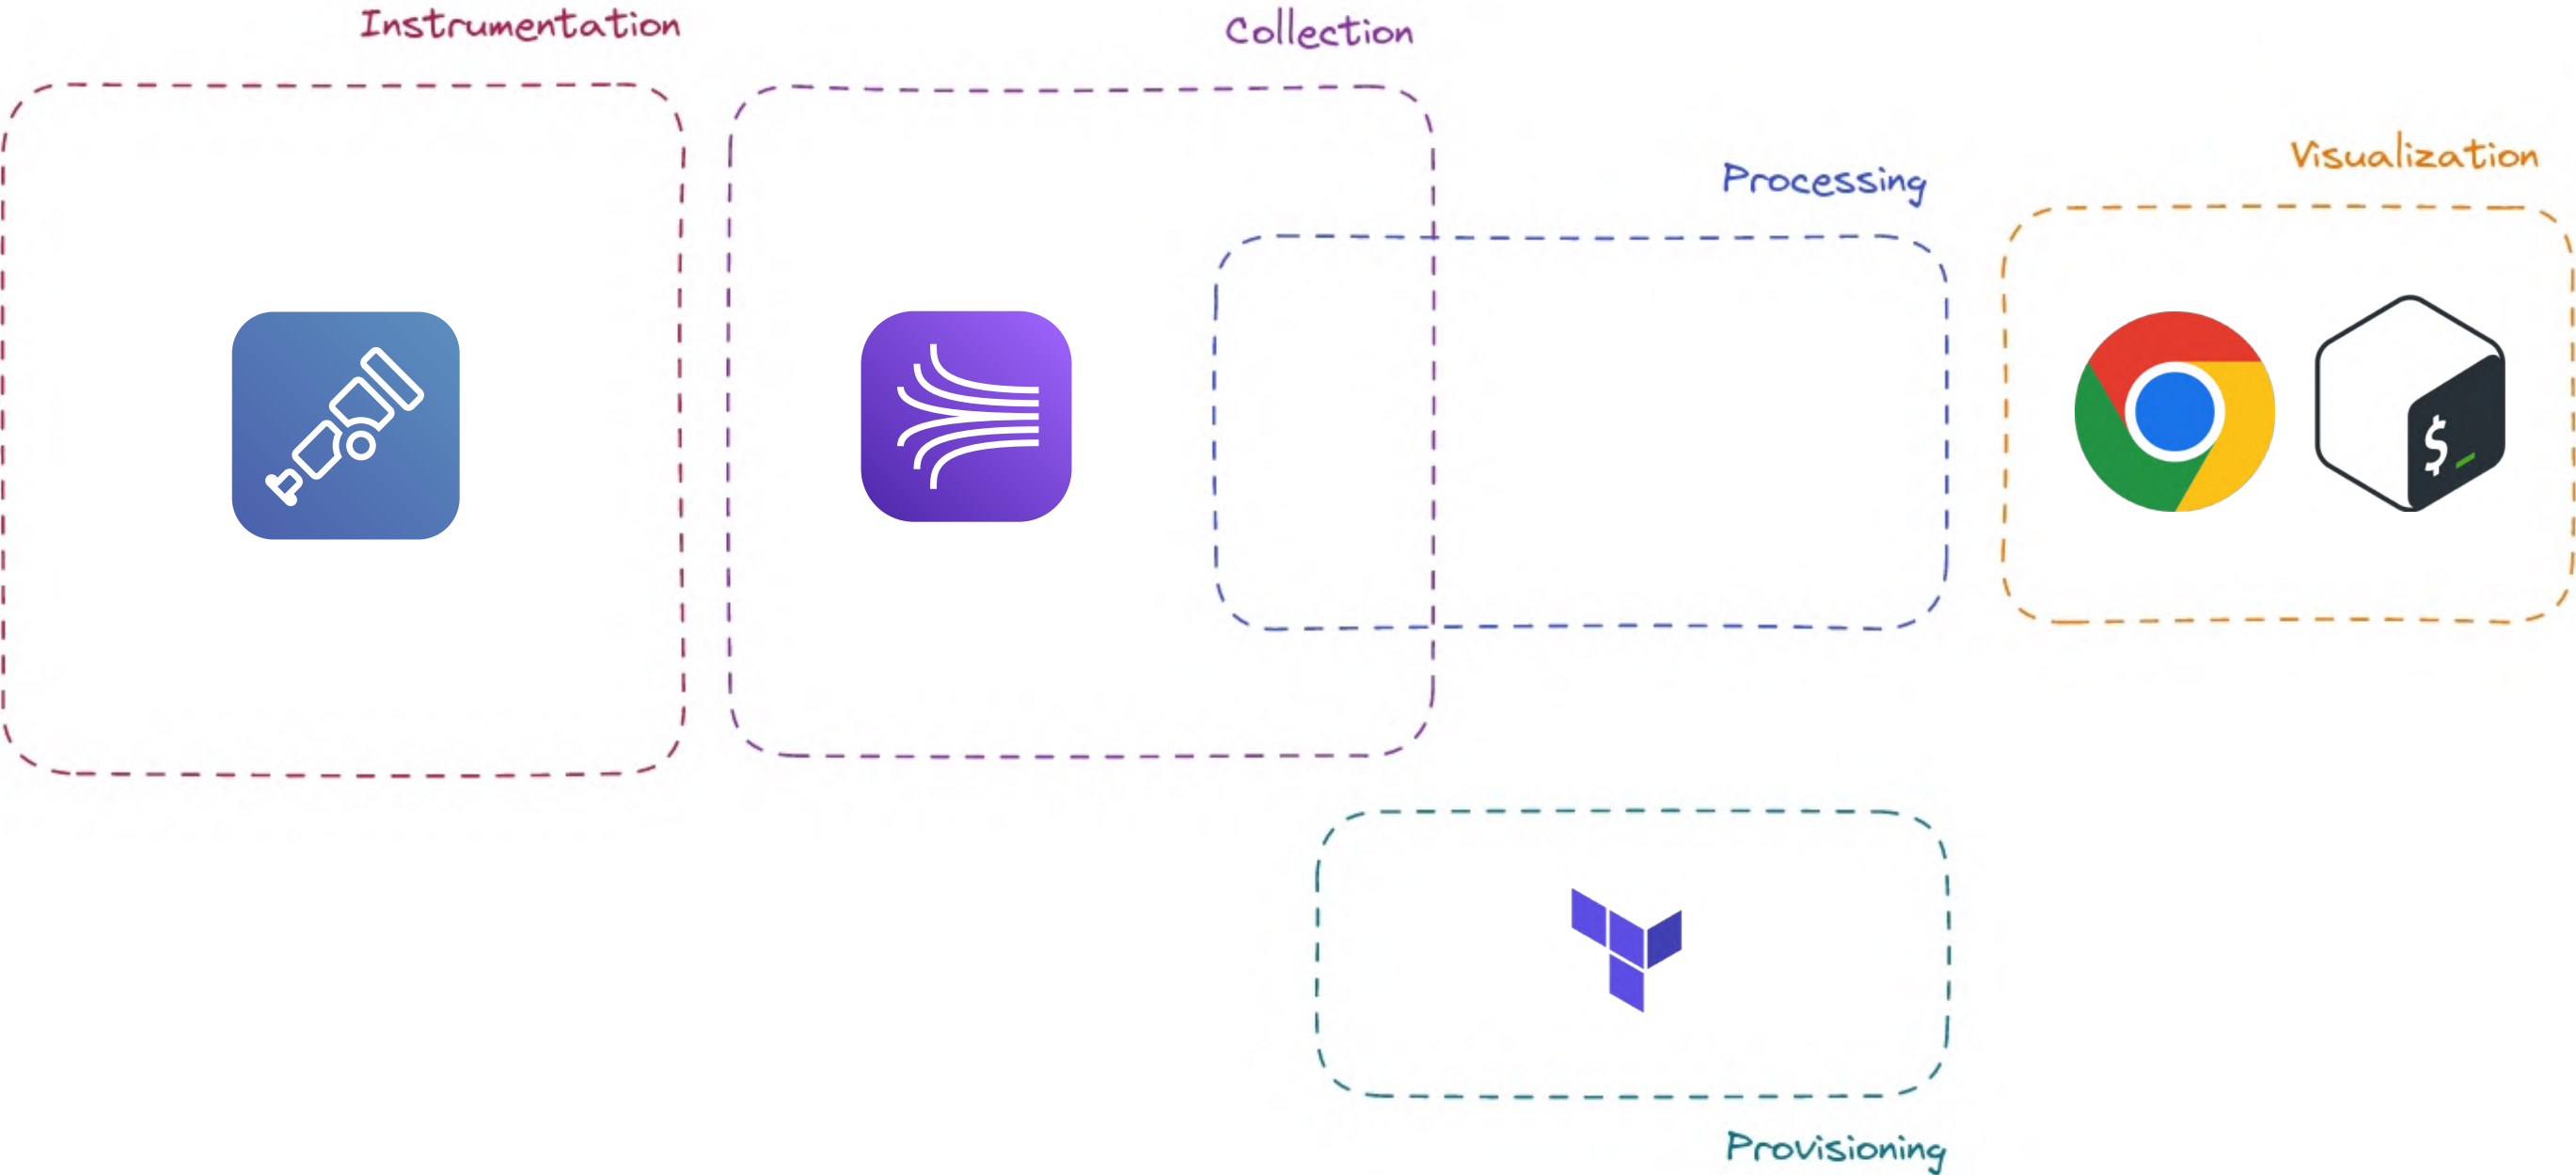
How do we get there?



Visualisation

Demo

How do we get there?



Processing



Why ClickHouse?

Why ClickHouse?

- Fast

Why ClickHouse?

- Fast
- Very Fast

Why ClickHouse?

- Fast
- Very Fast
- Very Blazing Fast

What's the data model?

- Schemaless

What's the data model?

- Schemaless
- High Cardinality

What's the data model?

- Schemaless
- High Cardinality
- Multi-tenant

What's the data model?

- Schemaless
- High Cardinality
- Multi-tenant
- Fast

What's the data model?

```
CREATE TABLE <table_name>
(
    _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
    _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
    _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
    _baselineId FixedString(26) NOT NULL CODEC(ZSTD(1)),
    _ttl DateTime NOT NULL,

    _traceId String CODEC(ZSTD(1)),
    _requestId String CODEC(ZSTD(1)),
    _spanId String CODEC(ZSTD(1)),
    _coldStart UInt32 Codec(T64, LZ4),
    _parentSpanId String CODEC(ZSTD(1)),
    _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
    _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
    _duration UInt32 Codec(T64, LZ4),

    _source String NOT NULL CODEC(ZSTD(1)),

    `string.names` Array(String) CODEC(ZSTD(2)),
    `string.values` Array(String) CODEC(ZSTD(2)),
    `number.names` Array(String) CODEC(ZSTD(2)),
    `number.values` Array(Float64) CODEC(ZSTD(2)),
    `boolean.names` Array(String) CODEC(ZSTD(2)),
    `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

    INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
    ...
)
```

What's the data model?

```
CREATE TABLE <table_name>
(
    _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
    _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
    _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
    _baselimeId FixedString(26) NOT NULL CODEC(ZSTD(1)),
    _ttl DateTime NOT NULL,

    _traceId String CODEC(ZSTD(1)),
    _requestId String CODEC(ZSTD(1)),
    _spanId String CODEC(ZSTD(1)),
    _coldStart UInt32 Codec(T64, LZ4),
    _parentSpanId String CODEC(ZSTD(1)),
    _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
    _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
    _duration UInt32 Codec(T64, LZ4),

    _source String NOT NULL CODEC(ZSTD(1)),

    `string.names` Array(String) CODEC(ZSTD(2)),
    `string.values` Array(String) CODEC(ZSTD(2)),
    `number.names` Array(String) CODEC(ZSTD(2)),
    `number.values` Array(Float64) CODEC(ZSTD(2)),
    `boolean.names` Array(String) CODEC(ZSTD(2)),
    `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

    INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
    ...
)
```

LowCardinality(String)

What's the data model?

```
CREATE TABLE <table_name>
(
  _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
  _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
  _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
  _baselineId FixedString(26) NOT NULL CODEC(ZSTD(1)),
  _ttl DateTime NOT NULL,

  _traceId String CODEC(ZSTD(1)),
  _requestId String CODEC(ZSTD(1)),
  _spanId String CODEC(ZSTD(1)),
  _coldStart UInt32 Codec(T64, LZ4),
  _parentSpanId String CODEC(ZSTD(1)),
  _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _duration UInt32 Codec(T64, LZ4),

  _source String NOT NULL CODEC(ZSTD(1)),

  `string.names` Array(String) CODEC(ZSTD(2)),
  `string.values` Array(String) CODEC(ZSTD(2)),
  `number.names` Array(String) CODEC(ZSTD(2)),
  `number.values` Array(Float64) CODEC(ZSTD(2)),
  `boolean.names` Array(String) CODEC(ZSTD(2)),
  `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

  INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
  ...
)
```

LowCardinality(String)

CODEC

What's the data model?

```
CREATE TABLE <table_name>
(
  _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
  _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
  _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
  _baselimeId FixedString(26) NOT NULL CODEC(ZSTD(1)),
  _ttl DateTime NOT NULL,

  _traceId String CODEC(ZSTD(1)),
  _requestId String CODEC(ZSTD(1)),
  _spanId String CODEC(ZSTD(1)),
  _coldStart UInt32 Codec(T64, LZ4),
  _parentSpanId String CODEC(ZSTD(1)),
  _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _duration UInt32 Codec(T64, LZ4),

  _source String NOT NULL CODEC(ZSTD(1)),

  `string.names` Array(String) CODEC(ZSTD(2)),
  `string.values` Array(String) CODEC(ZSTD(2)),
  `number.names` Array(String) CODEC(ZSTD(2)),
  `number.values` Array(Float64) CODEC(ZSTD(2)),
  `boolean.names` Array(String) CODEC(ZSTD(2)),
  `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

  INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
  ...
)
```

LowCardinality(String)

CODEC

Frequently queried files

What's the data model?

```
CREATE TABLE <table_name>
(
  _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
  _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
  _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
  _baselineId FixedString(26) NOT NULL CODEC(ZSTD(1)),
  _ttl DateTime NOT NULL,

  _traceId String CODEC(ZSTD(1)),
  _requestId String CODEC(ZSTD(1)),
  _spanId String CODEC(ZSTD(1)),
  _coldStart UInt32 Codec(T64, LZ4),
  _parentSpanId String CODEC(ZSTD(1)),
  _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _duration UInt32 Codec(T64, LZ4),

  _source String NOT NULL CODEC(ZSTD(1)),

  `string.names` Array(String) CODEC(ZSTD(2)),
  `string.values` Array(String) CODEC(ZSTD(2)),
  `number.names` Array(String) CODEC(ZSTD(2)),
  `number.values` Array(Float64) CODEC(ZSTD(2)),
  `boolean.names` Array(String) CODEC(ZSTD(2)),
  `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

  INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
  ...
)
```

LowCardinality(String)

CODEC

Frequently queried files

<type>.name
<type>.value

<type>.name | <type>.value

```
{
  "message": "Sending to the backend",
  "extra": {
    "source": "vega",
    "sampled": true,
    "num": 143
  },
  "time": "2023-05-24T20:16:50.722Z",
  "level": "info",
  "@requestId": "b8398e09-030a-49bc-b19c-86c48d5a5b09"
}
```


<type>.name | <type>.value

```
{
  "message": "Sending to the backend",
  "extra": {
    "source": "vega",
    "sampled": true,
    "num": 143
  },
  "time": "2023-05-24T20:16:50.722Z",
  "level": "info",
  "@requestId": "b8398e09-030a-49bc-b19c-86c48d5a5b09"
}
```



```
string.names
["message", "extra.source", "time", "level", "@requestId"]

string.values
[
  "Sending to the backend",
  "vega",
  "2023-05-24T20:16:50.722Z",
  "info",
  "b8398e09-030a-49bc-b19c-86c48d5a5b09"
]
```

<type>.name | <type>.value

```
{
  "message": "Sending to the backend",
  "extra": {
    "source": "vega",
    "sampled": true,
    "num": 143
  },
  "time": "2023-05-24T20:16:50.722Z",
  "level": "info",
  "@requestId": "b8398e09-030a-49bc-b19c-86c48d5a5b09"
}
```

<type>.name | <type>.value

```
{
  "message": "Sending to the backend",
  "extra": {
    "source": "vega",
    "sampled": true,
    "num": 143
  },
  "time": "2023-05-24T20:16:50.722Z",
  "level": "info",
  "@requestId": "b8398e09-030a-49bc-b19c-86c48d5a5b09"
}
```




```
indexOf(`string.names`, 'extra.source')
```

<type>.name | <type>.value

```
{
  "message": "Sending to the backend",
  "extra": {
    "source": "vega",
    "sampled": true,
    "num": 143
  },
  "time": "2023-05-24T20:16:50.722Z",
  "level": "info",
  "@requestId": "b8398e09-030a-49bc-b19c-86c48d5a5b09"
}
```



```
indexOf(`string.names`, 'extra.source')
```



```
`string.values`[indexOf(`string.names`, 'extra.source')] = 'vega'
```

<type>.name | <type>.value

```
SELECT
    MAX(`number.values`[indexOf(`number.names`, 'extra.num')]) as `MAX(extra.num)`
FROM
    <table>
WHERE
    `string.values`[indexOf(`string.names`, 'extra.source')] = 'vega';
```


What's the data model?

```
CREATE TABLE <table_name>
(
  _tenantId LowCardinality(String) NOT NULL CODEC(ZSTD(1)),
  _timestamp DateTime64(3) NOT NULL CODEC(DoubleDelta, LZ4),
  _now DateTime64(3) NOT NULL CODEC(ZSTD(1)),
  _baselineId FixedString(26) NOT NULL CODEC(ZSTD(1)),
  _ttl DateTime NOT NULL,

  _traceId String CODEC(ZSTD(1)),
  _requestId String CODEC(ZSTD(1)),
  _spanId String CODEC(ZSTD(1)),
  _coldStart UInt32 Codec(T64, LZ4),
  _parentSpanId String CODEC(ZSTD(1)),
  _startTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _endTime DateTime64(3) CODEC(DoubleDelta, LZ4),
  _duration UInt32 Codec(T64, LZ4),

  _source String NOT NULL CODEC(ZSTD(1)),

  `string.names` Array(String) CODEC(ZSTD(2)),
  `string.values` Array(String) CODEC(ZSTD(2)),
  `number.names` Array(String) CODEC(ZSTD(2)),
  `number.values` Array(Float64) CODEC(ZSTD(2)),
  `boolean.names` Array(String) CODEC(ZSTD(2)),
  `boolean.values` Array(UInt8) CODEC(ZSTD(2)),

  INDEX _idx_tenantId TYPE bloom_filter GRANULARITY 4,
  ...
)
```

LowCardinality(String)

CODEC

Frequently queried files

<type>.name
<type>.value

bloom_filter

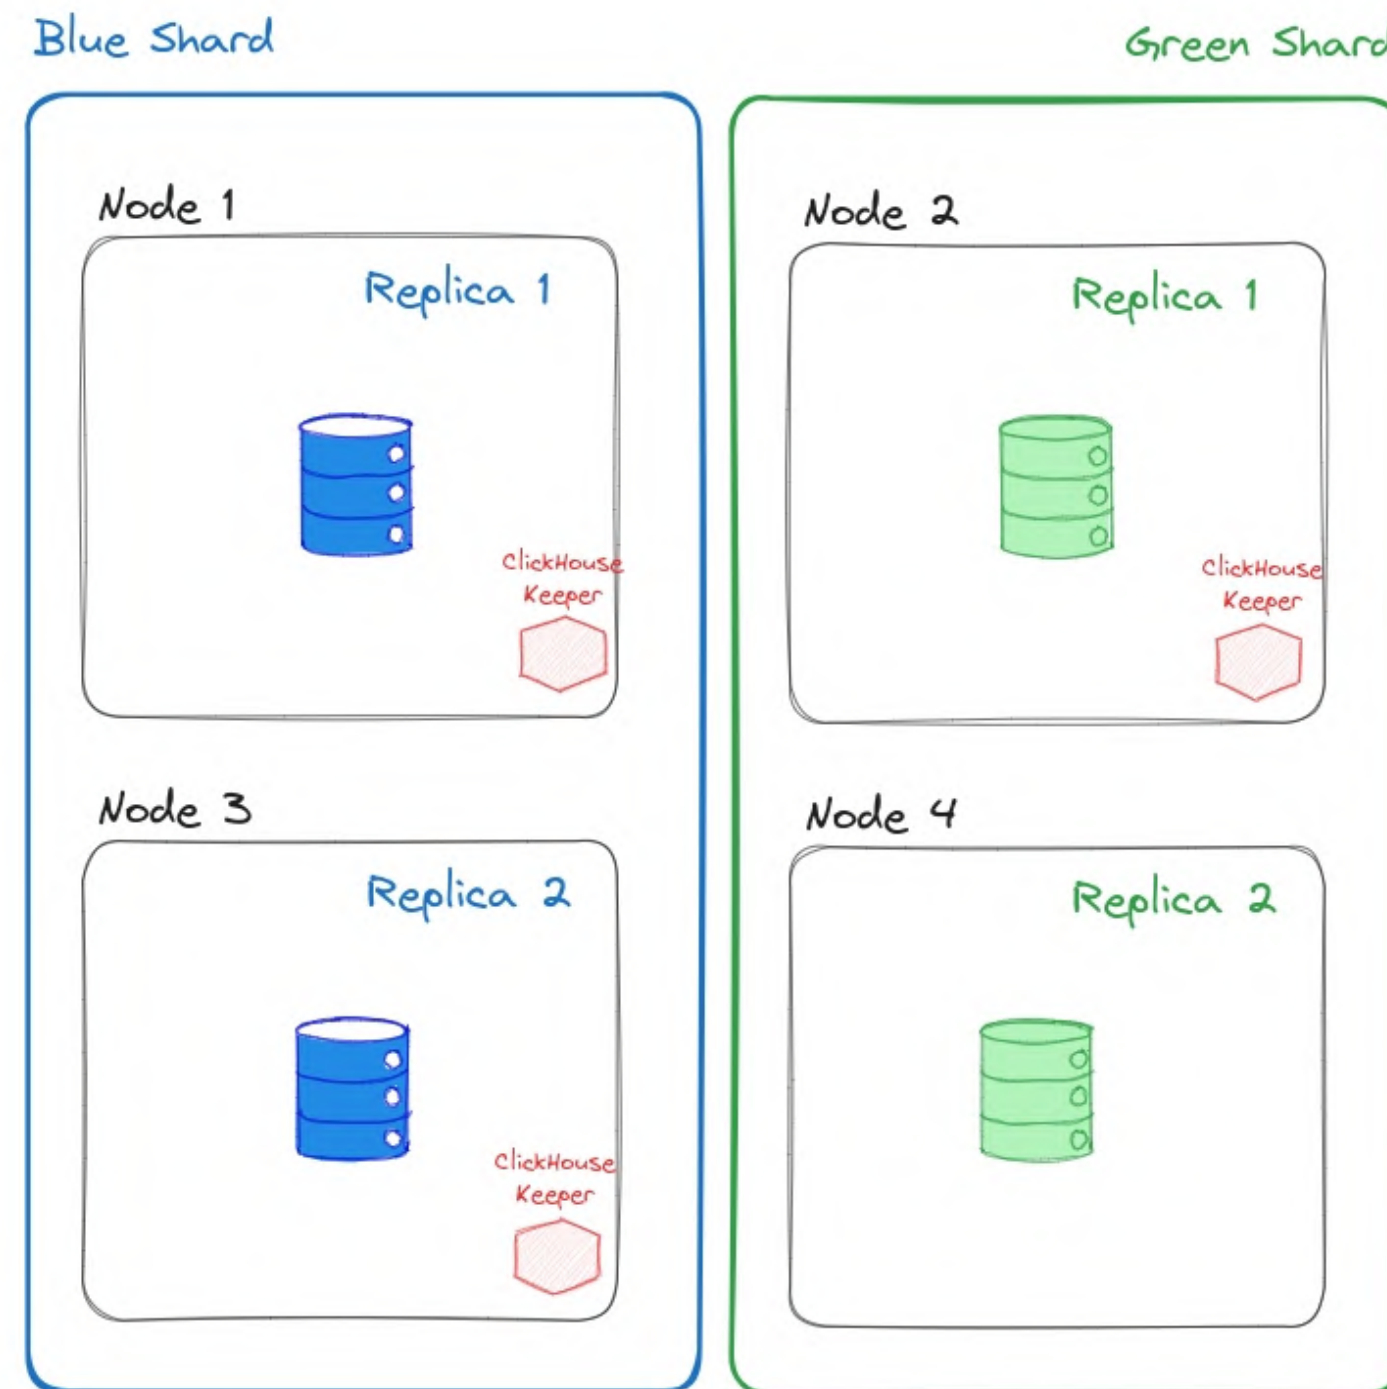
What's the engine?

Replicated Merge Tree

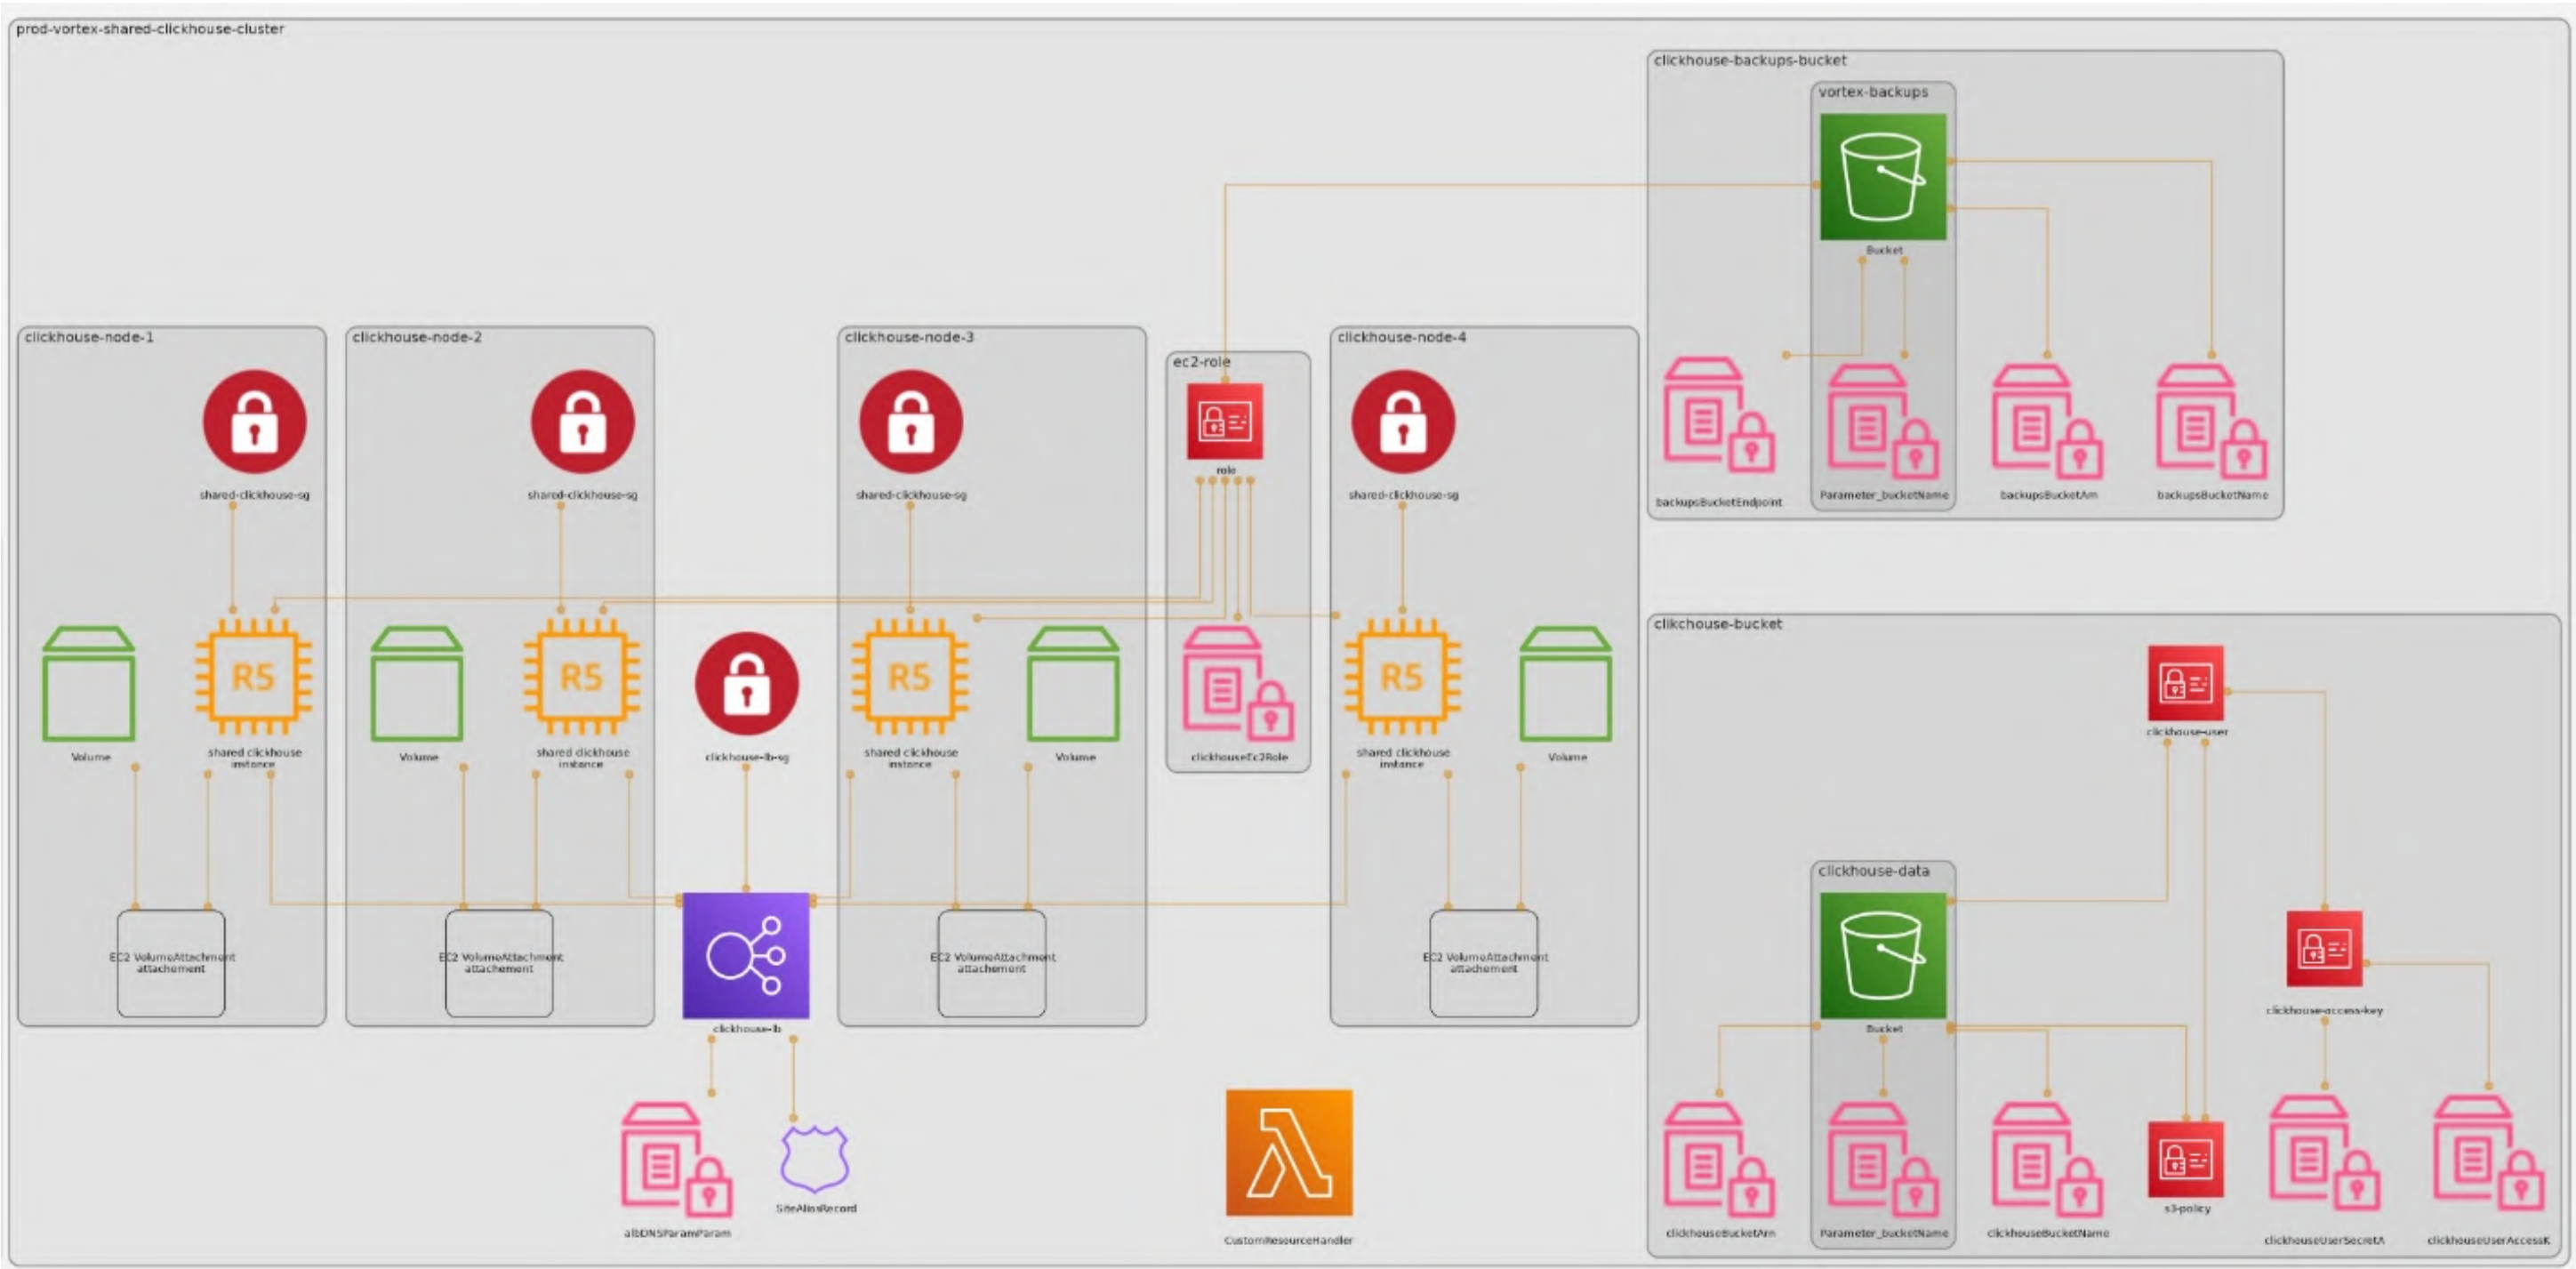
```
ENGINE = ReplicatedMergeTree('/clickhouse/tables/{cluster}/{shard}/{table}', '{replica}')  
PARTITION BY toDate(_timestamp)  
ORDER BY (_tenantId, _traceId, _timestamp)  
TTL _ttl DELETE;
```

Sort Key

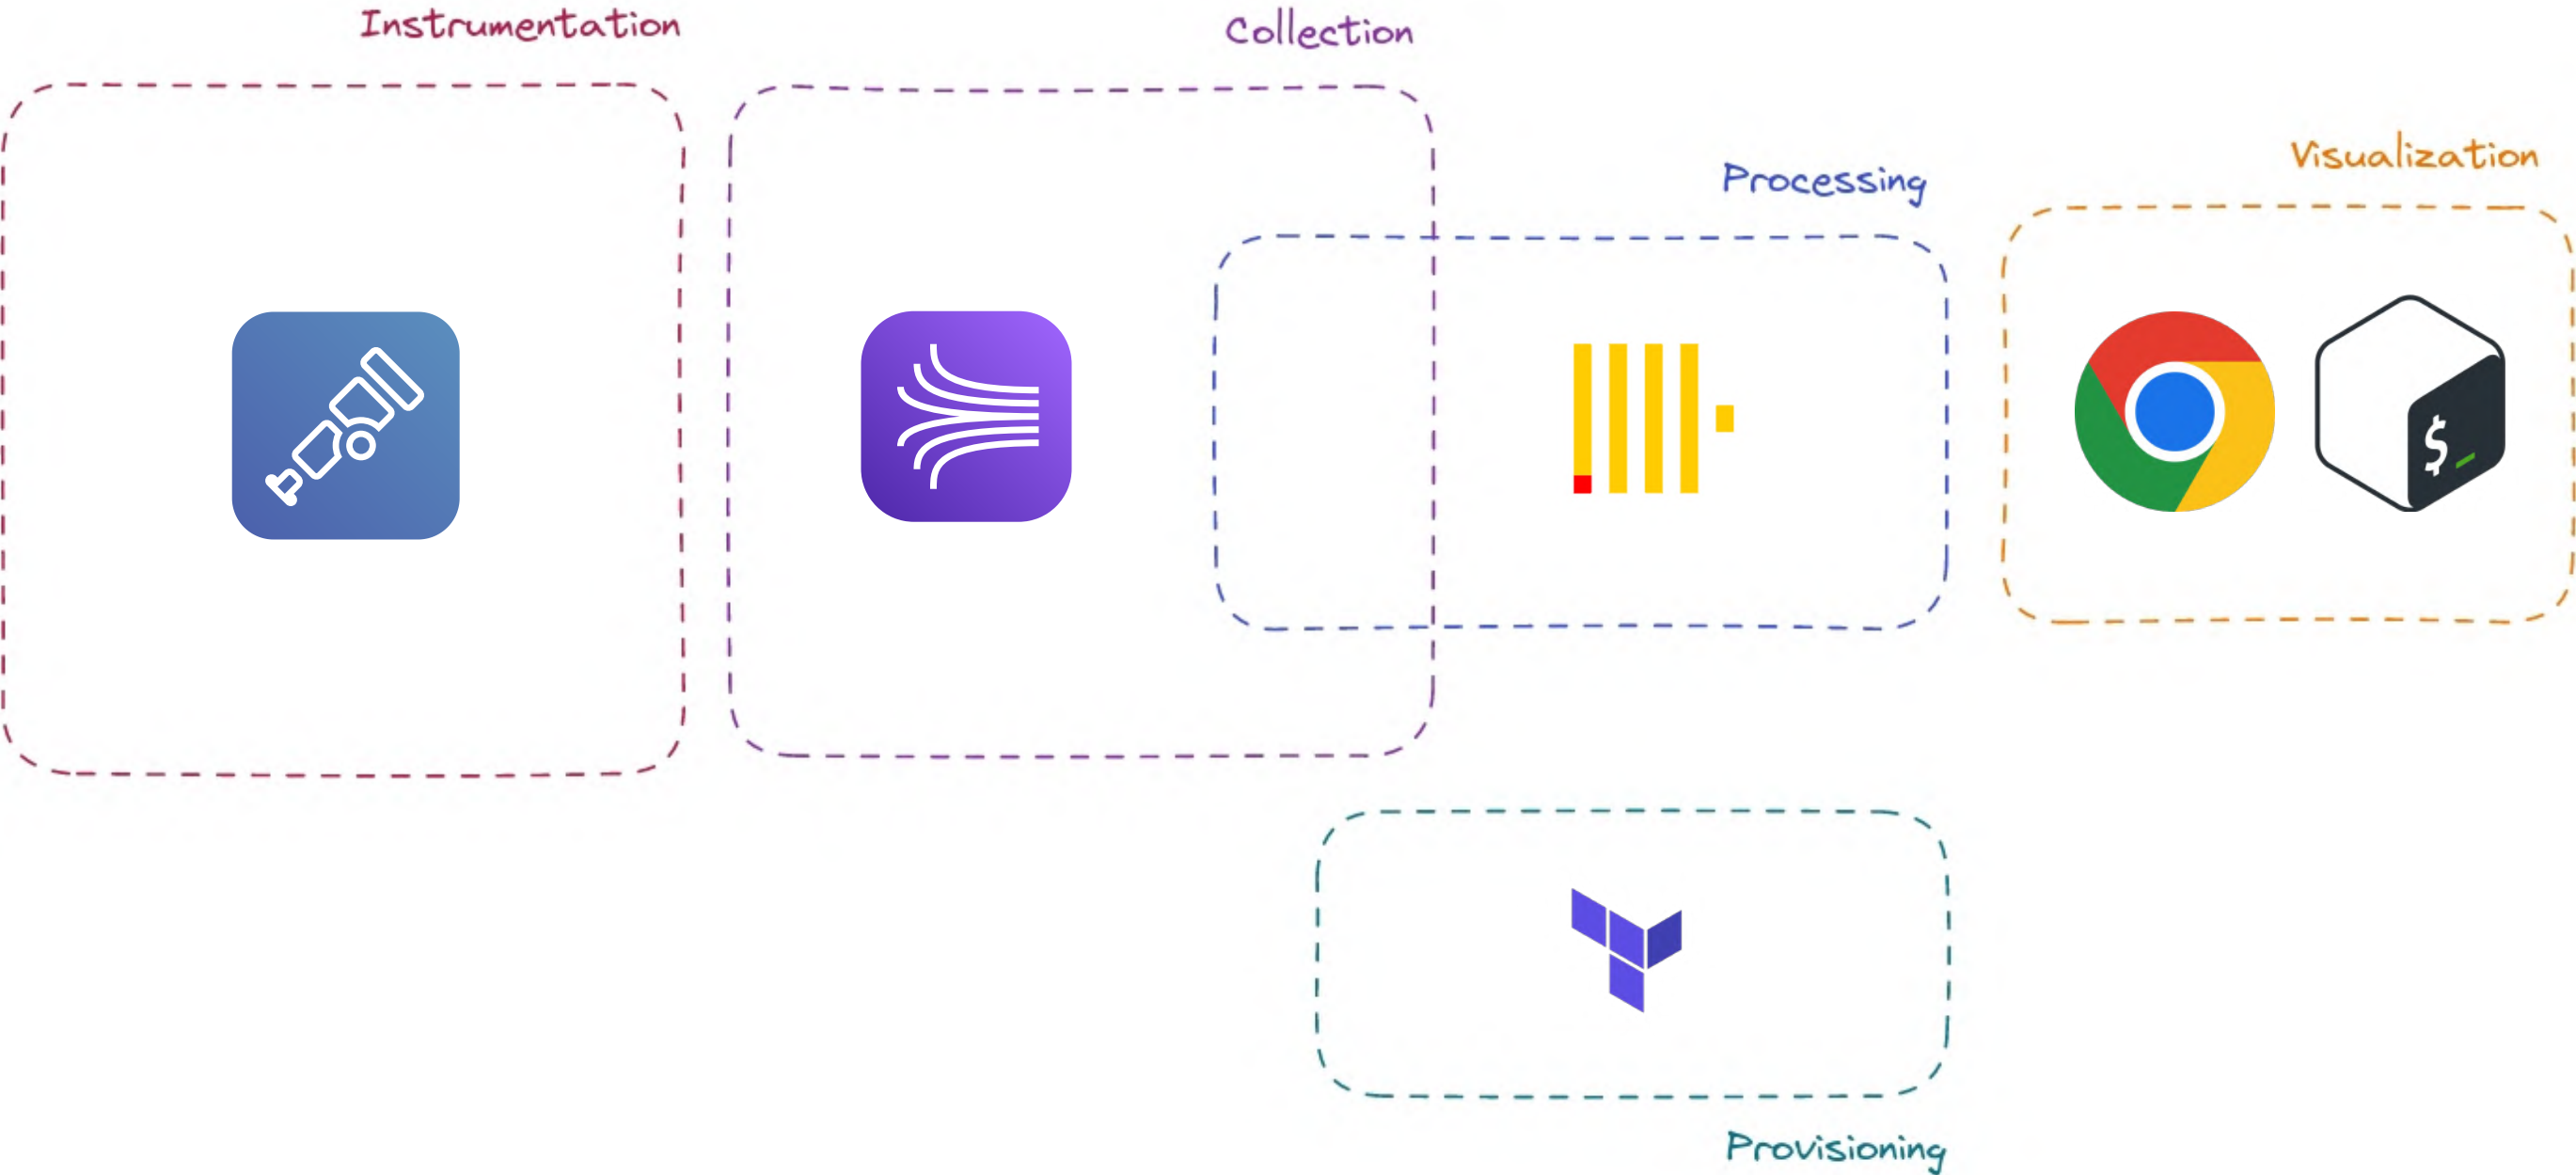
What's the architecture?



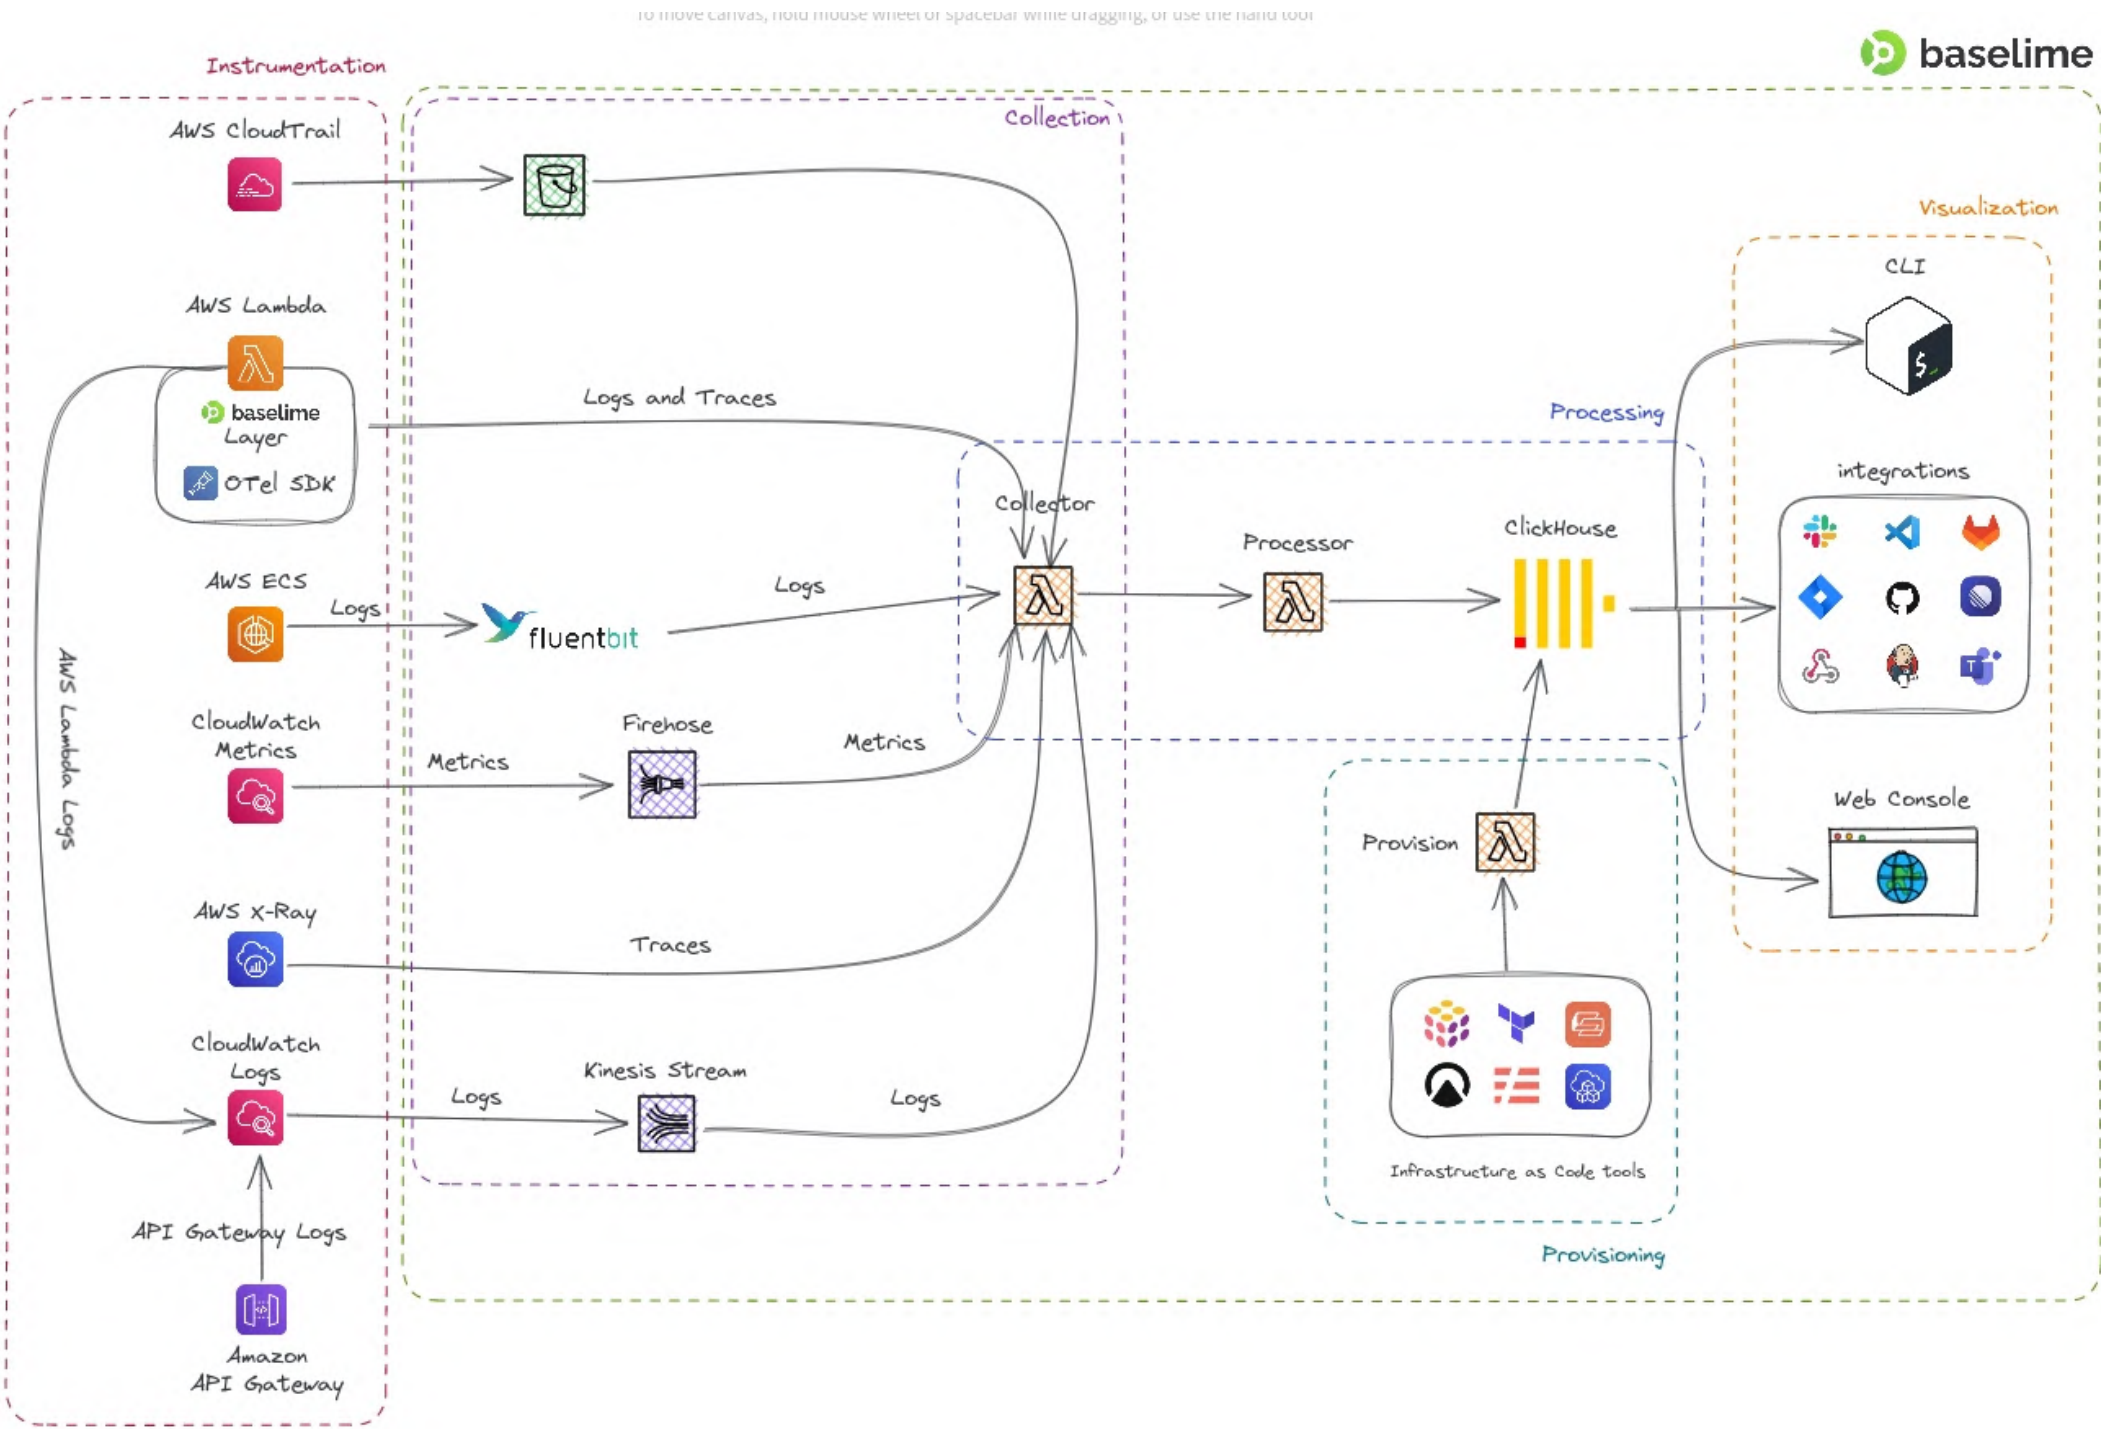
What's the architecture?



How do we get there?



Baseline architecture



Thanks for having me!

Boris Tane

baseline.io
@boristane