



Clickhouse Austin Meetup

Tony Burke



Clickhouse at SolarWinds

What are the characteristics of our Clickhouse use



SaaS Platform engineering

- Ingestion
 - Received Messages
 - Transformation
 - Queueing
- Services
 - Ingress Endpoint
 - Clickhouse Writes
 - Clickhouse Query
- Clickhouse
 - Schema
 - Management
 - Support

About our data

- Telemetry (Metrics, Logs and Traces)
 - Devices (hosts, firewalls, switches, logs)
 - Applications (traces, logs, and metrics)
 - Databases (SQL statements, metrics)
 - Kubernetes (metrics, logs)
 - ...
- Real-time customer-facing multi-tenant
- ~3M messages ingested per second (1 cluster)
 - ~550MB – 1G per second


Research: Metrics Primary Key

Can we change the primary key time boundary to improve hot queries



First, Understand How Clickhouse

- Stores and manages data
 - MergeTree table engine
 - Physical layout
- Resolve queries
 - Select Rows
 - Indexing
- What are the query characteristics
 - Frequent time-range queries

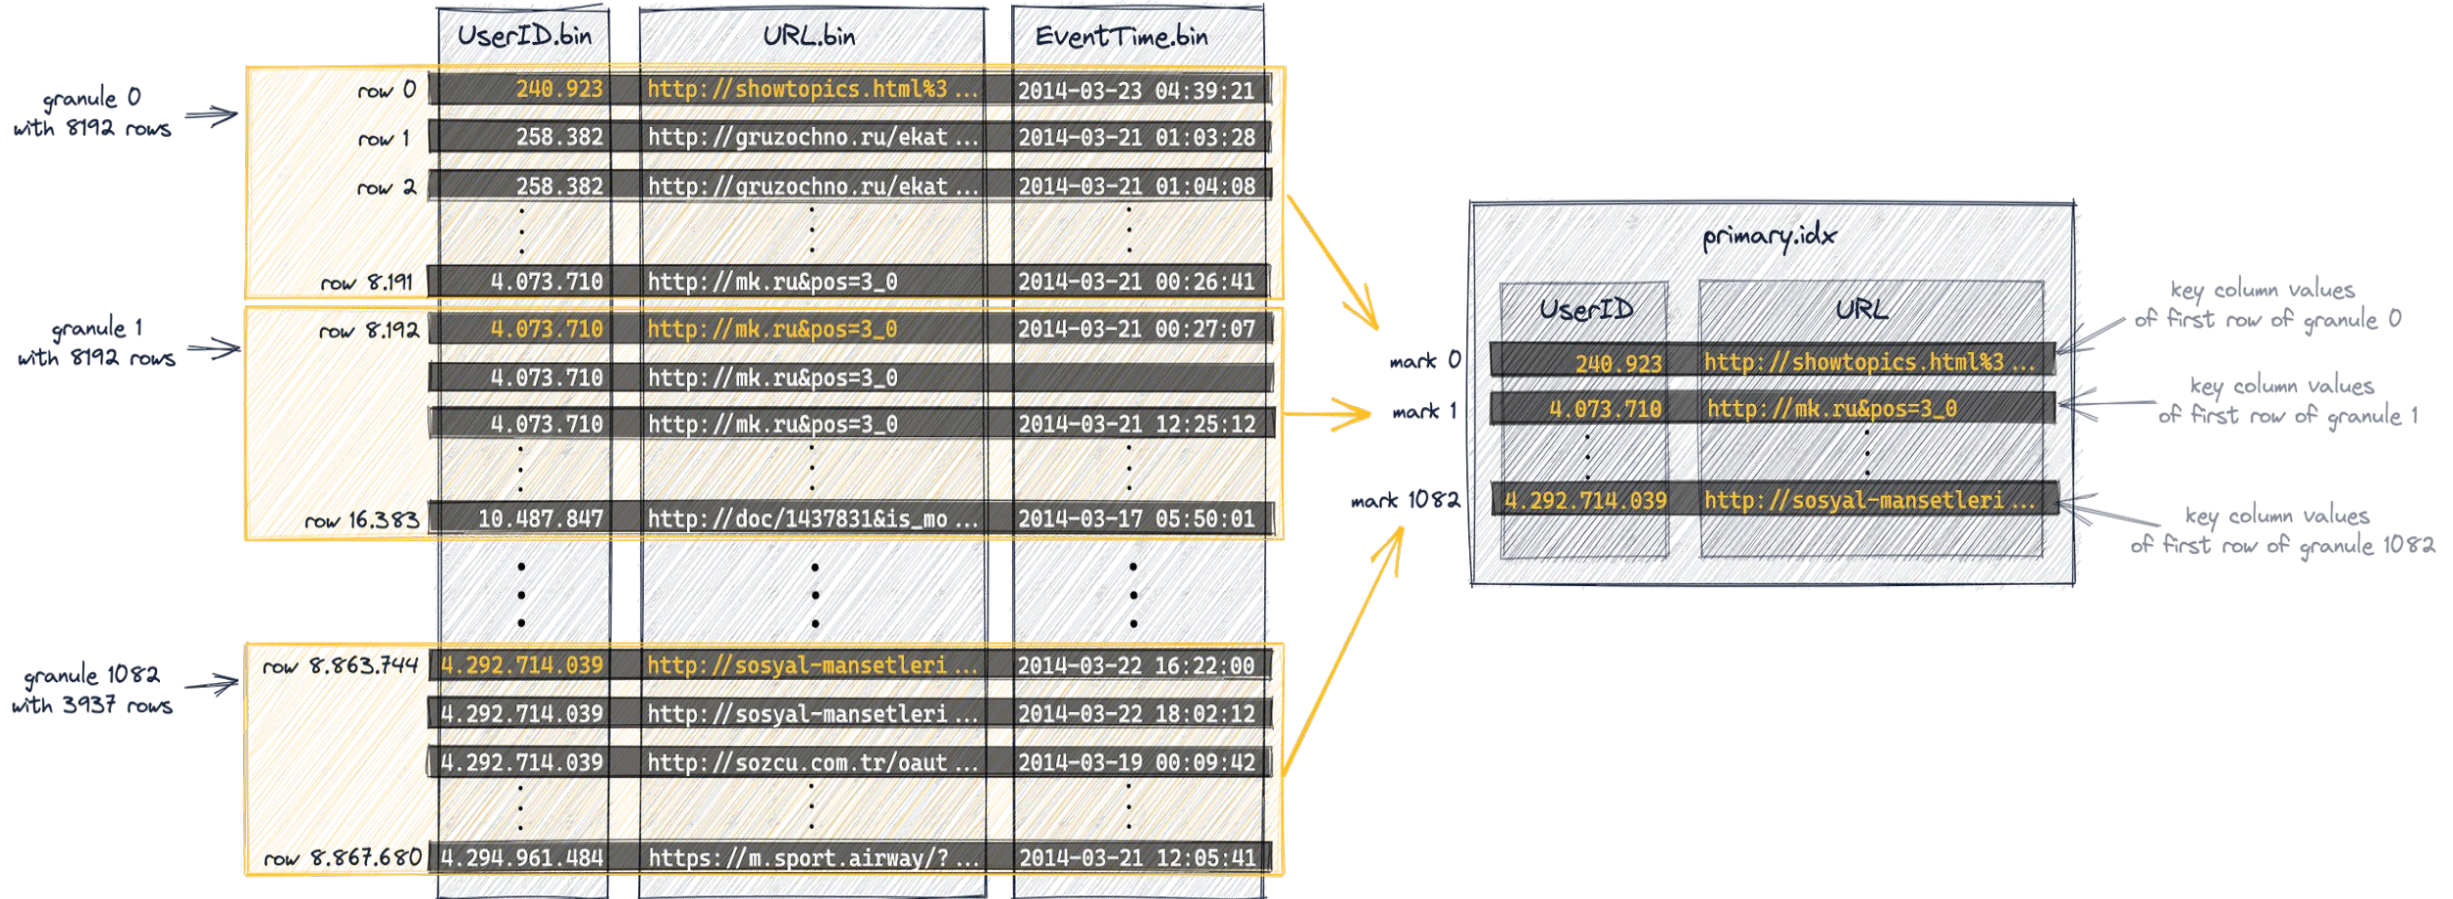


Most queries are for the latest 60 minutes

- Interactive users
- Alert evaluations

Clickhouse Indexing

Contents of a part with index_granularity of 8192 (default)



Primary Key Effectiveness

How would a change to the primary key (time-bound) affect performance



Analysis

- The last 60min is the hot query data
 - Not a lot of time for parts to merge
- The primary.idx
 - Few marks for the latest data
 - Cardinality of the first 2 columns are high.
(tenant, namespace, hour)
- Conclusion:
 - Any change to the primary key would have little effect on reducing the marks selected.

Reduce Index_granularity

- Use a Materialized View to mirror table data
 - The only difference is index_granularity
- Use our Clickhouse query service to
 - Alternate queries between the two tables
- Measure
 - Query Performance
 - Background Tasks (Merges)
 - Memory usage

query_log metrics

Measure queries for both tables and compare results



index_granularity=1024

queries	avg_cpu_seconds	avg_query_time_ms	avg_read_parts	avg_read_marks	avg_read_rows	time_window_in_mins
265602	0.07	76.52	9.41	9.45	9555.86	5
186923	0.09	98.38	9.53	10.31	10553.26	15
160668	0.09	92.2	9.43	9.84	10073.83	2
59489	0.08	79.8	9.37	10.3	10528.28	1
49987	0.09	82.16	9.21	11.2	12191.19	60
17793	0.1	117.82	9.78	13.19	14224.2	30

index_granularity=32

queries	avg_cpu_seconds	avg_query_time_ms	avg_read_parts	avg_read_marks	avg_read_rows	time_window_in_mins
200246	0.03	27.88	9.91	11.63	374.05	5
140138	0.07	59.05	10.13	32.15	1378.26	15
120856	0.05	44.88	9.93	23.64	957.41	2
45007	0.07	50.58	9.86	26.26	1069.82	1
38591	0.12	93.55	9.84	73.5	3542.48	60
15520	0.11	106.97	10.15	79.7	3126.16	30

Observations

The side effects of reducing the index_granularity



What should I watch?

- An increase in merge time
 - `peak_memory_usage` and `duration_ms` in **`system.part_log`**
- Primary index file size
 - `primary_key_bytes_in_memory` and `primary_key_bytes_in_memory_allocated` in **`system.parts`**
- Column mark size
 - `mark_cache_size` (default 5G) LRU cache
 - `MarkCacheHits` and `MarkCacheMisses` in **`system.events`**

Other recent tuning

- Physical file reads/writes
 - `local_filesystem_read_method`
 - The default setting is `pread_threadpool` is used to decrease the chances of exhausting the open files limit.
 - This negatively impacted filesystem reads in our environments. We have high number of concurrent queries and fast SSD drives. We set this value to `pread` and saw a nice performance improvement.

A vertical column of eight yellow dots on the left side of the slide.

THANK YOU