



From batch processing to streaming


Transitioning Private Cloud ClickHouse OSS to AWS Cloud Native

Intro

Software engineer for 7 years, recently started into the world of Data Engineering as our product grew more complex

Ran startup agency, brought over 30 new SaaS products to market

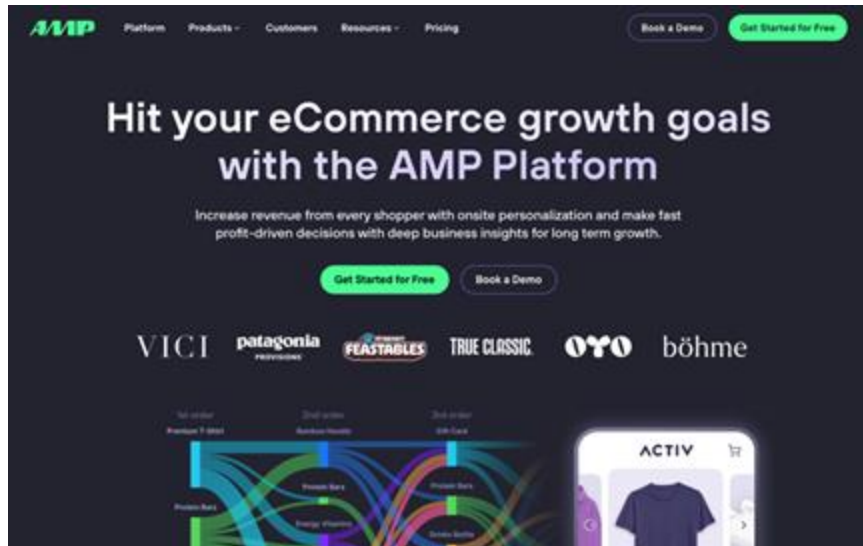
Love DnB, Psytrance, Scuba, MTB and Skiing!



AMP

AMP focuses on solving ecommerce problems

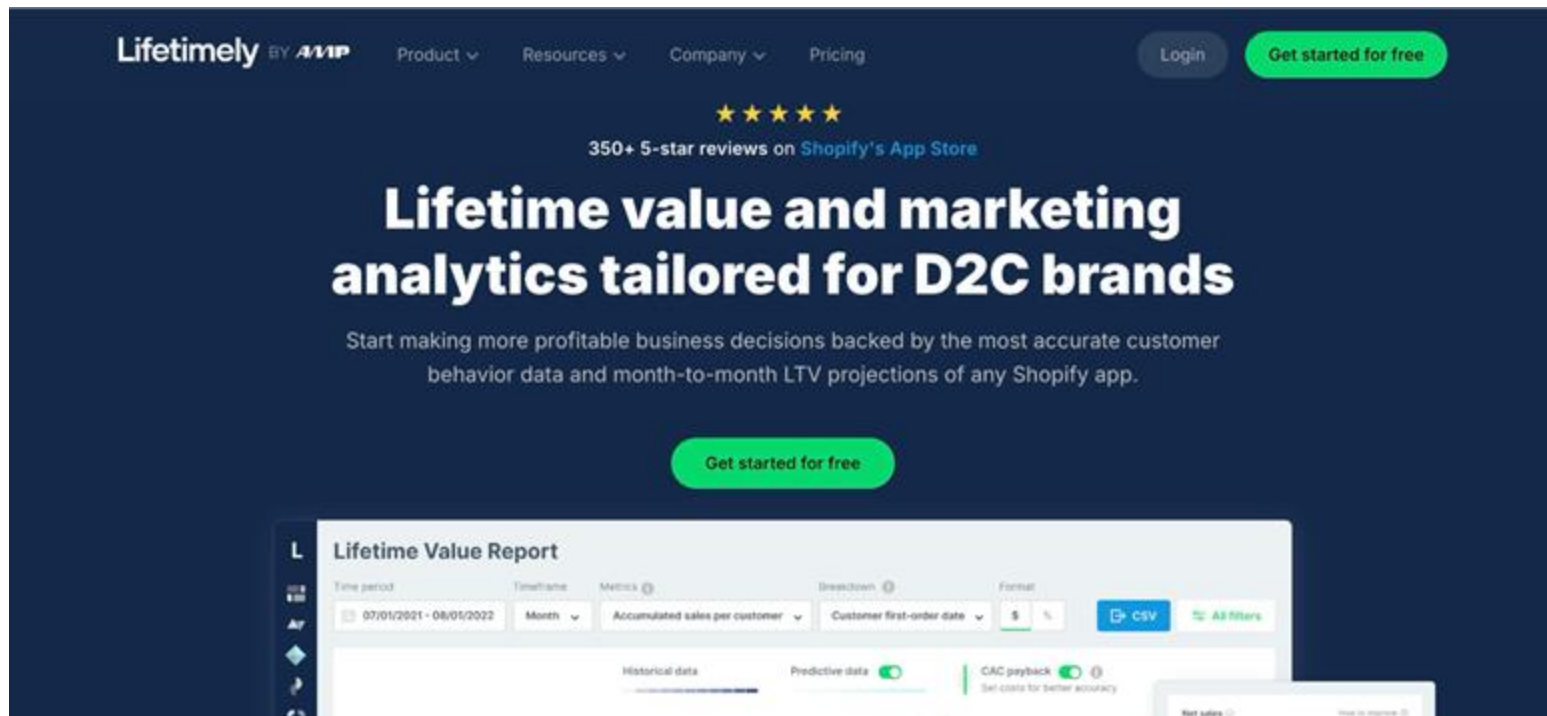
Team of ~30 based all over the world, mostly in APAC



Why ClickHouse?

Lifetimely

Our flagship analytics product



How we got here

PostgreSQL

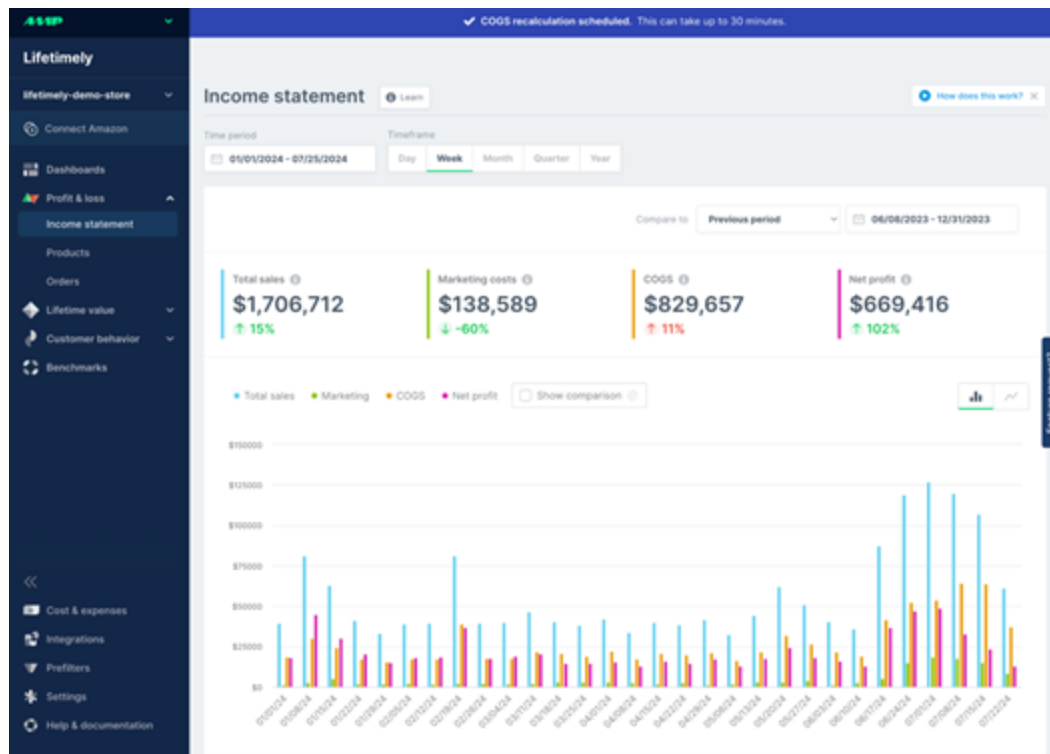
PG query performance degraded
once over ~100m rows per table

Filters

Dynamic filtering requires query
time computation for all reports

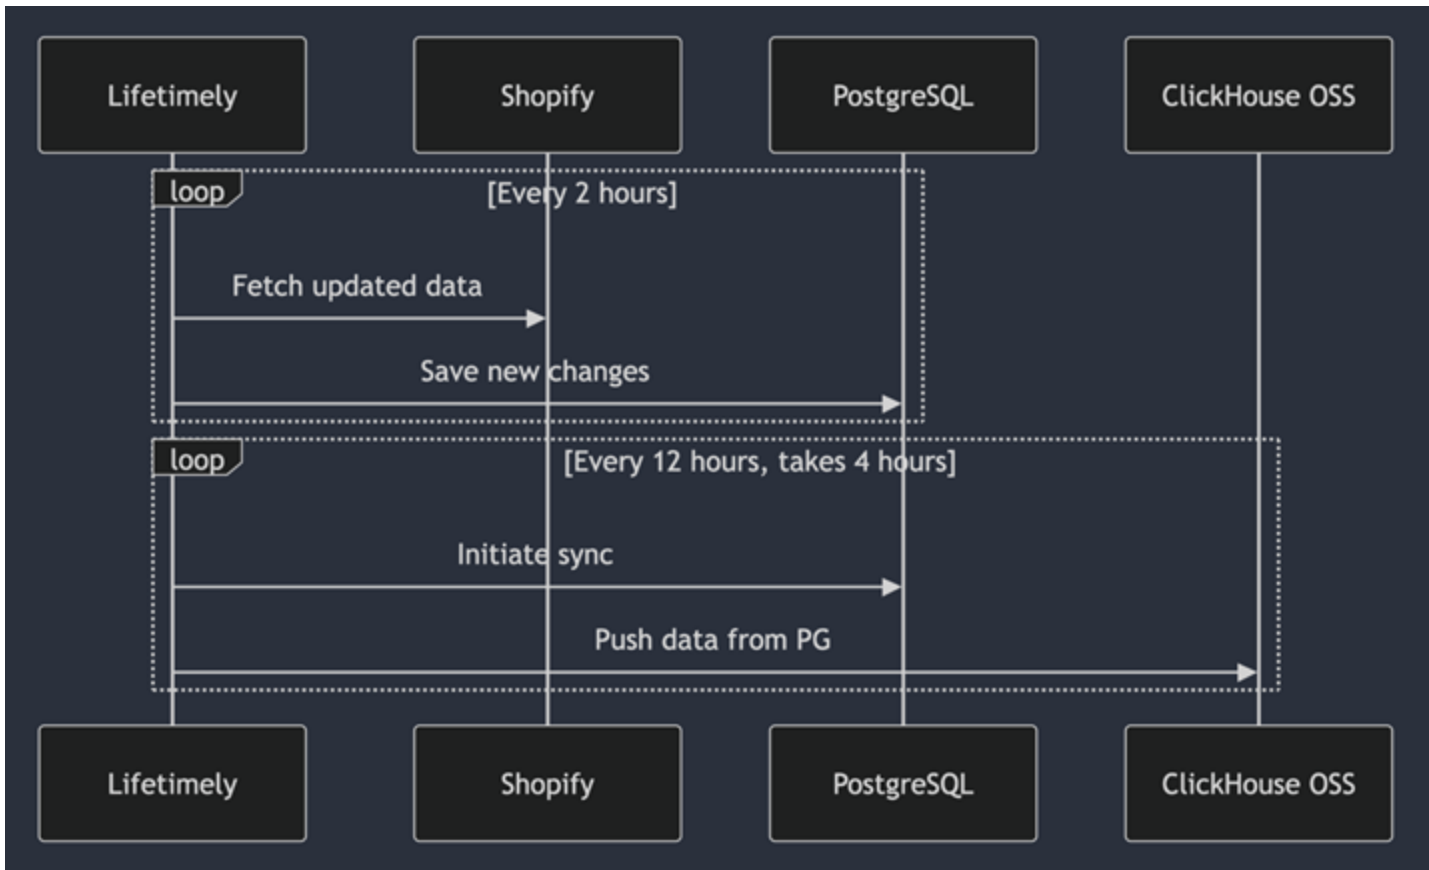
3,750+

Shopify stores ingesting and
querying their data



**“Lets clone our PG
database to ClickHouse
for queries”**

What problems are we facing?



Batch processing workflow



Problem 1 – Data Freshness



Problem 2 – Resource usage

Where do we want to go?

Our goals – Customer

Freshness

Provide customer data in a realistic and timely manner

Integrity

An observable and testable data pipeline

Flexibility

Enrich data with external integrations with ease

Our goals – Technical

Replay

Store events in a data lake for backup and rebuild

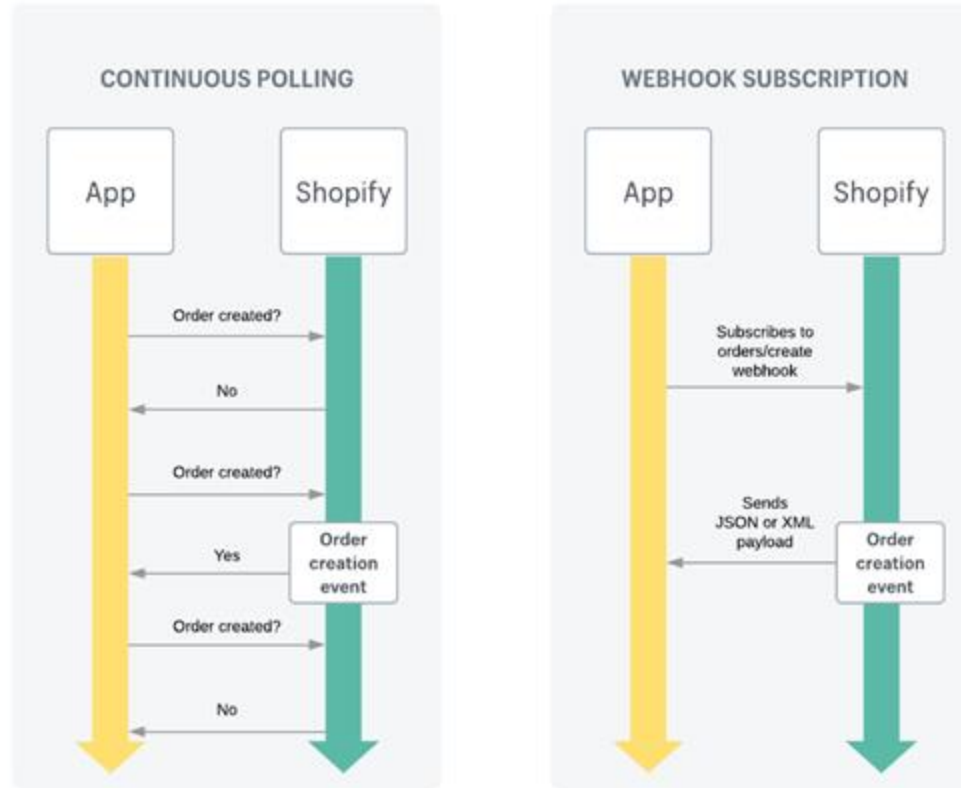
Webhooks

Handle streamed events sent at us from Shopify

IaC

Define and release our infrastructure with Terraform

Let's get streaming



Our data source – Shopify Webhooks



Shopify



Amazon EventBridge



AWS Lambda



Clickhouse Cloud

Solution 1: ReplacingMergeTree table engine

```
CREATE TABLE mySecondReplacingMT
(
    `key` Int64,
    `someCol` String,
    `eventTime` DateTime
)
ENGINE = ReplacingMergeTree(eventTime)
ORDER BY key;

INSERT INTO mySecondReplacingMT Values (1, 'first', '2020-01-01 01:01:01');
INSERT INTO mySecondReplacingMT Values (1, 'second', '2020-01-01 00:00:00');

SELECT * FROM mySecondReplacingMT FINAL;
```

key	someCol	eventTime
1	first	2020-01-01 01:01:01

**“We can insert every update
and let ClickHouse handle
deduplication”**

**“We can insert every update
and let ClickHouse handle
deduplication”**

Buuuut, we have a problem...

“Generally, we recommend inserting data in fairly large batches of at least 1,000 rows at a time, and ideally between 10,000 to 100,000 rows.”
– ClickHouse Docs

<https://clickhouse.com/docs/en/cloud/bestpractices/bulk-inserts>

Solution 2: AWS Data Firehose



Data Firehose allows buffering of JSON files into S3

Solution 2: AWS Data Firehose

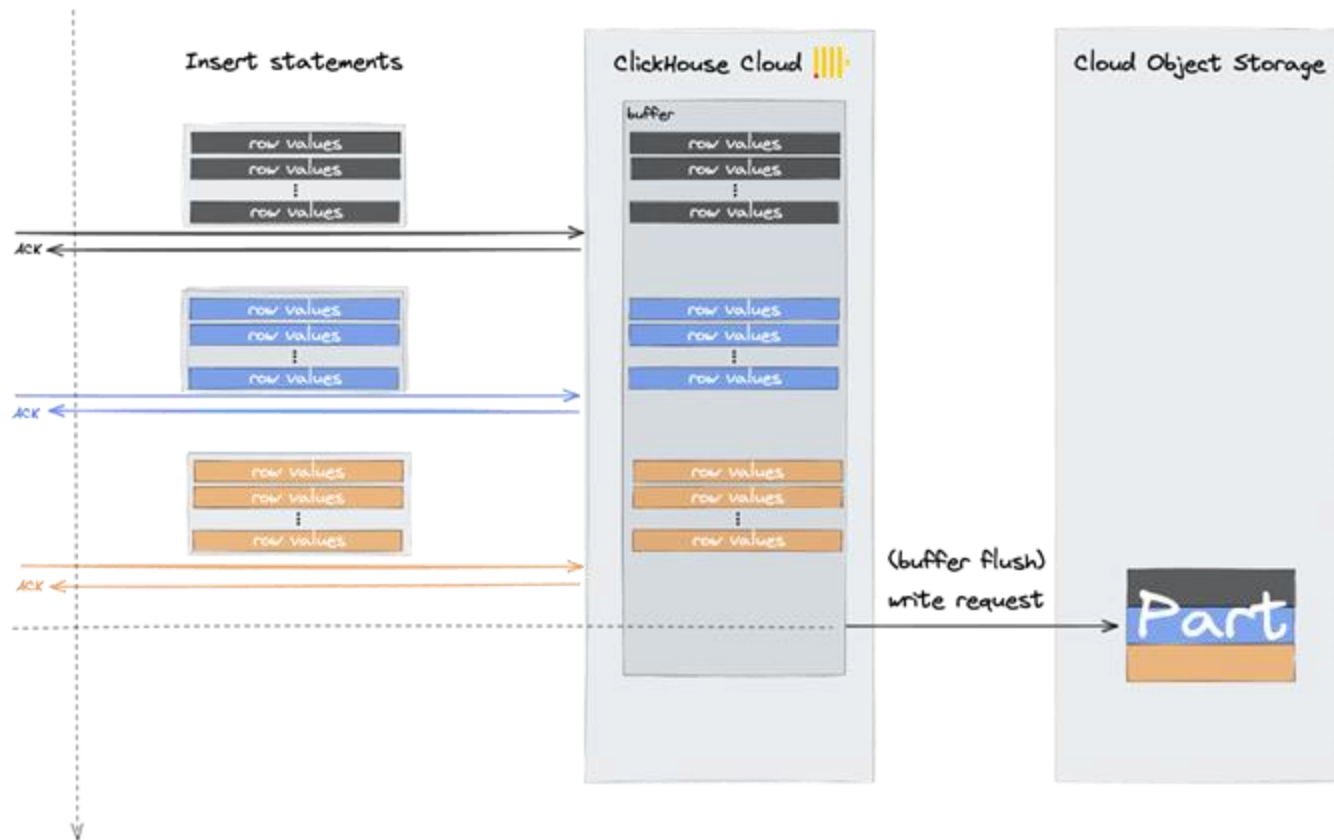


Data Firehose allows buffering of JSON files into S3

Buuuuut, we have another problem...

**“Our buffers are still under
1,000 rows”**

Solution 3: Async Inserts



Solution 3: Async Inserts

Inserts

Send insert of any size into ClickHouse

Memory

Inserts are stored in memory

Flush

Inserts flushed to disk when parameters are reached

```
const client : NodeClickHouseClient = createClient( config: {  
  url: process.env.CLICKHOUSE_URL,  
  username: process.env.CLICKHOUSE_USERNAME,  
  password: process.env.CLICKHOUSE_PASSWORD,  
  database: process.env.CLICKHOUSE_DATABASE,  
  application: 'clickhouse_insert_lambda',  
});  
  
export async function sendDataToClickHouse(table: Tables, stream: Readable) : Promise<InsertResult> {  
  try {  
    return await client.insert( params: {  
      table: table,  
      values: stream,  
      format: 'JSONEachRow',  
      clickhouse_settings: {  
        async_insert: 1,  
        async_insert_max_data_size: '104857600', // 100MiB  
        async_insert_busy_timeout_min_ms: 15 * 1000, // 15 seconds  
        async_insert_busy_timeout_max_ms: 45 * 1000, // 45 seconds  
        async_insert_use_adaptive_busy_timeout: 1,  
        async_insert_busy_timeout_increase_rate: 0.2,  
        wait_for_async_insert: 0,  
      },  
    });  
  } catch (error) {  
    logger.error( obj: { error }, msg: 'Error sending data to ClickHouse');  
    throw error;  
  }  
}
```

Solution 4: ClickHouse Cloud Terraform provider

- Create reproducible ClickHouse instances across all our environments
- Spin up services in CI/CD for testing
- Monitoring and automatic updates
- Autoscaling

```
resource "clickhouse_service" "this" {  
  name           = "clickhouse-cloud"  
  cloud_provider = "aws"  
  region         = "us-east-1"  
  tier           = "production"  
  idle_scaling   = false  
  
  max_total_memory_gb = 240  
  min_total_memory_gb = 72  
  
  password = var.service_password  
}
```



Demo!