

# ClickHouse and The One Billion Row Challenge

Derek Chia, Senior Support Engineer @ ClickHouse



||||· ClickHouse

<https://clickhouse.com/blog/clickhouse-one-billion-row-challenge>

# ClickHouse and The One Billion Row Challenge


1 🐛 - <https://github.com/gunnarmorling/1brc>

## The original task

- Write a *Java program* which reads the file, calculates the **min**, **mean**, and **max** temperature value per weather station, and emits the results on stdout like this:
  - ◆ **sorted** alphabetically by **station name**, and the result values per station
  - ◆ in the format `<min>/<mean>/<max>`, **rounded** to one fractional digit

```
Hamburg;12.0  
Bulawayo;8.9  
Palembang;38.8  
St. John's;15.2  
Cracow;12.6  
Bridgetown;26.9  
...
```

`<Station name>;<average temperature>`



```
{Abha=-23.0/18.0/59.2,  
Abidjan=-16.2/26.0/67.3,  
Abéché=-10.0/29.4/69.0,  
Accra=-10.1/26.4/66.4, Addis  
Ababa=-23.7/16.0/67.0,  
Adelaide=-27.8/17.3/58.5, ...}
```

`{<Station name>;<min>/<mean>/<max>, ...}`



# ClickHouse and The One Billion Row Challenge

Generating dataset using Java and Python

- The repo has a tool (create\_measurements.{sh,py}) to generate **one billion random points** by sampling a **Gaussian distribution with a mean and variance of 10** using a list of (413) distinct stations and their average temperatures

```
Hamburg;12.0  
Bulawayo;8.9  
Palembang;38.8  
St. John's;15.2  
Cracow;12.6  
Bridgetown;26.9  
...
```

```
<Station name>;  
<average  
temp.>
```

```
# clone and build generation tool. Output omitted.  
git clone git@github.com:gunnarmorling/1brc.git  
./mvnw clean verify  
./create_measurements.sh 1000000000
```

Created file with 1,000,000,000 measurements in **435900 ms**

Java - 435900 ms = 435.900 s = 7.27 mins (too slow!)

```
python3 create_measurements.py 1_000_000_000  
Estimated max file size is: 14.8 GiB.  
Building test data...  
[=====] 100%  
Test data successfully written to 1brc/data/measurements.txt  
Actual file size: 14.8 GiB  
Elapsed time: 13 minutes 36 seconds  
Test data build complete.
```

Python - ~13 mins (wayyy too slow!)

```
> head -n 3 data/measurements.txt  
Monatéle;65.2  
Glendora;-13.0  
Gundumāl;-67.2
```

```
> wc -l data/measurements.txt  
1000000000 data/measurements.txt
```

```
> ls -lh data/measurements.txt |  
awk '{print $5, $9}'  
15G data/measurements.txt
```

measurements.txt with 1B rows,  
taking up approx. 15 GB



# ClickHouse and The One Billion Row Challenge

## Generating dataset using ClickHouse

```
INSERT INTO FUNCTION file('measurements.csv', CustomSeparated)
WITH (
    SELECT groupArray((station, avg))
    FROM s3('https://datasets-documentation.s3.eu-west-3.amazonaws.com/1brc/stations.csv')
) AS averages
SELECT
    averages[floor(randUniform(1, length(averages))::Int64).1] as city,
    round(averages[floor(randUniform(1, length(averages))::Int64).2] +
    (10 * SQRT(-2 * LOG(randCanonical(1))) * COS(2 * PI() * randCanonical(2))), 2) as temperature
FROM numbers(1_000_000_000)
SETTINGS format_custom_field_delimiter=';', format_custom_escaping_rule='Raw'
```

```
derek-clickhouse :) INSERT INTO FUNCTION file('measurements.csv', CustomSeparated)
SETTINGS format_custom_field_delimiter = ';', format_custom_escaping_rule = 'Raw', max_threads = 20
WITH (
    SELECT groupArray((station, avg))
    FROM s3('https://datasets-documentation.s3.eu-west-3.amazonaws.com/1brc/stations.csv')
) AS averages
SELECT
    (averages[CAST(floor(randUniform(1, length(averages))), 'Int64')].1 AS city,
    round(((averages[CAST(floor(randUniform(1, length(averages))), 'Int64')].2) + ((10 * SQRT(-2 *
LOG(randCanonical(1)))) * COS((2 * PI()) * randCanonical(2))), 2) AS temperature
FROM numbers(1000000000)
SETTINGS format_custom_field_delimiter = ';', format_custom_escaping_rule = 'Raw'
```

Query id: 185b00d6-d815-4384-8ae4-b2703b217b1d

Ok.

0 rows in set. Elapsed: 69.459 sec. Processed 1.00 billion rows, 8.00 GB (14.40 million rows/s., 115.18 MB/s.)  
Peak memory usage: 49.44 MiB.

ClickHouse - ~1 min (fast!)

# ClickHouse and The One Billion Row Challenge

Generating dataset using ClickHouse, explained

file table function: writes data into the local (host) filesystem

INSERT  
SELECT

```
INSERT INTO FUNCTION file('measurements.csv', CustomSeparated)
WITH (
  SELECT groupArray((station, avg))
  FROM s3('https://datasets-documentation.s3.eu-west-3.amazonaws.com/1brc/stations.csv')
) AS averages
SELECT
  averages[floor(randUniform(1, length(averages))::Int64).1] as city,
  round(averages[floor(randUniform(1, length(averages))::Int64).2] +
    (10 * SQRT(-2 * LOG(randCanonical(1))) * COS(2 * PI() * randCanonical(2))), 2) as temperature
FROM numbers(1_000_000_000)
SETTINGS format_custom_field_delimiter=';', format_custom_escaping_rule='Raw'
```

groupArray(): creates an array of (station, avg). e.g. [('Abha',18),('Abidjan',26), ...]

numbers table function:  
table with the single 'number'  
column (UInt64) that contains  
integers from 0 to N-1

randCanonical function:  
returns a random Float64 number

p.s. We use the randCanonical function and use this to  
sample the Guassian distribution using a Muller transform.



# ClickHouse and The One Billion Row Challenge

The challenge, baseline using Java and ClickHouse

```
time ./calculate_average_baseline.sh
real    4m41.360s
user    4m38.427s
sys     0m4.728s
```

Java: *4m 41.360s (slow!)*

```
SELECT format('{ }={ }/{ }/{ }', city, min(temperature), round(avg(temperature), 2), max(temperature))
FROM file('measurements.csv', CSV, 'city String, temperature DECIMAL(8,1)')
GROUP BY city
ORDER BY city ASC
FORMAT CustomSeparated
SETTINGS
    format_custom_result_before_delimiter = '{',
    format_custom_result_after_delimiter = '}',
    format_custom_row_between_delimiter = ', ',
    format_custom_row_after_delimiter = '\n',
    format_csv_delimiter = ';';
```

412 rows in set. Elapsed: **23.827 sec**. Processed 1.00 billion rows, 14.68 GB (41.97 million rows/s., 616.08 MB/s.)  
Peak memory usage: 183.29 MiB.

Reading from CSV performs complete linear scan of the file. This is inefficient

ClickHouse local: *23.827s (faster, but can be better)*



# ClickHouse and The One Billion Row Challenge

The challenge, optimizing ClickHouse

- Using **Materialized View**, which acts as an insert trigger, we can compute the statistics during INSERT time
- This means that we shift the computation from SELECT time to INSERT time

1

```
CREATE TABLE weather (  
    `city` String,  
    `temperature` Decimal(8, 1)  
)  
ENGINE = Null;  
  
CREATE TABLE weather_results (  
    city String,  
    max AggregateFunction(max, Decimal(8, 1)),  
    min AggregateFunction(min, Decimal(8, 1)),  
    avg AggregateFunction(avg, Decimal(8, 1))  
)  
ENGINE = AggregatingMergeTree  
ORDER BY tuple();
```

2

```
CREATE MATERIALIZED VIEW weather_mv TO weather_results  
AS SELECT  
    city,  
    maxState(temperature) as max,  
    minState(temperature) as min,  
    avgState(temperature) as avg  
FROM weather GROUP BY city;
```

3

```
INSERT INTO weather  
SELECT city, temperature  
FROM (  
    SELECT  
        splitByChar(':', line) AS vals,  
        vals[1] AS city,  
        CAST(vals[2], 'Decimal(8, 1)') AS temperature  
    FROM file('measurements.csv', LineAsString));
```

0 rows in set. Elapsed: 24.398 sec. Processed 2.00 billion rows, 35.64 GB (81.97 million rows/s., 1.46 GB/s.)  
Peak memory usage: 242.30 MiB.



# ClickHouse and The One Billion Row Challenge

The challenge, optimizing ClickHouse

- How fast would it be?

```
SELECT
    format('{}={}/{}/{}', city, minMerge(min), round(avgMerge(avg), 2), maxMerge(max))
FROM weather_results
GROUP BY city
ORDER BY city ASC
FORMAT CustomSeparated
SETTINGS
    format_custom_result_before_delimiter = '{',
    format_custom_result_after_delimiter = '}',
    format_custom_row_between_delimiter = ', ',
    format_custom_row_after_delimiter = '',
    format_csv_delimiter = ';;';
```





# ClickHouse and The One Billion Row Challenge

The challenge, optimizing ClickHouse

- Using **Materialized View**, query duration lowered from **23s** to **0.014s (~1600 times** speedup)

```
SELECT
    format('{}={}/{}/{}', city, minMerge(min), round(avgMerge(avg), 2), maxMerge(max))
FROM weather_results
GROUP BY city
ORDER BY city ASC
FORMAT CustomSeparated
SETTINGS
    format_custom_result_before_delimiter = '{',
    format_custom_result_after_delimiter = '}',
    format_custom_row_between_delimiter = ', ',
    format_custom_row_after_delimiter = '',
    format_csv_delimiter = ';';
```

412 rows in set. Elapsed: 0.014 sec.



# Thank you!

Keep in touch!



[clickhouse.com/slack](https://clickhouse.com/slack)



[#clickhouseDB](#) [@clickhouseinc](#)



[clickhouse](#)