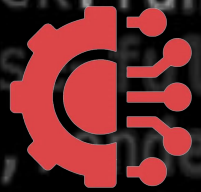

data + ClickHouse

Singapore Meetup

July 11, 2024

Thank you to our host!

INFOCOMM
MEDIA
DEVELOPMENT
AUTHORITY

Tech Talks

- **The State of SQL-based Observability**
  Pradeep Chhetri, Site Reliability Engineer @ ClickHouse
- **ClickHouse: Powering Coinhall's Real-Time Blockchain Data Platform**
  Aaron Choo, Co-Founder & CTO @ Coinhall
- **Panel Q&A**

The State of SQL-based Observability

July 11, 2024

ClickHouse

Pradeep Chhetri

Site Reliability Engineer @ ClickHouse

pradeep-chhetri-61a80935

p_chhetri

chhetripradeep

- I love playing with computers, trying out new softwares and databases.
- In my free time, i enjoy watching chess and football.

Table of contents

ClickHouse

Introduction

SQL

3rd most popular programming language

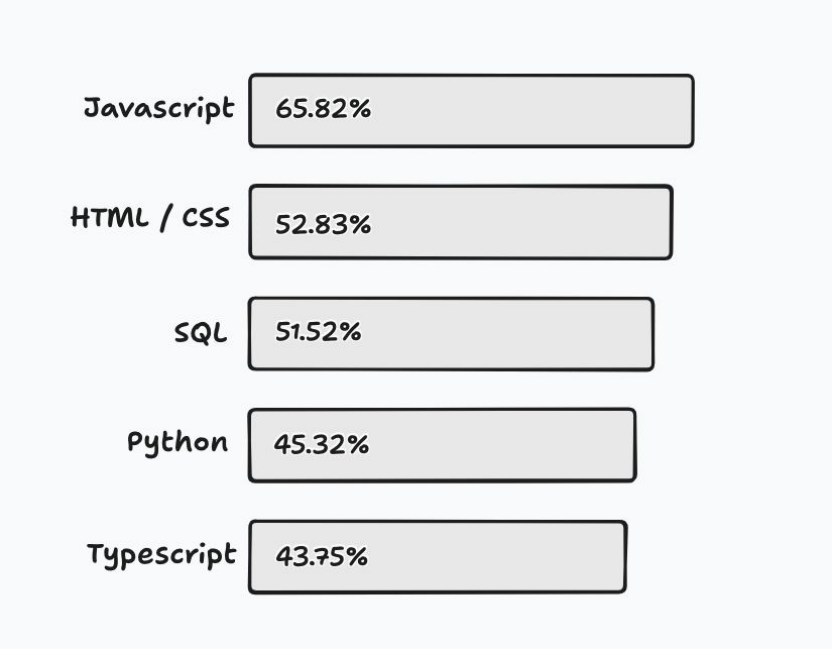

Observability

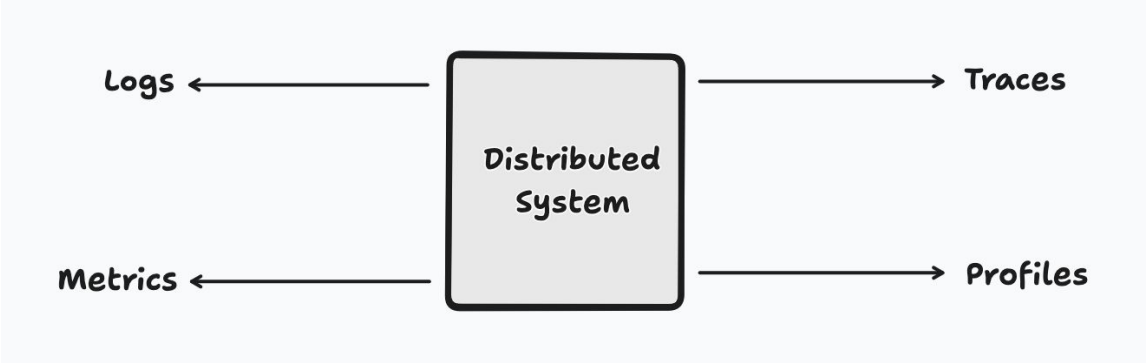

Evolution of SQL Databases

Evolution of Observability

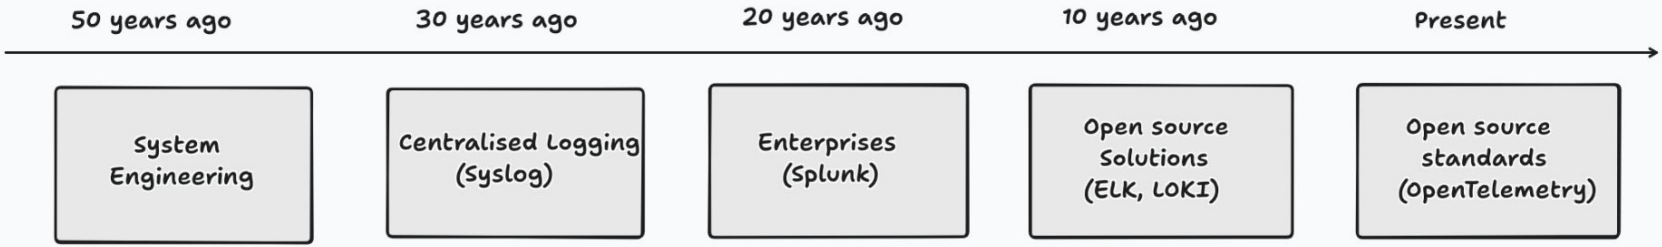

Overlap of SQL Databases & Observability

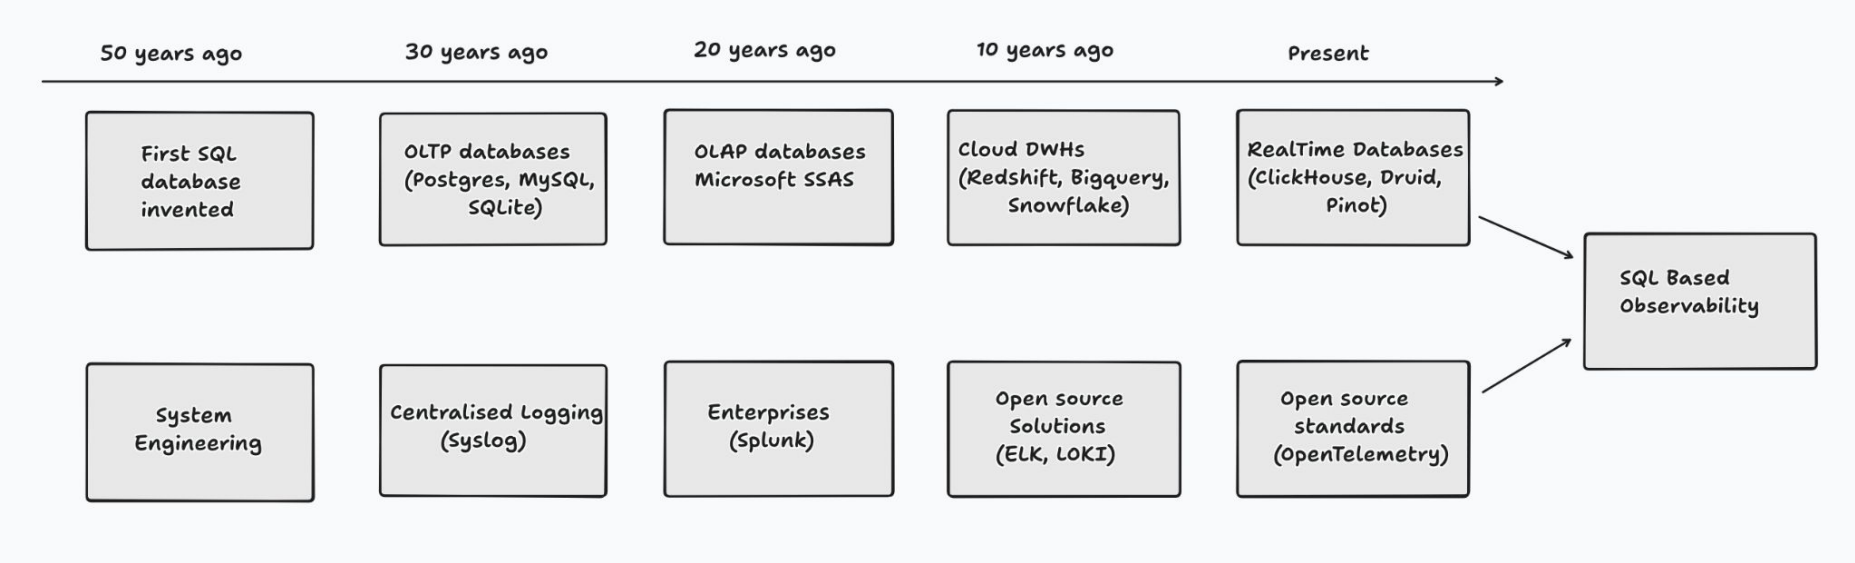

ClickHouse

ClickHouse

Challenges

Observability Pipeline

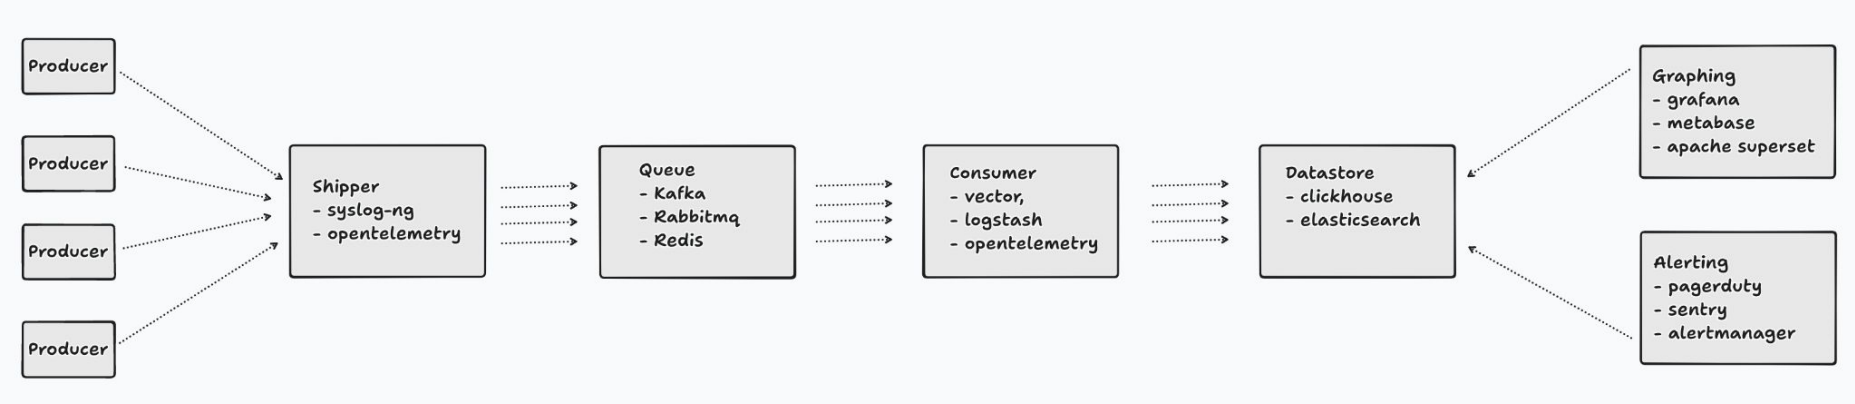

ClickHouse

Challenges with Observability

High Cardinality

Customer ID	Customer Plan
3A9D9780-0E89-4F3E-B299-459121D12ACC	startup
8FDAFF2F-6EBF-4C07-B1A9-0D893D868B11	enterprise
A61F1D6D-787C-434A-8B5C-69EF0D29A9FA	startup
9884F532-8F49-432C-9D4C-A815ACB6A0F2	enterprise
1F9566B5-E2BA-467F-8F18-EC0EFCE6E39C	startup

Challenges with Observability

Heavy Resource Utilization

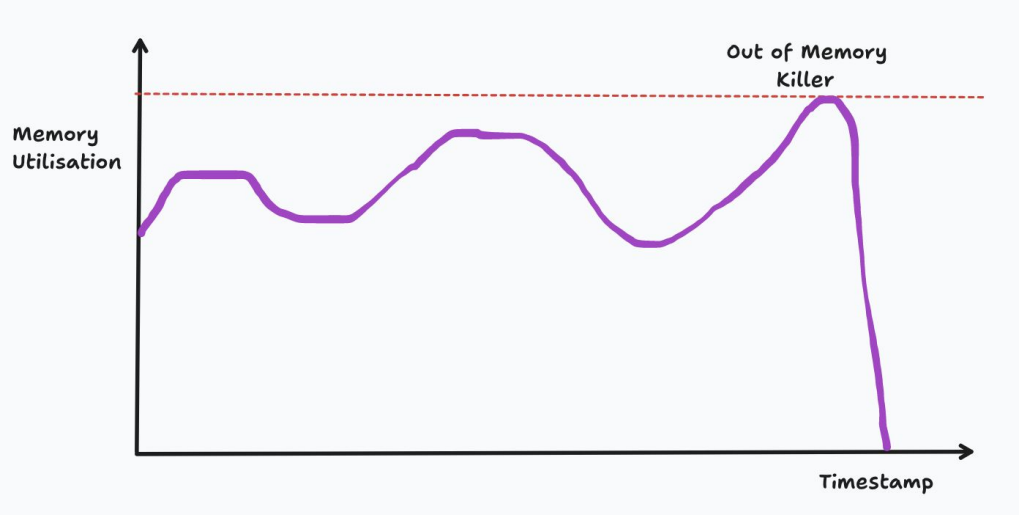

ClickHouse

Challenges with Observability

Multi Tenancy Issues

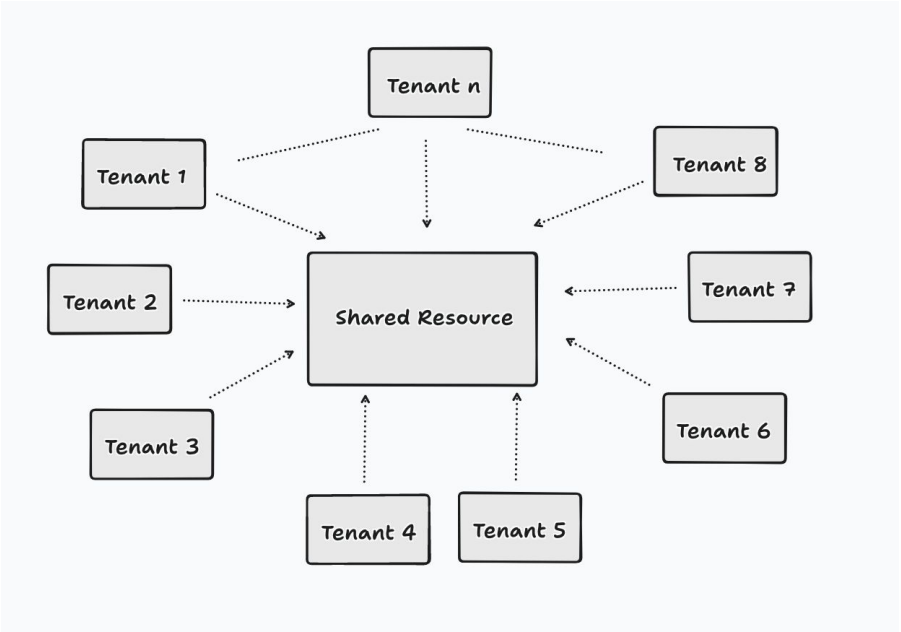

ClickHouse

Challenges with Observability

Resources Cost

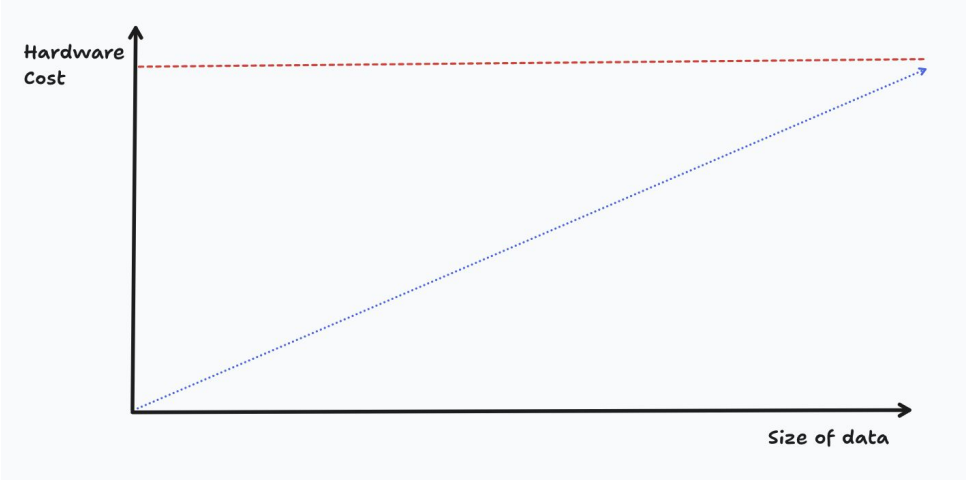

Solutions for Observability Challenges

Infinite Cardinality Support

Optimized Resource Utilization

Compute & Storage Separation

Support for Quota, Priority, Resource Management

Efficient Data Compression

Scale easily from 1 byte to 1000 petabytes

“Observability is just another data problem.”

What is ClickHouse?

open-source

- Developed since 2009
- Open sourced in 2016
- 35k+ Github stars
- 1k+ contributors
- 300+ releases

column-oriented

- Best for aggregations
- Files per column
- Sorting and indexing
- Background merges

distributed

- Replication
- Sharding
- Multi-master

OLAP database

- Analytics use cases
- Aggregations
- Visualizations
- Mostly immutable data

ClickHouse

Real-world Deployments

HTTP & DNS Analytics Platform

Architecture

Shippers → Kafka → ClickHouse

Wins

Efficient Ingestion & Compression

Improved Throughput & Latency

https://blog.cloudflare.com/http-analytics-for-6m-requests-per-second-using-clickhouse

ClickHouse

Log Analytics Platform

Architecture

Log shippers → Kafka → ClickHouse → Kibana

QueryBridge to convert lucene to sql queries

Wins

Speed of ingestion, cost control

Tradeoffs

Stack administration, UI development

https://www.uber.com/blog/logging/

https://presentations.clickhouse.com/meetup40/uber.pdf

Uber

ClickHouse

Log Analytics Platform

```
CREATE TABLE <table_name>
(
      // Common metadata fields.
      _namespace               String,
      _timestamp                Int64,
      hostname                 String,
      zone                     String,
      ...

      // Raw log event.
      _source                  String,

      // Type-specific field names and field values.
      string.names             Array(String),
      string.values            Array(String),
      number.names             Array(String),
      number.values            Array(Float64),
      bool.names               Array(String),
      bool.values              Array(UInt8),

      // Materialized fields
      bar.String               String,
      foo.Number               Float64,
      ...
)
...
```

Uber

ClickHouse

Adopting OLAP store for tracing

Architecture

OpenTelemetry → ClickHouse → Kibana

Wins

Data compression, open source licensing

Tradeoffs

Managing tiered OTel collectors

https://kccnceu2024.sched.com/event/1YeNu

ClickHouse

Adopting OLAP store for tracing

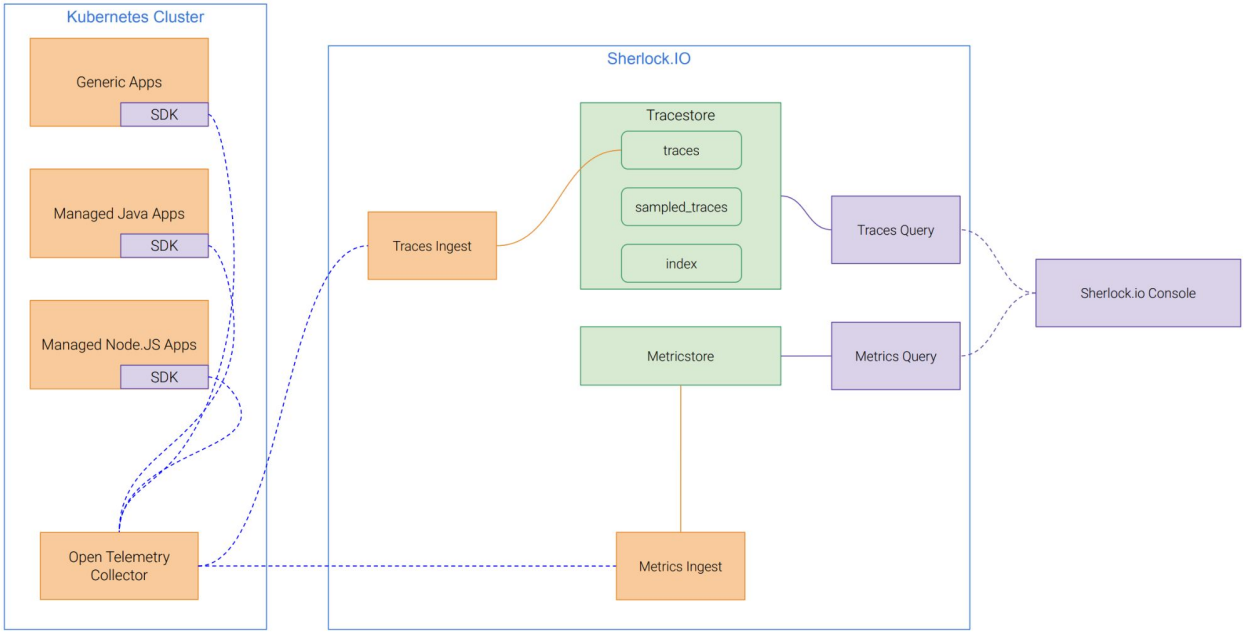

Dogfooding ClickHouse across o11y

Architecture

OpenTelemetry → ClickHouse → Grafana

Wins

Granular log retention, Saved money on Datadog

Tradeoffs

1.5 FTEs to build stack

https://clickhouse.com/blog/building-a-logging-platform-with-clickhouse-and-saving-millions-over-datadog

Dogfooding ClickHouse across o11y

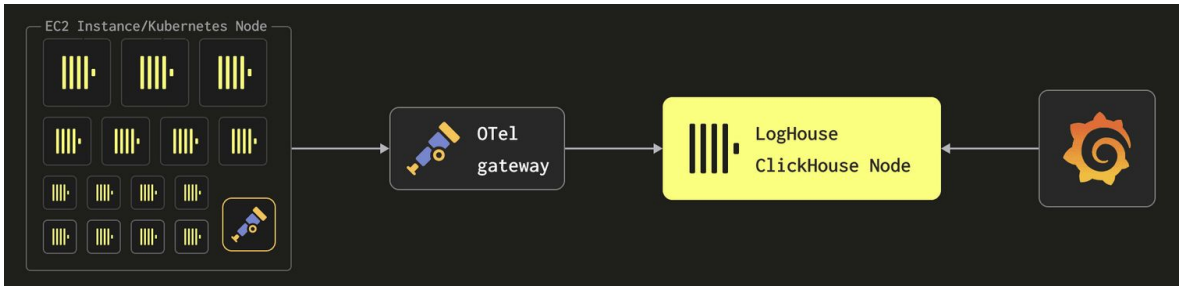

Uncompressed Size	Compressed Size	Ratio	Rows
19.11 PiB	1.13 PiB	17.0	36.96 trillion

ClickHouse

ClickHouse

ClickHouse

Common considerations

Query language considerations

“SQL is not compact enough compared to domain-specific query languages”

A simple query ?

```
source=events level=”warning”
| STATS avg(duration) BY level
| FIELDS level, avg(duration) AS avg_dur
| sort - avg_dur | head 10
```

```
GET events/_search
{
  "size": 0,
  "_source": false,
  "track_total_hits": -1,
  "aggregations": {
    "groupby": {
      "composite": {
        "size": 10,
        "sources": [
          {
            "4e8796da": {
              "terms": {
                "field": "level.keyword",
                "missing_bucket": true,
                "order": "asc"
              }
            }
          }
        ]
      },
      "aggregations": {
        "c3318afb": {
          "avg": {
            "field": "duration"
          }
        }
      }
    }
  }
}
```

In good old SQL

```
SELECT
    level,
    avg(duration) AS dur
FROM events
GROUP BY level
ORDER BY dur DESC
```

Schema considerations

It helps to stop thinking about *“metrics”, “logs” and “traces”* separately and just think them as *“wide events”*

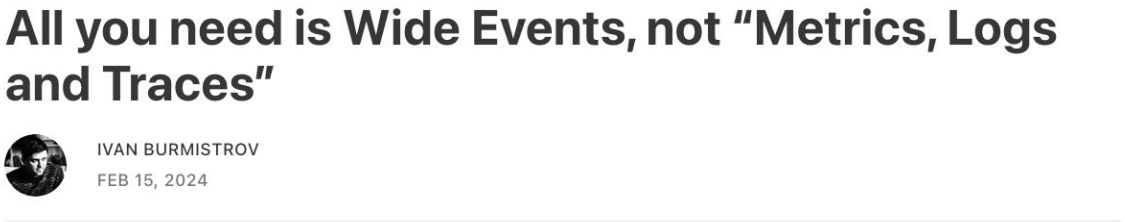

https://isburmistrov.substack.com/p/all-you-need-is-wide-events-not-metrics

https://news.ycombinator.com/item?id=39529775

Visualization considerations

Grafana

Apache Superset

Perses

Metabase

Build your own

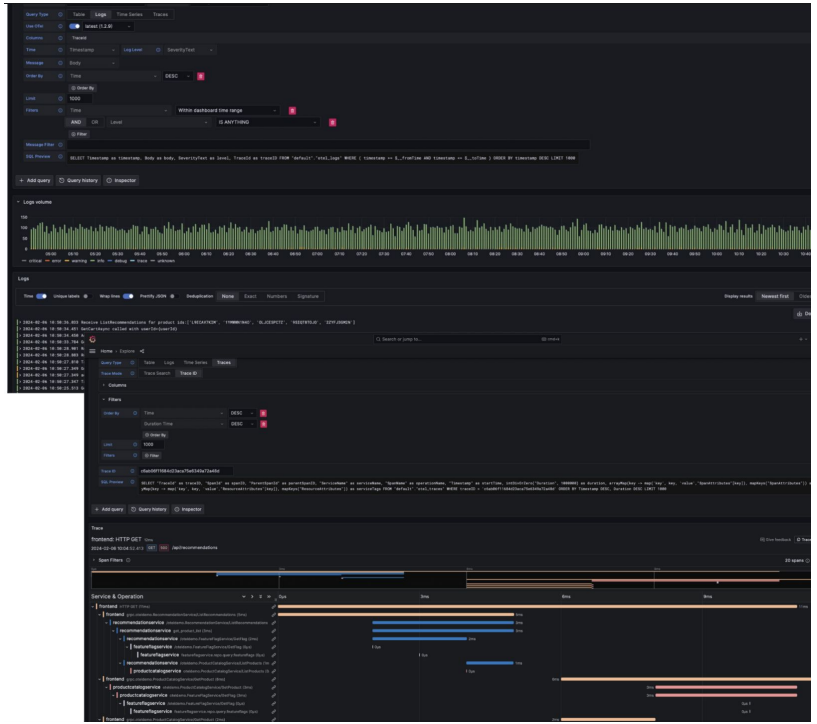

Multi-tenancy considerations

Example: Uber

Consider datastore ability to limit resources by table, user, session

Unified Multi-Tenant Storage Platform

- ClickHouse natively supports zero lock contention among concurrent reads and writes
- Service placement: single-tenant vs multi-tenant
  - Isolate heavy log producers, heavy log consumers
  - Co-locate everything else
  - Limit the impact of co-location, add service in order-by
- Workload isolation
  - Configure query parallelism per query
  - Eventually limit total query resource usage per node
  - Query cost accounting, defense against expensive queries

Multi-region deployment considerations

Example: ClickHouse Inc.

Per region data collection, cross-region queries

Resilient to AZ outage but not region outage

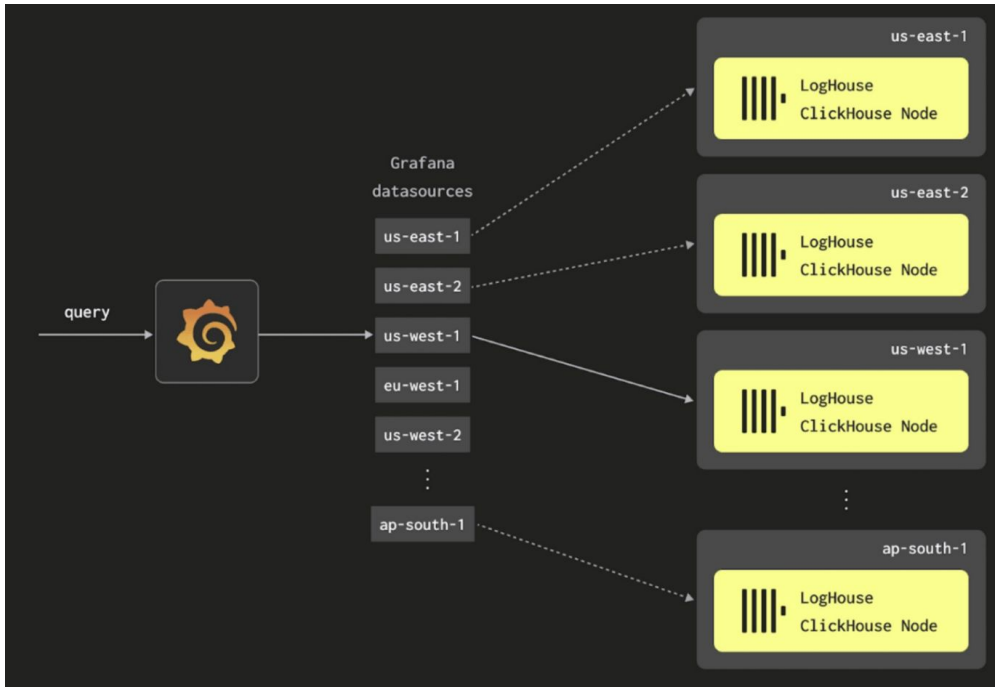

ClickHouse

Choice of analytical datastore matters

	ClickHouse	Druid	Pinot	BigQuery
Real-time speed	✓ Best	✓ Ok	✓ Ok	✗ Poor
Compression	✓ Best	✓ Ok	✓ Ok	✓ Better
Sep storage & compute	✓	✗	✗	✓
Interoperability - OTel - Grafana	✓ ✓	✗ ✓ (no logging & tracing support)	✗ ✗	✗ ✓ (no logging & tracing support)
SQL compliance	✓ Good	✗ Poor	✗ Poor	✓ Best
TCO improvements	✓ 10x	1-2x	1-2x	depends on how often you query

https://benchmark.clickhouse.com

ClickHouse

Demo

ClickHouse

ClickHouse Timeseries Engine

ClickHouse Timeseries Engine

Pull Request: https://github.com/ClickHouse/ClickHouse/pull/64183

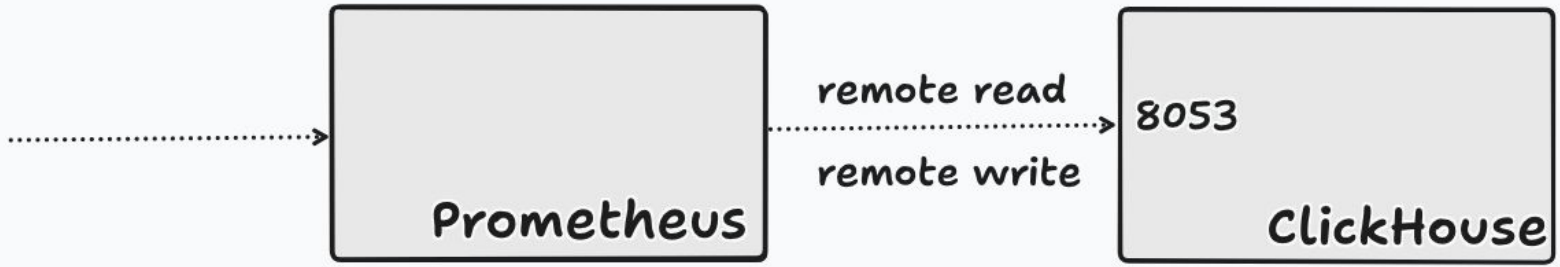

ClickHouse

Questions