

# DELIVERING ANALYTICS AT SCALE

---

Our Journey to Ingest and Analyse 5 Billion Events / Day

Sharat Madanapalli  
Head of Data and AI

# ABOUT



## Me

Started as a founding engineer @ Canopus  
(along with my PhD @ UNSW)

Worked on core product, data pipeline &  
produced some cool patents (applied AI)

Lead the Data and AI team:  
analysts, scientists, MLEs and SDEs.



\*Recent picture but not comprehensive

## Canopus Tech Team

Network Analytics startup born out of UNSW

Bunch of amazing tech folks who solve challenges!

Built an AI-powered network analytics platform  
and deployed around the world!

# NETWORK MONITORING FOOTPRINT

**250+**  
**Gbps**

Monitored  
Traffic Rates

**2M+**

Daily Active  
Users (DAU)

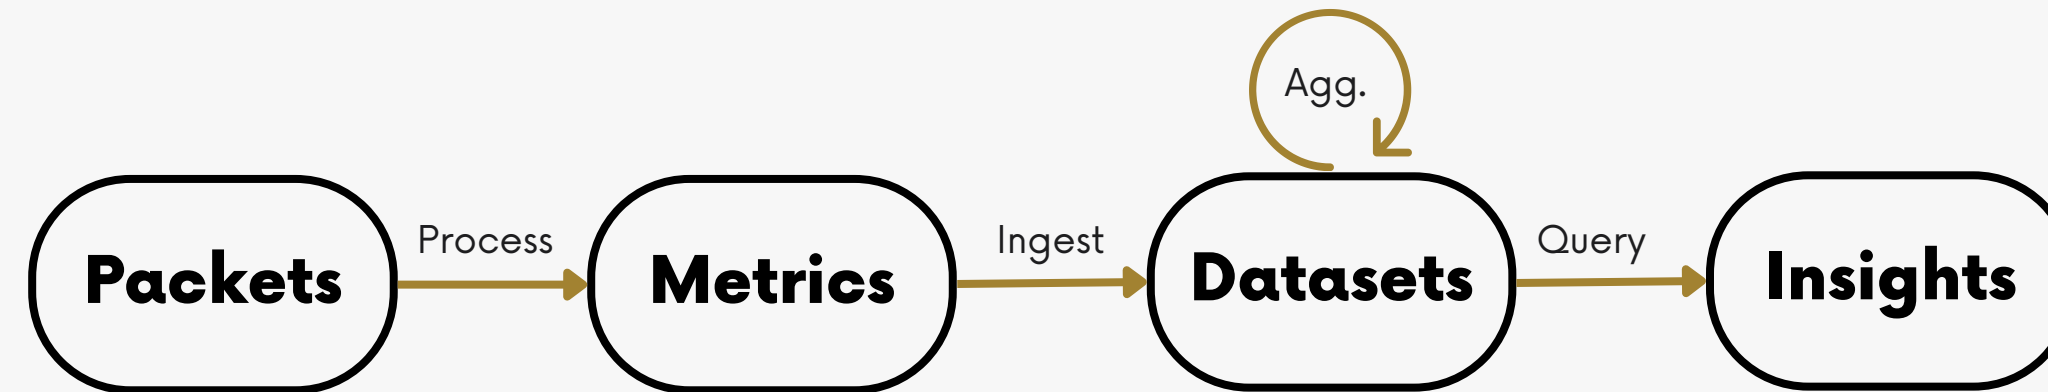
**5B+**

Daily Events  
Ingested &  
Analysed





# NETWORK ANALYTICS



## Packet Engine: Probe

Parses and extracts packet info

Tracks connections and hosts active on the network

Classifies applications using signatures and ML models

Measures performance in terms of RTT, Loss, Throughput etc.

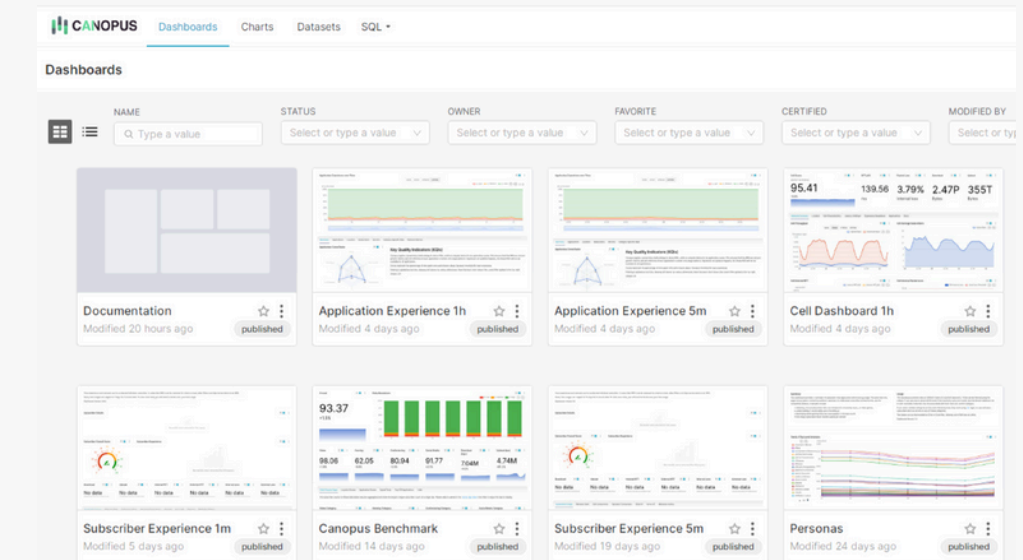
## Data pipeline (this talk!)

**Ingest:** Take the raw metric stream and dump into DB.

- Avro Serialization
- Kafka
- Write it to DB

**Aggregate:** Transform raw data into “analysis-ready” datasets

- Roll ups in time
- Group by
- Join with customer metadata



## Analysis and Insights

Address customer’s analytics needs

Query the data and show it on dashboards and GUI

- Query caching
- Creating appropriate “views”
- Modelling: Anomaly, Drilldown
- User permissions

# AN EXAMPLE RECORD

```
{  
  "timestamp": "2024-09-05T10:15:00Z",  
  "user_id": "user12345",  
  "client_ip": "192.168.1.5",  
  "server_ip": "203.0.113.10",  
  "download_bytes": 1256789,  
  "upload_bytes": 456789,  
  "internal_rtt": 15.6,  
  "external_rtt": 75.2,  
  "internal_loss": 0.01,  
  "external_loss": 0.05  
}
```



Timestamp	User ID	Client IP	Server IP	Download Bytes	Upload Bytes	Internal RTT (ms)	External RTT (ms)	Internal Loss (%)	External Loss (%)
2024-09-05T10:15:30Z	user12345	192.168.1.5	203.0.113.10	1,256,789	456,789	15.6	75.2	0.01	0.05
2024-09-05T10:20:45Z	user67890	192.168.1.12	203.0.113.11	3,456,123	789,123	22.3	80.5	0.03	0.07
2024-09-05T10:25:10Z	user54321	192.168.1.7	203.0.113.12	987,654	123,456	18.9	70.1	0.02	0.04

Metric


Dataset

# OUR JOURNEY TO SCALE

---

Focusing on Ingest, Aggregate and Analyse

# HUMBLE BEGINNINGS: CSV



## Ingest

Probe writes metrics to a CSV file.

Rotates the files (one per schema) every 24 hours.

## Aggregate

No aggregates whatsoever

## Analyse

Jupyter Notebooks + Pandas + Plotly


Wrote scripts that produced daily reports over our data that we mailed to our customers.

Good starting point: Getting insights to customers quickly. Pandas on CSV still counts as data science!

Writes aren't fast enough. Uncompressed so storage won't scale.

Querying requires loading all the files in memory.

# LET'S GET A SQL DB: TIMESCALE



## Ingest

Probe writes metrics to the DB.  
Each file becomes a new table.

## Aggregate

Probe pre-aggregated the data in memory and pushed down to DB.

## Analyse

Can now generate the same reports in an efficient way.  
Started to develop our full stack apps backed by Timescale.


Liked the convenience of SQL, multiple modules could query the data :)

Timescale recent "chunk" in RAM: Real-time dashboards  but historical dashboards 

As we scaled, Timescale was taking up half the RAM on our servers -> query timeouts!



# DB TO PIPELINE: KAFKA + CLICKHOUSE



## Ingest

Probe writes metrics to Kafka topics.

Use avro serialization

Clickhouse Kafka Engine

## Aggregate

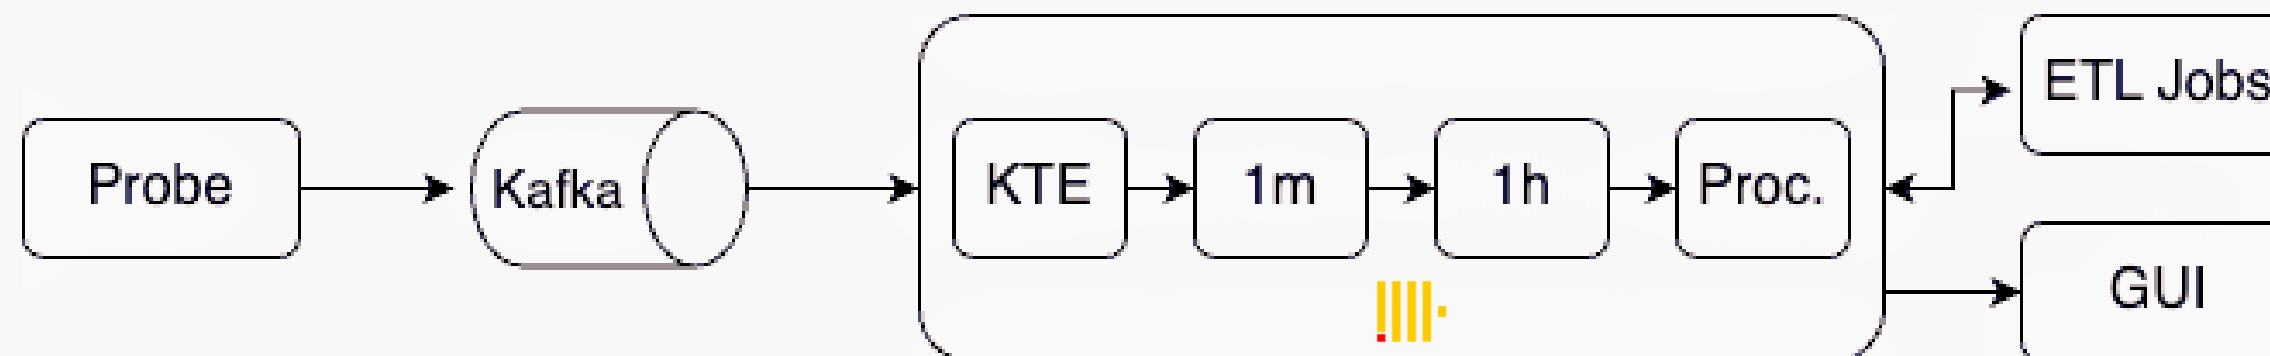
Clickhouse Materialized Views

- Materialize from kafka table into a merge tree table
- Roll up into time buckets

## Analyse

Introduced scheduled ETL jobs to transform the data via Prefect.


GUIs got faster.



PURE MAGIC! Clickhouse was fast at low RAM and disk usage! Served us good until 10-100 millions of rows / day.

The customers always want more: Won't scale to billions of rows!

# 5B+ EVENT ANALYTICS PIPELINE



## Scale Ingest

Implemented Kafka-Go-Connect instead of kafka table engine:

- ch-go lib
- convert from avro -> native format batches
- async write mode.

Use SSDs as landing zone and TTL move to HDDs.

## Scale Aggregate

Real-time Aggregates using Flink

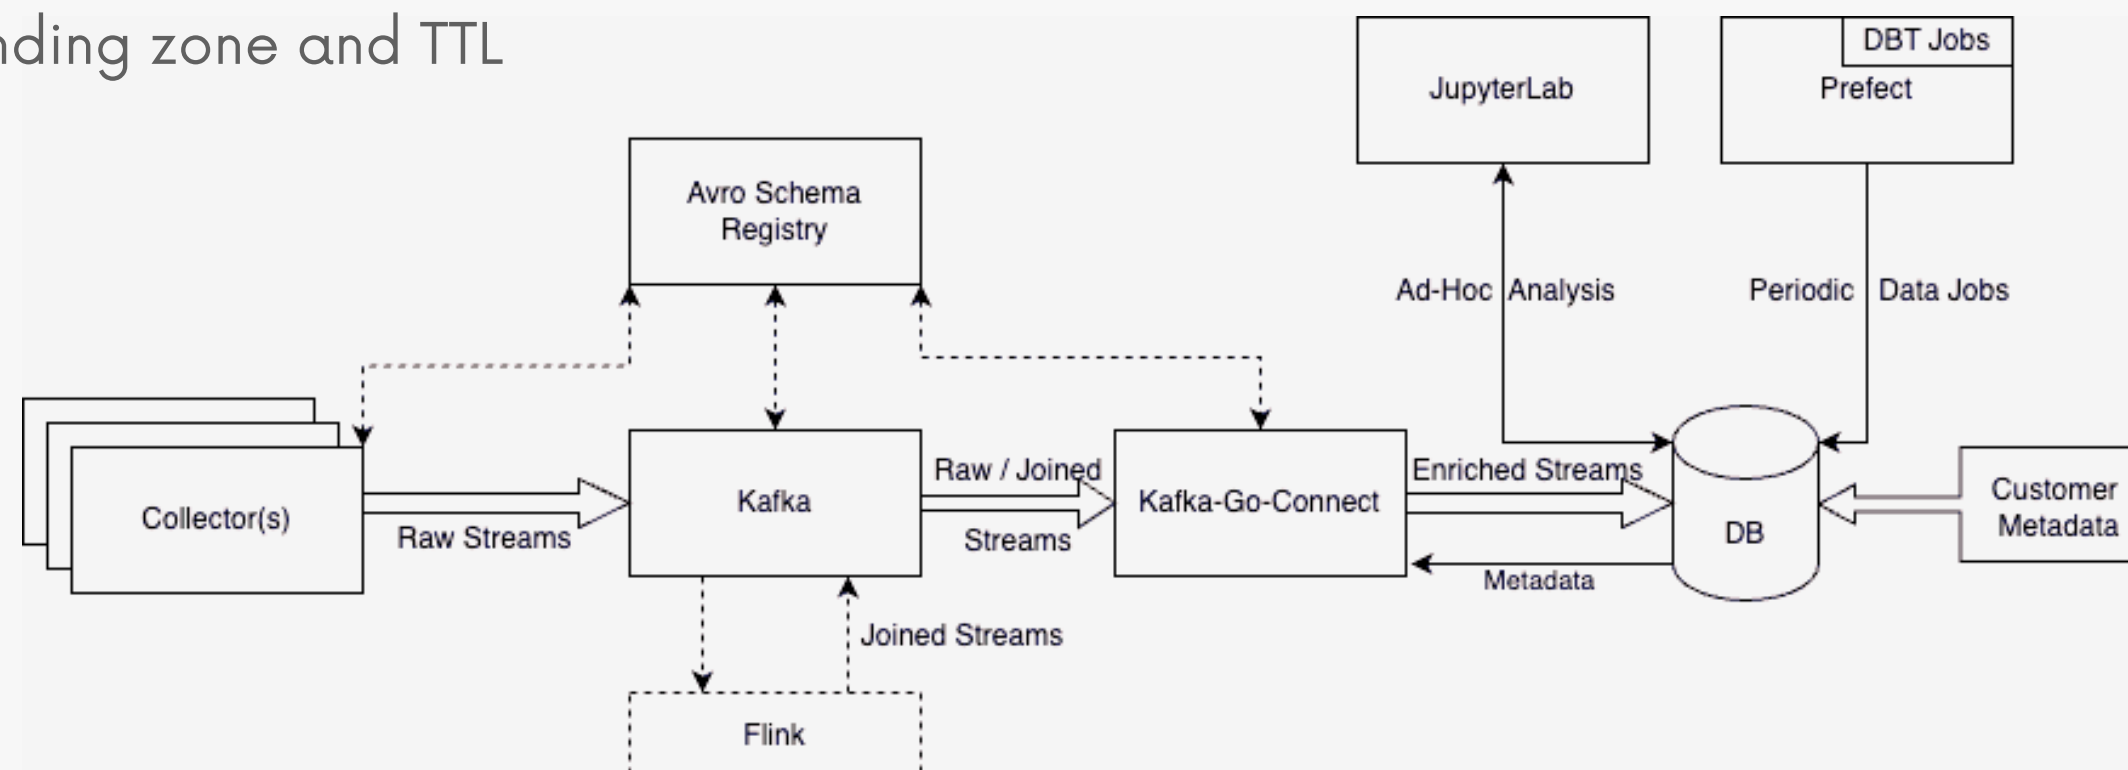
DBT to manage roll ups vs. MVs

Split up incoming data streams into two nodes. Agg independantly.

## Scale Analyse

Resource distribution: route query based on table between 2 servers.

Slow queries got converted into DBT models.



# LESSONS LEARNT

---

Tips and Takeaways

# MOVING FROM NATIVE MVs TO DBT

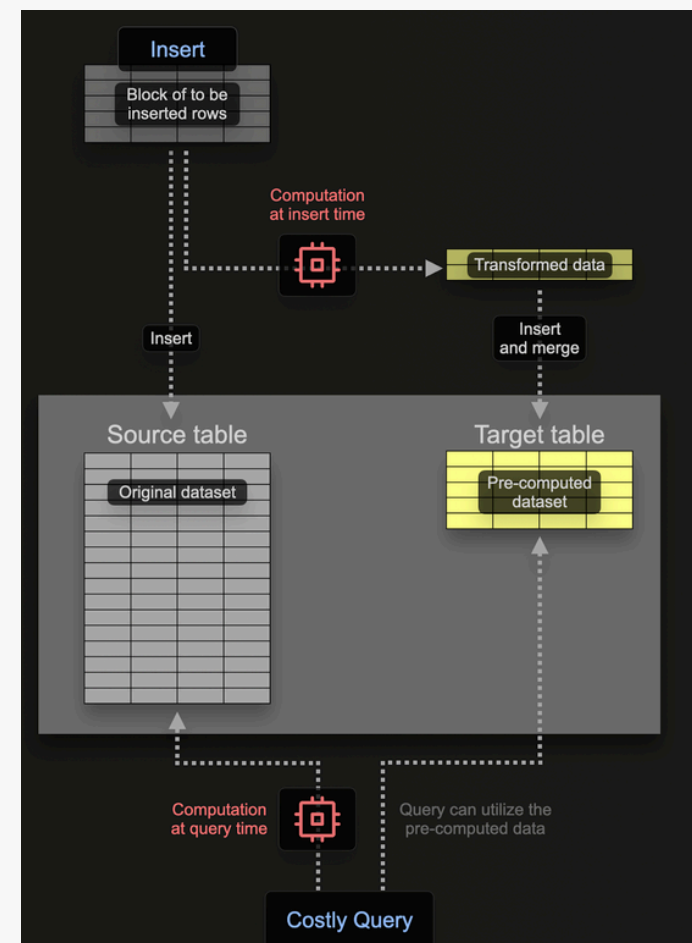
## Native MVs

Aggregate block of rows and store in a target table.

Good for Simple Aggregation (E.g. sum, count etc.) on a single stream via Plain SQL.


Complex Aggregation is a two stage materialization: state tables and "merge queries".

Joining two streams is challenging with MVs (there's no wait time)



Source: <https://clickhouse.com/docs/en/materialized-view>

## DBT



Source: <https://medium.com/@sagar.bhandge0310/dbt-data-build-tool-overview-67cc77e761ea>

General purpose data transformation tool via templated SQL

Write "Models" (basically SQL queries) => tables (via config)

Supports dependancies between models and incremental materialization only new rows added (for timeseries datasets)

For continuous data, `dbt run` has to be scheduled at regular intervals using schedulers like cron, prefect or airflow.

# ANALYZE AND MATERIALIZE LOOP

## Identify Slow Queries

Use Clickhouse Query Log table to observe repeated slow queries.


## Materialize

Use DBT to build an MV using the query. DBT will automatically handle table creation etc. (PK order given via config)

Schedule incremental updates to MVs using Prefect/Airflow

## Dashboard

Continue building more analytics for customers.





# MONITORING THE PIPELINE

“—

If you can't measure it, you can't manage it!

Peter Drucker

## Leading metrics: a potential future issue

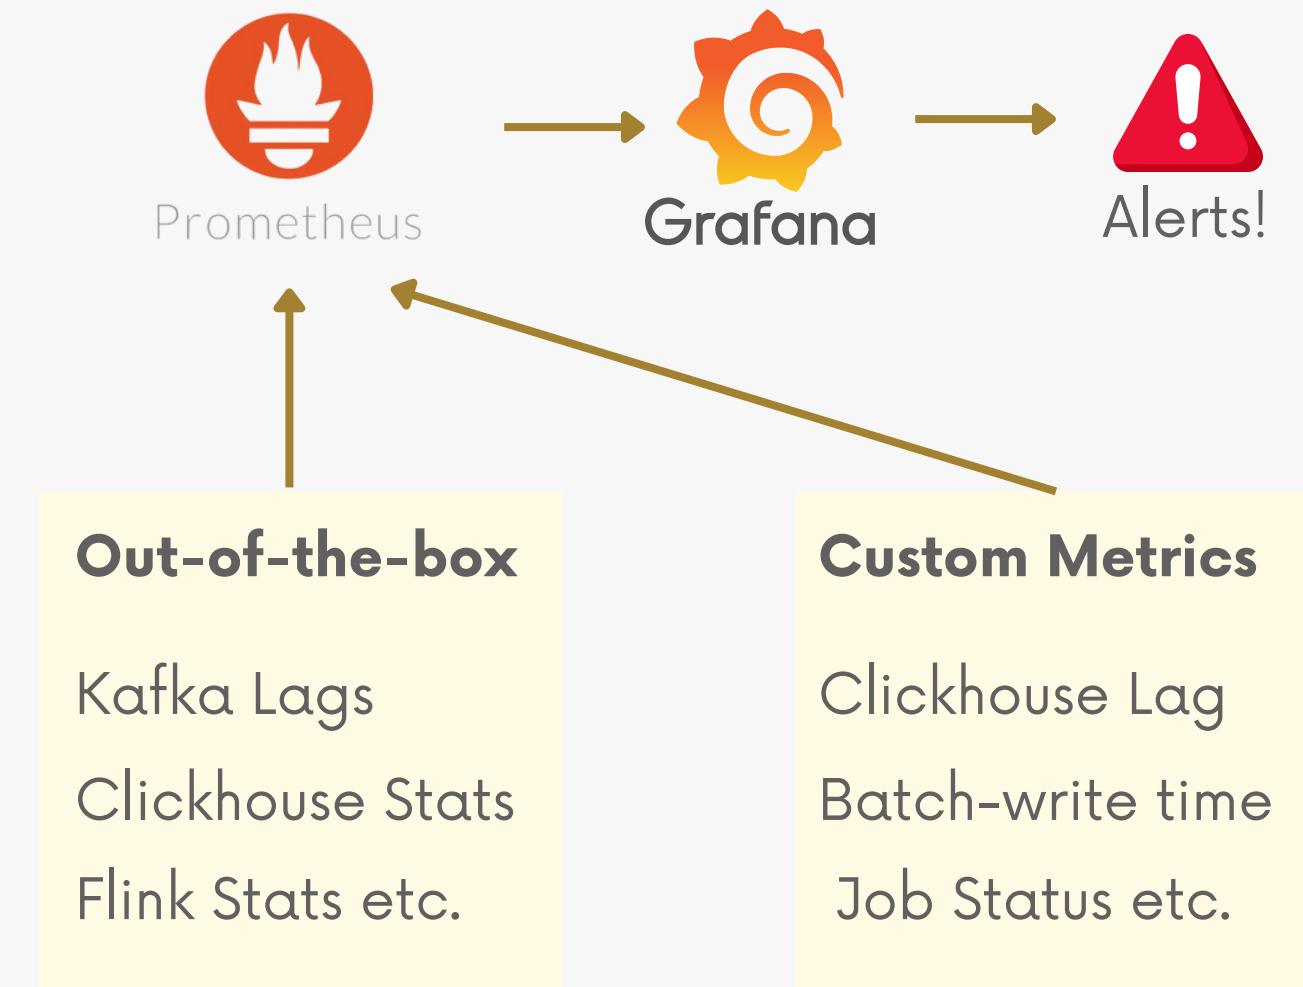
- Kafka lag increasing
- Clickhouse merges taking more time
- SSDs filling up (cascading effect)
- Batches taking more time to write (instrumented our kafka-connect)

## Alerting metrics: Indicate an immediate issue

- Data lag on the table
- Running out of memory, disk etc.
- ETL jobs failing

Clickhouse System Tables are an immensely useful source!

## Our Design





---

## My Top 3 takeaways

**Start small and scale up:** focus on delivering value through data first.

**Right Basics/Arch > Tools:** Get the tool/lib if you can't write the functionality yourself. E.g. ClickHouse :P

**Monitor is pre-req to scale:** Invest in measuring and monitoring the right metrics. Instrument via code if needed.

---

## Let's connect on LinkedIn!

