

Moving Druids to the house

ClickHouse Bengaluru Meetup on 8 Feb
2025




Speaker Info



- Platform engineer @ platformatory.io
- Confluent Community Catalyst, Kong Champion
- Occasional open source contributor to Cloud Native projects (k8s, ArgoCD, Tekton, Litmus, etc)
- Meetup organizer for Bengaluru Streams, Kong, Kafka, Grafana
- <https://www.linkedin.com/in/avinash-upadhyaya/>



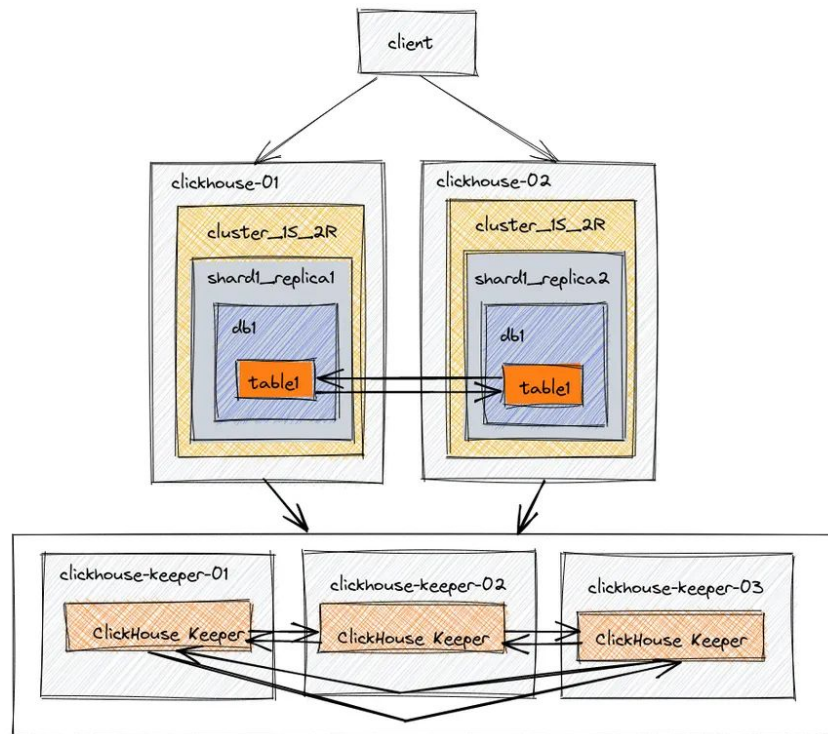



The Trifecta of Real-time OLAP



ONE DOES NOT JUST MIGRATE THEIR RTOLAP

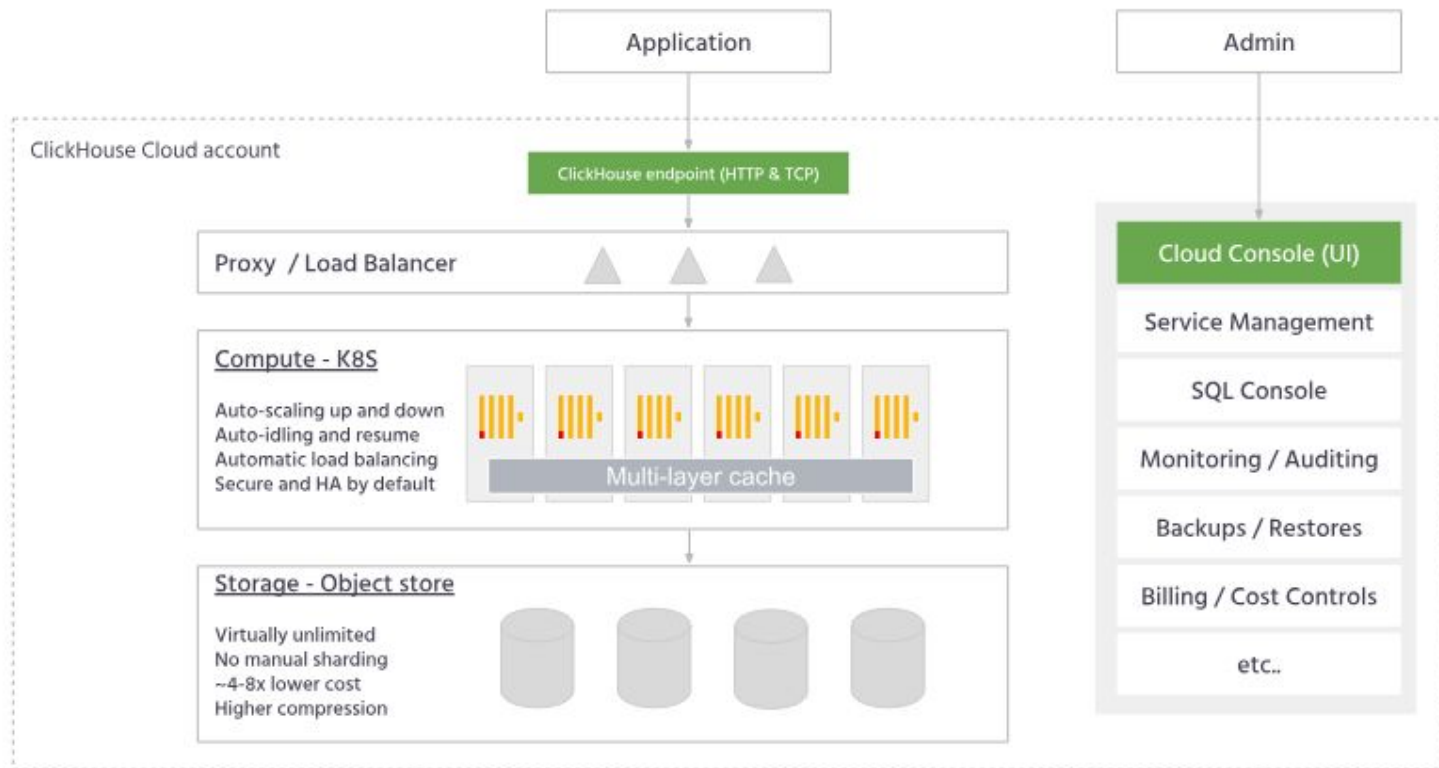




Apache Druid

ClickHouse





ClickHouse cloud is the best way to run ClickHouse.



Scoping your migration → API surface awareness

1. Querying
2. Ingestion
3. Management



Key gotchas

ClickHouse

- SQL over HTTP
- MySQL API
- Native binary API (with C++ client)
- gRPC (new)
- Table Engines are the most important detail
- Primary keys for ordering

Druid

- Druid SQL over HTTP
- JDBC (Avatica JDBC)
- Druid Native queries over HTTP
- Java APIs
- **Separate ingestion APIs → Task/submission API**
- No primary keys




Druid vs ClickHouse

Feature/Aspect	Apache Druid	Apache ClickHouse
Primary Query Interfaces	Dual: Native JSON query language and Druid SQL	Single unified SQL dialect with proprietary extensions
Query Expression	JSON for native queries; SQL for Druid SQL	Standard SQL (with additional functions for analytics, arrays, and distributed processing)
Specialized Functions	Time series-focused functions (e.g., TIME_FLOOR)	Array handling (<code>arrayJoin</code>), distributed query hints, MergeTree-specific options
Control and Flexibility	Native JSON offers detailed control over execution	SQL dialect optimized for columnar processing and high throughput
Ease of Use	SQL layer eases integration; native queries require more specific JSON structure	Unified SQL interface simplifies development and integration



Migration Overview

- Ingestion → Inventory of data sources, input formats, transformations and roll-ups
 - In ClickHouse, there's no separation between ingestion and query APIs
 - Carefully consider connectors, integrations and input formats
 - A table level mapping is crucial: choice of table family, schema parity and use of MVs
- Query → Inventory of user-facing applications, BI dashboards and clients at large
 - Druid SQL or JDBC offers scope for easier migration. Could save a lot of query rewriting time
 - Consider an API shim to perform incremental roll-out if using REST APIs
- Phased migration → across a single context of data sources and queries
- Develop a test and cut-over plan!
- Benchmark for tuning



Thank you

hello@platformatory.com

www.platformatory.io

