



# Supercharging Verihubs Data Warehousing with Clickhouse



# About Verihubs

Verihubs is a SaaS (software as-a-service) based company that provides a complete solution for compliance and security for companies' digital businesses in Indonesia. With Verihubs, the customer experience is smoother, more efficient operations, and well maintained data security.



Verihubs founded in 2019 on a mission to help companies reduce fraud risks through AI solution.



ISO/IEC:27001:2022 - Information Security Management System certified.



ISO/IEC:30107 - Presentation Attack Detection certified.



Recommended by OJK as ITSK (Inovasi Teknologi Sektor Keuangan) (S-42/D.07/2024).



Verihubs cutting edge computer vision is recognized by NIST for World-Class Technological Advancements in Computer Vision.

→ Awarded by NIST as #1 Face Recognition technology in Indonesia in terms of performance and accuracy rate and top 100 in the world.

→ Face Search technology is certified by NIST FRTE 1:N.

→ One and only NIST Presentation Attack Detection Certified from Indonesia (Liveness).



Processed over 50 million API/month.



Officially partnered with Meta as Business Solution Partner.



**400+ Clients** have seamlessly onboarded the right customers while preventing fraud - no matter their industry.

### Financial Institutions



### Digital App & Services



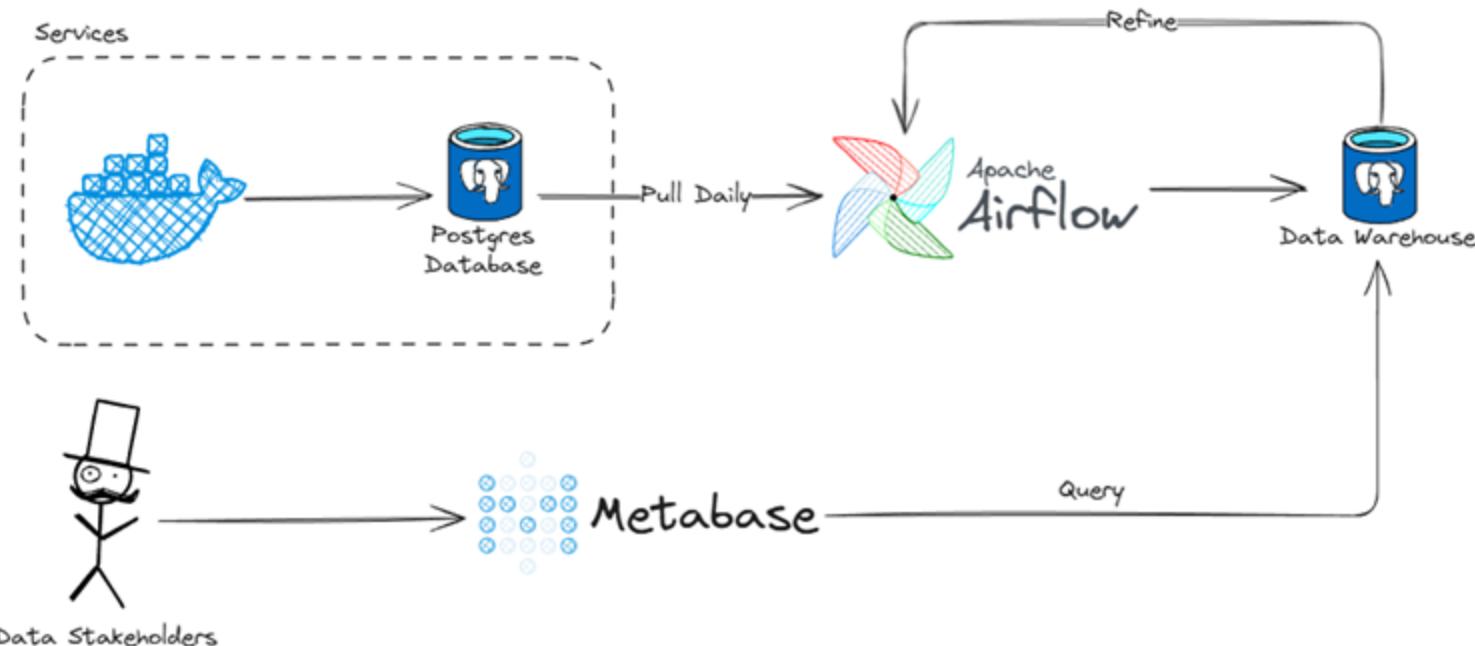
## Data Stored

- Non-sensitive API Transaction Logs (Traffic & Reconciliation)
- API Pricing Details, Client Company Details

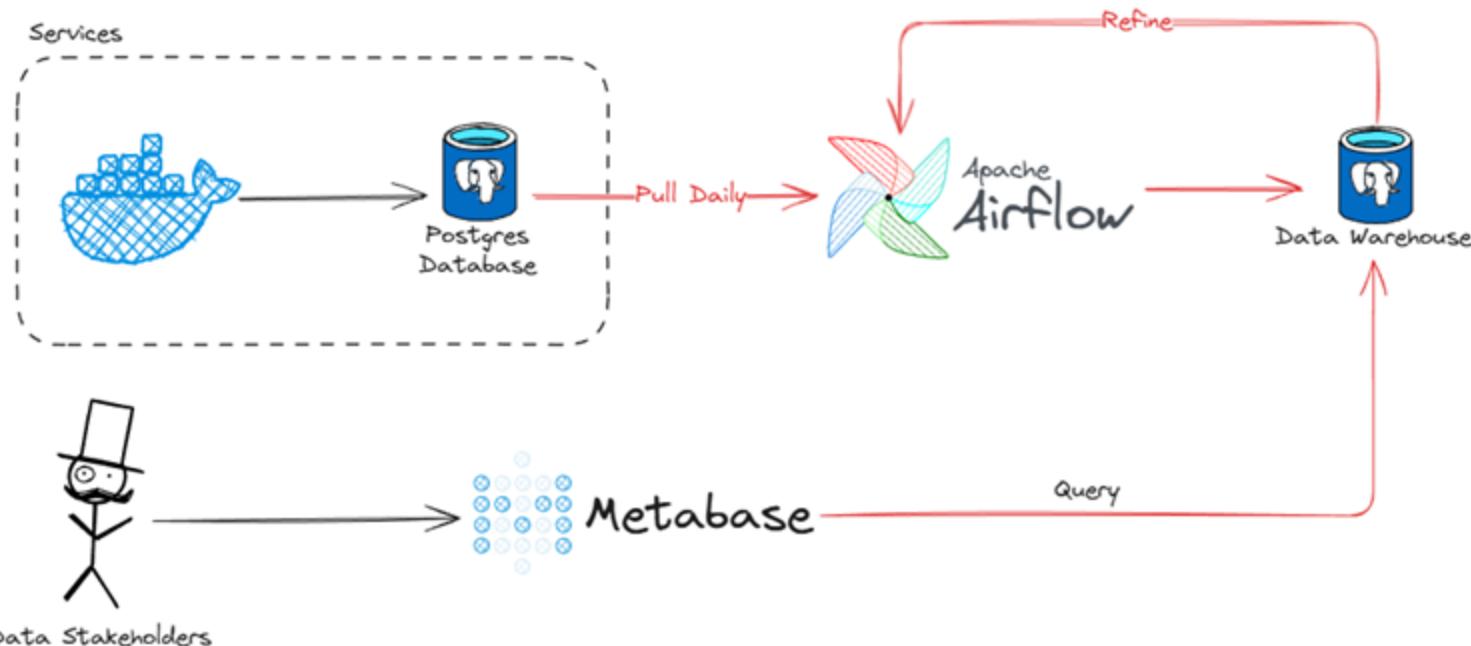
## Data Needed

- Traffic Logs Aggregated Daily and Monthly
- Revenue per API product
- Client Performance
- Invoicing for Clients
- Any ad-hoc business data requests

# Old Data Warehouse



# Bottlenecks



# PostgreSQL is fast up until some point

Usual daily data pull  
takes around **2-6 hours**  
per service

# Row table changes for the rows we already pulled for that day.

Created At	Updated At
November 30, 2025, 3:27 AM	December 6, 2025, 2:03 PM
November 20, 2025, 12:30 AM	December 7, 2025, 12:52 AM
November 9, 2025, 7:24 AM	December 5, 2025, 6:59 AM



# Why query take so long?

When can I get  
latest data for X?

Why is your data  
different?

# Data Stakeholder Complaints

User queries can take **minutes**  
and sometimes **timeout**

Data is **unusable** while backfill  
is being done

# Our Goals



Speed

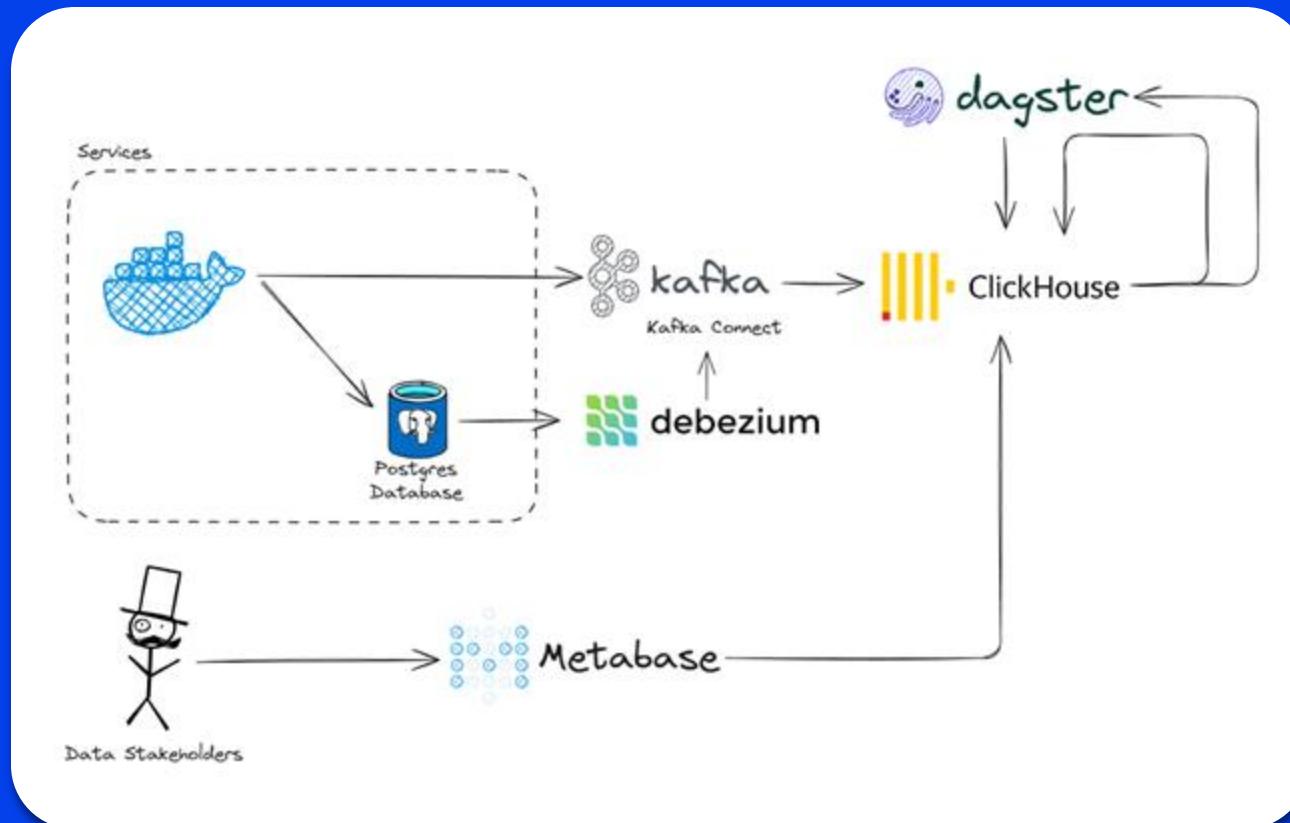


Cost

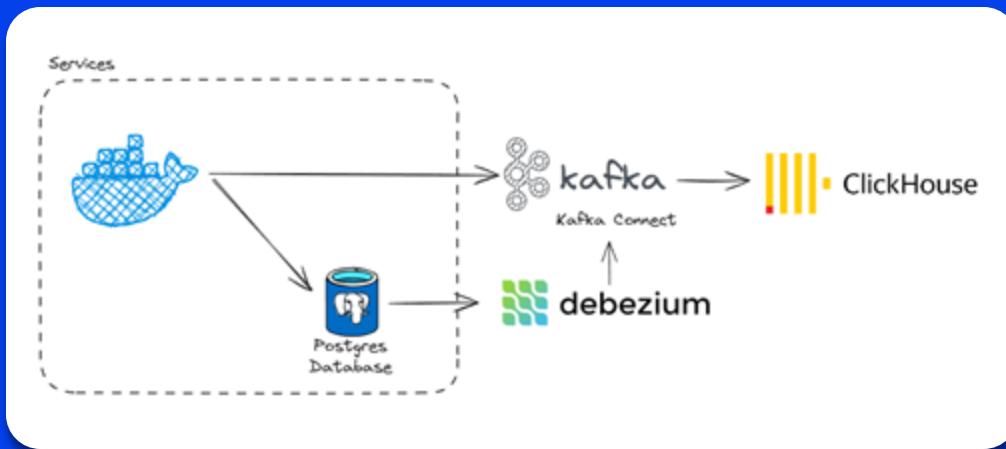


User Experience

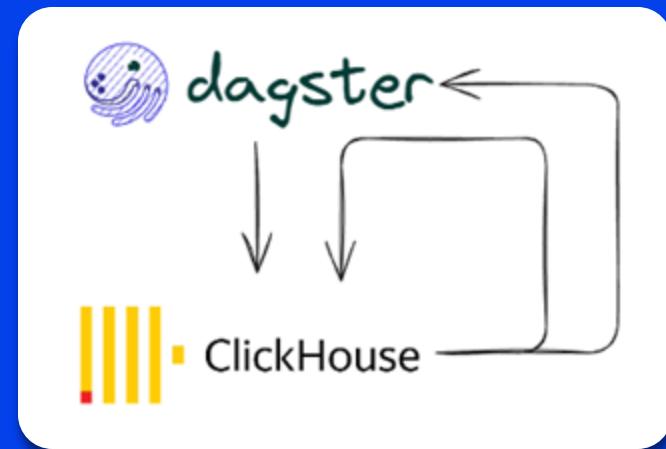
# New Data Warehouse



# Dissect

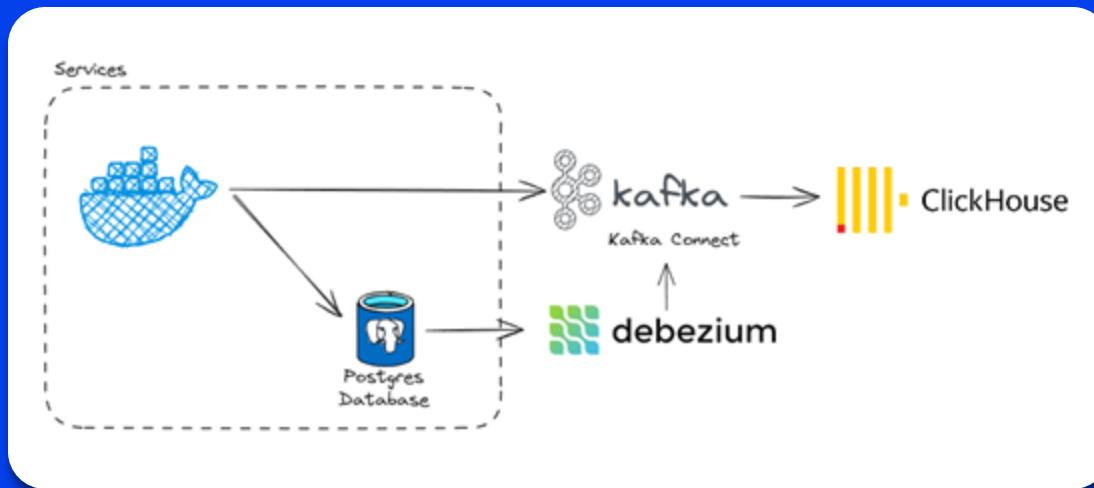


Data Ingestion



Data Post Processing

# Data Ingestion



- Data is now **realtime**
- Clickhouse **MergeTree** handles data storing optimization
- Clickhouse have a simple kafka integration via ClickPipes

# Data Ingestion



## Append Only Logs

Queued immediately  
into clickhouse.



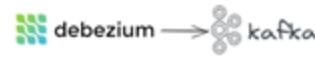
## Updatable Data

Captures every row  
upserts (CDC) and  
send the affected row.



## Append Only Logs

Queued immediately  
into clickhouse.



## Updatable Data

Captures every row  
upserts (CDC) and  
send the affected row.

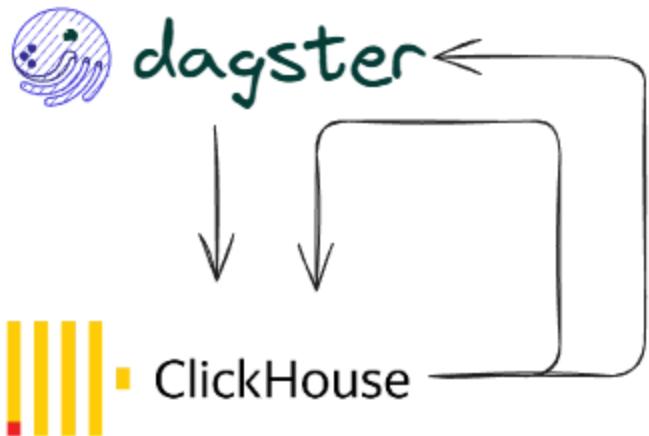


## MergeTree



## ReplacingMergeTree

# Data Post Processing



- Data is further refined in Clickhouse through **Materialized Views & Projections**
- Dagster is used for complex data processing (e.g. tracking suspicious behavior user, data mismatches between OLTP) and alert them to our slack



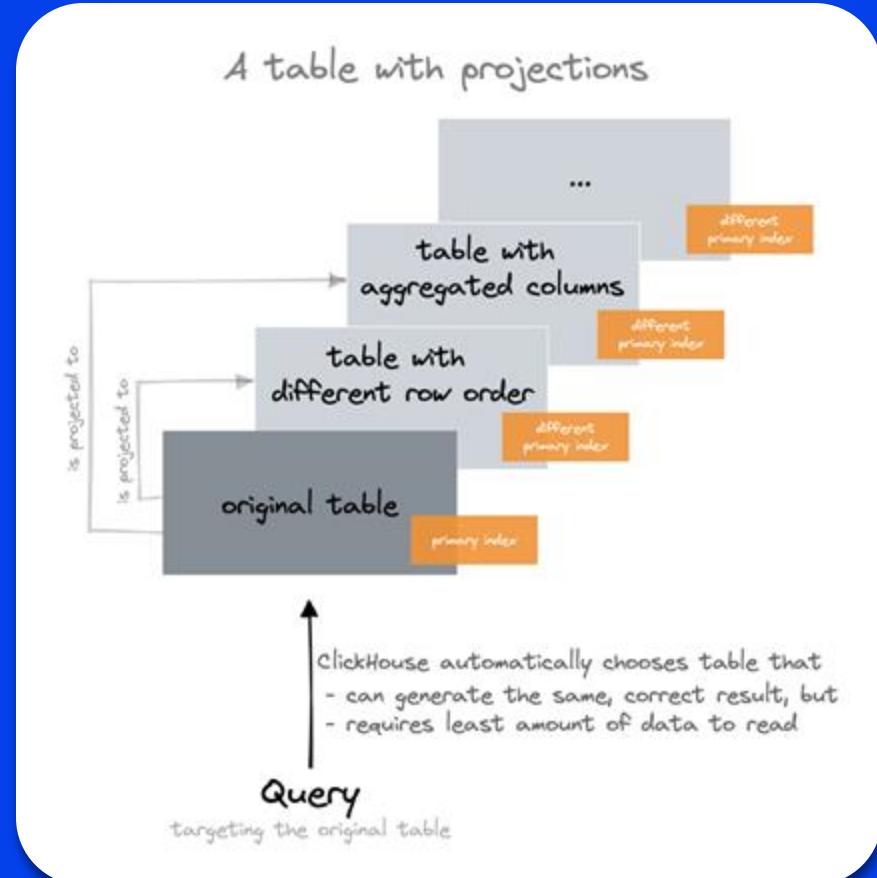
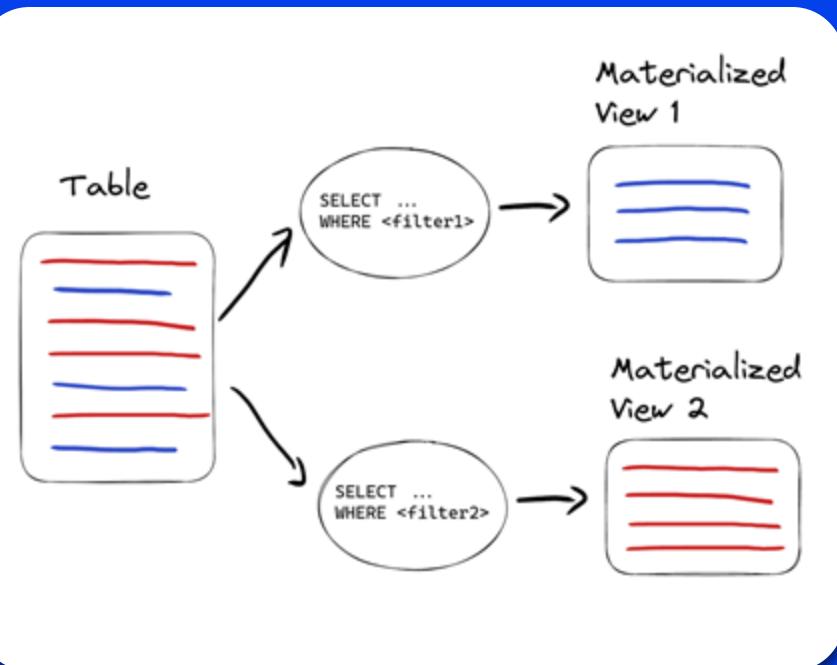
## Materialized View

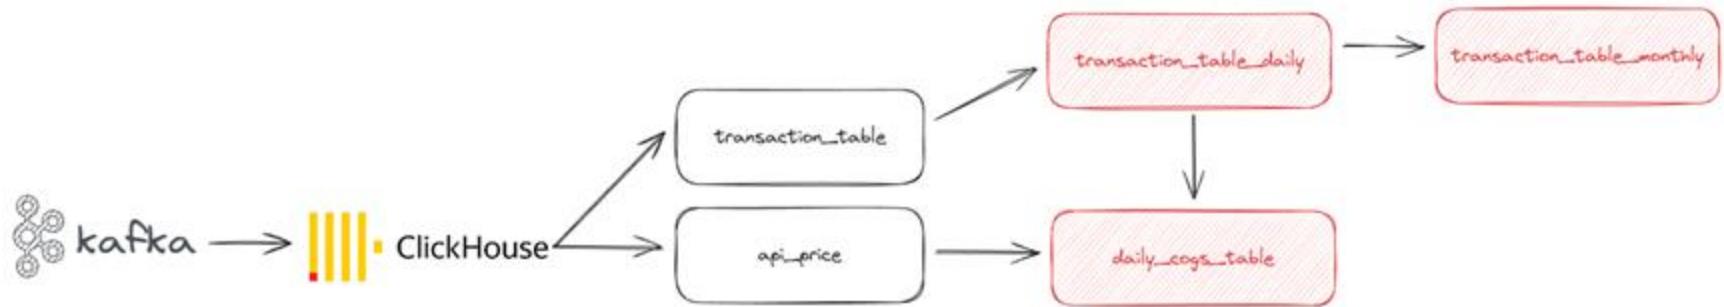
- Listens to a clickhouse table, transforms, and writes into **another physical clickhouse table**.
- Good for pre-aggregating **heavy metrics**



## Projection

- **Optimizes data layouts** that are physically stored **alongside the main table** data and are invisible to the user.
- Good for accelerating data querying





# Results:

Elapsed: 1.260s Read: 9,253,464 rows (433.00 MB)

Elapsed: 2.947s Read: 18,535,783 rows (982.27 MB)

# 98% Realtime 50%

Query speed  
improvement from  
minutes to seconds

Data warehouse  
ingestion

Cost saved from cloud  
expenses

# Important Things to Know Before Migrating to Clickhouse

- Your database design matters. A lot.
  - Denormalize schemas, fewer joins, and choose your sorting keys wisely.
  - Table Engines are commitments. Different engines behaves very differently.
- INSERTs are fast, but deletes & updates are not.
  - Deletes & Updates are labeled as “mutations” in Clickhouse
  - Clickhouse is Eventually Consistent (your updates might propagate later).
- Clickhouse SQL <> PostgreSQL
  - They have similar query syntax, but some differ and Clickhouse also have their own built-in functions
- Backfilling needs more planning
  - Materialized views do not backfill automatically from their source tables (projections do).
  - Think before you create cascading Materialized Views.

# Key Takeaways

- PostgreSQL is great until your workload outgrows it
  - Our workload couldn't scale with Postgres
- Clickhouse lets you redesign for speed while still being cost effective
- Streaming data is now possible thanks to Kafka
- Clickhouse offers many table engines to optimize your workload.
- Materialized Views and Projections do the heavy lifting inside Clickhouse.

# References & Further Reading/Watching

- <https://clickhouse.com/docs/engines/table-engines> [Table Engines]
- <https://clickhouse.com/blog/clickhouse-faster-queries-with-projections-and-primary-indexes> [Projections]
- <https://clickhouse.com/blog/using-materialized-views-in-clickhouse> [Materialized Views]
- <https://clickhouse.com/docs/managing-data/materialized-views-versus-projections> [Comparing materialized view and projections]
- <https://clickhouse.com/docs/sql-reference> [SQL Reference]
- [https://www.youtube.com/watch?v=6mmQU0mA-T0&list=PL0Z2YDlm0b3gtldcZI3B\\_8bMJclDOvY8s](https://www.youtube.com/watch?v=6mmQU0mA-T0&list=PL0Z2YDlm0b3gtldcZI3B_8bMJclDOvY8s) [Clickhouse “How To” Playlist]

# Thank you!

 rayantonius

 rayantonius

 Rocksus