

Petabyte-Scale Log Storage for Microservices

March 12, 2025

Daniel Muino

N

Problem

Ingest PBs/day of operational logs.

Fast queries

Reasonable costs

History

Elastic Search

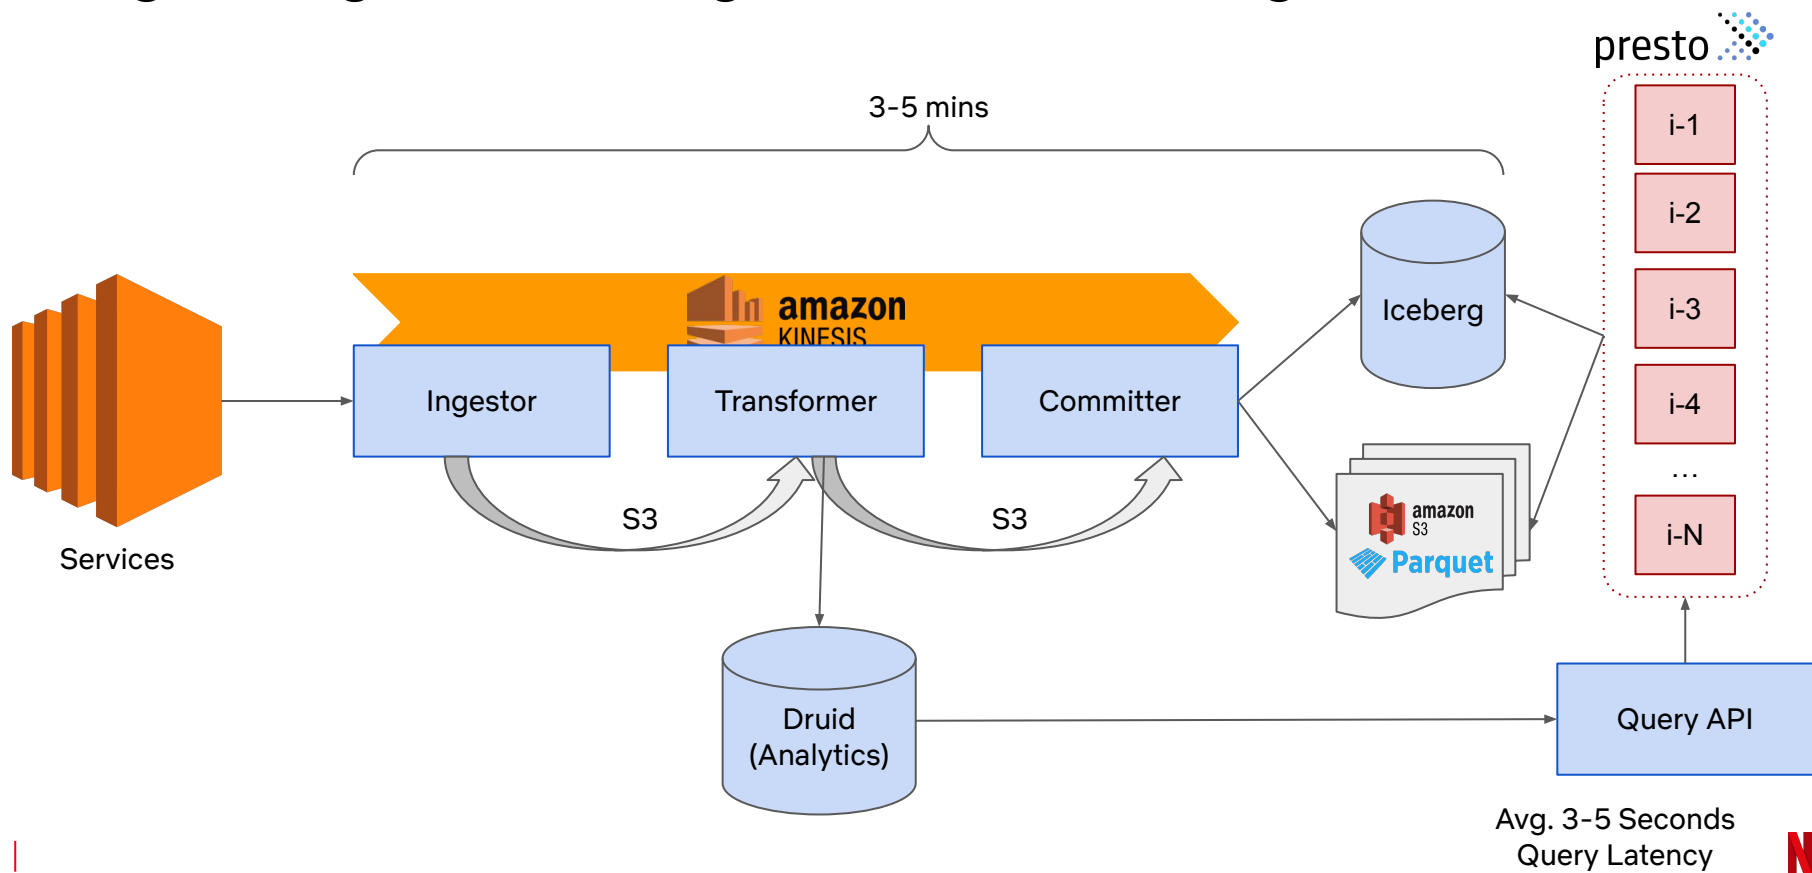
Fast but very expensive.

History

Trino + Iceberg

Cheap but slow queries + stale data

Insight Logs V1: Using Presto + Iceberg

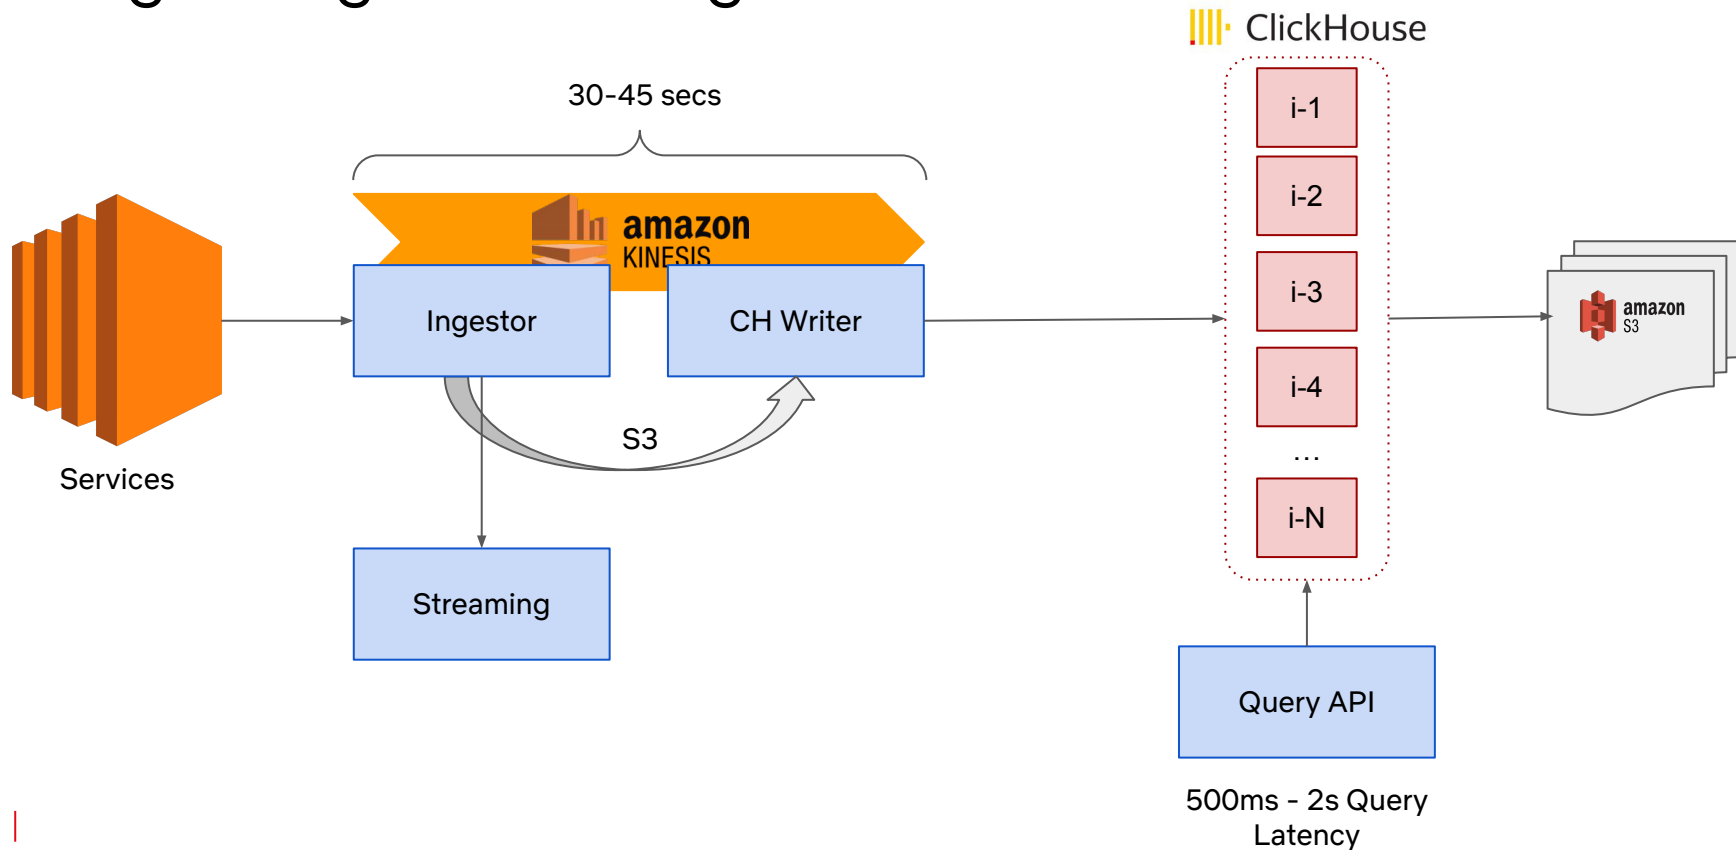


Maybe a different DB

ClickHouse

Fast queries, fresh data, not trivial to operate

Insight Logs V2: Using ClickHouse

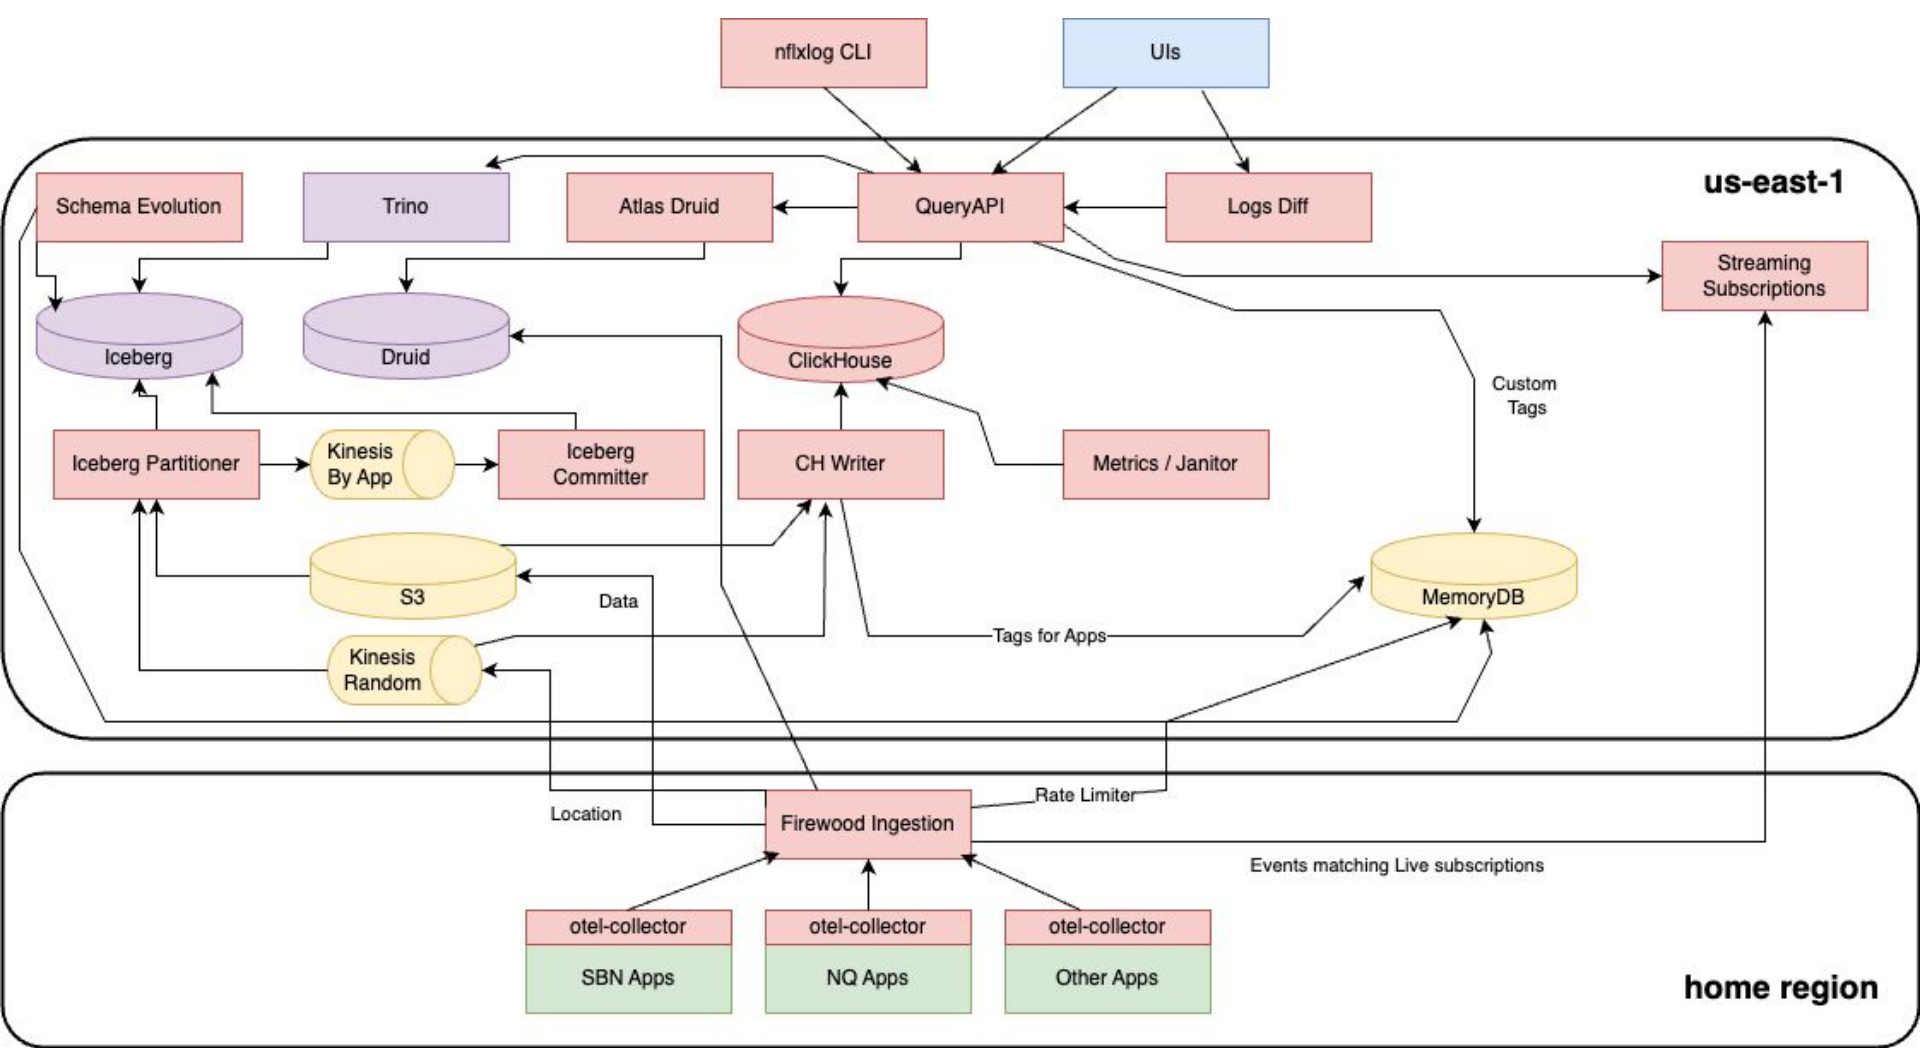


ClickHouse Challenges

- Stateful nodes
- Keeper
- Dealing with expensive queries

ClickHouse as Hot Tier

- Use CH for recent data
- Use Iceberg for the rest



Live Demo


[Alpha UI](#)

[Stable UI](#)

Thank You.

Daniel Muino

N



One of multiple CH deployments

