

Sparse data storage and query patterns in Clickhouse

Amit Sanjay Sadafule



About Me



Co-founder and Head of Technology @Manthan

Built and scaled platforms handling millions of vehicle data-points for India's largest automobile manufacturer. Real-time track & trace systems powering vehicle production plants across the country.

Amit Sanjay Sadafule

 amit-sadafule-366b7225

 @amit.sadafule

- Love challenges of distributed systems
- Pragmatic Engineer
- Treks and exploration

Happy to jam on distributed systems, code quality, or trekking stories.
Coffee ☕ invitations always welcome!

The Problem

Store and query data from thousands of vehicles

Each vehicle can transmit nearly 5,000 keys every minute, referred to as PIDs

Characteristics of these PIDs:

- PID Data is Time-Series Data
- Sparse Data
- 3 data types: Integer, Float and String
- Consistent Data Types
- Append-Only Data

Query Pattern

Retrieving PID Values for a group of vehicles within a specified time range

Aggregation Queries: average, count, sum on PID values.

But... What is Sparse Data?

| OBD data | | | | | |
|------------|-----------|----|----|-----|----|
| vehicle_id | timestamp | p1 | p2 | ... | pn |
| 1 | 123456 | x | x | | |
| 2 | 123456 | | x | | |
| 3 | 123456 | x | | | x |
| ... | ... | | | | |
| n | 123456 | | | | x |

How to efficiently store the pid data and retrieve it?

Choose right schema

Horizontal (Wide) Table Structure

Vertical Table Structure

| horizontal table structure | | | | | |
|----------------------------|---------------|----|----|-----|----|
| vehicle_id | timestamp | p1 | p2 | ... | pn |
| 4 | 1686830420666 | x | | ... | y |

| verticle table structure | | | |
|--------------------------|---------------|-----|-------|
| vehicle_id | timestamp | pid | value |
| 4 | 1686830420666 | p1 | x |
| ... | ... | ... | ... |
| 4 | 1686830420666 | pn | y |

Option 1 - Sparse Matrix

| sparse data | | | | | | |
|-------------|---------------|----|----|-----|----|---|
| vehicle_id | timestamp | p1 | p2 | ... | pn | |
| 1 | 1686830420666 | x | x | | | |
| 2 | 1686830420666 | | x | | | |
| 3 | 1686830420666 | x | | | | x |
| ... | ... | | | | | |
| n | 1686830420666 | | | | | x |

Pros:

- Different compression algorithms per column
- Primary key can consists of vehicle id and timestamp

Cons:

- High RAM requirements
- Slower querying on PIDs

Option 2 - Separate Tables for Each Data Type

| integer data | | | |
|--------------|---------------|-------|-------|
| vehicle_id | timestamp | pid | value |
| 1 | 1686830420666 | p1 | 22 |
| 2 | 1686830420666 | p2 | 36 |
| 3 | 1686830420666 | p44 | 1 |
| ... | ... | | |
| n | 1686830420666 | p1888 | 31 |

| string data | | | |
|-------------|---------------|-------|-------|
| vehicle_id | timestamp | pid | value |
| 1 | 1686830420666 | p50 | "r" |
| 2 | 1686830420666 | p60 | "an" |
| 3 | 1686830420666 | p61 | "do" |
| ... | ... | | |
| n | 1686830420666 | p1999 | "m" |

| float data | | | |
|------------|---------------|-------|-----------|
| vehicle_id | timestamp | pid | value |
| 1 | 1686830420666 | p40 | 12.344534 |
| 2 | 1686830420666 | p1490 | 34.2323 |
| 3 | 1686830420666 | p1600 | 1.00 |
| ... | ... | | |
| n | 1686830420666 | p876 | 43.12 |

Pros:

- Faster queries
- Significantly lower RAM usage during data insertion
- No sparseness

Cons:

- Caller code must distinguish between different tables
- Joins may be required to get PIDs of different data types

Option 3 - Data Type-Specific Columns in the Same Table

| data | | | | | |
|------------|---------------|-------|-----------|-------------|--------------|
| vehicle_id | timestamp | pid | int_value | float_value | string_value |
| 1 | 1686830420666 | p1 | 22 | NULL | NULL |
| 2 | 1686830420666 | p2 | 36 | NULL | NULL |
| 3 | 1686830420666 | p40 | NULL | 51.6 | NULL |
| 4 | 1686830420666 | p1490 | NULL | 413.25 | NULL |
| ... | ... | ... | ... | ... | ... |
| 1200 | 1686830420670 | p50 | NULL | NULL | "r" |
| 1201 | 1686830420670 | p60 | NULL | NULL | "an" |
| ... | ... | ... | ... | ... | ... |
| n | 1686830420670 | p369 | 31 | NULL | NULL |

Pros:

- Simplified caller code
- Faster queries

Cons:

- Potential sparseness in the data
- Higher RAM requirements during data insertion than Option 2

Optimizations and Best Practices

Avoid Nullable Columns

Using default values instead of nulls can improve query performance. This because without default values ClickHouse have to store and refer to extra storage required for null references.

Compression

Efficient compression algorithms like ZSTD can significantly reduce storage requirements. Ensure that the chosen compression method aligns with the characteristics of your data.

Query Performance

Optimize queries by ensuring the primary key includes the PID, which can significantly speed up query performance.

Option 3.1 - Data Type-Specific Columns in the Same Table - with default values

| data | | | | | |
|------------|---------------|-------|-----------|-------------|--------------|
| vehicle_id | timestamp | pid | int_value | float_value | string_value |
| 1 | 1686830420666 | p1 | 22 | 0.0 | "" |
| 2 | 1686830420666 | p2 | 36 | 0.0 | "" |
| 3 | 1686830420666 | p40 | 0 | 51.6 | "" |
| 4 | 1686830420666 | p1490 | 0 | 413.25 | "" |
| ... | ... | ... | ... | ... | ... |
| 1200 | 1686830420670 | p50 | 0 | 0.0 | "r" |
| 1201 | 1686830420670 | p60 | 0 | 0.0 | "an" |
| ... | ... | ... | ... | ... | ... |
| n | 1686830420670 | p369 | 31 | 0.0 | "" |

Pros:

- Simplified caller code
- Faster queries as null values are avoided

Cons:

- Higher RAM requirements during data insertion than Option 2

Experimental analysis

Data Generation: 2000 PIDs per vehicle, with 70% float values, 29% integer values, and 1% string values.

Data Volume: 3 billion PID values for 1000 vehicles over a 24-hour period.

Machine Setup: Mac OS Sonoma, Apple M3 MacBook Air, 16GB RAM, 8 cores.

ClickHouse Version: v24.8.1.

Experimental analysis - Query Performance

Get all PIDs and values for a vehicle and for entire 2hr duration

| Table | Option | Rows scanned | Max RAM used (MiB) | Time taken (Sec.) | Uncompressed data size (MB) |
|----------------------|------------|-----------------|--------------------|-------------------|-----------------------------|
| pid_float | Option 2 | 196.61 thousand | 4.23 | 0.007 | 7.22 |
| pid_integer | Option 2 | 73.73 thousand | 1.05 | 0.005 | 2.86 |
| pid_string | Option 2 | 16.38 thousand | 73.53 KiB | 0.007 | 425.91 KB |
| pid_without_defaults | Option 3 | 262.14 thousand | 4.08 | 0.253 | 14.57 |
| pid_with_defaults | Option 3.1 | 253.95 thousand | 1.08 | 0.169 | 10.88 |

References

[ClickHouse - Lightning Fast Analytics for Everyone]

<https://www.vldb.org/pvldb/vol17/p3731-schulze.pdf>

[Extending RDBMSs To Support Sparse Datasets]

<https://pages.cs.wisc.edu/~alanh/sparse.pdf>

[Too Wide or Not Too Wide | That is the ClickHouse Question]

<https://altinity.com/blog/too-wide-or-not-too-wide-that-is-the-clickhouse-question>

[Avoid Nullable Columns]

<https://clickhouse.com/docs/en/cloud/bestpractices/avoid-nullable-columns>