



Ankit Gupta

Founding Member
ThriveStack

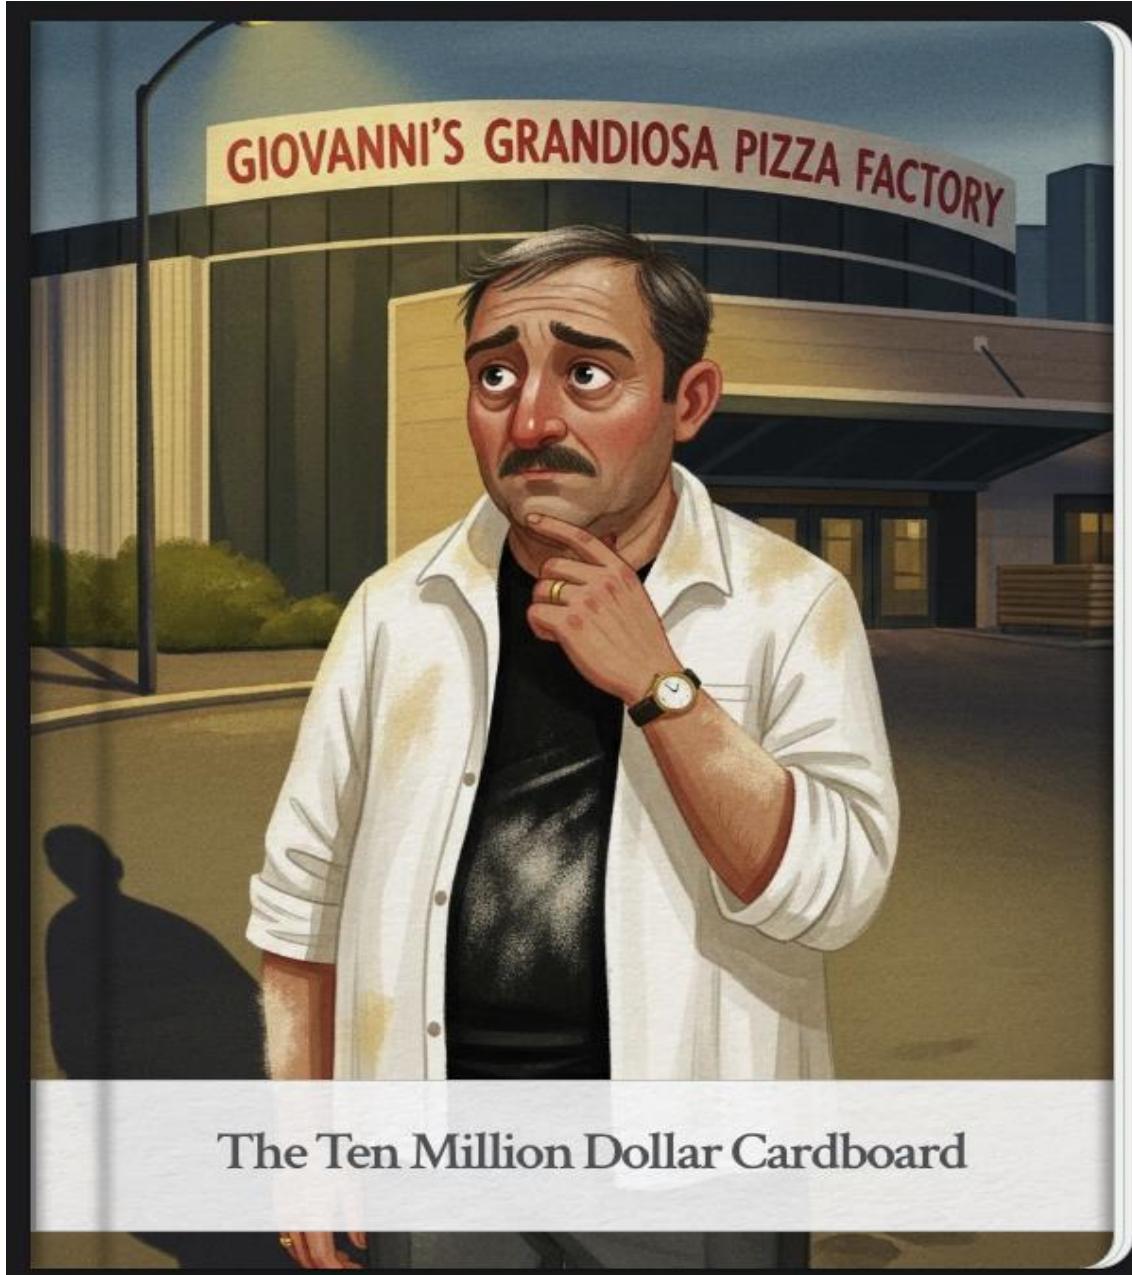
ankitg@thrivestack.in

Why We Chose at ThriveStack

<https://www.thrivestack.ai/>

Agenda

- 1** What is ThriveStack?
- 2** How we chose Clickhouse ?
- 3** Using ClickHouse Effectively
- 4** Our Next Steps



The Ten Million Dollar Cardboard



****1000 leads**
****\$10M budget**

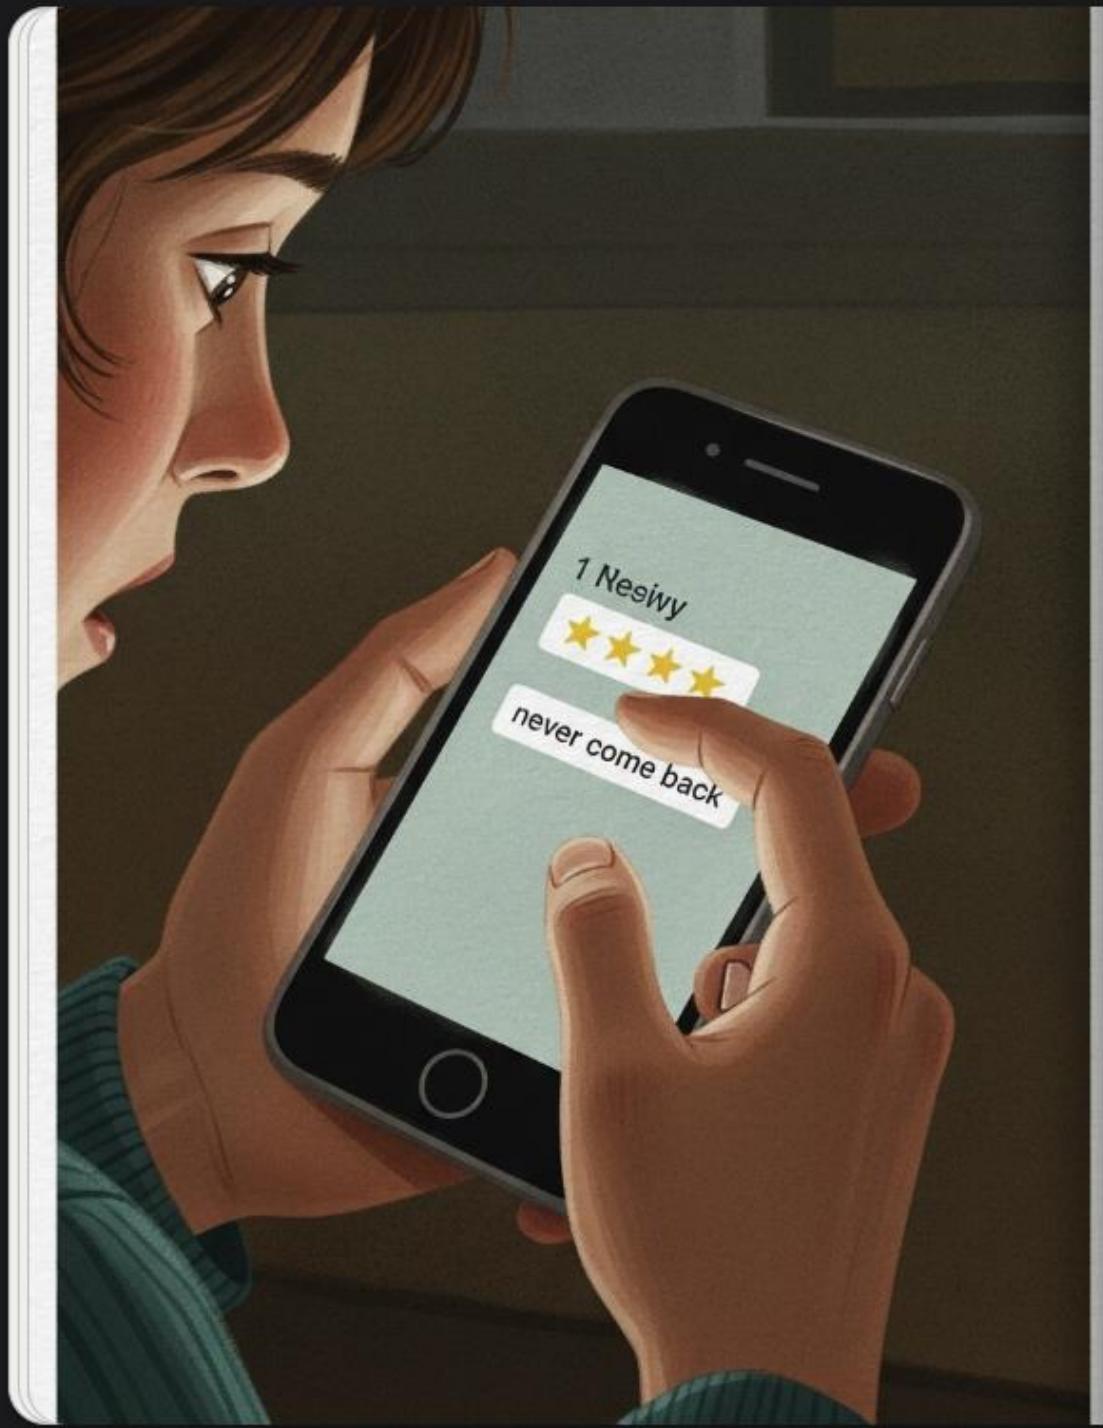
"**M**arketing promised you 1000 leads - you spent your \$10M budget."



"But you got 500 visitors walking in."



"Only 10 of them bought pizza."



"2 of them disliked it and vowed to never come back again."

Department Victory Laps

Marketing Team

500 website visits!

"Our campaigns are smashing it - look at this traffic surge!"

Sales Team

12 orders taken!

"CRM is performing brilliantly - pipeline's never been stronger!"

Operations Team

10 pizzas delivered!

"We're scaling delivery capacity like champions!"

Customer Service

2 complaints handled!

"Response times are excellent - customer satisfaction sorted!"



"You need to see the whole story. You
need to use ThriveStack Full Journey
Growth Intelligence."



"You must Measure your Growth,
Identify the Drivers, and Fix the Leaks."



The Universal Problem

This scenario plays out in countless businesses across every industry. Everyone has data, everyone has budget, but nobody has the story.

Unified Growth Intelligence

Full Funnel

Connected

Full Journey Analytics

Seamlessly integrate marketing, product, revenue, and customer success data into a single source of truth.

Marketing 

CRM 

Product 

Revenue 

Correlated

Growth Leaks Detector

250+  powered growth signals to automatically identify growth drivers and leaks.

Consolidated

Single platform for Growth

Eliminate data and tooling silos. Align all your teams with insights, customized views and scorecards

Connected, Correlated, Consolidated



Awareness

Visitors



Acquisition

Users
Accounts



Activation

Milestones
TTV



Monetization

Revenue
Growth



Engagement

Features
Depth



Retention

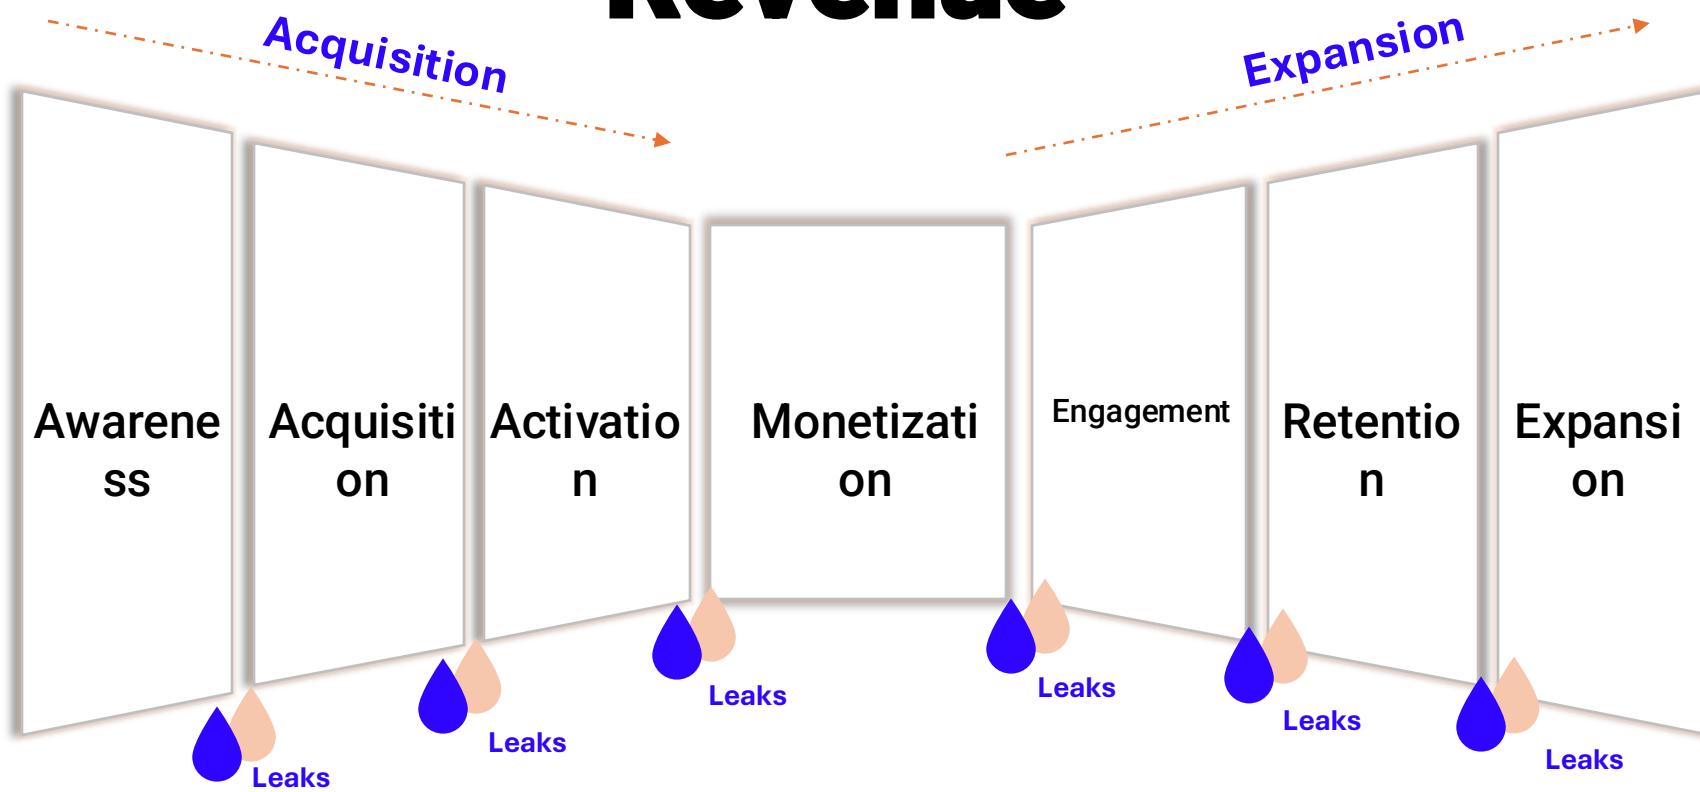
Churn
At Risk



Expansion

Accounts
Revenue
Up/X Sell

Every Company leaks Customers and Revenue



The questions are...

Where?

How Much?

Why?

Can I stop them?

How ThriveStack Chose ClickHouse

Parameters for Choosing OLAP:

1 Performance Fit

Query Latency, **Throughput Metrics**, Data Volume Handling

2 Cost/Pricing Fit

Predictable Costs, Operational Costs, Infrastructure Costs, Customer Scaling

3 Technical Architecture Fit

Stack Integration, Event Streaming, CDC Support, Batch Processing, Resource Isolation, Geographic Distribution

4 Customer Analytics Use-Cases Fit

Real-Time Customer Behavior, Customer Data Models, Cohort, Behavioral Segmentation, Customer Scores

Options We Considered OLAP:

1



2



3



4



BigQuery

5



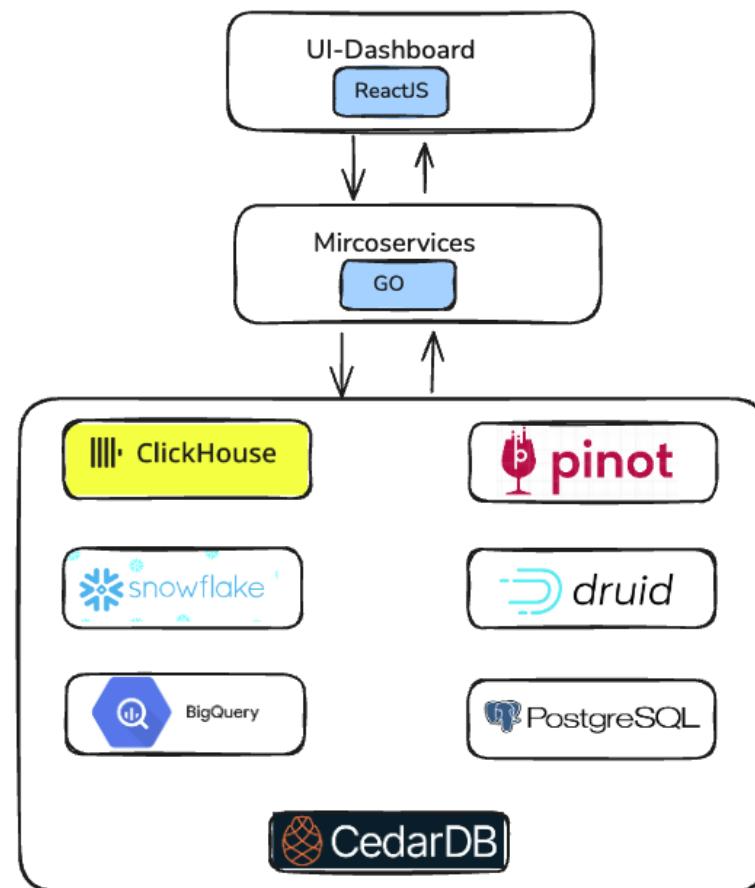
6



7



Test Setup



Total Metrics: **47** [Real Time + Score Computation + Correlation]

Data : **14** million rows + **12** tables + **12** GB Size

Use-Case & Metrics

This screenshot shows the ThriveStack dashboard for the 'Demo Product' under the 'Acme Labs' account. The interface features a top navigation bar with tabs for 'Production' and 'Demo Mode'. Below the navigation is a main dashboard area with a 7-day time range selector. The dashboard is organized into several sections: 'Awareness', 'Acquisition', 'Activation', 'Monetization', 'Engagement', 'Retention', and 'Expansion'. Each section contains key performance indicators (KPIs) and a brief description. Below these are detailed insights for 'Awareness' and 'Acquisition'.

Awareness

58K Visitors ↑ 12.0% vs last week

Acquisition

5K Users ↓ 8.0% vs last week

2K Accounts ↓ 2.0% vs last week

Activation

35% Activation Rate ↑ 15.7% vs last week

Monetization

\$166K MRR ↑ 12.0% vs last week

2K Customers ↑ 5.0% vs last week

Engagement

2K Users ↓ 8.0% vs last week

2K Accounts ↓ 5.0% vs last week

Retention

400 At risk Customers ↓ 5.0% vs last week

Expansion

\$7K MRR ↑ 5.0% vs last week

30 Customers ↑ 5.0% vs last week

Awareness Insights

Visitors ↗ 1295 ↑ 12.5%

Page Views ↗ 2758 ↑ 0.8%

Visitor Traffic by UTM

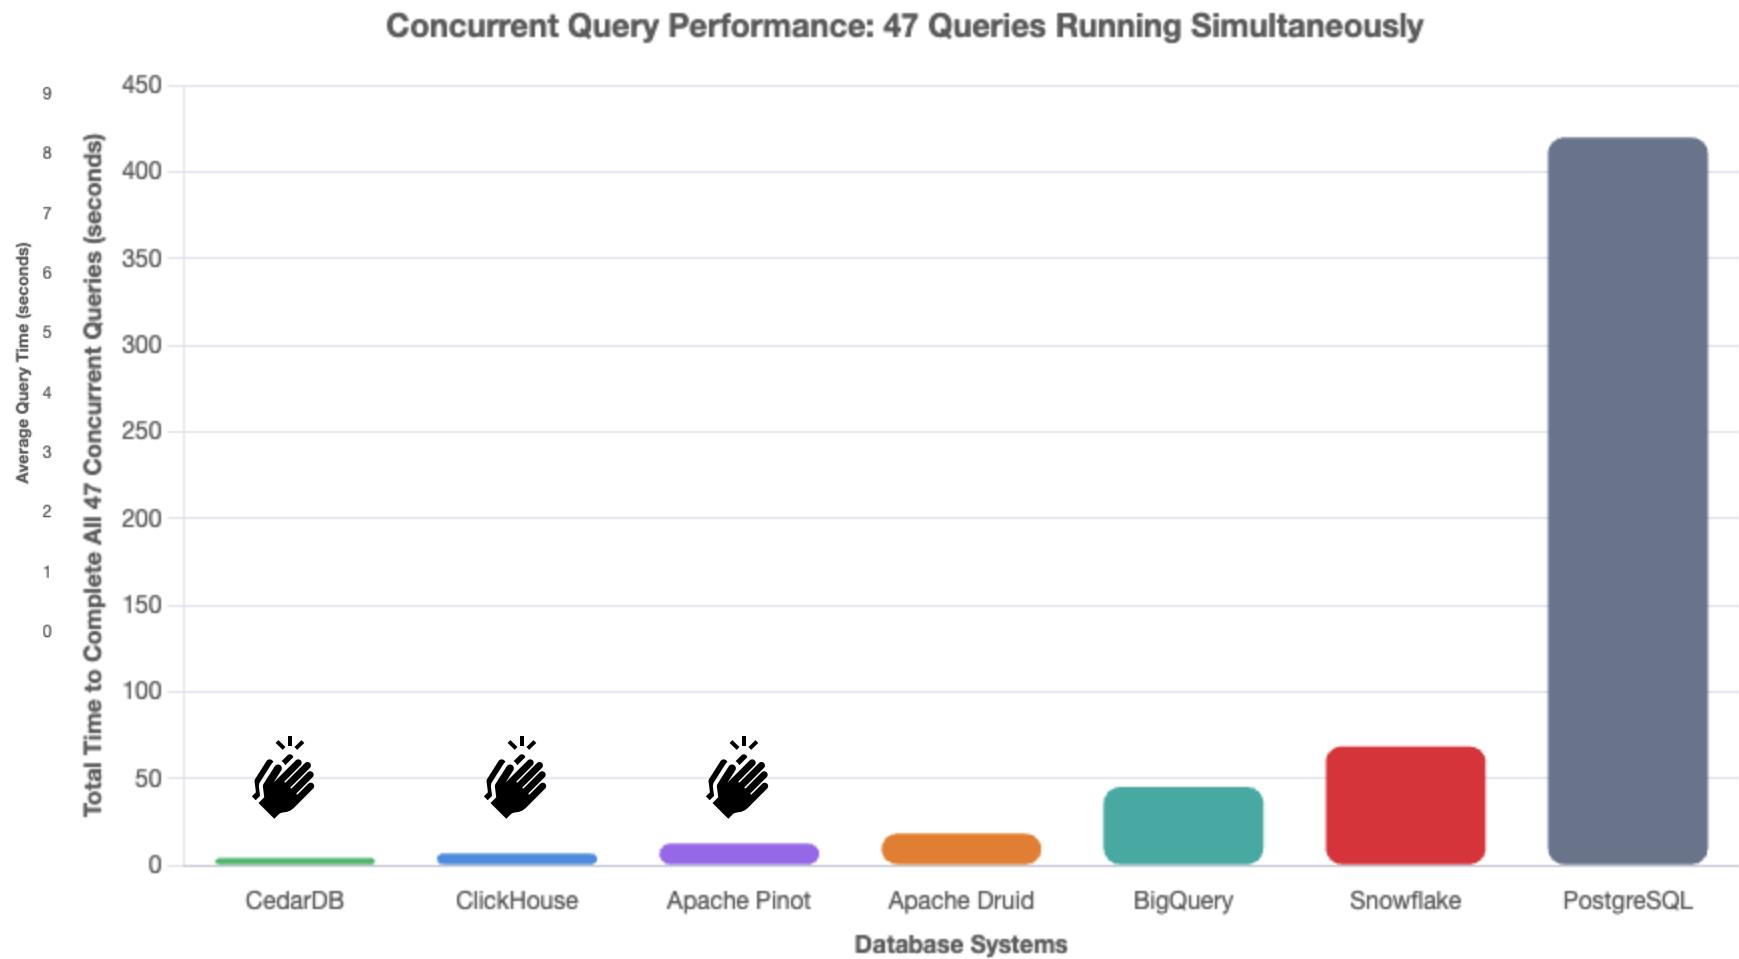
UTM source: Overall

Aug 4 Aug 5 Aug 6 Aug 7 Aug 8 Aug 9 Aug 10 Aug 11

Overall Instagram Direct Facebook LinkedIn Twitter

Acquisition Insights

Performance Results



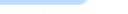
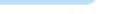
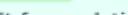
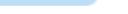
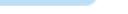
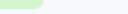
Database	Total Time for 47 Concurrent Queries
CedarDB	4.2s
ClickHouse	6.8s
Apache Pinot	12.5s
Apache Druid	18.3s
BigQuery	45.2s
Snowflake	1m 8s
PostgreSQL	7m 0s

<https://benchmark.clickhouse.com>

Total Metrics: **47** [Real Time + Score Computation + Correlation]
Data : **14 million rows + 12 tables + 12 GB Size**

Cost Fit Results

System	Predictable Costs	Operational Costs	Infrastructure Costs	Customer Scaling
● ClickHouse	 Good Per-minute billing with memory scaling controls. Recent 30% price increase in 2025.	Excellent Fully managed cloud service. Auto-scaling and maintenance included. Minimal ops overhead.	Good Compute: ~\$0.35-0.50/vCPU/hour. Storage: \$25.30/TiB/month. Data egress now charged.	Excellent Scales to petabytes. Auto-pause during inactivity. Linear cost scaling with usage.
● Apache Pinot	Excellent Open source = \$0 license. StarTree Cloud offers fixed monthly plans (\$25-299/month).	Poor Self-managed requires significant DevOps expertise. StarTree Cloud reduces this burden.	Variable Self-hosted: Your infrastructure costs. StarTree: \$0.162/vCPU-hour + storage costs.	Excellent Proven at Uber scale. Tiered storage reduces costs. Companies save \$2M+ annually vs alternatives.
● Apache Druid	Good Open source = \$0. AWS deployment: ~\$715-13,645/month depending on scale. Support: \$48k-96k/year.	Poor Complex to operate. Requires dedicated ops team. Imply provides managed service alternative.	Fair Higher infrastructure requirements vs competitors. Storage + compute tightly coupled.	Good Scales well but requires careful capacity planning. Non-linear cost increases at scale.
● CedarDB	 Excellent Community Edition = Free forever. Enterprise = Usage-based pricing, costs scale with activity.	Good PostgreSQL compatible. Auto-adapts to hardware. Quick deployment, minimal configuration needed.	Good Runs on your hardware/cloud. No data transfer costs. Efficient storage utilization.	Good New system (2024). HTAP capabilities. Automatic resource adaptation to workload.
● Snowflake	Poor Complex credit system. Costs can escalate quickly. Difficult to estimate without monitoring.	Excellent Fully managed service. Auto-scaling, maintenance, security handled. Zero operational overhead.	Poor Storage: \$23-40/TB/month. Compute: \$2-4/credit/hour. Data transfer and feature costs extra.	Fair Excellent scalability but expensive at scale. Enterprises spend \$15k-50k/month typically.
● PostgreSQL	Excellent Open source = \$0 license cost. Cloud managed: \$0.04-0.05/vCPU/hour, \$5-6/GB RAM/month.	Variable Self-managed: High ops burden. Managed services: Low ops but higher costs.	Excellent Lowest infrastructure costs. Runs on minimal hardware. Efficient resource utilization.	Fair OLTP-focused. Limited analytical performance at scale. Requires read replicas for scaling.
● BigQuery	 Good On-demand: \$5/TB processed. Flat-rate: \$8,500/month (500 slots). Storage: \$20-40/TB/month.	Excellent Serverless, fully managed. No infrastructure management. Auto-scaling included.	Good Pay per query or reserved capacity. Lower storage costs than Snowflake. Google Cloud only.	Excellent Petabyte scale. Automatic scaling. Cost-effective for large analytical workloads.

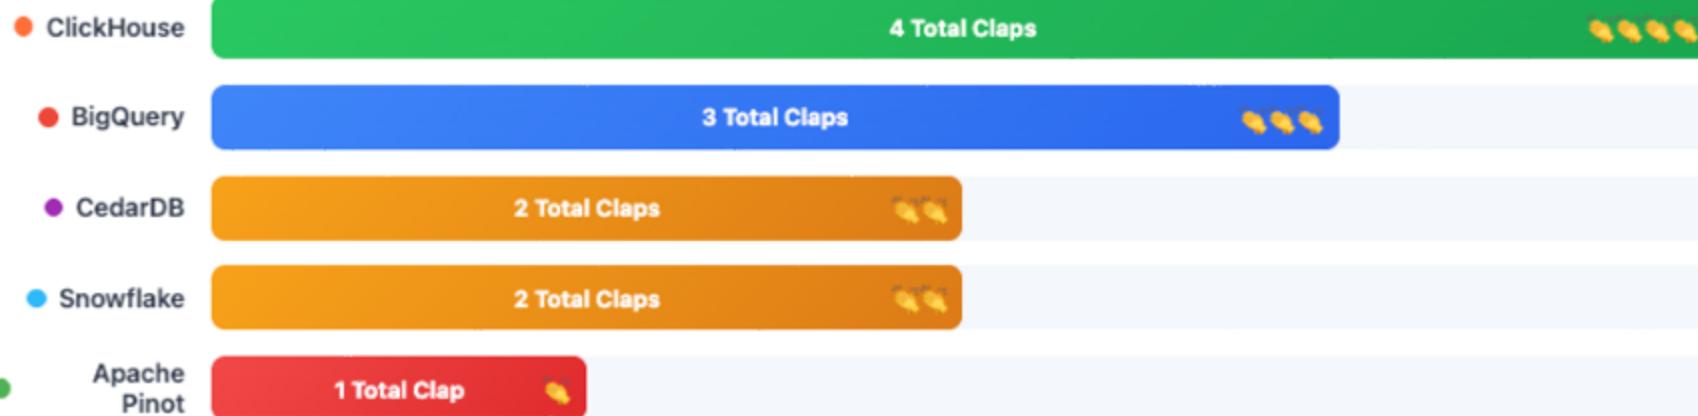
Business Use-case Fit						
System	Real-time Customer Behavior	Customer Data Models	Cohort Analysis	Behavioral Segmentation	Customer Scores	
● Snowflake	 Fair Batch processing focus. Snowpipe has latency. Not optimal for real-time dashboards.	 Excellent Best-in-class data modeling. JSON support. Time travel. Data sharing.	 Excellent Advanced window functions. Historical analysis via time-travel. Rich SQL analytics.	 Excellent Comprehensive analytics. ML integration. Secure data sharing for segments.	 Excellent Snowpark ML. Native ML functions. Statistical functions. ML platform integrations.	
● BigQuery	 Good Near real-time streaming. Better for batch processing. Some latency for real-time use.	 Excellent Flexible schema. Nested fields. JSON support. Table clustering.	 Excellent Advanced analytical functions. Pre-built cohort templates. BigQuery ML integration.	 Excellent Petabyte-scale analytics. Advanced SQL. ML integration. GCP ecosystem.	 Excellent BigQuery ML for scoring. Built-in algorithms. TensorFlow integration. AutoML.	
● ClickHouse	 Excellent Sub-second queries on billions of events. Real-time materialized views. Stream processing.	 Good Strong dimensional modeling. JSON support. Schema evolution. Window functions.	 Excellent Native time-series functions. Retention analysis. Custom cohort calculations via SQL.	 Excellent Advanced aggregations. ML functions. Real-time segment updates.	 Good Custom scoring via SQL. Statistical functions. ML pipeline integration.	
● Apache Pinot	 Perfect Built for real-time user-facing analytics. Sub-second latency. Direct Kafka ingestion.	 Excellent Star-tree indexes. JSON support. Pre-aggregation. Schema evolution.	 Good Time-based partitioning. Window functions. Requires complex SQL for cohorts.	 Excellent Real-time segmentation. Bitmap indexes. High-cardinality dimensions.	 Good Real-time scoring. Custom aggregations. ML serving integration.	
● Apache Druid	 Excellent Real-time Kafka ingestion. Sub-second queries. Time-series analytics. Event-driven.	 Good Dimensional model. JSON support. Roll-up capabilities. Schema flexibility.	 Good Time-based partitioning. Retention analysis. Complex calculations need custom code.	 Excellent Fast filtering. Bitmap indexes. Real-time segments. High-performance OLAP.	 Fair Basic aggregations. Limited ML. Requires external tools for complex scoring.	
● CedarDB	 Good HTAP capabilities. Real-time analytics on transactional data. New system (2024).	 Excellent PostgreSQL compatible. JSON, time-series, vector support. ACID compliance.	 Good Full SQL with window functions. Time-series analytics. PostgreSQL ecosystem.	 Good Advanced SQL analytics. Custom functions. HTAP advantages.	 Good Statistical functions. Vector operations. PostgreSQL ML extensions.	
● PostgreSQL	 Limited OLTP-focused. Not for real-time analytics. Performance degrades at scale.	 Excellent Mature relational model. JSON support. Extensive ecosystem. Custom data types.	 Good Full SQL. Window functions. TimescaleDB extension. Performance concerns at scale.	 Good Rich SQL. Custom aggregations. Performance issues with large segmentation.	 Excellent MADlib ML. PLPython. Statistical extensions. Rich ML ecosystem.	

Technical Architecture Fit

System	Stack Integration	Event Streaming	CDC Support	Batch Processing	Resource Isolation	Geographic Distribution
● ClickHouse	 Excellent ClickHouse Cloud managed integrations. Partner ecosystem.	Excellent ClickPipes for streaming. Native Kafka. Real-time ingestion.	Good ClickPipes CDC connectors. PostgreSQL CDC support.	Excellent Cloud-native batch processing. S3/object storage integration.	Excellent Compute-compute separation. Auto-scaling. Workload isolation.	Good Multi-region deployment. Cross-region data sharing.
● Snowflake	 Excellent Multi-cloud. 3000+ integrations. Native connectors.	Fair Snowpipe with latency. Better for micro-batch.	Good Streams & Tasks. Partner integrations (Fivetran, etc).	Excellent Built for batch. Auto-scaling. Time Travel.	Excellent Virtual warehouses. Complete workload isolation.	Excellent Multi-region. Cross-cloud replication. Data sharing.
● BigQuery	 Good GCP native. Good third-party support. Cloud-locked.	Good Pub/Sub integration. Dataflow streaming. Some latency.	Good Datastream for CDC. Third-party tools.	Excellent Petabyte batch processing. Serverless scaling.	Good Slot reservations. Project isolation. Fair queuing.	Good Multi-region datasets. Cross-region queries.
● Apache Pinot	Fair Kafka ecosystem. Limited BI integrations.	Excellent Built for streaming. Direct Kafka ingestion.	Limited No native CDC. Requires external CDC tools.	Good Hadoop/S3 batch ingestion. Offline tables.	Fair Tenant isolation. Resource management improving.	Fair Manual multi-region setup. Cross-DC replication.
● Apache Druid	Fair Kafka ecosystem. Limited modern integrations.	Excellent Native Kafka/Kinesis. Real-time ingestion.	Limited No native CDC. Requires external solutions.	Excellent Strong batch ingestion. Hadoop/S3 native.	Good Query tiers. Resource management. Multi-tenancy.	Fair Multi-region possible. Complex setup required.
● CedarDB	Excellent PostgreSQL compatible. Existing tool ecosystem.	Good HTAP capabilities. Real-time processing possible.	Good PostgreSQL CDC tools compatible. Logical replication.	Good HTAP batch processing. PostgreSQL tooling.	Fair Early stage. Resource management developing.	Limited New system. Geographic features under development.
● PostgreSQL	Excellent Mature ecosystem. Extensive tooling. Universal support.	Limited Not designed for streaming. Performance issues.	Excellent Native logical replication. WAL-based CDC.	Fair OLTP focused. Limited batch analytics performance.	Fair Database-level isolation. Read replicas.	Good Replication options. Multi-master with extensions.

Winner

Overall Winner Score (Claps 🙌 for each category win)



Overall Winner: ClickHouse 🙌🙌🙌

Winner in Cost Fit + Tech Architecture Fit + 2nd in Performance = Most claps earned across all categories!

What We Learnt About Effective Use Of ClickHouse

Effective Use Of Table engines

- 1
- 2
- 3
- 4

● MergeTree

Standard analytical storage

General-purpose OLAP queries, time-series data, event logging, raw data storage for ad-hoc analytics

SQL

```
-- Website clickstream data CREATE TABLE clicks ( timestamp DateTime,  
user_id UInt32, page_url String, session_id String ) ENGINE =  
MergeTree() ORDER BY (timestamp, user_id);
```

● ReplacingMergeTree

Automatic duplicate elimination

User profiles, configuration tables, slowly changing dimensions where you need the latest version of each record

SQL

```
-- User profiles with updates CREATE TABLE users ( user_id UInt32, name  
String, email String, updated_at DateTime ) ENGINE =  
ReplacingMergeTree(updated_at) ORDER BY user_id; -- Latest record per  
user_id is kept
```

● AggregatingMergeTree

Pre-computed metric aggregation

Real-time dashboards, pre-aggregated reports, KPI tables where you need instant aggregated results

SQL

```
-- Daily sales metrics CREATE TABLE daily_sales ( date Date, product_id  
UInt32, total_sales AggregateFunction(sum, Float64), order_count  
AggregateFunction(count, UInt32) ) ENGINE = AggregatingMergeTree()  
ORDER BY (date, product_id); -- Automatically sums sales and counts  
orders
```

Optimize Primary Key and Partitioning

1

2

3

4

✓ Best Practices

- Choose low-cardinality column as first part of primary key
- Use time-based field (e.g., event_date) for sorting advantage
- Use ORDER BY wisely—determines storage and query efficiency

✗ Avoid

- High-cardinality columns (e.g., UUIDs) in ORDER BY
- Overcomplicating primary keys with too many columns

✓ Best Practices

- Partition on date-based columns (e.g., toYYYYMM(event_date)) for time-series data
- Keep partition sizes between 10M - 100M rows to avoid small-file issues

✗ Avoid

- Excessive partitioning (too many small partitions slow down queries)
- Partitioning on high-cardinality columns like user IDs

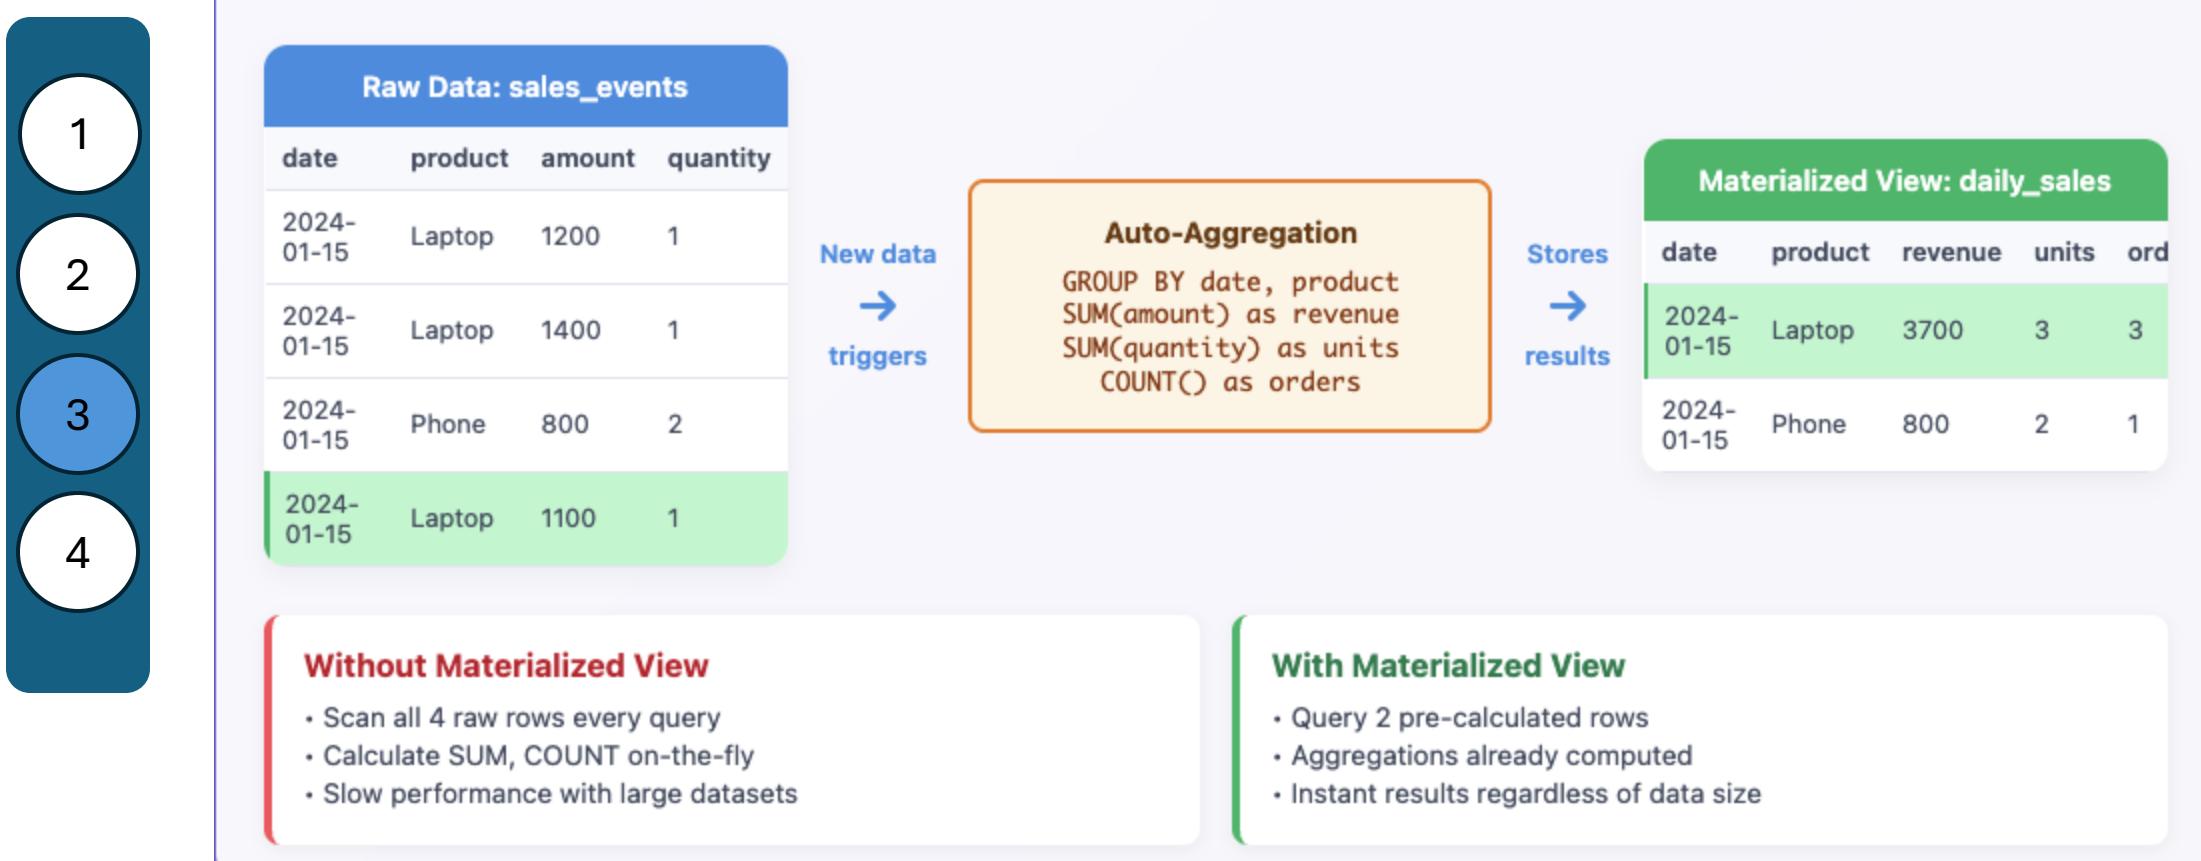
✓ Good Example:

```
CREATE TABLE events( event_date Date, user_id UInt32, event_type String, event_value Float32 ) ENGINE = MergeTree() PARTITION BY toYYYYMM(event_date) ORDER BY (event_date, user_id);
```

📅 Monthly Partitioning Example:

```
CREATE TABLE logs ( log_date Date, log_level String, message String ) ENGINE = MergeTree() PARTITION BY toYYYYMM(log_date) ORDER BY (log_date, log_level);
```

Precompute Aggregations using MV



- 1
- 2
- 3
- 4

ClickHouse MCP Integrations

The screenshot illustrates the ClickHouse MCP integration process:

- Code Editor:** Shows a file named `cloude_desktop_config.json` containing configuration for ClickHouse servers and a Python command to run.
- Query Interface:** A window titled "New User Signups in Last 7 Days" shows a series of steps:
 - Check tables for user signup information.
 - Identify events in the `events` table.
 - Get new users signed up in the last 7 days.
 - Count unique users signed up in the last 7 days.
 - Create a bar chart visualizing daily signups.
- Visualization:** A bar chart titled "New User Signups - Last 7 Days" displays the count of new users per day. The data is summarized as follows:

Date	New Users
05-09	8
05-10	2
05-11	16
05-12	1
05-13	12
05-14	14
05-15	9
05-16	8

The Results...

PRODUCT DEVELOPMENT

20% Reduction

in Engineering Resources

PRODUCT EXPERIMENTATION

3x faster

in release velocity

Develop → Release → Feedback

COGS COSTS

40% higher

More engineering resources now use
ClickHouse as opposed to only Backend /
Data engineers

Thank you



Give it a spin

Experience it yourself, at your own pace

Signup & try

<https://www.thrivestack.ai/>

Try the demo mode

After you signup, try the demo mode at the top of the screen



Schedule a demo/Onboarding

Bring in your team for a Demo/product onboarding

- Learn how to get the best out of the product
- Customize it to the stage of the company

[Schedule a call](#)

