

ClickHouse



# meetup

**Work Club @ 8 Chifley**

Sydney, Australia

June 19, 2025 at 5:30 PM AEST



# Tech Talks



## **Maximising Analytics with ClickHouse and Kafka Integration**

Johnny Mirza, ClickHouse Solution Architect, & Ravi Bhardwaj, Confluent Principal Technical Support Engineer



## **Nightwatch: How Laravel Monitors Billions of Events a Month with ClickHouse**

James Carpenter, Senior Infrastructure Engineer @ Laravel



## **Real-Time CDC at Scale: Building Secure Data Pipelines with ClickHouse Cloud and AWS PrivateLink**

Peter Hanssens, Founder of DataEngBytes & Founder and Principal Consultant @ Cloud Shuttle



# Maximising Analytics with ClickHouse and Kafka Integration

Powering Enterprise AI at Scale

2025

||||· ClickHouse

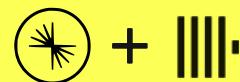
# Ravi Bhardwaj

## Technical Support Engineer, Confluent

 [ravi-bhardwaj](#)

 [ravib777](#)

- Full Stack Trace Engineer by trade
- 15 years in Middleware & Messaging space



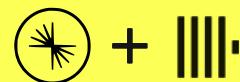
# **Johnny Mirza**

## **Solution Architect @ ClickHouse**

 [johnnym](#)

 [jmirza-debug](#)

- Enthusiastic about Big Data and Analytics
- Excited about the future of AI
- Marketing, Community Support, DevRel etc..



# Agenda

**01**

Introduction

**02**

Apache Kafka 101

**03**

Integrating Kafka with  
ClickHouse

**04**

Recent improvements to  
Kafka Support

**05**

Demo

**06**

Questions

# Introduction ClickHouse



# ClickHouse

2009  
*Prototype*

2012  
*Production*

2016  
*Open Source*

2021  
*ClickHouse Inc.*

2022  
*ClickHouse Cloud*

**The Most Popular Analytics Database on the Planet**

#1  
*Analytics DB on DB-Engines*

Over  
39,000  
*GitHub Stars*

Over  
200,000  
*Community Members*

# Key Features

## Some of the cool things ClickHouse can do

### 1 Speaks SQL

*Most SQL-compatible UIs, editors, applications, frameworks will just work!*

### 2 Lots of writes

*Up to several million writes per second - per server.*

### 3 Distributed

*Replicated and sharded, largest known cluster is 4000 servers.*

### 4 Highly efficient storage

*Lots of encoding and compression options - e.g. 20x from uncompressed CSV.*

### 5 Very fast queries

*Scan and process even billions of rows per second and use vectorized query execution.*

### 6 Joins and lookups

*Allows separating fact and dimension tables in a star schema.*



# Use cases



## Logs, events, traces

Monitor with confidence your logs, events, and traces. Detect anomalies, fraud, network or infrastructure issues, and more.



**zomato**

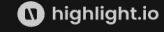


**runreveal**



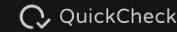
## Real-time Analytics

Power interactive applications and dashboards that analyze and aggregate large amounts of data on the fly. Run complex internal analytics in ms, not mins or hrs.



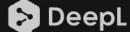
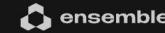
## Business intelligence

Interactively slice and dice your data for analysis, reporting, and building internal applications. Evaluate user behaviors, ad and media perf, market dynamics, and more.



## ML and Gen AI

Execute fast and efficient vector search. Plug-and-play Generative AI models from any provider. Use lightning-fast aggregations to power model training at petabyte scale.



# Row vs Column



**Row-based**

**Excess disk Reads** We need data in column, but it's stored in rows. So entire rows are read.



**Column-Based**

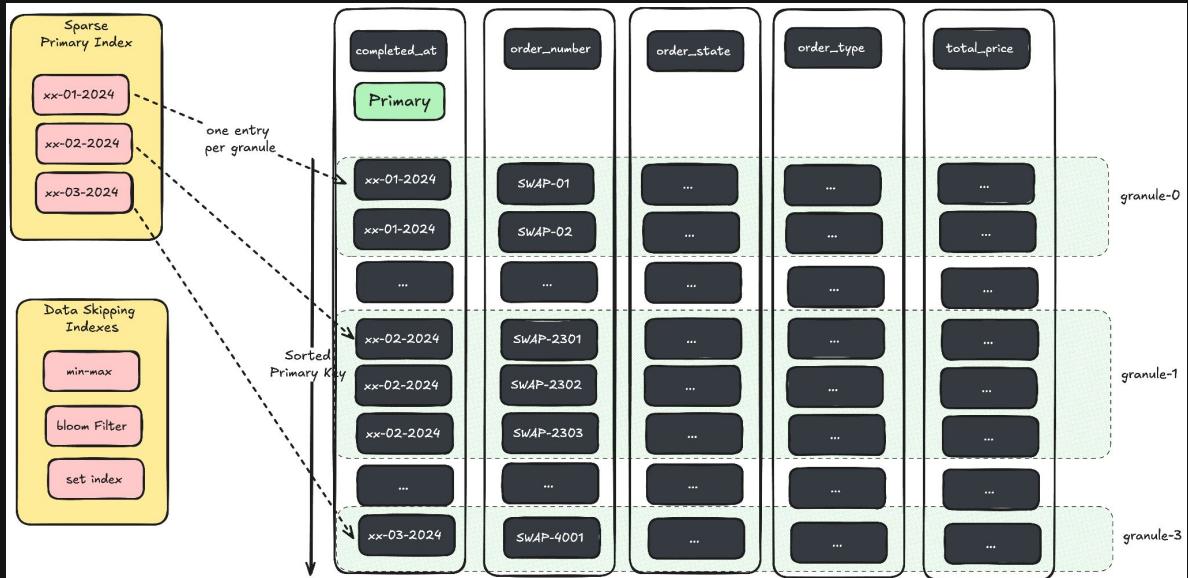
**Column Pruning** it only pick the subset of columns that are openly declared in the SQL query, thus eliminating 99% columns (990 of 1000 columns) from the disk read consideration.

## Vectorized Execution

data processed in **large arrays**, which could be the entire length of the column loaded into RAM.

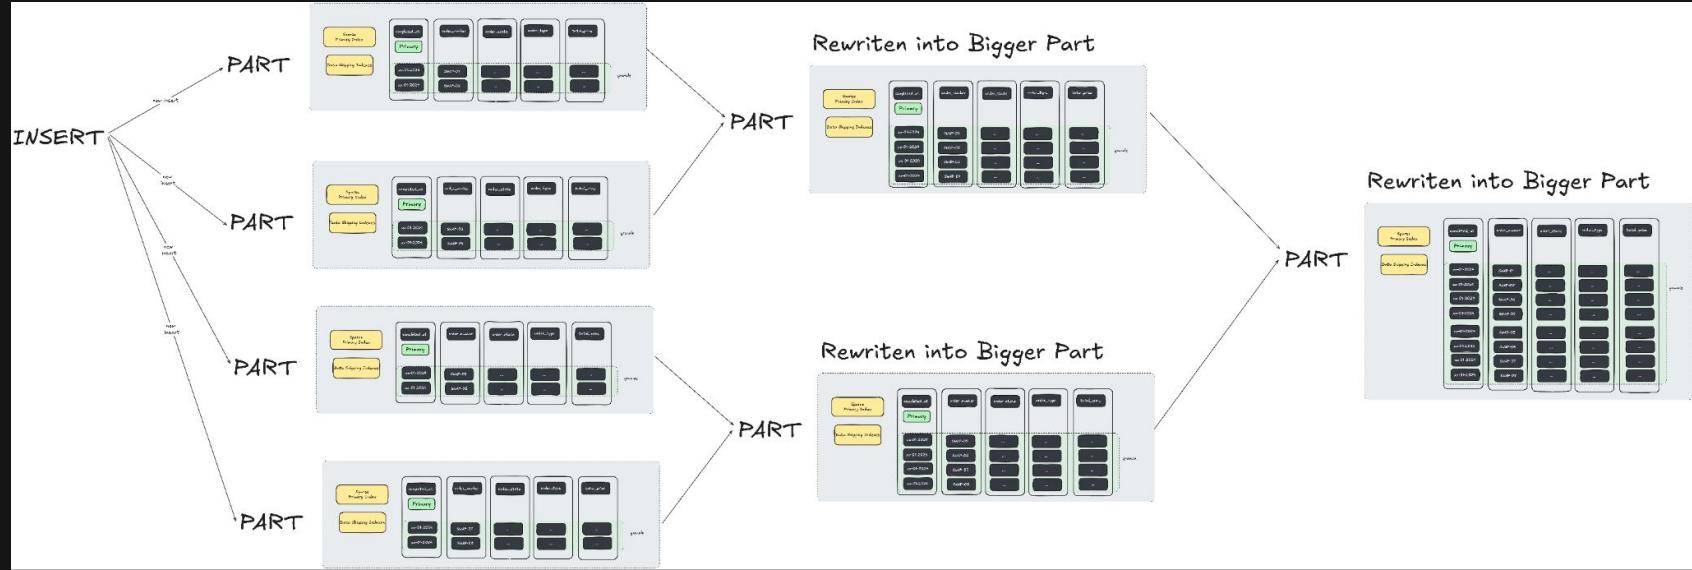
+ reduces CPU cache miss rates

# Multiple Layers of Indexes



- **Granules** are chunks of rows, and **ClickHouse** groups rows into granules based on the index\_granularity setting.
- **Sparse Primary Index:** **ClickHouse** uses the sparse primary index to quickly jump to the relevant granules, skipping over large portions of the data that don't need to be read.
- **Skip Indexes:** These allow **ClickHouse** to skip entire granules if it can determine from the index that no rows in the granule satisfy the query's filter.

# Efficient Merging

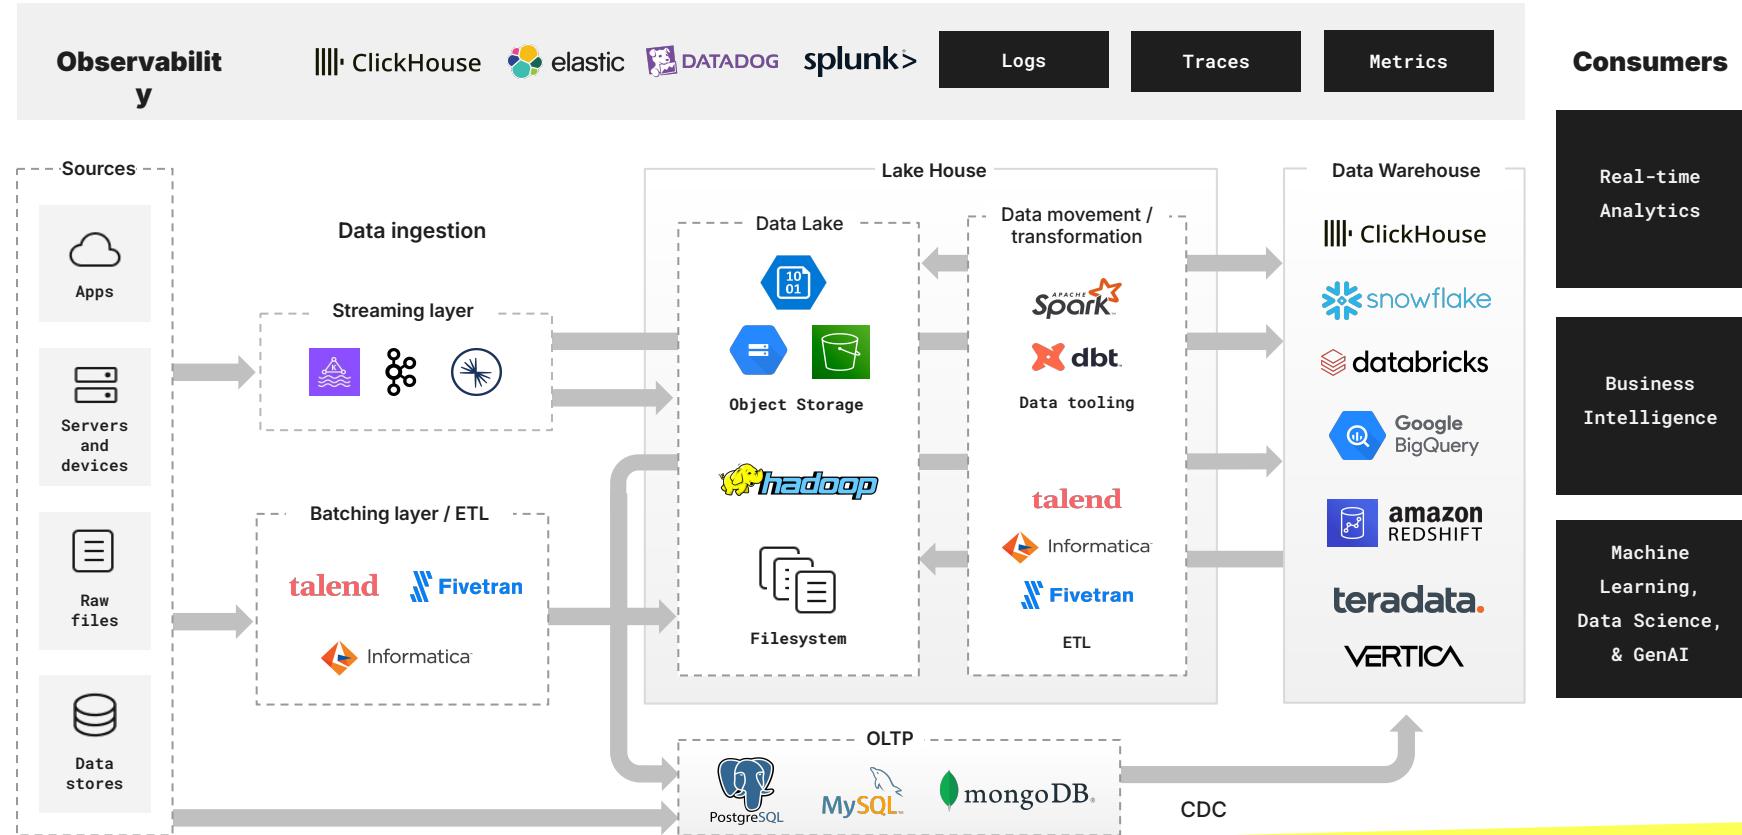


**Parts:** Each time data is inserted, a new part is created. Multiple parts can accumulate over time, which can slow down query performance.

To improve performance and reduce the number of parts, **ClickHouse** periodically runs merging operations.

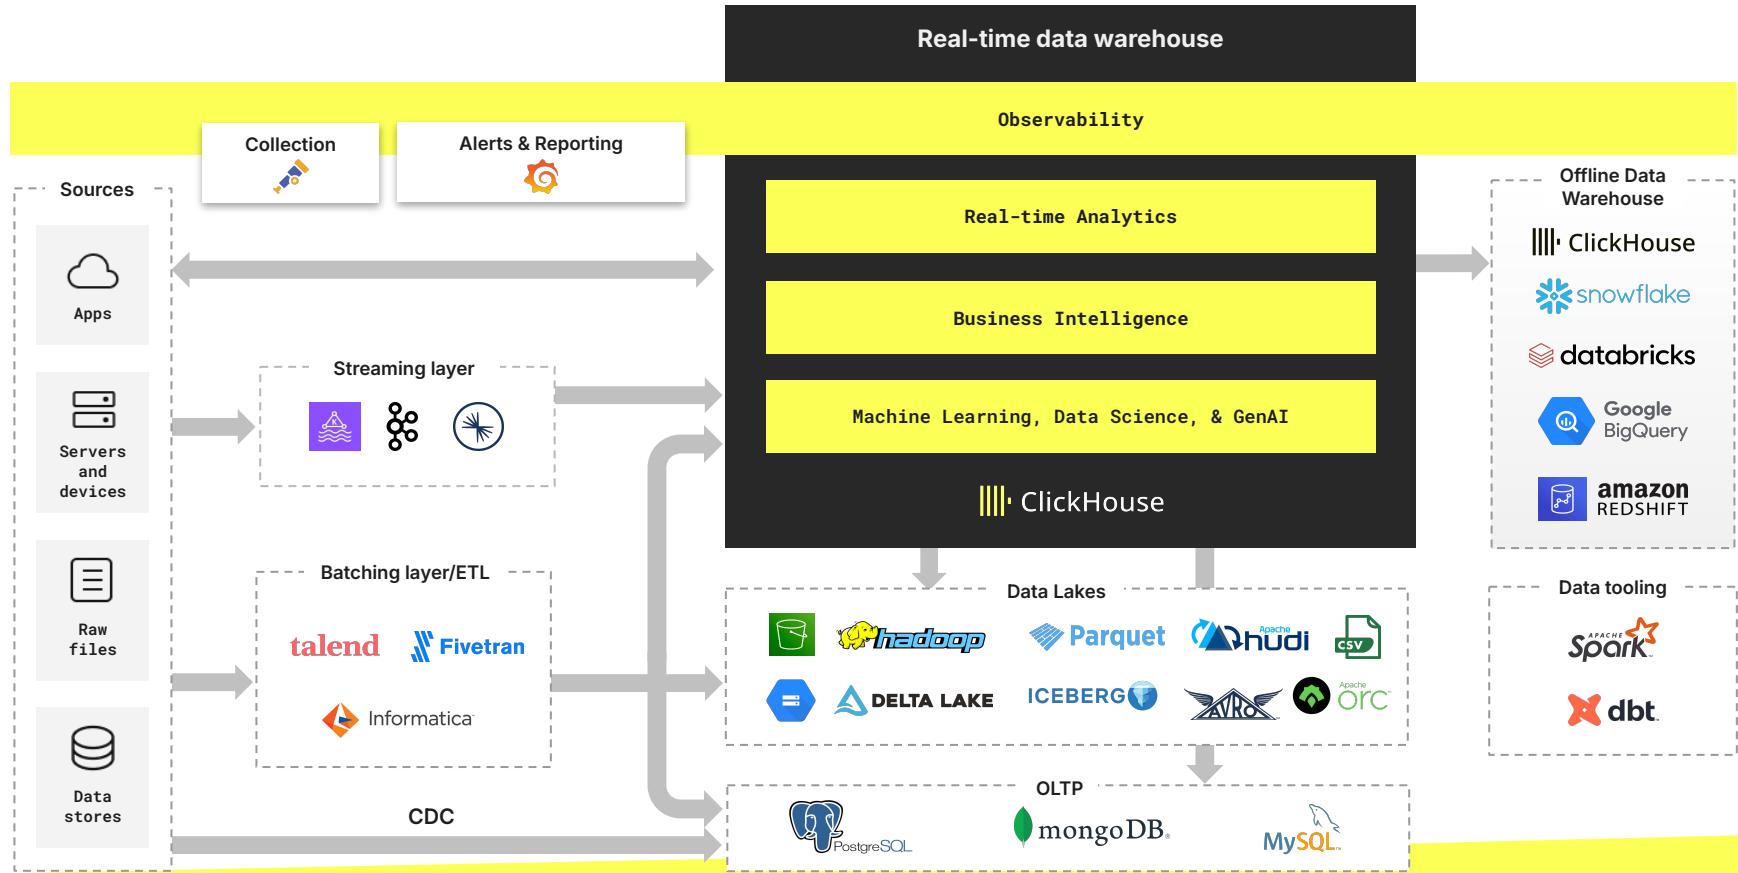
# Traditional Data Architecture

Operationally complex with skyrocketing expense



# Modern Data Architecture

Reduces system bloat and increases resource efficiency

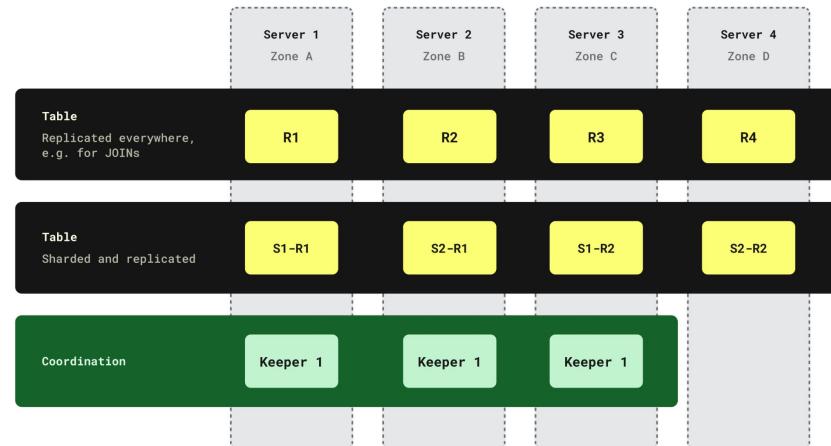


# ||| ClickHouse

## Self-managed

- ✓ Open-source
- ✓ Flexible architecture
- ✓ Efficient and robust
- ✓ Support contracts available

Sample self-managed architecture

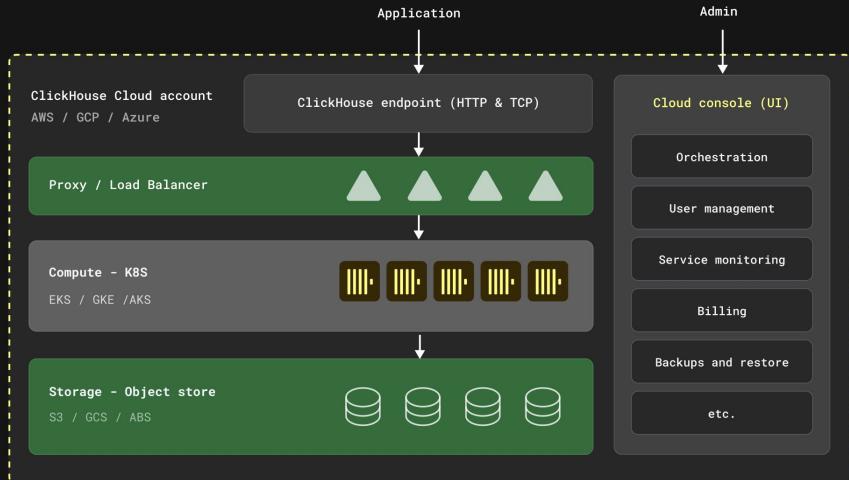


# ||| ClickHouse

## Cloud

- ✓ Easy to use
  - ✓ Feature-rich
  - ✓ Fast
  - ✓ Scalable
  - ✓ Reliable
  - ✓ PAYG
- Managed for you  
Cloud-first features & tooling  
Automatically maximizes performance/efficiency  
Scale seamlessly  
Ensure reliability  
SaaS usage and capacity based pricing

ClickHouse Cloud architecture





# Introduction to Confluent



# Introducing Apache Kafka®



# The Rise of Event Streaming





# Confluent Overview



Jay Kreps  
CEO

Neha  
Narkhede

Jun Rao  
VP  
Engineering



Founded by the original  
creators of **Apache Kafka**

Founded  
September 2014

Technology developed  
while at **LinkedIn**

BENCHMARK

Index  
Ventures

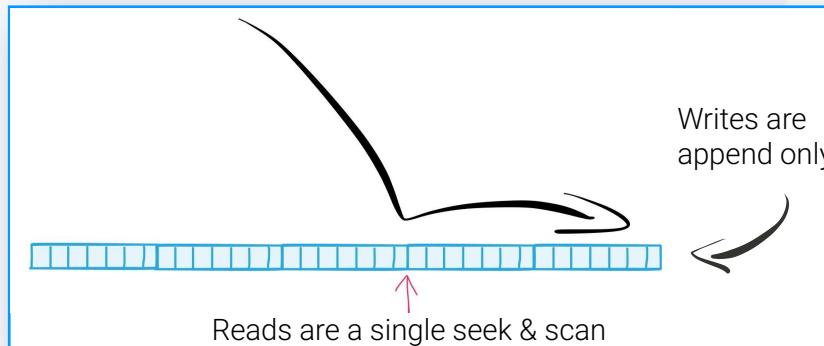
SEQUOIA



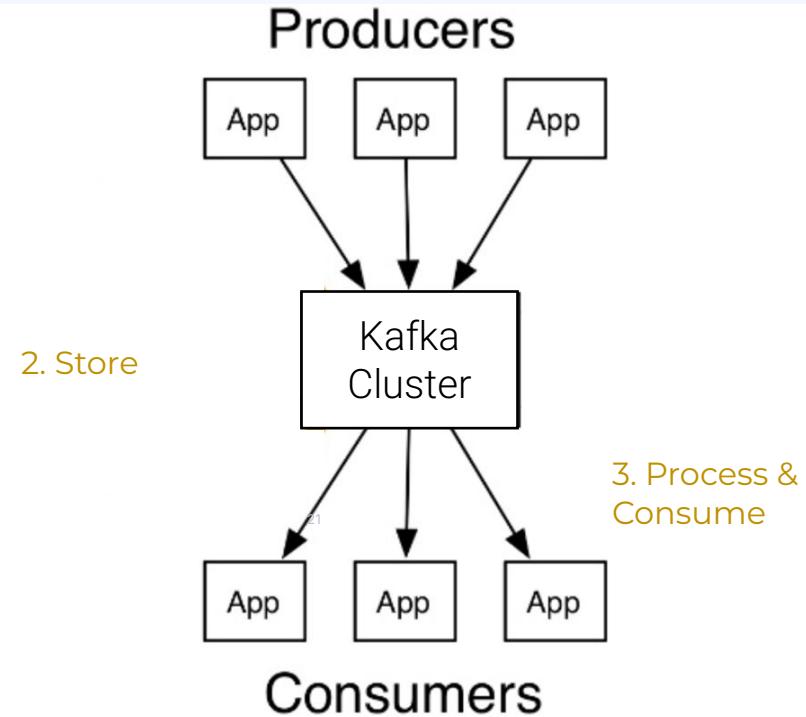
# Apache Kafka - Publish & Subscribe Reimagined

## Kafka

A **Distributed Commit Log**. Publish and subscribe to streams of records. Highly scalable, high throughput. Supports transactions. Persisted data.



### 1. Publish Stream Events



# Topics

Clicks



Orders



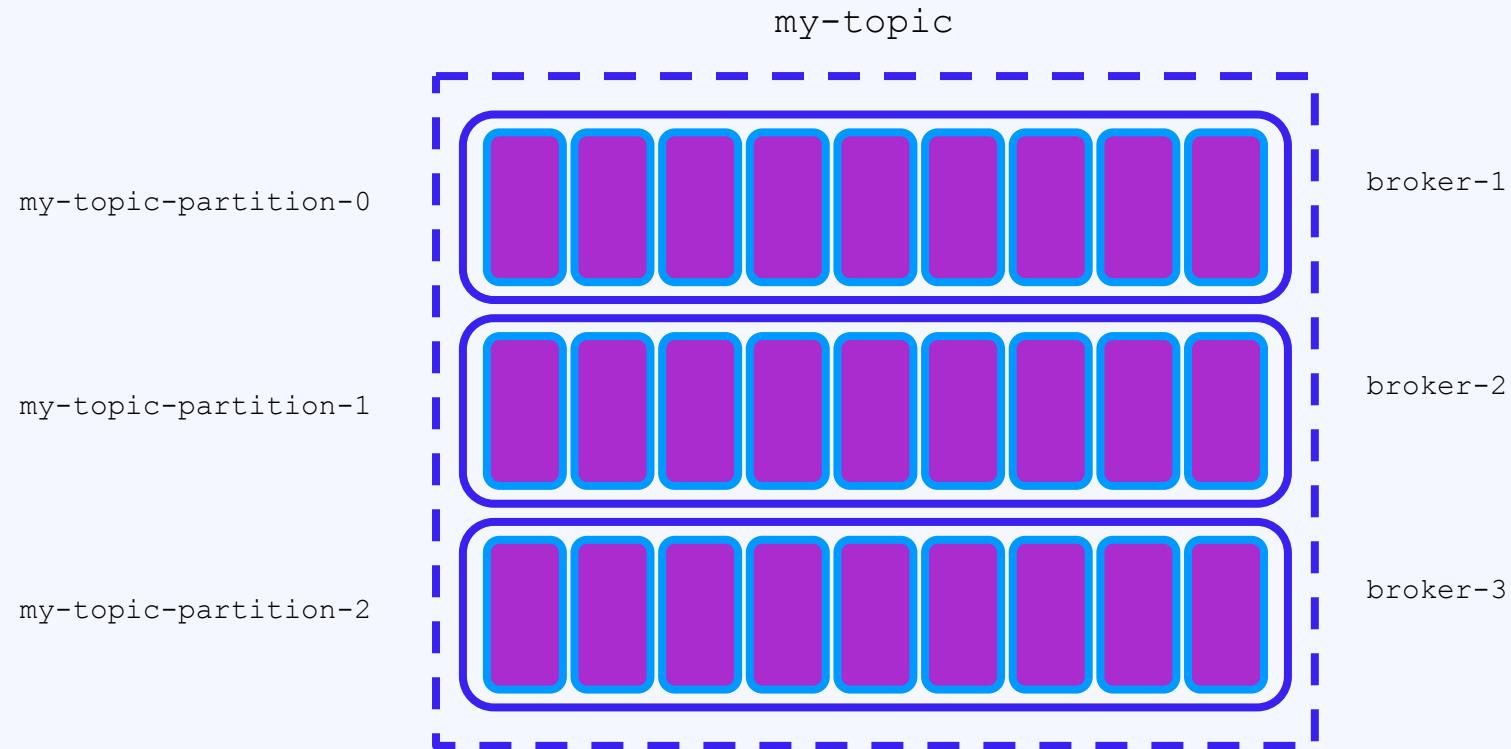
Customers



Topics are similar in concept  
to tables in a database

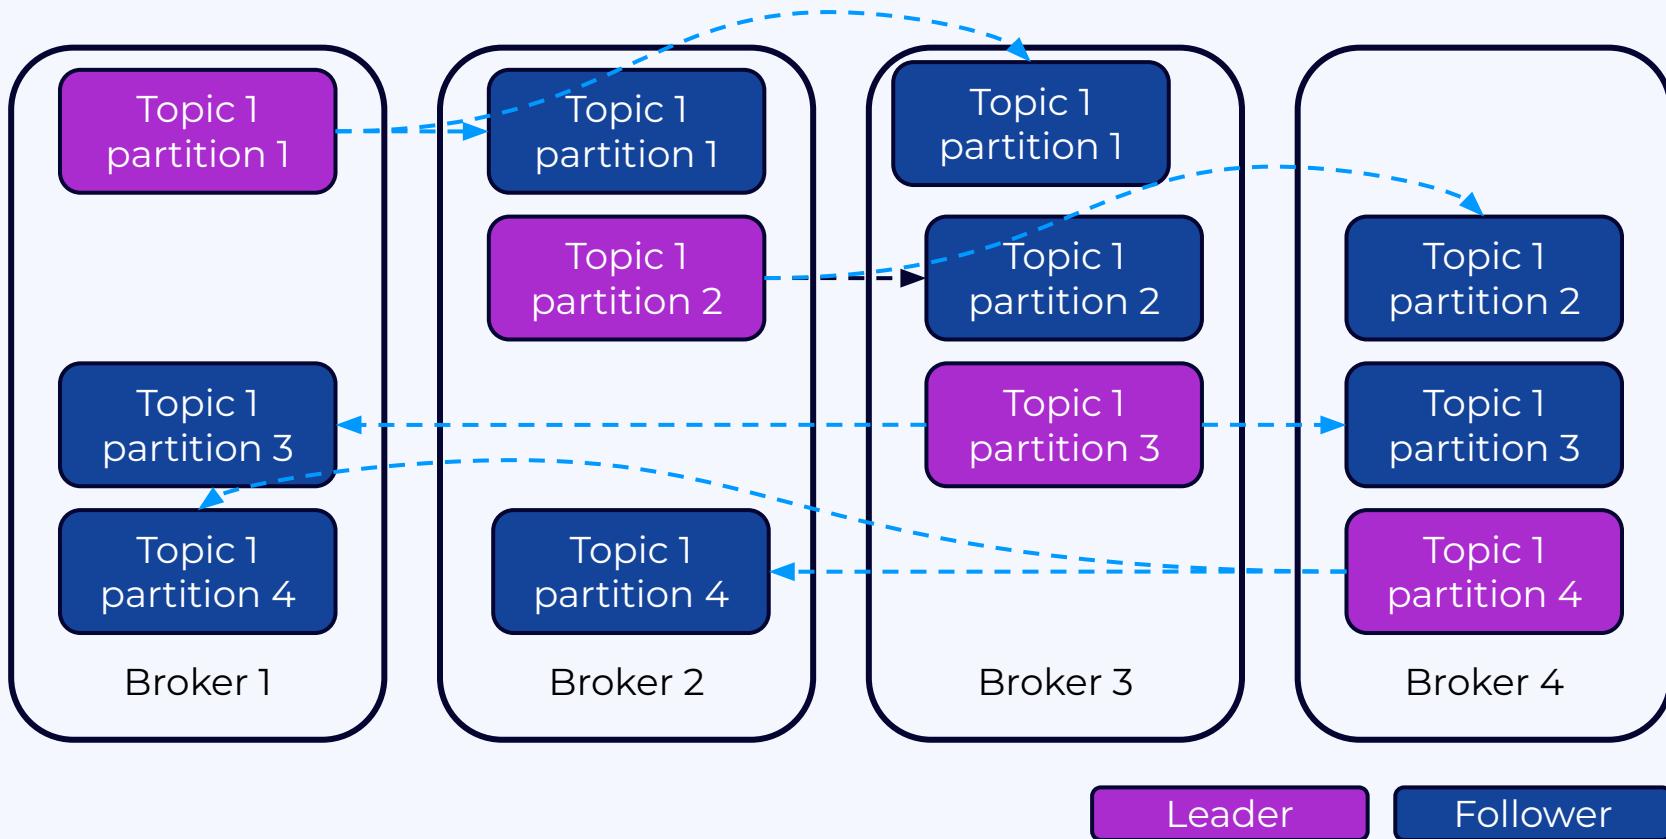


# Kafka Topics





# How Does Kafka Provide Resilience?

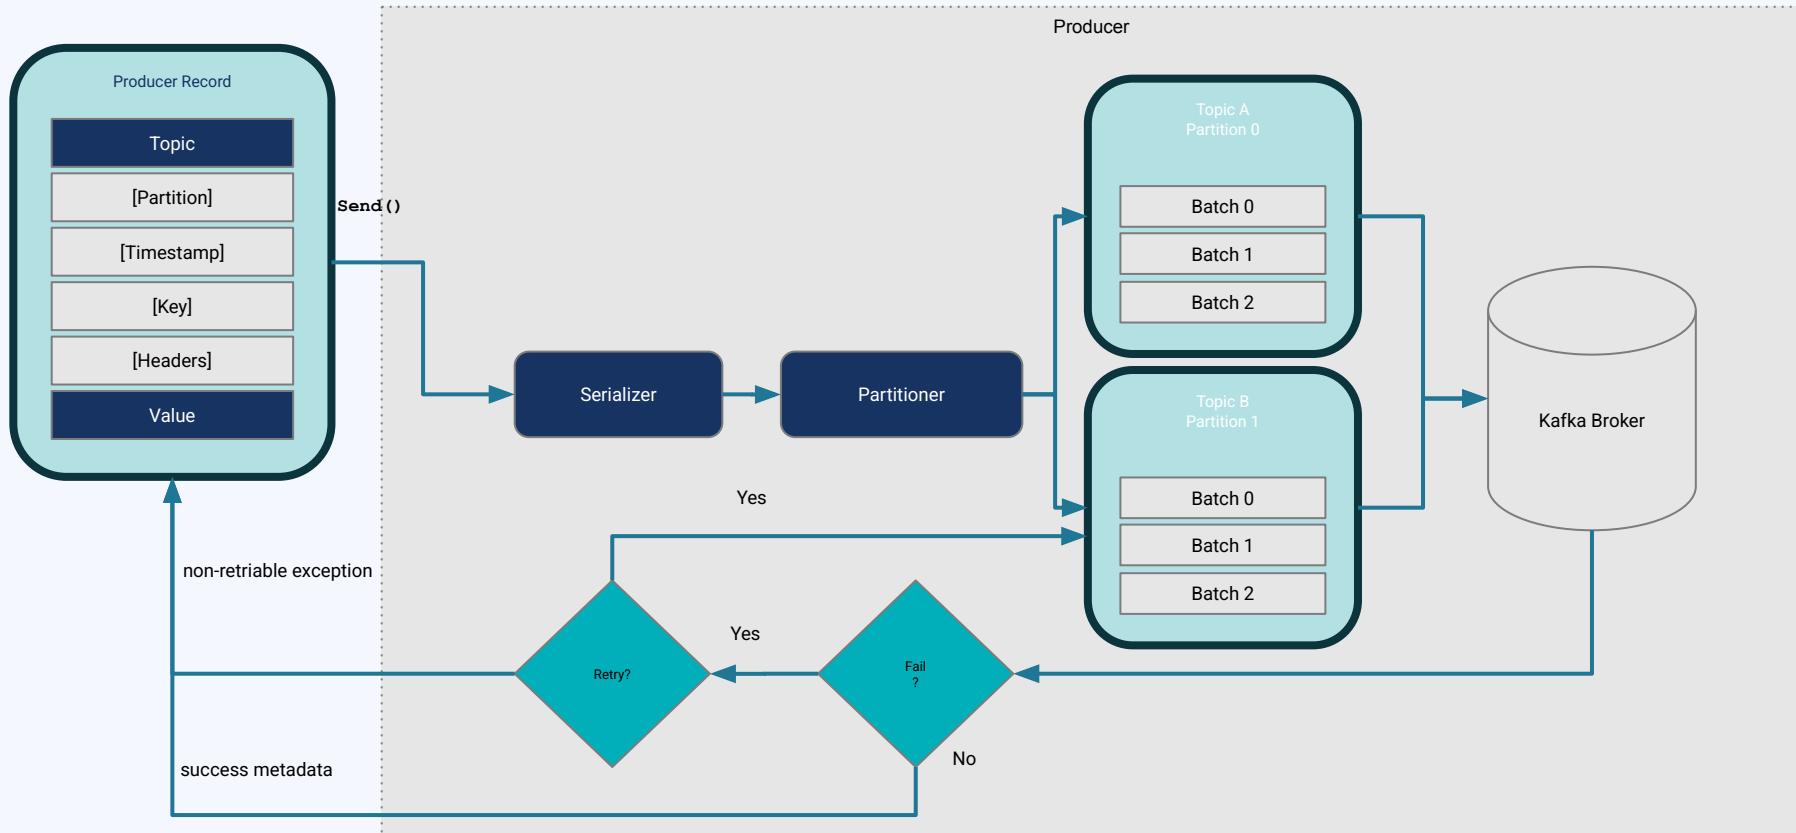




# Producing to Kafka



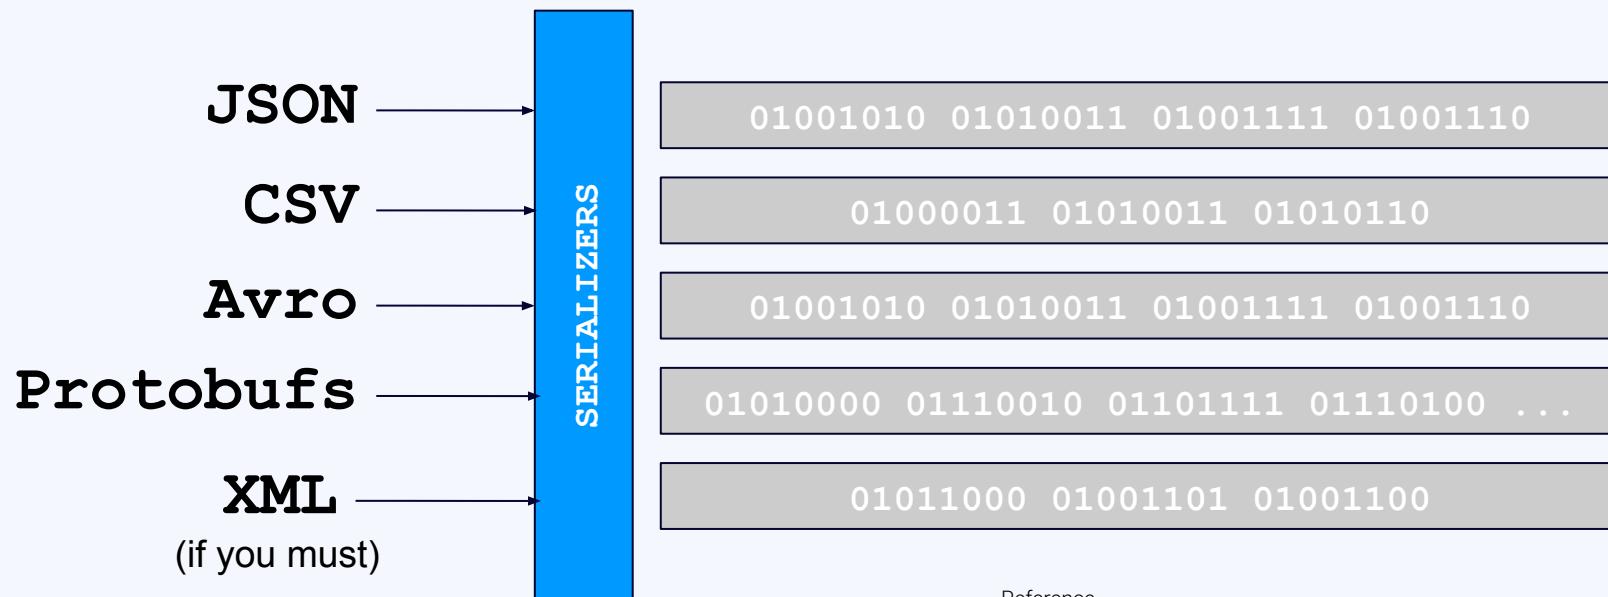
# What happens inside a producer?





# The Serializer

Kafka doesn't care about what you send to it as long as it's been converted to a byte stream beforehand.



Reference

<https://kafka.apache.org/10/documentationstreams/developer-guide/datatypes.html>



# The Serializer

```
private Properties kafkaProps = new Properties();

kafkaProps.put("bootstrap.servers", "broker1:9092,broker2:9092");
kafkaProps.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
kafkaProps.put("value.serializer", "io.confluent.kafka.serializers.KafkaAvroSerializer");
kafkaProps.put("schema.registry.url", "https://schema-registry:8083");

producer = new KafkaProducer<String, SpecificRecord>(kafkaProps);
```

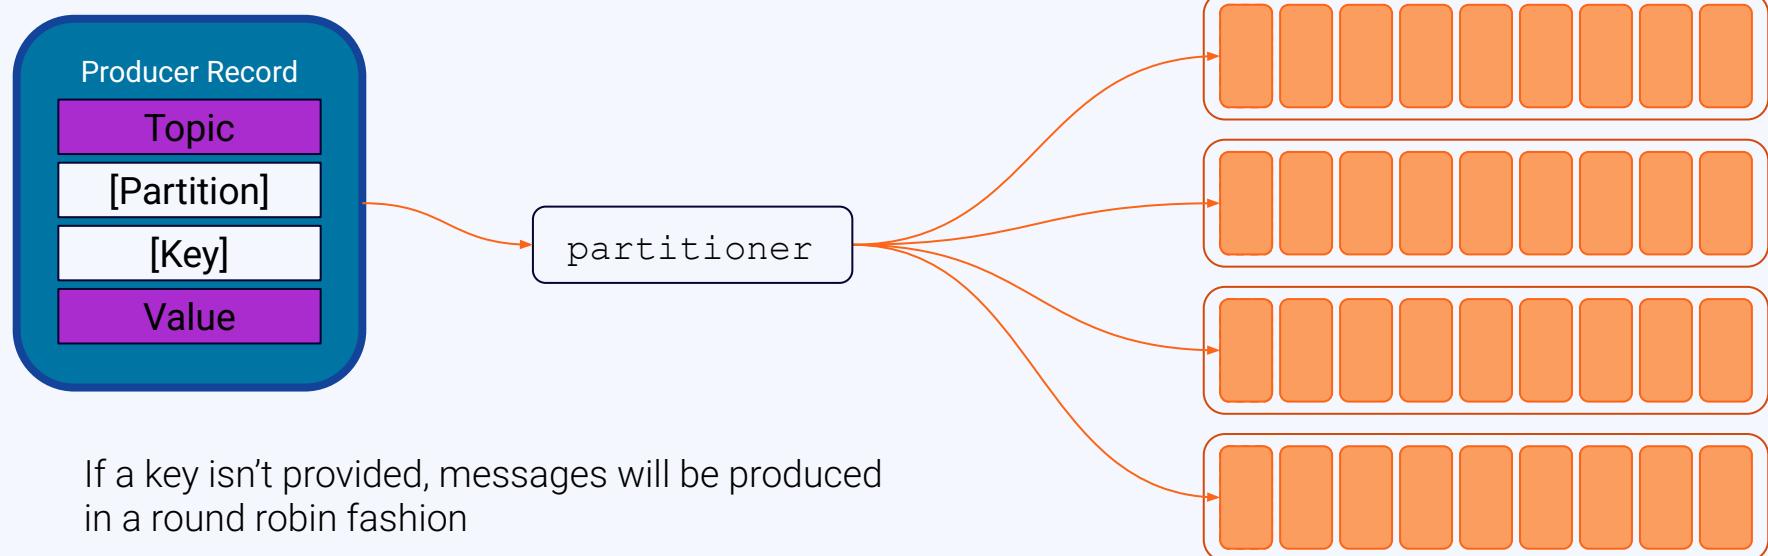
Reference

<https://kafka.apache.org/10/documentationstreams/developer-guide/datatypes.html>



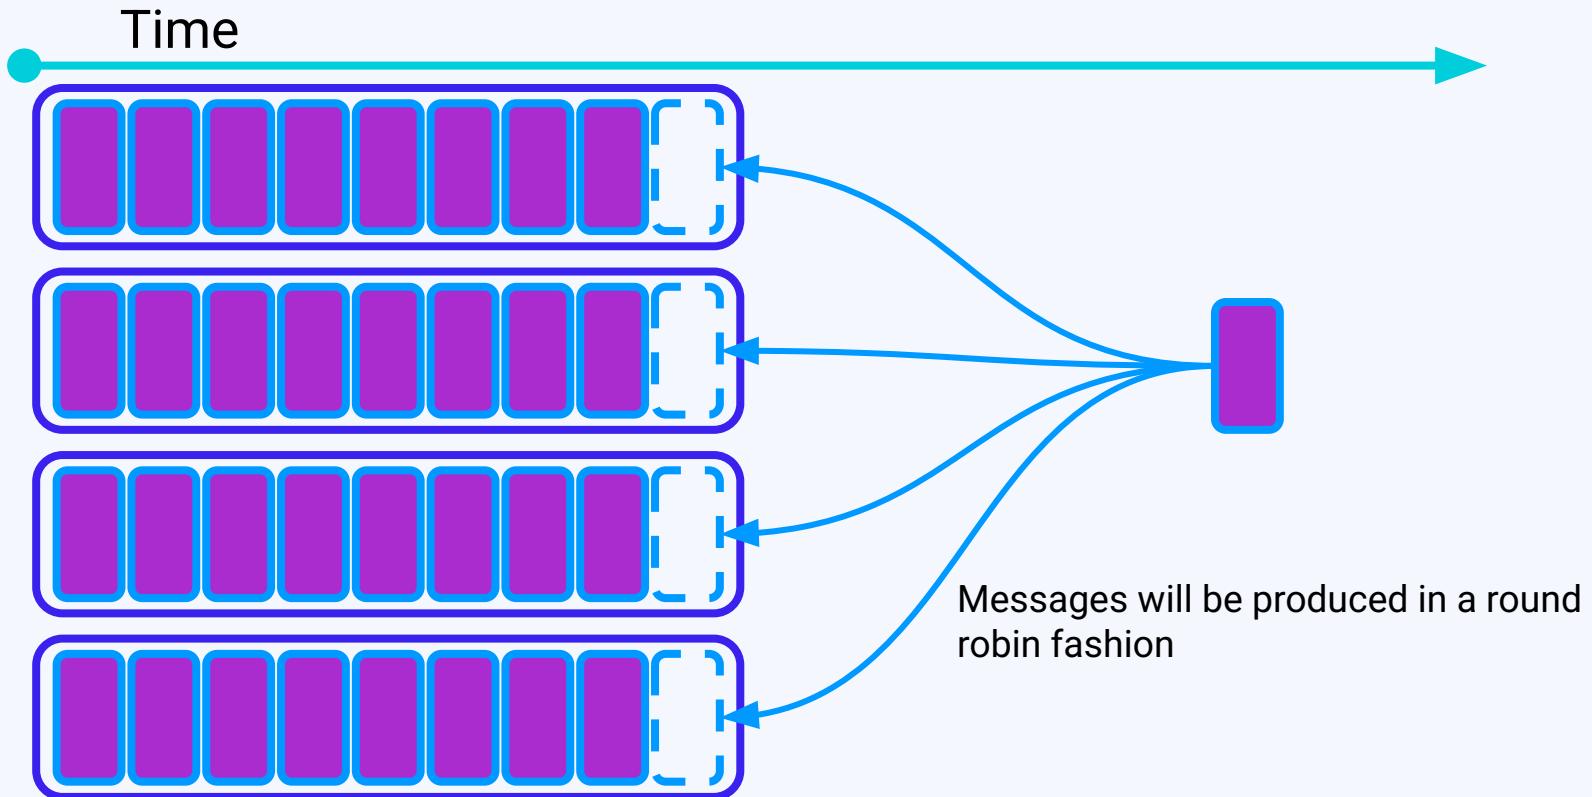
# Record Keys and why they're important - Ordering

Record keys determine the partition with the default kafka partitioner



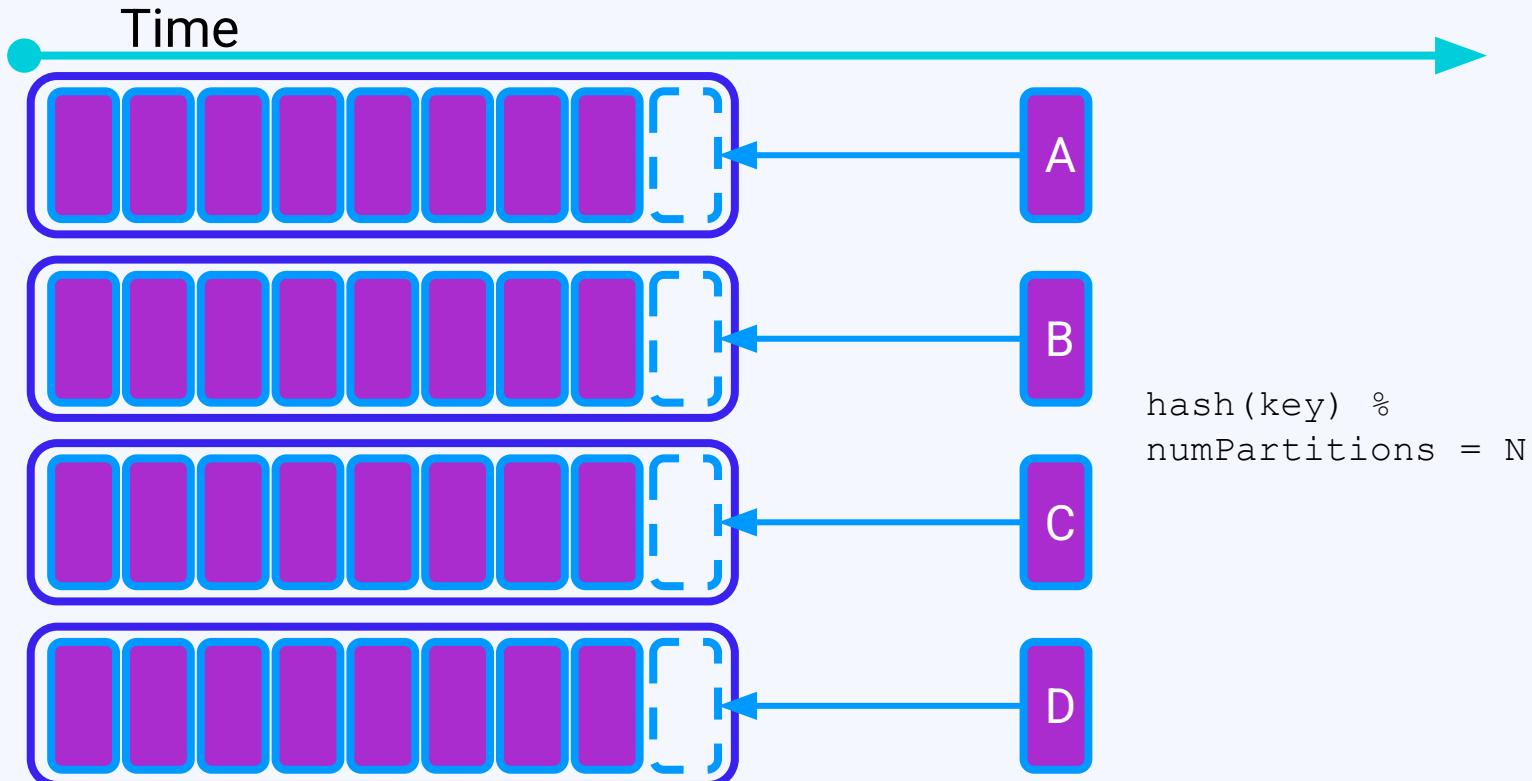


# Producing to Kafka - No Key





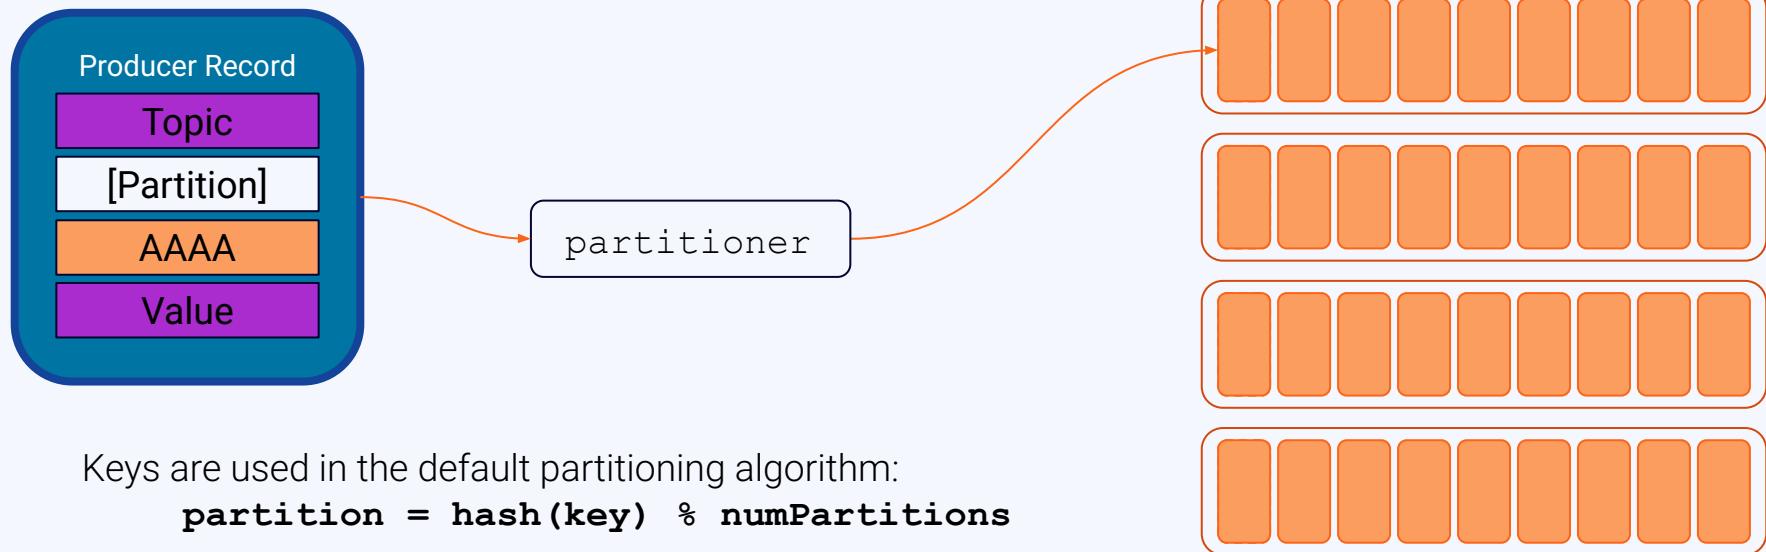
# Producing to Kafka - With Key





# Record Keys and why they're important - Ordering

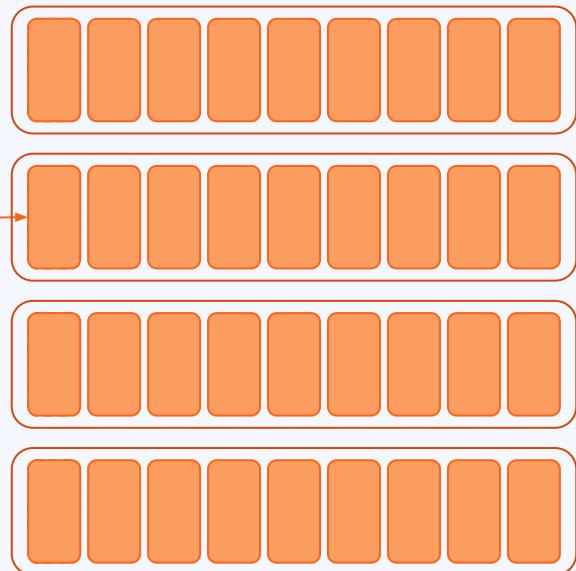
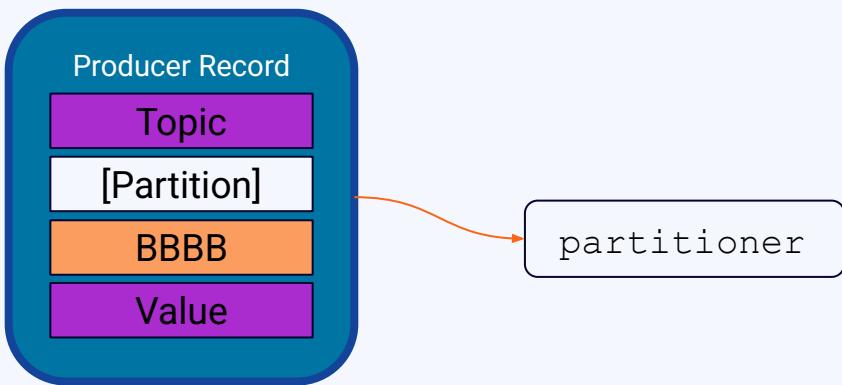
Record keys determine the partition with the default kafka partitioner, and therefore guarantee order for a key



# Record Keys and why they're important - Ordering



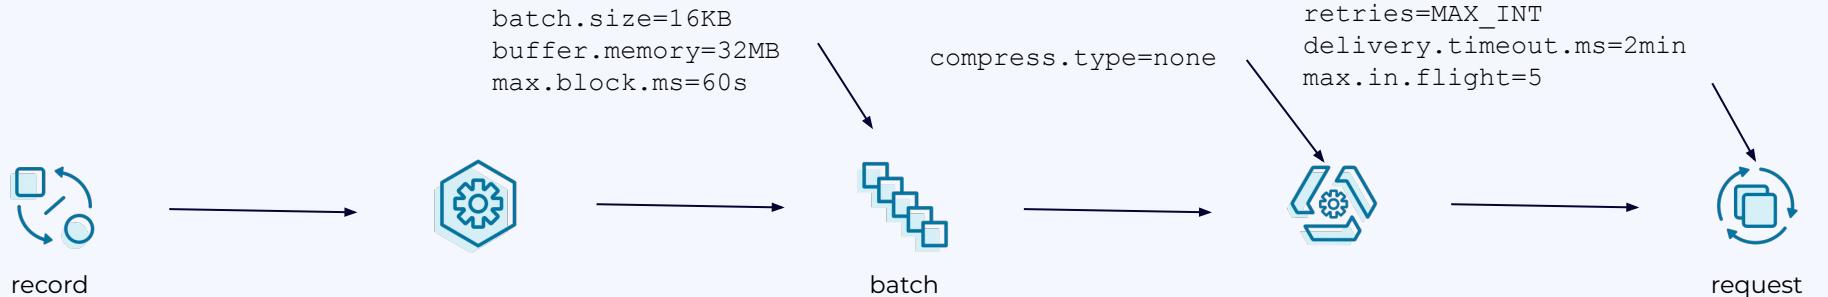
Record keys determine the partition with the default kafka partitioner, and therefore guarantee order for a key



Keys are used in the default partitioning algorithm:

```
partition = hash(key) % numPartitions
```

# Producer Summary



**Serializer**

- Retrieves and caches schemas from Schema Registry

**Partitioner**

- Java client uses murmur2 for hashing
- If key not provided performs round robin
- If keys unbalanced it will overload one leader

**Record accumulator**

- Buffer per partition, seldom used partitions may not achieve high batching due to this
- If many producers are in the same JVM, memory and GC could become important
- Sticky partitioner could be used to increase batches in the case of round robin (KIP-408)

**Compression**

- Allows faster transfer to the producer
- Reduces the inter broker replication load
- Uses less page cache and disk space at broker
- Compress the batch not the record
- Gzip is more CPU intensive, Snappy is lighter, LZ4/ZStd are a good balance\*

**Sender thread**

- Batches grouped in a producer request and sent by broker
- Multiple batches to different partitions could be in the same producer request



# Consuming from Kafka

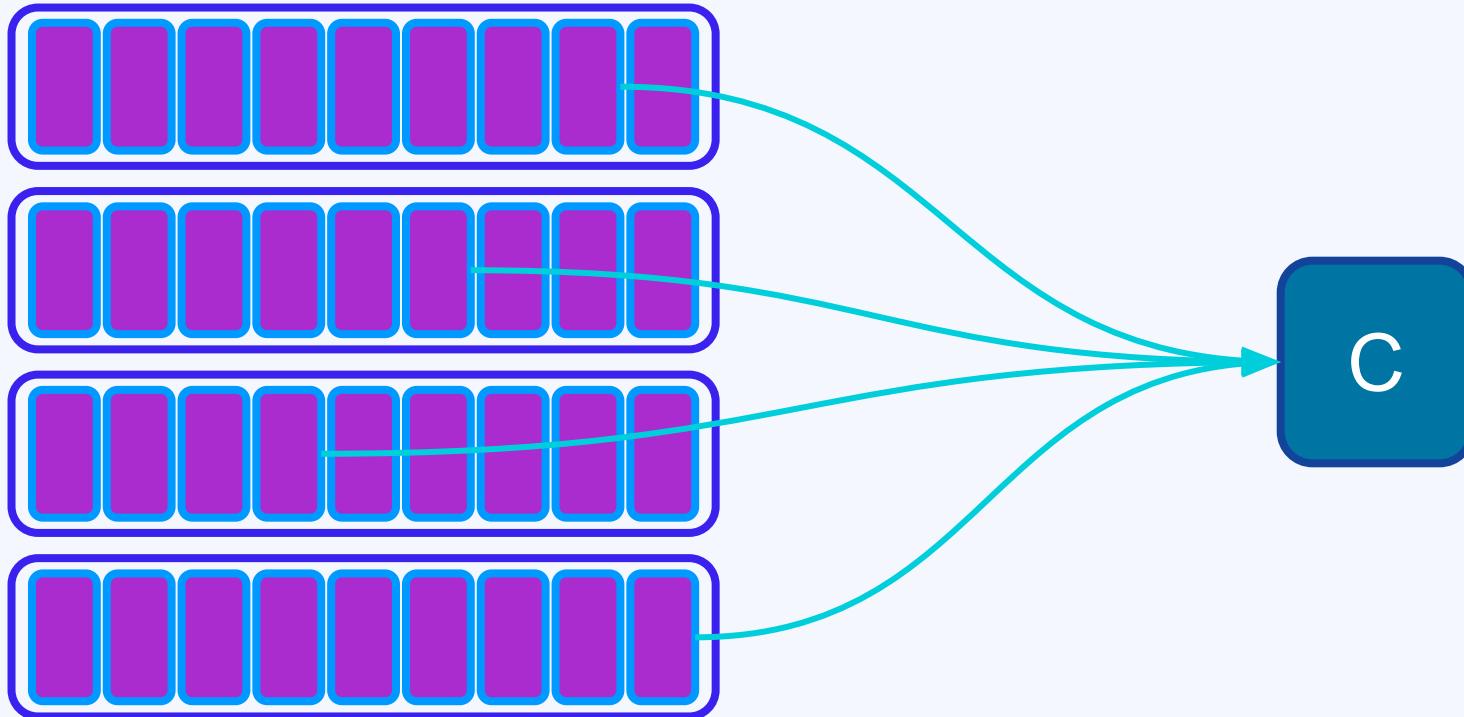


# A basic Java Consumer

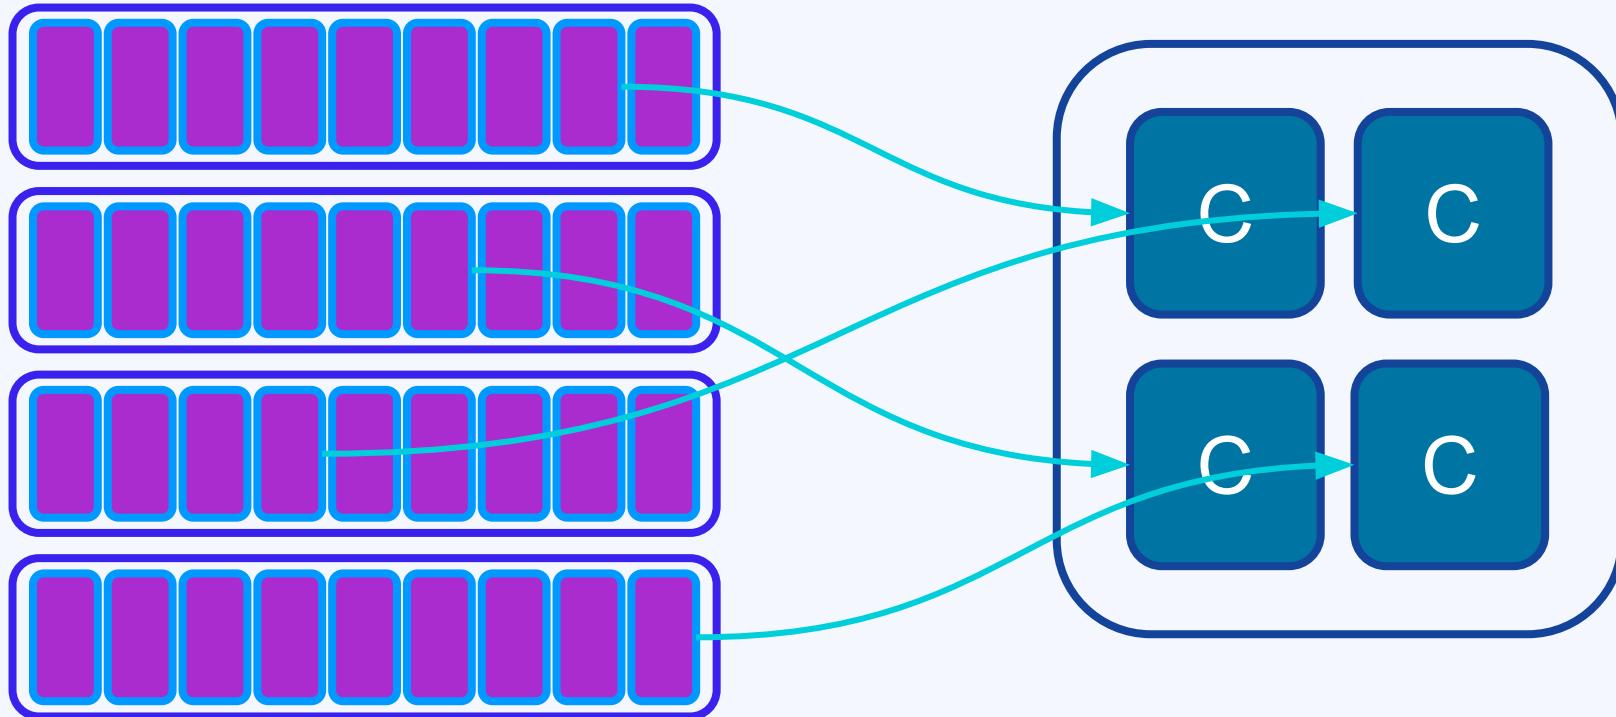
```
final Consumer<String, String> consumer = new KafkaConsumer<String, String>(props);
consumer.subscribe(Arrays.asList(topic));
try {
    while (true) {
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records) {
            -- Do Some Work --
        }
    }
} finally {
    consumer.close();
}
```



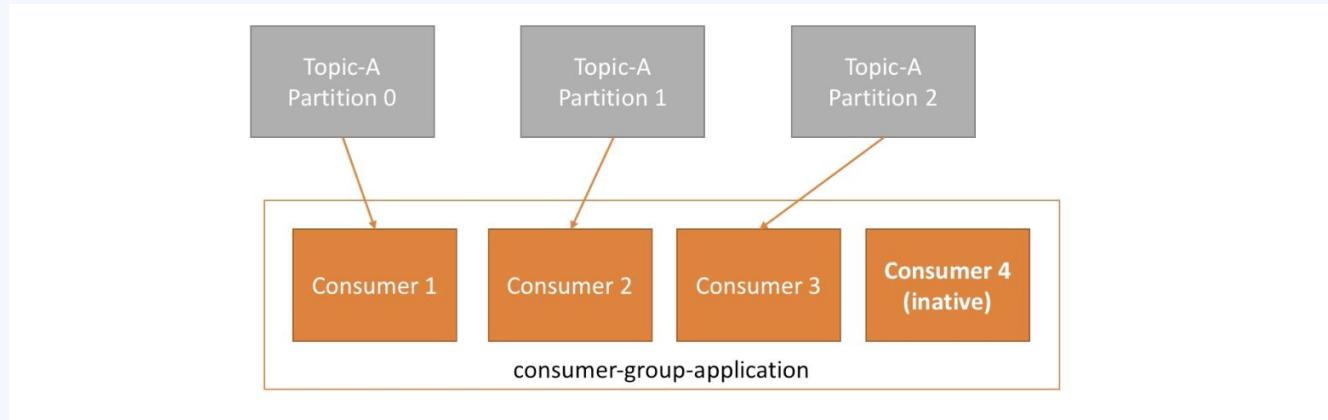
# Consuming From Kafka - Single Consumer



# Consuming From Kafka - Grouped Consumers

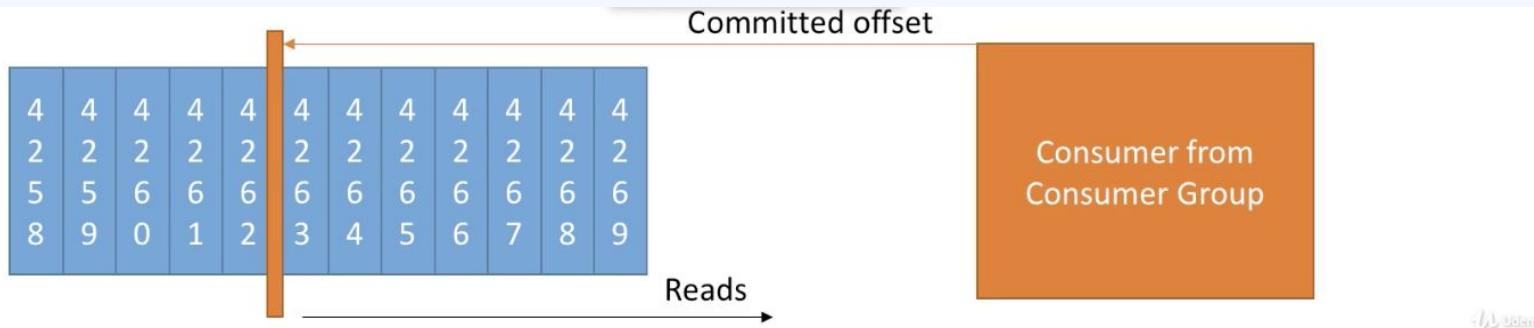


# What if you have more Consumers than partitions?



# Consumer Offsets

- Kafka stores the offsets at which a consumer group has been reading
- Offsets are committed in an internal **topic** called `_consumer_offsets`
- When a consumer in a group has processed data received from Kafka it should be committing the offsets
- If a consumer dies Kafka will be able to tell it where it was so it can start reading again





# Kafka Connect

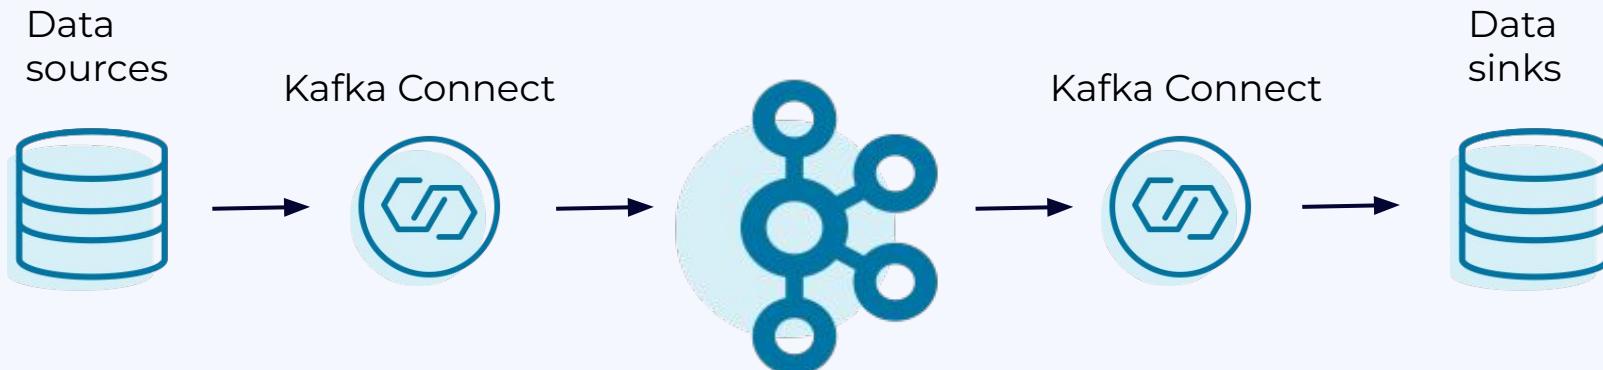
No/Low Code connectivity to many systems



# Kafka Connect

No-Code way of connecting known systems (databases, object storage, queues, etc) to Apache Kafka

Some code can be written to do custom transforms and data conversions though maybe out of the box Single Message Transforms and Converters exist

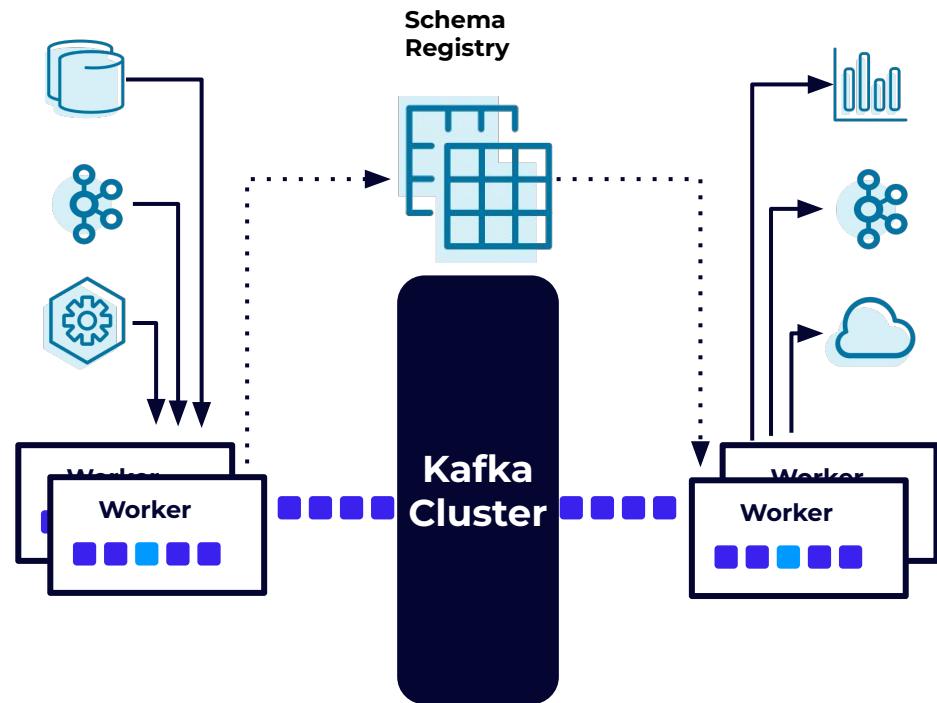


# Kafka Connect Durable Data Pipelines

Integrate upstream and downstream systems with Apache Kafka®

- **Capture** schema from sources, use schema to inform data sinks
- **Highly Available** workers ensure data pipelines aren't interrupted
- **Extensible** framework API for building custom connectors

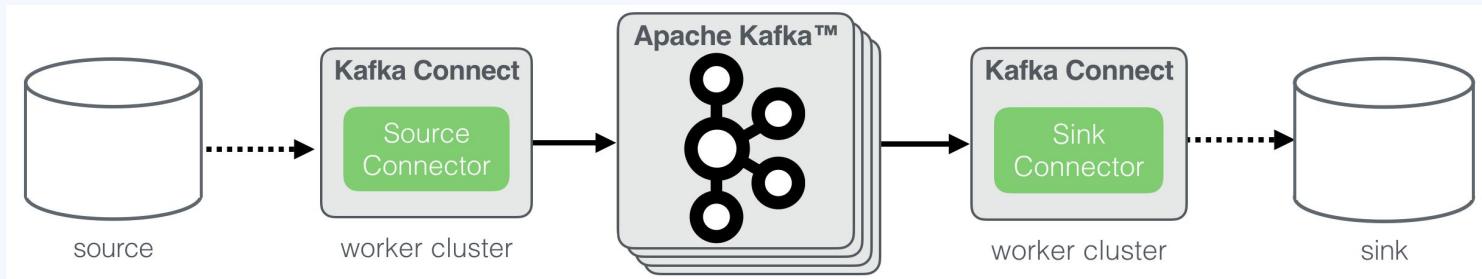
## Kafka Connect



# Kafka Connect



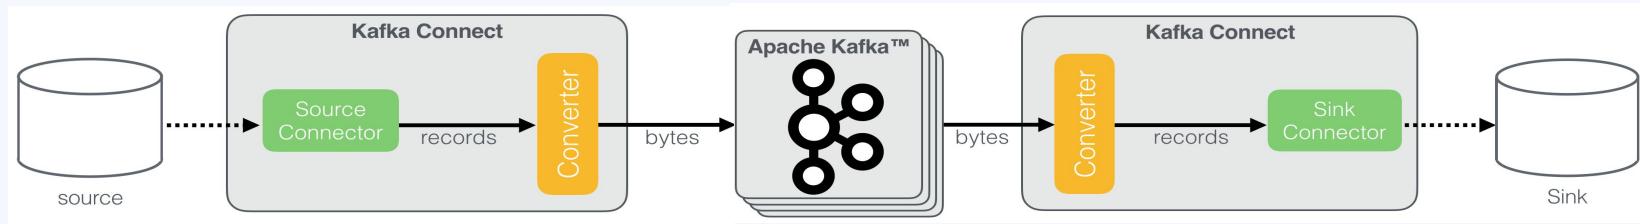
**Connectors** are reusable components that know how to talk to specific **sources and sinks**.



# Kafka Connect



Convert between the **source and sink record objects** and the **binary format** used to persist them in Kafka.

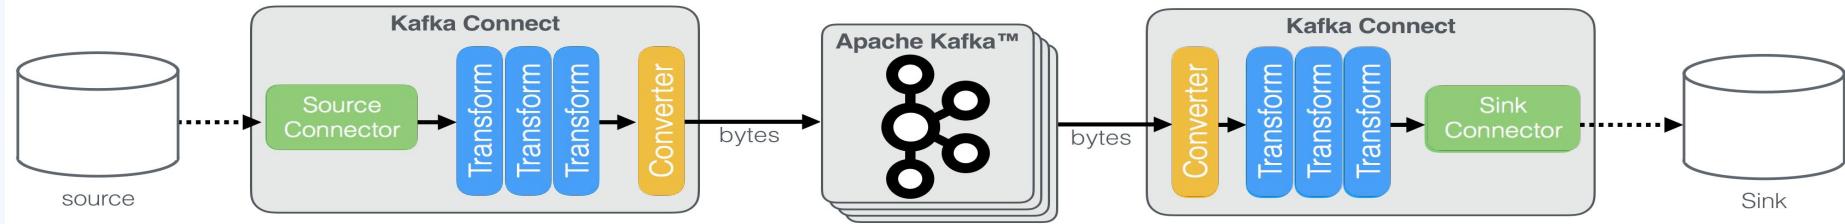


JSON, Avro, and others

# Kafka Connect



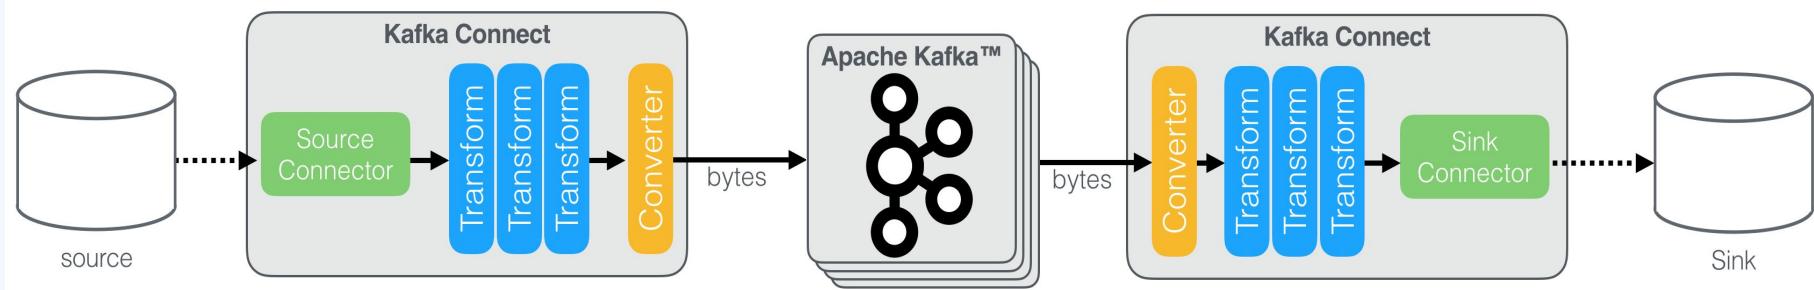
Connectors, transforms, and converters  
are **pluggable components**.



# Kafka Connect



**Modify the structure of keys and values, topics, and partition of source and sink record objects.**





# Your Apache Kafka journey begins here

[developer.confluent.io](https://developer.confluent.io)



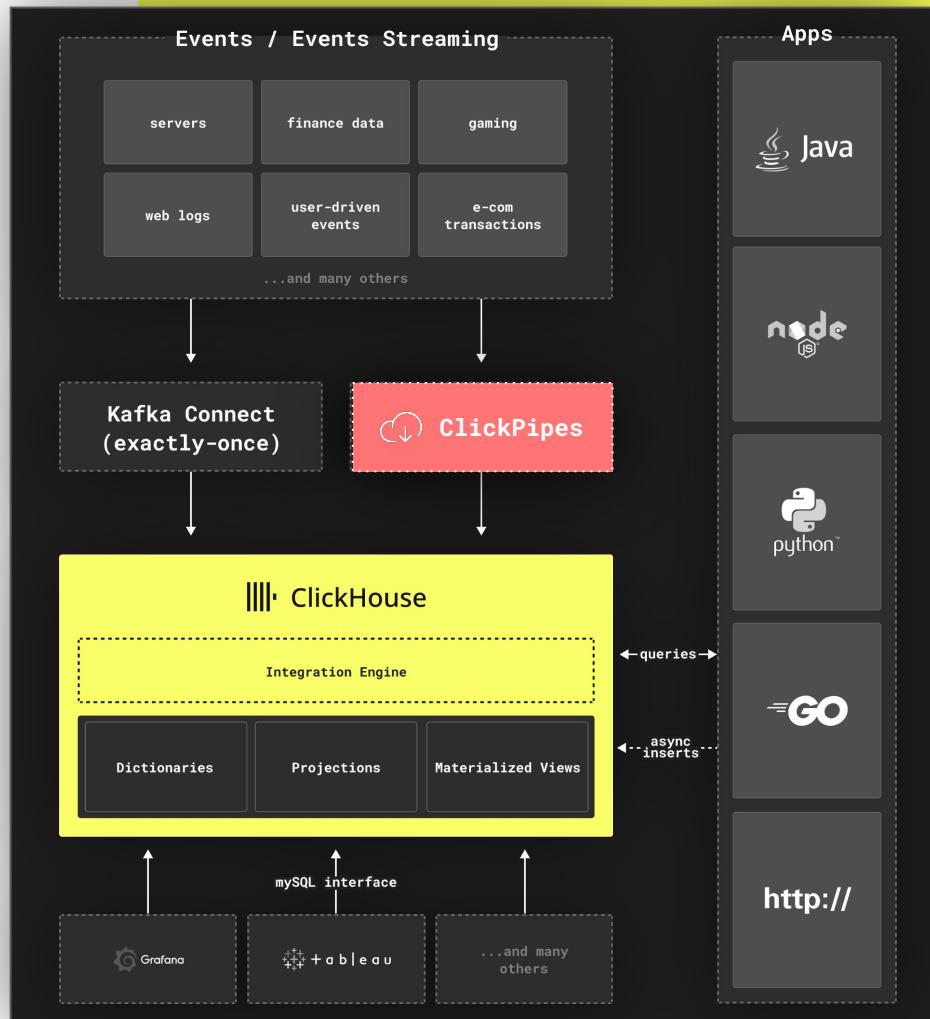
# Integrating Kafka with ClickHouse

$(\text{real-time})^2$



**ClickPipes** *noun*  
*/klik paips/*

**Cloud-Native experience**  
to ingest data from  
remote data sources to  
**ClickHouse Cloud**



# Existing Kafka Integration Options

## 01 Kafka Table Engine

```
CREATE TABLE queue (
    timestamp UInt64,
    level String,
    message String
)
ENGINE = Kafka('localhost:9092', 'topic', 'group1', 'JSONEachRow');
```

## 02 ClickHouse Kafka Connect Sink

The screenshot shows the ClickHouse Kafka Connect Sink interface. It includes sections for the JDBC Connector (Source and Sink) and the HTTP Sink Connector. Both sections provide download links for Java JDBC Version 10.7.4 and Confluent Cloud versions 1.7.4 and 0.0.17 respectively. The JDBC connector also indicates it is available fully managed on Confluent Cloud. The HTTP sink connector indicates it is available fully managed on Confluent Cloud. The ClickHouse Kafka Connect Sink section describes it as the official Kafka Connect Sink connector for ClickHouse, provides installation instructions using the Confluent Hub CLI, and includes a download link for the ZIP file.

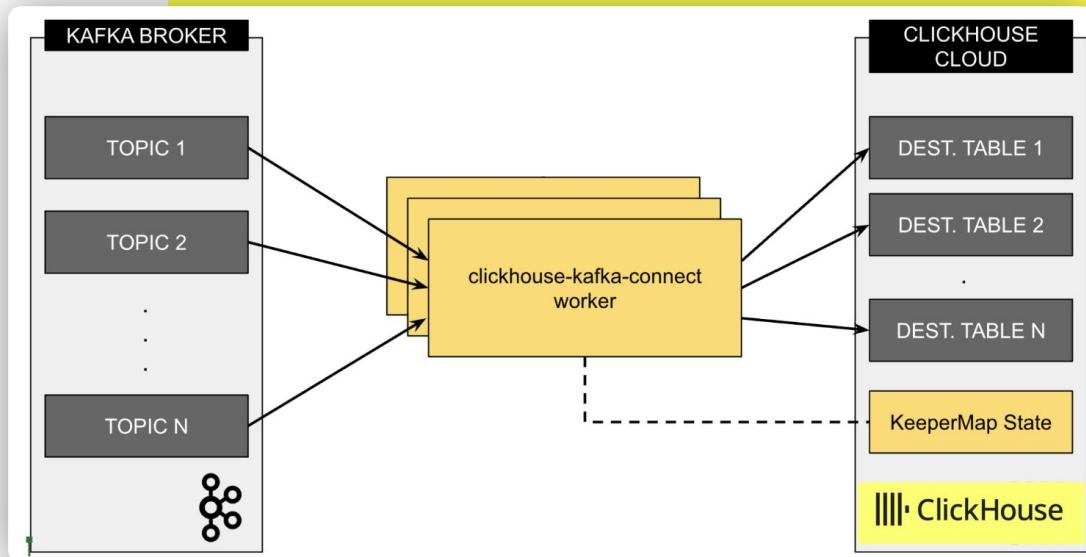
## 03 ClickPipes

The screenshot shows the ClickPipes integration engine interface. It features a grid of icons representing various data sources: Apache Kafka, Confluent Cloud, Amazon MSK, Amazon Kinesis, Redpanda, WarpStream, Azure Event Hubs, Amazon S3, Google Cloud Storage, and Postgres CDC. A prominent yellow bar chart icon is located in the top right corner. The interface includes sections for "Cloud / Data ingestion" and "ClickPipes", and a banner for the "Blazing-fast Postgres to ClickHouse CDC with our new ClickPipe connector — now in Private Preview".

The screenshot shows the ClickPipes setup wizard interface. It consists of a series of numbered steps: 1. Select the data source (Apache Kafka, Confluent Cloud, Amazon MSK, Amazon Kinesis, Redpanda, WarpStream, Azure Event Hubs, Amazon S3, Google Cloud Storage, Postgres CDC), 2. Setup your ClickPipe Connection, 3. Incoming Data, 4. Parse Information, and 5. Details and Settings. The interface also includes a "Data sources / ClickPipes / New" header and a "View documentation" button.

# ClickHouse Kafka Connect Sink

- Shipped with out-of-the-box **exactly-once semantics** powered by ClickHouse core feature named **KeeperMap** (table engine used as a state store by the connector) and allows for minimalistic architecture
- **Core integration:** Built, maintained, and **supported by ClickHouse**
- Tested **continuously** against ClickHouse Cloud.
- Data inserts with a **declared schema** (schema-based data, e.g. Avro, Protobuf, etc.) and **schemaless** (e.g. JSON)
- Support for all data types of ClickHouse



# Demo



# KEEP IN TOUCH!



[ClickHouse Sydney](#)  
[Meetup Group](#)



[clickhouse.com/slack](#)



[@clickhousedb](#)



[@clickhouseinc](#) #clickhouseDB



[clickhouse](#)

||||· ClickHouse

**MIGRATION FROM BIGQUERY TO  
CLICKHOUSE VIRTUAL WORKSHOP**  
July 9, 2025 - 11:30 AM SGT





# Questions