


Optimizing Time-Based Segment Evaluation with **ClickHouse**.



An overview of
segmentation.

01

An overview of segmentation.

User Profiles

klaviyo

← Profiles

Account plans

Support

Home

Campaigns

Flows

Sign-up forms

Audience

List & segments

Profiles

Content

Analytics

Conversations

Alice Goldsmith

agoldsmith@gmail.com • +1 (555) 555-5555 • Boston, MA (UTC-6)

Messages

Actions

Details

Metrics & insights

List & segments

Metrics

Last 30 days

All-time

Edit metrics

Highlighted

\$141

Revenue

6

Placed order

\$35

Average order value

Selected metrics

Metric	Value	Δ (last 30d)
Active on site, last 30 days	15	↑ 12%
Fulfilled order, last 30 days	1	
Checkout started, last 30 days	2	
Viewed product, last 30 days	5	
Clicked email, last 30 days	0	↑ 12%
Refunded order, last 30 days	0	

Activity log

All events

All-time

Sent review request by email

57 minutes ago

Shipment Delivered

2 hours ago

Ordered Eucalyptus Oil for \$29.99

4 days ago

Closed Ticket

4 days ago

"Yes the discount applies to our Eucalyptus oils"

Opened Ticket

4 days ago

"Does my VIP discount apply to the..."

Viewed Eucalyptus Oil product page

4 days ago

Clicked SMS

4 days ago

"Enjoy 15% off."

Segments

What someone has done (or not done)

Has

Started subscription

at least once

over all time

Add filter

An overview of segmentation.

What someone has done (or not done) ▼ 🗑️ 📋

Person has

Placed Order ▼

at least once ▼

between ▼

1

and

5

weeks ▼

ago

🔍

⊕ Add condition

👤 0

AND

OR

Properties about someone ▼ 🗑️ 📋

Birthday ▼

in the next ▼

3

days ▼

Type:

Date ▼

⊕ Add condition

👤 0

An overview of segmentation.

01

Initial Segment Creation

Initial population of a segment. When creating or updating a segment.

02

Event-Driven Real Time Updates

Update segment memberships for user profiles as we receive new data.

03

Time-Based Periodic Updates

Update segment memberships for segments with relative time conditions.

An overview of segmentation.

ClickHouse Cluster

- 192 hosts
- Bi-Level sharding
- 8 layers, each with 8 shards
- 3 replicas per shard

segment_defs_local

Segment definition data.

events_local

A time-series log of events,
partitioned by month.

traits_local

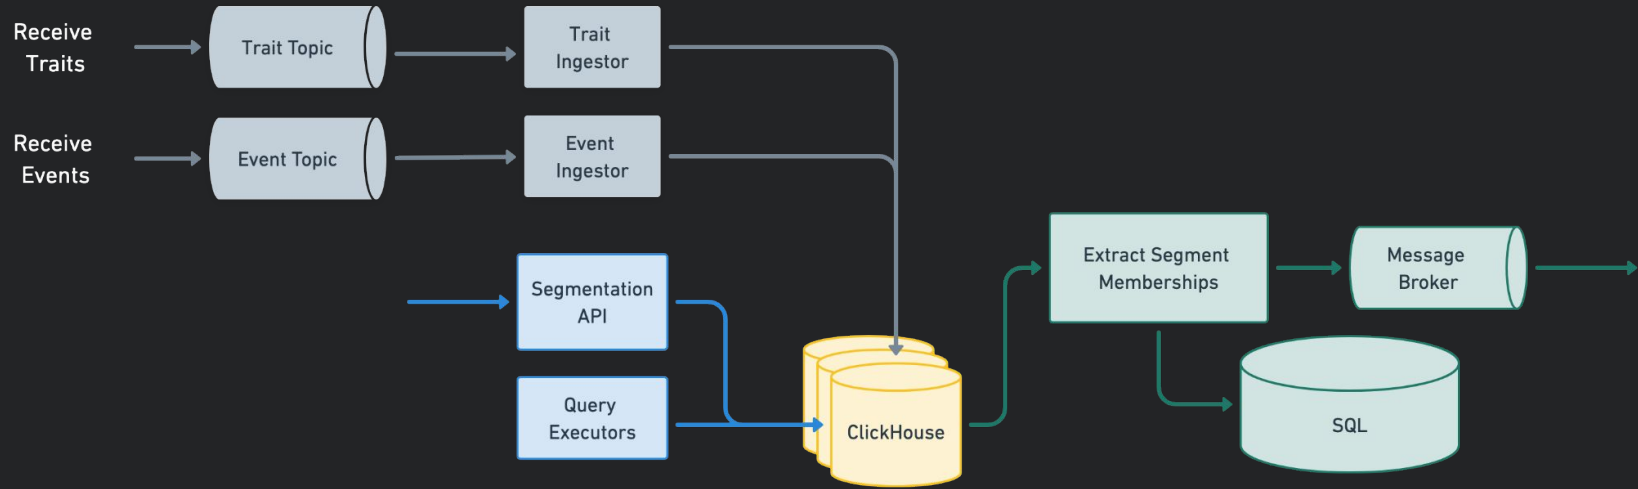
Stores snapshots of profile data.

Replacing merge tree, insert new
record when data changes.

segment_results_log_local

Queries write segment membership
results here.

An overview of segmentation.



A dark gray diagonal stripe runs from the top-left corner to the bottom-right corner of the slide.

Time Based
Periodic Updates

02

Properties about someone

Birthday

day is today

Type: Date

Add condition

0

“Birthday” property is today


23rd	24th	25th (Today)

23rd	24th	25th	26th (Today)

23rd	24th	25th	26th	27th (Today)

Time based periodic updates.

**53 million
req/s!**



Time based periodic updates.

```
select
  segment_id,
  company_id,
  profile_id,
  does_qualify,
  as_of
from segment_results_log_local
```

segment_id	company_id	profile_id	does_qualify	as_of
qwe123	abc123	1	true	2025-03-19 00:00:00
qwe123	abc123	1	true	2025-03-20 00:00:00
qwe123	abc123	1	true	2025-03-21 00:00:00

Time based periodic updates.

“Birthday” property is today

23rd	24th	25th (Today)
	Remove	Add

23rd	24th	25th	26th (Today)
		Remove	Add

23rd	24th	25th	26th	27th (Today)
			Remove	Add

Time based periodic updates.

Between 3 and 5 days ago

17th	18th	19th	20th	21st	22nd	23rd	24th	25th (Today)
		Remove			Add			

17th	18th	19th	20th	21st	22nd	23rd	24th	25th	26th (Today)
			Remove			Add			

17th	18th	19th	20th	21st	22nd	23rd	24th	25th	26th	27th (Today)
				Remove			Add			

Instead of updating segment memberships for every profile...

Run a lightweight query against the event and trait tables, to identify profiles with events or properties within specific time bounds.

Only update segment memberships for *these* profiles.

Time based periodic updates.

events_local

Time Series Aligned 

Since events are ordered by timestamp, we can easily identify profiles with events during specified time bounds.

traits_local

Not Ordered By Timestamp 

Traits are snapshots of profile data, and thus aren't ordered by timestamp.

Read every trait row to identify impacted profiles 🙄.

Time based periodic updates.

Properties about someone

Birthday

day is today

Type: Date

⊕ Add condition

👤 0

```
source_id:                profile_property
company_id:               abc123
profile_id:               1
created_at:               2025-03-25 00:00:00.000
property_keys:            ['Email',      'First Name',  'Favorite Color', 'Birthday'      ]
property_vals_str:        ['patrick@example.com', 'Patrick',    'Red',          '1980-03-25 00:00:00.000']
property_vals_numeric:    [0,          0,          0,          322808400      ]
```

Time based periodic updates.

```
create table traits_eav_local
(
  source_id LowCardinality(String),

  company_id String,
  profile_id UInt32 CODEC(Delta(4), LZ4),
  created_at DateTime64(3) CODEC(T64, LZ4),

  property_name String,
  value_str String,
  value_float Float64
)
ENGINE = ReplicatedReplacingMergeTree(created_at)
PARTITION BY source_id
ORDER BY (company_id, property_name, profile_id)
SETTINGS index_granularity = ...
```

```
CREATE MATERIALIZED VIEW mv_traits_to_traits_eav
TO traits_eav_local
AS
SELECT
  source_id,

  company_id,
  profile_id,
  created_at,

  property_name,
  value_str,
  value_float
FROM traits_local
ARRAY JOIN
  property_keys as property_name,
  property_vals_str as value_str,
  property_vals_numeric as value_float
```

Time based periodic updates.

source_id	company_id	profile_id	created_at	property_name	value_str	value_float
profile_property	abc123	1	2025-03-25 00:00:00.000	Birthday	1980-03-25 00:00:00.000	322808400
profile_property	abc123	1	2025-03-25 00:00:00.000	Email	patrick@example.com	0
profile_property	abc123	1	2025-03-25 00:00:00.000	Favorite Color	Red	0
profile_property	abc123	1	2025-03-25 00:00:00.000	First Name	Patrick	0

Time based periodic updates.

What
(*profile_id*, *segment_id*)
pairs need to be
updated?

```

with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local

    -- Filter to profiles with a qualifying property!
    where (source_id, company_id, profile_id) in (
        select source_id, company_id, profile_id
        from traits_eav_local
        join trait_criteria using (source_id, company_id, property_name)
        where (source_id, company_id, property_name) in (
            select distinct source_id, company_id, property_name
            from trait_criteria
        ) and date_time_in_window(value_float, ...)
    )
)

group by source_id, company_id, profile_id
)

```

```

with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local

    -- Filter to profiles with a qualifying property!
    where (source_id, company_id, profile_id) in (
        select source_id, company_id, profile_id
        from traits_eav_local
        join trait_criteria using (source_id, company_id, property_name)
        where (source_id, company_id, property_name) in (
            select distinct source_id, company_id, property_name
            from trait_criteria
        ) and date_time_in_window(value_float, ...)
    )
)

group by source_id, company_id, profile_id
)

```

```

with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local

    -- Filter to profiles with a qualifying property!
    where (source_id, company_id, profile_id) in (
        select source_id, company_id, profile_id
        from traits_eav_local
        join trait_criteria using (source_id, company_id, property_name)
        where (source_id, company_id, property_name) in (
            select distinct source_id, company_id, property_name
            from trait_criteria
        ) and date_time_in_window(value_float, ...)
    )
)

group by source_id, company_id, profile_id
)

```

```

with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local

    -- Filter to profiles with a qualifying property!
    where (source_id, company_id, profile_id) in (
        select source_id, company_id, profile_id
        from traits_eav_local
        join trait_criteria using (source_id, company_id, property_name)
        where (source_id, company_id, property_name) in (
            select distinct source_id, company_id, property_name
            from trait_criteria
        ) and date_time_in_window(value_float, ...)
    )
)

group by source_id, company_id, profile_id
)

```



```
with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local
)

-- Filter to profiles with a qualifying property!
where (source_id, company_id, profile_id) in (
    select source_id, company_id, profile_id
    from traits_eav_local
    join trait_criteria using (source_id, company_id, property_name)
    where (source_id, company_id, property_name) in (
        select distinct source_id, company_id, property_name
        from trait_criteria
    ) and date_time_in_window(value_float, ...)
)

group by source_id, company_id, profile_id
)
```

```

with trait_criteria as (
    select ... from segment_defs_local
    where ...
)
select profile_id, segment_id
from traits_eav_local
join trait_criteria using (source_id, company_id, property_name)
where date_time_in_window(value_float, ...)

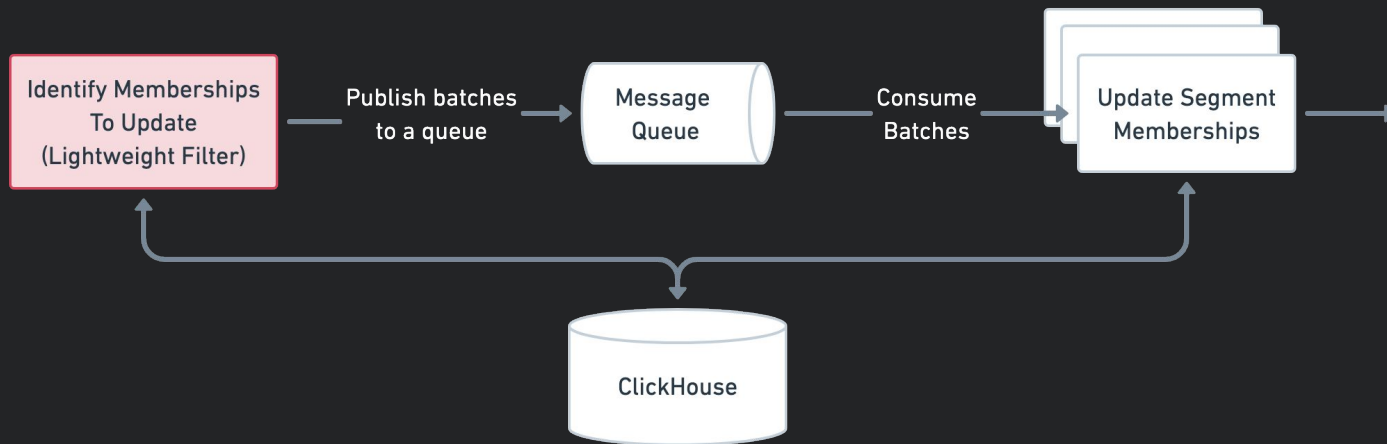
-- Filter to properties that have relative time filters.
and (source_id, company_id, property_name) in (
    select distinct source_id, company_id, property_name from trait_criteria
)

-- Filter on created_at to ensure we read the latest records.
and (source_id, company_id, profile_id, created_at) in (
    select
        source_id, company_id, profile_id, max(created_at) as created_at
    from traits_local

    -- Filter to profiles with a qualifying property!
    where (source_id, company_id, profile_id) in (
        select source_id, company_id, profile_id
        from traits_eav_local
        join trait_criteria using (source_id, company_id, property_name)
        where (source_id, company_id, property_name) in (
            select distinct source_id, company_id, property_name
            from trait_criteria
        ) and date_time_in_window(value_float, ...)
    )

    group by source_id, company_id, profile_id
)

```



Benefits of This Approach

Enables easy horizontal-scaling.

The best-effort filtering step is fast & lightweight. Updating segment memberships is more compute heavy, so parallelize this.



Josh Bradt



Vinay Garg



Isabel Yap



Mike Burton




Jeff Johansen



Himnish Hunma



Sarah Fida



Kanav Bengani



Andrew Volkmann



Kavya Sri
Thadakamalla

klaviyo[®]



careers.klaviyo.com