



Hotpatching ClickHouse in production with XRay

Pablo Marcos Oltra

FOSDEM 2026

\$ whoami



- Pablo Marcos Oltra
- Software Engineer @ ClickHouse's Core Team
- Systems Developer
- Passionate about videogame development
- Interested in compilers, build systems and new programming languages



**Who knows what
hotpatching/hot reloading
is?**



Who knows what LLVM's XRay is?



Who's had to debug something in production and wished they had added an extra log to know what's happening?



01

What is hotpatching

Hotpatching: The basics

- Change the code while the process is running
- Allows faster iteration
- Often done in videogame development
 - Hot reloading rather than hotpatching
 - Client-server architecture where the lightweight client loads a new version of the game (in a shared library)
 - Examples
 - <http://runtimecompiledcplusplus.blogspot.com.es>
 - <https://fungos.github.io/cr-simple-c-hot-reload>
 - <https://liveplusplus.tech>

Hotpatching: Adding instrumentation

- `-finstrument-functions` : Generates instrumentation code for every function entry/exit


```
void __cyg_profile_func_enter(void *this_fn, void *call_site);  
void __cyg_profile_func_exit(void *this_fn, void *call_site);
```



<https://gcc.gnu.org/onlinedocs/gcc/Instrumentation-Options.html#index-finstrument-functions>

Hotpatching: Adding instrumentation

- `-fpatchable-function-entry=N[,M]:`
Generate M NOPs before function entry
and N-M NOPs after function entry



<https://gcc.gnu.org/onlinedocs/gcc/Instrumentation-Options.html#index-fpatchable-function-entry>



02 XRay

XRay

XRay consists of three main parts:

1. Compiler-inserted instrumentation points.
2. A runtime library for enabling/disabling tracing at runtime.
3. A suite of tools for analysing the traces.



As of July 25, 2018, XRay is only available for the following architectures running Linux:

x86_64, arm7 (no thumb), aarch64, powerpc64le, mips, mipsel, mips64, mips64el, NetBSD: x86_64, FreeBSD: x86_64 and OpenBSD: x86_64.

<https://llvm.org/docs/XRay.html>

XRay: Example

```
15  int main()
16  {
17      __xray_init();
18      __xray_patch();
19      __xray_set_handler(my_handler);
20
21      foo();
22
23      __xray_remove_handler();
24      __xray_unpatch();
25
26      return 0;
27 }
```

XRay: Example

```
1 #include <stdint.h>
2 #include <stdio.h>
3 #include <xray/xray_interface.h>
4
5 [[clang::xray_never_instrument]] void my_handler(int32_t fid, XRayEntryType type)
6 {
7     printf("FuncID: %d, Type: %d\n", fid, type);
8 }
9
10 void foo()
11 {
12     printf("Foo!\n");
13 }
```

XRay: Example

```
$ clang++ -o test_xray test_xray.cpp -fxray-instrument  
-fxray-instruction-threshold=1  
  
$ ./test_xray  
FuncID: 1, Type: 0  
Foo!  
FuncID: 1, Type: 1
```

XRay: xray_instr_map

```
$ objdump -h -j xray_instr_map test_xray
```

```
test_xray:      file format elf64-x86-64
```

```
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Align
17	xray_instr_map	00000080	000000000002c6c4	000000000002c6c4	0002c6c4	2**0
CONTENTS, ALLOC, LOAD, READONLY, DATA						

XRay: Before patching foo

```
$ (lldb) dis -C 10
test_xray`foo:
-> 0x555555578360 <+0>: jmp    0x55555557836b ; <+11>
  0x555555578362 <+2>: nopw   0x200(%rax,%rax)
  0x55555557836b <+11>: pushq  %rbp
  0x55555557836c <+12>: movq   %rsp, %rbp
  0x55555557836f <+15>: leaq    0x8348(%rip), %rdi
  0x555555578376 <+22>: movb   $0x0, %al
  0x555555578378 <+24>: callq  0x555555556040 ; symbol stub for: printf
  0x55555557837d <+29>: popq   %rbp
  0x55555557837e <+30>: retq
  0x55555557837f <+31>: nopw   %cs:0x200(%rax,%rax)
```

entry sled



exit sled



XRay: After patching foo

```
$ (lldb) dis -C 10
test_xray`foo:
-> 0x555555578370 <+0>: movl $0x1, %r10d
  0x555555578376 <+6>: callq 0x555555574340 ; __xray_FunctionEntry
    0x55555557837b <+11>: pushq %rbp
    0x55555557837c <+12>: movq %rsp, %rbp
    0x55555557837f <+15>: leaq 0x8338(%rip), %rdi
    0x555555578386 <+22>: movb $0x0, %al
    0x555555578388 <+24>: callq 0x555555556040 ; symbol stub for: printf
    0x55555557838d <+29>: popq %rbp
  0x55555557838e <+30>: movl $0x1, %r10d
  0x555555578394 <+36>: jmp 0x555555574480 ; __xray_FunctionExit
```



03

Integration into ClickHouse

What's ClickHouse

- Fast OLAP DBMS
- Started in 2009 at Yandex.Metrica by Alexey Milovidov
- Open sourced in 2016 under Apache 2.0 license
- Developed in C++
- Strong focus in performance
- Lots of dogfooding
 - CI results, observability, pastila, c.house, etc

<https://clickhouse.com>

<https://benchmark.clickhouse.com>



Integrating XRay into ClickHouse

- Use XRay's instrumentation to have negligible overhead in production when disabled
- Use XRay's runtime library to enable in production only the symbols we want to instrument
- Use SQL statements to enable/disable instrumentation points at runtime
- Kudos to our intern Alina Badakhova for doing the POC to prove it was feasible



<https://github.com/ClickHouse/ClickHouse/issues/74249>

Integrating XRay into ClickHouse

- Landed in ClickHouse 25.12 (December 25)
- Minor improvements since then, all backported

<https://github.com/ClickHouse/ClickHouse/pull/89173>



Hotpatching using SQL: Cheatsheet

```
-- Check symbols to add instrumentation points
SELECT * FROM system.symbols WHERE function_id IS NOT NULL

-- Add or remove instrumentation points
SYSTEM INSTRUMENT [ADD|REMOVE] 'SYMBOL' HANDLER [ENTRY|EXIT] [PARAMS]

-- Check the enabled instrumentation points
SELECT * FROM system.instrumentation

-- Check the instrumented points hit
SELECT * FROM system.trace_log WHERE trace_type = 'Instrumentation'
```

Hotpatching using SQL: Inspecting symbols

```
SET allow_introspection_functions=1

-- Check symbols to add instrumentation points
SELECT
    function_id, symbol_demangled
FROM
    system.symbols
WHERE
    symbol_demangled ILIKE '%QueryMetricLog::start'
```

Hotpatching using SQL: Inspecting symbols

Row 1:

```
function_id: 202915
symbol_demangled: DB::QueryMetricLog::startQuery(std::__1::__basic_string<char,
std::__1::__char_traits<char>, std::__1::__allocator<char>> const&,
std::__1::__chrono::__time_point<std::__1::__chrono::__system_clock,
std::__1::__chrono::__duration<long long, std::__1::__ratio<11, 10000001>>>, unsigned long)
```

Row 2:

```
function_id: 202933
symbol_demangled: void std::__1::__function::__policy_func<void
()>::__call_func[abi:ne210105]<DB::QueryMetricLog::startQuery(std::__1::__basic_string<char,
std::__1::__char_traits<char>, std::__1::__allocator<char>> const&,
std::__1::__chrono::__time_point<std::__1::__chrono::__system_clock, std::__1::__chrono::__duration<long
long, std::__1::__ratio<11, 10000001>>>, unsigned
long)::__$_0>(std::__1::__function::__policy_storage const*)
```

Hotpatching using SQL: Adding instrumentation

```
-- Add a log with stacktrace at function entry
SYSTEM INSTRUMENT ADD 'QueryMetricLog::startQuery'  LOG ENTRY 'FOSDEM 26';

-- Add a random sleep [0, 0.5] secs at function exit
SYSTEM INSTRUMENT ADD 'QueryMetricLog::startQuery'  SLEEP EXIT 0 0.5;

-- Profile deterministically
SYSTEM INSTRUMENT ADD 'QueryMetricLog::startQuery'  PROFILE;
```

Hotpatching using SQL: Adding instrumentation

```
SELECT * FROM system.instrumentation
```

Row 1:

```
id:          0
function_id: 202915
function_name: QueryMetricLog::startQuery
handler:      log
entry_type:   Entry
symbol:      DB::QueryMetricLog::startQuery(std::__1::basic_string<char,
std::__1::char_traits<char>, std::__1::allocator<char>> const&,
std::__1::chrono::time_point<std::__1::chrono::system_clock,
std::__1::chrono::duration<long long, std::__1::ratio<11, 10000001>>>,
unsigned long)
parameters:  ['FOSDEM 26']
```

Hotpatching using SQL: Adding instrumentation

Row 2:

```
id:          1
function_id: 202915
function_name: QueryMetricLog::startQuery
handler:      sleep
entry_type:   Exit
symbol:
DB::QueryMetricLog::startQuery(std::__1::basic_string<char,
std::__1::char_traits<char>, std::__1::allocator<char>> const&,
std::__1::__chrono::time_point<std::__1::__chrono::system_clock,
std::__1::__chrono::duration<long long, std::__1::ratio<1l,
1000000l>>>, unsigned long)
parameters:  [0,0.5]
```

Hotpatching using SQL: Adding instrumentation

Row 3:

```
id:          2
function_id: 202915
function_name: QueryMetricLog::startQuery
handler:      profile
entry_type:   EntryAndExit
symbol:
DB::QueryMetricLog::startQuery(std::__1::basic_string<char,
std::__1::char_traits<char>, std::__1::allocator<char>> const&,
std::__1::__chrono::time_point<std::__1::__chrono::system_clock,
std::__1::__chrono::duration<long long, std::__1::ratio<1l,
1000000l>>>, unsigned long)
parameters:  []
```

Hotpatching using SQL: Removing instrumentation

```
-- Remove all instrumented points  
SYSTEM INSTRUMENT REMOVE ALL;  
  
-- Remove a specific instrumented point with id 2  
SYSTEM INSTRUMENT REMOVE 2;  
  
-- Remove all entry points for function 'QueryMetricLog::startQuery'  
SYSTEM INSTRUMENT REMOVE 'QueryMetricLog::startQuery';
```

Hotpatching using SQL: Checking trace_log

```
SELECT
    event_time_microseconds,
    function_id,
    function_name,
    handler,
    entry_type,
    duration_nanoseconds
FROM system.trace_log
WHERE trace_type = 'Instrumentation'
```

Hotpatching using SQL: Checking trace_log

Row 1:

```
event_time_microseconds: 2026-01-29 11:55:20.553718
function_id:          202915
function_name:
DB::QueryMetricLog::startQuery(std::__1::basic_string<char,
std::__1::__1::char_traits<char>, std::__1::__1::allocator<char>> const&,
std::__1::__1::chrono::time_point<std::__1::__1::chrono::system_clock,
std::__1::__1::chrono::duration<long long, std::__1::__1::ratio<11,
10000001>>>, unsigned long)
handler:              profile
entry_type:           Entry
duration_nanoseconds: 00000
```

Hotpatching using SQL: Checking trace_log

Row 2:

```
event_time_microseconds: 2026-01-29 11:55:20.553758
function_id:          202915
function_name:
DB::QueryMetricLog::startQuery(std::__1::basic_string<char,
std::__1::__1::char_traits<char>, std::__1::__1::allocator<char>> const&,
std::__1::__1::chrono::time_point<std::__1::__1::chrono::system_clock,
std::__1::__1::chrono::duration<long long, std::__1::__1::ratio<11,
10000001>>>, unsigned long)
handler:           profile
entry_type:        Exit
duration_nanoseconds: 40059
```

Hotpatching using SQL: Visualizing the profile

- Use Chrome's Trace Event Format with different visualizers

```
[  
  {"name": "Asub", "cat": "PERF", "ph": "B", "pid": 22630,  
  "tid": 22630, "ts": 829},  
  {"name": "Asub", "cat": "PERF", "ph": "E", "pid": 22630,  
  "tid": 22630, "ts": 833}  
]
```




[https://docs.google.com/document/d/1CvACIvFfyA5R-PhYUmn5OOQtYMH4h6l0nSsKchNAySU/p
review?tab=t.0](https://docs.google.com/document/d/1CvACIvFfyA5R-PhYUmn5OOQtYMH4h6l0nSsKchNAySU/preview?tab=t.0)

Hotpatching using SQL: Converting the data

```
WITH traces AS (
    SELECT * FROM system.trace_log
    WHERE event_date >= today() AND trace_type = 'Instrumentation' AND handler = 'profile'
    ORDER BY event_time, entry_type
)
SELECT
    format(
        '{"traceEvents": [{}\\n]}',
        arrayStringConcat(
            groupArray(
                format(
                    '\\n{{"name": "{}", "cat": "clickhouse", "ph": "{}", "ts": {}, "pid": 1, "tid": {}, "args": {"query_id": "{}", "cpu_id": {}}, "stack": [{}]}},',
                    function_name, if(entry_type = 0, 'B', 'E'), timestamp_ns/1000,
                    toString(thread_id), query_id, cpu_id,
                    arrayStringConcat(arrayMap((x, y) -> concat('"', x, ': ', y, '", '), lines, symbols))
                )
            )
        )
    )
FROM traces;
```

Hotpatching using SQL: Visualizing in Perfetto



Future work

- Add more handlers as we see the need for them
- Add some VM to allow scripting? Wren, Lua
- Educate others that this exists so that they use it. Collect their feedback

Caveats

- Can't mix XRay and sanitized builds due to symbols clashing in libclang_rt

```
ld.lld-19: error: duplicate symbol: __sanitizer::internal_allocator()
>>> defined at
sanitizer_allocator.cpp.o:(__sanitizer::internal_allocator()) in
archive
/usr/lib/llvm-19/lib/clang/19/lib/linux/libclang_rt.ubsan_standalone-
x86_64.a
>>> defined at
sanitizer_allocator.cpp.o:(.text._ZN11__sanitizer18internal_allocator
Ev+0x0) in archive
/usr/lib/llvm-19/lib/clang/19/lib/linux/libclang_rt.xray-x86_64.a
```

- If you mess something with this, you do it big time cause it's in production
- Profiling a function doesn't profile all functions underneath. That's very difficult



Connect with ClickHouse



**ClickHouse
Community Dinner**

**Try ClickHouse for
your use case**

- ClickHouse Cloud
- Download open source

Learn

- Academy / certifications
- Blogs / YouTube

**Engage with
our community**

- Community Slack
- Monthly Community calls
- Meetups / events

**We are Hiring. Come
Work with Us!**



Q&A

All links from the slides



<https://c.house/links>