

chDB

Blazing Fast SQL Engine for Data Science

July, 2024 @Prague

 ClickHouse



Auxten: about me



Experience in RecSys, Database

- Technical Director of ClickHouse core team
- Principal Engineer in Shopee (DB for RecSys)



Love Open Source!

- Contributed to ClickHouse, Jemalloc, K8s, Memcached, CockroachDB, Superset
- Creator of chDB, CovenantSQL
- auxten.com

What is chDB?

in-process


SQL



OLAP Engine

powered by ClickHouse

 PostgreSQL ? but no Server

Python dict ? with SQL support

 SQLite ? but Columnar

 ClickHouse +  python™

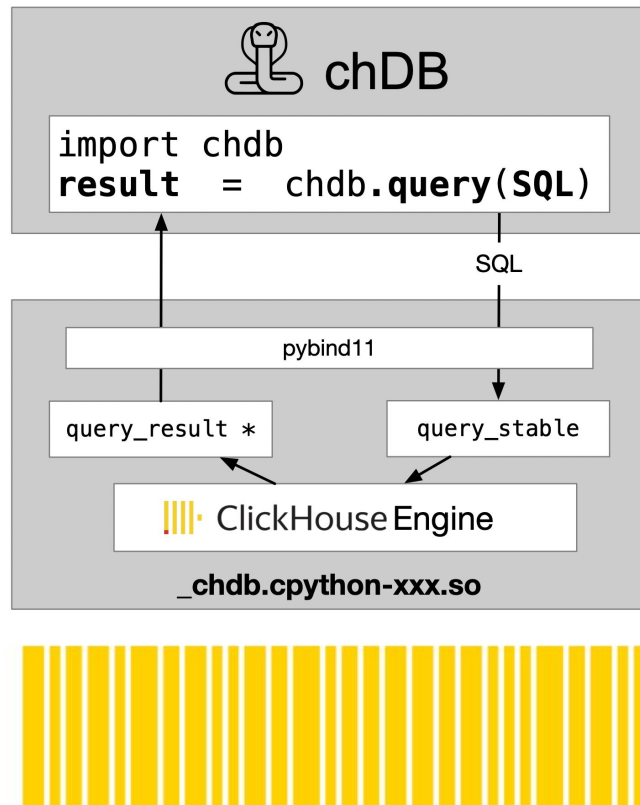
Rocket Engine on a Bicycle

★ chDB 🚲 + 🚀

- In-process SQL OLAP Engine, powered by ClickHouse
- Serverless. No need to install/run ClickHouse
- Supports all ClickHouse Functions & Formats (24.5)
- Support for Python DB API 2.0 and Dataframes
- Support for Stateful Query Sessions w/ Autoclean
- Minimized data copy from C++ to Library binding
- Bindings for Python, Go, Rust, NodeJS, Bun, .NET.

★ Project Background

- Read about the birth of chDB auxten.com/the-birth-of-chdb
- Apache 2.0 Software License



Okay, in-process Database 🤔
What can I do with it?

Everything as a Table

SQLite PostgreSQL MySQL

Parquet

NumPy

CSV

DataFrame

JSON



Pyarrow

[80+ formats](#)

SQL Dump

PyReader

HTTP

S3

HDFS

Query on Python Objects

```
import chdb
import pandas as pd
import pyarrow as pa

data = {
    "a": [1, 2, 3, 4, 5, 6],
    "b": ["tom", "jerry", "auxten", "tom", "jerry", "auxten"],
}

chdb.query("SELECT b, sum(a) FROM Python(data) GROUP BY b").show()

arrow_table = pa.table(data)
chdb.query("SELECT b, sum(a) FROM Python(arrow_table) GROUP BY b").show()

df = pd.DataFrame(data)
chdb.query("SELECT b, sum(a) FROM Python(df) GROUP BY b").show()
```

Only Numerical and String column type supported on v2.0.0b1

Join multiple data sources



chDB



```
chdb.query("""
SELECT name, age, sex, some new tag, ...
FROM url('some_http_data.parquet') big
LEFT JOIN file('some_new_data.csv') local_csv
ON big.uid = local_csv.uid
LEFT JOIN Python(some_processed_df) df
ON df.uid = big.uid
LIMIT 1000
""").show()
```



 pandas



Your own Table Engine in Python

```
import chdb

class myReader(chdb.PyReader):
    def __init__(self, data):
        # do some init

        ...
        self.cursor = 0


    def read(self, col_names, count):
        # return block like data[cursor:cursor+count]

        ...
        self.cursor += count
        return ret

reader = myReader()

chdb.query("SELECT b, sum(a) FROM Python(reader) GROUP BY b").show()
```





★ chDB in Python/Golang/Rust/NodeJS/Bun/.NET



```
import chdb
res = chdb.query('select version()', 'CSV');
```



_chdb.cpython-xxx.so



```
const chdb = require("chdb-node");
var result = chdb.Execute("SELECT version()", "CSV");
```



Golang

```
package main

import (
    "github.com/chdb-io/chdb-go/chdb"
)

func main() {
    result := chdb.Query("SELECT version()", "CSV")
}
```



libchdb.so





★ **Use chDB** (almost) **anywhere**




On

Data Science / LLM / Lambda / Mobile Phone



Okay, Database 🤔
Is it fast?

Benchmark on Parquet




Detailed Comparison

	chDB (Parquet, partitioned) (c6a.metal, 500gb gp2)	ClickHouse (Parquet, partitioned) (c6a.metal, 500gb gp2)	DuckDB (Parquet, partitioned) (c6a.4xlarge, 500gb gp2)	chDB (Parquet, partitioned) (c6a.4xlarge, 500gb gp2)	ClickHouse (Parquet, partitioned) (c6a.4xlarge, 500gb gp2)	DataFusion (Parquet, partitioned) (c6a.4xlarge, 500gb gp2)
Load time:	0	0	0	0	0	0
Data size:	13.73 GiB (×1.00)	13.73 GiB (×1.00)	13.73 GiB (×1.00)	13.73 GiB (×1.00)	13.73 GiB (×1.00)	13.76 GiB (×1.00)
Q0.	0.035s (×2.35)	0.085s (×5.00)	0.043s (×2.78)	0.020s (×1.57)	0.040s (×2.63)	0.009s (×1.00)
Q1.	0.071s (×2.31)	0.114s (×3.54)	0.061s (×2.03)	0.069s (×2.24)	0.085s (×2.71)	0.025s (×1.00)
Q2.	0.115s (×1.58)	0.129s (×1.76)	0.104s (×1.45)	0.104s (×1.44)	0.144s (×1.95)	0.069s (×1.00)
Q3.	0.111s (×1.46)	0.181s (×2.30)	0.093s (×1.24)	0.101s (×1.33)	0.110s (×1.45)	0.073s (×1.00)
Q4.	1.183s (×3.04)	0.382s (×1.00)	0.539s (×1.40)	0.448s (×1.17)	0.429s (×1.12)	0.782s (×2.02)
Q5.	1.399s (×2.85)	0.485s (×1.00)	0.753s (×1.54)	0.640s (×1.31)	0.646s (×1.33)	1.172s (×2.39)
Q6.	0.104s (×2.78)	0.104s (×2.78)	0.128s (×3.37)	0.086s (×2.34)	0.099s (×2.66)	0.031s (×1.00)
Q7.	0.087s (×2.62)	0.104s (×3.08)	0.064s (×1.99)	0.074s (×2.27)	0.087s (×2.62)	0.027s (×1.00)
Q8.	0.410s (×1.00)	0.463s (×1.13)	0.666s (×1.61)	0.643s (×1.56)	0.600s (×1.45)	1.389s (×3.33)
Q9.	0.430s (×1.00)	0.476s (×1.10)	0.899s (×2.07)	0.783s (×1.80)	0.696s (×1.60)	0.964s (×2.21)
Q10.	0.221s (×1.05)	0.246s (×1.16)	0.210s (×1.00)	0.295s (×1.39)	0.301s (×1.41)	0.274s (×1.29)
Q11.	0.237s (×1.00)	0.313s (×1.31)	0.246s (×1.04)	0.343s (×1.43)	0.355s (×1.48)	0.308s (×1.29)
Q12.	0.350s (×1.00)	0.481s (×1.36)	0.633s (×1.79)	0.716s (×2.02)	0.792s (×2.23)	1.237s (×3.47)
Q13.	0.398s (×1.00)	0.559s (×1.39)	1.014s (×2.51)	1.024s (×2.54)	1.088s (×2.69)	2.509s (×6.17)
Q14.	0.392s (×1.00)	0.483s (×1.23)	0.688s (×1.74)	0.833s (×2.10)	0.923s (×2.32)	1.387s (×3.47)
Q15.	0.264s (×1.00)	0.390s (×1.46)	0.598s (×2.22)	0.570s (×2.12)	0.639s (×2.37)	0.899s (×3.32)
Q16.	0.707s (×1.00)	0.823s (×1.16)	1.392s (×1.96)	1.969s (×2.76)	2.013s (×2.82)	2.619s (×3.67)
Q17.	0.621s (×1.00)	0.671s (×1.08)	1.323s (×2.11)	1.207s (×1.93)	1.276s (×2.04)	2.555s (×4.06)
Q18.	1.239s (×1.00)	1.545s (×1.24)	2.332s (×1.88)	3.762s (×3.02)	4.364s (×3.50)	5.596s (×4.49)
Q19.	0.077s (×1.15)	0.137s (×1.93)	0.087s (×1.28)	0.091s (×1.33)	0.101s (×1.46)	0.066s (×1.00)
Q20.	0.414s (×1.00)	0.658s (×1.57)	1.841s (×4.36)	1.230s (×2.92)	1.838s (×4.36)	1.558s (×3.70)
Q21.	0.419s (×1.00)	0.790s (×1.86)	1.689s (×3.96)	1.742s (×4.08)	2.318s (×5.42)	1.855s (×4.34)
Q22.	0.891s (×1.00)	1.320s (×1.48)	3.460s (×3.85)	4.073s (×4.53)	5.124s (×5.70)	4.159s (×4.63)
Q23.	4.408s (×1.00)	4.840s (×1.10)	11.130s (×2.52)	15.832s (×3.59)	18.346s (×4.15)	11.146s (×2.53)
Q24.	0.180s (×1.00)	0.297s (×1.61)	0.479s (×2.57)	0.427s (×2.29)	0.462s (×2.48)	0.488s (×2.62)
Q25.	0.192s (×1.00)	0.230s (×1.19)	0.357s (×1.81)	0.385s (×1.95)	0.352s (×1.79)	0.422s (×2.14)

Benchmark on DataFrame – 4.6x Faster

DataFrame Benchmark Results on 1000000 rows of ClickBench



Okay, But WHY?

Why chDB is Fast

Mostly, ClickHouse is Fast



Why chDB is Fast

Just make sure

Python

does not slow it down

1. Hold this →



2. Do everything with C++ in Parallel




Okay, But why is ClickHouse Fast?

Why ClickHouse is Fast


- Column-oriented storage
- Data compression
- Vectorized query execution
- JIT(Just In Time) & Dynamic Dispatch
 - Compile SQL into native cpu instruction
 - Runtime check CPU spec and dispatch to AVX, AVX2, AVX512 specialized function
-
- Keep benchmark and optimization for 15 years

<https://clickhouse.com/docs/en/faq/general/why-clickhouse-is-so-fast>

Row-oriented databases



Column-oriented databases



Recap



chDB

Features & Use-Cases



Pure Performance

In-process chDB eliminates overhead communication between clients and servers accessing cloud datasets



Seamless Integration

Full ClickHouse OLAP functionality included, no need to change query style or renounce any advanced feature



Reduced Consumption

chDB runs alongside your code on-demand, with no need to maintain any costly backend server infrastructure



Real-time Analytics

In-process chDB enables OLAP analytics directly off cloud storage, S3, Parquet files or ClickHouse services



Quick Prototype

Develop and showcase your prototype directly in your Notebook on Laptop



chDB

Just try chDB

Docs

- For chDB specific examples and documentation refer to clickhouse.com/docs/en/chdb
- For SQL syntax, please refer to [ClickHouse SQL Reference](#)
- The birth of chDB auxten.com/the-birth-of-chdb

Demos

- [Project Documentation](#) and [Usage Examples](#)
- [Colab Notebooks](#) and [other Script Examples](#)
- [ClickBench of embedded engines](#)

Contact

- Email: auxten@clickhouse.com
- Twitter: [@auxten](#)



Thank You!

Check it out → chdb.io