

Using SQLancer to test ClickHouse and other database systems

Manuel Rigger



@RiggerManuel

Ilya Yatsishin



qoega

Plan

- | What is ClickHous and why do we need good testing?
- | How do we test ClickHouse and what problems we have to solve?
- | What is SQLancer and what are the ideas behind them?
- | How to add yet another DBMS support to SQLancer?



ClickHouse

Open Source analytical DBMS for BigData with SQL interface.

- Blazingly fast
- Scalable
- Fault tolerant



2013 Project started

2016 Open Sourced

2021
★15K GitHub

<https://github.com/ClickHouse/clickhouse>

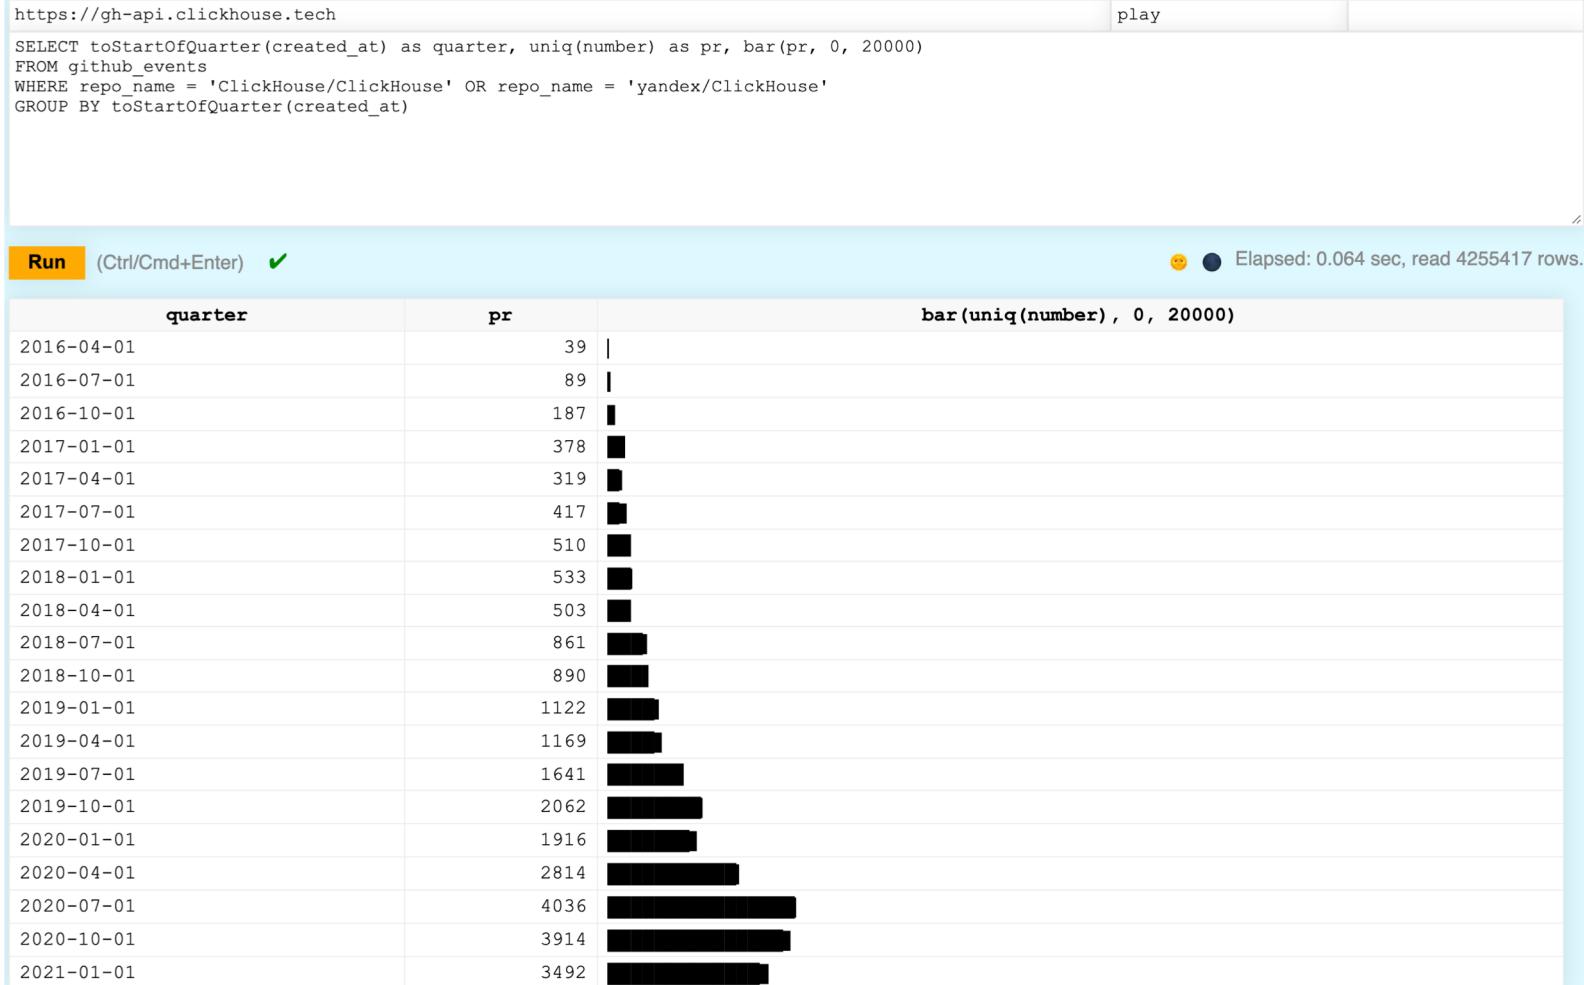
Why do we need good CI?

In 2020:

361	4081	261	11	<15
-----	------	-----	----	-----

Contributors	Merged Pull Requests	New Features	Releases	Core Team
--------------	-------------------------	--------------	----------	-----------

Pull Requests exponential growth



All GitHub data available in ClickHouse

```
https://gh-api.clickhouse.tech
```

play

```
SELECT
    repo_name,
    uniq(actor_login),
    uniq(number) as pull_requests,
    sum(merged) as merged_prs,
    sum(state='closed' and not merged) as closed_prs,
    round(sum(merged) / uniq(number) * 100) AS merged_share
FROM github_events
WHERE repo_name = 'ClickHouse/ClickHouse' and created_at BETWEEN '2020-01-01 00:00:00' AND '2021-01-01 00:00:00'
    AND event_type = 'PullRequestEvent' AND base_ref = 'master'
GROUP BY repo_name
ORDER BY uniq(actor_login) DESC
```

Run (Ctrl/Cmd+Enter) ✓

Elapsed: 0.044 sec, read 156163 rows.

repo_name	uniq(actor_login)	pull_requests	merged_prs	closed_prs	merged_share
ClickHouse/ClickHouse	361	4711	4021	578	85

<https://gh.clickhouse.tech/explorer/>
<https://gh-api.clickhouse.tech/play>

How is ClickHouse tested?

 Open

Disk S3 possibility to migrate to restorable schema #22070

Jokser wants to merge 4 commits into [ClickHouse:master](#) from [Jokser:disk-s3-migration](#) 



Some checks were not successful

3 failing and 48 successful checks

[Hide all checks](#)

  Integration tests (thread) — fail: 27, passed: 411, flaky: 9 [Details](#)

  Performance — 5 faster, 8 slower, 65 unstable [Details](#)

  AST fuzzer (ASan) — OK [Details](#)

  AST fuzzer (MSan) — OK [Details](#)

  AST fuzzer (TSan) — OK [Details](#)

 AST fuzzer (UBSan) — OK [Details](#)



This branch has no conflicts with the base branch

Merging can be performed automatically.

[Merge pull request](#) ▾

You can also [open this in GitHub Desktop](#) or view [command line instructions](#).

ClickHouse Testing

- Style
- Unit
- Functional
- Integrational
- Performance
- Stress
- Static analysis (clang-tidy, PVSStudio)
- Compatibility with OS versions
- Flaky check for new or changed tests
- All tests are run in with sanitizers (address, memory, thread, undefined behavior)
- Thread fuzzing (switch running threads randomly)
- Coverage
- Fuzzing

<https://rc1a-ity5agjmuhyu6nu9.mongodb.yandexcloud.net:8443>

```
SELECT DISTINCT check_name
FROM "gh-data".checks
WHERE commit_sha = '0675f9403cbd8bf17c0d114113dbafb8f911be69'
```

Run

(Ctrl/Cmd+Enter) ✓

	check_
Compatibility check	
Docs release	
Fast test	
Unit tests release gcc	
Yandex synchronization (only for Yandex employees)	
Functional stateful tests (address)	
Functional stateful tests (debug)	
Functional stateful tests (memory)	
Functional stateful tests (release)	
Functional stateful tests (release, DatabaseOrdinary)	
Functional stateful tests (thread)	
Functional stateful tests (ubsan)	
Functional stateless tests (release, wide parts enabled)	
Functional stateless tests (ubsan)	
PVS check	
Push to dockerhub	
	Push to dockerhub
	SQLancer test
	Split build smoke test
	Stress test (debug)
	Style Check
	Functional stateless tests (ANTLR debug)
	Unit tests ASAN
	Unit tests release clang
	Functional stateless tests (address)
	Functional stateless tests (debug)
	Functional stateless tests (memory)
	Functional stateless tests (release)
	Functional stateless tests (release, DatabaseOrdinary)
	Functional stateless tests (thread)
	Integration tests (memory)
	Integration tests (release)
	Stress test (address)
	Stress test (memory)
	Stress test (thread)
	Stress test (undefined)
	Testflows check
	Unit tests MSAN
	Unit tests TSAN
	Unit tests UBSAN
	Integration tests (asan)
	Integration tests (thread)

100K tests is not enough

<https://rc1a-ity5agjmuhyu6nu9.mdb.yandexcloud.net:8443> qoega ······

```
SELECT COUNT(DISTINCT check_name), COUNT(DISTINCT (check_name, test_name))
FROM "gh-data".checks
WHERE commit_sha = '0675f9403cbd8bf17c0d114113dbafb8f911be69'
```

Run

(Ctrl/Cmd+Enter)



Elapsed: 2.268 sec, read 315706649 rows.

uniqExact(check_name)	uniqExact(tuple(check_name, test_name))
41	106714

Fuzzing

- | libFuzzer to test generic data inputs – formats, schema etc.
- | Thread Fuzzer – randomly switch threads to trigger races and dead locks
 - > https://presentations.clickhouse.tech/cpp_siberia_2021/
- | AST Fuzzer – mutate queries on AST level.
 - > Use SQL queries from all tests as input. Mix them.
 - > High level mutations. Change query settings
 - > <https://clickhouse.tech/blog/en/2021/fuzzing-clickhouse/>

Test development steps



Unit,
Functional,
Integration,
Stress,
Performance
etc.

Find bugs earlier.
Sanitizers:
• Address
• Memory
• Undefined Behavior
• Thread

Improve
coverage.

???

In search for correctness

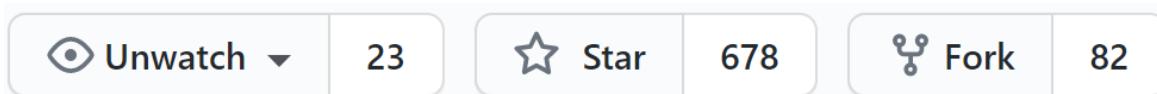
- | Test cases for correctness require developer or QA specialist to write it manually
- | Additional tests mostly check if DBMS crashes or stops working
- | It is hard to find reference system to validate results with:
 - SQL syntax differs between systems
 - No SQL conformance testing suite
 - All SQL standard extensions can't be tested this way

How we can check correctness?

- | SQL has some invariants and basic assumptions
- | DBMS can help us!

There should be some solutions that already exploit that

SQLancer



<https://github.com/sqlancer>

SQLancer is an effective and widely-used
automatic testing tool to find **logic bugs** in DBMSs

Facts

- Written in Java



```
$ git clone https://github.com/sqlancer/sqlancer
$ cd sqlancer
$ mvn package -DskipTests
$ cd target
$ java -jar sqlancer-*.jar sqlite3
```

You can quickly **try out** SQLancer on the embedded DBMSs **SQLite, H2, and DuckDB** without setting up a connection

Facts

- Written in Java
- Permissive license (MIT License)

master ▾ [sqlancer / LICENSE.md](#) Go to file ...

 sqlancer/sqlancer is licensed under the MIT License	Permissions ✓ Commercial use ✓ Modification ✓ Distribution ✓ Private use	Limitations ✗ Liability ✗ Warranty	Conditions (i) License and copyright notice
A short and simple permissive license with conditions only requiring preservation of copyright and license notices. Licensed works, modifications, and larger works may be distributed under different terms and without source code.			
This is not legal advice. Learn more about repository licenses.			

Facts

- Written in Java
- Permissive license (MIT License)
- 43,000 LOC

Facts

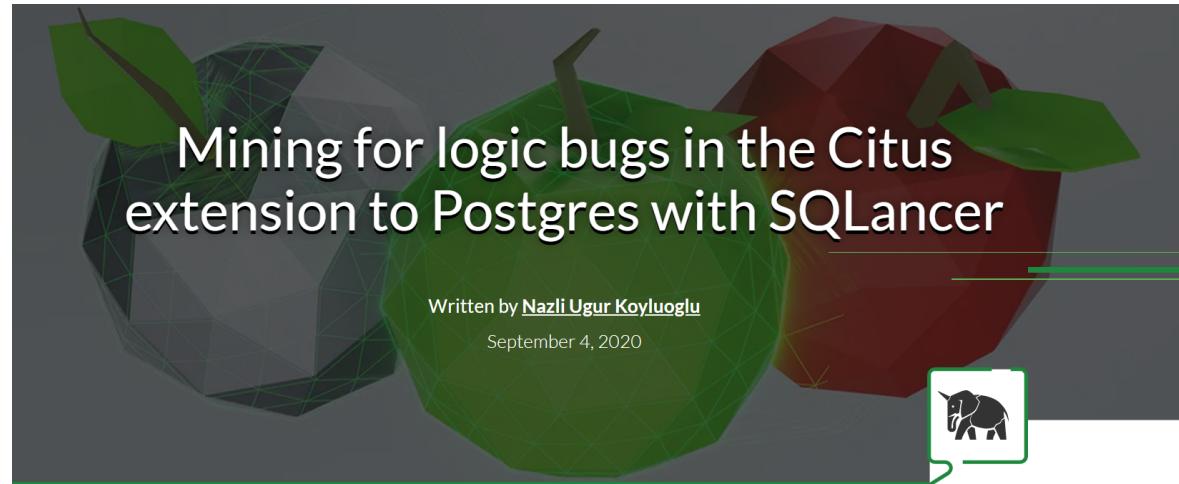
- Written in Java
- Permissive license (MIT License)
- 43,000 LOC

Facts

- Written in Java
- Permissive license (MIT License)
- 43,000 LOC
- 9 contributors

Contributors

Nazli Ugur Koyluoglu



<https://www.citusdata.com/blog/2020/09/04/mining-for-logic-bugs-in-citus-with-sqlancer/>

Contributors

Patrick Stäuble



Contributors

Ilya Yatsishin



Yandex



Facts

- Written in Java
- Permissive license (MIT License)
- 43,000 LOC
- 9 contributors
- >= 10 supported DBMSs

Supported DBMSs

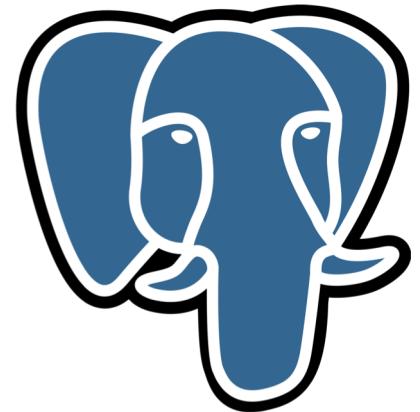


TiDB



PostgreSQL

● DuckDB



SQLancer



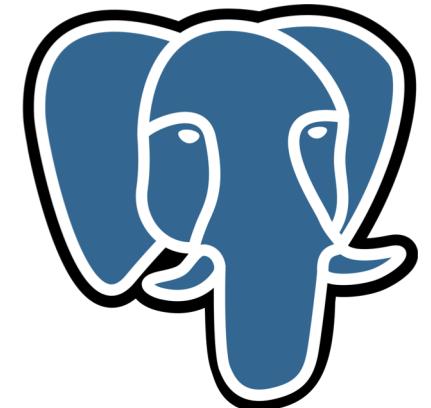
<https://github.com/sqlancer>

SQLancer is an **effective** and widely-used
automatic testing tool to find **logic bugs** in DBMSs

More than 450 Bugs



PostgreSQL



• DuckDB

I used SQLancer to find **over 450 unique, previously unknown bugs** in widely-used DBMSs

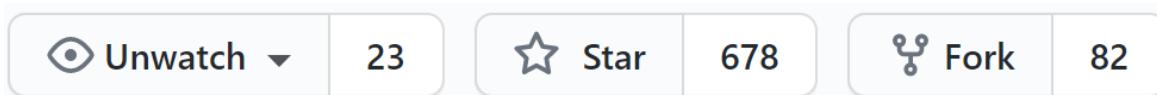


More than 450 Bugs

DBMS	# Bugs		
	Logic	Error	Crash
CockroachDB	16	46	5
DuckDB	29	13	31
H2	2	15	1
MariaDB	5	0	1
MySQL	21	9	1
PostgreSQL	1	11	5
SQLite	93	44	42
TiDB	30	27	4
Sum	197	165	90

<https://github.com/sqlancer/bugs>

SQLancer



<https://github.com/sqlancer>

SQLancer is an effective and **widely-used**
automatic testing tool to find **logic bugs** in DBMSs

Adoption



DuckDB



ClickHouse



Adoption



DuckDB



ClickHouse



Yandex uses SQLancer to test
every commit



Adoption



DuckDB



ClickHouse

✓	# 5677.15	AMD64	Bionic	SQLancer	40 min 52 sec
✓	# 5677.16	AMD64	Bionic	SQLancer (with Address Sanitizer)	24 min 34 sec

DuckDB runs SQLancer on every pull request

Adoption

“With the help of SQLancer, an automatic DBMS testing tool, we have been able to identify **>100 potential problems** in corner cases of the SQL processor.”



<https://www.monetdb.org/blog/faster-robuster-and-feature-richer-monetdb-in-2020-and-beyond>

Adoption

Watch 8 Star 43 Fork 4

chaos-mesh / go-sqlancer

Code Issues Pull requests Actions Projects Wiki Security Insights

go-sqlancer

fuzzing tidb

58 commits	4 branches	0 packages	0 releases	4 contributors
Branch: master	New pull request	Create new file	Upload files	Find file
Clone or download				
Latest commit 3c755aa yesterday				
rename flags and update readme (#50)	yesterday			
cmd	rename flags and update readme (#50)	yesterday		
pkg	rename flags and update readme (#50)	yesterday		
.gitignore	remove sqldsmith	2 months ago		
Makefile	add transformer and support TLP (#48)	yesterday		
README.md	rename flags and update readme (#50)	yesterday		
go.mod	add transformer and support TLP (#48)	yesterday		
go.sum	add transformer and support TLP (#48)	yesterday		



PingCAP implemented a tool go-sqlancer

<https://github.com/chaos-mesh/go-sqlancer>

Adoption

```
original query:  
SELECT ALL t1.c1, t1.c0, t2.c1, t4.c1, t2.c0  
FROM t1  
CROSS JOIN t4  
FULL OUTER JOIN t2 ON (((((((['832125354,1134163512')::int4range)*('0,2106623281-...  
  
semantically equivalent to partitioned query:  
SELECT ALL t1.c1, t1.c0, t2.c1, t4.c1, t2.c0  
FROM t1  
CROSS JOIN t4  
FULL OUTER JOIN t2 ON (((((((['832125354,1134163512')::int4range)*('0,2106623281-...  
WHERE ((t4.c0)<(CAST(upper(t1.c0) AS VARCHAR(893)) AS VARCHAR)))  
  
UNION ALL SELECT t1.c1, t1.c0, t2.c1, t4.c1, t2.c0  
FROM t1  
CROSS JOIN t4  
FULL OUTER JOIN t2 ON (((((((['832125354,1134163512')::int4range)*('0,2106623281-...  
WHERE NOT ((t4.c0)<(CAST(upper(t1.c0) AS VARCHAR(893)) AS VARCHAR)))  
  
UNION ALL SELECT ALL t1.c1, t1.c0, t2.c1, t4.c1, t2.c0  
FROM t1  
CROSS JOIN t4  
FULL OUTER JOIN t2 ON (((((((['832125354,1134163512')::int4range)*('0,2106623281-...  
WHERE ((t4.c0)<(CAST((upper(t1.c0)::VARCHAR(893) AS VARCHAR)))) ISNULL
```



SQLancer



<https://github.com/sqlancer>

SQLancer is an effective and widely-used
automatic testing tool to find **logic bugs** in **DBMSs**

Database Management Systems (DBMS)

**Structured Query
Language (SQL)**



Interact with

Database Management System

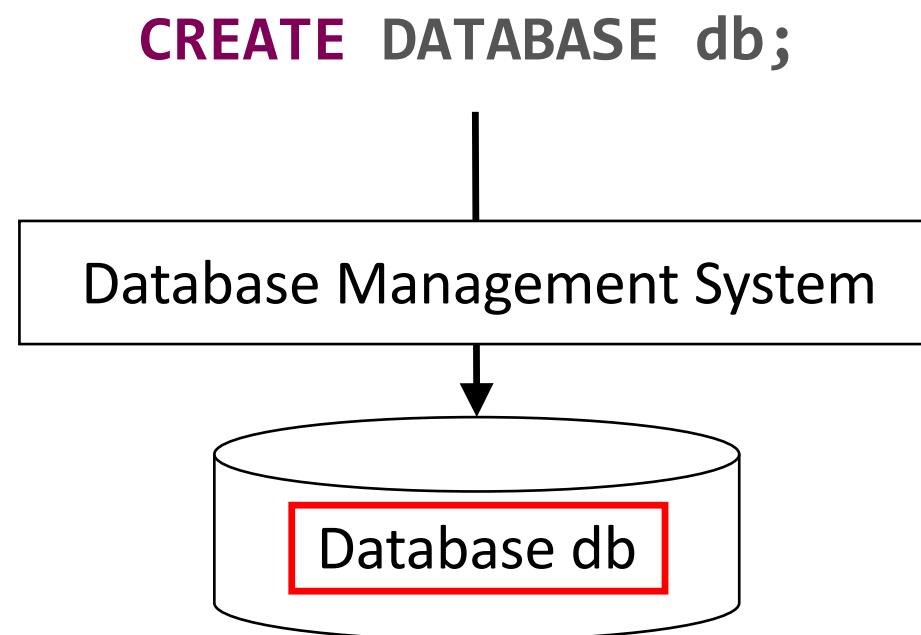
Database Management Systems (DBMS)

CREATE DATABASE db;



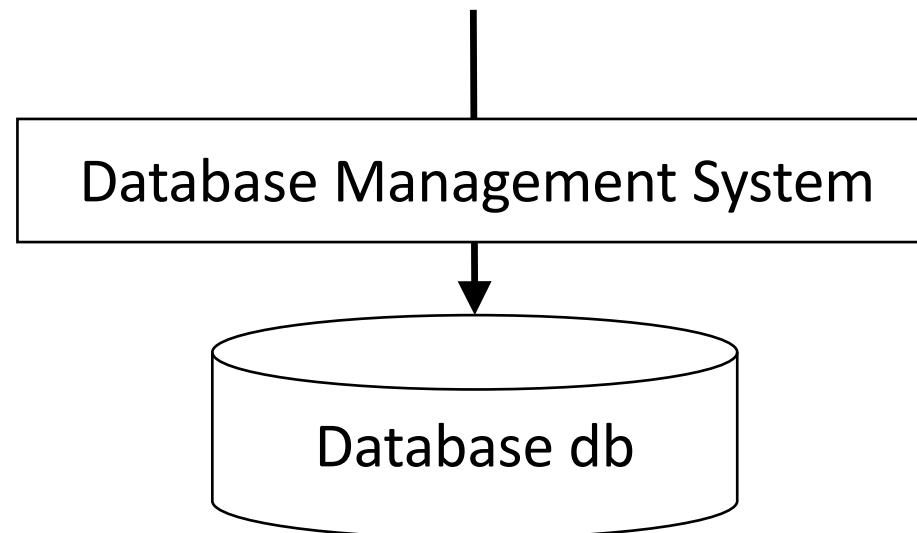
Database Management System

Database Management Systems (DBMS)



Database Management Systems (DBMS)

CREATE TABLE t0(c0 INTEGER);



Relational Model

Relational Model

Table

Relational Model

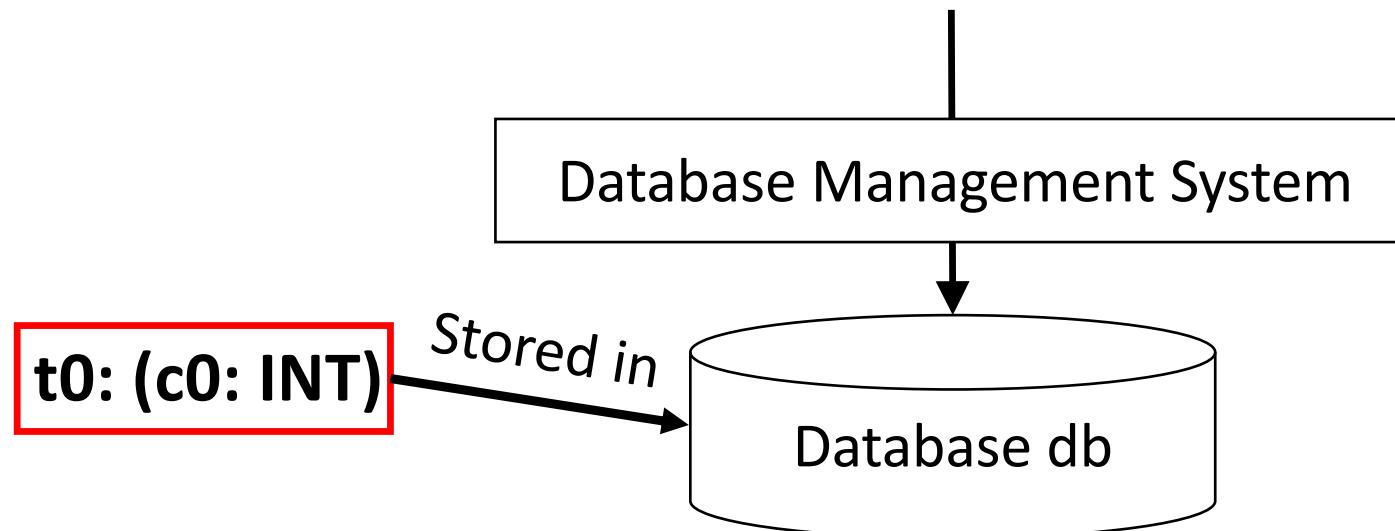
Column

Relational Model

Row

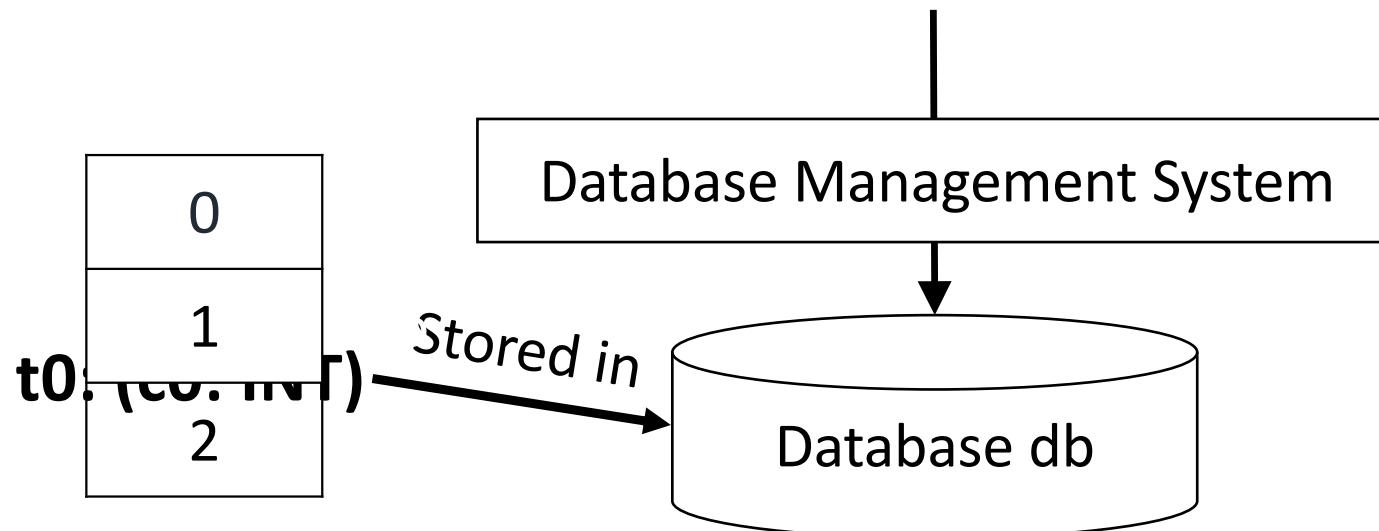
Database Management Systems (DBMS)

```
CREATE TABLE t0(c0 INTEGER);
```

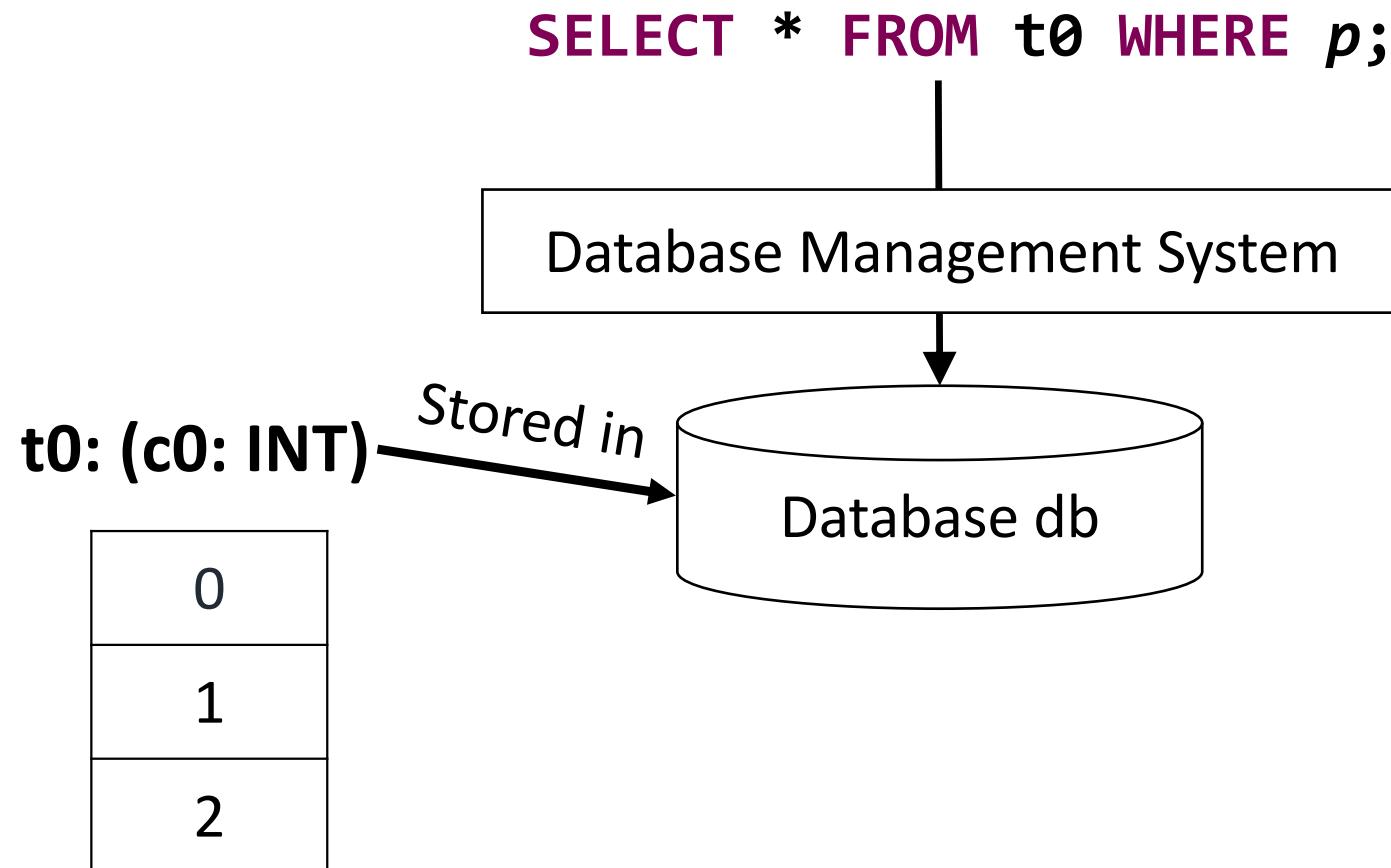


Database Management Systems (DBMS)

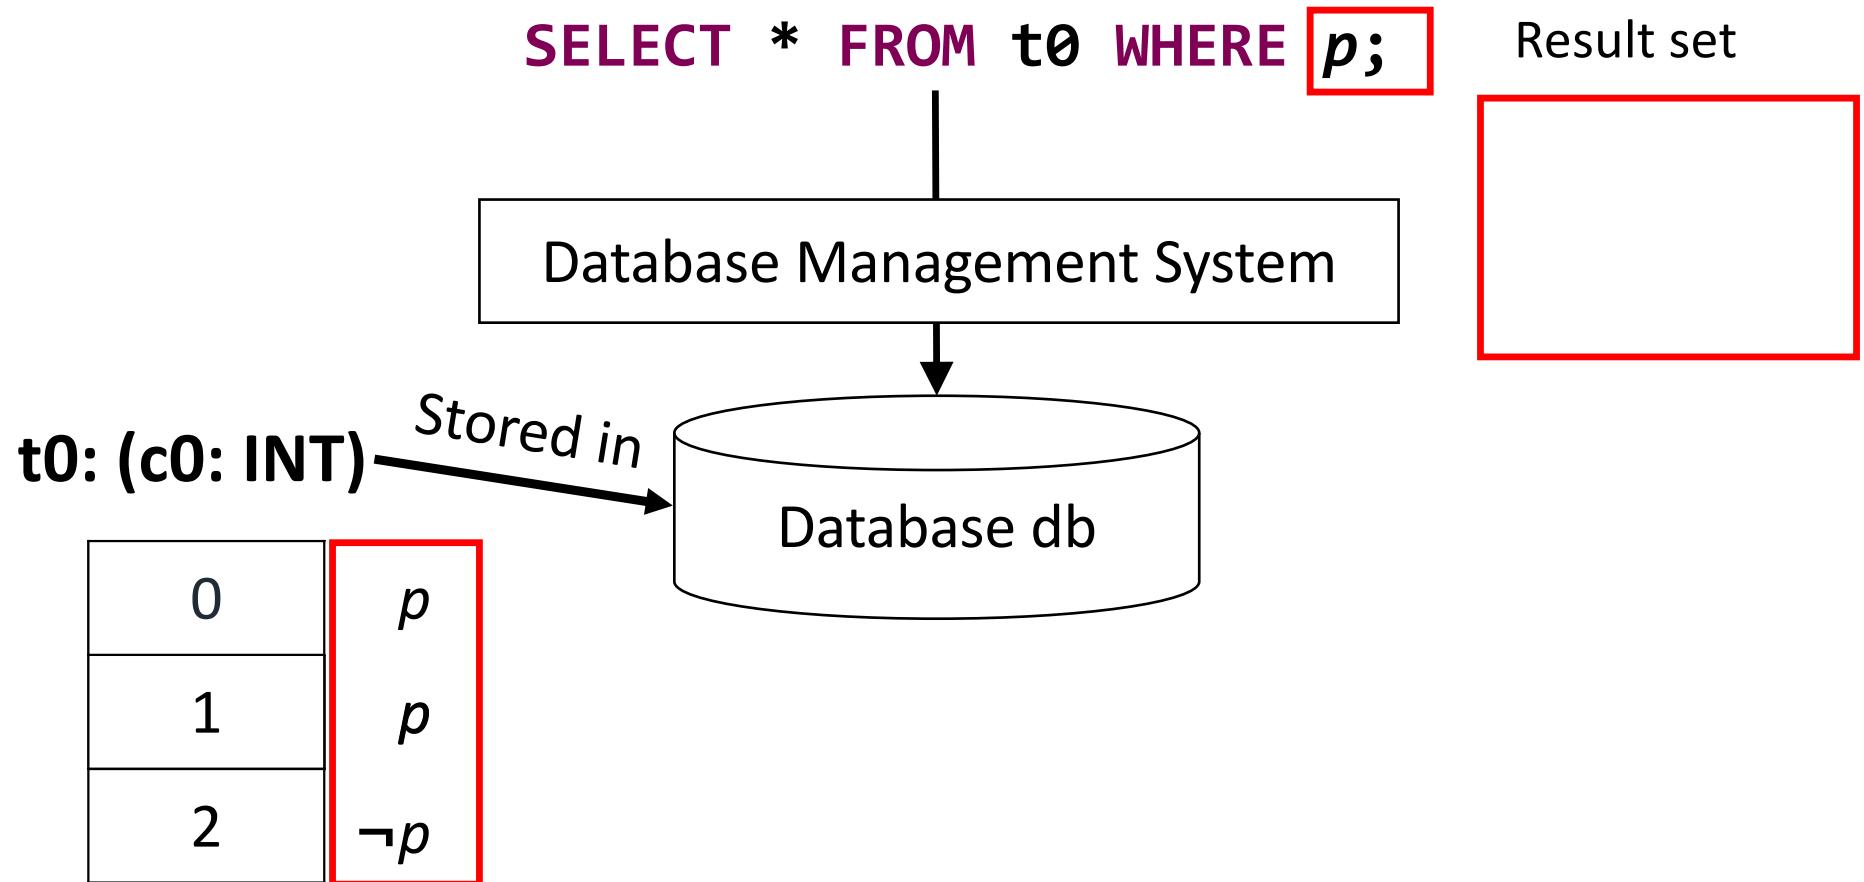
```
INSERT INTO t0(c0) VALUES (0), (1), (2);
```



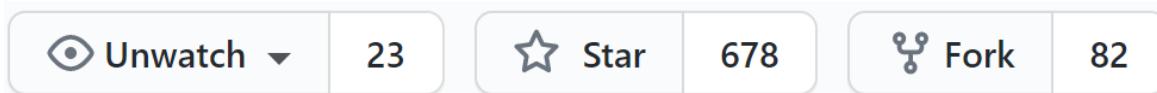
Database Management Systems (DBMS)



Database Management Systems (DBMS)



SQLancer



<https://github.com/sqlancer>

SQLancer is an effective and widely-used
automatic testing tool to find **logic bugs** in DBMSs

Goal: Find Logic Bugs



Bugs that cause the DBMS to return an **incorrect result set**

Motivating Example

t0	t1
c0	c0
0.0	-0.0



```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```

→ ?

It might seem **disputable** whether the predicate should evaluate to true

Motivating Example

t0	t1
c0	c0
0.0	-0.0

false?

```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



0	0	0	0	...	0	0	0	0
-0	1	0	0	0	...	0	0	0

Different binary representation

t0.c0	t1.c0
-------	-------

Motivating Example

t0	t1
c0	c0
0.0	-0.0

true?

```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



Evaluates to true for **most** programming languages

t0.c0	t1.c0
0.0	-0.0

Motivating Example

t0	t1
c0	c0
0.0	-0.0



```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



t0.c0	t1.c0
0.0	-0.0

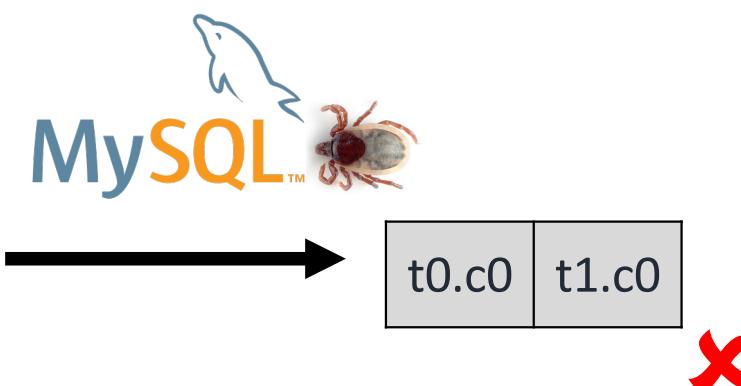


Motivating Example

t0	t1
c0	c0
0.0	-0.0

The latest version of MySQL that we tested failed to fetch the row

```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



Motivating Example

t0

c0
0.0

t1

c0
-0.0

[3 Apr 2020 13:07] Jon Stephens

Documented fix as follows in the MySQL 8.0.21 changelog:

A query whose predicate compared 0 with -0 where at least one of these was a floating-point value returned incorrect results.



```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



t0.c0	t1.c0
-------	-------

✗

Motivating Example

t0	t1
c0	c0
0.0	-0.0

We could automatically find the bug without having an accurate understanding ourselves



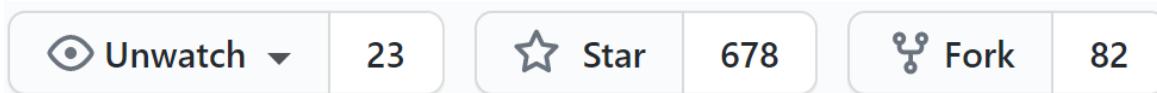
```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



t0.c0	t1.c0
-------	-------



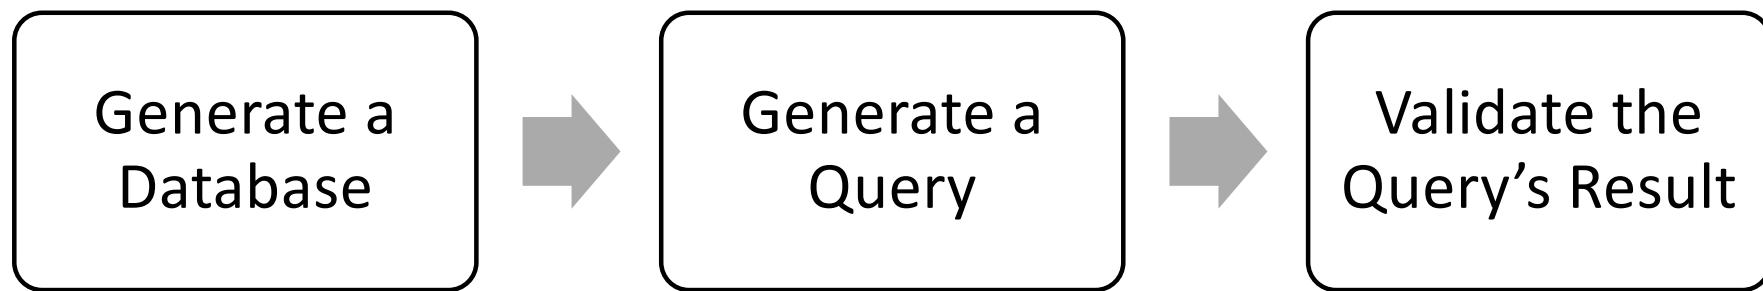
SQLancer



<https://github.com/sqlancer>

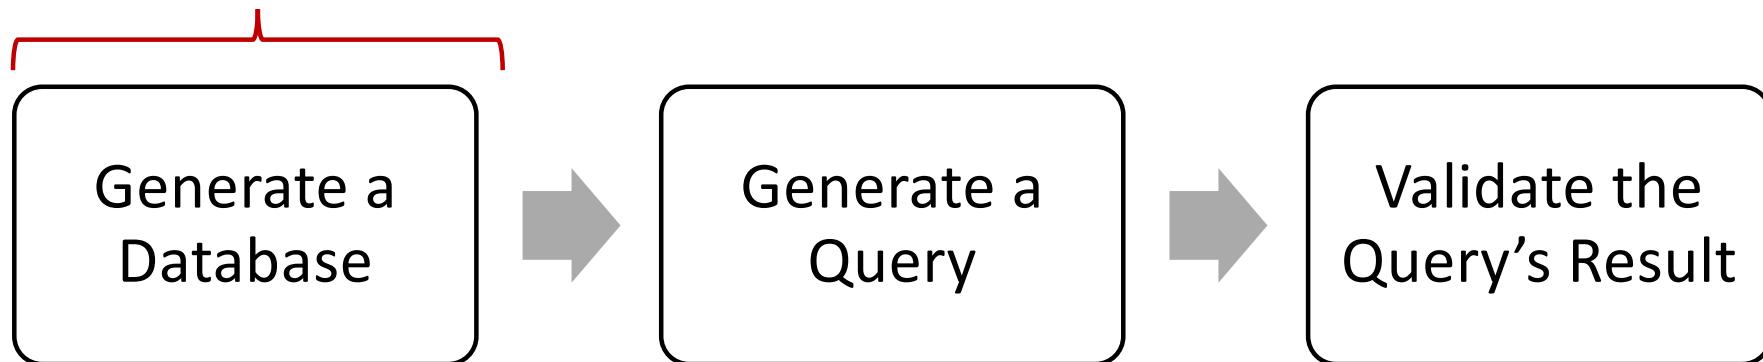
SQLancer is an effective and widely-used
automatic testing tool to find **logic bugs** in DBMSs

Automatic Testing Core Challenges



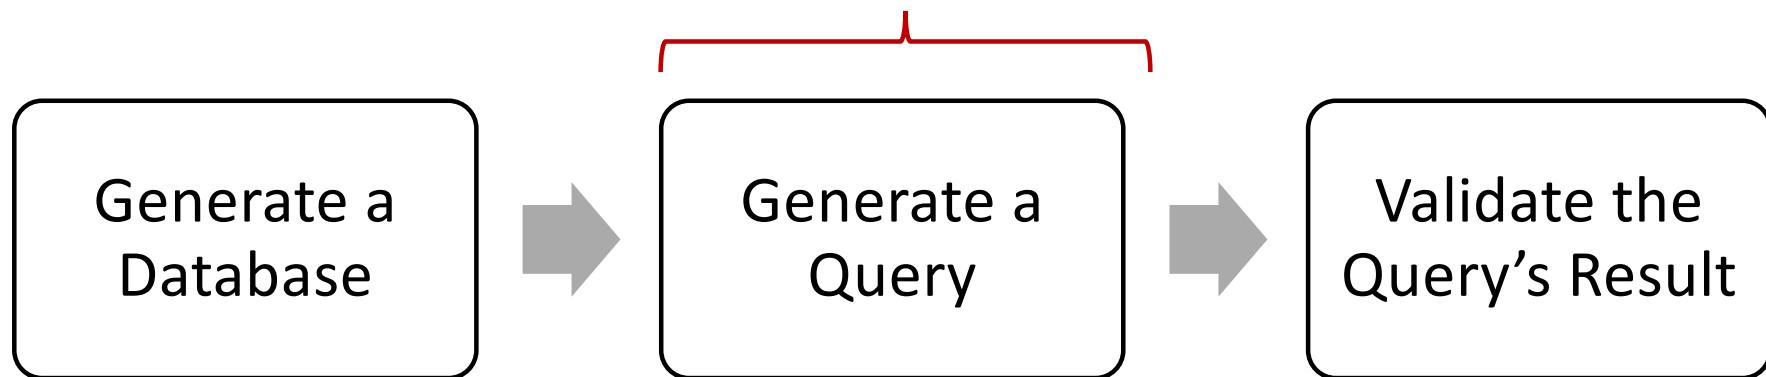
Automatic Testing Core Challenges

```
CREATE TABLE t0(c0 DOUBLE);  
CREATE TABLE t1(c0 DOUBLE);  
INSERT ...
```

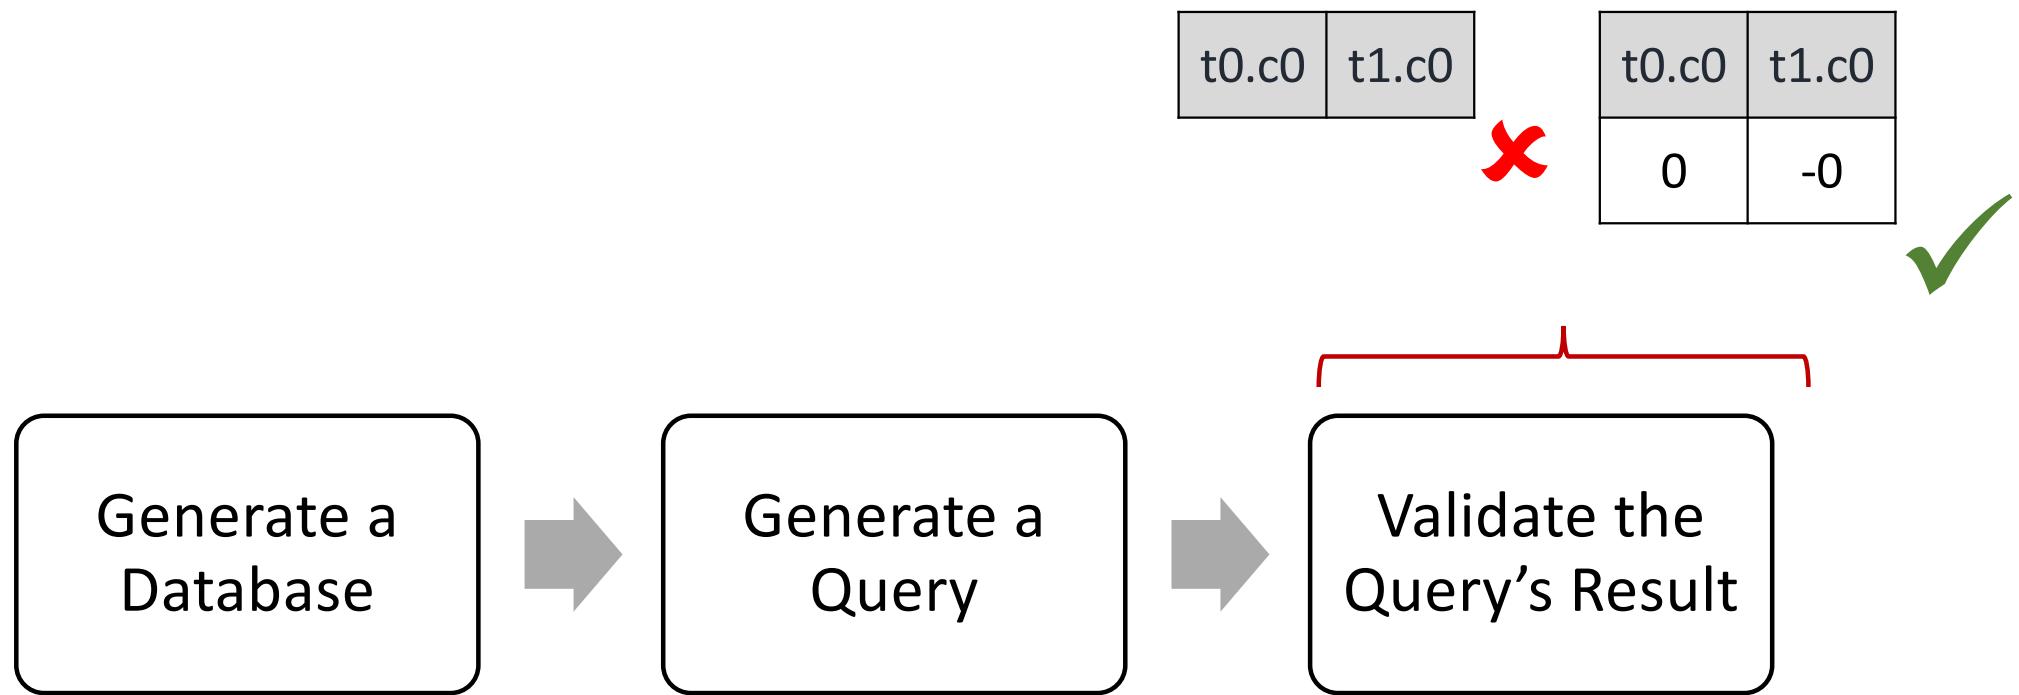


Automatic Testing Core Challenges

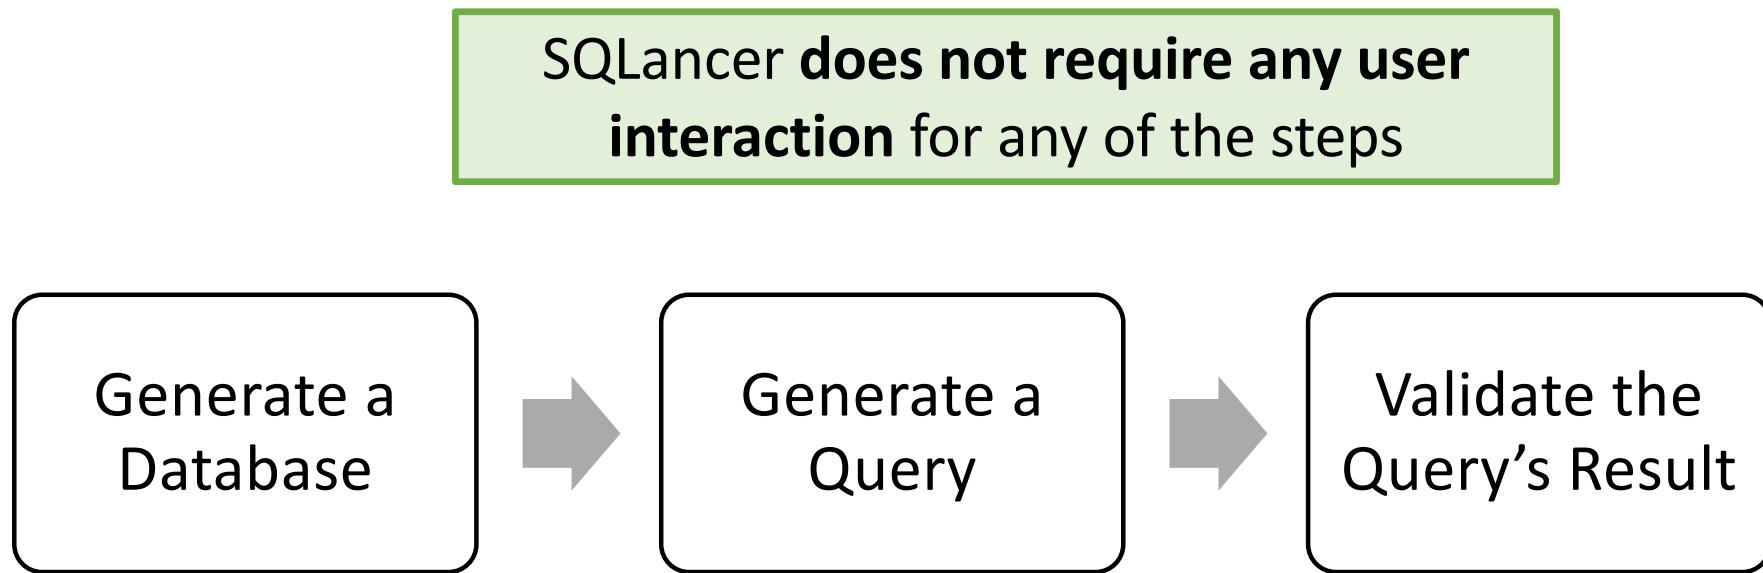
```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



Automatic Testing Core Challenges

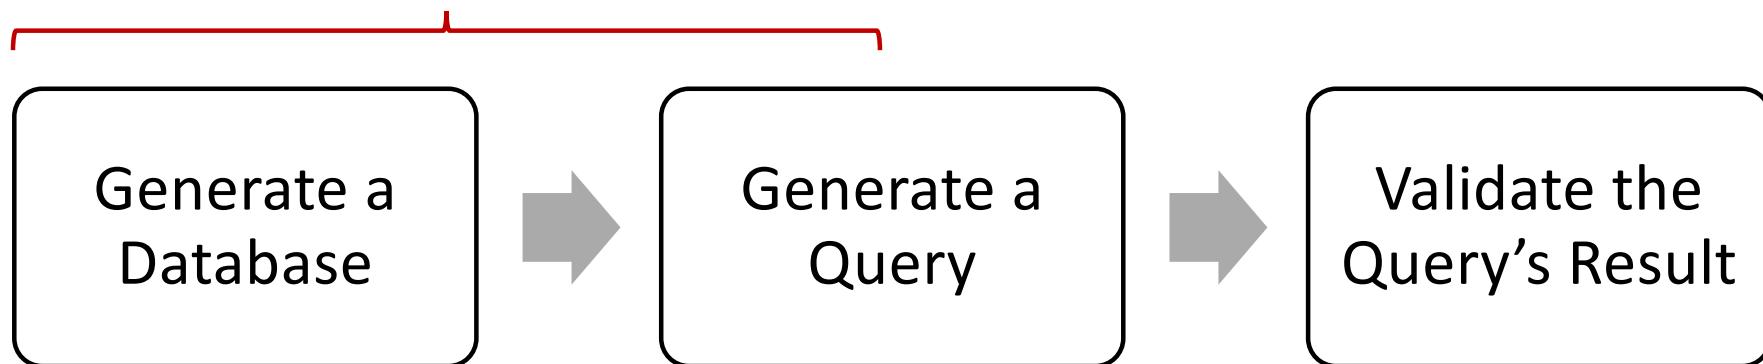


Automatic Testing Core Challenges



Automatic Testing Core Challenges

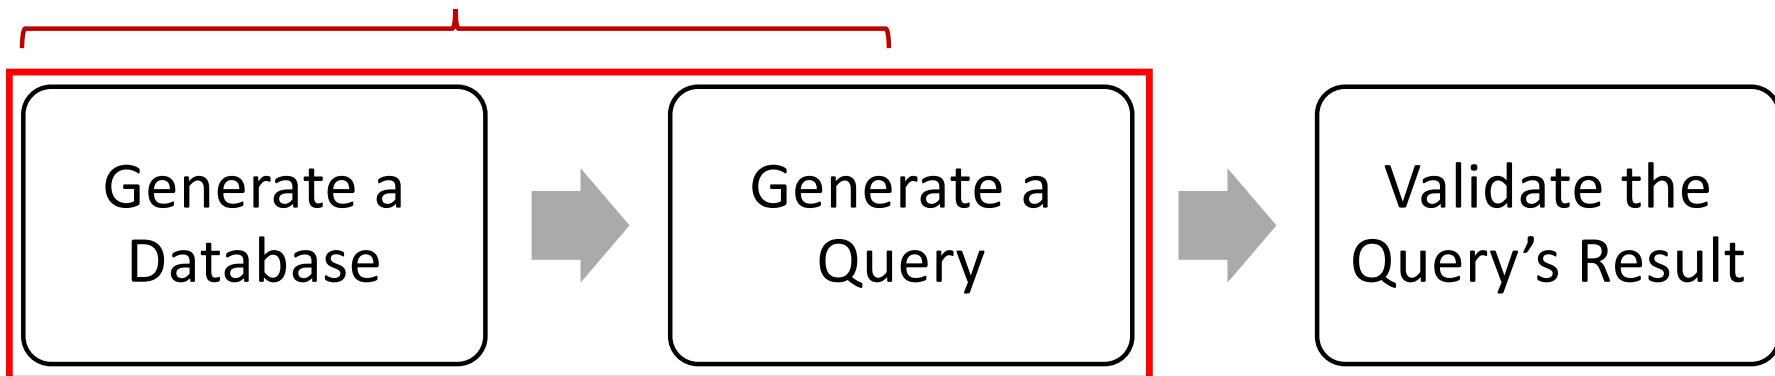
1. Effective test case



Automatic Testing Core Challenges

Manually implemented heuristic
database and query generators

1. Effective test case



SQLsmith

 Watch ▾ 29  Star 348  Fork 61


SQLsmith
`<mba> "I love the smell of core dumps in the morning"`

SQLancer's random database and query generation works similar to existing tools

Description

SQLsmith is a random SQL query generator. Its paragon is [Csmith](#), which proved valuable for quality assurance in C compilers.

It currently supports generating queries for PostgreSQL, SQLite 3 and MonetDB. To add support for another RDBMS, you need to implement two classes providing schema information about and connectivity to the device under test.

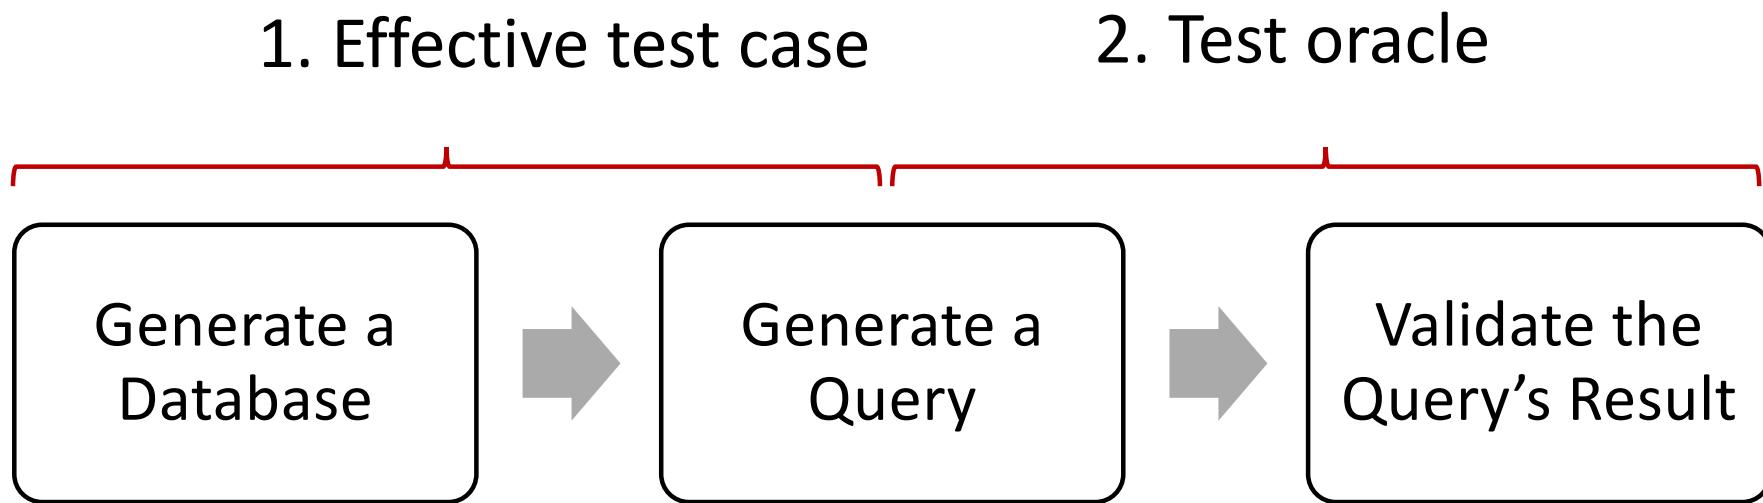
Besides developers of the RDBMS products, users developing extensions might also be interested in exposing their code to SQLsmith's random workload.

Since 2015, it found 118 bugs in alphas, betas and releases in the aforementioned products, including security vulnerabilities in released versions. Additional bugs were squashed in extensions and libraries such as orafce and glibc.

<https://github.com/anse1/sqlsmith/wiki#score-list>

<https://github.com/anse1/sqlsmith>

Automatic Testing Core Challenges

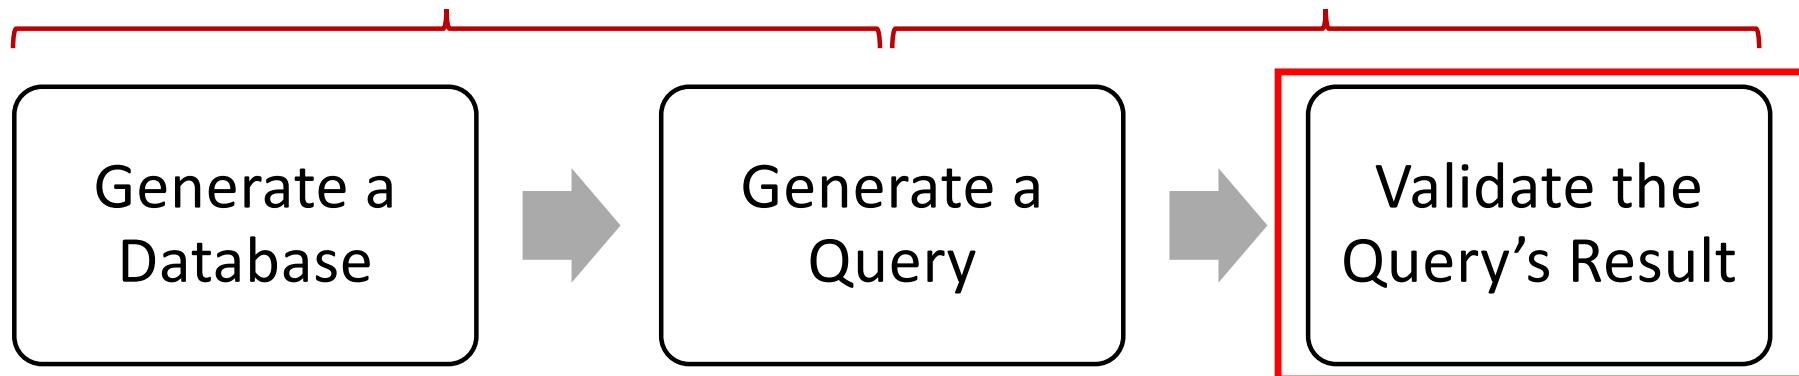


Automatic Testing Core Challenges

Test oracles are the “secret sauce”
of SQLancer

1. Effective test case

2. Test oracle



Test Oracle



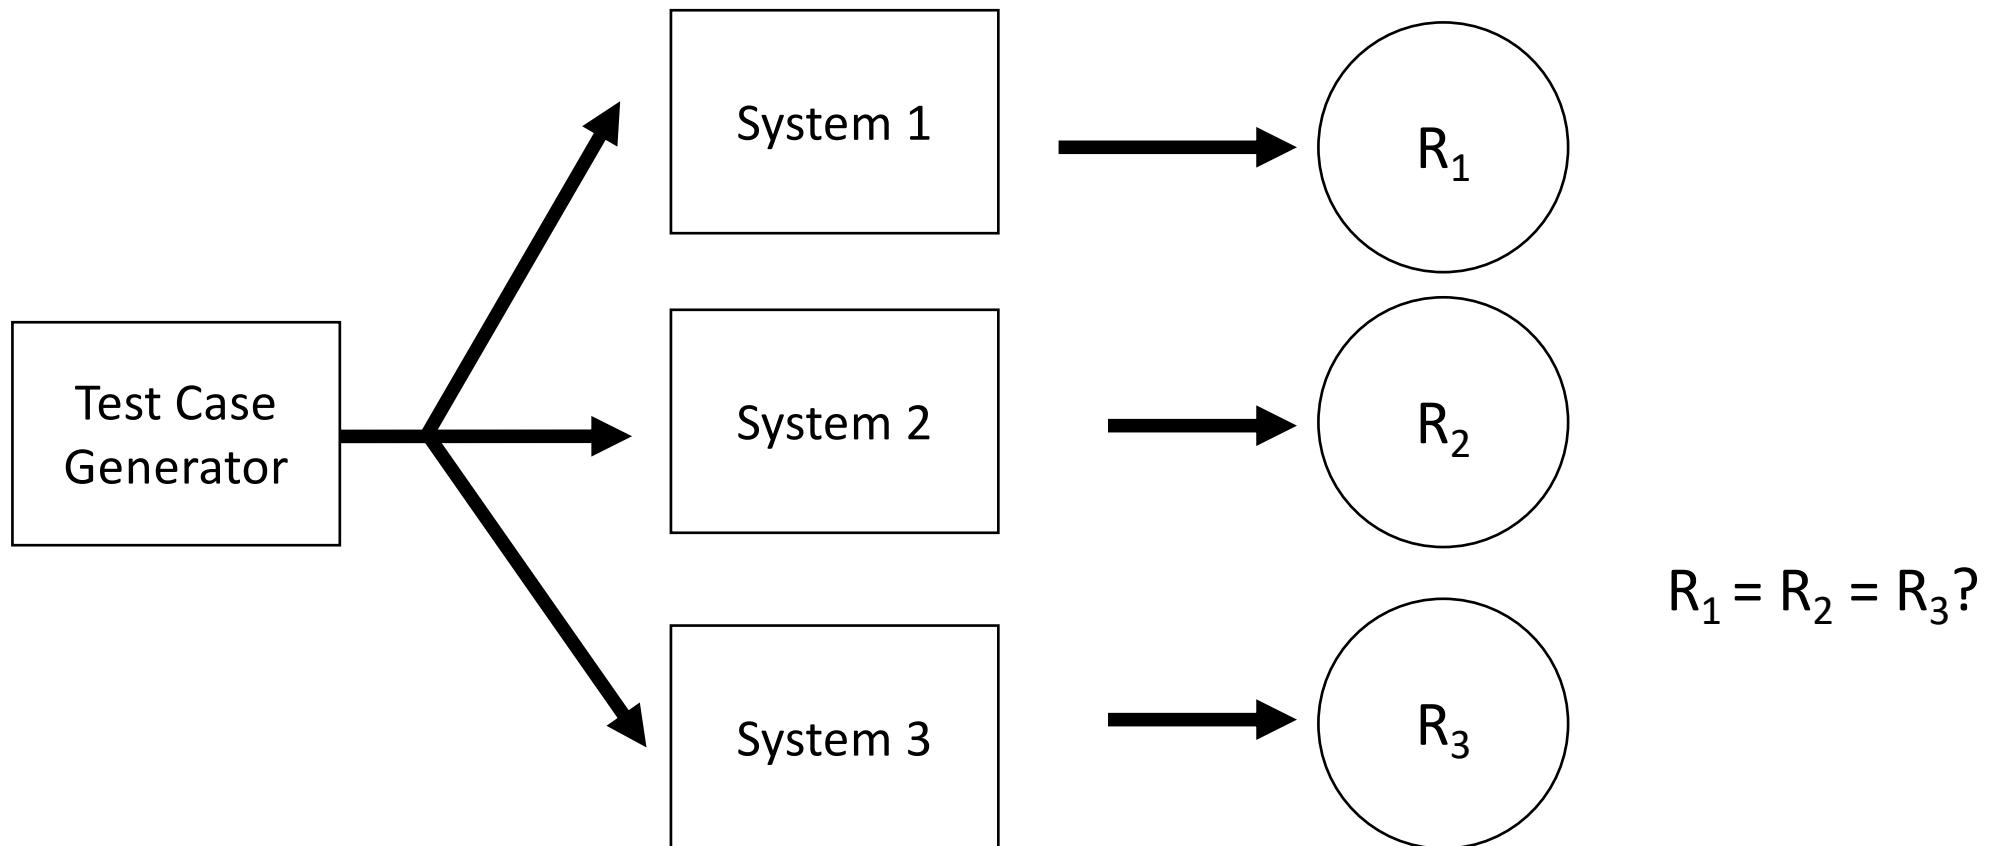
Incorrect result!

“a test oracle (or just oracle) is a mechanism for determining whether a test has passed or failed”

https://en.wikipedia.org/wiki/Test_oracle

**What existing test oracles could we use
to find bugs in DBMSs?**

Differential Testing



Differential Testing: RAGS (Slutz, VLDB 1998)

Massive Stochastic Testing of SQL

Don Slutz
Microsoft Research
dslutz@Microsoft.com

Abstract

Deterministic testing of SQL database systems is human intensive and cannot adequately cover the SQL input domain. A system (RAGS), was built to stochastically generate valid SQL statements 1 million times faster than a human and execute them.

1 Testing SQL is Hard

Good test coverage for commercial SQL database systems is very hard. The *input domain*, all SQL statements, from any number of users, combined with all states of the database, is gigantic. It is also difficult to verify output for positive tests because the semantics of SQL are complicated.

Software engineering technology exists to predictably improve quality ([Bei90] for example). The techniques involve a software development process including unit tests and final system validation tests (to verify the absence of bugs). This process requires a substantial investment so commercial SQL vendors with tight schedules tend to use a more ad hoc pro-

distribute the SQL statements in useful regions of the input domain. If the distribution is adequate, stochastic testing has the advantage that the quality of the tests improves as the test size increases [TFW93].

A system called RAGS (Random Generation of SQL) was built to explore automated testing. RAGS is currently used by the Microsoft SQL Server [MSS98] testing group. This paper describes RAGS and some illustrative test results.

Figure 1 illustrates the test coverage problem. Customers use the hexagon, bugs are in the oval, and the test libraries cover the shaded circle.

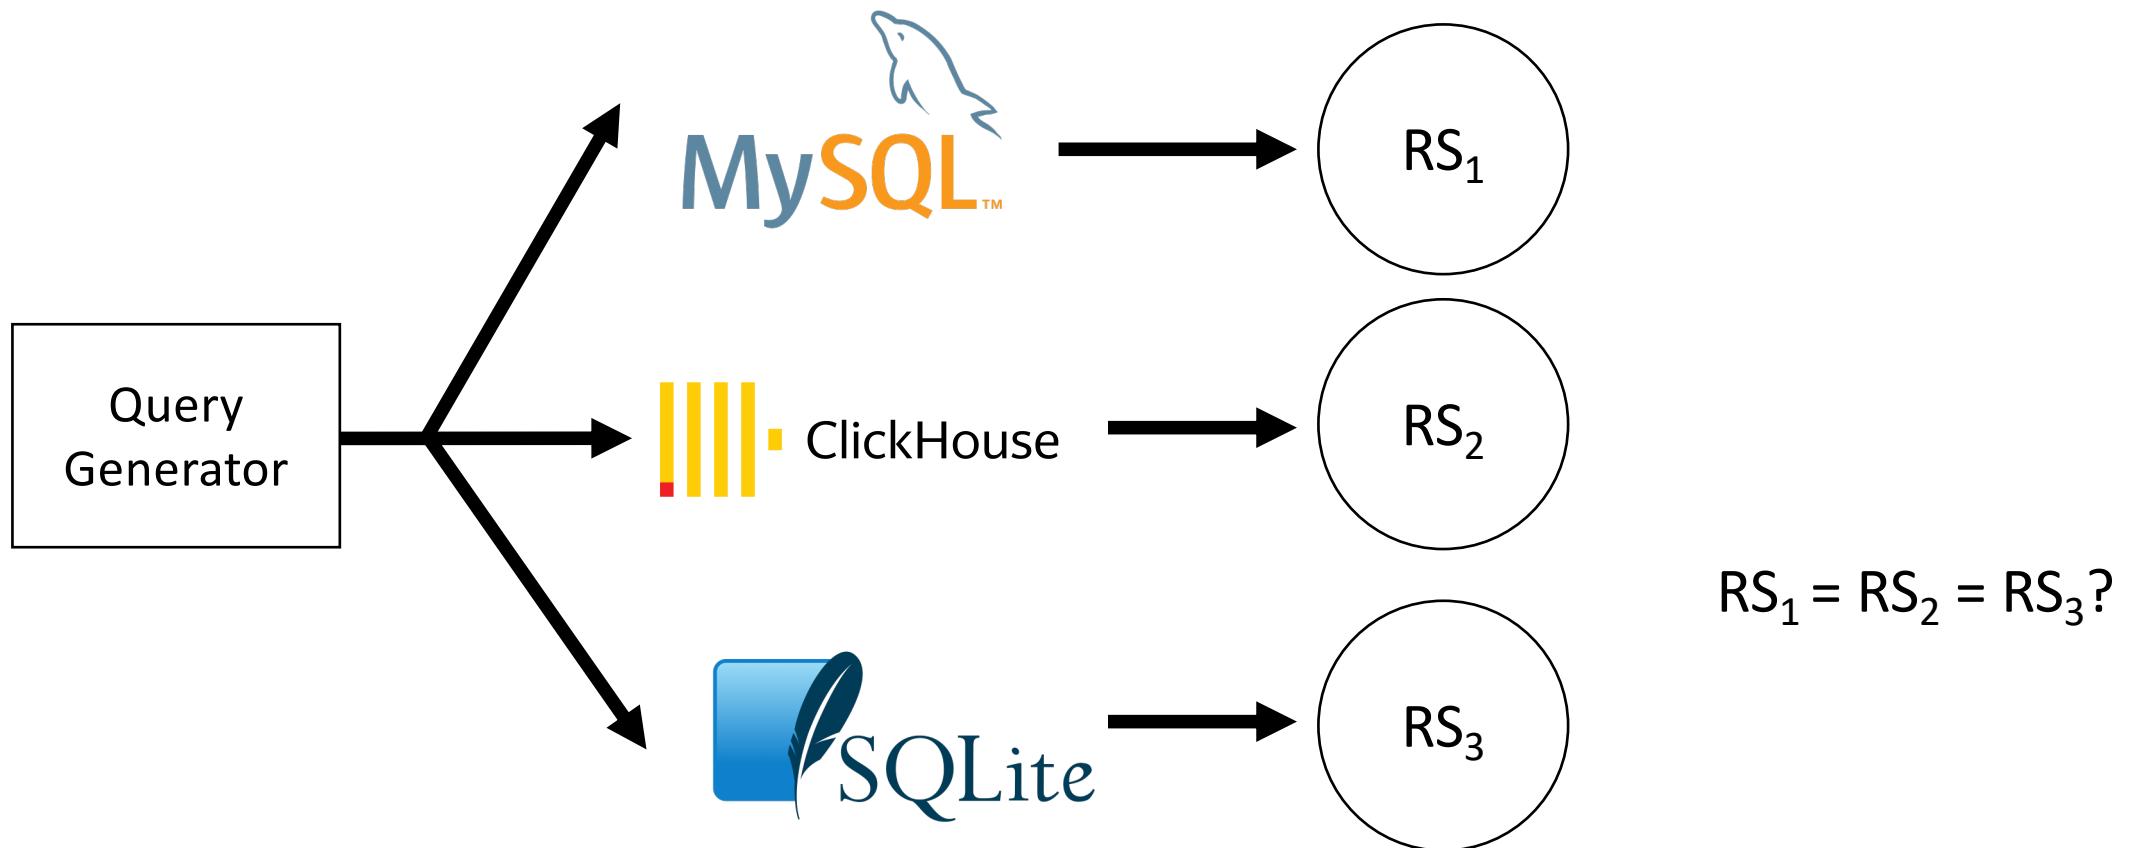
Input Domain

All possible SQL statements and database states

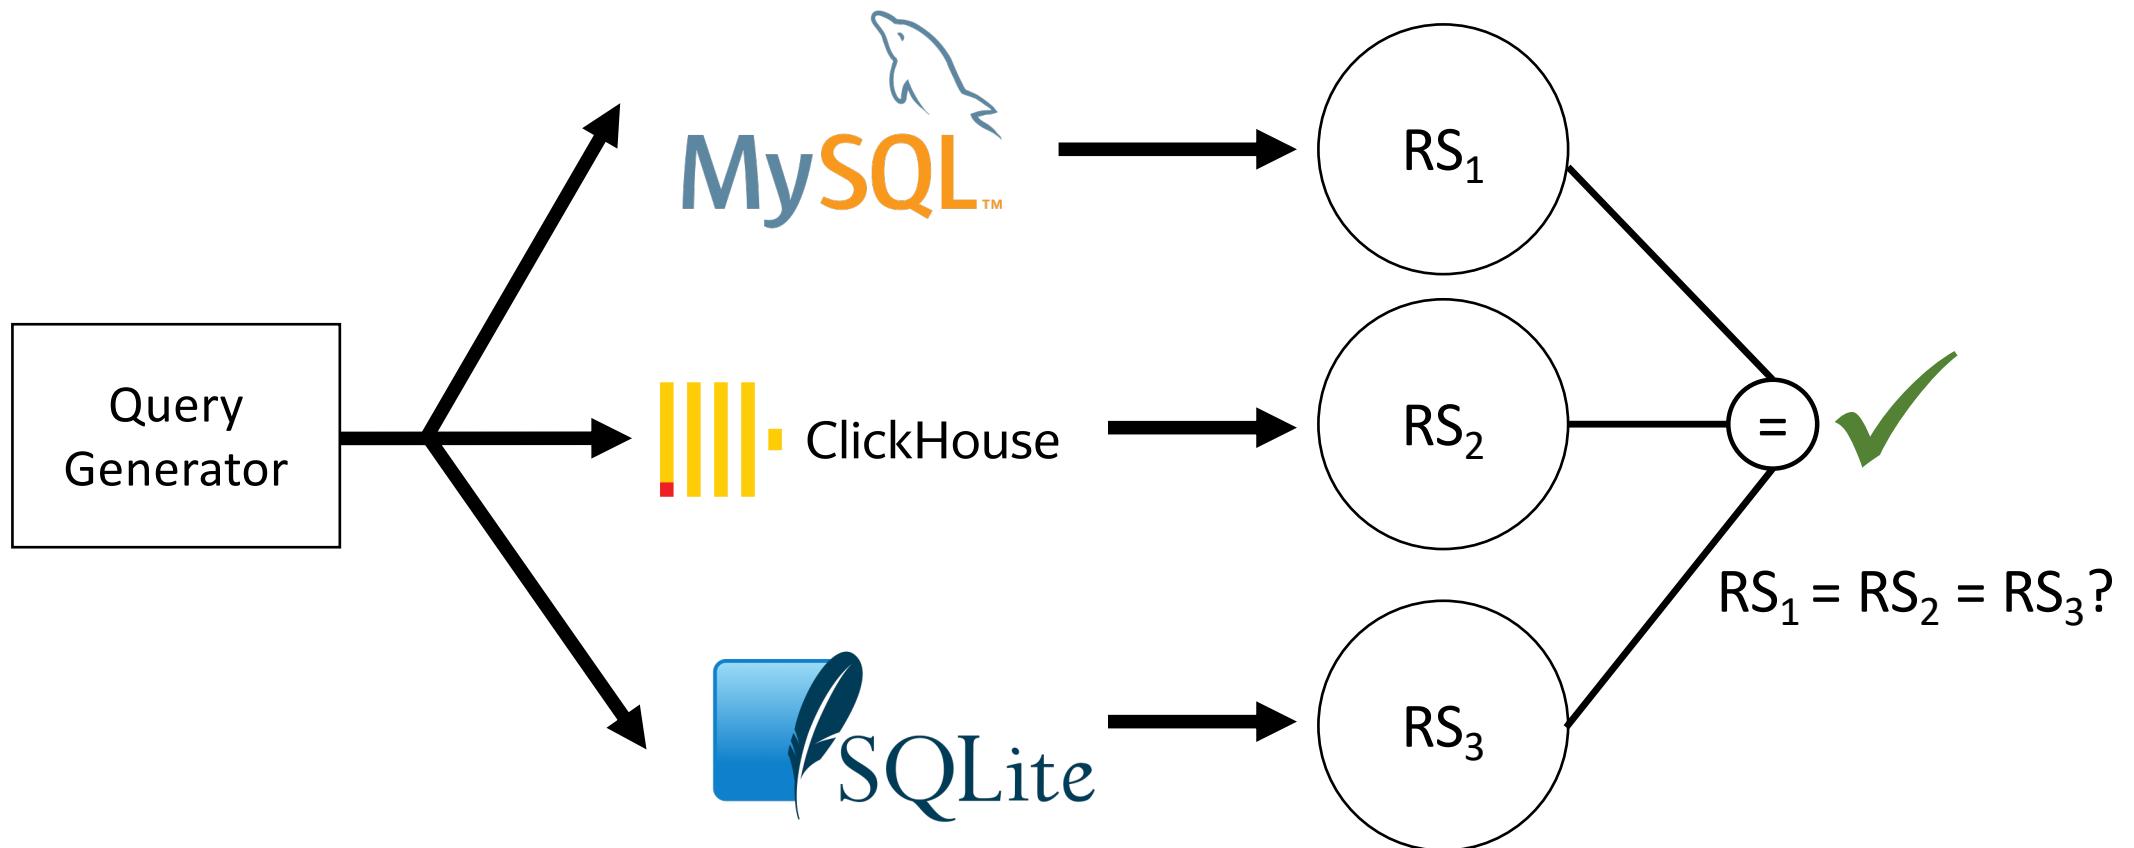
Used by customers

Bugs

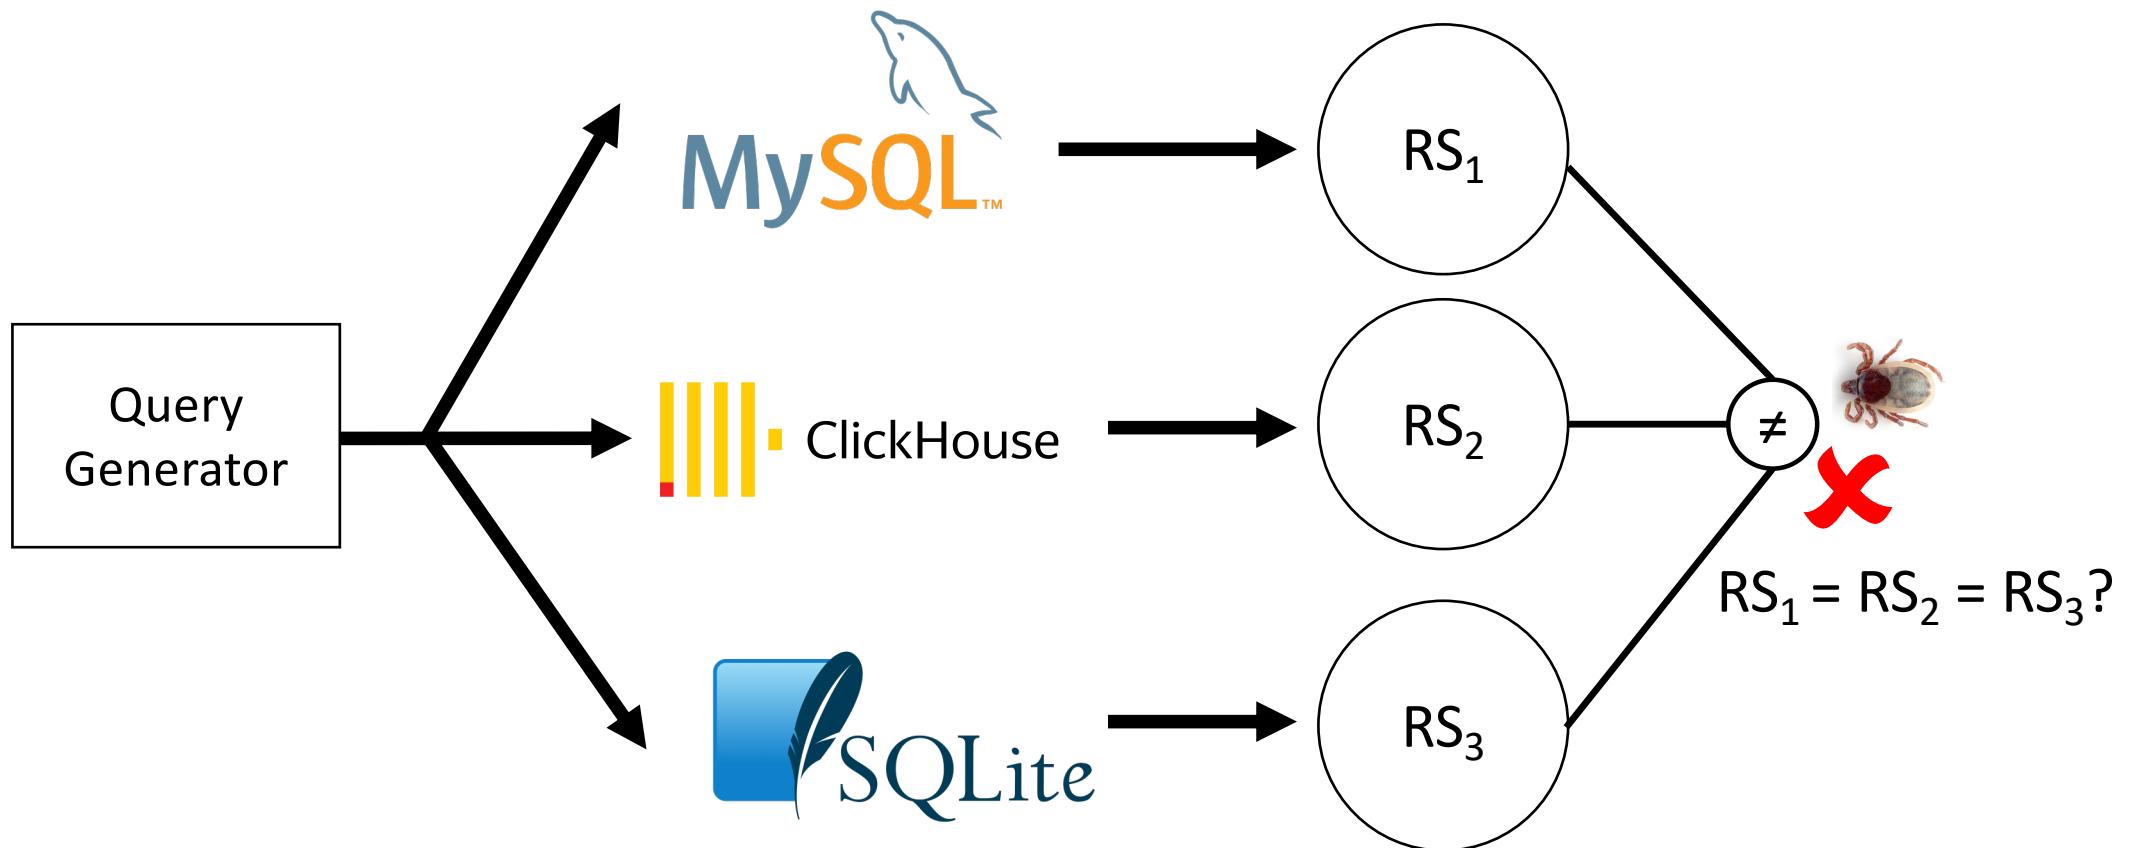
Differential Testing: RAGS (Slutz, VLDB 1998)



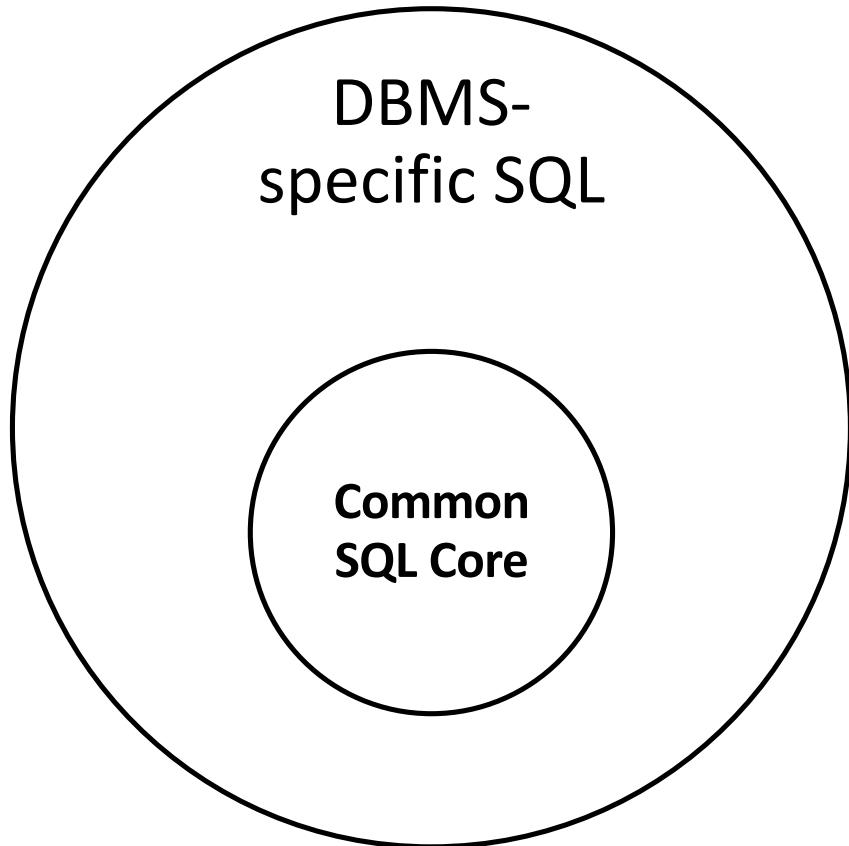
Differential Testing: RAGS (Slutz, VLDB 1998)



Differential Testing: RAGS (Slutz, VLDB 1998)



Differential Testing: RAGS (Slutz, VLDB 1998)



“We are **unable to use Postgres** as an oracle because CockroachDB has **slightly different semantics** and SQL support, and generating queries that execute identically on both is tricky [...].” – Cockroach Labs

What oracles were introduced in SQLancer?

Finding Logic Bugs in DBMSs

Ternary Logic
Partitioning

OOPSLA '20

SQLancer provides
three test oracles

Non-optimizing
Reference Engine
Construction

ESEC/FSE '20

Pivoted Query
Synthesis

OSDI '20

Finding Logic Bugs in DBMSs

Ternary Logic
Partitioning

OOPSLA '20

Non-optimizing
Reference Engine
Construction

ESEC/FSE '20

Pivoted Query
Synthesis

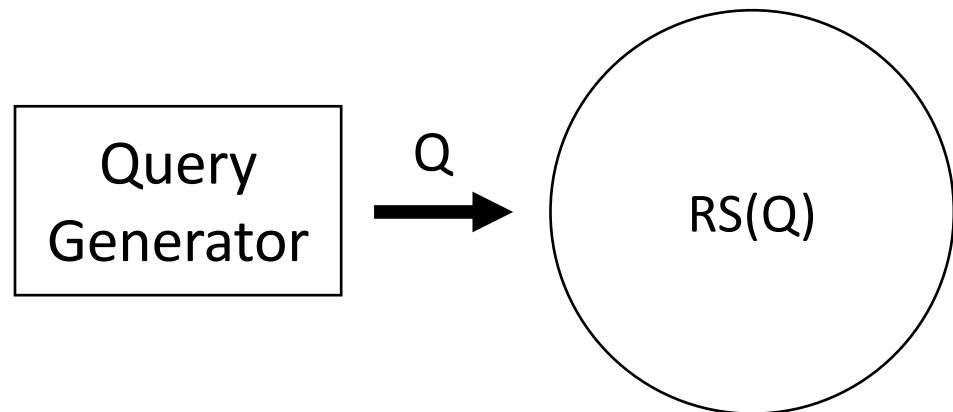
OSDI '20

Query Partitioning

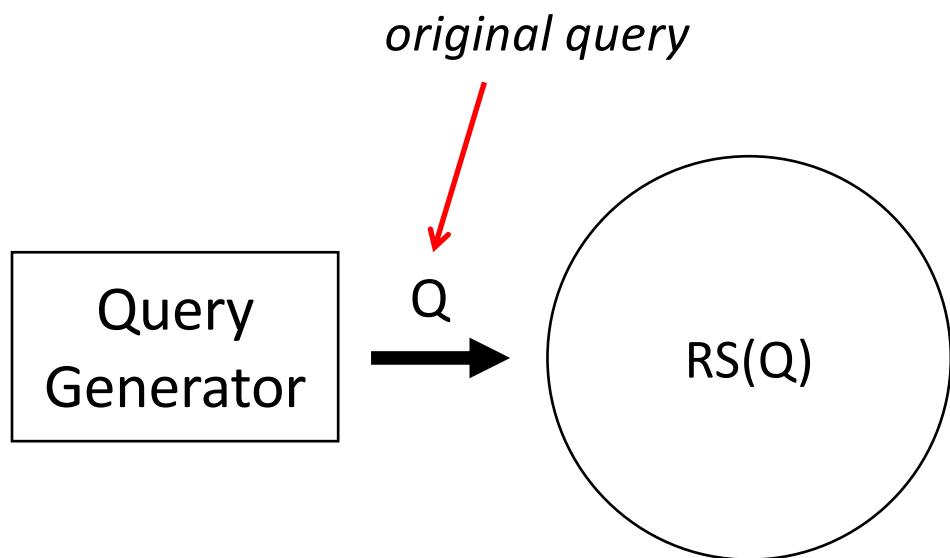


Ternary Logic Partitioning (TLP) is based on a conceptual framework called **Query Partitioning**

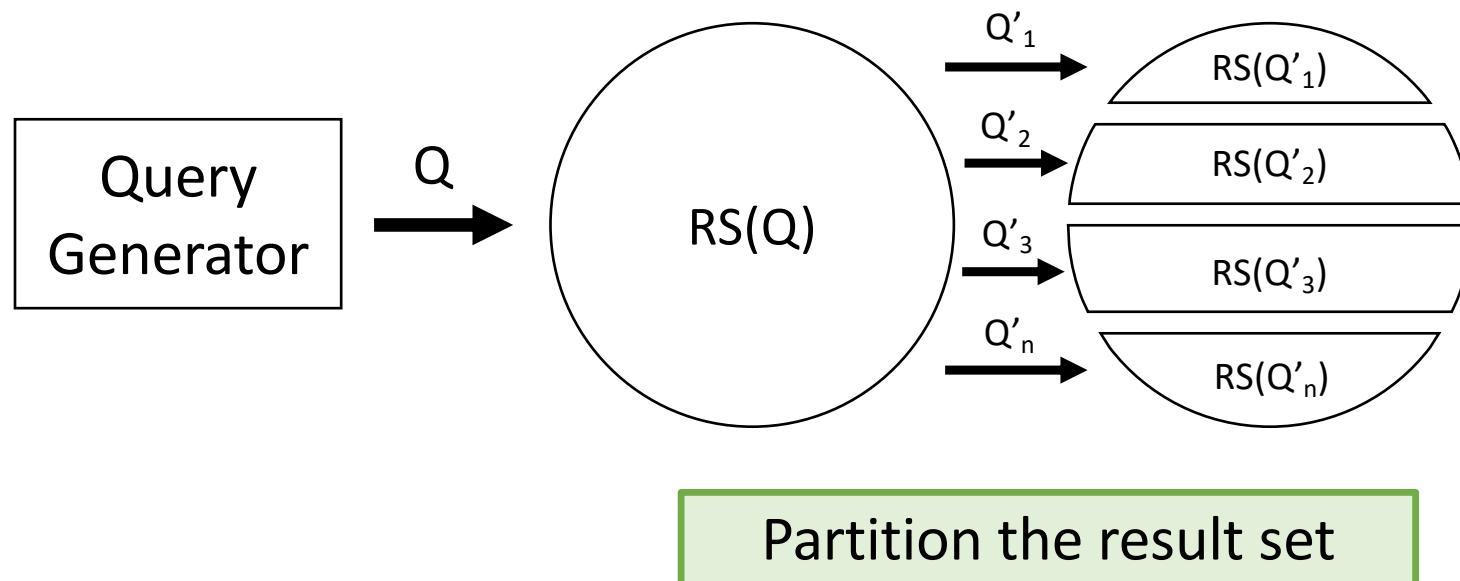
Query Partitioning



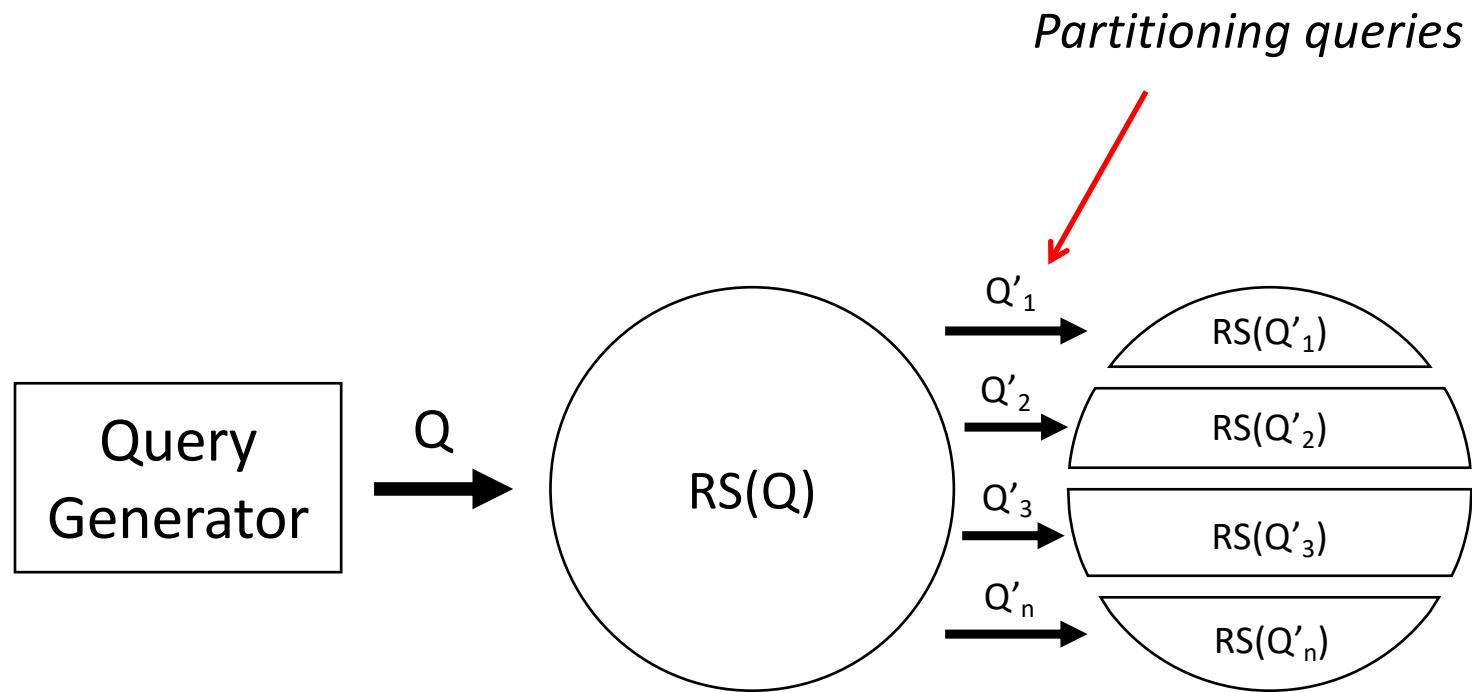
Query Partitioning



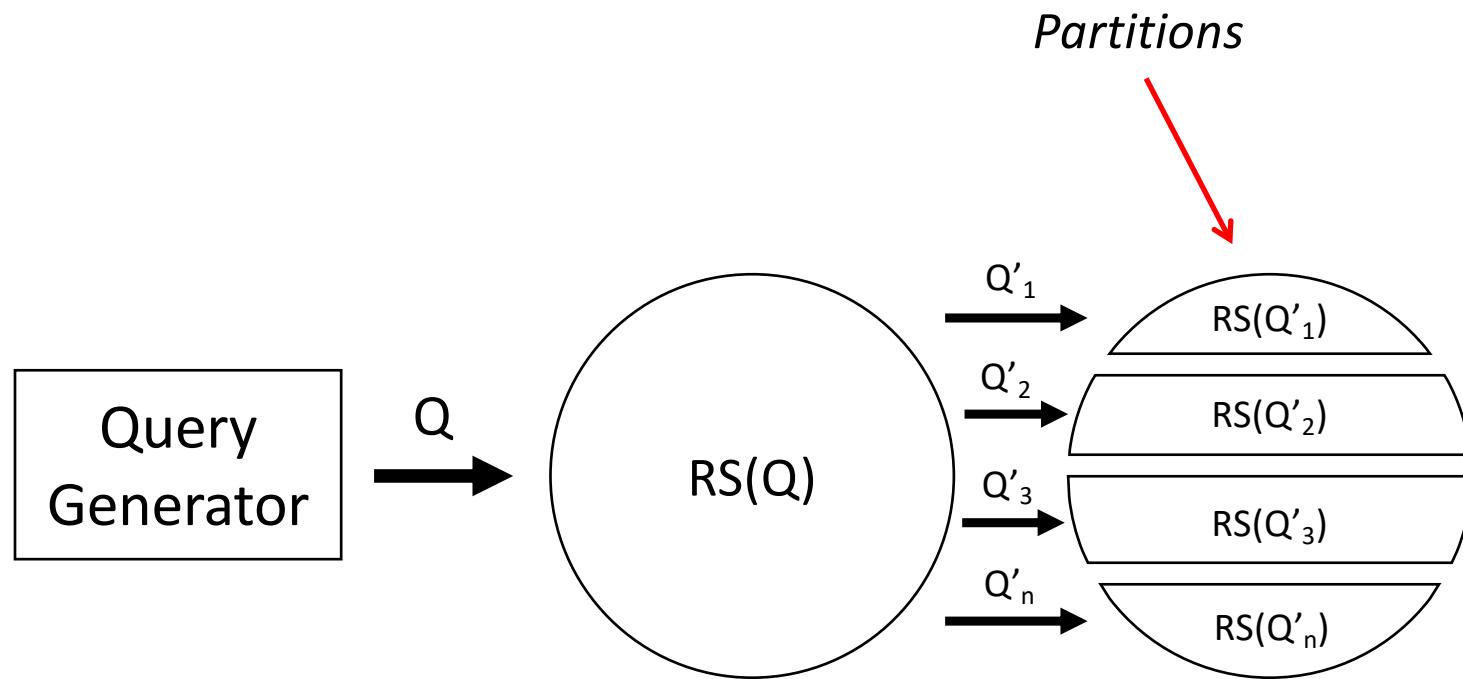
Query Partitioning



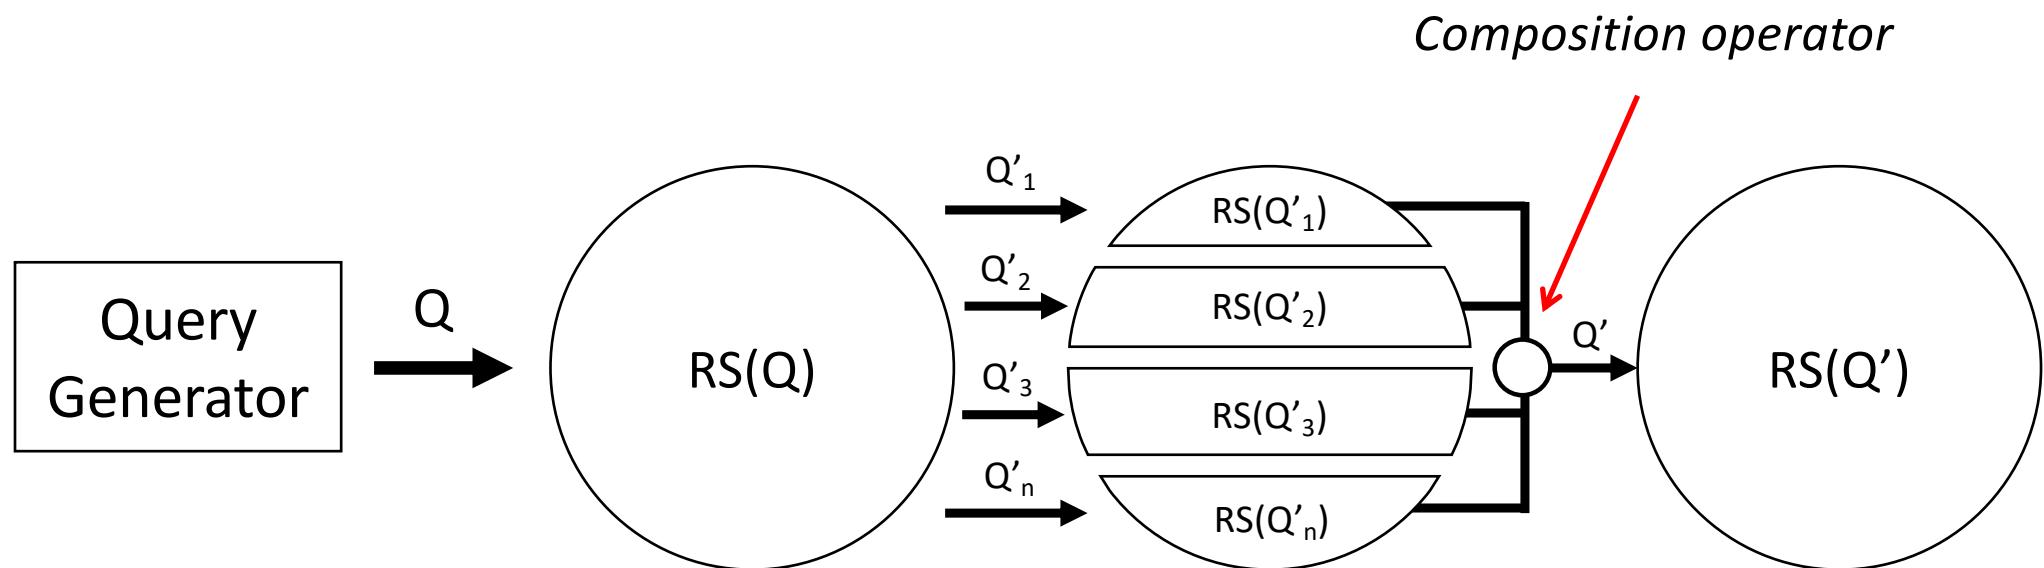
Query Partitioning



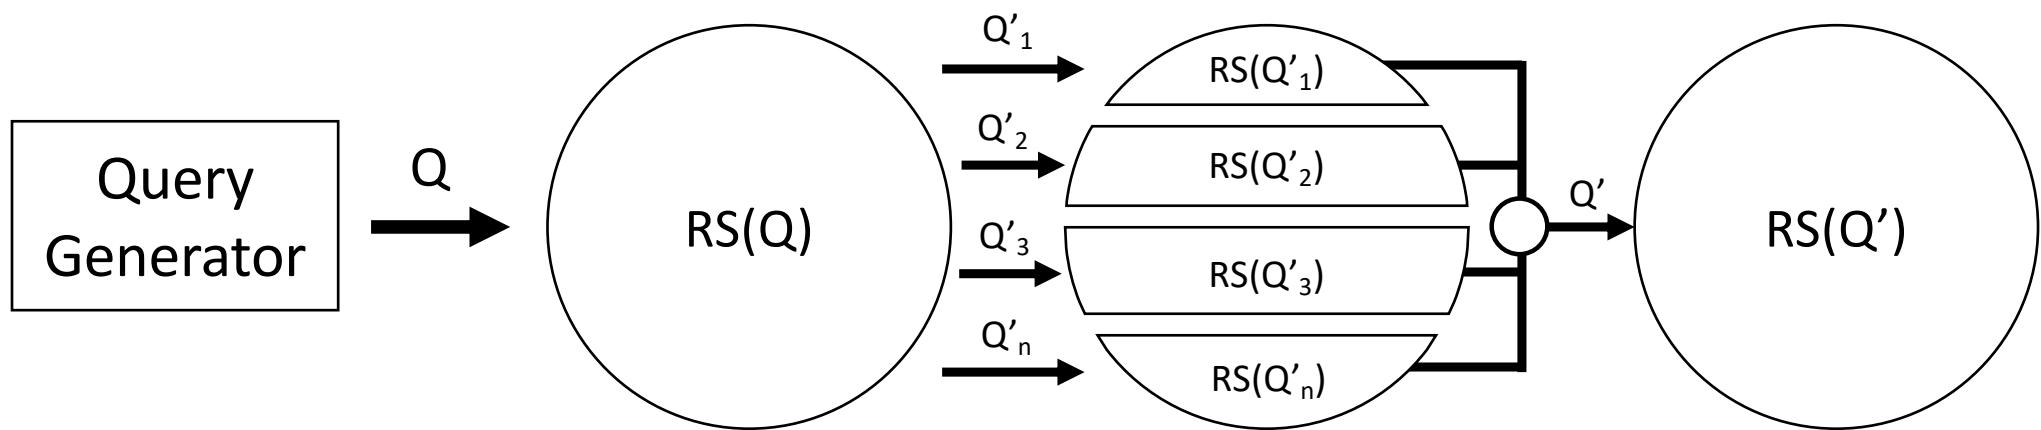
Query Partitioning



Query Partitioning

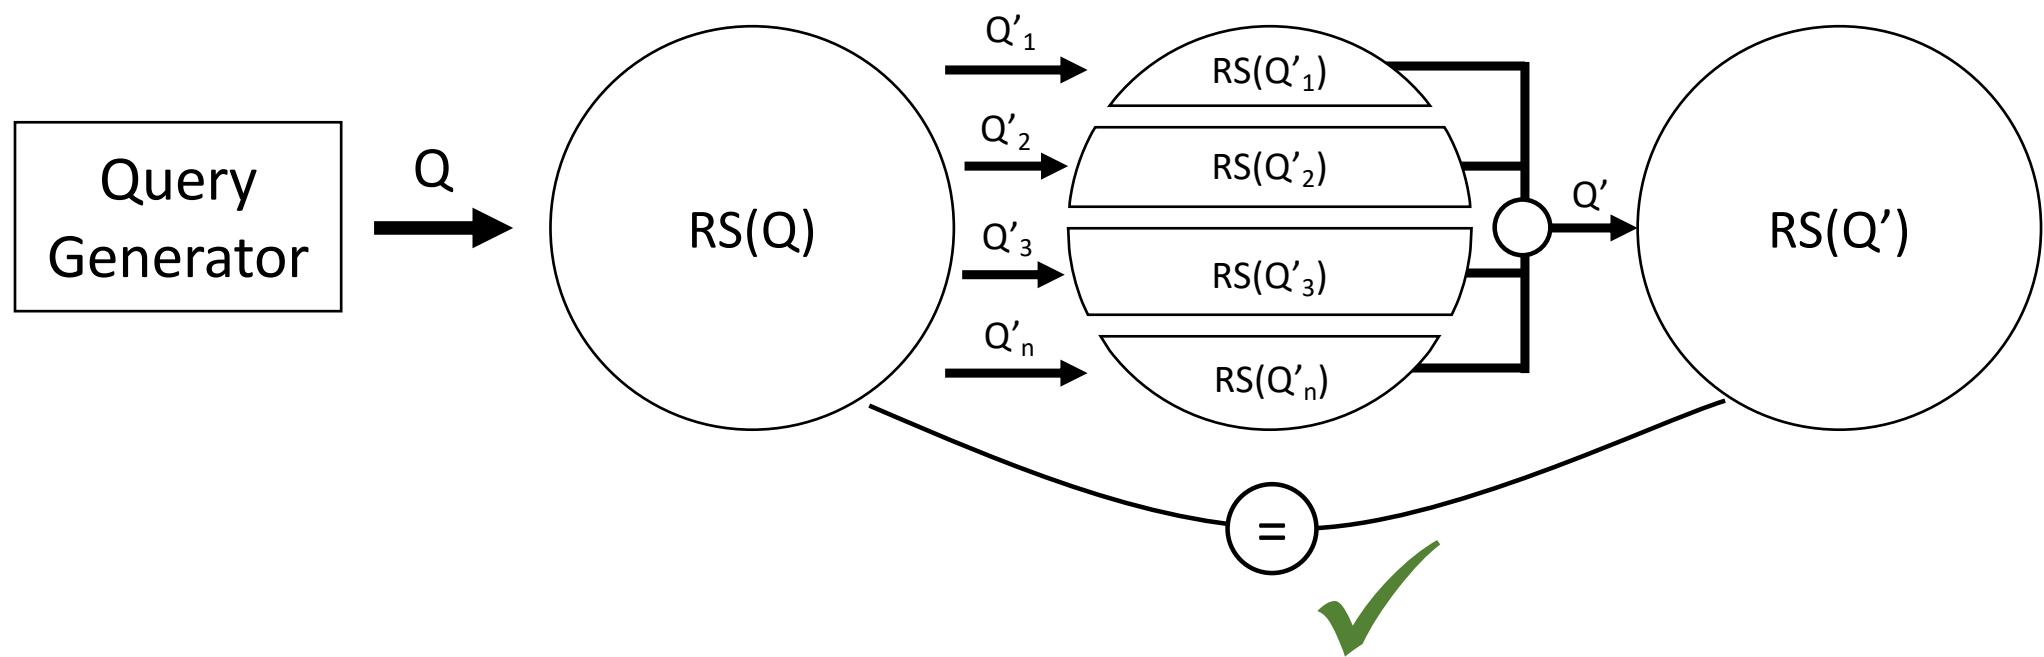


Query Partitioning

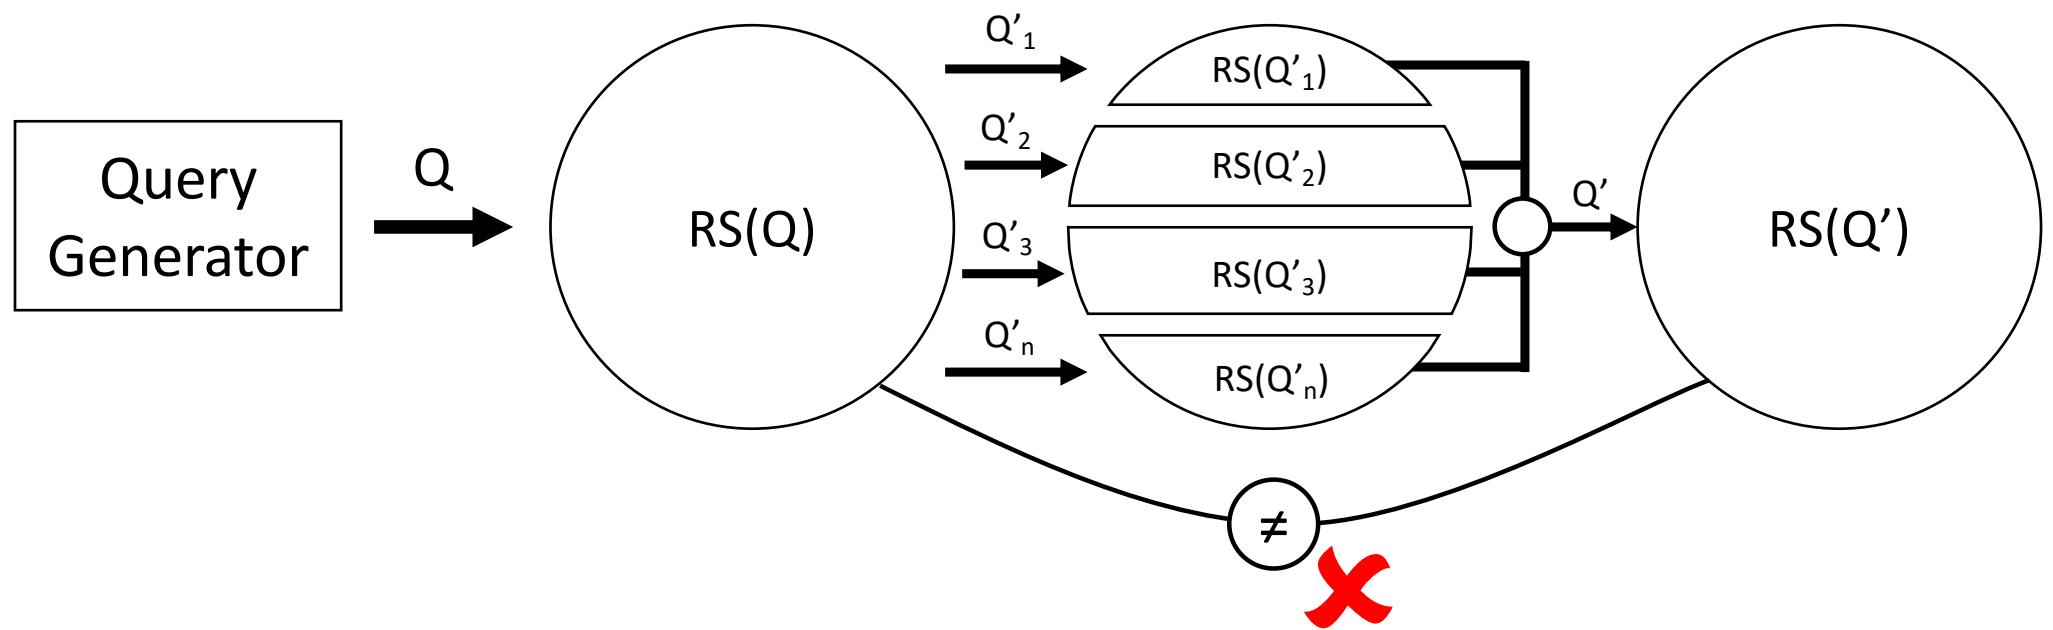


Combine the results so that
 $RS(Q)=RS(Q')$

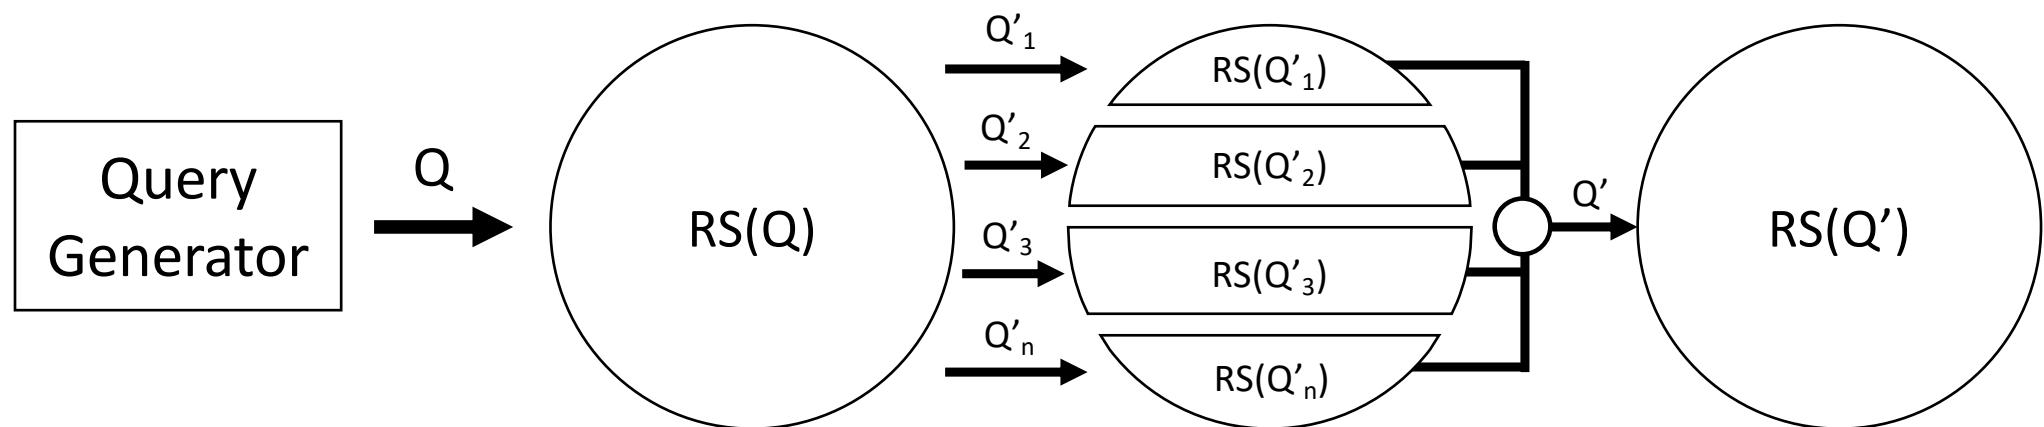
Query Partitioning



Query Partitioning



Query Partitioning



In contrast to differential testing, we
need **only a single DBMS**

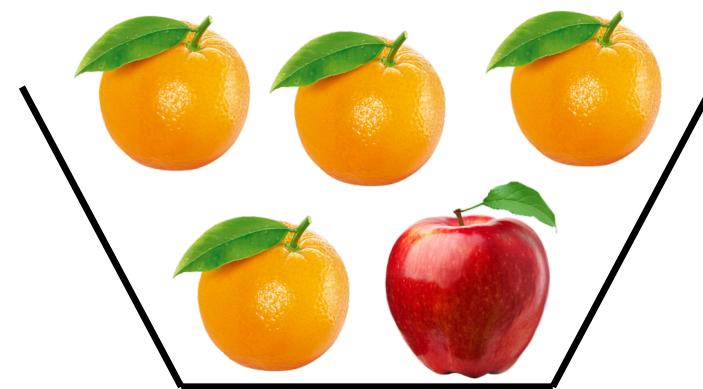
How to Realize This Idea?

How can we partition the result set?

Scenario: Coffee Kitchen



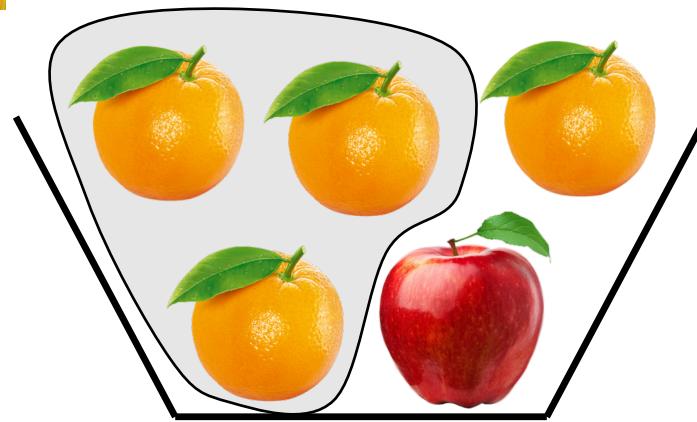
Tangerines vs. Clementines



Tangerines vs. Clementines



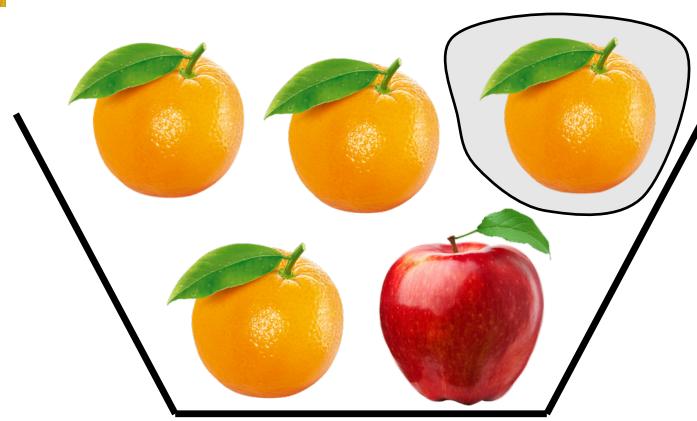
tangerines



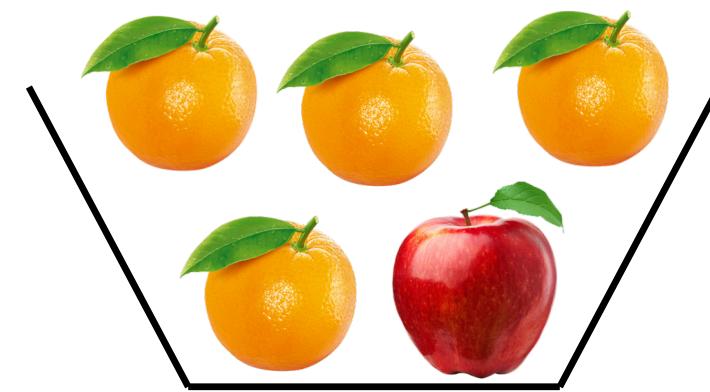
Tangerines vs. Clementines



clementine

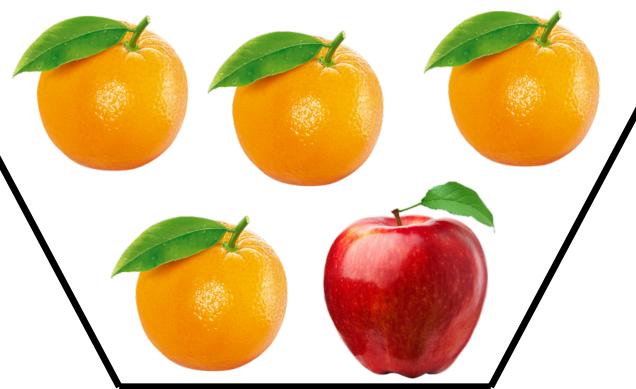


Tangerines vs. Clementines

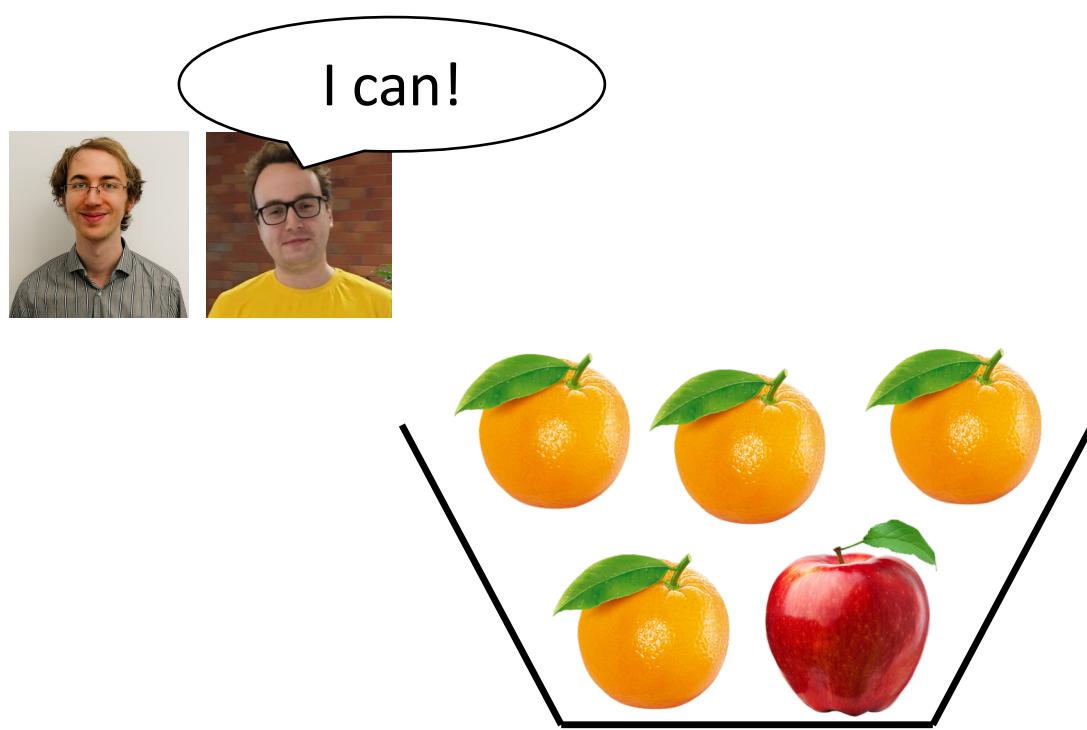


Tangerines vs. Clementines

I can never tell different citrus
fruits apart

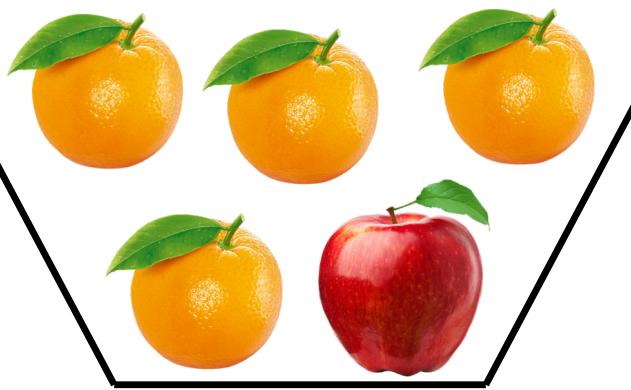


Tangerines vs. Clementines

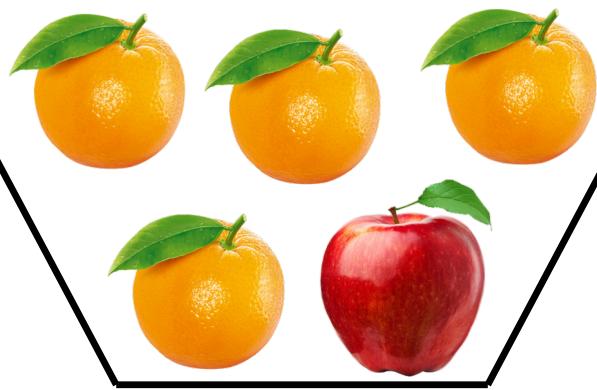


Tangerines vs. Clementines

Show me!



How I test Ilya's understanding without
knowing the differences myself?

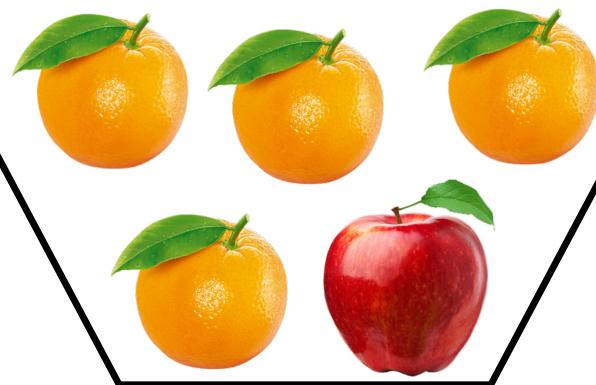


Tangerines vs. Clementines

Please bring me all
clementines



2 fruits



Tangerines vs. Clementines

Please bring me all
clementines

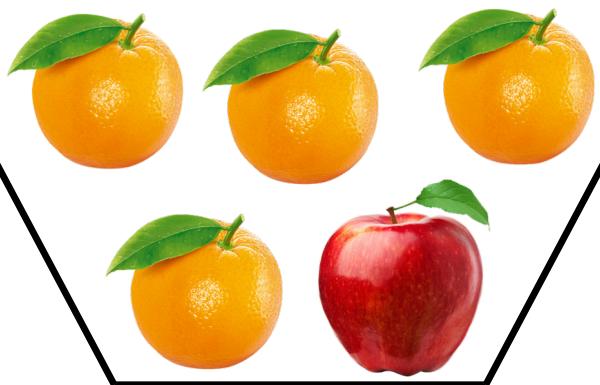


Tangerines vs. Clementines

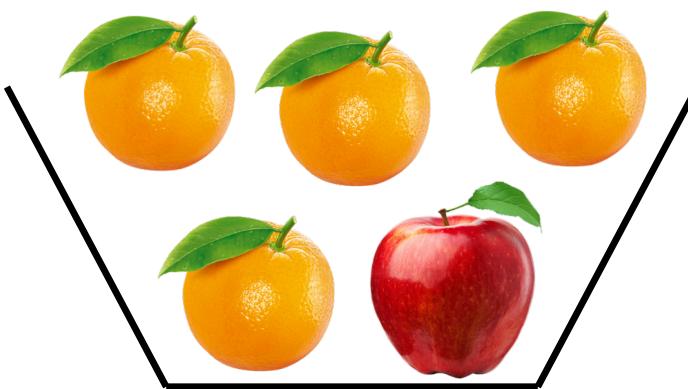
Please bring me all fruits
that are **not clementines**



4 fruits



Tangerines vs. Clementines



2 fruits
4 fruits

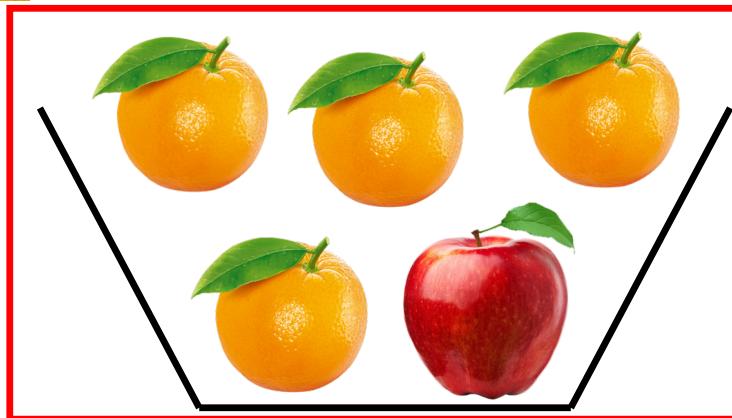
6 fruits

Tangerines vs. Clementines

You likely classified a fruit as **both**
a tangerine and a clementine!



5 fruits



2 fruits

4 fruits

6 fruits

Insight



Insight: Every object in a (mathematical) universe
is either **a tangerine or not a tangerine**

**How can we apply this idea
to find bugs in DBMSs?**

Ternary Logic

Consider a predicate p and a given row r .
Exactly **one** of the following must hold:

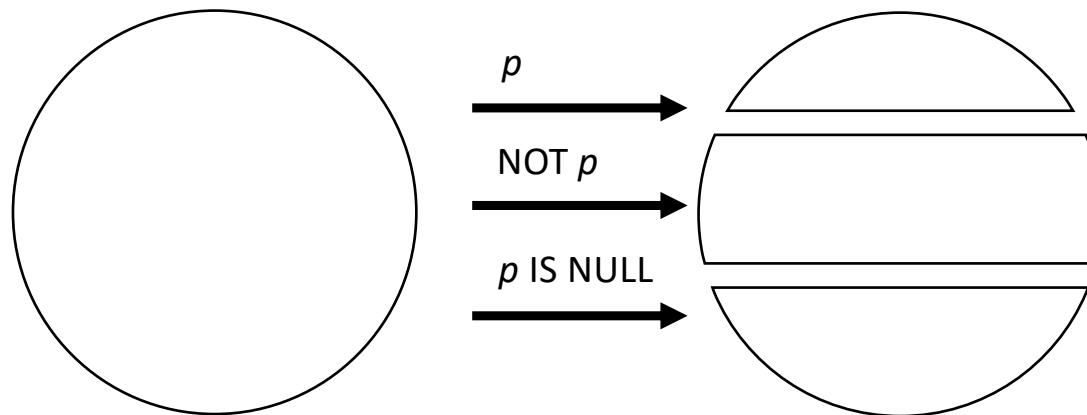
- p
- NOT p
- p IS NULL

Ternary Logic

Consider a predicate p and a given row r .

Exactly **one** of the following must hold:

- p
- $\text{NOT } p$
- $p \text{ IS NULL}$



Motivating Example

t0	t1
c0	c0
0	-0

How did this insight allow us to detect this bug?



```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



t0.c0	t1.c0
-------	-------

X

Example: MySQL

```
SELECT * FROM t0, t1;
```



t0.c0	t1.c0
0	-0

```
SELECT * FROM t0, t1 WHERE t0.c0=t1.c0
```

```
UNION ALL
```

```
SELECT * FROM t0, t1 WHERE NOT (t0.c0=t1.c0)
```

```
UNION ALL
```

```
SELECT * FROM t0, t1 WHERE (t0.c0=t1.c0) IS NULL;
```

p
↓



t0.c0	t1.c0
0	-0



Insight



The DBMS is **more likely** to process the **partitioning queries incorrectly** due to their higher complexity

**What kind of features we can apply
Query Partitioning testing oracle?**

Scope

- WHERE
- GROUP BY
- HAVING
- DISTINCT queries
- Aggregate functions

Testing WHERE Clauses

Q	Q'_{ptern}	Composition Operator
SELECT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE ptern	UNION ALL

Testing WHERE Clauses

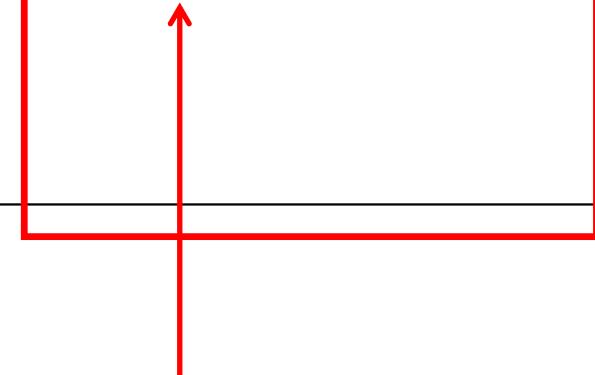
Q	Q'_{ptern}	Composition Operator
SELECT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE ptern	UNION ALL

Testing WHERE Clauses

Q	Q'_{ptern}	Composition Operator
SELECT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE ptern	UNION ALL

Testing WHERE Clauses

Q	Q'_{ptern}	Composition Operator
SELECT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE p_{tern}	UNION ALL



UNION ALL keeps duplicate rows

Scope

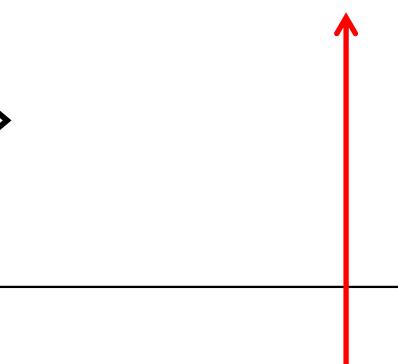
- WHERE
- GROUP BY
- HAVING
- DISTINCT queries
- Aggregate functions

Testing DISTINCT Clauses

Q	Q'_{ptern}	Composition operator
SELECT DISTINCT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE p_{tern};	UNION

Testing DISTINCT Clauses

Q	Q'_{ptern}	Composition operator
SELECT DISTINCT <columns> FROM <tables> [<joins>]	SELECT <columns> FROM <tables> [<joins>] WHERE p_{tern};	UNION



UNION removes duplicate rows

Scope

- WHERE
- GROUP BY
- HAVING
- DISTINCT queries
- Aggregate functions

Testing Self-decomposable Aggregate Functions

Q	Q'_{ptern}	Composition operator
SELECT MAX(<e>) FROM <tables> [<joins>]	SELECT MAX(<e>) FROM <tables> [<joins>] WHERE p_{tern};	MAX

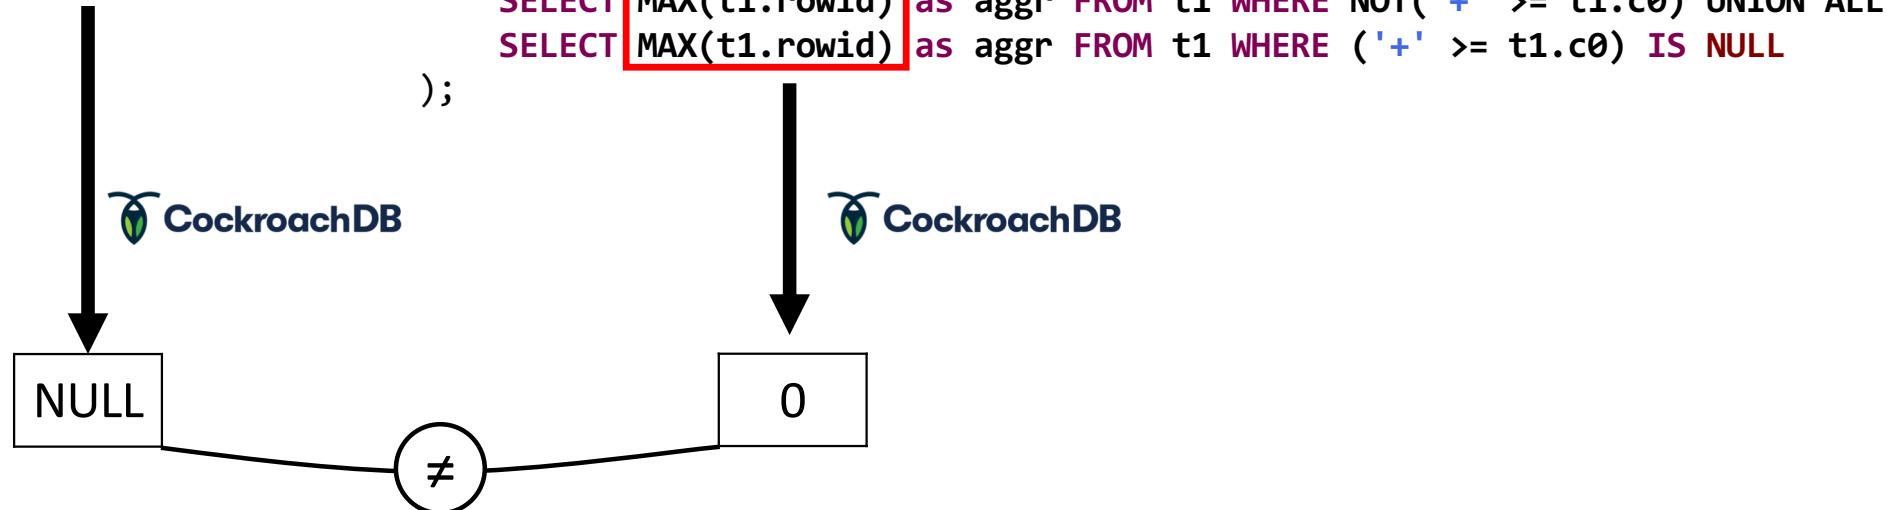


A partition is an **intermediate result**, rather than a subset of the result set

Bug Example: CockroachDB

```
SET vectorize=experimental_on;
CREATE TABLE t0(c0 INT);
CREATE TABLE t1(c0 BOOL) INTERLEAVE IN PARENT t0(rowid);
INSERT INTO t0(c0) VALUES (0);
INSERT INTO t1(rowid, c0) VALUES(0, TRUE);

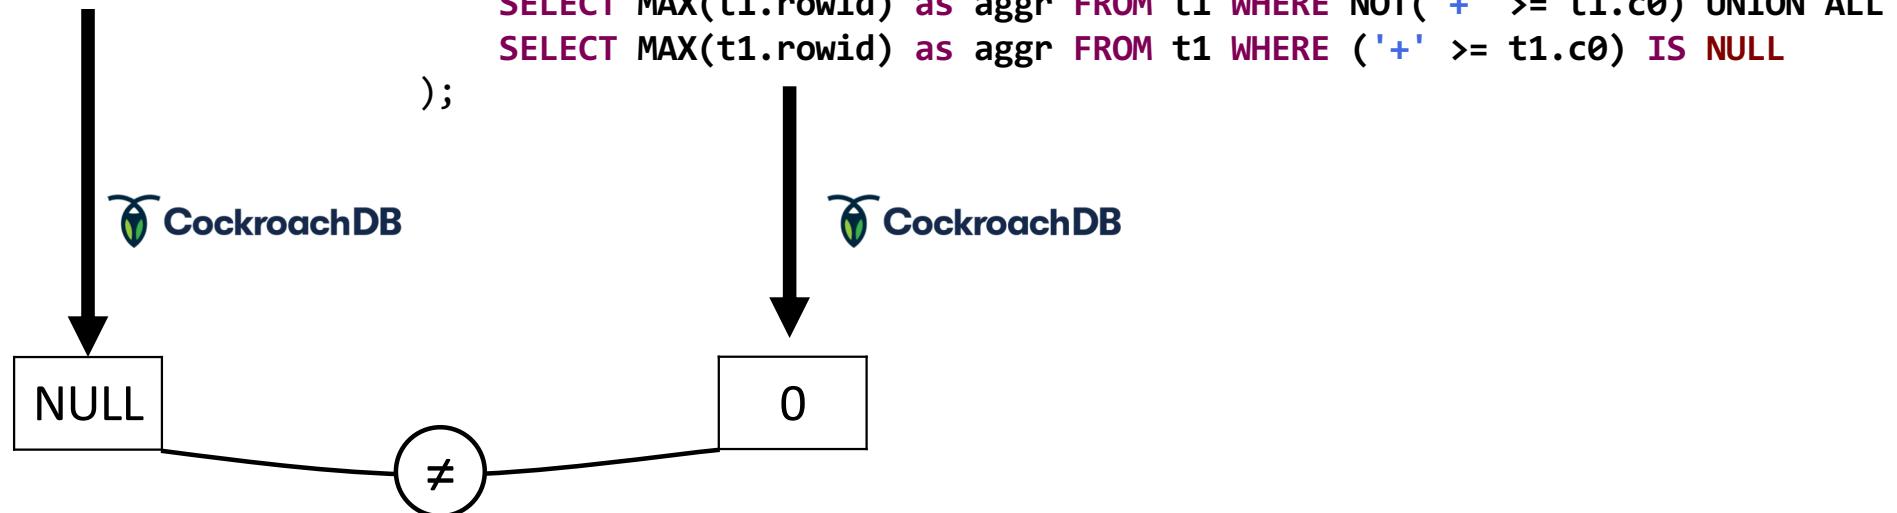
SELECT MAX(t1.rowid) FROM t1;
SELECT MAX(aggr) FROM (
    SELECT MAX(t1.rowid) as aggr FROM t1 WHERE '+' >= t1.c0 UNION ALL
    SELECT MAX(t1.rowid) as aggr FROM t1 WHERE NOT('+' >= t1.c0) UNION ALL
    SELECT MAX(t1.rowid) as aggr FROM t1 WHERE ('+' >= t1.c0) IS NULL
);
```



Bug Example: CockroachDB

```
SET vectorize=experimental_on;
CREATE TABLE t0(c0 INT);
CREATE TABLE t1(c0 BOOL) INTERLEAVE IN PARENT t0(rowid);
INSERT INTO t0(c0) VALUES (0);
INSERT INTO t1(rowid, c0) VALUES(0, TRUE);

SELECT MAX(t1.rowid)    SELECT MAX(aggr) FROM (
FROM t1;                  SELECT MAX(t1.rowid) as aggr FROM t1 WHERE '+' >= t1.c0 UNION ALL
                                SELECT MAX(t1.rowid) as aggr FROM t1 WHERE NOT('+' >= t1.c0) UNION ALL
                                SELECT MAX(t1.rowid) as aggr FROM t1 WHERE ('+' >= t1.c0) IS NULL
);
```



Testing Decomposable Aggregate Functions

Q	Q'_{ptern}	Composition operator
SELECT AVG(<e>) FROM <tables> [<joins>];	SELECT SUM(<e>) as s, COUNT (<e>) as c FROM <tables> [<joins>];	$\frac{\text{SUM}(s)}{\text{SUM}(c)}$

A single value to represent a partition is insufficient

Bug example: ClickHouse

```
CREATE TABLE t3 (`c0` Int32, `c1` Int32, `c2` String) ENGINE = Log()
INSERT INTO t3(c0,c1,c2) VALUES (1,10,'1'), (1,0,'2');
```

```
SELECT *
FROM t3
```

c0	c1	c2
1	10	1
1	0	2

```
SELECT MIN(c2)
FROM t3
GROUP BY c0
```

MIN(c2)
1



```
SELECT MIN(t3.c2)
FROM t3
GROUP BY t3.c0
HAVING NOT t3.c1
UNION ALL
SELECT MIN(t3.c2)
FROM t3
GROUP BY t3.c0
HAVING NOT (NOT t3.c1)
UNION ALL
SELECT MIN(t3.c2)
FROM t3
GROUP BY t3.c0
HAVINGisNull(NOT t3.c1)
```

MIN(c2)
2
1

Bug example: ClickHouse

```
SELECT MIN(t3.c2)
FROM t3
GROUP BY t3.c0
HAVING NOT t3.c1
```

Incorrect Query should return an error,
but not incorrect answer.

Each column reference directly contained in the search condition shall be one of the following:
a) An unambiguous reference to a column that is functionally dependent on the set consisting of
every column referenced by a column reference contained in group by clause.

...

**How general is the technique?
Can I apply it to other domains?**

Metamorphic Testing

```
SELECT * FROM t0, t1  
WHERE t0.c0 = t1.c0;
```



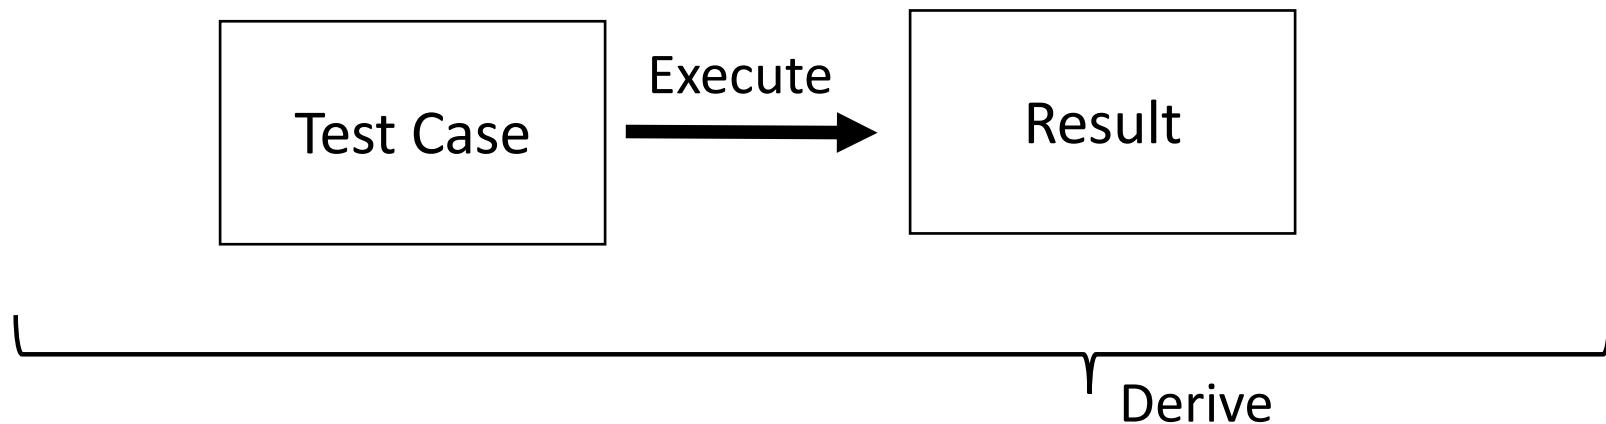
t0.c0	t1.c0
0.0	-0.0

Derive

```
SELECT * FROM t0, t1 WHERE t0.c0=t1.c0  
UNION ALL  
SELECT * FROM t0, t1 WHERE NOT (t0.c0=t1.c0)  
UNION ALL  
SELECT * FROM t0, t1 WHERE (t0.c0=t1.c0) IS NULL;
```



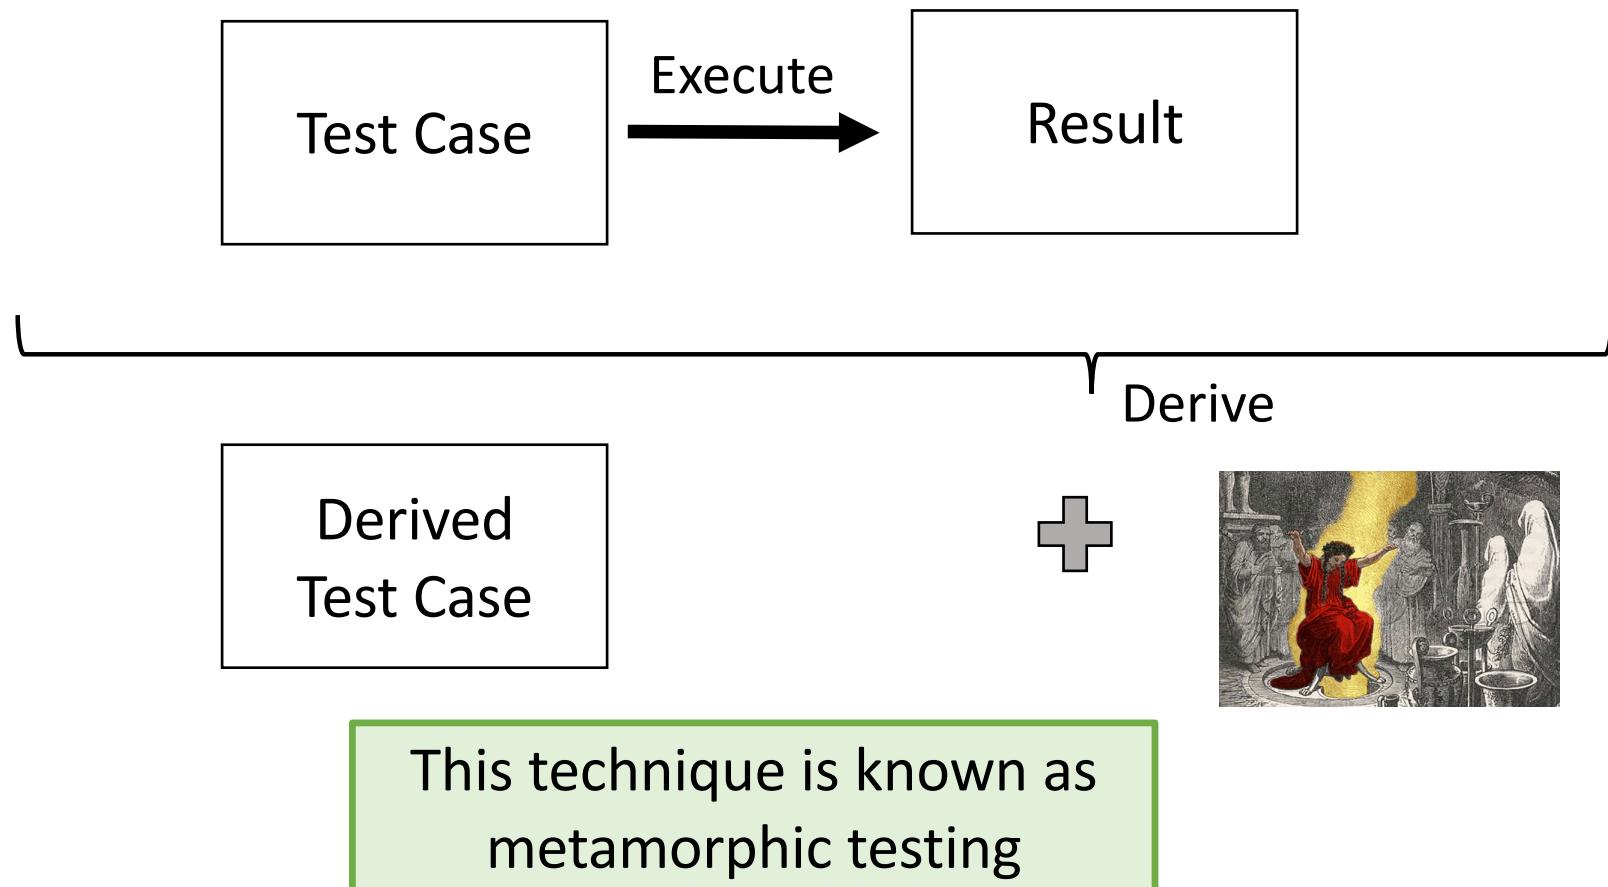
Metamorphic Testing



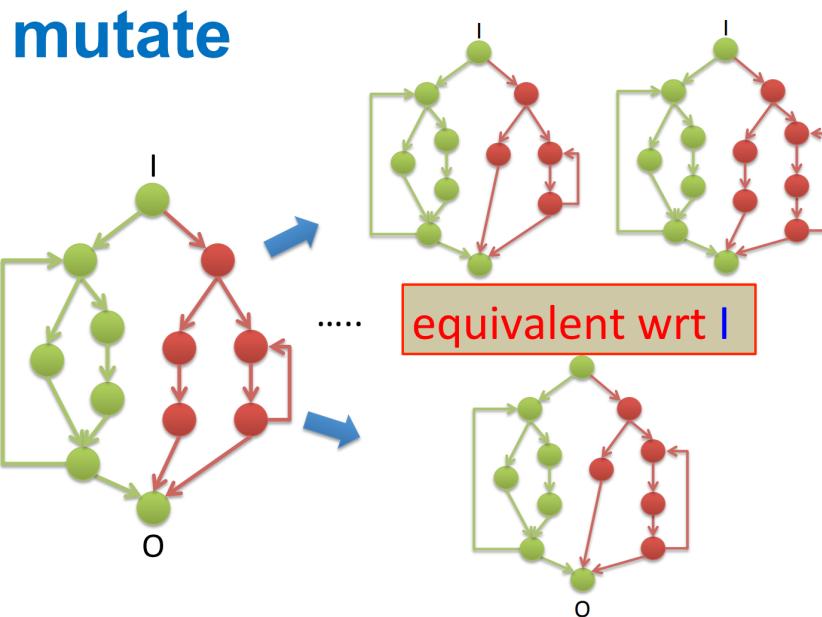
```
SELECT * FROM t0, t1 WHERE t0.c0=t1.c0  
UNION ALL  
SELECT * FROM t0, t1 WHERE NOT (t0.c0=t1.c0)  
UNION ALL  
SELECT * FROM t0, t1 WHERE (t0.c0=t1.c0) IS NULL;
```



Metamorphic Testing



Equivalence Modulo Inputs (EMI) for Testing Compilers



EMI's idea is to create programs that produce
the same output for a given input

Can we test not only DBMS?

Approach is quite generic.

We can try to use it in

1. Generic API with filtering or grouping
2. Regular expressions library
3. Event processing
4. Image recognition
5. Your ideas

**What about other
metamorphic test oracles for DBMSs?**

Finding Logic Bugs in DBMSs

Ternary Logic
Partitioning

OOPSLA '20

Non-optimizing
Reference Engine
Construction

ESEC/FSE '20

Pivoted Query
Synthesis

OSDI '20

Goal: Find Logic Bugs



Optimization bugs: logic
bugs in the query optimizer

Motivating Example

t0

c0
-1

```
CREATE TABLE t0(c0 UNIQUE);
INSERT INTO t0 VALUES (-1) ;
SELECT * FROM t0 WHERE t0.c0 GLOB '-*';
```



-1



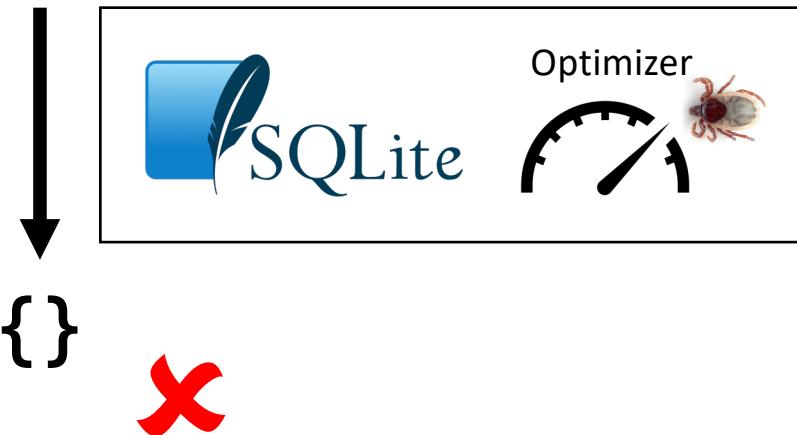
<https://www.sqlite.org/src/tktview?name=0f0428096f>

Motivating Example

t0

c0
-1

```
CREATE TABLE t0(c0 UNIQUE);
INSERT INTO t0 VALUES (-1) ;
SELECT * FROM t0 WHERE t0.c0 GLOB '-*';
```



<https://www.sqlite.org/src/tktview?name=0f0428096f>

Motivating Example

t0

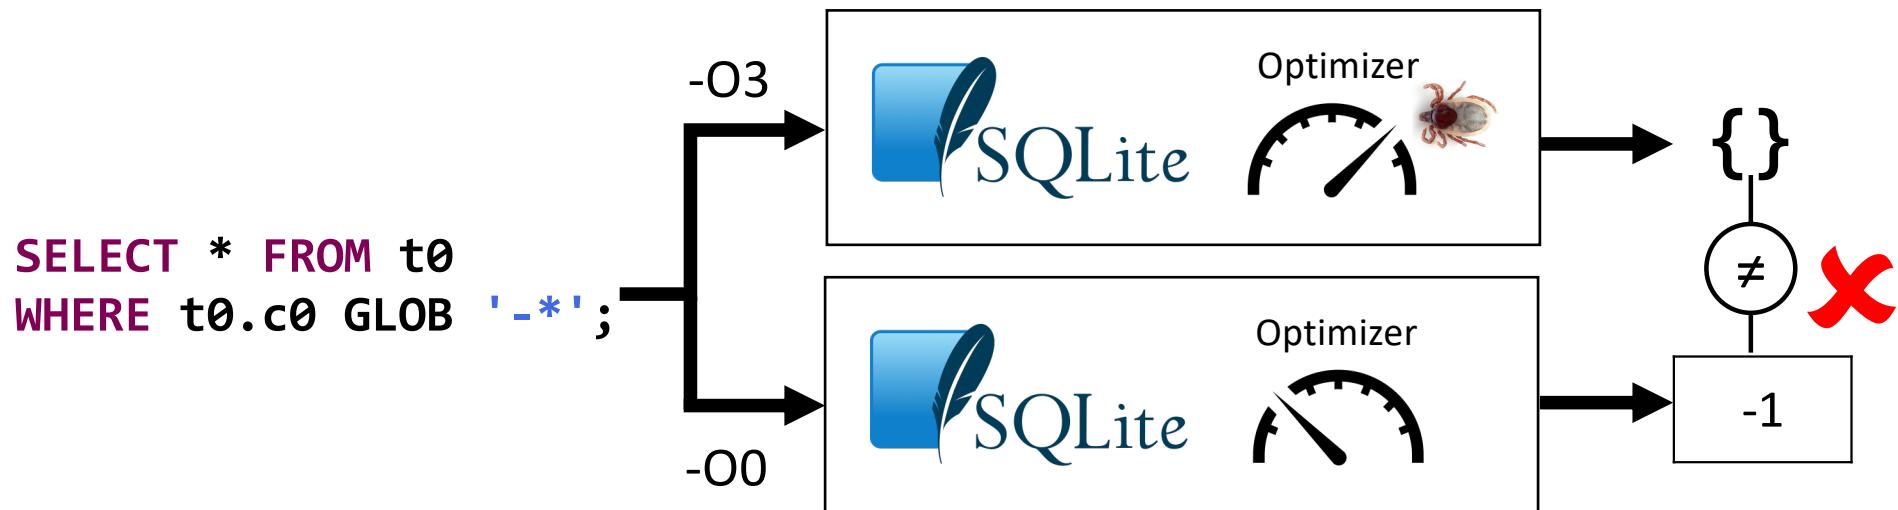
c0
-1

```
CREATE TABLE t0(c0 UNIQUE);
INSERT INTO t0 VALUES (-1) ;
SELECT * FROM t0 WHERE t0.c0 GLOB '-*';
```

The *LIKE optimization* malfunctioned for non-text columns and a pattern prefix of “-”



Differential Testing



Differential Testing

List Of PRAGMAs

[analysis_limit](#)
[application_id](#)
[auto_vacuum](#)
[automatic_index](#)
[busy_timeout](#)
[cache_size](#)
[cache_spill](#)
[case_sensitive_like](#)
[cell_size_check](#)
[checkpoint_fullfsync](#)
[collation_list](#)
[compile_options](#)
[count_changes¹](#)
[data_store_directory¹](#)
[data_version](#)
[database_list](#)
[default_cache_size¹](#)
[defer_foreign_keys](#)
[empty_result_callbacks¹](#)
[encoding](#)
[foreign_key_check](#)
[foreign_key_list](#)
[foreign_keys](#)
[freelist_count](#)
[full_column_names¹](#)

[fullfsync](#)
[function_list](#)
[hard_heap_limit](#)
[ignore_check_constraints](#)
[incremental_vacuum](#)
[index_info](#)
[index_list](#)
[index_xinfo](#)
[integrity_check](#)
[journal_mode](#)
[journal_size_limit](#)
[legacy_alter_table](#)
[legacy_file_format](#)
[locking_mode](#)
[max_page_count](#)
[mmap_size](#)
[module_list](#)
[optimize](#)
[page_count](#)
[page_size](#)
[parser_trace²](#)
[pragma_list](#)
[query_only](#)
[quick_check](#)
[read_uncommitted](#)

[recursive_triggers](#)
[reverse_unordered_selects](#)
[schema_version³](#)
[secure_delete](#)
[short_column_names¹](#)
[shrink_memory](#)
[soft_heap_limit](#)
[stats³](#)
[synchronous](#)
[table_info](#)
[table_xinfo](#)
[temp_store](#)
[temp_store_directory¹](#)
[threads](#)
[trusted_schema](#)
[user_version](#)
[vdbe_addoptrace²](#)
[vdbe_debug²](#)
[vdbe_listing²](#)
[vdbe_trace²](#)
[wal_autocheckpoint](#)
[wal_checkpoint](#)
[writable_schema³](#)

Differential Testing

List Of PRAGMAS

[analysis_limit](#)
[application_id](#)
[auto_vacuum](#)
[automatic_index](#)
[busy_timeout](#)
[cache_size](#)
[cache_spill](#)
[case_sensitive_like](#)
[cell_size](#)
[check](#)
[collation_sequence](#)
[compile_options](#)
[count_changes¹](#)
[data_store_directory](#)
[data_version](#)
[database_list](#)
[default_cache_size](#)
[defer_foreign_keys](#)
[empty_result_callbacks¹](#)
[encoding](#)
[foreign_key_check](#)
[foreign_key_list](#)
[foreign_keys](#)
[freelist_count](#)
[full_column_names¹](#)

[fullfsync](#)
[function_list](#)
[hard_heap_limit](#)
[ignore_check_constraints](#)
[incremental_vacuum](#)
[index_info](#)
[index_list](#)
[index_xinfo](#)
[indexable](#)
[legacy_alter_table](#)
[legacy_file_format](#)
[page_count](#)
[page_size](#)
[parser_trace²](#)
[pragma_list](#)
[query_only](#)
[quick_check](#)
[read_uncommitted](#)

[recursive_triggers](#)
[reverse_unordered_selects](#)
[schema_version³](#)
[secure_delete](#)
[short_column_names¹](#)
[shrink_memory](#)
[soft_heap_limit](#)
[stats³](#)

[temp_store](#)
[temp_store_directory¹](#)

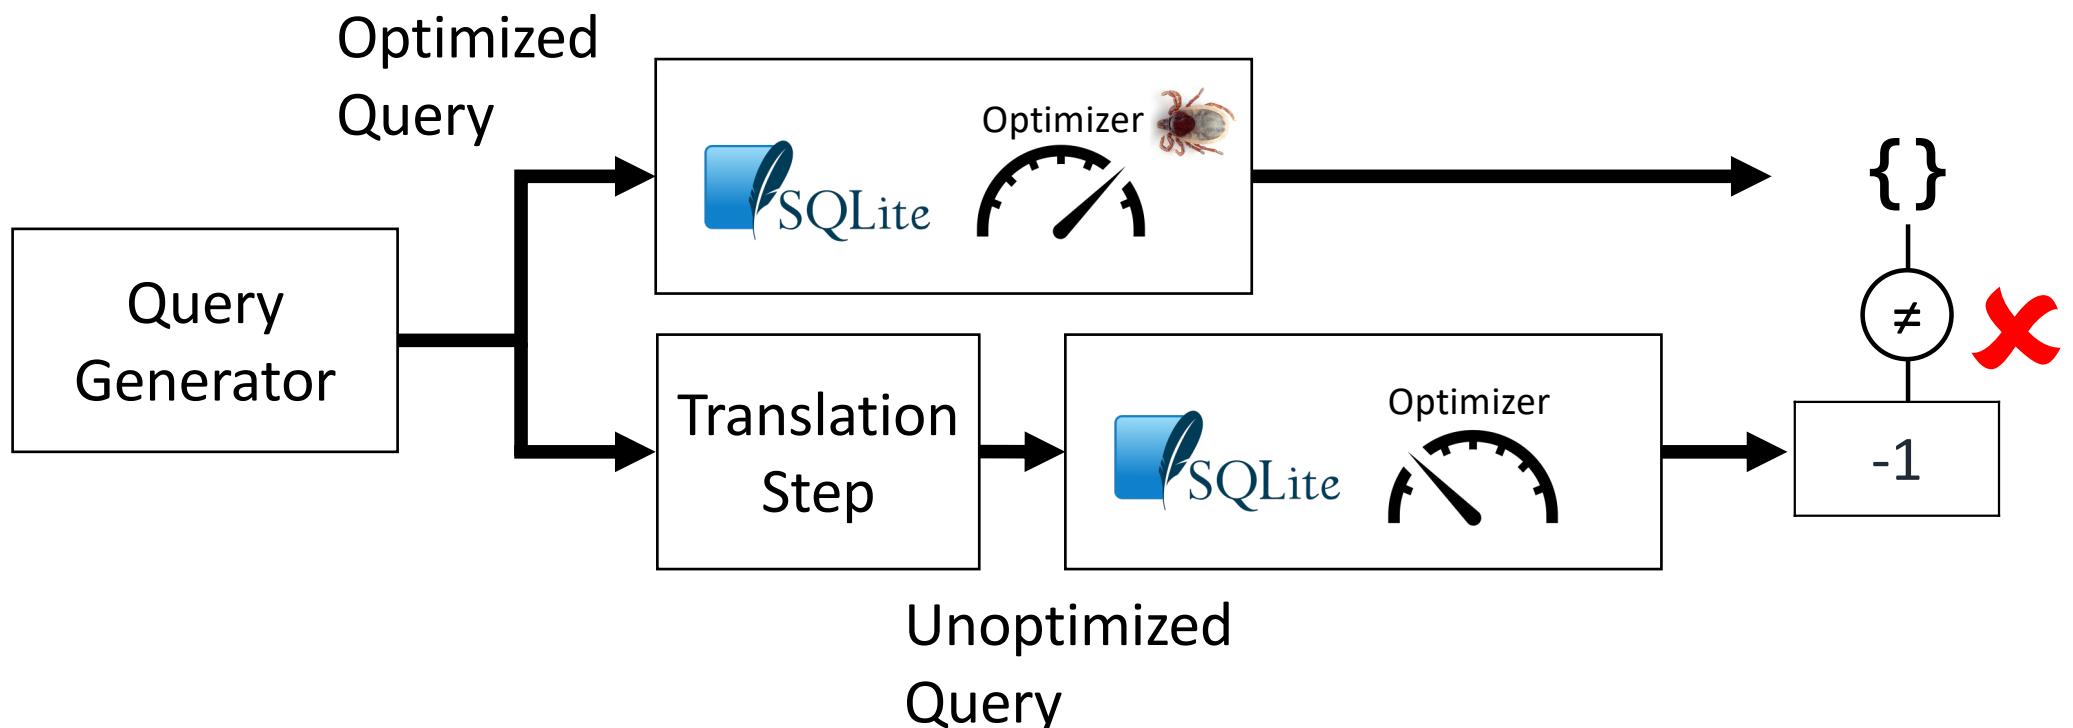
[vdbe_listing²](#)
[vdbe_trace²](#)
[wal_autocheckpoint](#)
[wal_checkpoint](#)
[writable_schema³](#)

DBMSs typically provide only very
limited control over optimizations

NoREC

Idea: **Rewrite** the query so that
the DBMS **cannot optimize it**

Idea



We want to create a “non-optimizing reference engine”

Given Query

Consider the following format for the optimized query:

```
SELECT * FROM t0  
WHERE p;  
  
↑  
t0.c0 GLOB '_*' 
```

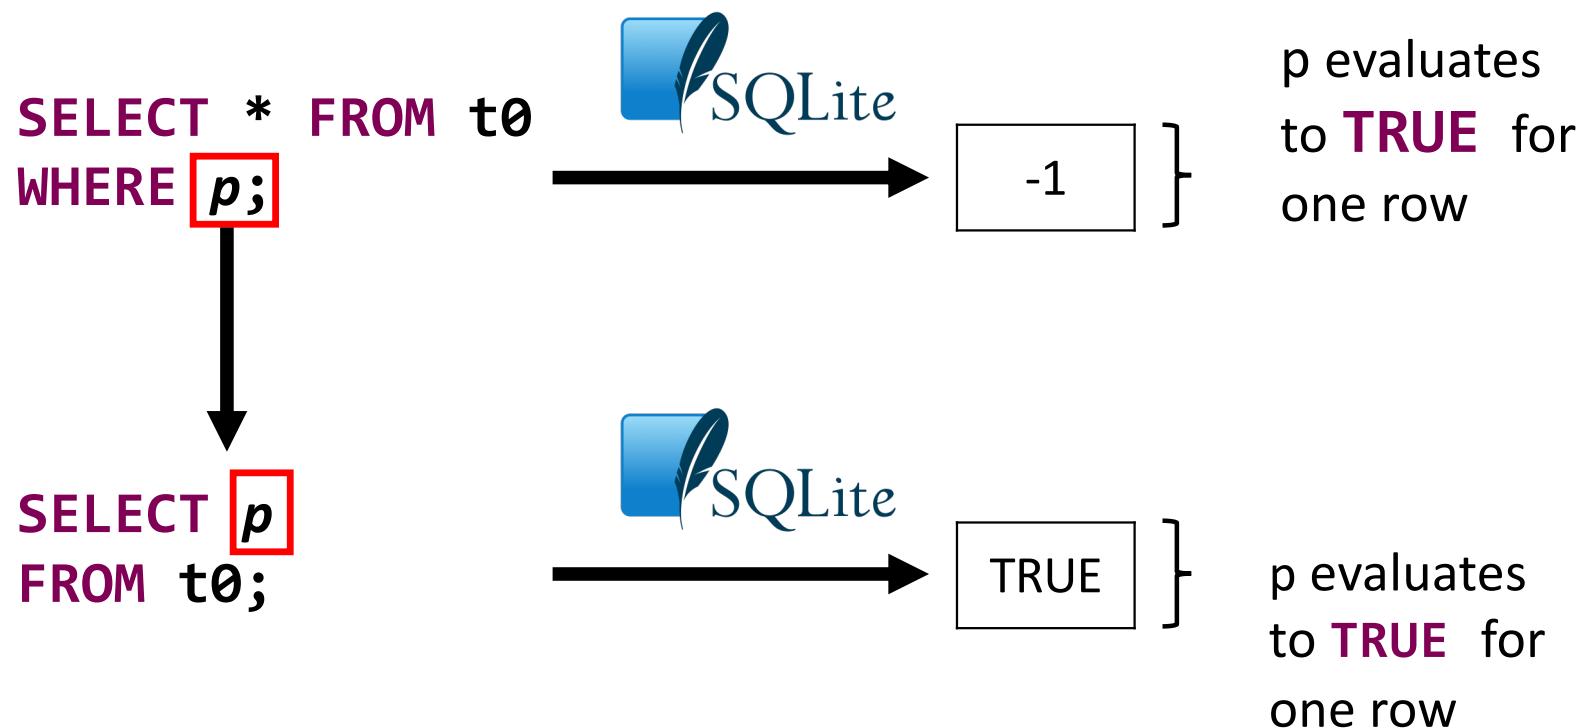
It is unobvious how we could derive an unoptimized query

Insight



First Insight: The predicate p must
always evaluate to the same value,
irrespective of its context

Translation Step (Correct Case)

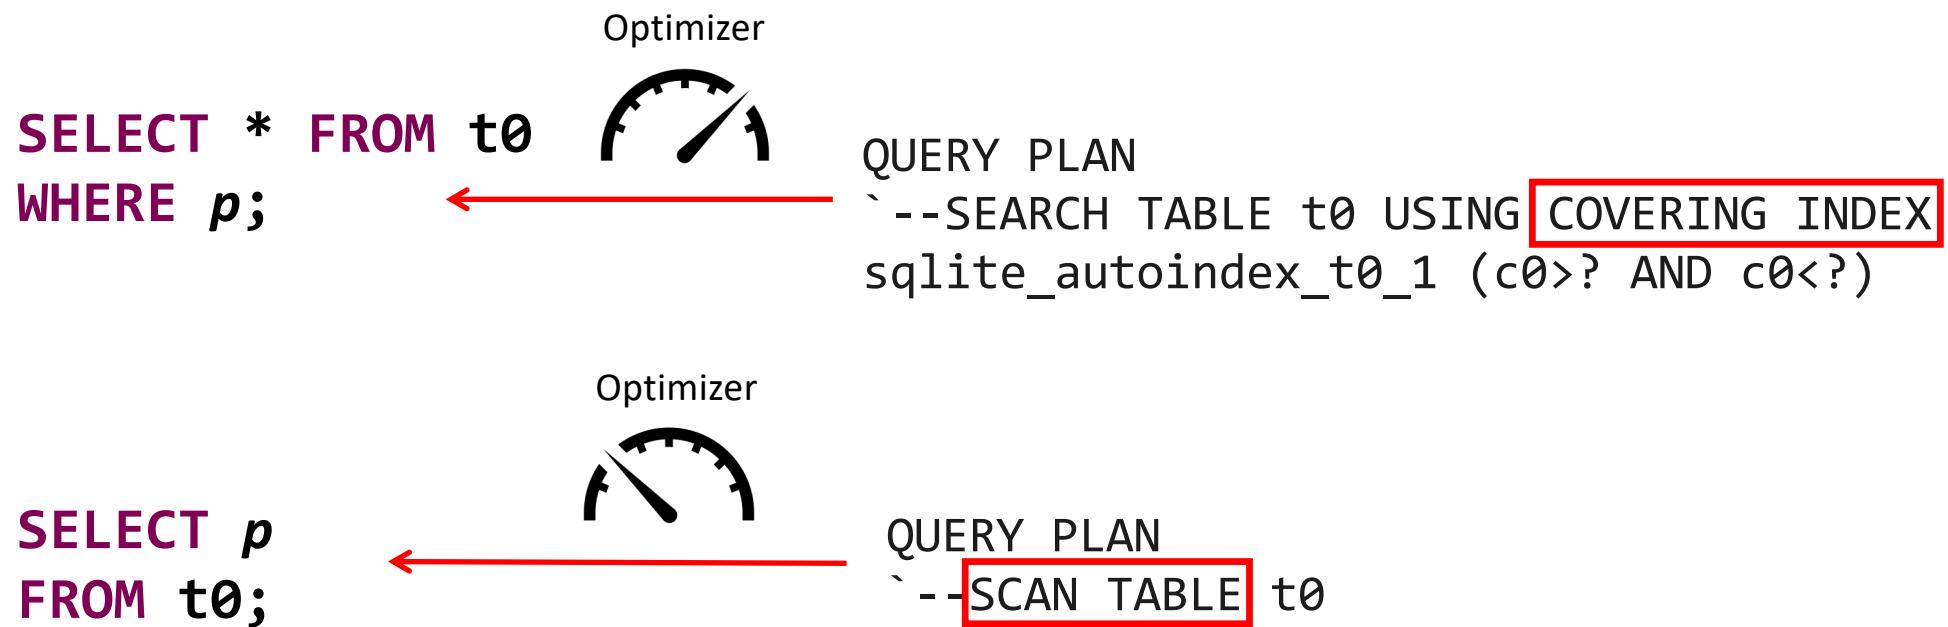


Insights



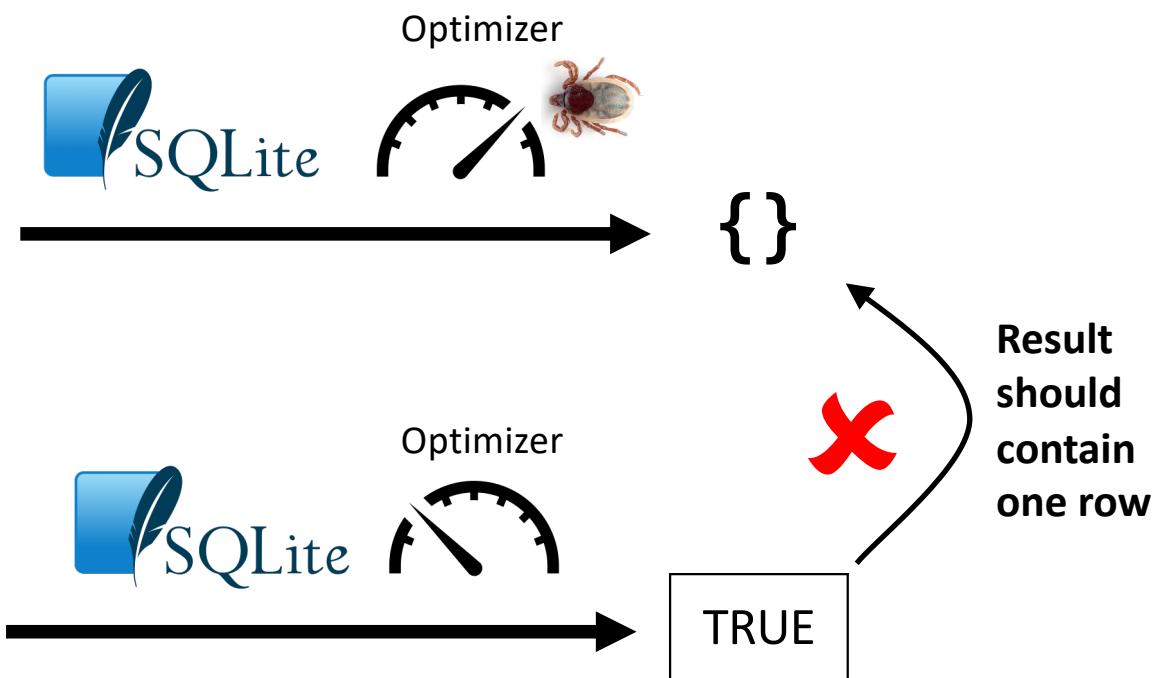
Second Insight: DBMSs focus their optimizations on reducing the amount of data that is processed

Translation Step



Translation Step

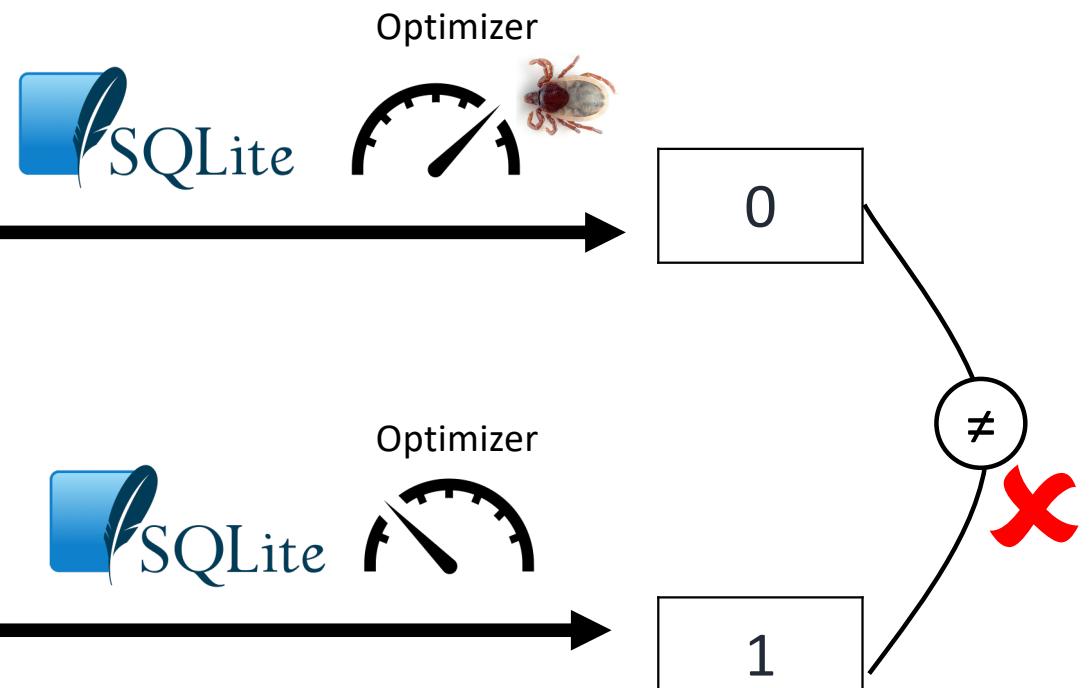
```
SELECT * FROM t0  
WHERE p;
```



```
SELECT p  
FROM t0;
```

Counting Implementation

```
SELECT COUNT(*)  
FROM ...  
WHERE p
```



```
SELECT SUM(count) FROM (  
    SELECT p IS TRUE  
        as count  
    FROM <tables>  
);
```

What about other, non-metamorphic test oracles?

Finding Logic Bugs in DBMSs

Ternary Logic
Partitioning

OOPSLA '20

Non-optimizing
Reference Engine
Construction

ESEC/FSE '20

Pivoted Query
Synthesis

OSDI '20

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

IS NOT is a “null-safe”
comparison operator

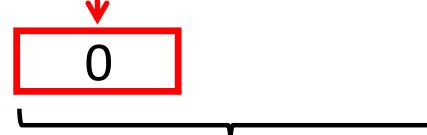
Example: SQLite3 Bug

t0
c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



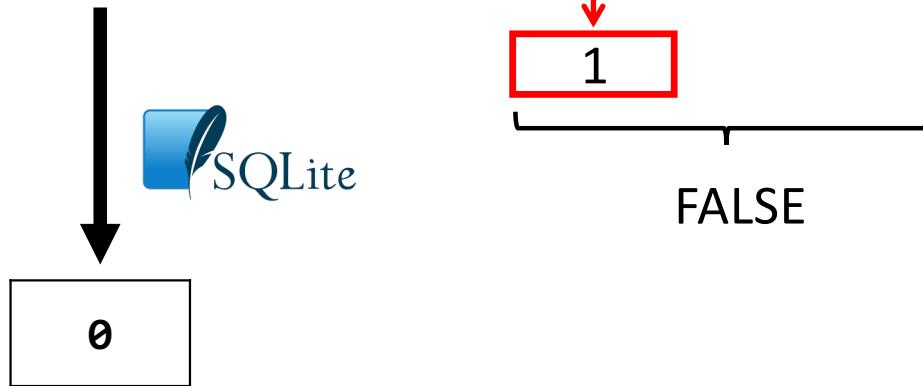
0



Example: SQLite3 Bug

t0
c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

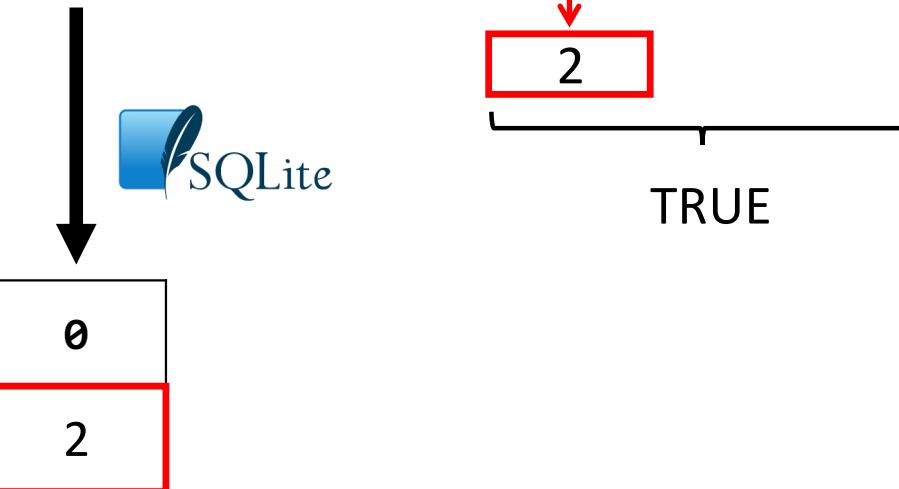


Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0
2

NULL was not contained
in the result set!



PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

← Pivot row

Validate the result set based on
one randomly-selected row

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

↓
NULL

TRUE

Generate a query that is
guaranteed to at least fetch
the pivot row

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

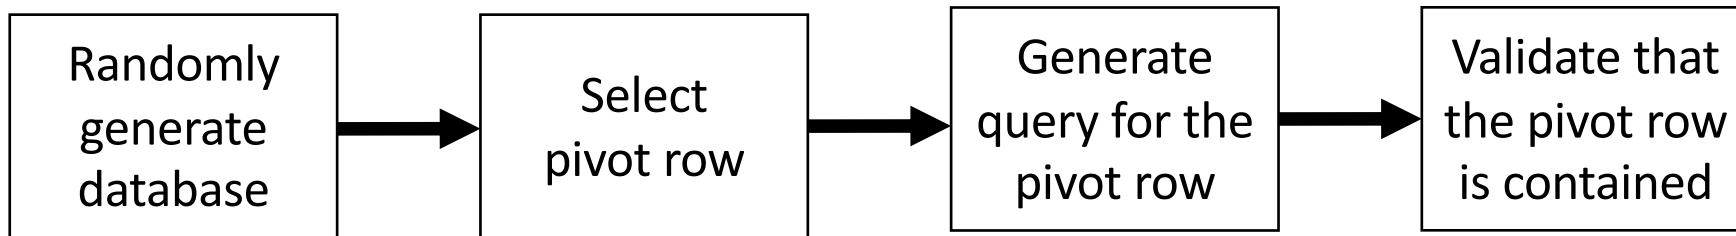


0
2

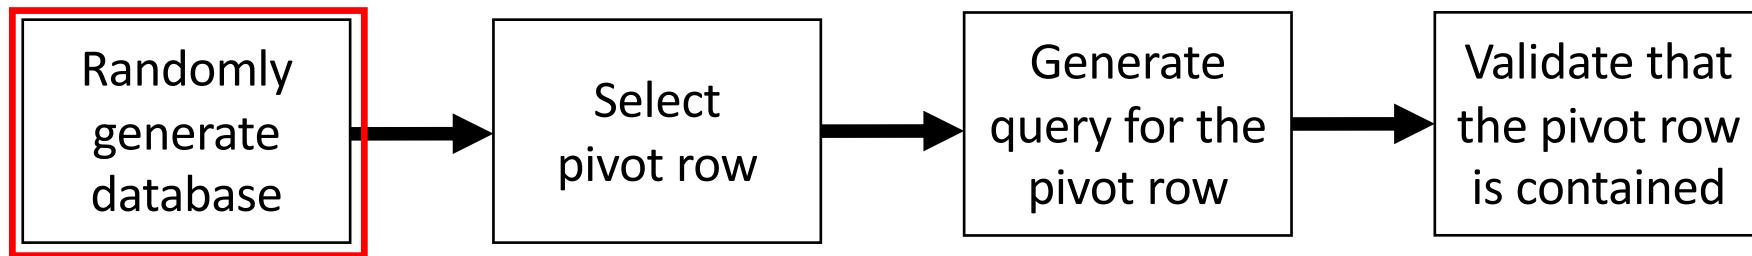


If the **pivot row is missing** from the result set a **bug** has been detected

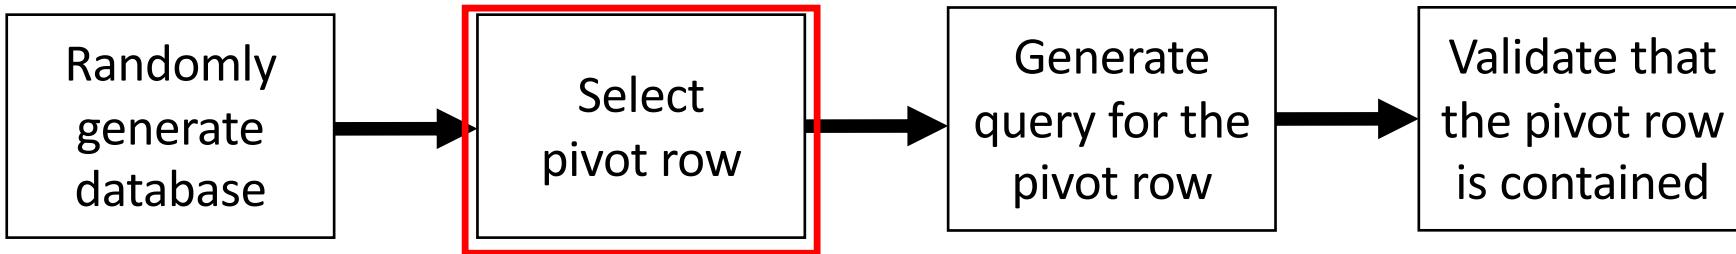
Approach



Approach

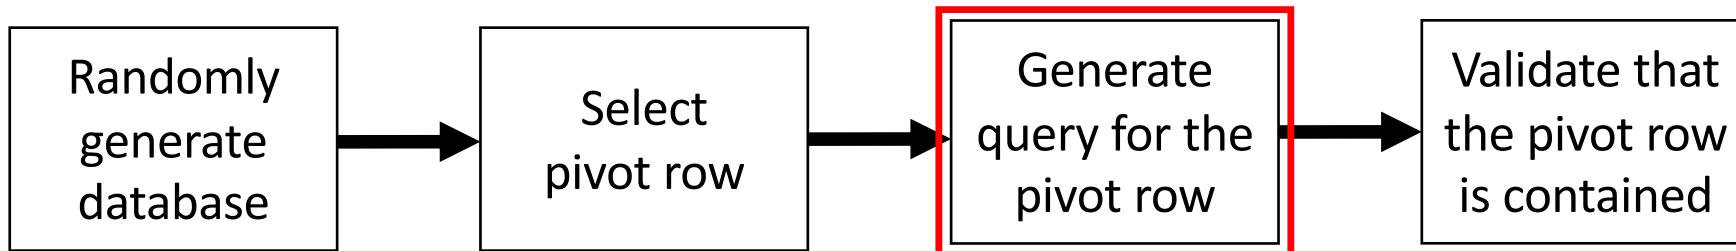


Approach



One **random row** from multiple tables and views

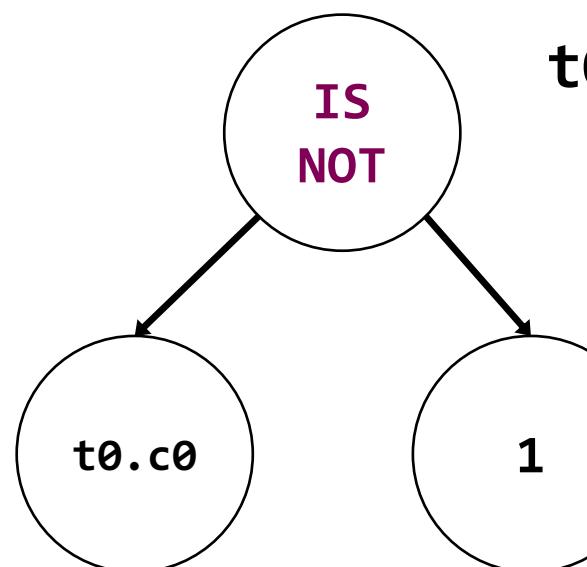
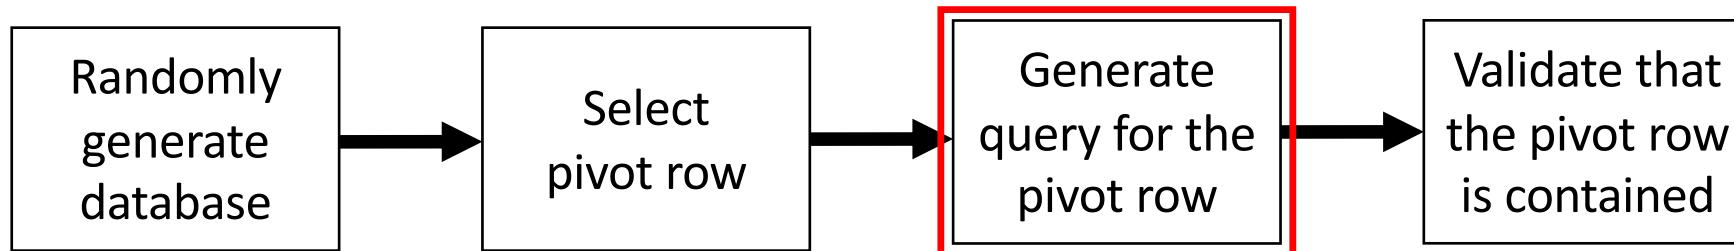
Approach



**SELECT c0 FROM t0
WHERE**

Generate **predicates** that **evaluate to TRUE** for the pivot row and use them in JOIN and WHERE clauses

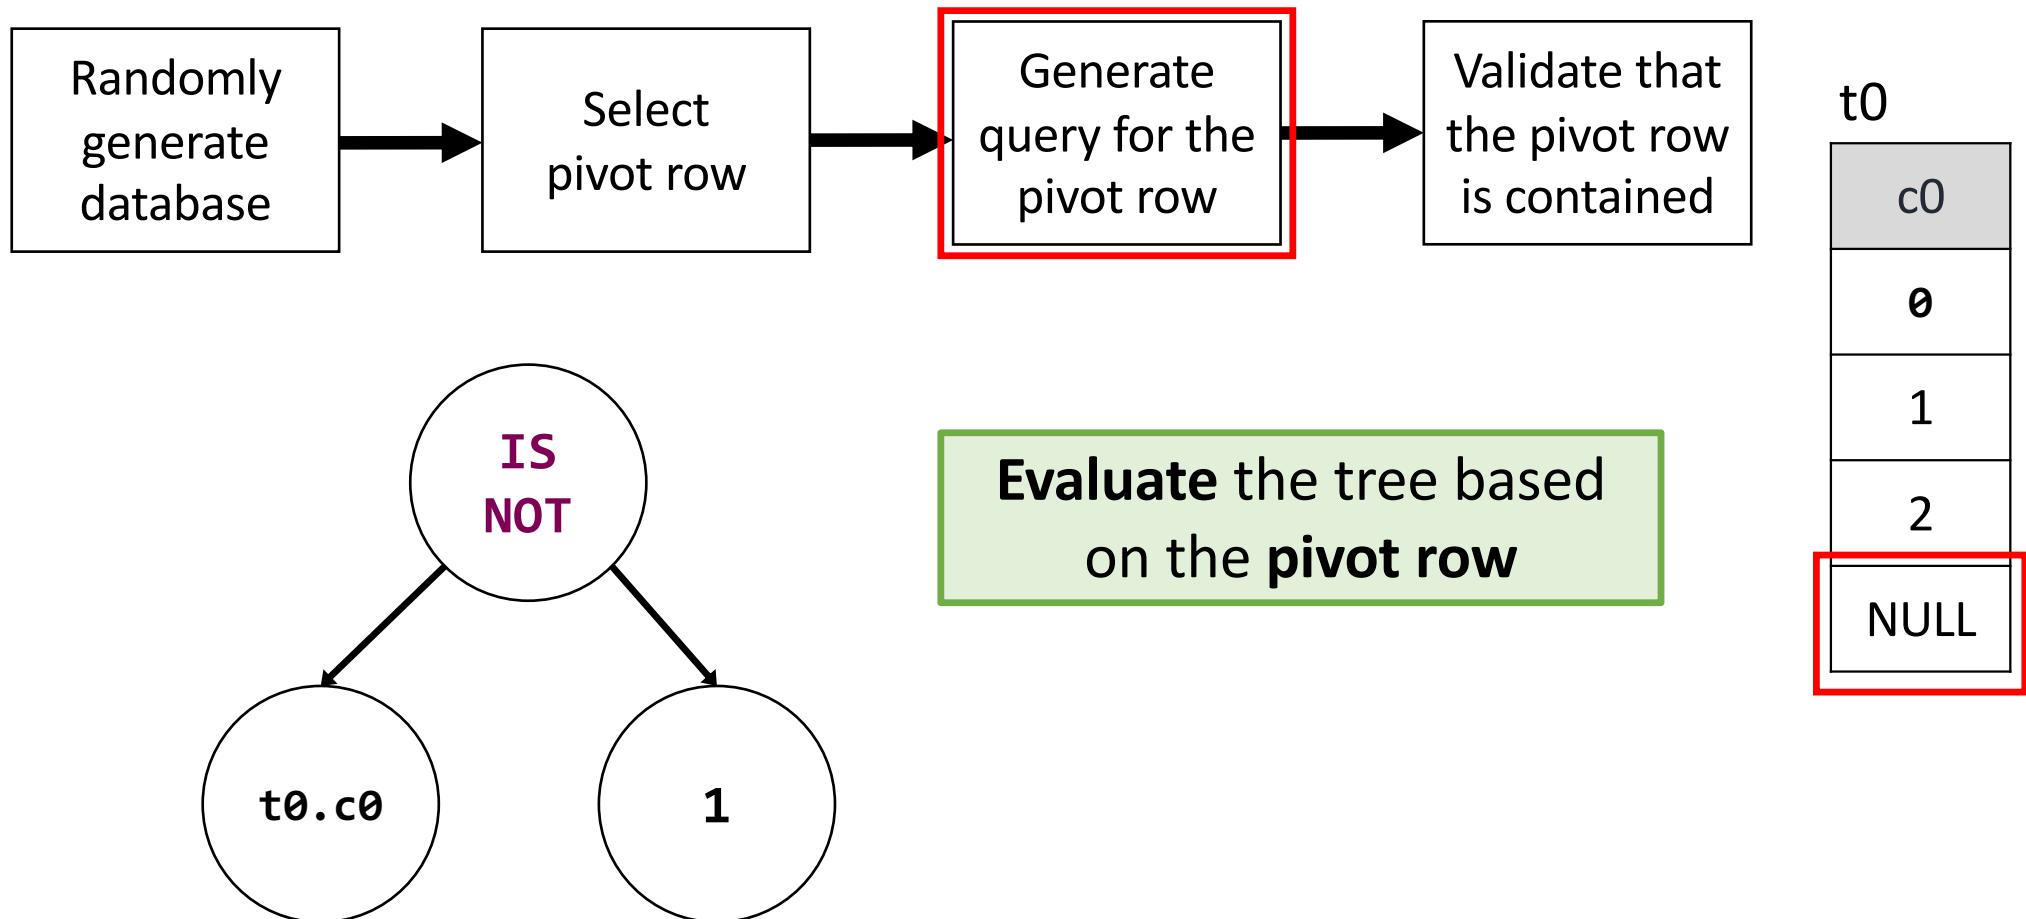
Random Expression Generation



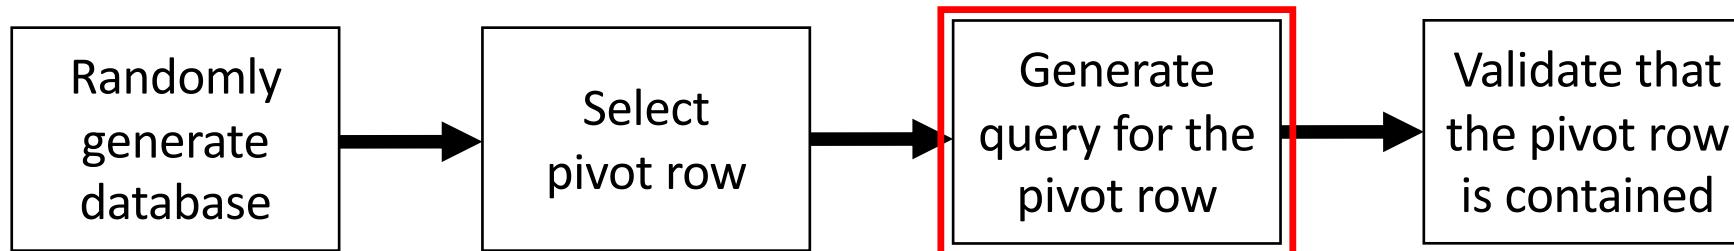
$t0.c0 \text{ IS NOT } 1;$

We implemented an
expression evaluator for
each node

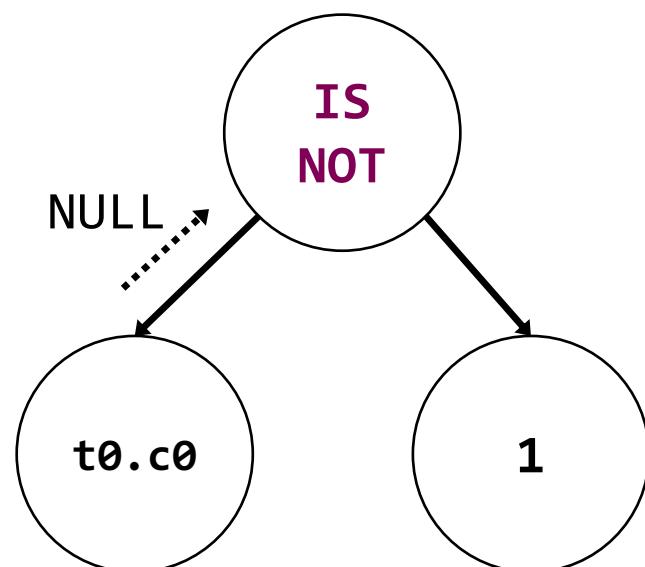
Random Expression Generation



Random Expression Generation

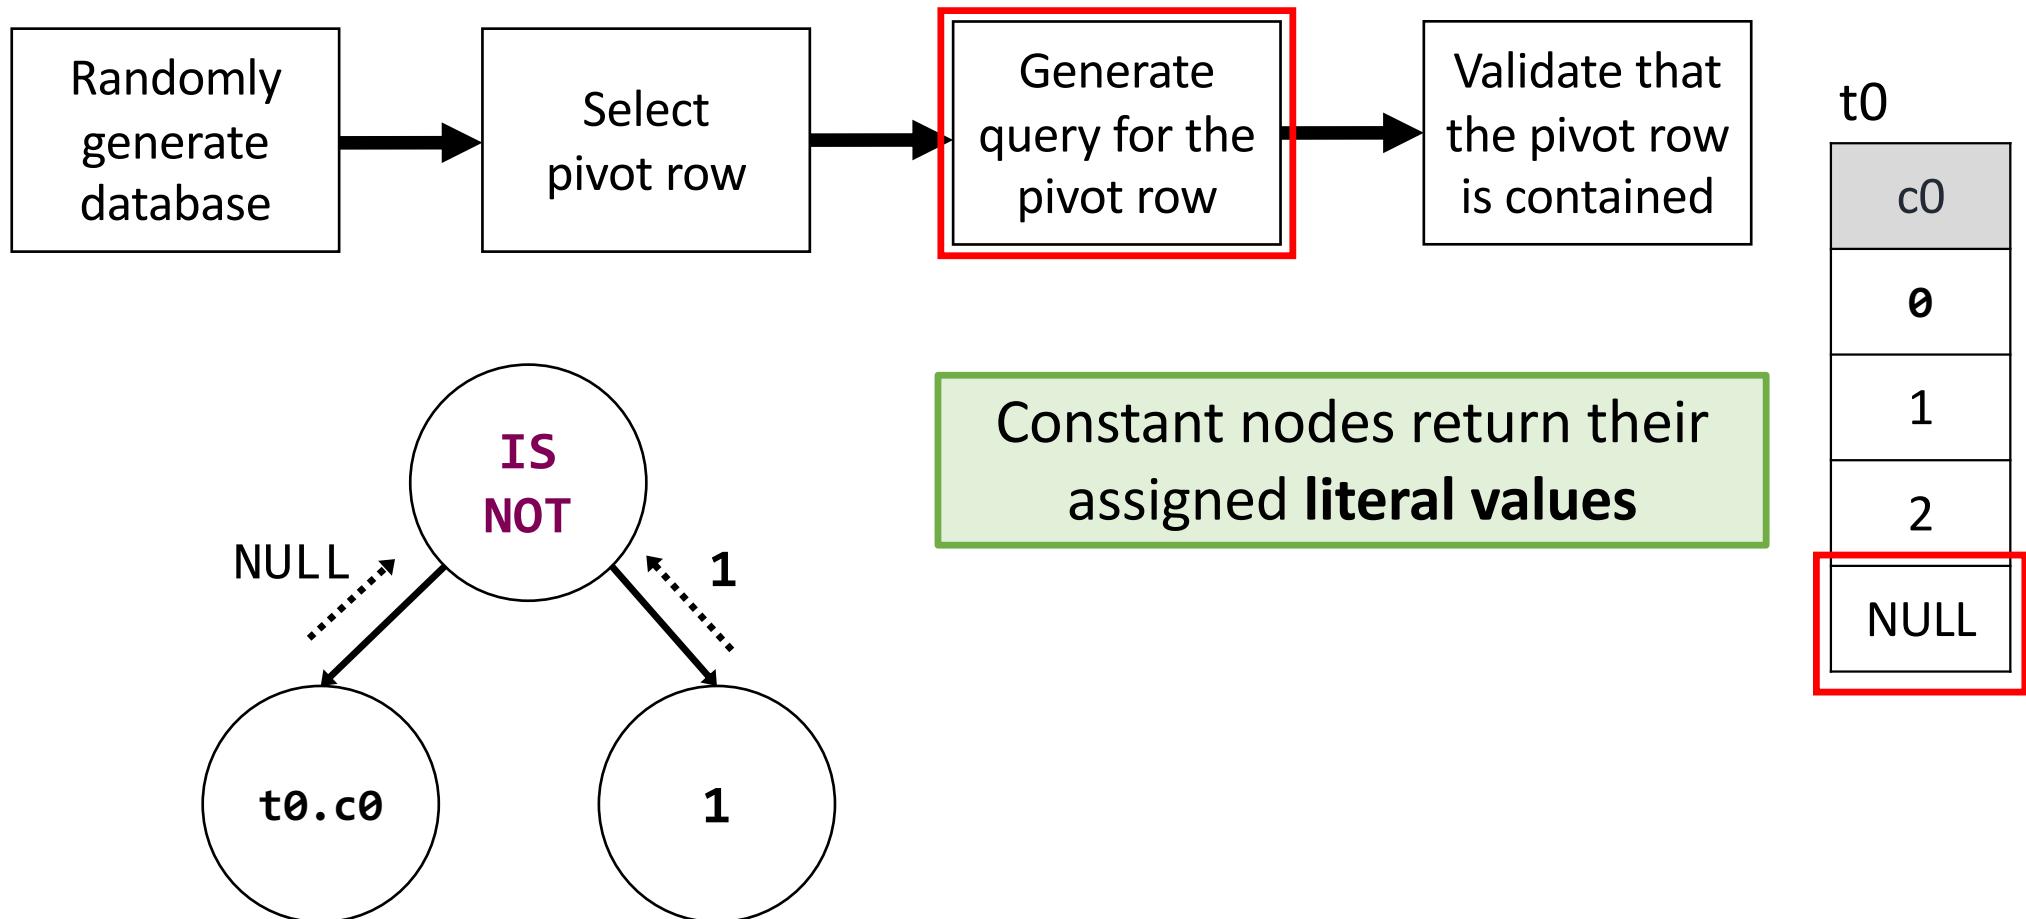


Column references return the values from the pivot row

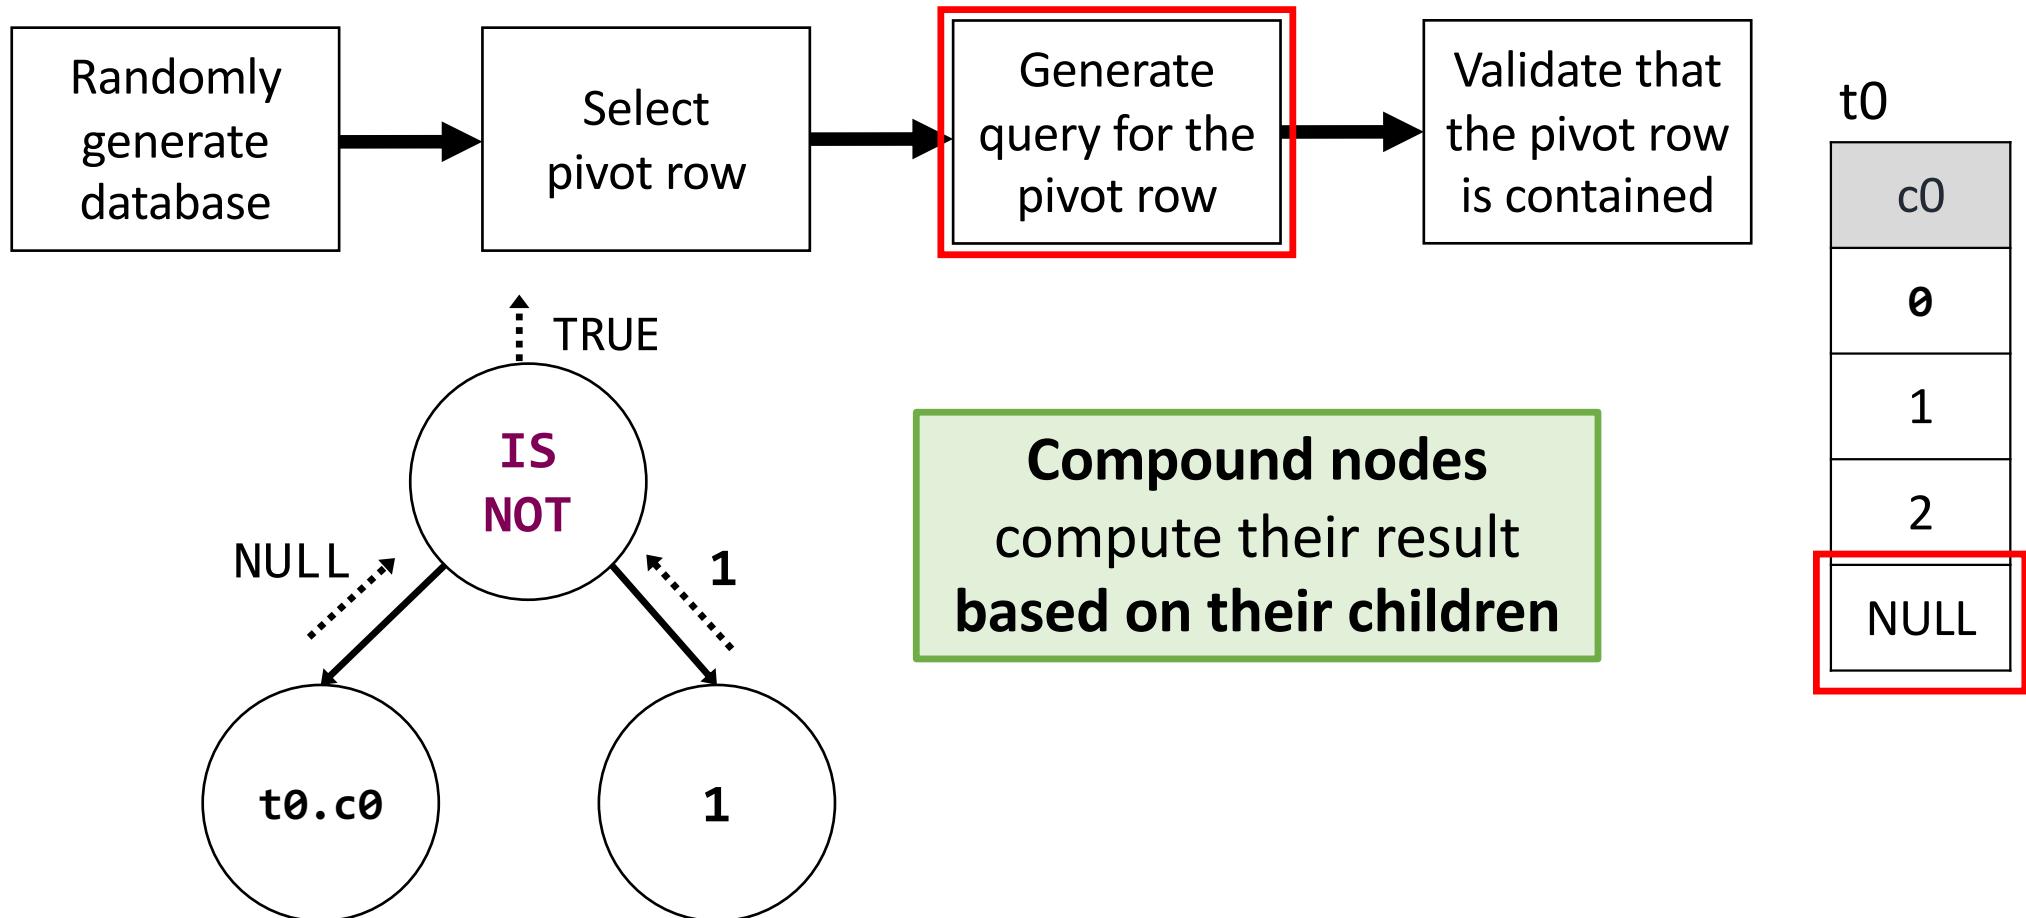


t0	c0
0	
1	
2	
	NULL

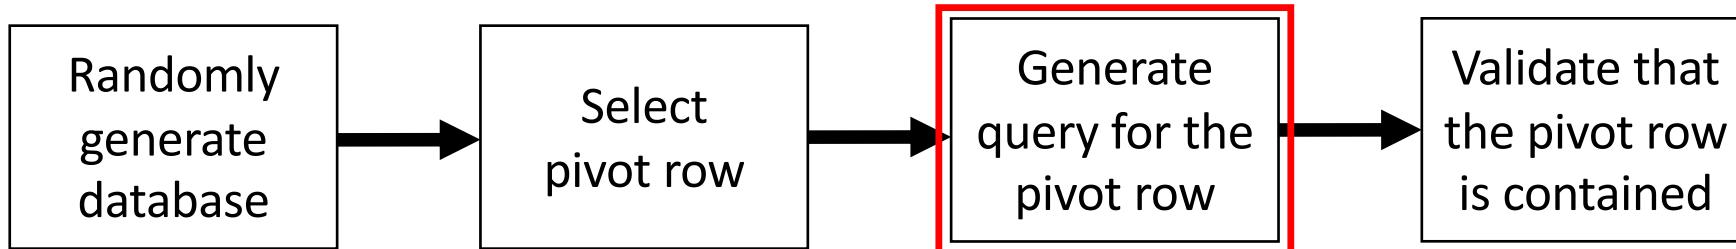
Random Expression Generation



Random Expression Generation



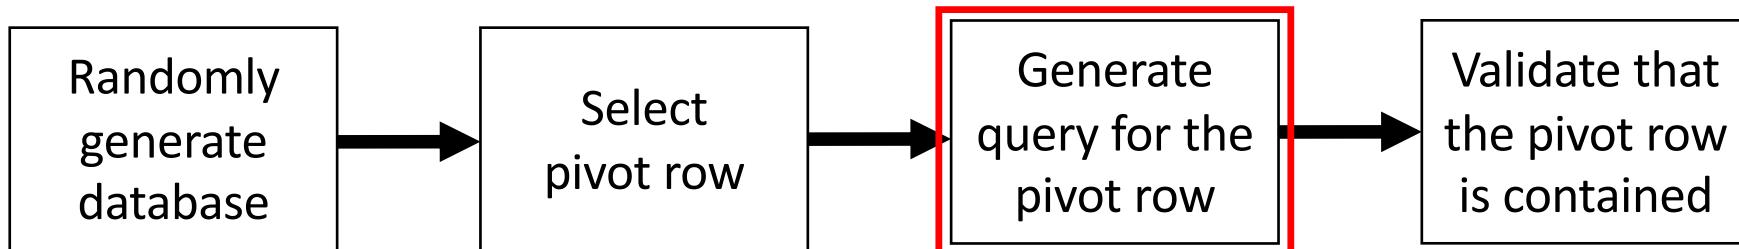
Query Synthesis



```
SELECT c0 c0 FROM t0  
WHERE t0.c0 IS NOT 1;
```

What if the expression **does not evaluate to TRUE?**

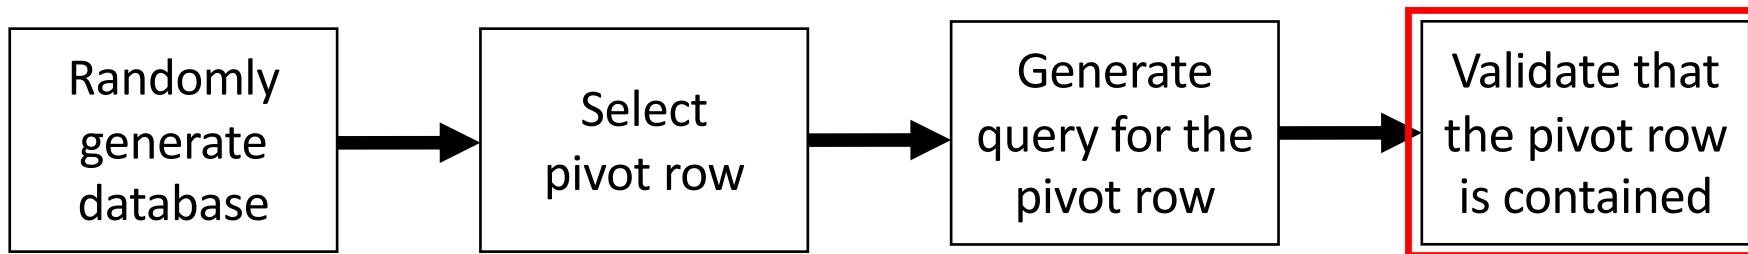
Random Expression Rectification



```
switch (result) {  
    case TRUE:  
        result = randexpr;  
    case FALSE:  
        result = NOT randexpr;  
    case NULL:  
        result = randexpr IS NULL;  
}
```

Alternatively, we could validate that the pivot row is expectedly **not fetched**

Approach



```
SELECT (NULL) INTERSECT  
SELECT c0 FROM t0 WHERE NULL IS NOT 1;
```

Rely on the DBMS to check whether the row is contained

**How do the techniques
compare to each other?**

Comparison

Property	PQS	NoREC	TLP
WHERE	✓	✓	✓
Additional SQL features	✗	✗	✓
Ground truth	✓	✗	✗
No domain knowledge required	✗	✓	✓
Implementation effort	Moderate	Very Low	Low

Comparison

Property	PQS	NoREC	TLP
WHERE	✓	✓	✓
Additional SQL features	✗	✗	✓
Ground truth	✓	✗	✗
No domain knowledge required	✗	✓	✓
Implementation effort	Moderate	Very Low	Low

TLP is applicable to testing a wider range of features

Comparison

Property	PQS	NoREC	TLP
WHERE	✓	✓	✓
Additional SQL features	✗	✗	✓
Ground truth	✓	✗	✗
No domain knowledge required	✗	✓	✓
Implementation effort	Moderate	Very Low	Low

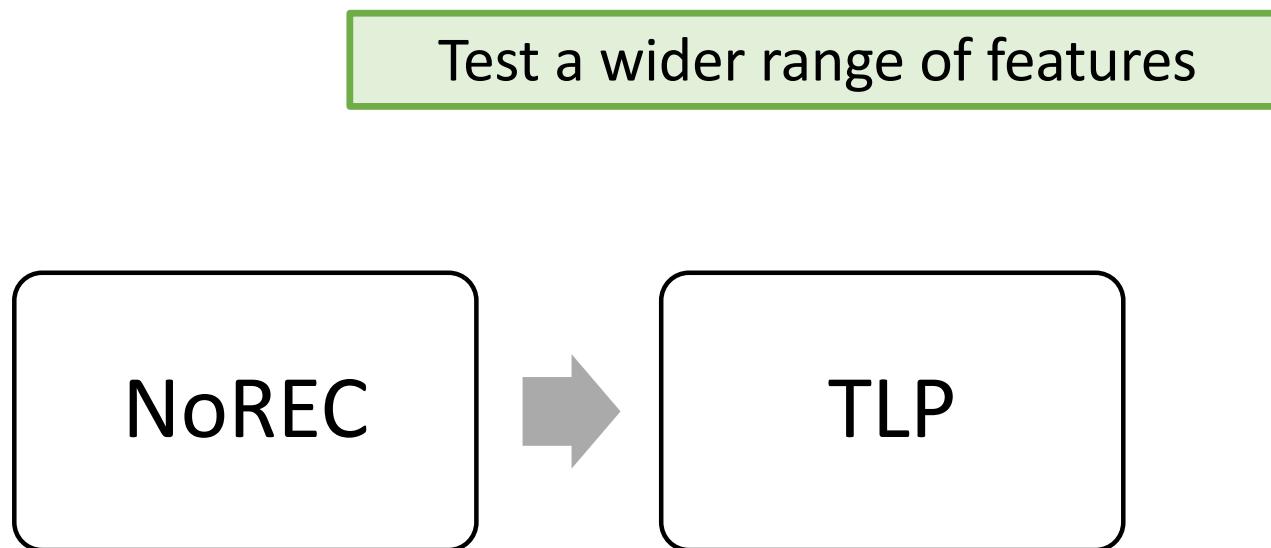
Both NoREC and TLP are
metamorphic testing approaches

Proposed Testing Strategy

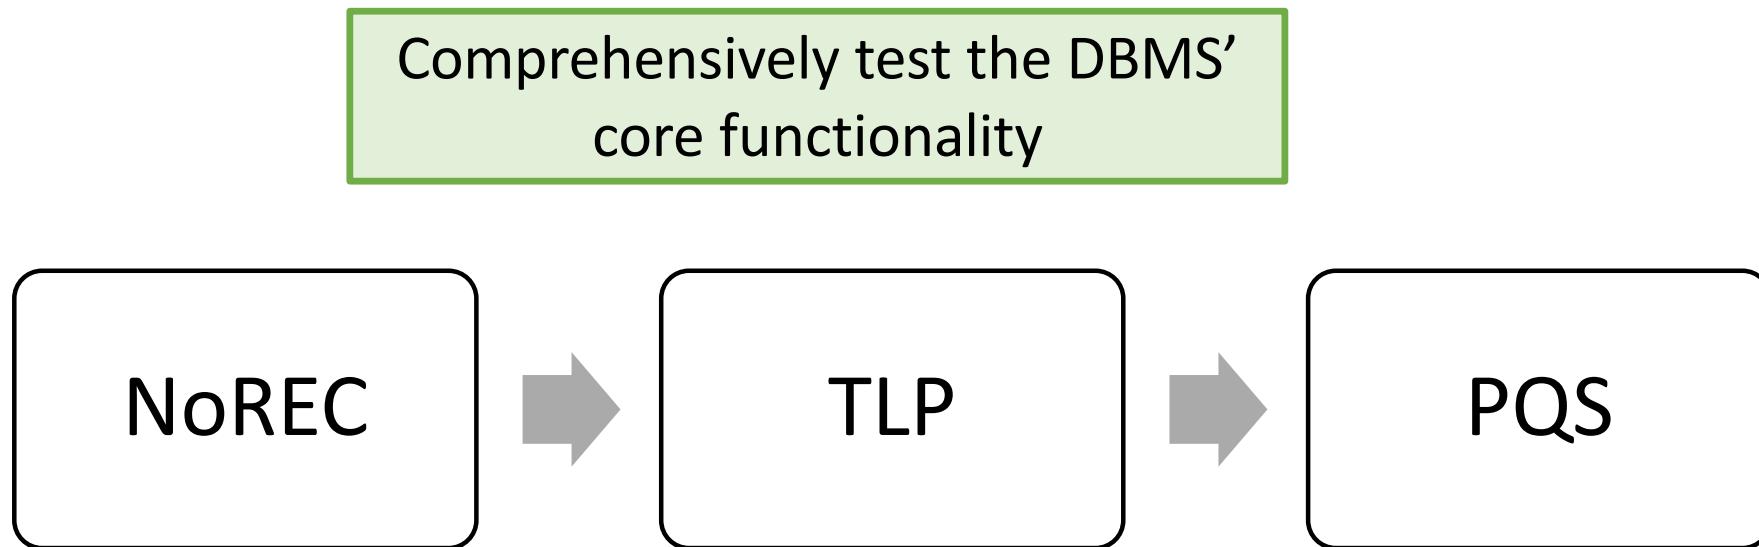
Quickly find the optimization bugs

NoREC

Proposed Testing Strategy



Proposed Testing Strategy



**What implementation strategy
did you use for ClickHouse?**

SQLancer from developer point of view

- | Natural and intuitive approach
- | Fuzzing meets correctness check
- | Pluggable (but you need to write code on Java)
- | You can incorporate approach in other software

How to make integration with your DBMS

1. Import Java or ODBC driver if you have one
2. Teach SQLancer to make connection, create database and make a generic query
3. A bit more to prepare: Create table with schema and insert generator
4. Implement expression generator
5. Implement oracle
6. Expected error handling
7. RUN!

```
254 <dependency>
255   <groupId>org.slf4j</groupId>
256   <artifactId> slf4j-simple</artifactId>
257   <version>1.7.30</version>
258 </dependency>
259 <dependency>
260   <groupId>ru.yandex.clickhouse</groupId>
261   = <artifactId>clickhouse-jdbc</artifactId>
262   <version>0.2.5</version>
263 </dependency>
264 <dependency>
265   <groupId>com.h2database</groupId>
266   <artifactId>h2</artifactId>
267   <version>1.4.200</version>
268 </dependency>
269 </dependencies>
```

Add your
driver

```
① @  
②     public SQLConnection createDatabase(ClickHouseGlobalState globalState) throws SQLException  
③     {  
④         ClickHouseOptions clickHouseOptions = globalState.getDbmsSpecificOptions();  
⑤         globalState.setClickHouseOptions(clickHouseOptions);  
⑥         String url = "jdbc:clickhouse://localhost:8123/default";  
⑦         String databaseName = globalState.getDatabaseName();  
⑧         Connection con = DriverManager.getConnection(url, globalState.getOptions().getUse  
⑨             globalState.getOptions().getPassword());  
⑩         String dropDatabaseCommand = "DROP DATABASE IF EXISTS " + databaseName;  
⑪         globalState.getState().logStatement(dropDatabaseCommand);  
⑫         String createDatabaseCommand = "CREATE DATABASE IF NOT EXISTS " + databaseName;  
⑬         globalState.getState().logStatement(createDatabaseCommand);  
⑭         try (Statement s = con.createStatement()) {  
⑮             s.execute(dropDatabaseCommand);  
⑯             Thread.sleep(millis: 1000);  
⑰         } catch (InterruptedException e) {  
⑱             e.printStackTrace();  
⑲         }  
⑳         try (Statement s = con.createStatement()) {  
⑳             s.execute(createDatabaseCommand);  
⑳             Thread.sleep(millis: 1000);  
⑳         } catch (InterruptedException e) {  
⑳             e.printStackTrace();  
⑳         }  
⑳     }
```

Make
connection,
database
and query

```
210     List<ClickHouseColumn> columns = new ArrayList<>();
211     try (Statement s = con.createStatement()) {
212         try (ResultSet rs = s.executeQuery(sql: "DESCRIBE " + tableName)) {
213             while (rs.next()) {
214                 String columnName = rs.getString(columnLabel: "name");
215                 String dataType = rs.getString(columnLabel: "type");
216                 String defaultType = rs.getString(columnLabel: "default_type");
217                 boolean isAlias = "ALIAS".compareTo(defaultType) == 0;
218                 boolean isMaterialized = "MATERIALIZED".compareTo(defaultType) == 0;
219                 ClickHouseColumn c = new ClickHouseColumn(columnName, getColumnType(dataType)
220                     isMaterialized);
221                 columns.add(c);
222             }
223         }
```

Create table with schema

Your DBMS can has specific SQL extensions.
Implement subset of them

```
@Override
protected ClickHouseExpression generateExpression(ClickHouseLancerDataType type, int de
    if (allowAggregateFunctions && Randomly.getBoolean()) {
        return generateAggregate();
    }
    if (depth >= globalState.getOptions().getMaxExpressionDepth() || Randomly.getBoolean())
        return generateLeafNode(type);
    }
    Expression expr = Randomly.fromOptions(Expression.values());
    ClickHouseLancerDataType leftLeafType = ClickHouseLancerDataType.getRandom();
    ClickHouseLancerDataType rightLeafType = ClickHouseLancerDataType.getRandom();
    if (Randomly.getBoolean()) {
        rightLeafType = leftLeafType;
    }

    switch (expr) {
        case UNARY_PREFIX:
            return new ClickHouseUnaryPrefixOperation(generateExpression(leftLeafType, dept
                ClickHouseUnaryPrefixOperation.ClickHouseUnaryPrefixOperator.getRandom(
        case UNARY_POSTFIX:
            return new ClickHouseUnaryPostfixOperation(generateExpression(leftLeafType, dept
                ClickHouseUnaryPostfixOperation.ClickHouseUnaryPostfixOperator.getRandom(
    }
}
```

Expression generator

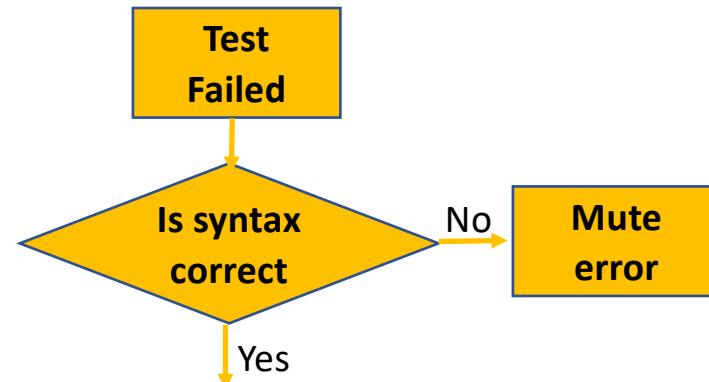
Implement Oracle

Almost the same for different DBMS

```
13  public class ClickHouseTLPWhereOracle extends ClickHouseTLPBase {  
14        
15      public ClickHouseTLPWhereOracle(ClickHouseProvider.ClickHouseGlobalState  
16          state) {  
17          super(state);  
18      }  
19      @Override  
20      public void check() throws SQLException {  
21          super.check();  
22           if (Randomly.getBooleanWithRatherLowProbability()) {  
23              select.setOrderByExpressions(gen.generateOrderBys());  
24          }  
25          String originalQueryString = ClickHouseVisitor.asString(select);  
26          List<String> resultSet = ComparatorHelper.getResultSetFirstColumnAsS  
27            
28          boolean orderBy = Randomly.getBooleanWithRatherLowProbability();  
29          if (orderBy) {  
30              select.setOrderByExpressions(gen.generateOrderBys());  
31          }  
32          select.setWhereClause(predicate);  
33          String firstQueryString = ClickHouseVisitor.asString(select);  
34          select.setWhereClause(negatedPredicate);  
35          String secondQueryString = ClickHouseVisitor.asString(select);  
36          select.setWhereClause(isNullPredicate);  
37          String thirdQueryString = ClickHouseVisitor.asString(select);  
38          List<String> combinedString = new ArrayList<>();  
39          List<String> secondResultSet = ComparatorHelper.getCombinedResultSet  
40              thirdQueryString, combinedString, !orderBy, state, errors);  
41          ComparatorHelper.assertResultSetsAreEqual(resultSet, secondResultSet  
42              state);  
43      }  
44  }
```

|| How to deal with results?

Query can be incorrect



Mute errors

You don't need to implement ideal query generator – bad queries also test your system

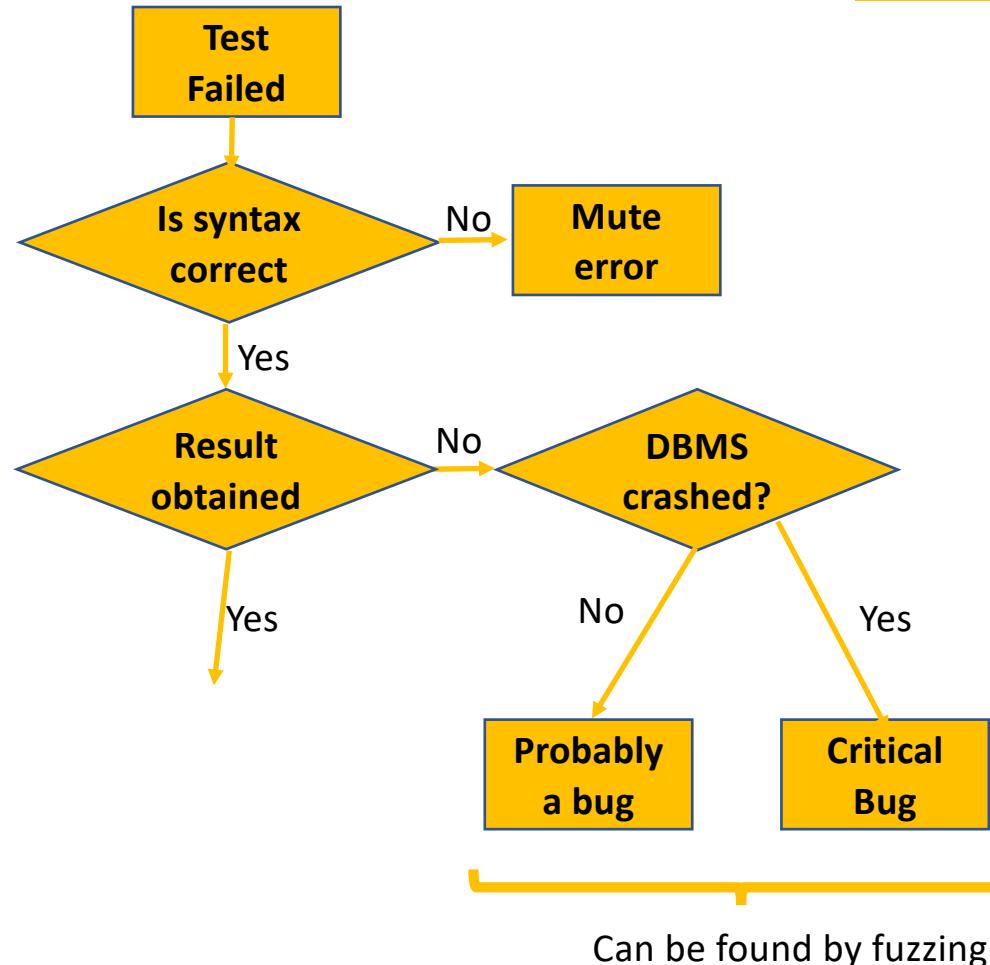
You can mute known bugs and find more

```
1 package sqlancer.clickhouse;
2
3 import sqlancer.common.query.ExpectedErrors;
4
5 public final class ClickHouseErrors {
6
7     private ClickHouseErrors() {
8     }
9
10    @
11    public static void addExpectedExpressionErrors(ExpectedErrors errors) {
12        errors.add("Illegal type");
13        errors.add("Argument at index 1 for function like must be constant");
14        errors.add("Argument at index 1 for function notLike must be constant");
15        errors.add("does not return a value of type UInt8");
16        errors.add("invalid escape sequence");
17        errors.add("invalid character class range");
18        errors.add("Memory limit");
19        errors.add("There is no supertype for types");
20        errors.add("Bad get: has Int64, requested UInt64");
21        errors.add("Cannot convert string");
```

|| How to deal with results?

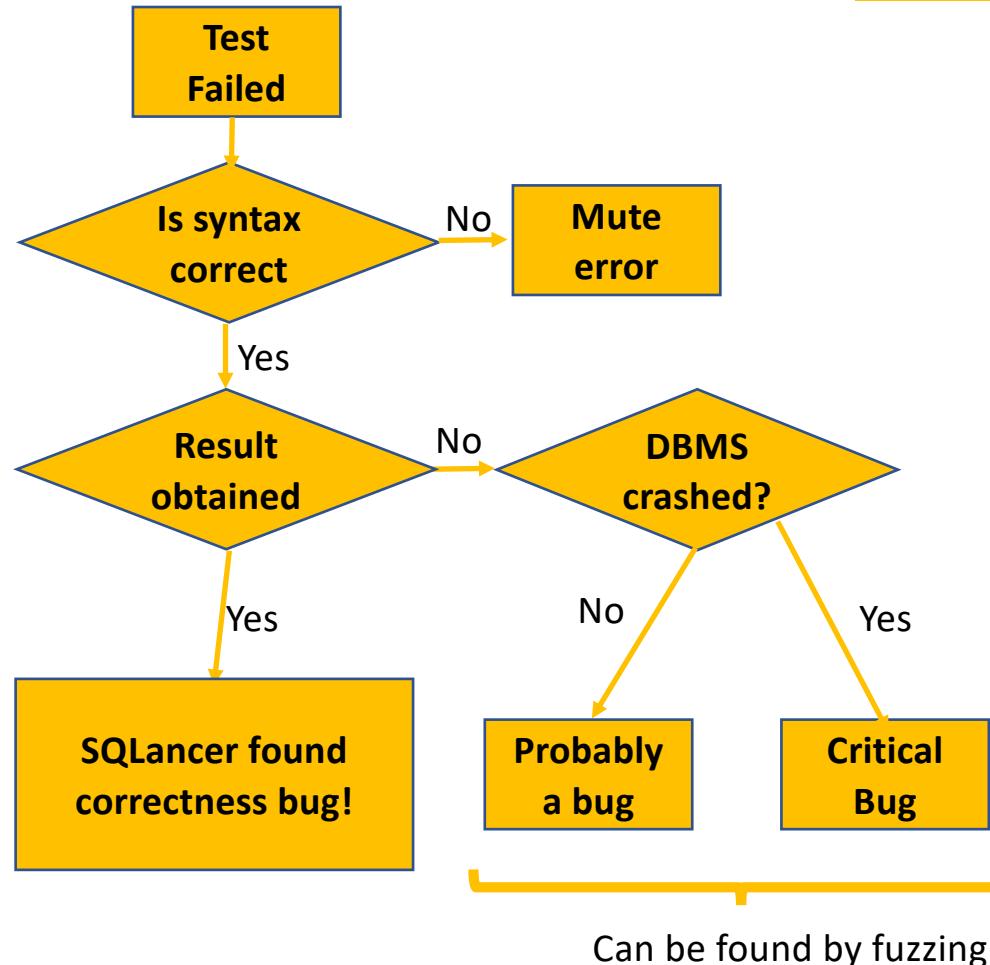
| Query can be incorrect

| DMBS can crash



How to deal with results?

- Query can be incorrect
- DBMS can crash
- If results differ – you found a correctness bug



Add Oracles to your Fuzzer

If you already have a query fuzzer or expression generator you can implement oracles trivially.

You can try implementing NoREC or TLP Where logic by yourself:

- 1) Take *expr* from your generator
- 2) Add some logic to compare results:

```
SELECT * FROM t0  
WHERE expr;
```

```
SELECT expr  
FROM t0;
```

```
SELECT * FROM t;  
  
SELECT * FROM t WHERE expr  
UNION ALL  
SELECT * FROM t WHERE NOT expr  
UNION ALL  
SELECT * FROM t WHERE expr IS NULL;
```

How to contribute?

- Add a new DBMS
- Automatic reduction of test cases (#333)
- Plugin system for DBMS support (#8)
- New test oracles
- Blog posts, tutorials, ...

Summary

ClickHouse

Open Source analytical DBMS for BigData with SQL interface.

- Blazingly fast
- Scalable
- Fault tolerant

2013 Project started → 2016 Open Sourced → 2021 ★15K GitHub

2

SQLancer

Unwatch 23 Star 678 Fork 82



<https://github.com/sqlancer>

SQLancer is an effective and widely-used automatic testing tool to find logic bugs in DBMSs

13

Finding Logic Bugs in DBMSs

Ternary Logic Partitioning OOPSLA '20

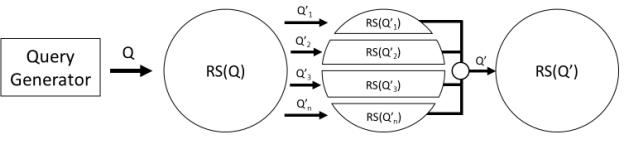
Non-optimizing Reference Engine Construction ESEC/FSE '20

Pivoted Query Synthesis OSDI '20

SQLancer provides three test oracles

64

Query Partitioning



In contrast to differential testing, we need only a single DBMS

76

Metamorphic Testing

Test Case → Execute → Result

Derived Test Case + Derive

This technique is known as metamorphic testing

115

How to make integration with your DBMS

- Import Java or ODBC driver if you have one
- Teach SQLancer to make connection, create database and make a generic query
- A bit more to prepare: Create table with schema and insert generator
- Implement expression generator
- Implement oracle
- Expected error handling
- RUN!

163

Q&A Session

- By the way, you can contribute to SQLancer and ClickHouse



<https://github.com/sqlancer/sqlancer>



<https://github.com/ClickHouse/ClickHouse>