



Async Inserts and Deduplication

Fei.Han@clickhouse.com



Inserts in ClickHouse

- 1 insert will generate at least 1 part.
- We recommend a insert contains at least 1k rows
 - ◇ Ideally 10k - 100k
- If inserts are too small, it's easy to get 'TOO_MANY_PARTS' errors.

Work around for small inserts

<https://clickhouse.com/docs/en/optimize/bulk-inserts>

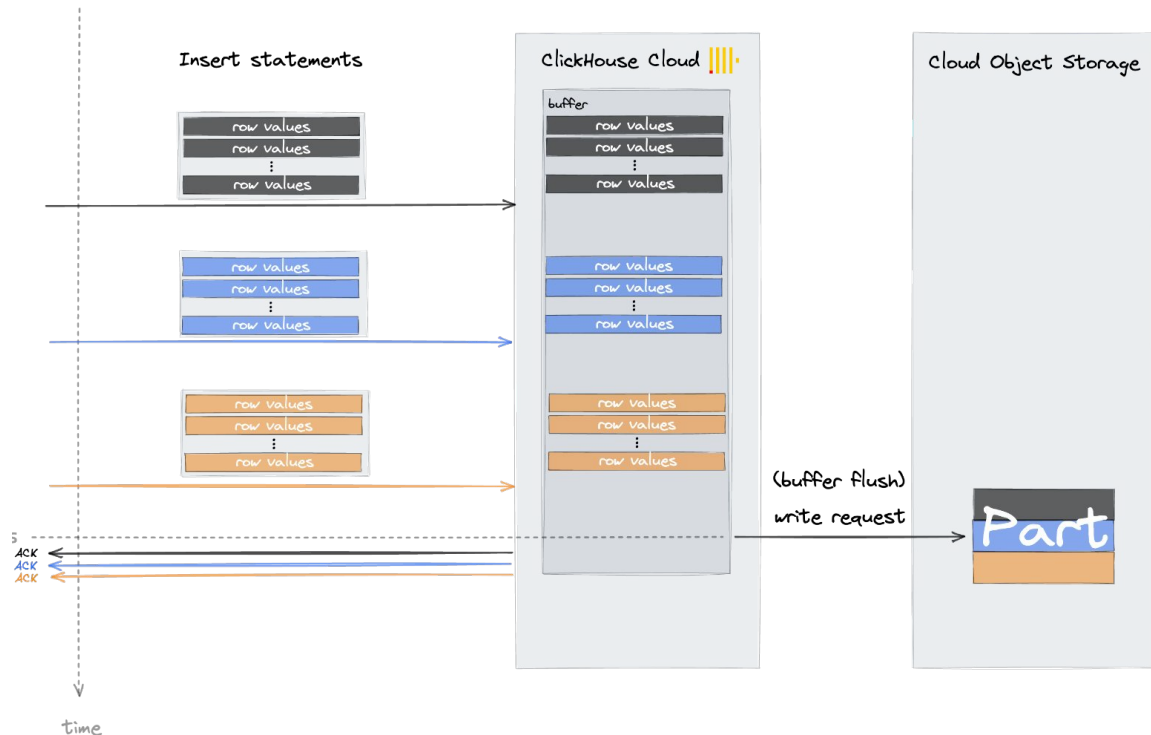
- Use buffer table linking to another table
 - ◇ It's possible to lose data if the buffer is not flushed and server crashes
- Batch data on client side
 - ◇ Need extra user code
 - ◇ Retries costs more



What is async insert ?

We launched “Async Insert” in ClickHouse v2021.11

- “Async insert”s *buffers* inserts inside server to make a larger block.
- Return ACK if and only if the insert is flushed to the storage.
- Support deduplicate since ClickHouse v2023.01



How do we guarantee Idempotence?

Some possible solutions in other systems

- Streaming services
 - ◇ Assign unique id for network packets
- RDBMS
 - ◇ Primary keys / Unique Keys

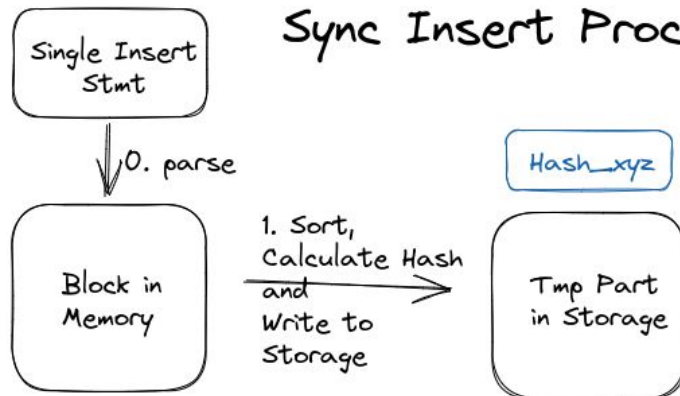
How do we guarantee Idempotence?

Some possible solutions in other systems ClickHouse resolves the problem by:

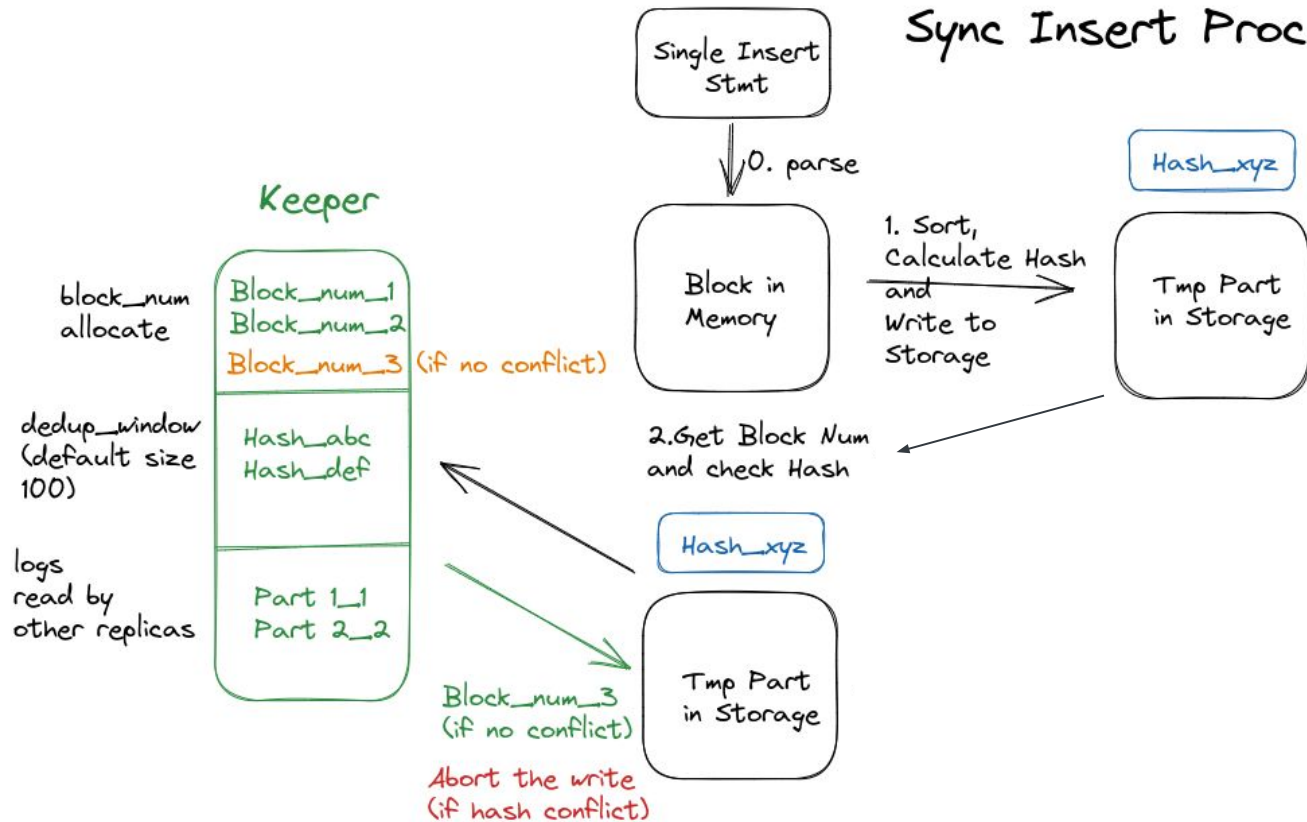
- Streaming services
 - ◇ Assign unique id for network packets
- RDBMS
 - ◇ Primary keys / Unique Keys
- Calculating hash id for an insert
- Check if it conflicts with hash ids in the “dedup_window”
 - ◇ Commit write and write hash id to “dedup_window”, or
 - ◇ Abort write if conflicts



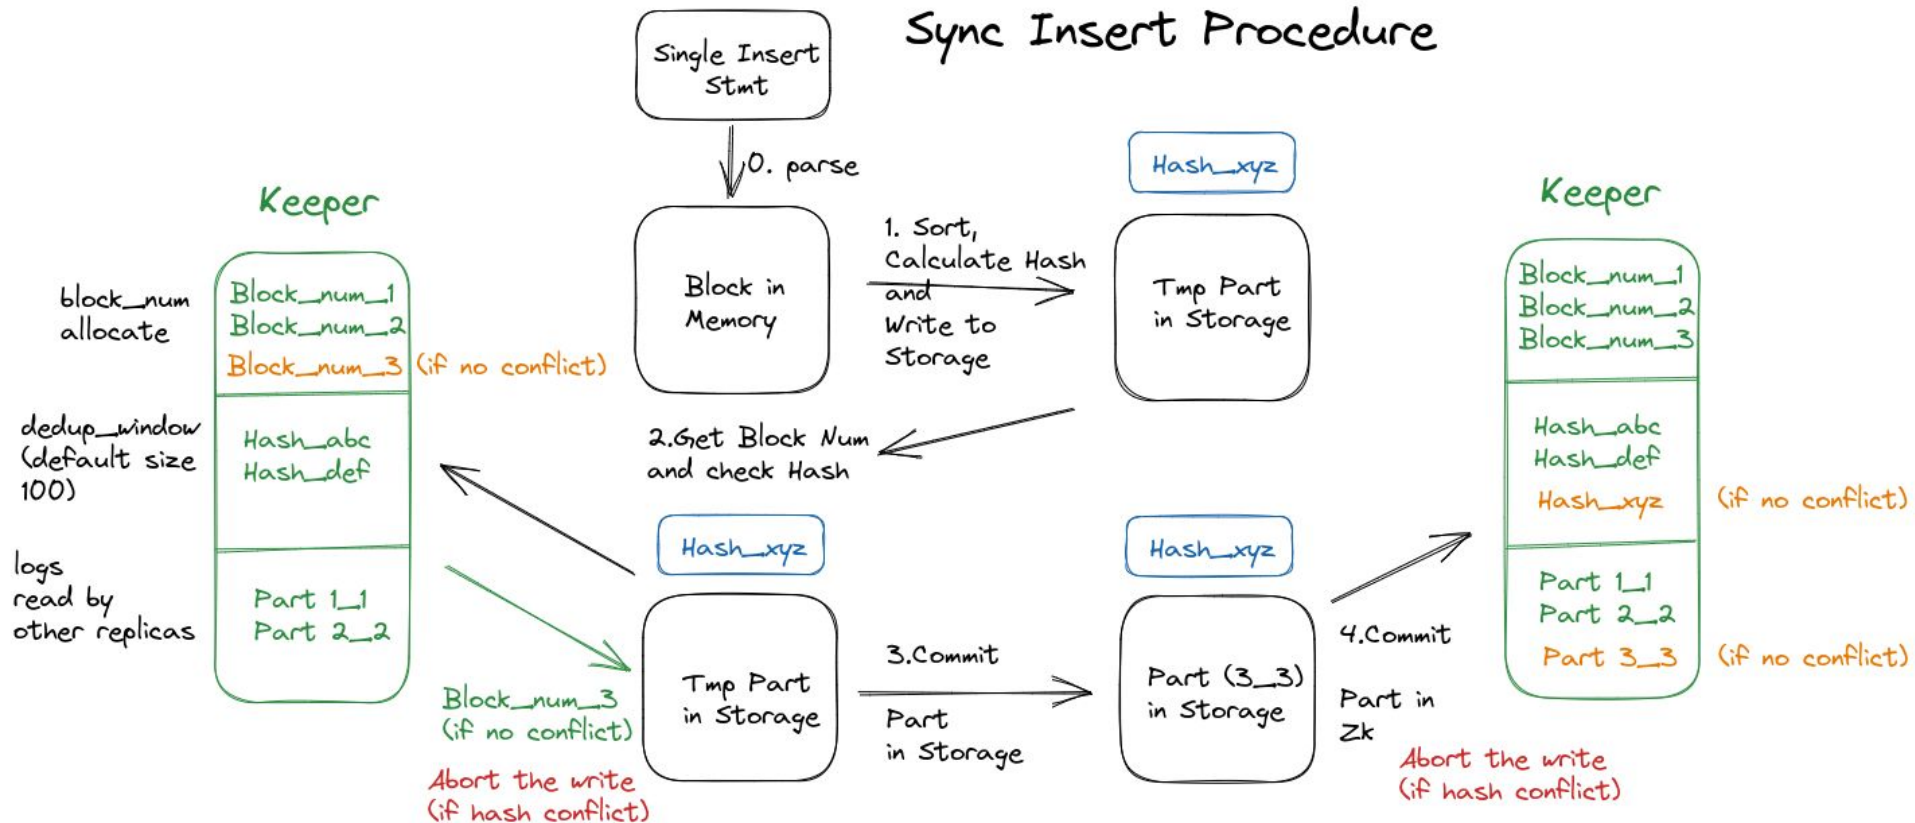
Sync Insert Procedure



Sync Insert Procedure



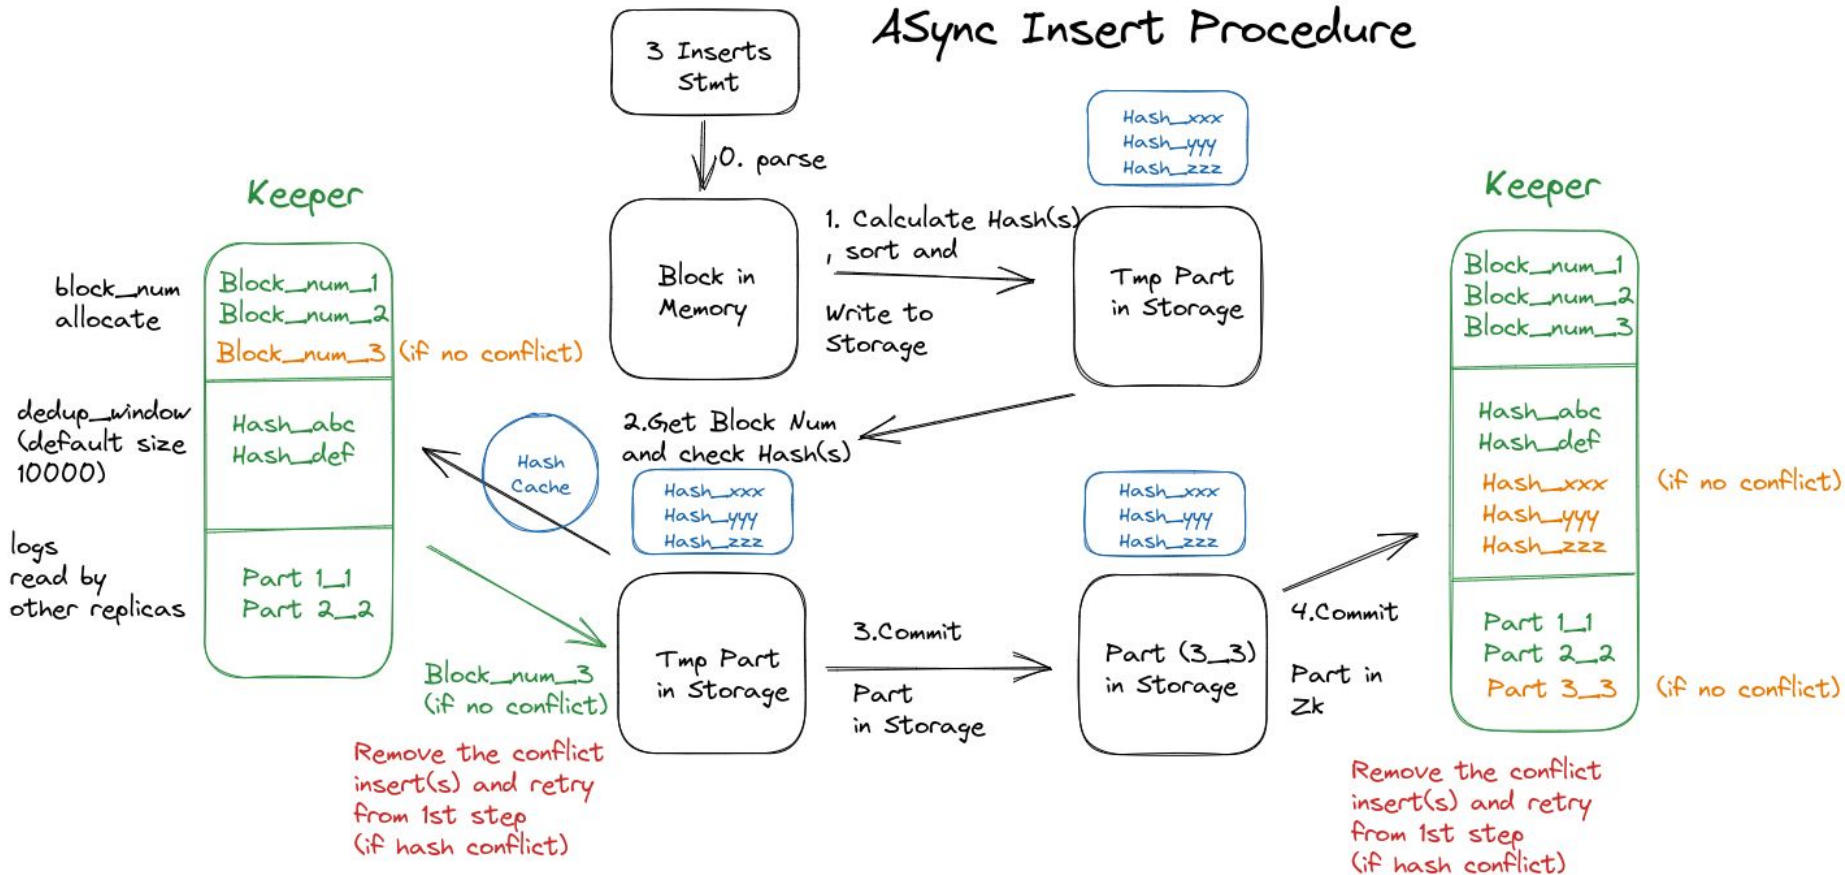
Sync Insert Procedure



Problems with async inserts' deduplication

- Multiple inserts in one block
 - ◇ And we need larger duplication window
- Should retry (partially) when hash id(s) conflict

ASync Insert Procedure



Simple Benchmark

Insert 1,000,000 integers, 200 per insert, 100 threads

benchmark

Percentage of deduplication	Async (deduplicate=0)	Async (deduplicate=1)	Async (wait_for_async_inser=0 deduplicate=1)
0%	20.070	20.116	8.780
5%		21.221	9.359
10%		22.858	9.790
20%		24.092	10.614
40%		27.972	12.435
80%		36.483	15.918



The End

Thank you!