

# Vector Search in ClickHouse

Do you really need a Vector database?

Dale McDiarmid - Product@ClickHouse



**May, 2023**

 ClickHouse

# ***A very* brief history of search**

(more [here](#))

# Bag of Words + Inverted Indices

Documents

doc\_1

The moonlight danced on the surface of the calm lake as the night creatures began to stir.

Tokenization

(split on whitespace ->  
remove punctuation ->  
lowercase ->  
remove plurals)

doc\_2

The sound of rain splashing against the lake was a comforting lullaby for the sleepless night.

doc\_3

As the rain pattered softly on the surface of the tranquil lake under the glow of the moonlight, the night came alive with the soothing sounds of nature.

Bag of Words

the moonlight danced on the surface of the calm lake as the night creatures began to stir the sound of rain splashing against the lake was a comforting lullaby for the sleepless night as the rain pattered softly on the surface of the tranquil lake under the glow of the moonlight the night came alive with the soothing sounds of nature

terms to  
index

Term Index

term frequency doc ids (doc\_id: frequency)

the	14	[ (doc_1, 4), (doc_2, 3), (doc_3, 7) ]
moonlight	2	[ (doc_1, 1), (doc_3, 1) ]
danced	1	[ (doc_1, 1) ]
on	2	[ (doc_1, 1), (doc_3, 1) ]
surface	2	[ (doc_1, 1), (doc_3, 1) ]
of	5	[ (doc_1, 1), (doc_2, 1), (doc_3, 3) ]
calm	1	[ (doc_1, 1) ]
lake	3	[ (doc_1, 1), (doc_2, 1), (doc_3, 1) ]
as	2	[ (doc_1, 1), (doc_2, 1) ]
night	3	[ (doc_1, 1), (doc_2, 1), (doc_3, 1) ]
sound	2	[ (doc_2, 1), (doc_3, 1) ]
rain	2	[ (doc_2, 1), (doc_3, 1) ]
etc...		

# Scoring - Measuring Relevancy - TF-IDF/BM25

TF is the frequency of a term in a document

$$\text{TF-IDF} = \text{TF}(t, d) \times \text{IDF}(t)$$

IDF - The inverse document frequency  
the ratio of documents that include the term

Term Index

term frequency doc ids (doc\_id: frequency)

the	14	[ (doc_1, 4), (doc_2, 3), (doc_3, 7) ]
moonlight	2	[ (doc_1, 1), (doc_3, 1) ]
danced	1	[ (doc_1, 1) ]
on	2	[ (doc_1, 1), (doc_3, 1) ]
surface	2	[ (doc_1, 1), (doc_3, 1) ]
of	5	[ (doc_1, 1), (doc_2, 1), (doc_3, 3) ]
calm	1	[ (doc_1, 1) ]
lake	3	[ (doc_1, 1), (doc_2, 1), (doc_3, 1) ]
as	2	[ (doc_1, 1), (doc_2, 1) ]
night	3	[ (doc_1, 1), (doc_2, 1), (doc_3, 1) ]
sound	2	[ (doc_2, 1), (doc_3, 1) ]
rain	2	[ (doc_2, 1), (doc_3, 1) ]
etc...		

Search: "The sound of the lake" (assume a logical AND)

doc_2			
term	TF	IDF	TF x IDF
the	3	1/14	3/14
sound	1	1/2	1/2
of	1	1/5	1/5
the	3	1/14	3/14
lake	1	1/3	1/3
Total (score)			1.46

doc_3			
term	TF	IDF	TF x IDF
the	7	1/14	7/14
sound	1	1/2	1/2
of	3	1/5	3/5
the	7	1/14	7/14
lake	1	1/3	1/3
Total (score)			2.43

doc\_3 is more relevant



# What is a vector/embedding?

(more than just text)

# Vector vs Embedding

Example 10-dimensional vector

$(0.0318, 0.0356, 0.0693, -0.0147, -0.0417, 0.0014, -0.005, -0.0034, -0.0683, -0.0432)$

Embedding for the word "moonlight"

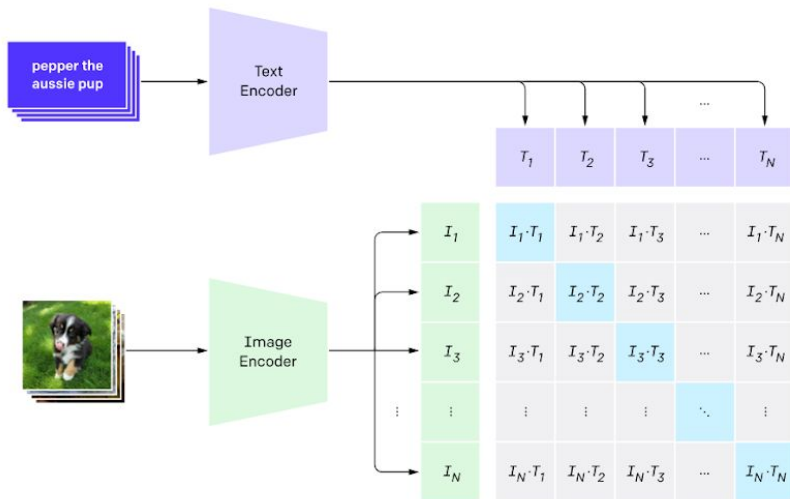
$(0.0318, 0.0356, 0.0693, 0.567, -0.0417, 0.9675)$



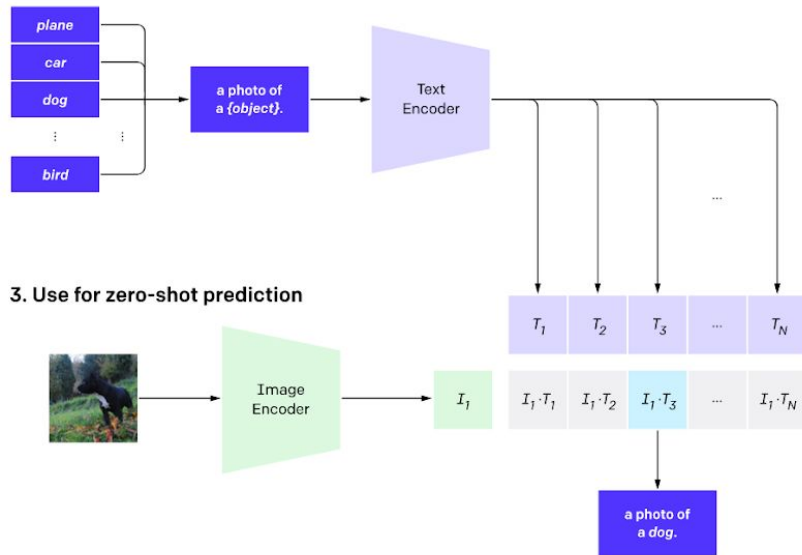
$(\text{color temperature, directionality, mood, brightness, warmth, nighttime})$

# Producing Embeddings - a lot more than text


## 1. Contrastive pre-training



## 2. Create dataset classifier from label text



## 3. Use for zero-shot prediction



CLIP pre-trains an image encoder and a text encoder to predict which images were paired with which texts in our dataset. We then use this behavior to turn CLIP into a zero-shot classifier. We convert all of a dataset's classes into captions such as "a photo of a dog" and predict the class of the caption CLIP estimates best pairs with a given image.

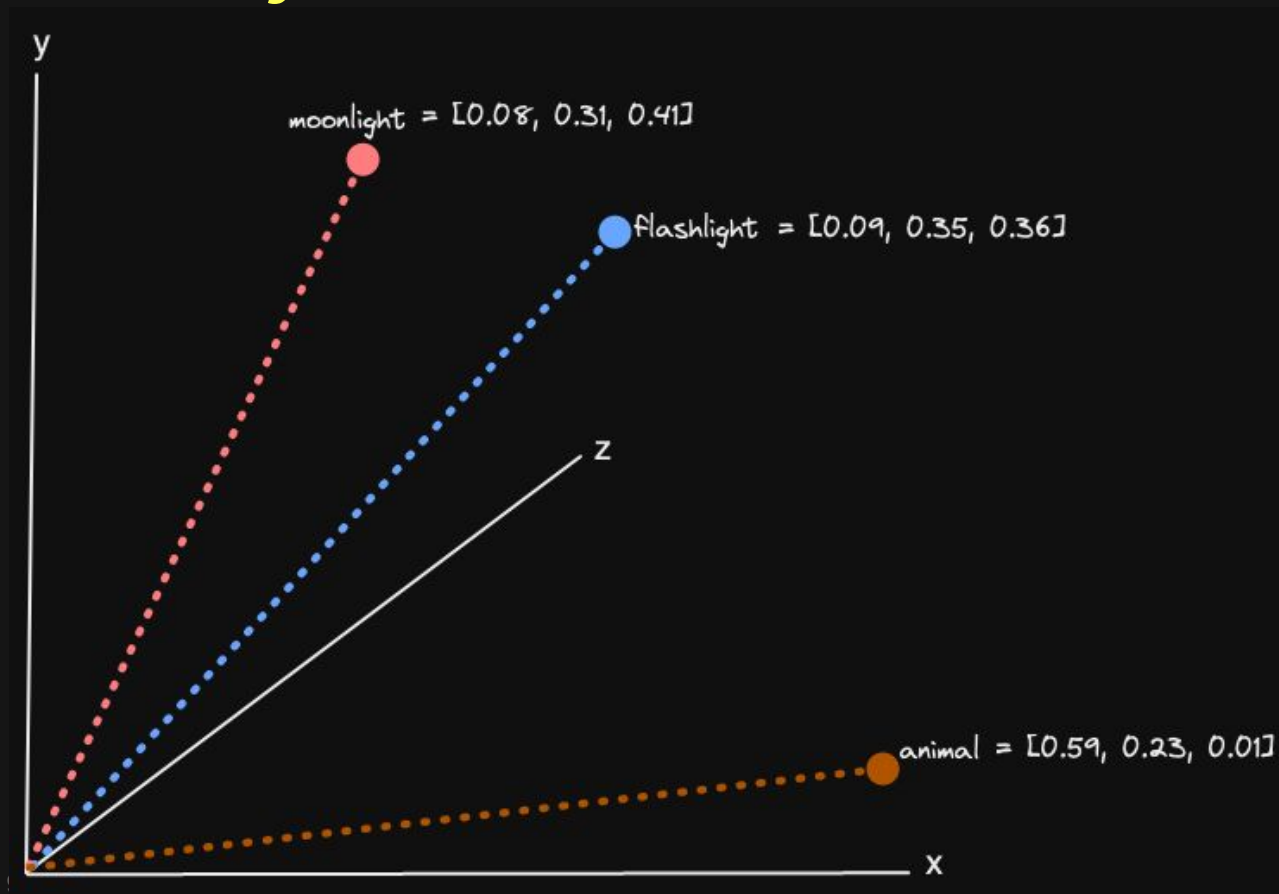
Credit: CLIP -

<https://openai.com/research/clip>

# The importance of distance




# Similarity measured with distance



# Searching vectors

# Exact linear scan





## Approximate Techniques

- Approximate Nearest Neighbour (ANN)
- Trade accuracy for performance
- Image is for Annoy algorithm
- Annoy builds a tree structure by partitioning the space into smaller and smaller hyperplanes for recursive evaluation

# What problems can we solve?

# Possible Applications

**01**

## Recommendations

Relevant to e-commerce websites. Beyond embedding text meaning into vectors, page views and past purchases can be encoded.

**02**

## Question Answering

Equivalent meanings can, however, be encoded with vectors that are close, e.g., X and Y.

**03**

## Image and video search

Search for images and videos based on text and vice versa

**04**

## Fraud Detection

Find similar or dissimilar transactions by encoding users' behaviors or log-in patterns into vectors. These can be anomalous behaviors and prevent fraud.

**05**

## Providing Context

Leveraging vector search to provide contextual content to chat applications powered by APIs like ChatGPT.

**06**

## Multilingual search

Allow cross-language searching with the same concept in two languages encoded to the same vector.



# Vector Search in ClickHouse

- **Vectors are just Array(Float32) - 1.5x to 2x compression**
- **Matching is just a distance function - mainly Euclidean or Cosine but others supported**
- **Full SQL support and aggregations**
- **Approximate Nearest Neighbour supported through Annoy. Used to optimize distance functions.\***
- **UDFs are useful for integrating embedding generation**

**\*Experimental.**



```
SELECT
    url,
    caption,
    L2Distance(image_embedding, [0.5736801028251648...0.2516217529773712]) AS score
FROM laion WHERE similarity >= 0.2
ORDER BY score ASC
LIMIT 2
FORMAT Vertical
```

10 rows in set. Elapsed: 1.603 sec. Processed 10.00 million rows, 32.70 GB (**6.24 million rows/s.**, **20.40 GB/s.**)



# When to use ClickHouse for Vector Search?

# ClickHouse for Vector Search

You have/need:

- High performance and scalable linear matching using many cores
- Metadata with your vectors for which you'll benefit from high compression and extremely fast querying/filtering
- Full SQL support in your querying inc. Joins and aggregations
- A very large dataset and do not wish to be memory bound or a small dataset where linear matching is fine (~millions)
- Data in ClickHouse and wish to complement with vectors
- An existing pipeline for generating vectors from a model and don't need tight coupling



# Demo

# LAIION dataset

- [5 billion embeddings](#) created with performance analysis in mind
- Created using a [CLIP \(Contrastive Language–Image Pre-training\)](#) multi-modal model. Allows images to be classified with a relevant category for previously unknown images not seen during the training process
- Trained on 400m image/caption pairs
- Each row has image metadata and a text + image embedding (768 dimensions)
- Allows searching for images with text and vice versa using Euclidean distance for matching
- We used a [2.2b english subset](#) (4 billion vectors in total)
- We combined embeddings with metadata and made available in Parquet\*



# bfloat16

## Floating Point Formats

bfloat16: Brain Floating Point Format

Range:  $\sim 1e^{-38}$  to  $\sim 3e^{38}$



fp32: Single-precision IEEE Floating Point Format

Range:  $\sim 1e^{-38}$  to  $\sim 3e^{38}$



fp16: Half-precision IEEE Floating Point Format

Range:  $\sim 5.96e^{-8}$  to 65504



```
INSERT INTO laion_bfloat16
SELECT
    arrayMap(x -> reinterpretAsFloat32(bitAnd(reinterpretAsUInt32(x),
4294901760)), image_embedding) AS image_embedding,
    arrayMap(x -> reinterpretAsFloat32(bitAnd(reinterpretAsUInt32(x),
4294901760)), text_embedding) AS text_embedding
FROM laion
```



**END**



# Speakers



**Aaron Katz**

CEO @ ClickHouse



**Alexey Milovidov**

CTO @ClickHouse



**Yury Izrailevsky**

President and  
VP of Engineering

**01**

# Section title with number

Author or subtitle

# Content heavy slide (2 columns)

# **Content slide - less text heavy**



# What is ClickHouse?

## Open source

Developed since 2009  
OSS 2016  
28k+ Github stars  
1k+ contributors  
300+ releases

## column-oriented

Best for aggregations  
Files per column  
Sorting and indexing  
Background merges

## distributed

Replication  
Sharding  
Multi-master  
Cross-region

## OLAP database

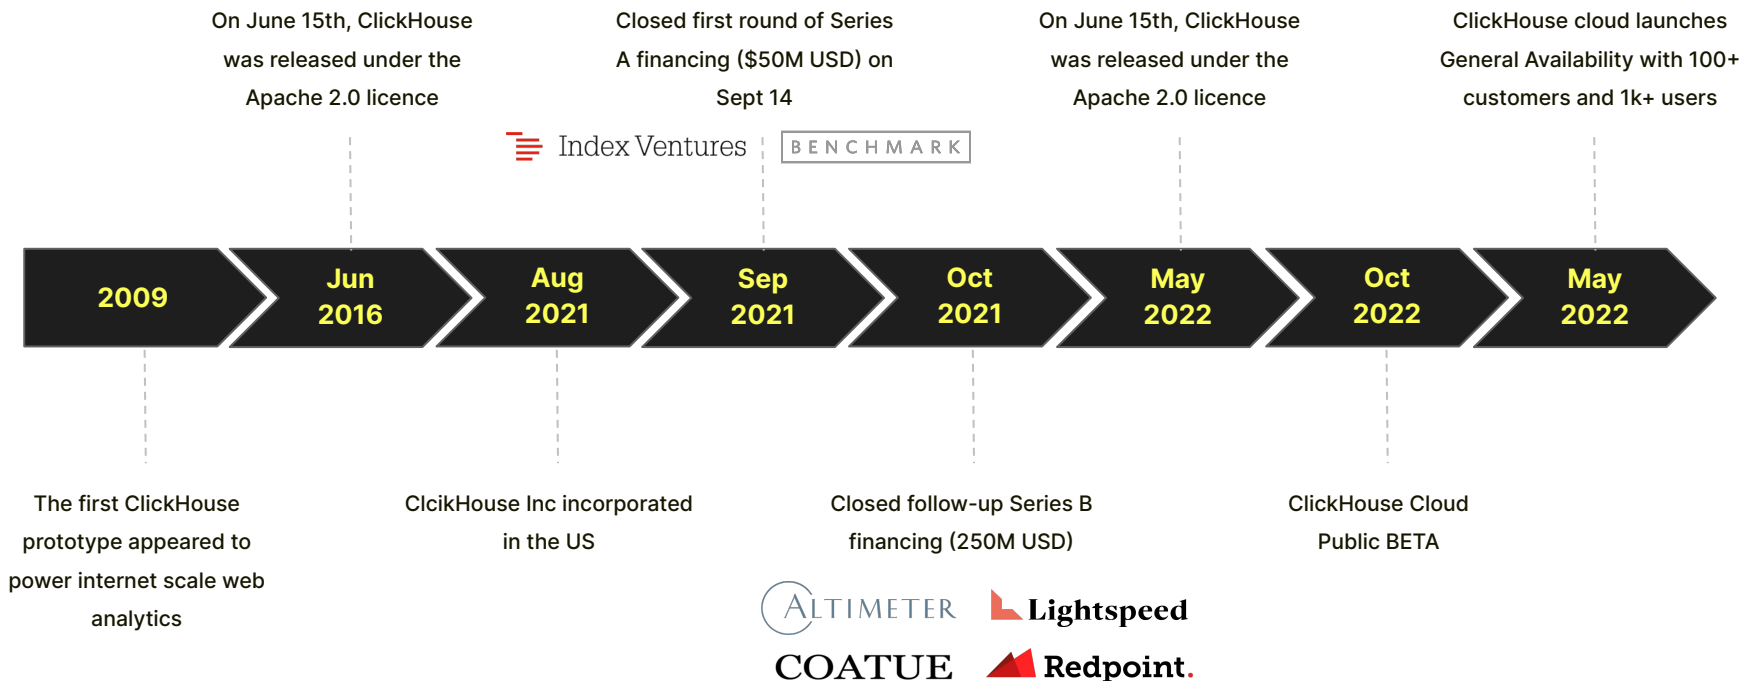
Analytics use cases  
Aggregations  
Visualizations  
Mostly immutable data



# What is ClickHouse? (highlighted)

Open source	column-oriented	distributed	OLAP database
<ul style="list-style-type: none"><li>Developed since 2009</li><li>OSS 2016</li><li>28k+ Github stars</li><li>1k+ contributors</li><li>300+ releases</li></ul>	<ul style="list-style-type: none"><li>Best for aggregations</li><li>Files per column</li><li>Sorting and indexing</li><li>Background merges</li></ul>	<ul style="list-style-type: none"><li>Replication</li><li>Sharding</li><li>Multi-master</li><li>Cross-region</li></ul>	<ul style="list-style-type: none"><li>Analytics use cases</li><li>Aggregations</li><li>Visualizations</li><li>Mostly immutable data</li></ul>

# ClickHouse Journey



## Feature highlight

- list item 1
- list item 2
- list item 3

another one	Running
monitoring-internal	Directly assigned permissions
operator-internal	Directly assigned permissions
sql-console	default_role
default	default_role

My service	Running
monitoring-internal	Directly assigned permissions
operator-internal	Directly assigned permissions
default	default_role
sql-console	default_role
sql-console:cristina.albu@clickhouse.com	default_role



# Two cards highlight

## Title here

Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

## Title here

Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

# Three cards highlight

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

# Use this slide for code (dark)

```
SELECT
    toStartOfMonth(upload_date) AS month,
    sum(view_count) AS `Youtube Views`,
    bar(sum(has_subtitles) / count(), 0.55, 0.7, 100) AS `% Subtitles`
FROM youtube
WHERE (month >= '2020-08-01') AND (month <= '2021-08-01')
GROUP BY month
ORDER BY month ASC
```

13 rows in set. Elapsed: 0.823 sec **Processed 1.07 billion** rows, 11.75 GB (1.30 billion rows/s., 14.27 GB/s.)



# ClickHouse Cloud - from 0 to 1 in under a year



## Private preview

May - July 2022

### Serverless hosted ClickHouse for key design partners

- AWS with 3 regions initially
- Limited scalability within predefined limits
- Basic cloud console for ClickHouse service and user management
- Strong security and privacy with SOC 2 Type I compliance

## Public Beta

Oct 2022

### Serverless hosted ClickHouse for early adopters

- Ecosystem of first-party connectors to onboard and work with data
- Automatic scaling up and down for compute
- Integrated billing using pay as you go pricing model
- Enhanced security features such as Private Link, IP Filtering, Auditing

## General Availability

Dec 2022

### Serverless hosted ClickHouse for the broader market

- Enhanced cloud console for analytics and operational controls
- Support for more AWS regions and AWS marketplace billing
- Uptime SLA and additional operations tools
- Advanced security features with SOC 2 Type II compliance



“

*Rokt has been an eager partner of ClickHouse as we modernize our analytics stack. By offloading operations to the experts our developers are focused on delivering the best experience possible while the business scales. We we are thrilled to see the path ClickHouse is forging.*

Attributor, role, company

**ROKT**

# Use this slide for code (light)

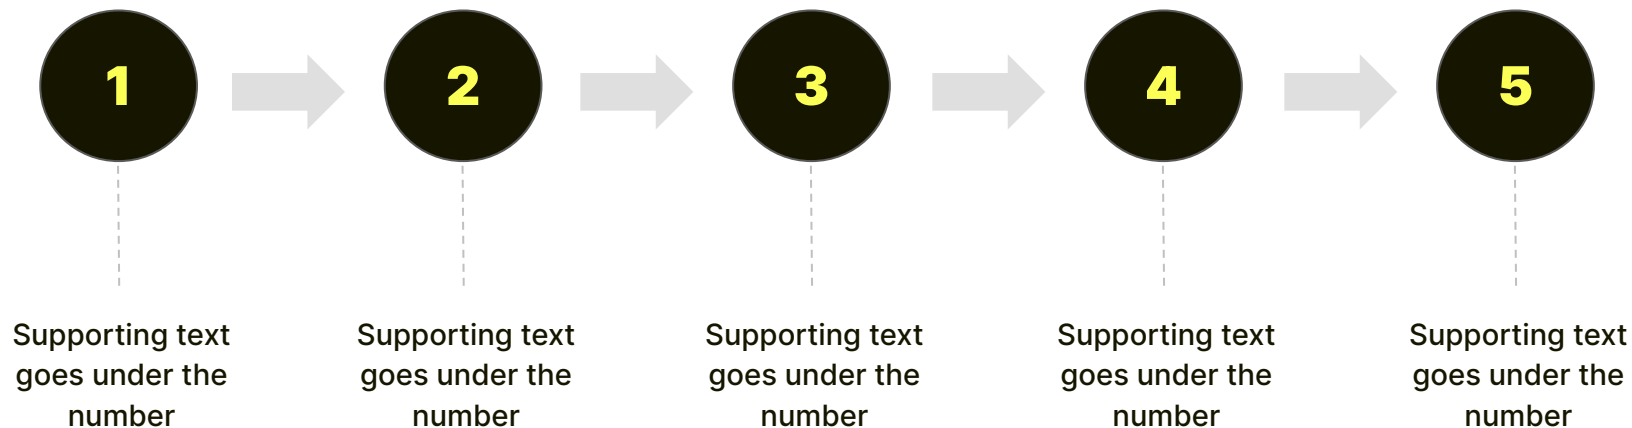
```
SELECT
    toStartOfMonth(upload_date) AS month,
    sum(view_count) AS `Youtube Views`,
    bar(sum(has_subtitles) / count(), 0.55, 0.7, 100) AS `% Subtitles`
FROM youtube
WHERE (month >= '2020-08-01') AND (month <= '2021-08-01')
GROUP BY month
ORDER BY month ASC
```

13 rows in set. Elapsed: 0.823 sec **Processed 1.07 billion rows**, 11.75 GB (1.30 billion rows/s., 14.27 GB/s.)



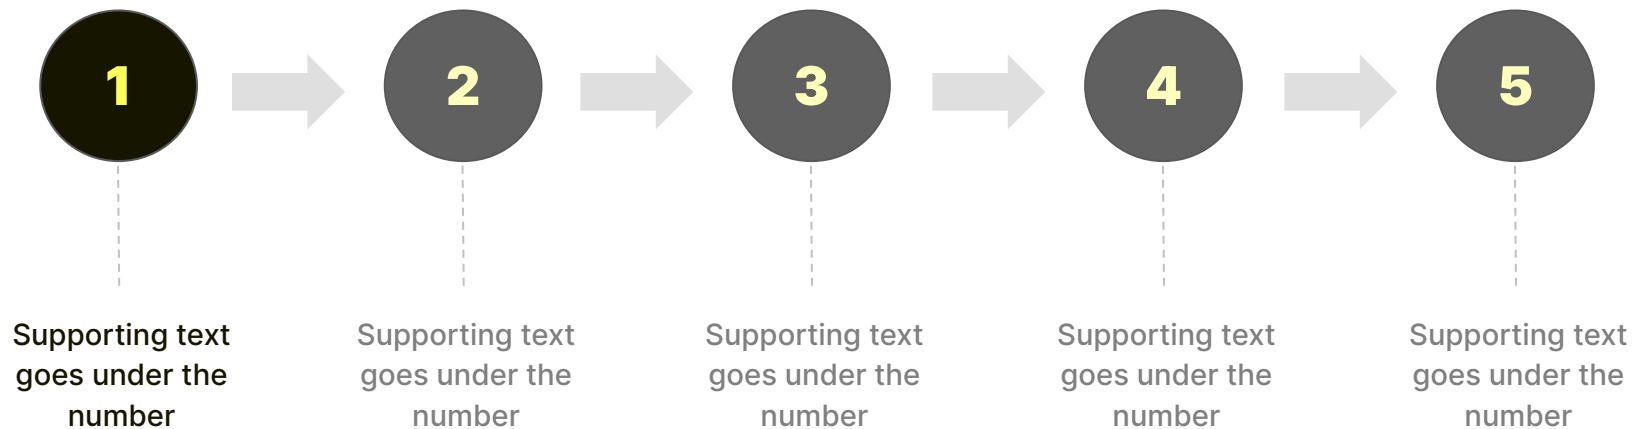
# Process diagram, 5 ideas

Here you could describe the topic of the section, or not



# Process diagram, 5 ideas

Highlighting one of the steps





# Big number treatment

**2.3k**

**Header here**

Supporting text goes here,  
under the header

**1.2M**

**Header here**

Supporting text goes here,  
under the header

**45**

**Header here**


Supporting text goes here,  
under the header

# Features



## Blazing fast

Uses all available hardware to its full potential to process each query as fast as possible. Peak processing performance for a single query stands at more than 2 terabytes per second.



## Fault tolerant

Supports async replication and can be deployed across multiple datacenters. All nodes are equal, which allows avoiding having single points of failure.



## Easy to use

ClickHouse is simple and works out-of-the-box. Simplifies data processing by instantly processing structured data using a user-friendly SQL dialect and eliminating non-standard API requirements.



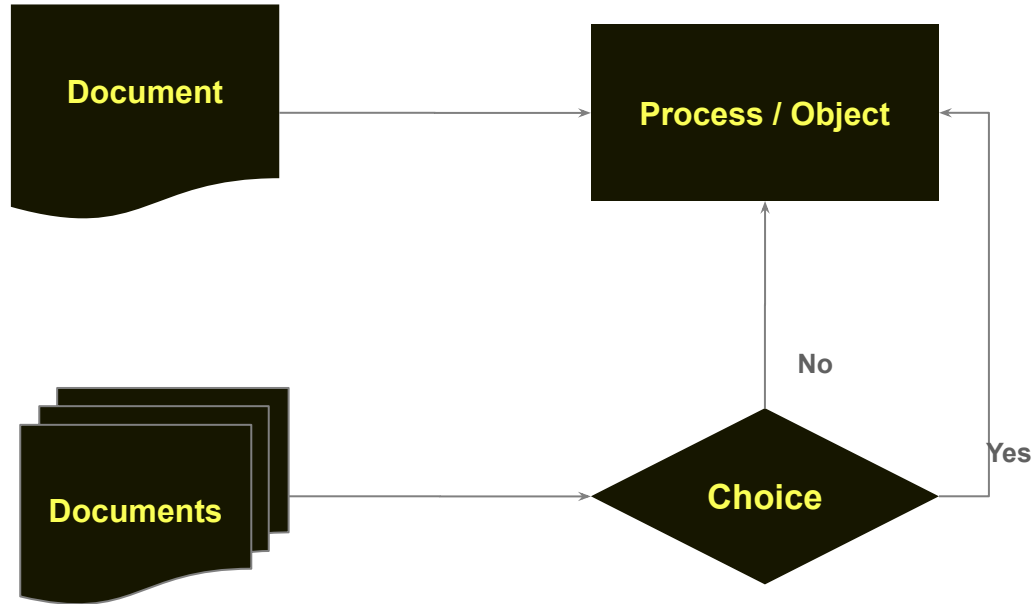
## Highly reliable

Can be configured as a purely distributed system located on independent nodes, without any single points of failure. It also includes a lot of enterprise-grade security features and fail-safe mechanisms against human errors.


# Table format

Header 1	Header 2	Header 3	Header 4	Header 5

# For architecture/business process drawings



# Icons for light background



## Speakers (dark mode)



**Aaron Katz**

CEO @ ClickHouse



**Alexey Milovidov**

CTO @ClickHouse



**Yury Izrailevsky**

President and  
VP of Engineering



# Table of contents - many topics (dark mode)

**01**

**Content title**

Presenter / brief description

**04**

**Content title**

Presenter / brief description

**02**

**Content title**

Presenter / brief description

**05**

**Content title**

Presenter / brief description

**03**

**Content title**

Presenter / brief description

**06**

**Content title**

Presenter / brief description



# Table of contents - fewer topics(dark mode)

**01**

**Content title**

Presenter / brief content description

**02**

**Content title**

Presenter / brief content description

**03**

**Content title**

Presenter / brief content description

**04**

**Content title**

Presenter / brief content description





# Section title

Subtitle or description of the section.  
Remove if not needed

01

# Section title

Subtitle or description of the section.  
Remove if not needed

# Content heavy slide



# **Content slide - less text heavy**



## Feature highlight

- Item 1
- Item 2
- Item 3

another one	Running
monitoring-internal	Directly assigned permissions
operator-internal	Directly assigned permissions
sql-console	default_role
default	default_role
My service	Running
monitoring-internal	Directly assigned permissions
operator-internal	Directly assigned permissions
default	default_role
sql-console	default_role
sql-console:cristina.albu@clickhouse.com	default_role

# Two cards highlight

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

# Three cards highlight

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet her

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet here

## Title here

### Subtitle or a brief description

- Another bullet here
- Another bullet here
- Another bullet here
- Another bullet her



“

*Rokt has been an eager partner of ClickHouse as we modernize our analytics stack. By offloading operations to the experts our developers are focused on delivering the best experience possible while the business scales. We we are thrilled to see the path ClickHouse is forging.*

Attributor, role, company

**ROKT**



# Features



## Blazing fast

Uses all available hardware to its full potential to process each query as fast as possible. Peak processing performance for a single query stands at more than 2 terabytes per second.



## Fault tolerant

Supports async replication and can be deployed across multiple datacenters. All nodes are equal, which allows avoiding having single points of failure.



## Easy to use

ClickHouse is simple and works out-of-the-box. Simplifies data processing by instantly processing structured data using a user-friendly SQL dialect and eliminating non-standard API requirements.



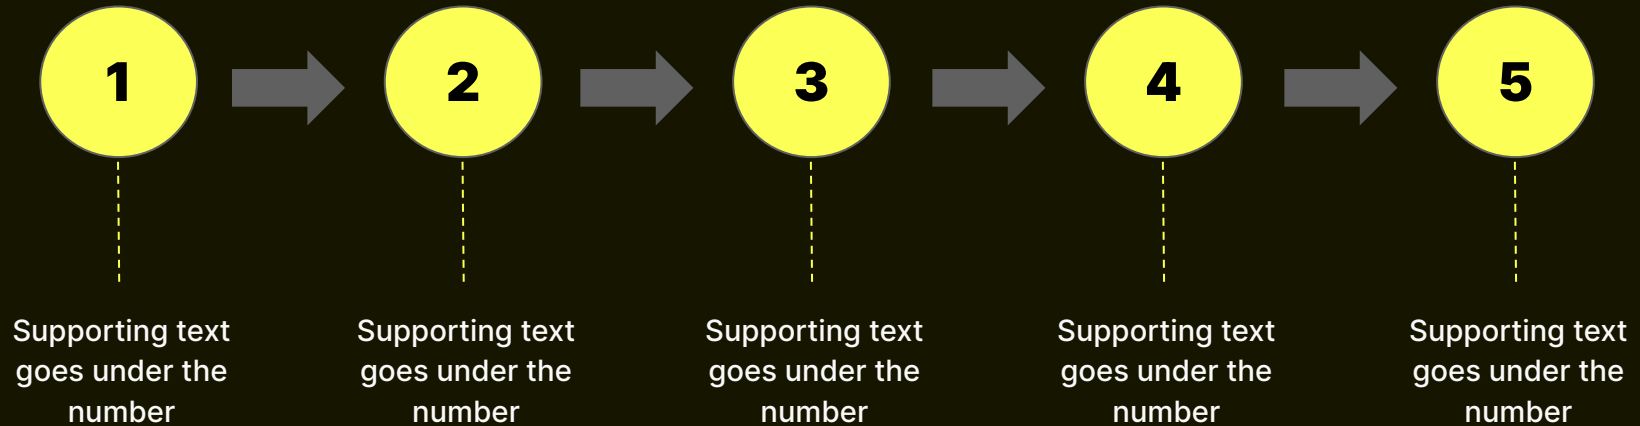
## Highly reliable

Can be configured as a purely distributed system located on independent nodes, without any single points of failure. It also includes a lot of enterprise-grade security features and fail-safe mechanisms against human errors.




# Process diagram, 5 ideas

Here you could describe the topic of the section, or not



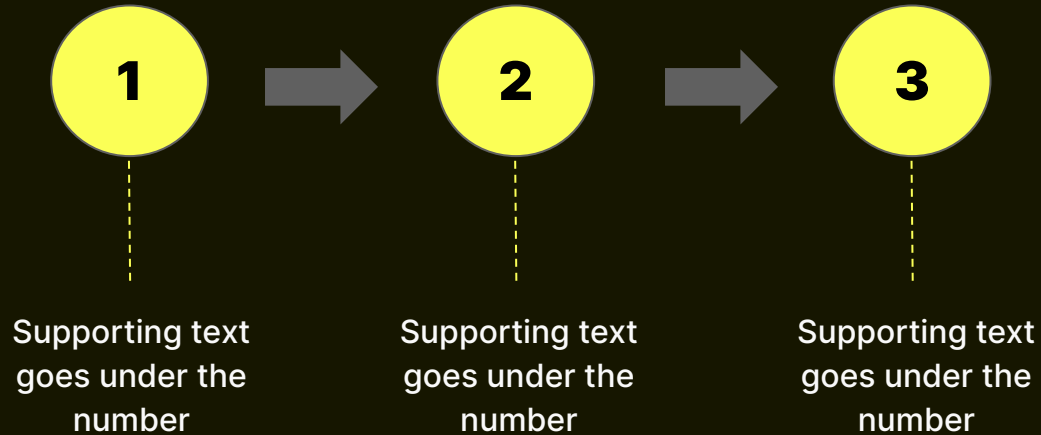
# Process diagram, 5 ideas

Highlighting one of the steps



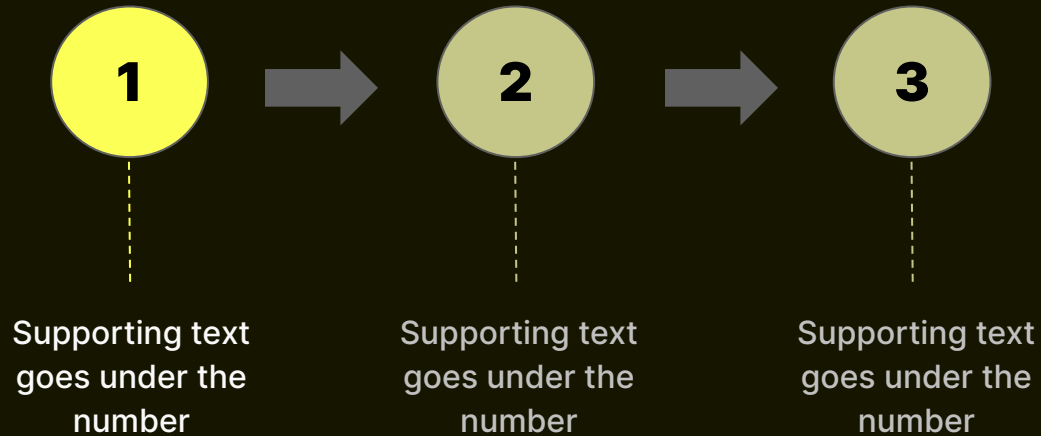
# Process diagram, 3 ideas

Here you could describe the topic of the section, or not



# Process diagram, 3 ideas

Highlighting one of the ideas



# Big number treatment

**1.2k**

**Header here**

Supporting text goes here,  
under the header

**1M**

**Header here**

Supporting text goes here,  
under the header

**45**

**Header here**

Supporting text goes here,  
under the header



# Use this slide for code

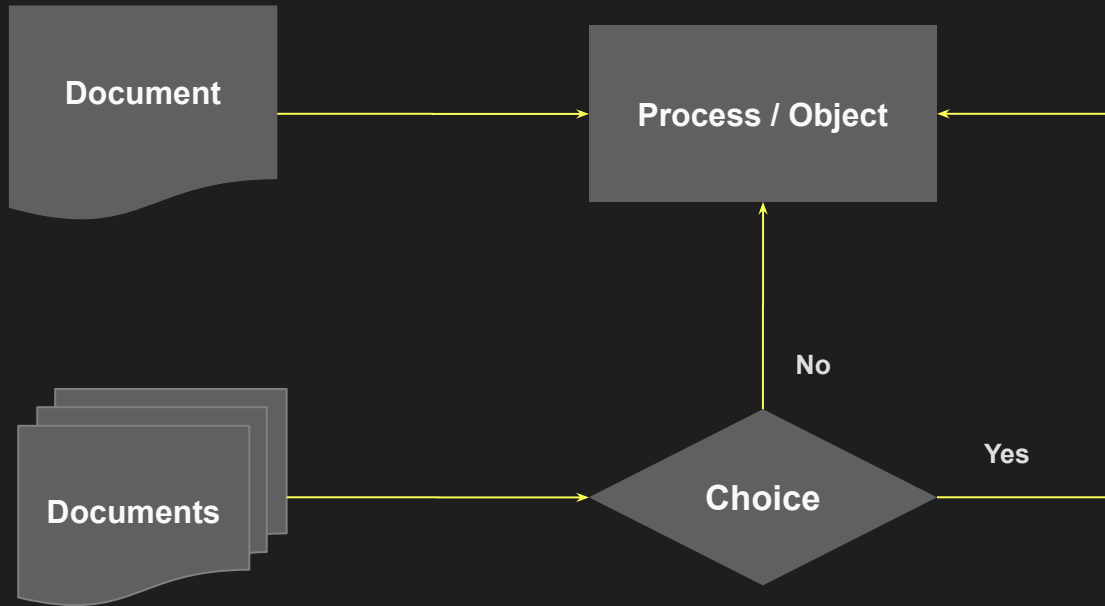
Use template colors to highlight code. Feel free to remove this paragraph.

```
SELECT
    toStartOfMonth(upload_date) AS month,
    sum(view_count) AS `Youtube Views`,
    bar(sum(has_subtitles) / count(), 0.55, 0.7, 100) AS `% Subtitles`
FROM youtube
WHERE (month >= '2020-08-01') AND (month <= '2021-08-01')
GROUP BY month
ORDER BY month ASC
```

13 rows in set. Elapsed: 0.823 sec **Processed 1.07 billion rows**, 11.75 GB (1.30 billion rows/s., 14.27 GB/s.)



# For architecture / business process drawings





# Icons for dark background

