

# Migrating 100 Billion rows of financial data

- Existing Setup
- Challenges with Postgres and need for migration
- Migration Strategy and the nuances
- Single node Postgres to Multi-node Clickhouse(s)
- Benefits

# WhoAml

- Ritesh Shrivastav
- Software Engineer at Zerodha
- Over 7 years of building and scaling Console
- @johnnybravo-xyz <ritesh@shrivastav.xyz>

# Existing Setup

- Console stores and processes reports for over 10 million clients
- Over 100 billion rows across 8 tables using ~40TB of storage
- Single node Postgres that serves live traffic
- One more supplementary Postgres to aid ETL-pipeline
- This talk is about migrating data from this supplementary Postgres.

# Challenges with Postgres

- Pushed single-node Postgres to its limits
- Hardware requirements for maintaining these nodes are prohibitively expensive
- Storage was also becoming a bottleneck.
- ETL pipelines became choke point
  - Ingesting data took 3.5 hours on average
  - Low throughput while ingesting data into multiple tables parallelly

# Why Clickhouse?

- SQL support with more goodies
- Queries are lightning fast
- Data ingestion too is lightning fast
- Data Compression
- Materialized Views
- Easy to scale on commodity server(s)
- Cons?
  - Update/Deletes have to be planned well
  - Data merge timelines are non-definite

# Migration

- Make sure existing queries can be changed/adapted to CH's SQL
- Finding the right schema
  - Table Engines
  - Data Types
  - Compression Codecs

# Sample Table Schema

- Postgres

```
CREATE TABLE orders (  
  user_id          varchar  
  order_date       date  
  order_punch_time datetime  
  order_exec_time  datetime  
  trade_type       varchar  
  quantity         decimal  
  price            decimal  
  exchange         varchar  
  symbol           varchar  
  secondary_id     varchar  
  order_id         varchar  
)  
PARTITION BY RANGE(order_date)
```



- Clickhouse

```
CREATE TABLE orders (  
  user_id          String  
  order_date       Date  
  order_time       Array(DateTime)  
  trade_type       LowCardinality(String)  
  qty_price        Array(Decimal)  
  exch_symbol      Array(String)  
  order_ids        Array(String)  
)  
ENGINE = MergeTree  
PARTITION BY toYYYYMM(order_date)  
ORDER BY (order_date, user_id)  
SETTINGS index_granularity = 365
```

# Migration

- Convert data into CSVs with format that fits the new schema
- Move them via block storages or use HTTP interface

```
$ psql> \COPY (  
  SELECT  
    user_id, order_date,  
    trade_type,  
    JSON_BUILD_ARRAY(order_punch_time, order_exec_time) AS order_time,  
    JSON_BUILD_ARRAY(price, quantity) AS qty_price  
    JSON_BUILD_ARRAY(exchange, symbol) AS exch_symbol,  
    JSON_BUILD_ARRAY(secondary_id, order_id) AS order_ids  
  FROM orders  
  WHERE  
    order_date >= '2017-01-01'  
    AND order_date <= '2017-12-31'  
) TO '/tmp/ch-orders/partition-1-2017.csv' CSV HEADER
```

```
$ cat /tmp/ch-orders/partition-1-2017.csv | clickhouse-client -d console --query "INSERT INTO orders FORMAT CSVWithNames"
```



# Sharding Data into multiple nodes

- Don't repeat mistake of single-node setup like Postgres
- Clickhouse scales linearly so, its easy to deploy multiple instances.
- Physically shard data into multiple nodes
- Deal with multiple shards in app logic than to get into menace of zookeeper and all.
  - CH has improved with CH-Keeper and we plan to take a look at it again in future

# Benefits

- We have six EC2 instances for the price of one huge EC2 instance used for single-node Postgres
- 50% savings on storage, can be improved further with use of codecs
- Data ingestion is 10x faster
- Extracting data in bulk for ETL-pipeline takes 10 mins compared to 140 mins from previous setup
  - Ingest data in parallel to multiple tables