

# History of ClickHouse at Ahrefs

Yasunari Watanabe

# Agenda

- **About us**
- Road to ClickHouse adoption
- Upstream patch: In-memory marks compression

# About me

- Joined Ahrefs ~2 years ago
- Backend developer
- Working on our web crawler



# About Ahrefs: Who we are

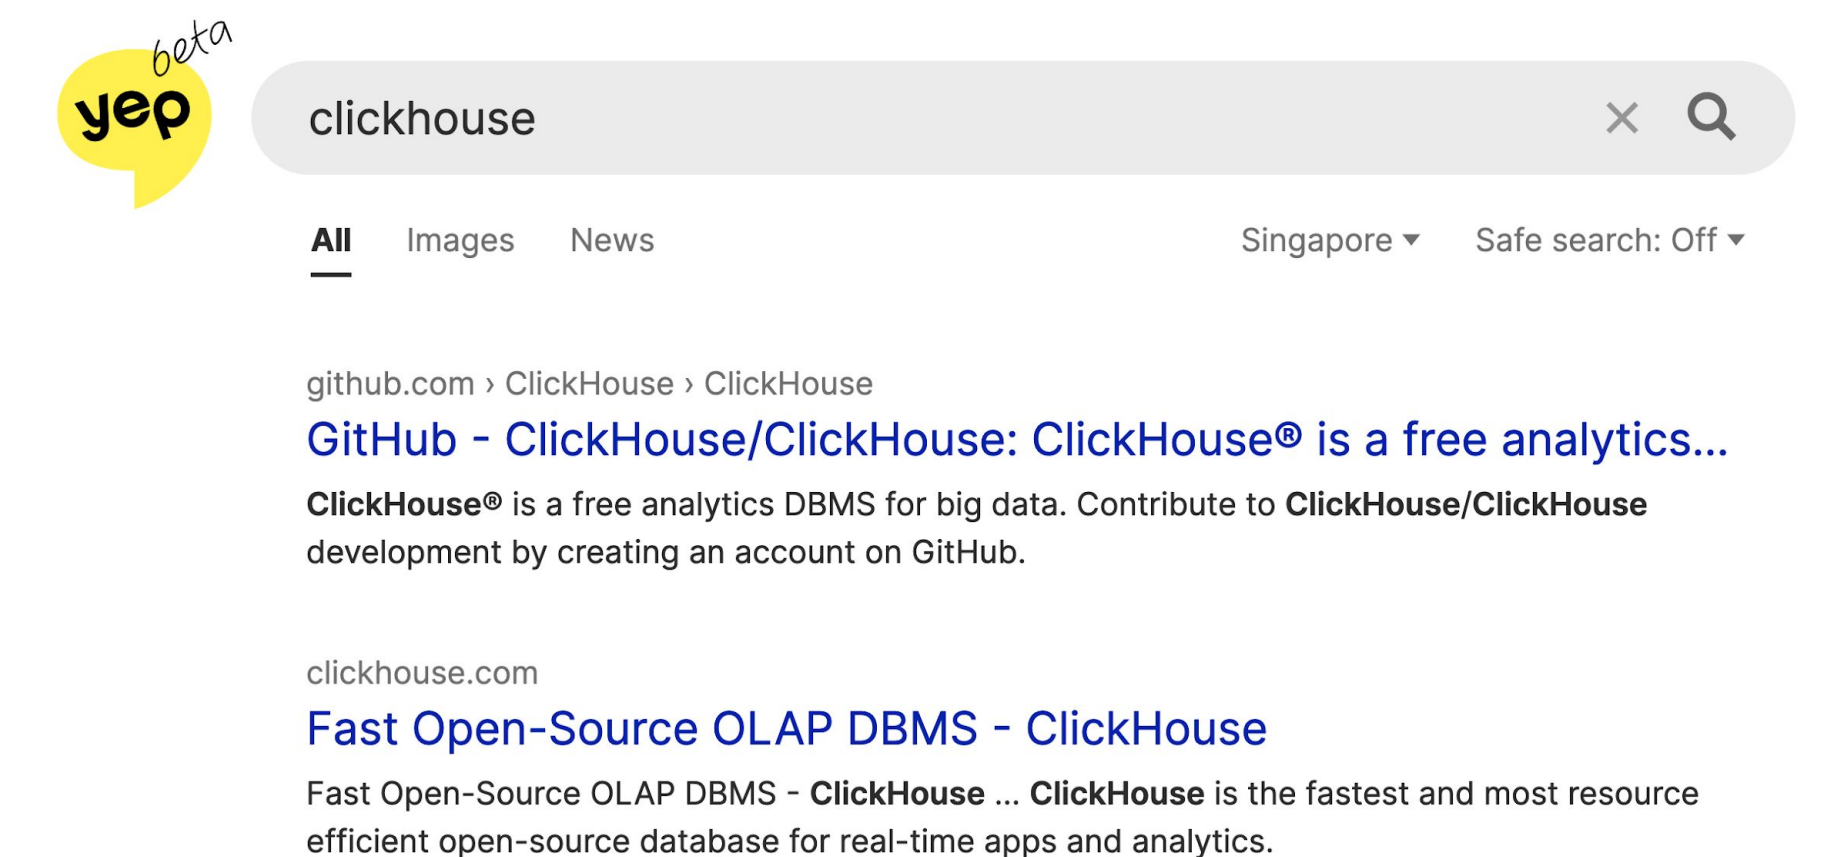
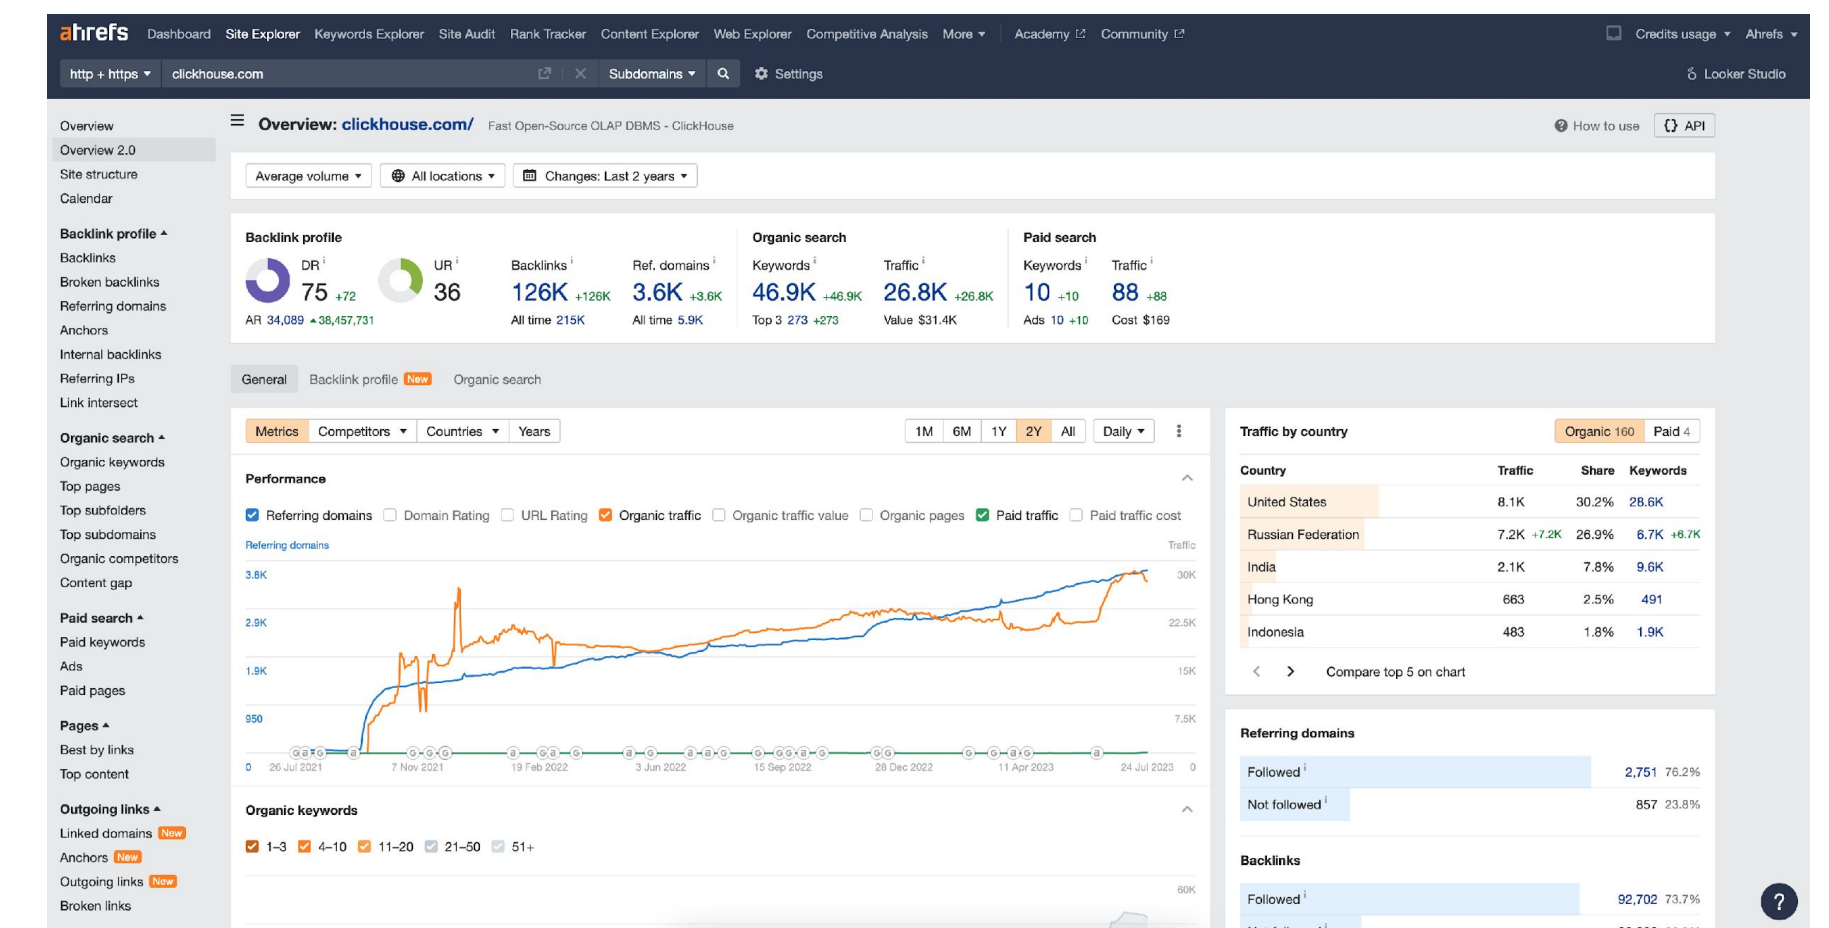
- Founded in 2010, HQ in Singapore
- ~120 people, two thirds remote
- 50% engineers/data scientists
  - 19 backend developers, including 1 ClickHouse developer
  - 6 DevOps engineers






# About Ahrefs: What we do

- Process big data about the Web and transform into relevant indicators
- SEO metrics
- Our new search engine, Yep



# About Ahrefs: Data flow

Web pages



Metrics





Search results

**crawl**

**process**

**store**

**query**



# Quick facts

- Most active web crawler in SEO industry
- World's largest index of live backlinks
- Running on our own hardware

ahrefs.com/big-data

Every minute, we crawl

**5M**

Pages

**3300**

Servers

**596k**

CPU cores

**33PB**

HDD

There are

**170T**

Rows in our key-value database

**4PB**

RAM


**346PB**

SSD

External backlinks history

**35.0T**

Records



Pages in index

**357.4B**

Domains (post-vetting)

**162.7M**

Internal backlinks

**24.7T**

# Agenda

- About us
- **Road to ClickHouse adoption**
- Upstream patch: In-memory marks compression



# Early days

- Tried all available solutions (Cassandra, Hypertable); none fast enough
- Developed custom solution optimized for crawling the web with limited resources (tens of servers, circa 2010)
- QFS to handle unevenness among shards
- Elasticsearch for various non-crawler use cases

# Worked well, but...

- Limitations of custom storage
  - Lack of versatile querying API
  - Lack of bells and whistles
- Evolving feature requirements, size of the Web, our infra...

# Enter ClickHouse

- Started experimenting with ClickHouse in 2019
- Similarities in design, but ClickHouse has dedicated time + resources to features and optimizations
  - SQL interface
  - Many input/output formats
- Notable difference: column-based

# Migration to ClickHouse

- Custom storage
  - Migrated part of the web crawler a few years ago
  - Rest of migration ongoing
  - Challenge: need to keep everything running during migration
- Elasticsearch
  - Most indices rebuilt in ClickHouse



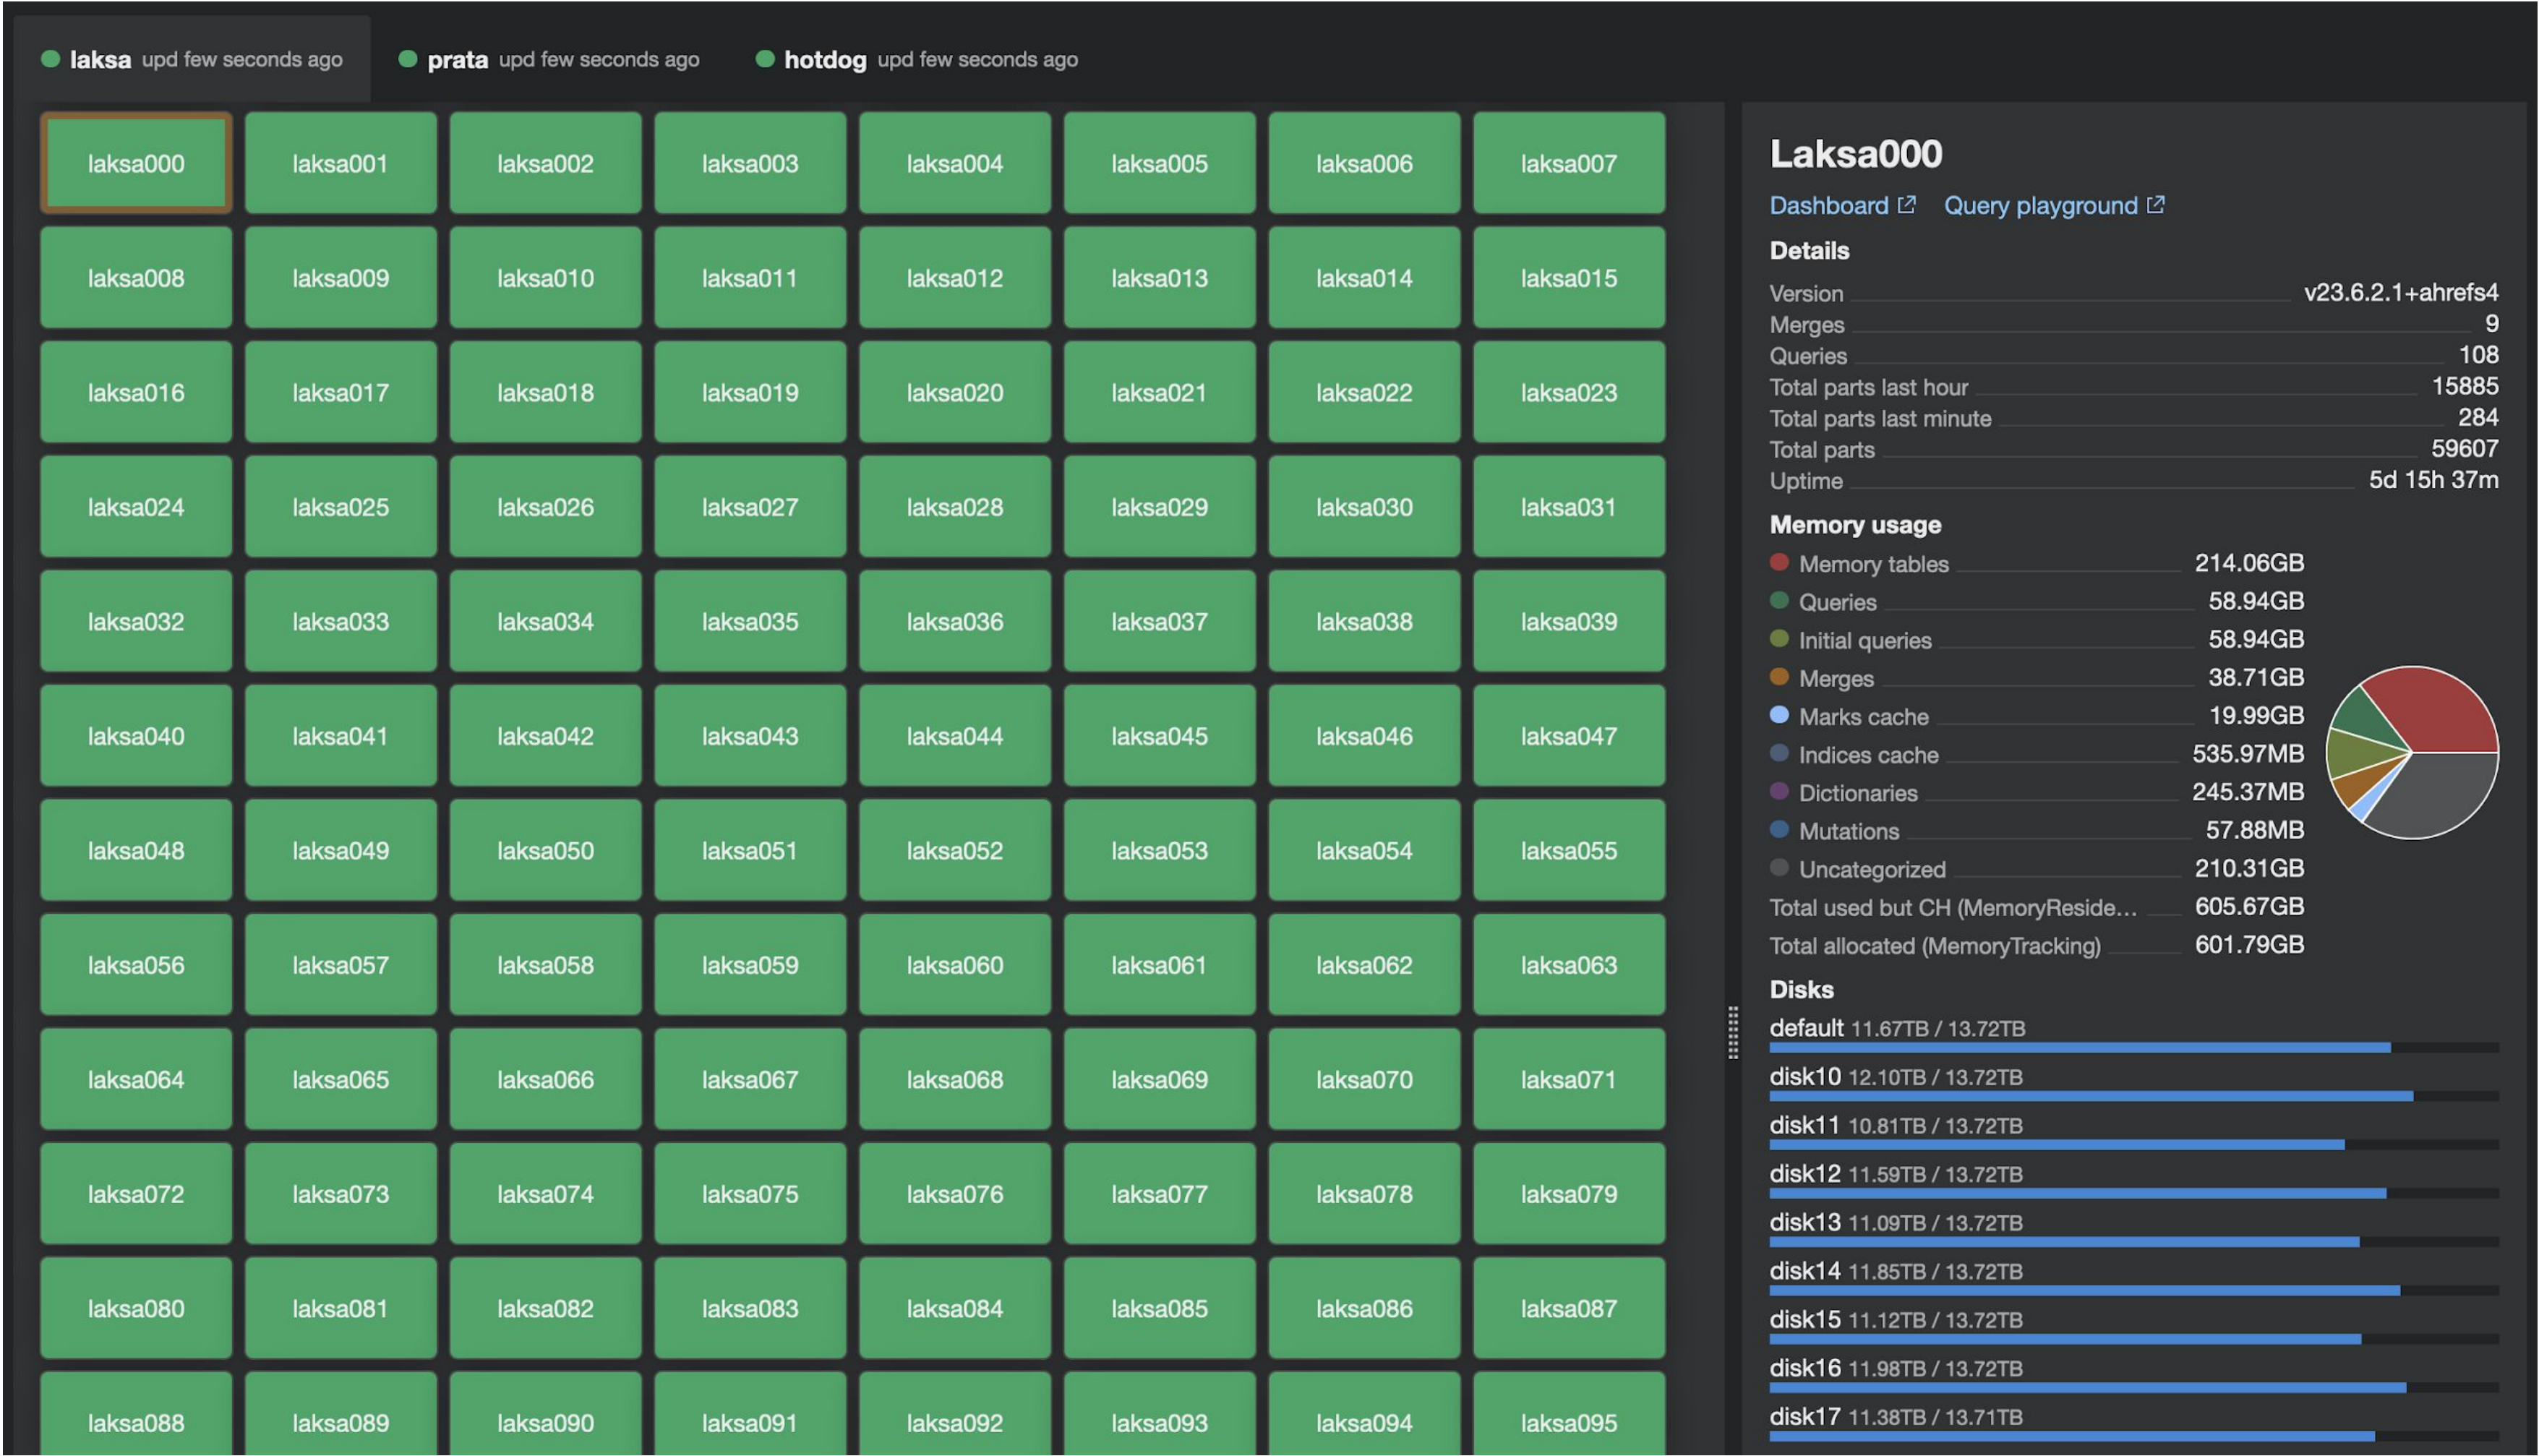
# ClickHouse deployment

- Multiple clusters each with hundreds of hosts
- Geo-replicated main cluster, designated read/write replicas
- Large tables with trillions of rows, tens of columns
- More capacity to come

# ClickHouse in our workflow

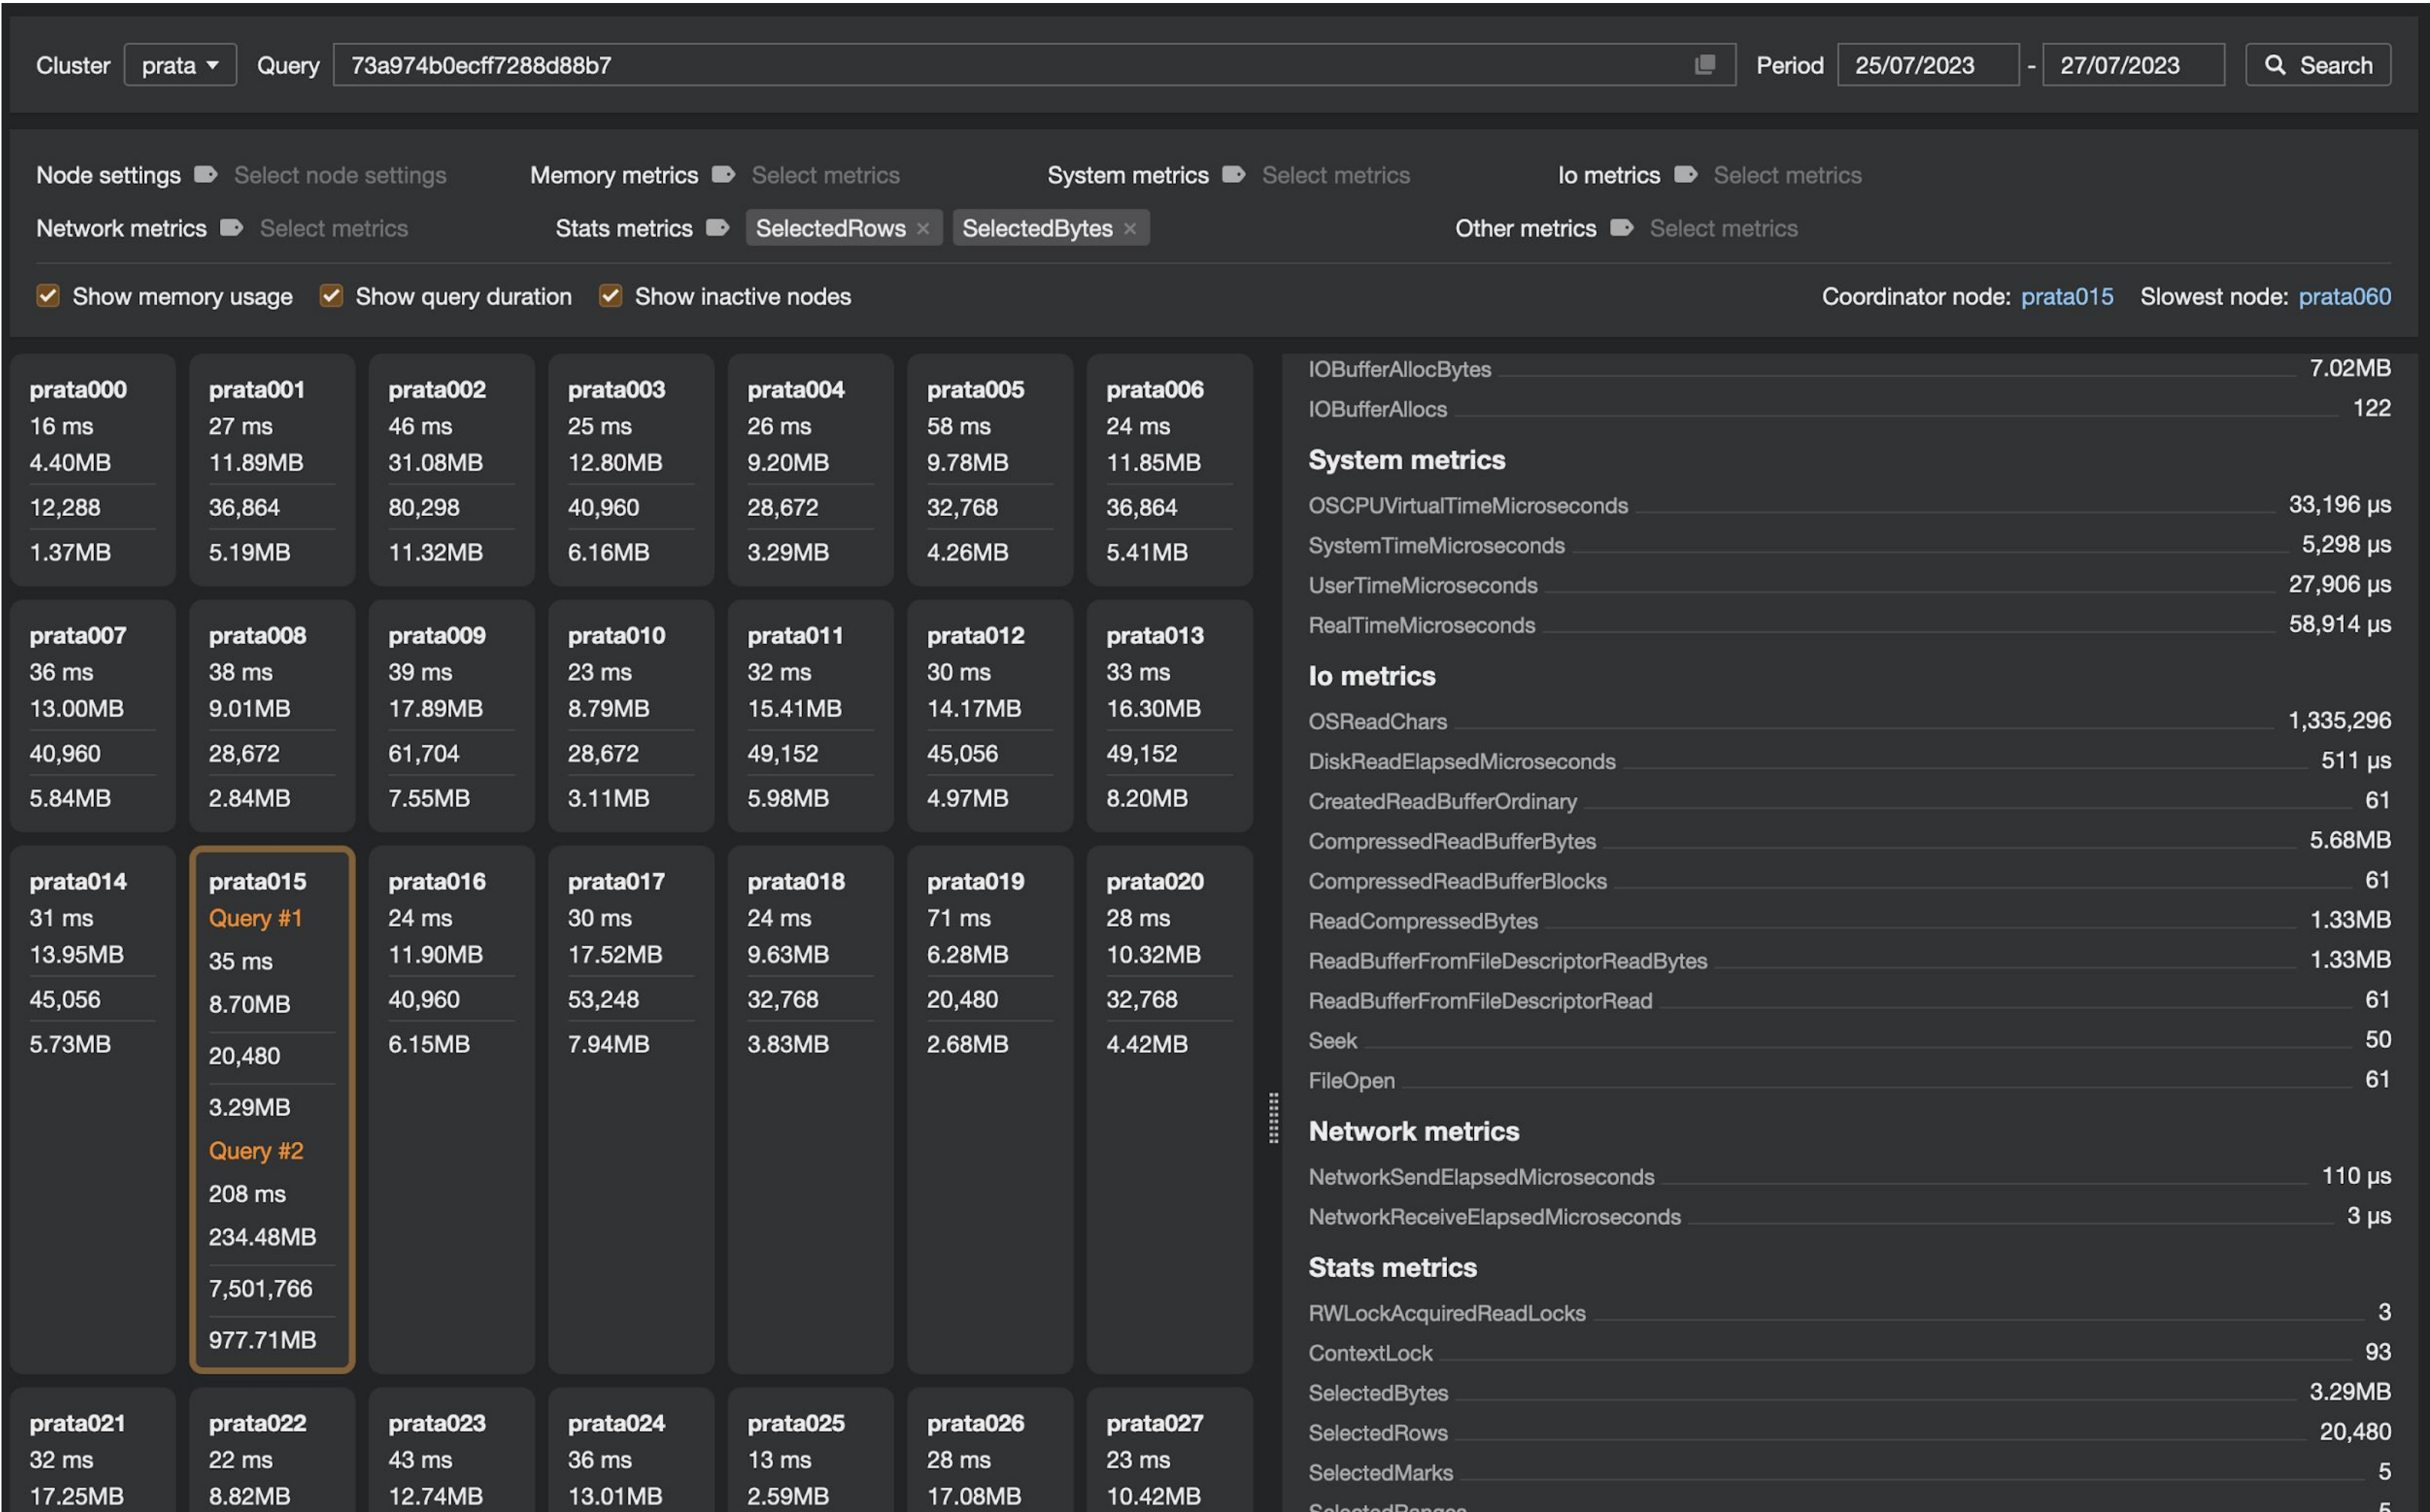
- Techniques
  - Optimizing insert: custom code for partitioned buffers
  - Fetch/attach to move parts between servers efficiently
- Heavy tooling on top of SQL interface
  - Metaprogramming for tight integration with application code (in OCaml)

# Tools: Cluster monitoring





# Tools: Query analyzer



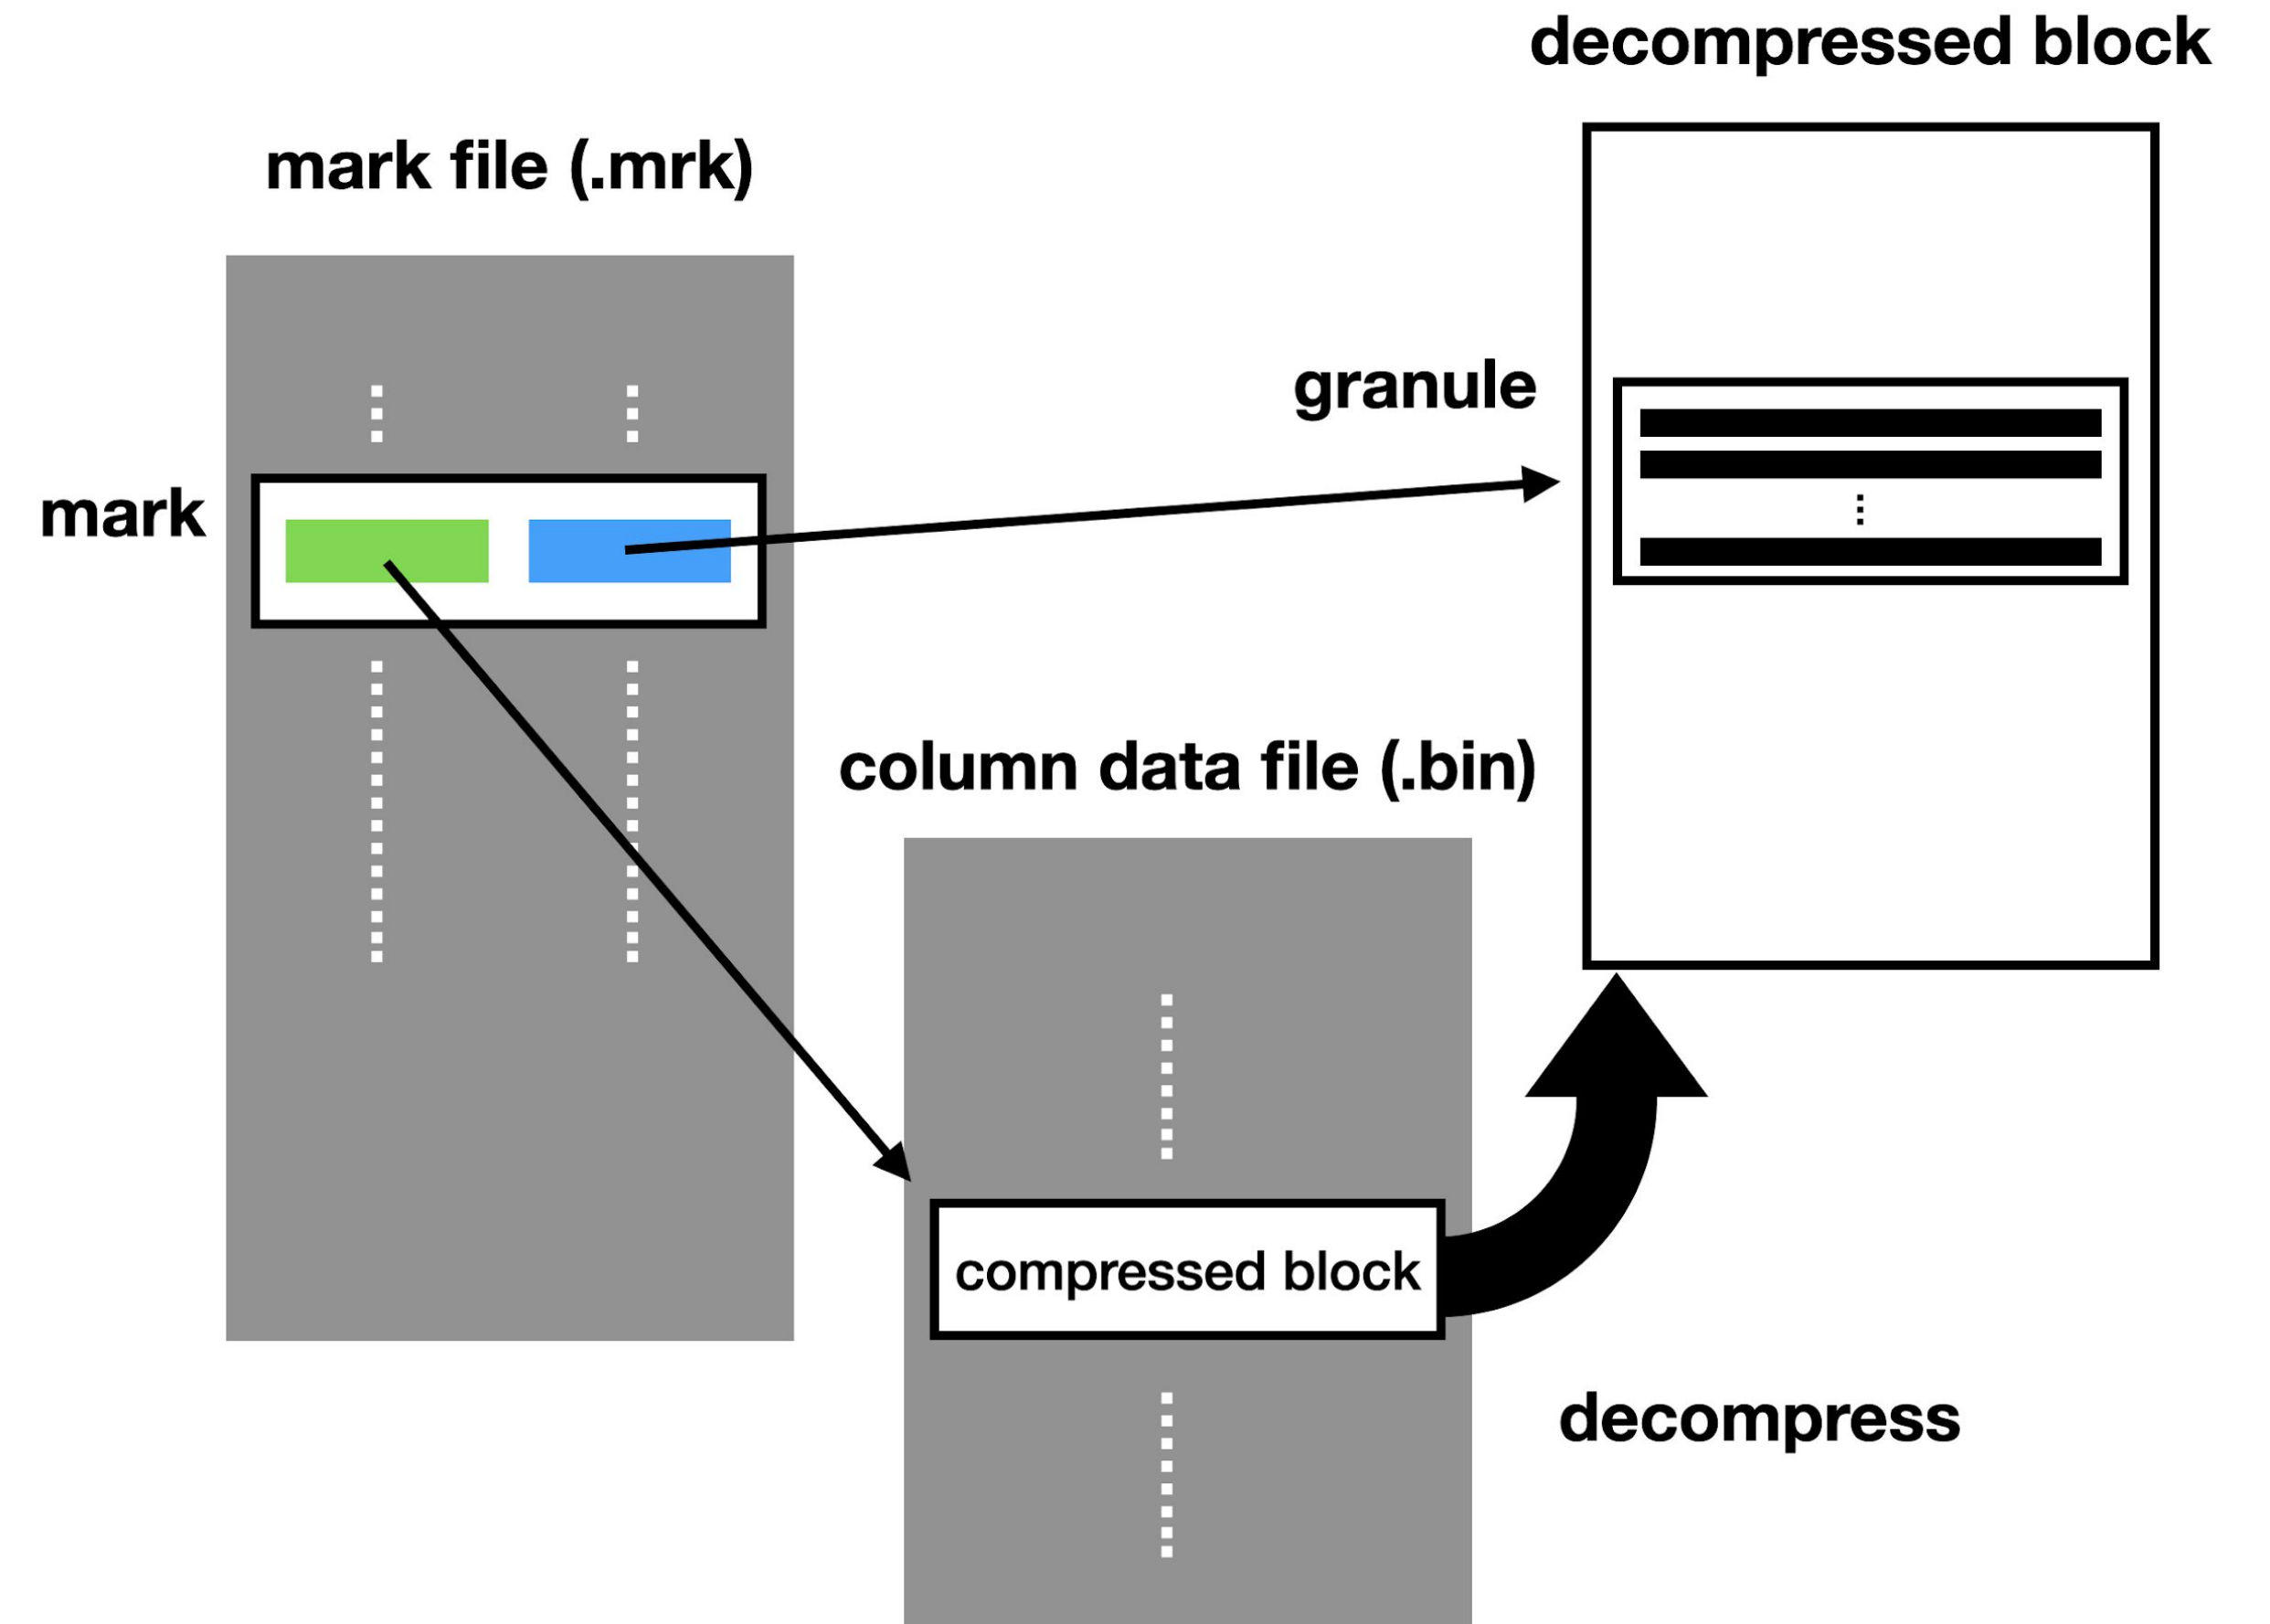


# Agenda

- About us
- Road to ClickHouse adoption
- **Upstream patch: In-memory marks compression**

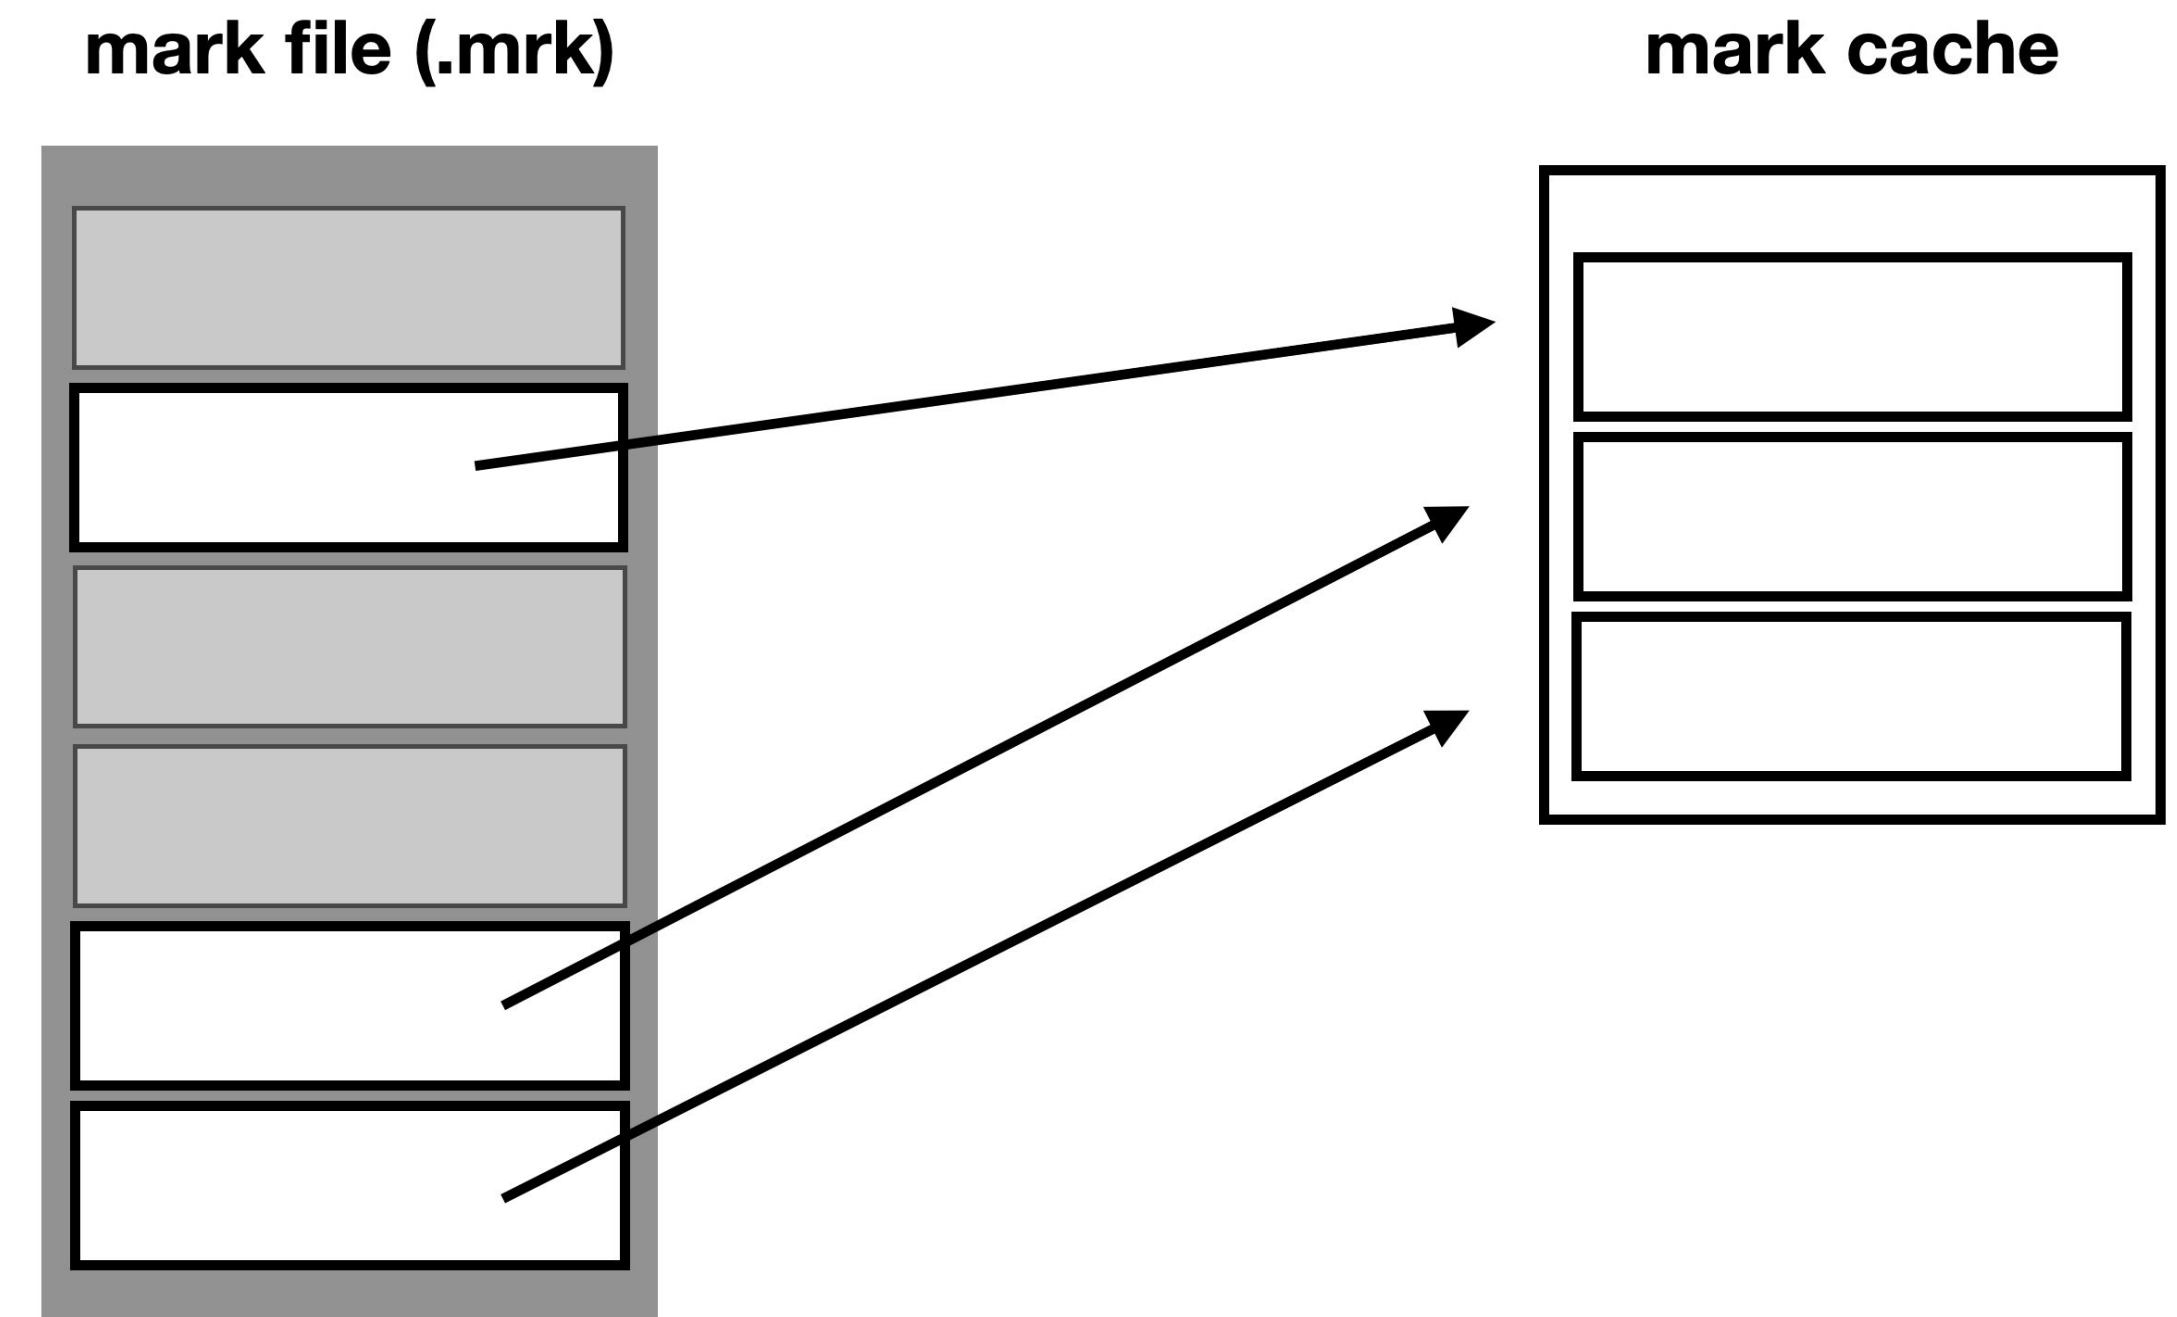
# Background

- ClickHouse uses offset pairs (“marks”) to locate groups of rows in compressed data files
- Marks are stored in “mark files” and cached in memory



# Problem

- Large queries read rows from many marks  
→ poor cache performance
- At Ahrefs, want all marks to fit in memory for max performance
- Need more memory-efficient marks representation



# Insight

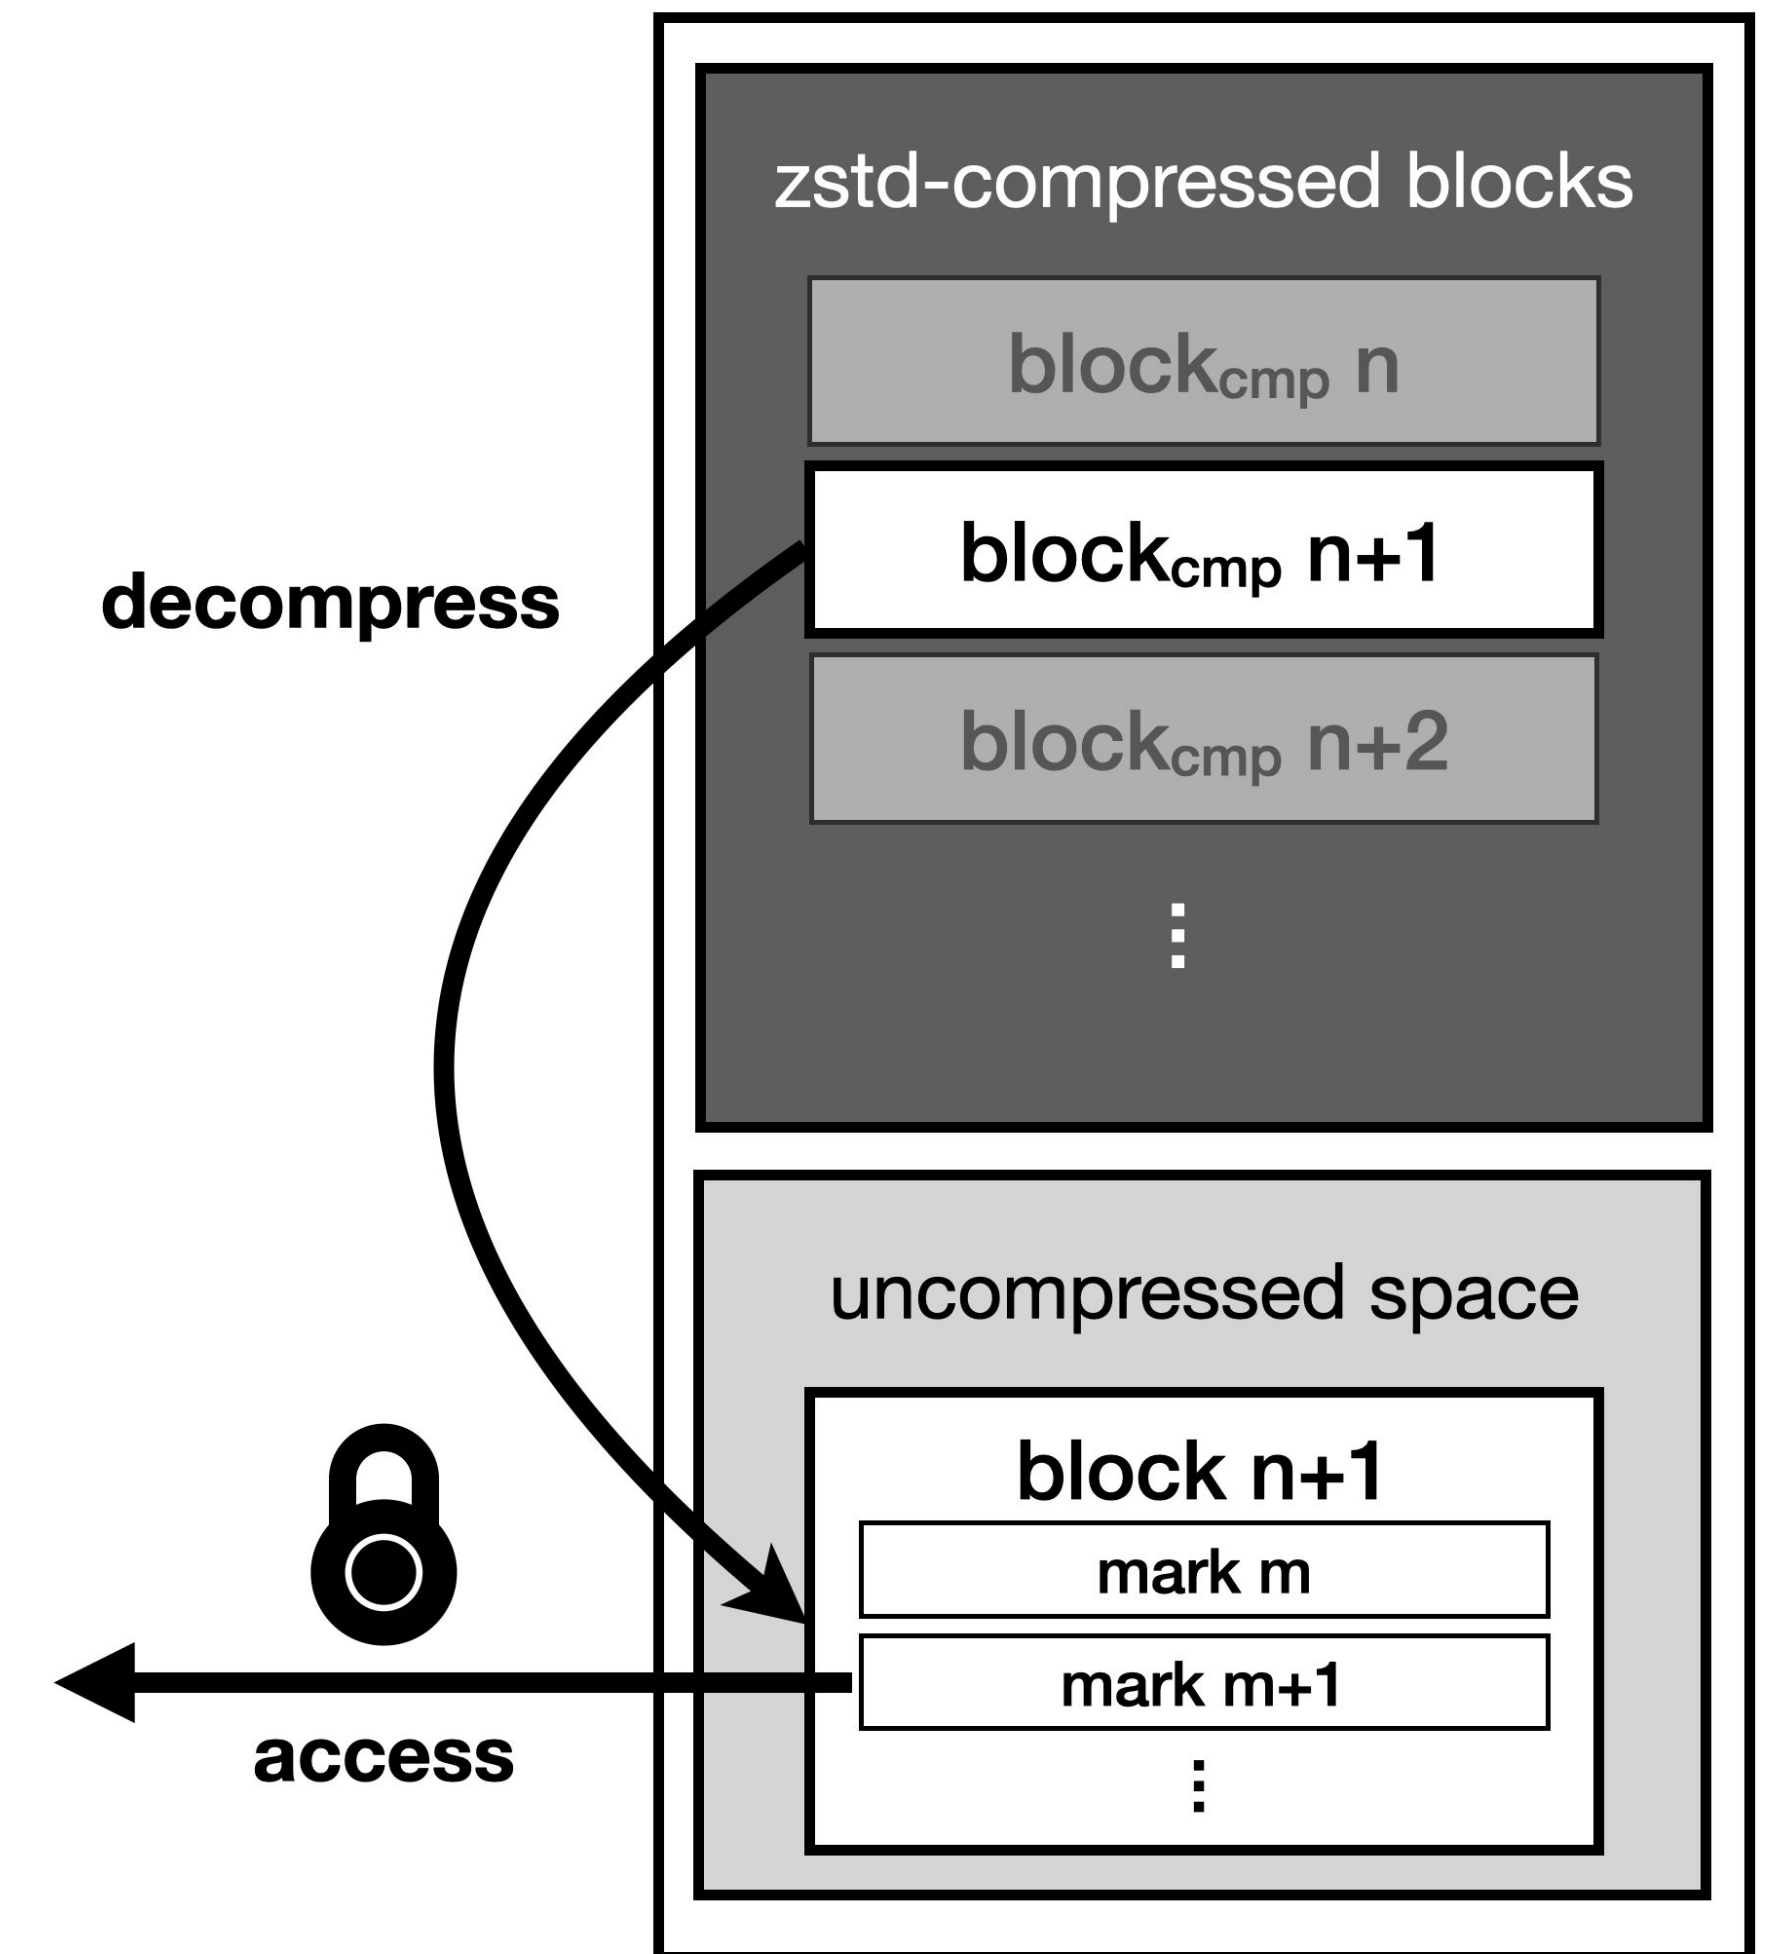
- The marks are suitable for compression
  - First offset is monotonically increasing
  - Second offset is almost always (but sometimes not) zero
  - But both stored as 64-bit integers (wasteful)





# Early implementation at Ahrefs

- Split marks into blocks, store blocks compressed with zstd in memory
- Allocate space for one uncompressed block
- Guard access with mutex
  - Downside: contention if many threads
- In practice, okay



# Early implementation at Ahrefs

Oct 2021

RFC: save memory and improve mark cache hit rate by compressing marks in memory #30434

< > Code

Closed

jorisgio wants to merge 1 commit into ClickHouse:master from jorisgio:compress\_mark\_cache

Conversation 10

Commits 1

Checks 5

Files changed 4

+111 -9

jorisgio

commented on Oct 20, 2021

Contributor

When using tables with hundreds of billions of rows and 10s of columns, storing all marks in memory can easily use several hundreds of gigabytes. Although mark cache is officially a cache, running queries with high miss rate has a sizeable performance impact. One solution would be to increase granularity, but this sacrifices performances of small queries.  
But actually, marks data representation is very wasteful. Offset in compressed block is usually 0, so it takes 8 bytes to store zeros, and offset field is monotonic. Data can

Reviewers

azat

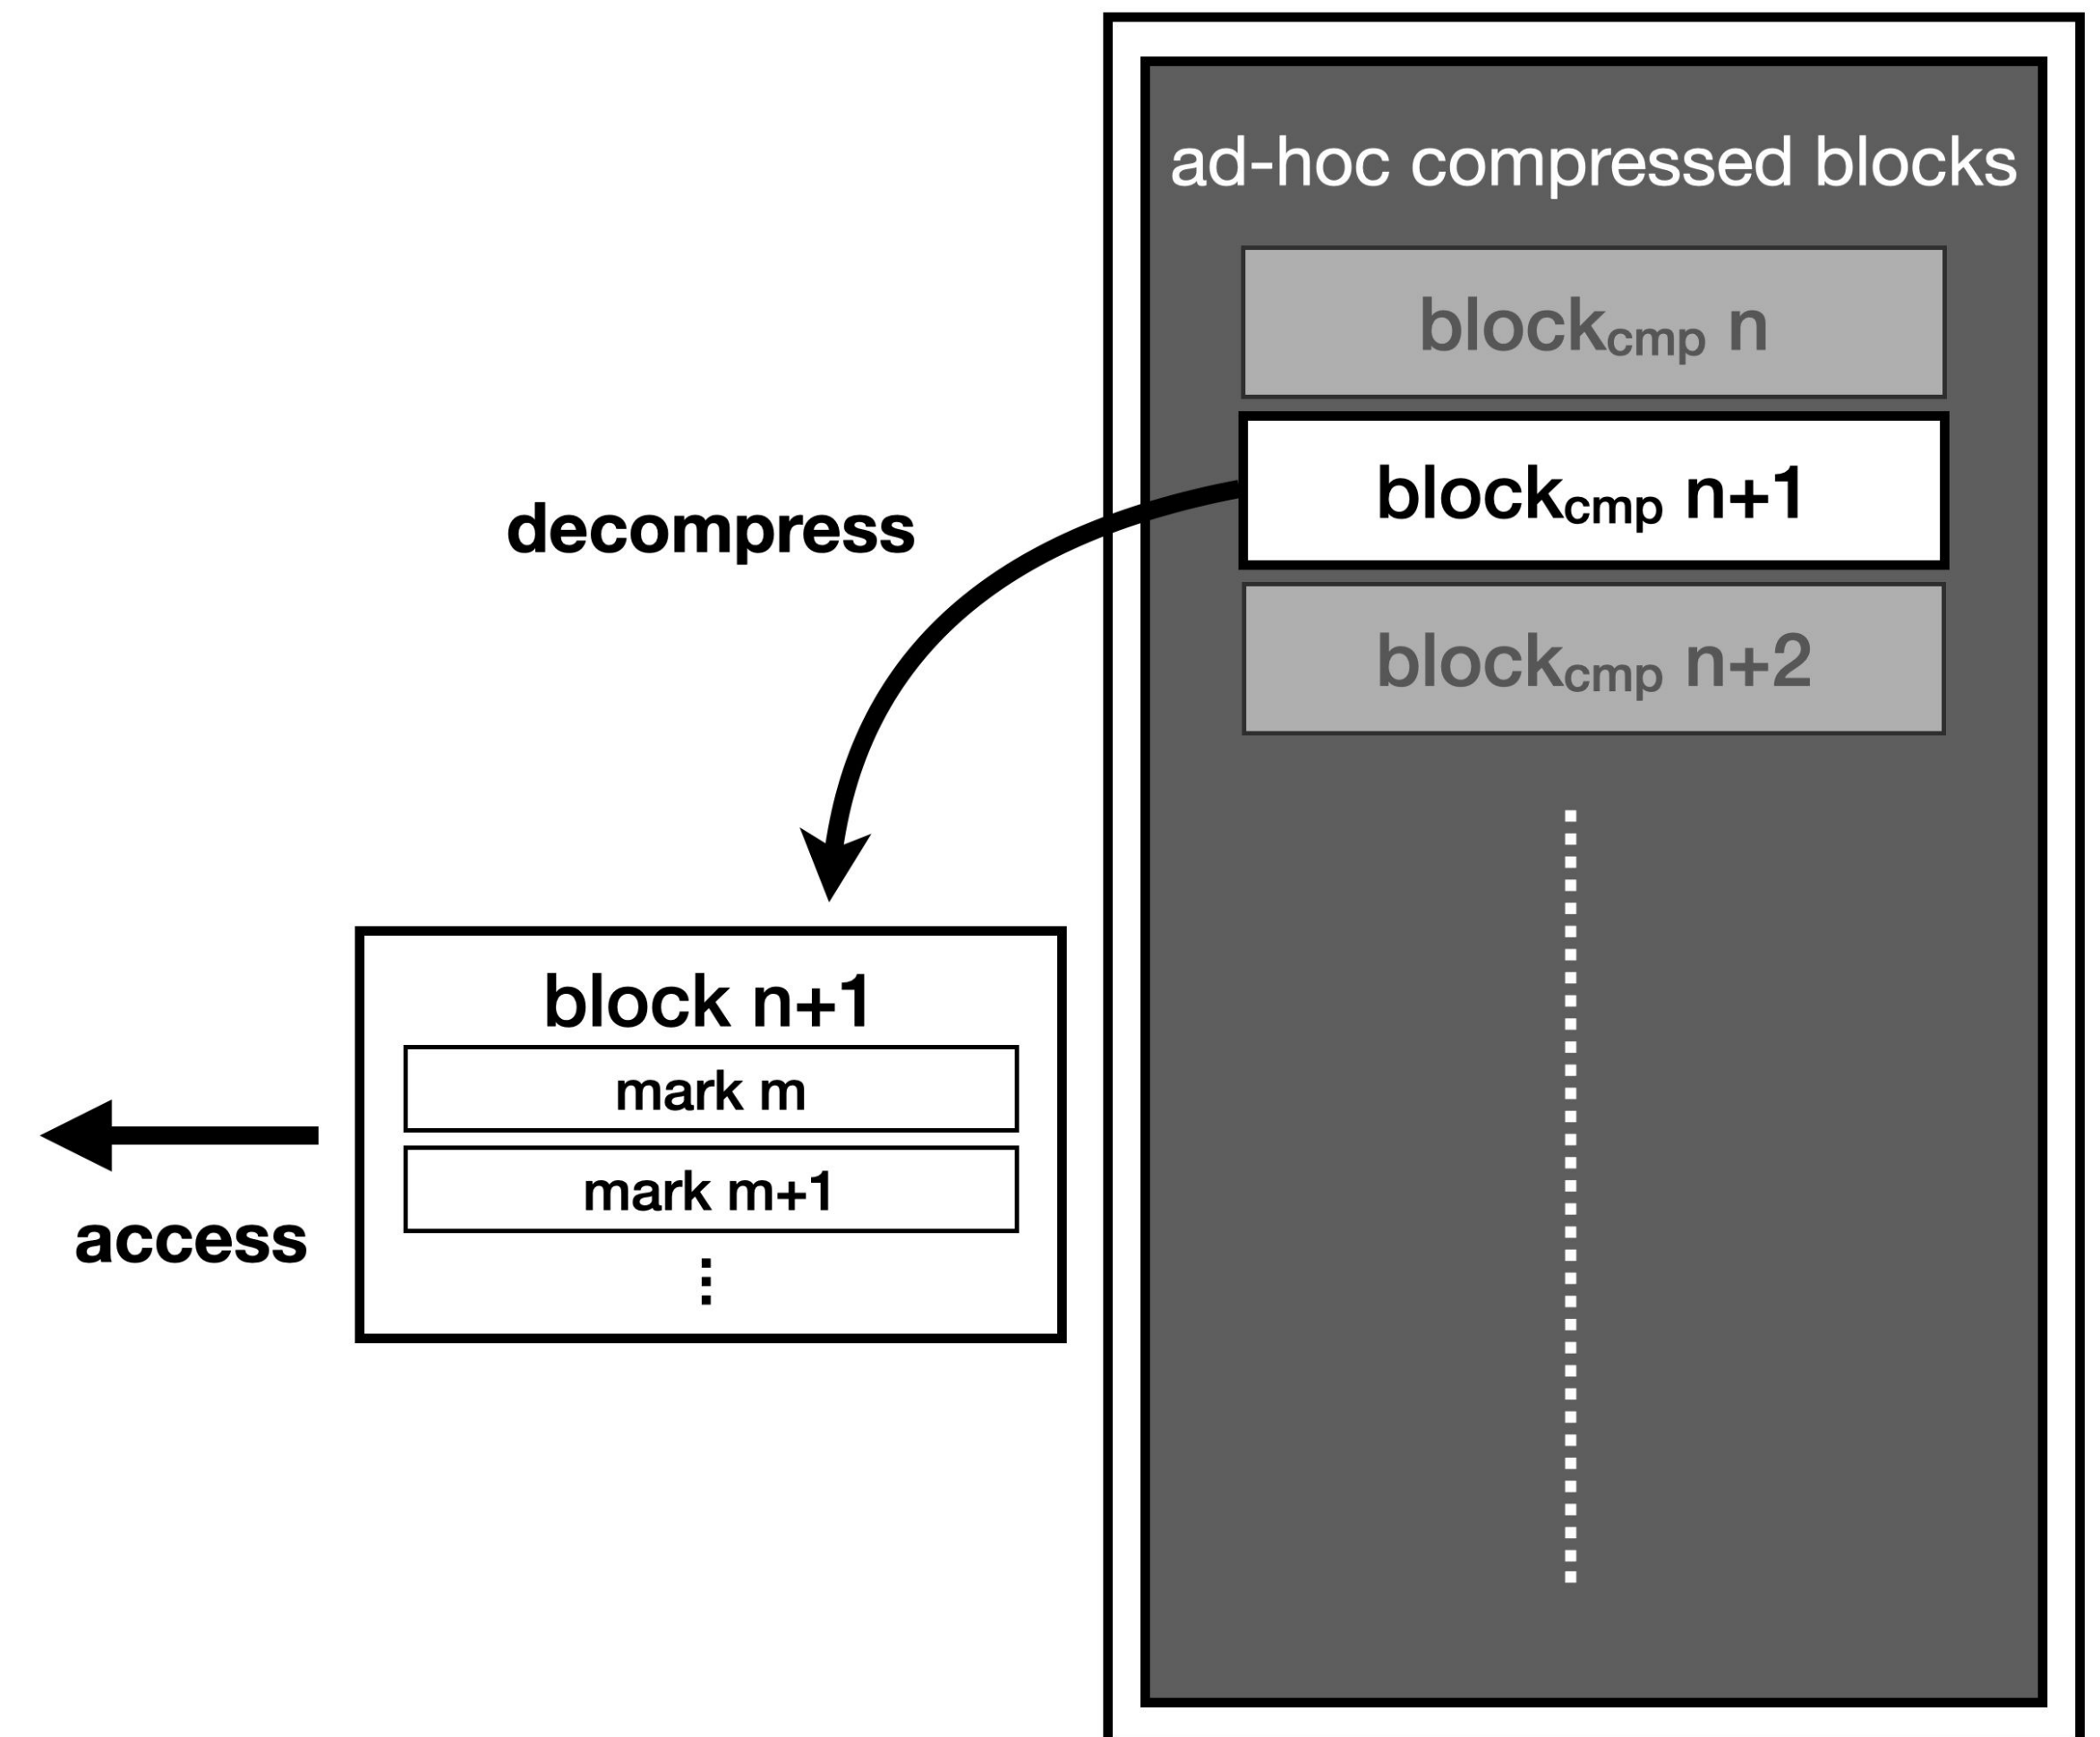
Assignees

No one assigned

Labels

# Eventual upstream solution

- Also split marks into blocks, but use ad-hoc compression
  - Store various deltas, bitpacked
- All in-memory marks compressed → no mutex
- Downside: decompression for each mark access
- Result: 3-6x less memory than without compression





# Eventual upstream solution

March 2023

Compress marks in memory #47290

<> Code ▾

Merged

al13n321 merged 1 commit into master from compress-marks-in-memory on Mar 14

Conversation 14

Commits 1

Checks 139

Files changed 7

+275 -41

al13n321 commented on Mar 7 • edited by alexey-milovidov ▾

Member

...

Changelog category (leave one):

- Performance Improvement

Changelog entry (a user-readable short description of the changes that goes to CHANGELOG.md):

Marks in memory are now compressed, using 3-6x less memory.

Reviewers

pufit

✓

Assignees

pufit

Labels

can be tested

pr-performance

SG ClickHouse Meetup · July 27, 2023

24 / 27



# New: Key-value interface

## Feature discussion: fast key-value query over http #52194

 Open canhld94 opened this issue last week · 3 comments



canhld94 commented last week

Contributor



This is more like a continuation of [#33581](#). Some have been done since [#33581](#):

- An abstract class for key-value entity `IKeyValueEntity`
- `IKeyValueEntity` methods has been implemented for `EmbeddedRocksDB` and `Dictionary`

### Assignees

No one assigned

### Labels

feature

# ClickHouse at Ahrefs

- Great success overall
- Performance that meets our usage demands
- Active feature development and bug fixes, regular monthly releases

# We're hiring!

[ahrefs.com/jobs](https://ahrefs.com/jobs)

- D/C++ developers
- OCaml developers
- Data scientists

