

ClickHouse Roadmap & Favourite Features

Tyler Hannan

July 2023

 ClickHouse

ClickHouse Core

ClickHouse Core Roadmap 2023

<https://github.com/ClickHouse/ClickHouse/issues/44767>

- **Query analysis and optimization**
 - Enable Analyzer by default
 - Correlated subqueries (with decorrelation)
 - Automatically select optimal JOIN algorithms
- **Formats and integrations**
 - Support for Apache Iceberg
 - Optimization of reading for Parquet
- **Security and access control**
 - Bcrypt or other PBKDF
 - Data masking in row-level security
 - JWT support
- **Separation of storage and compute**
 - Parallel replicas (production readiness)
 - Parallel replicas with custom sharding key
 - Optimization of high-cardinality operations
- **Data Storage**
 - Asynchronous inserts by default
 - Lightweight deletes (production readiness)
 - Lightweight updates
 - Transactions for ReplicatedMergeTree
- **Testing and hardening**
 - Green CI: <https://aretestsgreenyet.com/>
 - More complete fuzzing coverage
 - Fully-static ClickHouse builds
 - Embedded documentation



Faster Data Onboarding

Core recent features and roadmap

Asynchronous inserts

Ensure deduplication when using `async_insert`

- ✓ `async_insert` setting
- ✓ dedup with `async_insert`

Datalakes integrations

Integrate natively with popular datalakes formats

- ✓ Hudi, Deltalake, Iceberg format support
- ✓ Optimized Parquet reading

Semi-structured data

Add support for natively storing semi-structured data

- PR JSON Object type (experimental)
- PR SQL/JSON support



More flexible analytics

Core recent features and roadmap

Analyzer by default

Enabling Analyzer by default opens many possibilities in ClickHouse

- ✓ Implement Analyzer in ClickHouse
- PR Enable Analyzer by default
- PR Simplify existing code
 - Implement cost based optimizer

Extended JOIN support

Faster JOINS for more use cases; automatic JOIN algorithm selection

- ✓ New JOIN algorithms
 - Dynamic JOIN algorithm selection and JOIN rewrite (dep. Analyzer)
 - TPC-H and TPC-DS benchmarks

Indexing / Search

More functionality and seamless migrations / ecosystem integrations

- ✓ Inverted indices (experimental)
- PR More vector search functions
- PR More standard SQL support



Extended data management capabilities

Core recent features and roadmap

Lightweight operations

Delete data frequently without impacting performance

- ✓ Lightweight DELETES
(production-ready in 23.3)
- PR Lightweight UPDATES

Transactions

Bring ACID properties to use cases important to ClickHouse users

- PR BEGIN TRANSACTION,
COMMIT and ROLLBACK
statements (experimental)

Query cache

Speed up repeated queries with built-in query cache

- ✓ Query cache
(production-ready in 23.5)
- ✓ Enable in ClickHouse Cloud



ClickHouse Integrations

Ecosystem Integrations status and roadmap

	Available	Recent/Next
Data Ingestion / Orchestration	<ul style="list-style-type: none">✓ S3✓ Kafka✓ DBT✓ Airbyte✓ Vector (community)✓ PostgreSQL (Supabase webinar 5/2)✓ MySQL✓ Decodable (partner)✓ Metaplane (partner)	<ul style="list-style-type: none">✓ Kafka Connect✓ Confluent Cloud (partner)✓ RedPanda (partner)✓ OpenTelemetry (community)✓ FluentBit (community)🔥 Apache Beam (community)◆ Google Dataflow (partner)◆ Azure Event Hub (partner)✓ EMQX (partner)
Data Visualization	<ul style="list-style-type: none">✓ Superset / Preset✓ Deepnote (partner)✓ HEX (partner) (HEX webinar 5/10)✓ Mode (partner)🔥 Grafana (partner)🔥 Metabase	<ul style="list-style-type: none">✓ Tableau (community)● GCP Looker (partner)● AWS QuickSight◆ Azure Power BI (partner)
Language Clients	<ul style="list-style-type: none">🔥 Go🔥 Python🔥 Node.js✓ Java✓ C# (community)	<ul style="list-style-type: none">◆ Ruby (community)◆ Rust (community)
SQL Clients	<ul style="list-style-type: none">✓ DBeaver✓ Datagrip	

Legend

- ✓ Completed
- In development
- ◆ Coming Next / Evaluating
- 🔥 Recently improved

See 75+ integrations on

<https://clickhouse.com/docs/en/integrations>



Favourite Features

session_timezone

```
:) SELECT now()
```

```
┌──────────────────now()──┐  
| 2023-06-29 13:11:16 |  
└────────────────────────┘
```

```
:) SELECT now('America/Los_Angeles')
```

```
┌now('America/Los_Angeles')┐  
|          2023-06-29 04:11:21 |  
└────────────────────────┘
```

session_timezone

```
:) SET session_timezone = 'America/Los_Angeles';
```

```
:) SELECT now();
```

```
|-----now()-----|  
| 2023-06-29 04:11:29 |  
|-----|
```



Faster Parquet Reading

23.4

From S3 and URLs.
100 times faster.

23.6

clickhouse-local over partitioned
Parquet files improved 2 times 🔥
clickhouse-local over a single Parquet
file improved 4 times 🔥

A Deep Dive into Apache Parquet with ClickHouse - Part 1



Dale McDiarmid
Apr 18, 2023

System & Machine

clickhouse-local (single, new) (c6a.4xlarge, 500gb gp2):

clickhouse-local (single) (c6a.4xlarge, 500gb gp2):

Detailed Comparison

	clickhouse-local (single, new) (c6a.4xlarge, 500gb gp2)	clickhouse-local (single) (c6a.4xlarge, 500gb gp2)
Load time:	0	0
Data size:	13.76 GiB (x1.00)	13.76 GiB (x1.00)
Q0.	0.15s (x1.00)	0.25s (x1.61)
Q1.	0.17s (x1.13)	0.15s (x1.00)
Q2.	0.17s (x1.00)	0.40s (x2.32)
Q3.	0.18s (x1.00)	0.36s (x1.94)
Q4.	0.51s (x1.00)	2.18s (x4.26)
Q5.	0.84s (x1.00)	6.05s (x7.13)

ClickHouse

has become almost

SOUNDEX - Developed in 1918, obsolete

```
SELECT soundex('Hello, world!') AS x
```

```
┌──x──┐
```

```
| H464 |
```

```
└────┘
```

```
SELECT soundex('Hallo, wereId!') AS x
```

```
┌──x──┐
```

```
| H464 |
```

```
└────┘
```

Integration with MongoDB 5.1+

```
CREATE TABLE test_mongo ... ENGINE = MongoDB(  
    'mongo1:27017', 'test', 'simple_table', 'testuser', 'clickhouse');  
  
SELECT * FROM test_mongo;  
  
SELECT * FROM mongodb(  
    'mongo1:27017', 'test', 'simple_table', 'testuser', 'clickhouse',  
    'x String, y UInt64, ...');
```



Describing A Release

```
$ sudo docker run --rm clickhouse/clickhouse-server:23.5 clickhouse-local --query  
"SELECT * FROM system.contributors ORDER BY name" > contributors_23.5.txt
```

```
$ sudo docker run --rm clickhouse/clickhouse-server:23.6 clickhouse-local --query  
"SELECT * FROM system.contributors ORDER BY name" > contributors_23.6.txt
```



Describing A Release

numbersix at thevillage in ~/clickhouse-mac

```
$ ./clickhouse local --query "  
  SELECT arrayStringConcat(groupArray(line), ', ')  
  FROM file('contributors_23.6.txt', LineAsString)  
  WHERE line NOT IN (  
    SELECT *  
    FROM file('contributors_23.5.txt', LineAsString))  
  FORMAT TSVRaw"
```

Chang Chen, Dmitry Kardymon, Hongbin Ma, Julian Maicher, Thomas Panetti, YalalovSM,
kevinyhzou, tpanetti, 郭小龙



Thank You!