

ClickHouse (クリックハウス) Meetup Tokyo

6月5日

Hosts & Speakers



Paul Davis
ポール・デビス

Area Vice President for APAC
ClickHouse アジア太平洋地域統
括担当



Johnny Mirza
ジョニー・ミルザ

Solutions Architect
ClickHouse ソリューションアーキテ
クト



Satomi Takebe
竹部智実

Senior Manager, Technology
Partnerships
テクノロジーパートナー本部 本部長
@ AWS Japan



Harunobu Kameda
亀田治伸

Evangelist
エヴァンジェリスト
@ Cloudflare Japan



Ryosuke Minami
南 良介

Senior Software Engineer
シニア ソフトウェア エンジニア
@ Backpack Exchange



Takafumi Kataoka
片岡崇史

Site Reliability Engineering Team
AI Marketing solution統括本部テク
ノロジー戦略本部 SREチーム
@ GENIEE



当日のアジェンダ:




ClickHouse ソリューションアーキテクト,
ジョニー ミルザ「ClickHouseとは？」



アマゾン ウェブ サービス ジャパン合同会社テクノロジー
パートナー本部 本部長, 竹部智実
「AWS Marketplaceによるデジタル調達」



Cloudflare Japan株式会社 エヴァンジェリスト, 亀田治伸
「CloudflareによるClickHouseの活用について」



Backpack Exchange シニア ソフトウェア エンジニア, 南
良介「ClickHouse への数十億行におよぶデータベース移
行のオーケストレーション」



GENIEE AI Marketing Solution統括本部テクノロジー戦
略本部SREチーム, 片岡崇史「ClickHouseで実現する広
告成果レポートの高速化と安定化」



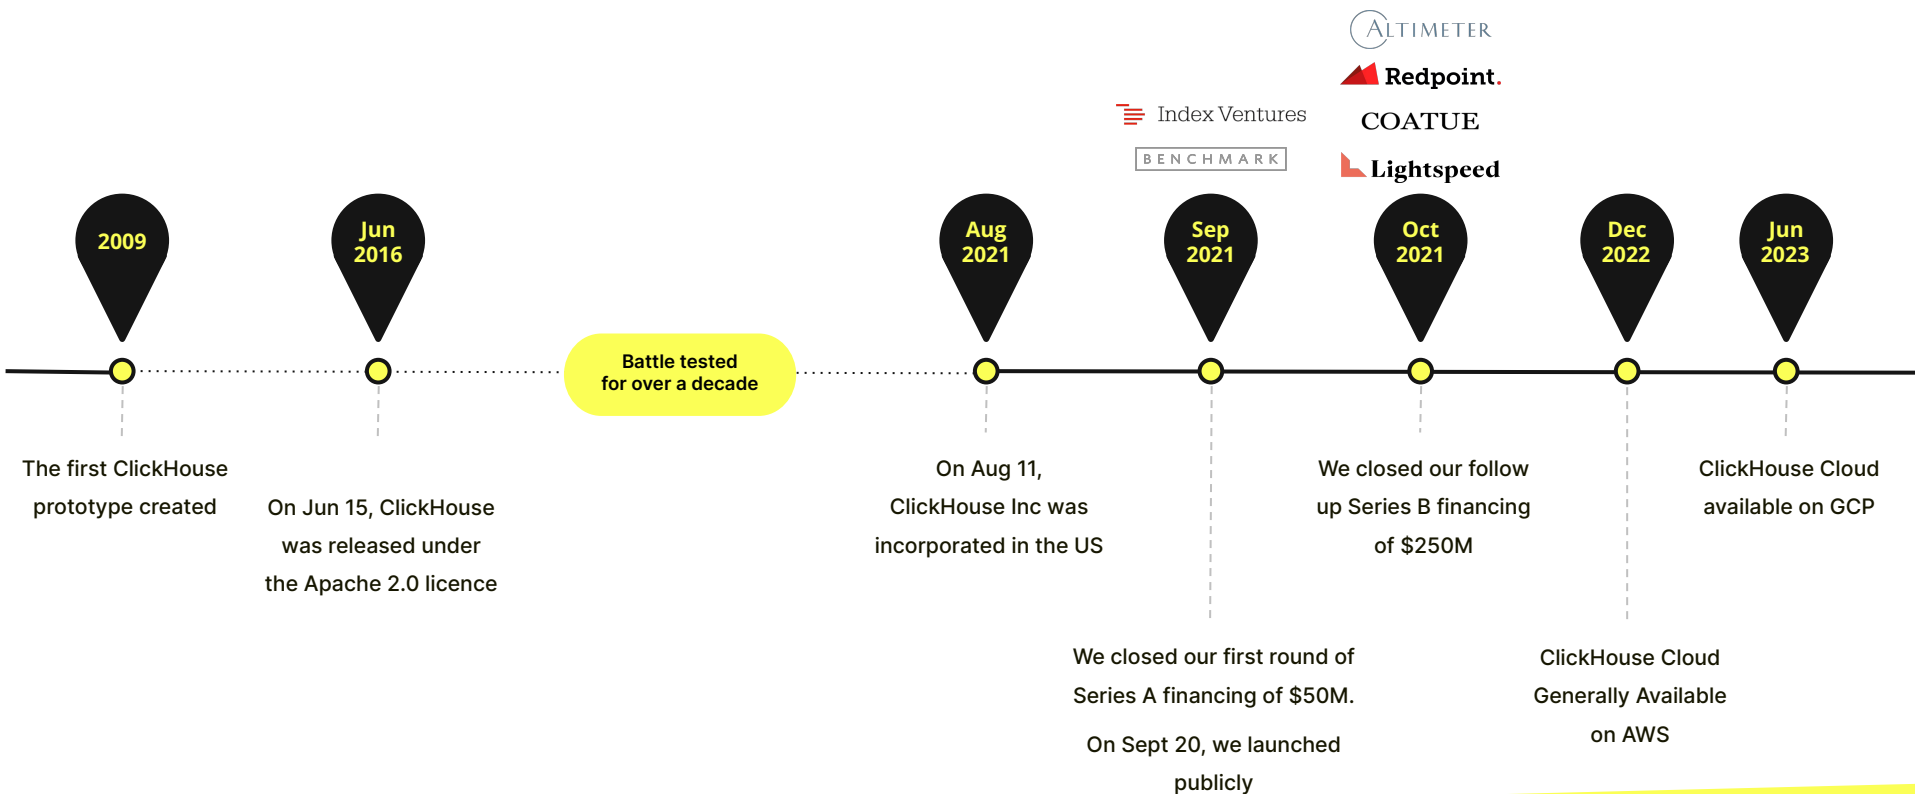
ClickHouse

A Quick Overview

2024


||||· ClickHouse

The ClickHouse Journey



ClickHouse Open Source

- ✓ 34.9k Stars
- ✓ 6.1k Forks
- ✓ 1.3k Contributors
- ✓ 100k+ Commits
- ✓ 114k active community members



ClickHouse Cloud



Bloomberg

- ✓ 300+ customers
- ✓ 150 active trials
- ✓ Processed 64+ billion queries, 28 quadrillion records, and 580 PiB of data since launch

NETFLIX

Uber



Deutsche Bank



NGINX



ROKT



What is ClickHouse?

Your (soon-to-be) favorite database!

Open source **column-oriented** **distributed** **OLAP** database

Since 2009
31,000+ GitHub stars
1300+ contributors
500+ releases

Best for aggregations
Files per column
Sorting and indexing
Background merges

Replication
Sharding
Multi-master
Cross-region


Analytics use cases
Aggregations
Visualization
Mostly immutable data

Use cases





Real-time Analytics

Power interactive applications and dashboards that analyze and aggregate large amounts of data on the fly. Run complex internal analytics in ms, not mins or hrs.



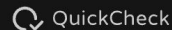

Logs, events, traces

Monitor with confidence your logs, events, and traces. Detect anomalies, fraud, network or infrastructure issues, and more.



Business intelligence

Interactively slice and dice your data for analysis, reporting, and building internal applications. Evaluate user behaviors, ad and media perf, market dynamics, and more.



ML and Gen AI

Execute fast and efficient vector search. Plug-and-play Generative AI models from any provider. Use lightning-fast aggregations to power model training at petabyte scale.



Key Features

Some of the cool things ClickHouse can do

1 Speaks SQL

Most SQL-compatible UIs, editors, applications, frameworks will just work!

2 Lots of writes

Up to several million writes per second - per server.

3 Distributed

Replicated and sharded, largest known cluster is 4000 servers.

4 Highly efficient storage

Lots of encoding and compression options - e.g. 20x from uncompressed CSV.

5 Very fast queries

Scan and process even billions of rows per second and use vectorized query execution.

6 Joins and lookups

Allows separating fact and dimension tables in a star schema.

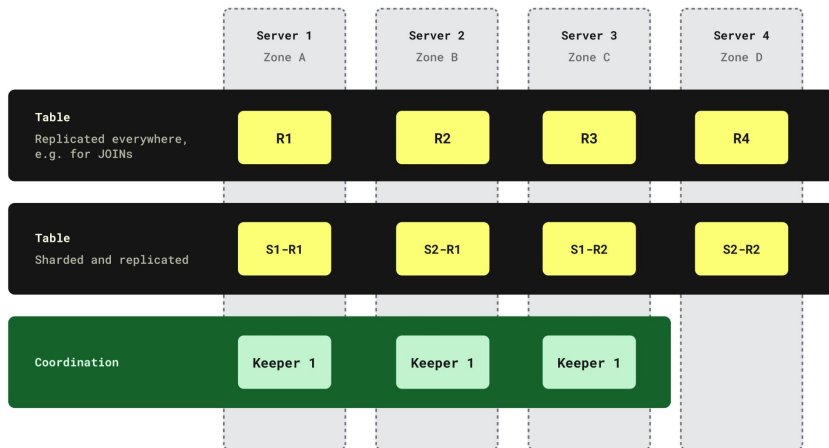


ClickHouse

Self-managed

- ✓ Open-source
- ✓ Flexible architecture
- ✓ Efficient and robust
- ✓ Support contracts available

Sample self-managed architecture



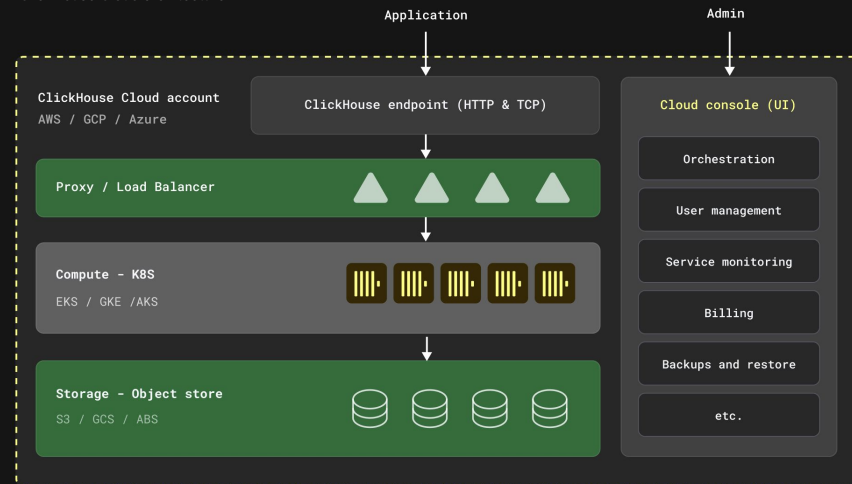
ClickHouse

Cloud



- | | |
|----------------|--|
| ✓ Easy to use | Managed for you |
| ✓ Feature-rich | Cloud-first features & tooling |
| ✓ Fast | Automatically maximizes performance/efficiency |
| ✓ Scalable | Scale seamlessly |
| ✓ Reliable | Ensure reliability |
| ✓ PAYG | SaaS usage and capacity based pricing |

ClickHouse Cloud architecture



ありがとう ございます
Thank you

Keep in touch!



clickhouse.com/slack



[#clickhouseDB @clickhouseinc](#)



[clickhouse](#)