

One Feature to Win it All

Materialized View a low key winning feature

About: Nahwin Rajan

- CTO & Co-founder In Stealth Mode
- Software Architect
 - High frequency Trading System
 - (+-) Reverse Engineered 1 million Request/minute with Go worker
 - Several Tokopedia's Hero Feature (Hourly Flash Sale, Shop Home, etc)
- Head of Engineering
- Motogp Fans

Outline

- ClickHouse Materialized View
 - Clickhouse Magical Feature
 - You can skip “TL” in “ETL” (extract, transform, load)
- ClickHouse Remote Ingestion
 - Send your system native data to clickhouse without extra pre-processing
 - Come with vast acceptable input format
- “Automatic” Schema
 - Easily create and sync your schema through code

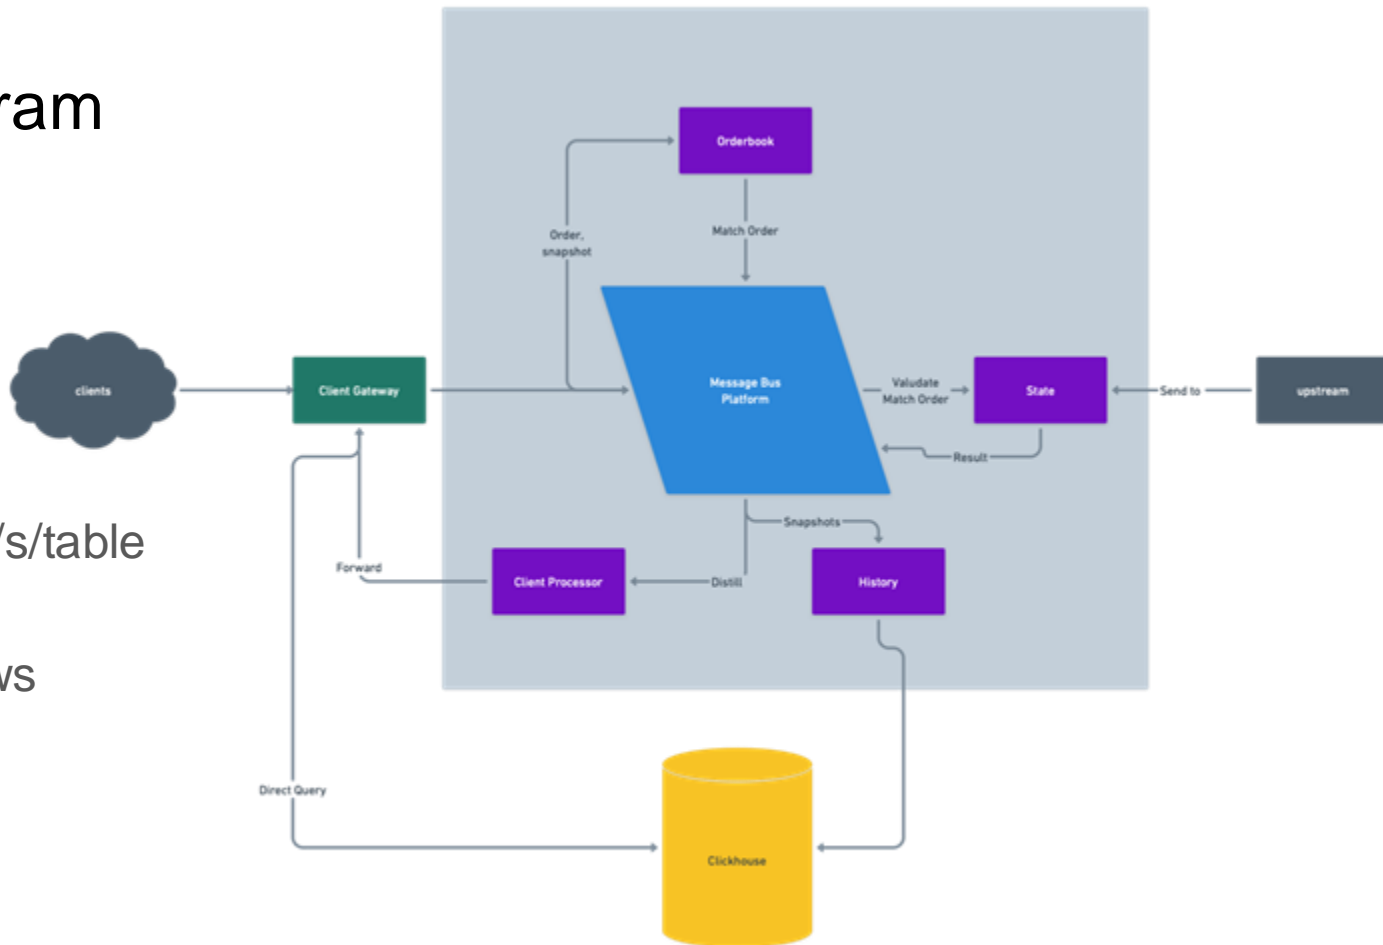
Technical Spec

- Stock Exchange Platform from Scratch
- Support High Frequency Trading
- 5ns to finish instruction once it reach our system
 - $1s/5ns \Rightarrow 200$ million instruction per se second
- Everything on Memory
- 24/7
 - No Downtime
 - Stock Exchange does have Office Hours
 - Order for the day is written to Persistent Storage during Off time

System Diagram

System Diagram

- +- 200K rows/s/table
- 18 tables
- 3,6 Million rows



Why time series

- The right tools for the job
- Postgresql Ingestion Max Capacity +- 100K Rows
- Requirement
 - 200 Million Rows per Second
 - Near Real Time: Read (high Read and Write)
 - Processing Data for User View
 - 24/7

Why ClickHouse

Available Options:

InfluxDB, QuestDB, TimescaleDB, Clickhouse, Druid, Prometheus

Niche winning factor each db:

- Druid & Prometheus: yeah, lets stick them for monitoring purpose
- **InfluxDB**: is highly specialized for time-series data but lacks in relational or complex query use cases
- **TimescaleDB**: benefits from SQL compatibility and PostgreSQL's robust ecosystem

Why ClickHouse (2)

- **Clickhouse:** excels in real-time, high-performance analytics
- **QuestDB:** extreme fast ingestion

Basically Everyone Claim they are #1

- Uber's case with mysql and postgre

The differences are minutes

Clickhouse wasn't leading the selection 🙌

- When engineer fully focused working on work

INSIDER HELP!

The (original) Task

- Work on snapshot of the order
 - Create task to dump into disk
 - Create service to read the data
 - Transfer the data to database
- Creating schema for each service output
 - If payload structure changed, schema must be changed
 - Rippled to user view schema's
- Creating useful schema for end-user (how trader will see their data)

Quotes

I will always choose a lazy person to do a difficult job because a lazy person will find an easy way to do it ~ Bill Gate (unsolicited)



Materialized View to the RESCUE

- Each TimeseriesDB does some sort of aggregation
- Clickhouse Materialized View is Low Key SUPER FEATURE
 - CLICKHOUSE GOT IT RIGHT WITH:
 - “AUTOMATIC INCREMENTAL DATA TRANSFORMATION”
- Pattern 3 table/views
 - **Source -> Materialized View -> Target Table**
 - Automate the Transformation of your Data using Materialized View



MergeTree

```
CREATE TABLE uk_price_paid
(
  date Date,
  town String,
  price UInt32
)
ENGINE = MergeTree
ORDER BY (town, date)
```

date	town	price
2023-01-03	London	172000
...

date	town	price
2023-01-04	London	180000
...

date	town	price
2023-01-03	London	172000
2023-01-04	London	180000
...

Replacing MergeTree

```
CREATE TABLE uk_listings
(
  id UInt32,
  date Date,
  town String,
  price UInt32
)
ENGINE = ReplacingMergeTree
ORDER BY id
```

id	date	town	price
23	2023-01-03	London	172000
...

id	date	town	price
23	2023-01-04	London	180000
...

id	date	town	price
23	2023-01-04	London	180000
...

Aggregating MergeTree

```
CREATE TABLE uk_price_paid_aggregates
(
  town String,
  max_price
    SimpleAggregateFunction
      (max, UInt32),
  avg_price
    AggregateFunction
      (avg, UInt32)
)
ENGINE = AggregatingMergeTree
ORDER BY town
```

town	max_price	avg_price
London	1900000	783000 13
...

town	max_price	avg_price
London	3000000	1576000 25
...

town	max_price	avg_price
London	3000000	2359000 38
...

REMOTE INGESTION !!

- Turn out Clickhouse have REMOTE INGESTION Feature!
- And Clickhouse can work with multiple Format!!
 - Saving time to manually adjust the data into certain format
 - We use Binary format to directly inject to Clickhouse




remoteSecure
table function

clickhouse-local

integration
table function



Your current
database system



A ClickHouse cluster



clickhouse-server

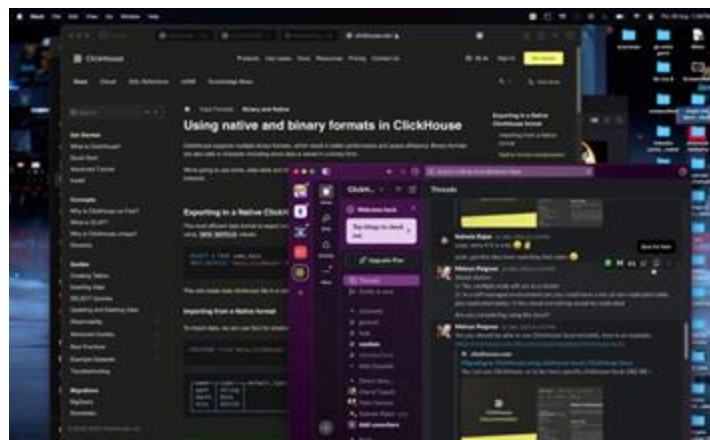
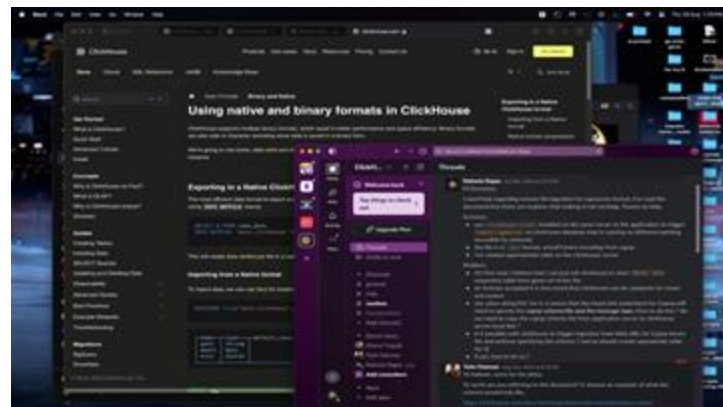
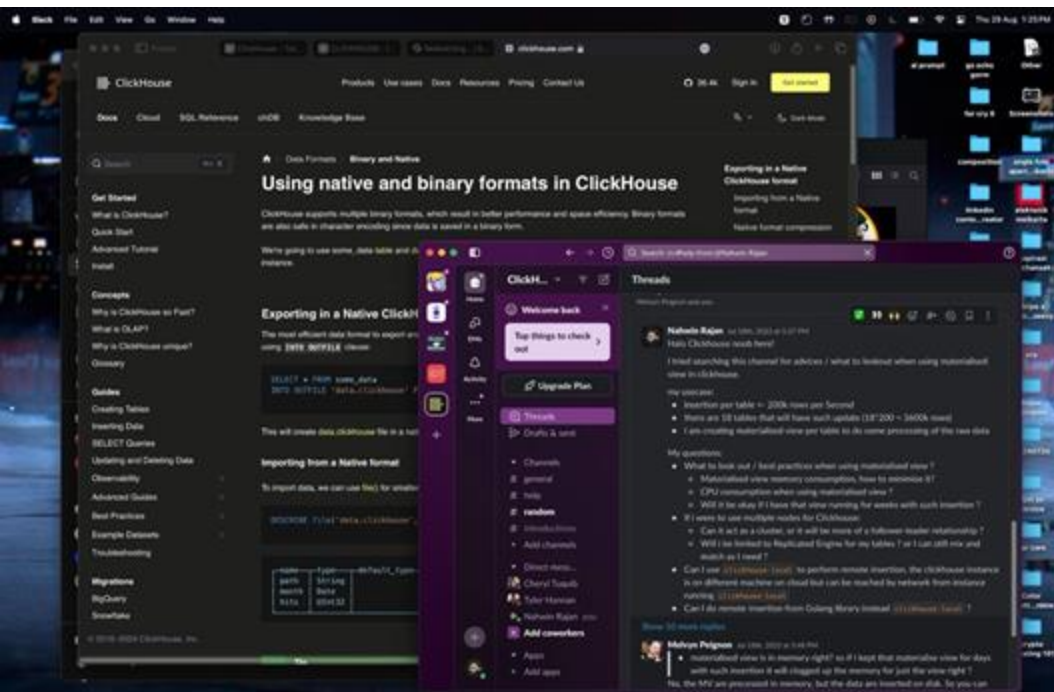


clickhouse-local

The (original) Task

- Noticed the thing with Materialized View
 - Bye bye manual data processing
- Noticed the Remote ingestion & multiple Input Format
 - Bye bye labour intensive data loading
- Can be mix with flexible schema
 - Schema can be included with the data dump
 - Schema can be created from code (golang)
- Using the Self Hosting (open source option)
- Helpful community and tech affiliate in Slack Channel

Helpful Document and Community



Possible Use case for Indonesia Tech Scene

- “Event Driven System” crave
 - Persist it to ClickHouse instead of Kafka
 - OR you no longer need all those intermediary service / process to persist your data
 - Save lot of work and pain when Persisted to Clickhouse with 3 tables Pattern
- Trading Fintech can also benefits
 - Play nieces with existing trading tech stacks and format
 - Can use managed service for ease of mind