



ClickHouse in Action:
Powering **Electrum**'s
Electric Mobility
Data Infrastructure

Andi Pangeran



- Over **10 years of experience** in the IT Industry
- Currently, **Principal Software Engineer** at **Electrum.id**



@apangeran



@a_pangeran



@A_Pangeran



We aim to create a sustainable future where electric mobility is **accessible, efficient, and integral** to urban life. We offer electric motorcycles & mobility solutions with **convenient charging battery swap technology.**

3,000+

electric
motorcycles

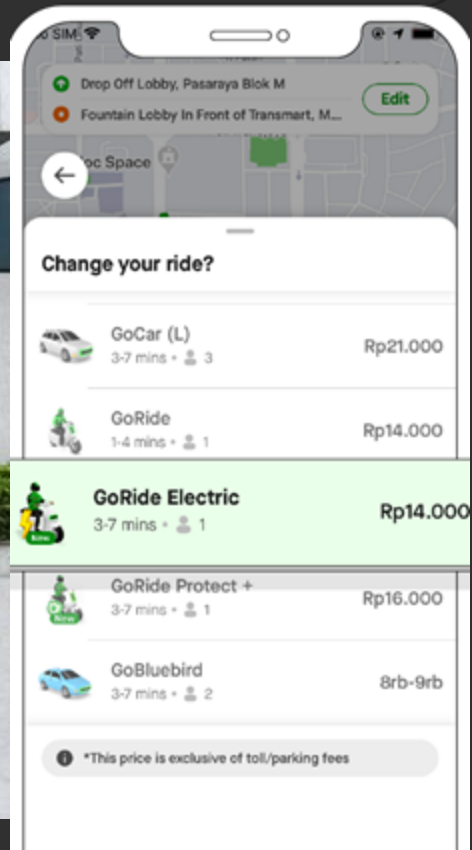
250+

swap stations
across Jakarta

220k

km travelled
per day





electrum 

Paket Beli Baterai

Motor + Baterai + Charger

Rp 23,9 Jt

OTR Jabodetabek



Battery



Portable Charger

DP Mulai 0%

Cicilan Mulai 600 ribuan
Tenor hingga 60 bulan

Garansi Baterai
5 tahun/
50ribu KM

1+1 | 20.000 km
Tahun | Garansi Part*

2+1 | 30.000 km
Tahun | Garansi Dynalene*

1 | Gratis
Tahun | As Service
1x Part Service

1 | Gratis
Tahun | Home Services

Gratis
Helm Electrum
& Toolkit

1 | Emergency
Tahun | Road Assistance

*Syarat dan ketentuan berlaku

Pembelian Terbatas

electrum 

Paket Sewa Baterai

Motor + Sewa Baterai + Swap

Rp 19,9 Jt

OTR Jabodetabek



Swap Baterai
di BSS Terdekat



Sewa Baterai Rp 230ribu /bulan
GRATIS Selama 18 Bulan

DP Mulai 0%

Cicilan Mulai 500 ribuan
Tenor hingga 60 bulan

1 | 10.000 km
Tahun | Garansi Part*

2 | 30.000 km
Tahun | Garansi Dynalene*

1 | Gratis
Tahun | Home Services


Additional
Portable Charger
Rp 1,2 jt

Gratis
Helm Electrum
& Toolkit

*Syarat dan ketentuan berlaku

Pembelian Terbatas

titipan
kantor





Data at **Electrum**

Data: *Internet of Things* (IoT)



Battery

Battery

Management System:

Voltage, current, temperature, state of charge (SOC), state of health (SOH), Cycle Count, Cells Information

GPS:

GPS coordinates

Vehicle:

Speed and distance traveled



Swap Station


General Data:

State, Input Power, voltage, current

Doors Data:

State (lock/unlock), status (charging/idle), All batteries data

Data: *Transactions* (Order)




- **Top-Up Transactions:**
Data on account recharges, including amounts, payment methods, and status.
- **Swap Transactions:**
Details on battery swaps (station ID, door-number, battery ID, completion status).
- **Rental Transactions:**
 - Onboarding: Start data (user, vehicle, battery).
 - Offboarding: End data (user, vehicle, vehicle condition).
- **Operational Activity:**
Maintenance, logistics, and fleet management data (e.g., technician actions, station repairs).

So what do we look for?

Flexible, Scalable, & Low-Cost
Data Stack

Need: *Quick & Flexible* Query



Question: Why do one of our battery swap stations was charging slower than others?

Answer: Query showed that the station's temperature is 4°C hotter than others, which caused slower charging to avoid overheating. Later, we learned the station was positioned under direct sunlight.

Need: Scalable & Low-Cost

Resource Efficiency

Efficiently use resources so we're not **over-provisioning**

Expandable Infrastructure

Able to **scale linearly** with resources; as datasets grow, adding more resources will keep response times fast.



Why ClickHouse?



Why: *It's columnar based*

completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price

Row-based

Excess disk Reads We need data in column, but it's stored in rows. So entire rows are read.

completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price

Column-Based

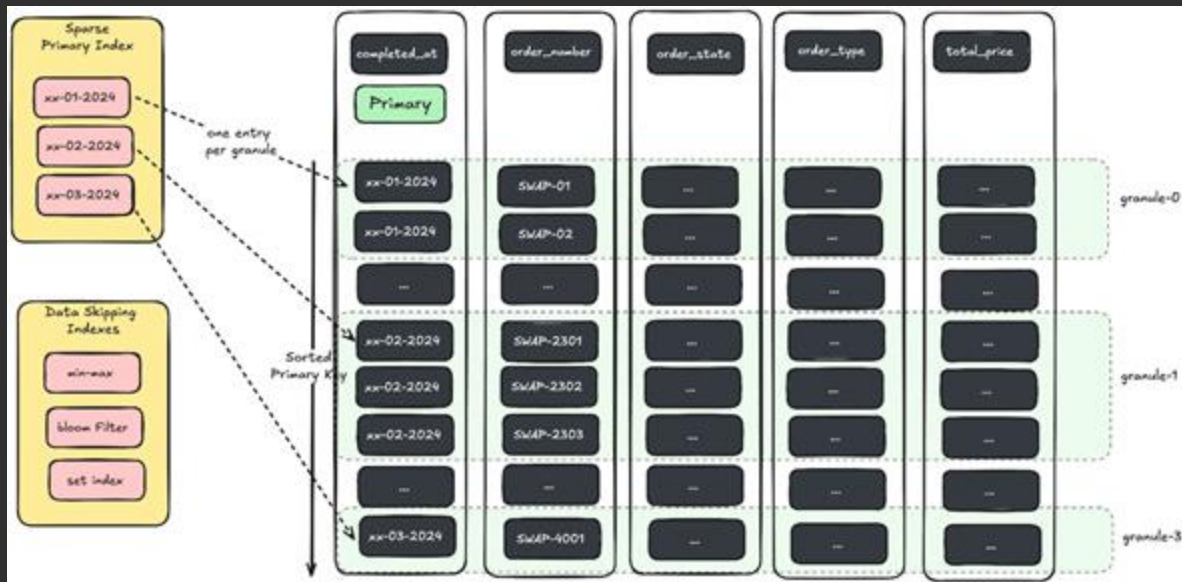
Column Pruning it only pick the subset of columns that are openly declared in the SQL query, thus eliminating 99% columns (990 of 1000 columns) from the disk read consideration.

Vectorized Execution

data processed in **large arrays**, which could be the entire length of the column loaded into RAM.

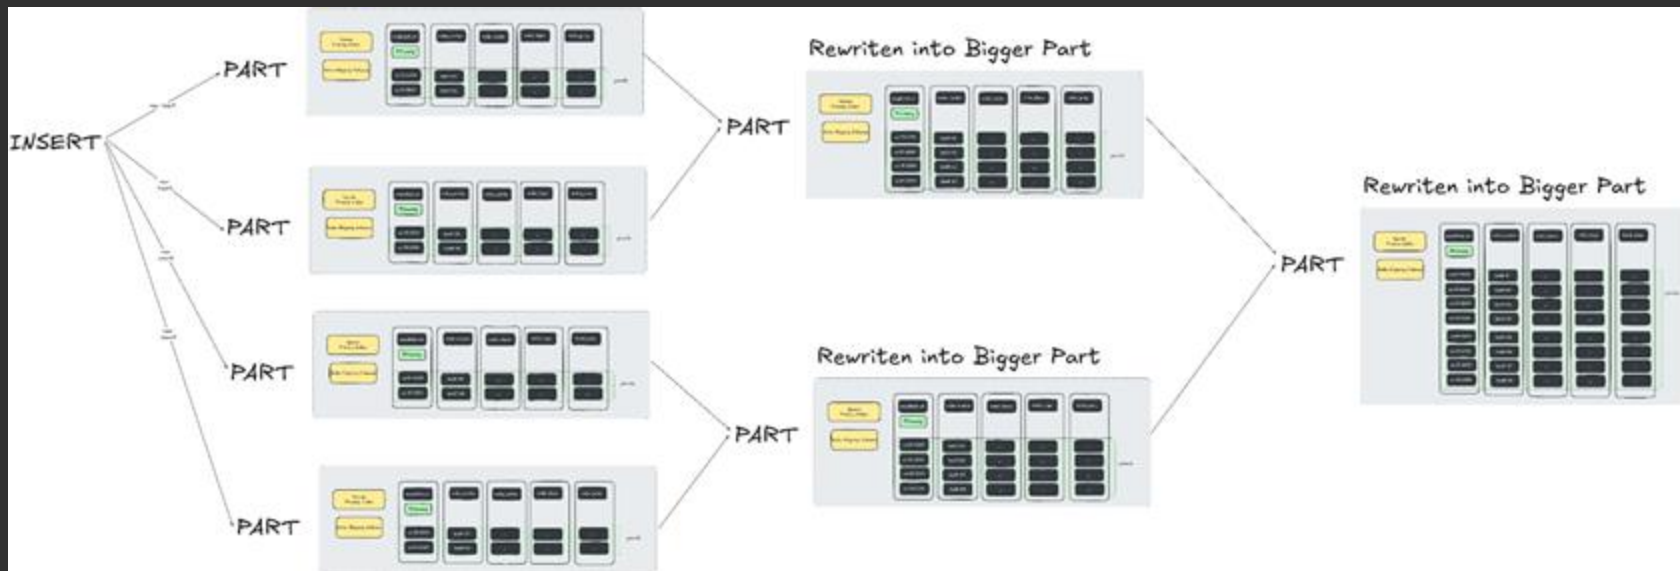
+ **reduces CPU cache miss rates**

Why: Multiple Layers of Indexes



- **Granules** are chunks of rows, and **ClickHouse** groups rows into granules based on the `index_granularity` setting.
- **Sparse Primary Index:** **ClickHouse** uses the *sparse primary index* to quickly jump to the relevant granules, skipping over large portions of the data that don't need to be read.
- **Skip Indexes:** These allow **ClickHouse** to skip entire granules if it can determine from the index that no rows in the granule satisfy the query's filter.

Why: *Efficient Merging*



Parts: Each time data is inserted, a new part is created. Multiple parts can accumulate over time, which can slow down query performance.

To improve performance and reduce the number of parts, **ClickHouse** periodically runs merging operations.

What Happens During a Merge?



```
CREATE TABLE uk_price_paid
(
    date Date,
    town String,
    price UInt32
)
ENGINE = MergeTree
ORDER BY (town, date);
```

MergeTree:

- Combines parts without deduplication or transformations.
- Simply reduces the number of parts but keeps all rows (including duplicates).
- Ideal for append-only data.

What Happens During a Merge?



```
CREATE TABLE uk_listings
(
  id UInt32,
  date Date,
  town String,
  price UInt32
)
ENGINE = ReplacingMergeTree(date)
ORDER BY id;
```

ReplacingMergeTree:

- Combines parts and deduplicates based on a key (or version).
- Replaces older rows with newer ones during a merge.
- Suitable for upserts and scenarios where data is frequently updated.

What Happens During a Merge?



```
CREATE TABLE uk_price_paid_aggregates
(
  town String,
  max_price SimpleAggregateFunction(max, UInt32),
  avg_price AggregateFunction(avg, UInt32)
)
ENGINE = AggregatingMergeTree
ORDER BY town;
```

AggregatingMergeTree:

- Combines parts by aggregating rows with the same key using predefined functions like sum, count, avg, etc.
- Ideal for pre-aggregating data to reduce query time on large datasets.

Another thing we love about ClickHouse is
materialized view.

Let's deep-dive

Materialized view 101

```
CREATE TABLE IF NOT EXISTS trx_orders
(
  completed_at DateTime64(3, 'UTC'),
  order_number String,
  order_type LowCardinality(String),
  total_price Int64
)
ENGINE = MergeTree
PARTITION BY toYYYYMM(completed_at)
ORDER BY (completed_at, order_type, order_number)
SETTINGS index_granularity = 8192;
```

Raw-Table

```
CREATE TABLE IF NOT EXISTS trx_agg_daily_orders
(
  order_date String,
  order_type LowCardinality(String),
  count_orders AggregateFunction(count),
  total_price AggregateFunction(sum, Int64)
)
ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(order_date)
ORDER BY (order_date, order_type);
```

Aggregated-Table

SYNC ?

Materialized view 101

```
CREATE TABLE IF NOT EXISTS trx_orders
(
  completed_at DateTime64(3, 'UTC'),
  order_number String,
  order_type LowCardinality(String),
  total_price Int64
)
ENGINE = MergeTree
  PARTITION BY toYYYYMM(completed_at)
  ORDER BY (completed_at, order_type, order_number)
  SETTINGS index_granularity = 8192;
```

```
CREATE TABLE IF NOT EXISTS trx_agg_daily_orders
(
  order_date String,
  order_type LowCardinality(String),
  count_orders AggregateFunction(count),
  total_price AggregateFunction(sum, Int64)
)
ENGINE = AggregatingMergeTree()
  PARTITION BY toYYYYMM(order_date)
  ORDER BY (order_date, order_type);
```

Raw-Table —→ **Materialized View** —→ **Aggregated-Table**

```
CREATE MATERIALIZED VIEW trx_agg_daily_orders_mv
TO trx_agg_daily_orders
AS

SELECT toDate(completed_at) AS order_date,
       order_type,
       countState()          AS total_orders,
       sumState(total_price) AS total_price
FROM trx_orders
GROUP BY order_date, order_type;
```

Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256



Raw-Table



Materialized View



Aggregated-Table

Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256



Raw-Table



Materialized View



Aggregated-Table

Insert new
part on disk



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Raw-Table

Materialized View

Aggregated-Table

Insert new
part on disk

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

GROUP BY
date, order-type

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD963128	SWAP	59
2024-09-28 13:48:26.249282	ORD967709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Raw-Table

Insert new
part on disk

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD963128	SWAP	59
2024-09-28 13:48:26.249282	ORD967709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Materialized View

GROUP BY
date, order-type

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Aggregated-Table

STORE as new part
of target table

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Materialized view 101 - insert merge scenario

Raw-Table → **Materialized View** → **Aggregated-Table**

Existing part
on disk



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Existing part
of target table



	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624



Raw-Table



Materialized View



Aggregated-Table

Merge with
existing part



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624



	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

Raw-Table

Merge with
existing part

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

Materialized View

GROUP BY
date, order-type

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	2	4016
2024-09-28	TOPUP	2	4958

Aggregated-Table

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

Raw-Table

Merge with
existing part

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

Materialized View

GROUP BY
date, order-type

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	2	4016
2024-09-28	TOPUP	2	4958

Aggregated-Table

Merge with
existing part
Of target table

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	4	6355
2024-09-28	TOPUP	4	9669

Materialized view 101 - Query Comparison

Raw-Table → Materialized View → Aggregated-Table



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

```
SELECT toDate(completed_at) AS order_date,  
       order_type,  
       count(*)              AS total_orders,  
       sum(total_price)      AS total_price  
FROM   trx_orders  
GROUP BY order_date, order_type;
```

Processed 8 rows



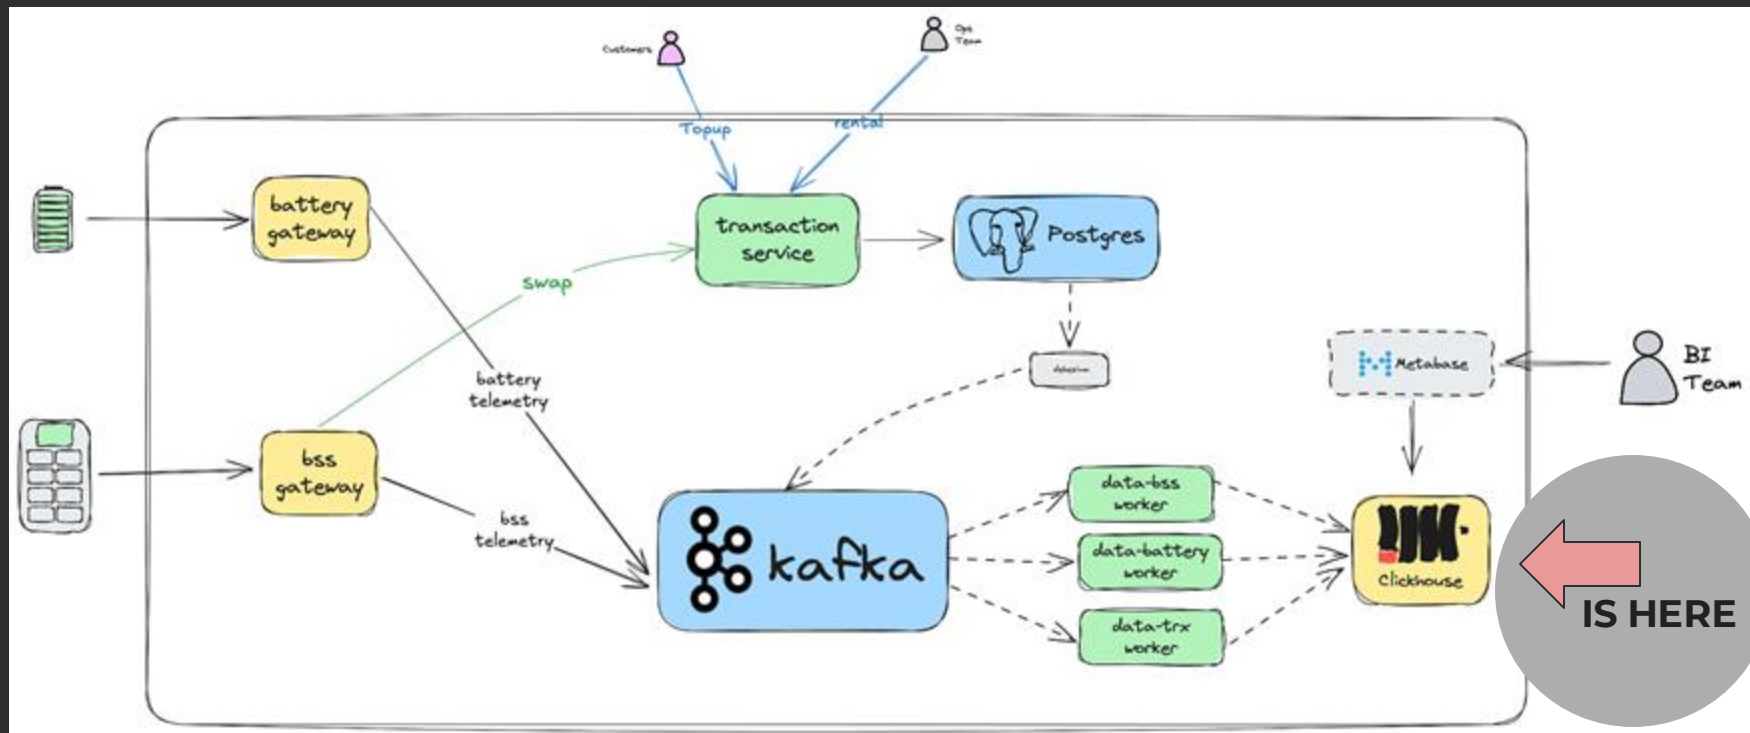
completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	4	6355
2024-09-28	TOPUP	4	9669

```
SELECT  
       order_date,  
       order_type,  
       countState() AS total_orders,  
       sumState(total_price) AS total_price  
FROM   trx_agg_daily_orders  
GROUP BY order_date, order_type;
```

Processed 2 rows

How ClickHouse
Powers Our Data Platform

How ClickHouse Powers Our Data Platform



DEMO

<https://github.com/meong1234/clickhouse-meetup-10-2024>

Where are we?

50 million+

Records / Day

≤\$500/month

*"**ClickHouse** enables us to have a data analytics capability that grows with our needs, while keeping costs low and predictable" - Electrum.id*

Key Takeaways

Columnar Format:

ClickHouse stores data in **columns**, optimizing read performance for analytical queries by reading only the relevant columns, reducing I/O and improving query speed.

Sparse Primary Index:

ClickHouse uses a **sparse index** to store pointers to data blocks (granules), allowing it to **skip irrelevant data** during queries, improving performance when filtering by the primary key.

Skip Index:

The **skip index** allows **ClickHouse** to **bypass granules** based on non-primary key column values, helping to avoid unnecessary data scanning and speeding up queries on non-primary key columns.

Merging Process in Different Engines:

- **MergeTree:** Performs basic merging of data parts without deduplication or aggregation.
- **ReplacingMergeTree:** Deduplicates data during merging, keeping only the most recent version of rows with the same primary key.
- **AggregatingMergeTree:** Aggregates data during merging based on predefined aggregation functions, storing the results in summarized form.

ClickHouse Materialized Views:

- **Triggered on Inserts:** Automatically updates the materialized view whenever new data is inserted into the source table.
- **Real-Time Pre-Aggregation:** Performs pre-aggregation or data transformation in real time, allowing for faster query performance by reducing the need for complex computations at query time.

Ref

- 1. Why ClickHouse Is So Fast**
ChistaDATA Inc. (2023). Why ClickHouse Is So Fast.
<https://chistadata.com/why-clickhouse-is-so-fast>
- 2. How to Use ClickHouse Indexes: A Practical Guide**
ChistaDATA Inc. (2023). How to Use ClickHouse Indexes: A Practical Guide.
<https://chistadata.com/how-to-use-clickhouse-indexes-practical-guide>
- 3. ClickHouse Skipping Indexes**
ClickHouse Documentation. (2023). Skipping Indexes.
<https://clickhouse.com/docs/en/optimize/skipping-indexes>
- 4. ClickHouse Overview**
YouTube. (2023). ClickHouse Overview.
<https://www.youtube.com/watch?v=QDAJTKZT8y4>
- 5. ClickHouse Architecture Deep Dive**
YouTube. (2023). ClickHouse Architecture Deep Dive.
<https://www.youtube.com/watch?v=iLXXoDaFoxs>
- 6. Optimizing ClickHouse Performance**
YouTube. (2023). Optimizing ClickHouse Performance.
<https://www.youtube.com/watch?v=BHcIEszF6Fk>
- 7. Supercharge Your ClickHouse Data Loads - Part 1**
ClickHouse Blog. (2023). Supercharge Your ClickHouse Data Loads - Part 1.
<https://clickhouse.com/blog/supercharge-your-clickhouse-data-loads-part1>

