

ClickHouse

Top Highlights

H2 2022



Reliability

ClickHouse should always work.

Asynchronous Initialization Of Tables

If ZooKeeper is **unavailable** at server startup, the **ReplicatedMergeTree** tables will start in **read-only** mode and **initialize asynchronously** in background as soon as ZooKeeper will be available.

- the same applicable for ClickHouse Keeper as well;
- especially useful for embedded ClickHouse Keeper;

Developers: Antonio Andelic. Since 22.9.

Retries On INSERT

Session expired. Table is in readonly mode 😞

Never again:

```
SET insert_keeper_max_retries = 10;
```

INSERT will survive restarts of ClickHouse Keeper or ZooKeeper and reconnections.

Developer: Igor Nikonov. Since 22.11.

Faster Reconnection to Keeper

Table is in readonly mode 😞

Was: around one minute for reconnection, for no reason.

Now: milliseconds 😊

Developer: Raul Marin. Since 22.10.

No More "Too many parts"

DB::Exception: Too many parts (300). Merges are processing significantly slower than inserts 😞

Now: relaxing the "too many parts" threshold:

Allow large number of parts if they are large in average.

Developer: Alexey Milovidov. Since 22.10.

Flexible Memory Limits

Was: strict per-query limit, **max_memory_usage** = 10 GB by default.

Now: query can use the memory if it's available;
in case of memory shortage, the most "overcommitted" query is terminated
with exception.

Developer: Dmitry Novik.

ClickHouse Keeper

No need to use ZooKeeper anymore!

Always use ClickHouse Keeper instead of ZooKeeper.

Developer: Alexander Sapin, Antonio Andelic.

SQL Compatibility

ClickHouse should support everything you expect.

Window Functions Inside Expressions

```
SELECT
  y::String || '.'
  || (y < toYear(today()) - 2000 - 1 ? '*' : m::String) AS Version,
  (n <= 3 OR (is_lts AND lts_n <= 2)) ? '✓' : '✗' AS Supported
FROM (
  SELECT y, m,
    count() OVER (ORDER BY y DESC, m DESC) AS n,
    m IN (3, 8) AS is_lts,
    countIf(is_lts) OVER (ORDER BY y DESC, m DESC) AS lts_n
  FROM (
    WITH extractGroups(version, 'v(\d+)\.(\d+)') AS v,
         v[1]::INT AS y, v[2]::INT AS m
    SELECT y, m
    FROM file('version_date.tsv', TSV, 'version String, date String')
    ORDER BY y DESC, m DESC
    LIMIT 1 BY y, m)
  )
LIMIT 1 BY Version
FORMAT Markdown
```

Extended Range For Date&Time

Supported range for **DateTime64** and **Date32** data types:

Was: **1925–2283**.

Now: **1900–2300**.

Uses proleptic Gregorian calendar.

Motivation: store dates of birth of the customers.

Developer: Roman Vasin. Since 22.8.

DELETE Query

```
SET mutations_sync = 1;  
ALTER TABLE hits  
DELETE WHERE Title LIKE '%Mongo%';
```

— 205 sec (for a table with 100 million records).

```
DELETE FROM hits  
WHERE Title LIKE '%Mongo%';
```

— ??? sec.

Developers: Alexander Gololobov, Jianmei Zhang.

GROUPING Function

Used with ROLLUP, CUBE or GROUPING SETS.

To distinguish different sets.

```
SELECT k, GROUPING(k) FROM table GROUP BY k WITH ROLLUP
```

Developer: Dmitriy Novik.

Non-Constant LIKE and match

```
SELECT DISTINCT repo_name, title
FROM github_events
WHERE title ILIKE (
    repo_name LIKE '%ClickHouse%' ? '%fast%' : '%slow%')
AND repo_name IN ('ClickHouse/ClickHouse', 'elastic/elasticsearch')
```

Now I can put LIKE inside LIKE and looks like you're going to like it.

Developer: Robert Schulze. Since 22.6.

Composite Time Intervals

Examples:

```
SELECT now() + INTERVAL 1 MONTH;  
SELECT now() + INTERVAL '1 MONTH';  
SELECT now() + INTERVAL 1 MONTH - INTERVAL 2 DAY;  
  
SELECT now() + INTERVAL '1 MONTH -2 DAY';  
SELECT now() + (INTERVAL 1 MONTH - INTERVAL 2 DAY);  
  
SELECT INTERVAL '1 MONTH -2 DAY';  
SELECT (INTERVAL 1 MONTH - INTERVAL 2 DAY);  
  
SELECT INTERVAL '1 MONTH 1 MONTH';
```

Developer: Nikolai Degterinsky. Since 22.11.

Performance

ClickHouse never slows down!

Performance Optimizations

Improvement of ORDER BY, insert and merge in MergeTree, and window functions.

```
SELECT WatchID FROM hits_100m_obfuscated ORDER BY Age
```

Before:

Elapsed: **4.154** sec. (24.07 million rows/s., 216.64 MB/s.)

After:

Elapsed: **0.482** sec. (207.47 million rows/s., 1.87 GB/s.)

Developer: Maksim Kita.

Performance Optimizations

Speed-up of SELECT with **FINAL** modifier.

It "simply" improves performance up to 4 times.

Especially for complex transforms like Collapsing and Replacing.

Developer: Nikita Taranov.

Performance Performance

Optimize **ORDER BY** with **LIMIT**.

Optimize **ORDER BY** with single column.

Optimize **INSERT** into MergeTree with composite **ORDER** key.

Optimize **dictGetChildren**, **dictGetDescendants**.

Optimize cleanup stage of queries with large **GROUP BY**.

Optimize **background CPU usage** of large number of tables.

Developer: Maksim Kita, Nikita Taranov.

More Performance

Optimize **COUNT(DISTINCT ...)** for low number of GROUP BY keys.

Optimize **GROUP BY** with CPU prefetcher.

Optimize **GROUP BY** with better block sizes.

Developer: Nikita Taranov.

New JOIN algorithms

- **"direct"** algorithm:
to join with key-value tables by direct lookups a.k.a. nested loops.
Good if the left table is small, but the right table is like a large dictionary.
Good to use in MATERIALIZED VIEW queries.
- **"parallel_hash"** algorithm:
speed-up if the right hand table is large.
- **"full_sorting_merge"** algorithm:
when right hand side is large
and does not fit in memory and does not allow lookups.
- **"grace_hash"** algorithm:
since in 22.12.

More Performance

... and we are reading from object storage 100 times faster.

Developer: Ksenia Sumarokova.

Updated Benchmark

ClickBench — a Benchmark For Analytical DBMS



Methodology | Reproduce and Validate the Results | Add a System | Report Mistake | Hardware Benchmark

System: All Athena (partitioned) Athena (single) Aurora for MySQL Aurora for PostgreSQL ByteHouse Citus clickhouse-local (partitioned) clickhouse-local (single) ClickHouse ClickHouse (zstd) CrateDB Databend Druid DuckDB Elasticsearch Greenplum HeavyAI Infobright MariaDB ColumnStore MariaDB MonetDB MySQL (MyISAM) MySQL Pinot PostgreSQL QuestDB Redshift SingleStore Snowflake SQLite StarRocks (tuned) StarRocks TimescaleDB (compression) TimescaleDB

Type: All stateless managed Java column-oriented C++ MySQL compatible row-oriented C PostgreSQL compatible ClickHouse derivative embedded Rust search time-series

Machine: All serverless 16acu L M S XS c6a.4xlarge, 500gb gp2 c6a.metal, 500gb gp2 c6a.4xlarge, 1500gb gp2 ra3.4xlarge ra3.xplus S24 S2 2XL 3XL 4XL XL

Cluster size: All 1 2 4 8 12 16 32 64 128 serverless undefined

Metric: Cold Run Hot Run Load Time Storage Size

System & Machine	Relative time (lower is better)
ClickHouse (c6a.metal, 500gb gp2):	×1.41
SingleStore (12×S24) [†] :	×2.08
Snowflake (64×3XL):	×3.04
Snowflake (32×2XL):	×3.16
Snowflake (128×4XL):	×3.36
StarRocks (tuned) (c6a.4xlarge, 500gb gp2):	×3.61
Snowflake (16×XL):	×3.70
ClickHouse (c6a.4xlarge, 500gb gp2):	×3.84
ClickHouse (zstd) (c6a.4xlarge, 500gb gp2):	×4.06

Integrations

ClickHouse integrates with everything!

Integrations

ClickHouse can work **as a server** (clickhouse-server) or **as a tool** without installation (clickhouse-local).

ClickHouse can **store the data** or process externally stored data **on the fly**.

External data:

- remote databases: MySQL, PostgreSQL, MongoDB, ODBC, JDBC...
- object storages: S3, HDFS, Azure, COSN, OSS...
- from URL and local files;

All possible data formats:

- text: CSV, TSV, JSON, Values, MySQLDump, Regexp...
- binary: Parquet, Arrow, ORC, Avro, Protobuf, MsgPack...
- schemaful and schemaless;

Data Lakes

Now ClickHouse supports **Apache Hudi** and **Delta Lake** for SELECT queries.

TODO: Apache Iceberg.

Advantages:

- these formats are somewhat resembling **MergeTree** allowing incremental data insertion, approaching to ClickHouse data formats;

Disadvantages:

- alien data formats from **Apache/Hadoop/Java** world;
- nothing works out of the box unless you really know how to deal with it;

Data Lakes

Now ClickHouse supports **Apache Hudi** and **Delta Lake** for SELECT queries.

```
SELECT count() FROM deltaLake(  
    'https://clickhouse-public-datasets.s3.amazonaws.com/delta_lake/hits/')  
WHERE URL LIKE '%google%'  
  
-- 4.396 sec.
```

Developers: Daniil Rubin, Ksenia Sumarokova, Flynn ucasfl. Since 22.11.

Integrations

- Querying MongoDB and Meilisearch with table functions.
- Streaming data consumption from NATS.

Developer: Anastasia Petrenko, Ksenia Sumarokova. Since 22.7.

Integrations

Visualizations:

- official ClickHouse plugin for **Grafana**;
- official support for **Superset**;
- **HEX** and **Deepnote** support.

Data ingestion and processing:

- **Kafka Connect** integration;
- **Airflow**, **dbt** support.

Language drivers:

- official **Node.JS** driver;
- optimized **Go** driver;
- a new **Python** client.

Operations

ClickHouse is easy to configure for your needs.

Self-Extracting Executable

The most simple way to install ClickHouse:

```
curl https://clickhouse.com/ | sh
```

Single binary package. Installs the latest version. Includes debug info.

Works on every **Linux** (x86_64, aarch64, powerpc64le),
macOS (x86_64, M1), **FreeBSD** and **Windows** (WSL2).

Was: **2.1 GB**.

Now: **446 MB**, takes ~5 seconds to decompress on first run.

Developer: Arthur Filatenkov, Yakov Olkhovskiy.

Composable Protocols

So, ClickHouse supports a lot of protocols:

- HTTP
- HTTPs
- Native TCP
- Native TCP wrapped in PROXYv1
- Native TCP with TLS
- MySQL (with TLS support)
- PostgreSQL (with TLS support)
- GRPC (with TLS)
- Replication protocol over HTTP
- Replication protocol over HTTPs
- Keeper client-server protocol;
- Keeper consensus protocol;
- ...

Composable Protocols

So, ClickHouse supports a lot of protocols.

How to configure all of them? What if:

- server has multiple network interfaces?
- enable one protocol on multiple ports?
- I want native TCP for localhost only and HTTPs from everywhere?
- I want different TLS certificates for different protocols?
- I want to wrap one protocol in another?

```
<protocols>
  <tcp>
    <type>tcp</type>
    <host>::</host>
    <port>9000</port>
    <description>native protocol</description>
  </tcp>
  <tcp_secure>
    <type>tls</type>
    <impl>tcp</impl>
    <port>9440</port>
    <description>secure native protocol</description>
  </tcp_secure>
  <tcp_endpoint>
    <impl>tcp</impl>
    <host>0.0.0.0</host>
    <port>9001</port>
    <description>native protocol, another</description>
  </tcp_endpoint>
  <tcp_proxy>
    <type>proxy1</type>
    <impl>tcp</impl>
    <port>9100</port>
    <description>native protocol with PROXYv1</description>
  </tcp_proxy>
```

Composable Protocols

The case: ClickHouse under proxy:

Envoy Proxy / HAProxy / CloudFlare.

ClickHouse server will receive connections from the proxy.

But it needs to know the source IP address for quotas, ACL and logging.

Solution: enable PROXYv1 protocol in the proxy
and configure it as a protocol wrapper in ClickHouse.

ClickHouse will read the header and unwrap the network packets.

Developer: Yakov Olkhovskiy. Since 22.10.



Compliance



✓ CCPA



✓ GDPR



Security
Standards Council

✓ PCI DSS



✓ SOC 2

<https://trust.clickhouse.com/>

+ penetration testing, bug bounty program, audit reports...

ClickHouse Cloud Beta

— available since Oct 4th;

— free 14-day trial up to 10 TB of data;

Try it! <https://clickhouse.cloud/>.

Q&A

