



Read from object storage 100 times faster

Kseniia Sumarokova, ClickHouse Core SWE



ClickHouse: Compute and Storage

Compute – blazing fast.

Storage – depends...

- Local SSD, HDD, NBS
- HTTP server with static files
- Object storages:
 - AWS S3
 - MinIO
 - Google Cloud Storage
 - R2 (Cloudflare)
 - Azure Blob Storage




ClickHouse: Compute and Storage

Compute – blazing fast.

Storage – depends...

- Local SSD, HDD, NBS
- Remote HTTP server
- **Object storages:**
 - AWS S3
 - MinIO
 - Google Cloud Storage
 - R2 (Cloudflare)
 - Azure Blob Storage



What is Object Storage?

- A distributed storage
- Usually provides HTTP API
- Stores blobs with some path (URL)
- Blobs are immutable
- Storage has flat structure
- Very different from traditional filesystems:
 - No seeks
 - No hardlinks
 - No appends
 - etc



Object Storage: Pros and Cons


- + Infinite data storage
- + Highly reliable
- + Scales automatically
- + Storage is cheap
- + High throughput
- Big latency
- Requests (read, writes) are not free
- API is limited compared to FS

Towards Object Storage support in MergeTree engine

- ✓ Design FS abstraction
- ✓ Make MergeTree work with this abstraction
- ✓ Implement the abstraction for Object Storages
- ? Optimize... as much as possible
 - Use benefits of the API
 - Utilize high throughput
 - Reduce high latency

Reading data in MergeTree: the basics

- Column oriented storage
- Index is sparse (fits in memory)
- Columns are compressed
- Marks allow to read compressed data according to index
- Marks are located in .mrk file for each column
- Data is located in .bin file
- Granule is a range of rows between two marks
- Read data by granules
- Read different granules in parallel



Reading data in MergeTree


primary.idx

OrderID BannerID

| | |
|-------|-------|
| 0 | 10000 |
| 8192 | 18192 |
| 16384 | 26384 |
| ... | |


| | |
|---------|---------|
| 1998848 | 2008848 |
|---------|---------|

Sparse index




| OrderID | BannerID |
|------------|-----------------------|
| 3861813205 | n#0=n)SNW%-(?clG%[hN |
| 1356609667 | ui5vptJlh+mKS]1}eLL& |
| 3265981873 | [D/,i5G\6F5hP=i]<qI |
| 2267935489 | Hk.Qo78?jM9,,\ -M,Rj^ |
| 1150892606 | g&C]A&@2U/b- BU,c;d% |
| 2836790991 | lfGmuuauNz"He0jH1unN |
| 2749022404 | 7AZ&5LY`X\ydCAhT9LJh |
| 2393185453 | \T`6R-?q}:6CtQJ2f9rS |
| 1211724375 | `=6P&_ieD^ 1: <?s?(M |
| 3109046926 | Y&B1I\N'@C.A/4/a"Y'Ia |

Reading data in MergeTree: in Parallel




Reading data in MergeTree

- Each thread reads a non-intersecting set of granule(s).
- Each granule (in the picture: red, blue or green) is a range of rows between two marks and includes multiple columns
- Columns in a granule are read **one by one**



What if the number of columns is huge?

- Huge number of columns + big latency results in too slow query execution



Optimization #1: Prefetch all the columns

- In each query thread for each column in granules make asynchronous request to read the data
- Optimization #2: Using Object Storage range requests if supported



What can be improved:

- Marks are still loaded synchronously
- Only max_threads workers can read the granules and prefetch the data


Optimization #3: Asynchronous Loading of Marks

- Mark cache already exists but size is limited
- Preload all marks for query into cache on query start
- Use separate pool for marks preload


Dealing with limited number of query threads




Dealing with limited number of query threads




Read tasks:





Dealing with limited number of query threads




Read tasks:





Dealing with limited number of query threads




Read tasks:



Start prefetching for all read tasks



Read tasks:



Optimization #4: Prefetch all the tasks

- Read task consists of:
 - Part to read
 - Marks ranges within part to read
- Can prefetch them independently from query threads:
 - Start asynchronously with the start of query
 - Use single threadpool per server with a fixed number of threads
 - Additionally limit the number of concurrent prefetches (avoid OOM)

Optimization #5: Caching on local disk

- Cache has limited size
- Cache consists of file segments - limited sized segments of files
- Removal from cache follows eviction policy
- Caching on reads, writes, background merges
- Files are immutable, so cached file segments validation is not required.
- On file deletion, corresponding cache files are also removed.

Performance comparison

```
:) CREATE TABLE test_s3 (p UInt16, s String, PRIMARY KEY s) ENGINE=MergeTree() PARTITION BY p SETTINGS storage_policy='s3_cache';
```

```
:) INSERT INTO test SELECT number DIV 100000, toString(number) FROM numbers(100000000);
```

Performance comparison

:) SET enable_filesystem_cache=0, prefer_prefetched_read_pool=0, remote_filesystem_read_method='read', max_threads='8'

ip-172-31-4-213.eu-west-1.compute.internal :) SELECT count() FROM test_s3 WHERE s IN (SELECT toString(arrayJoin(range(1000)) * 100000 + 55555));

```
SELECT count()
FROM test_s3
WHERE s IN (
  SELECT toString((arrayJoin(range(1000)) * 100000) + 55555)
)
```

Query id: 52c954b7-e897-47c4-8db0-59ba5f9a028f

| |
|-----------------|
| count() 1000 |
|-----------------|

1 row in set. Elapsed: 16.804 sec. Processed 15.88 million rows, 260.41 MB (945.24 thousand rows/s., 15.50 MB/s.)



Performance comparison

:) SET enable_filesystem_cache=0, prefer_prefetched_read_pool=1, remote_filesystem_read_method='threadpool', max_threads='8'

ip-172-31-4-213.eu-west-1.compute.internal :) SELECT count() FROM test_s3 WHERE s IN (SELECT toString(arrayJoin(range(1000)) * 100000 + 55555));

```
SELECT count()
FROM test_s3
WHERE s IN (
  SELECT toString((arrayJoin(range(1000)) * 100000) + 55555)
)
```

Query id: 87e7c909-5358-4293-9040-284d73a75036

| |
|---------|
| count() |
| 1000 |

1 row in set. Elapsed: 2.792 sec. Processed 15.88 million rows, 260.41 MB (5.69 million rows/s., 93.26 MB/s.)



Thanks for attention!

Questions?