

Self-repairing processing using ClickHouse for unreliable data sources

Ingmar Poesse

05/12/2022

About BENOCS

Location: Berlin, Germany

Industry: Telecommunications,
Network Telemetry

Mission: To improve network traffic
conditions and visibility



2013

Founded as a Spin-off of Deutsche
Telekom

2015

First live customers for BENOCS
Director

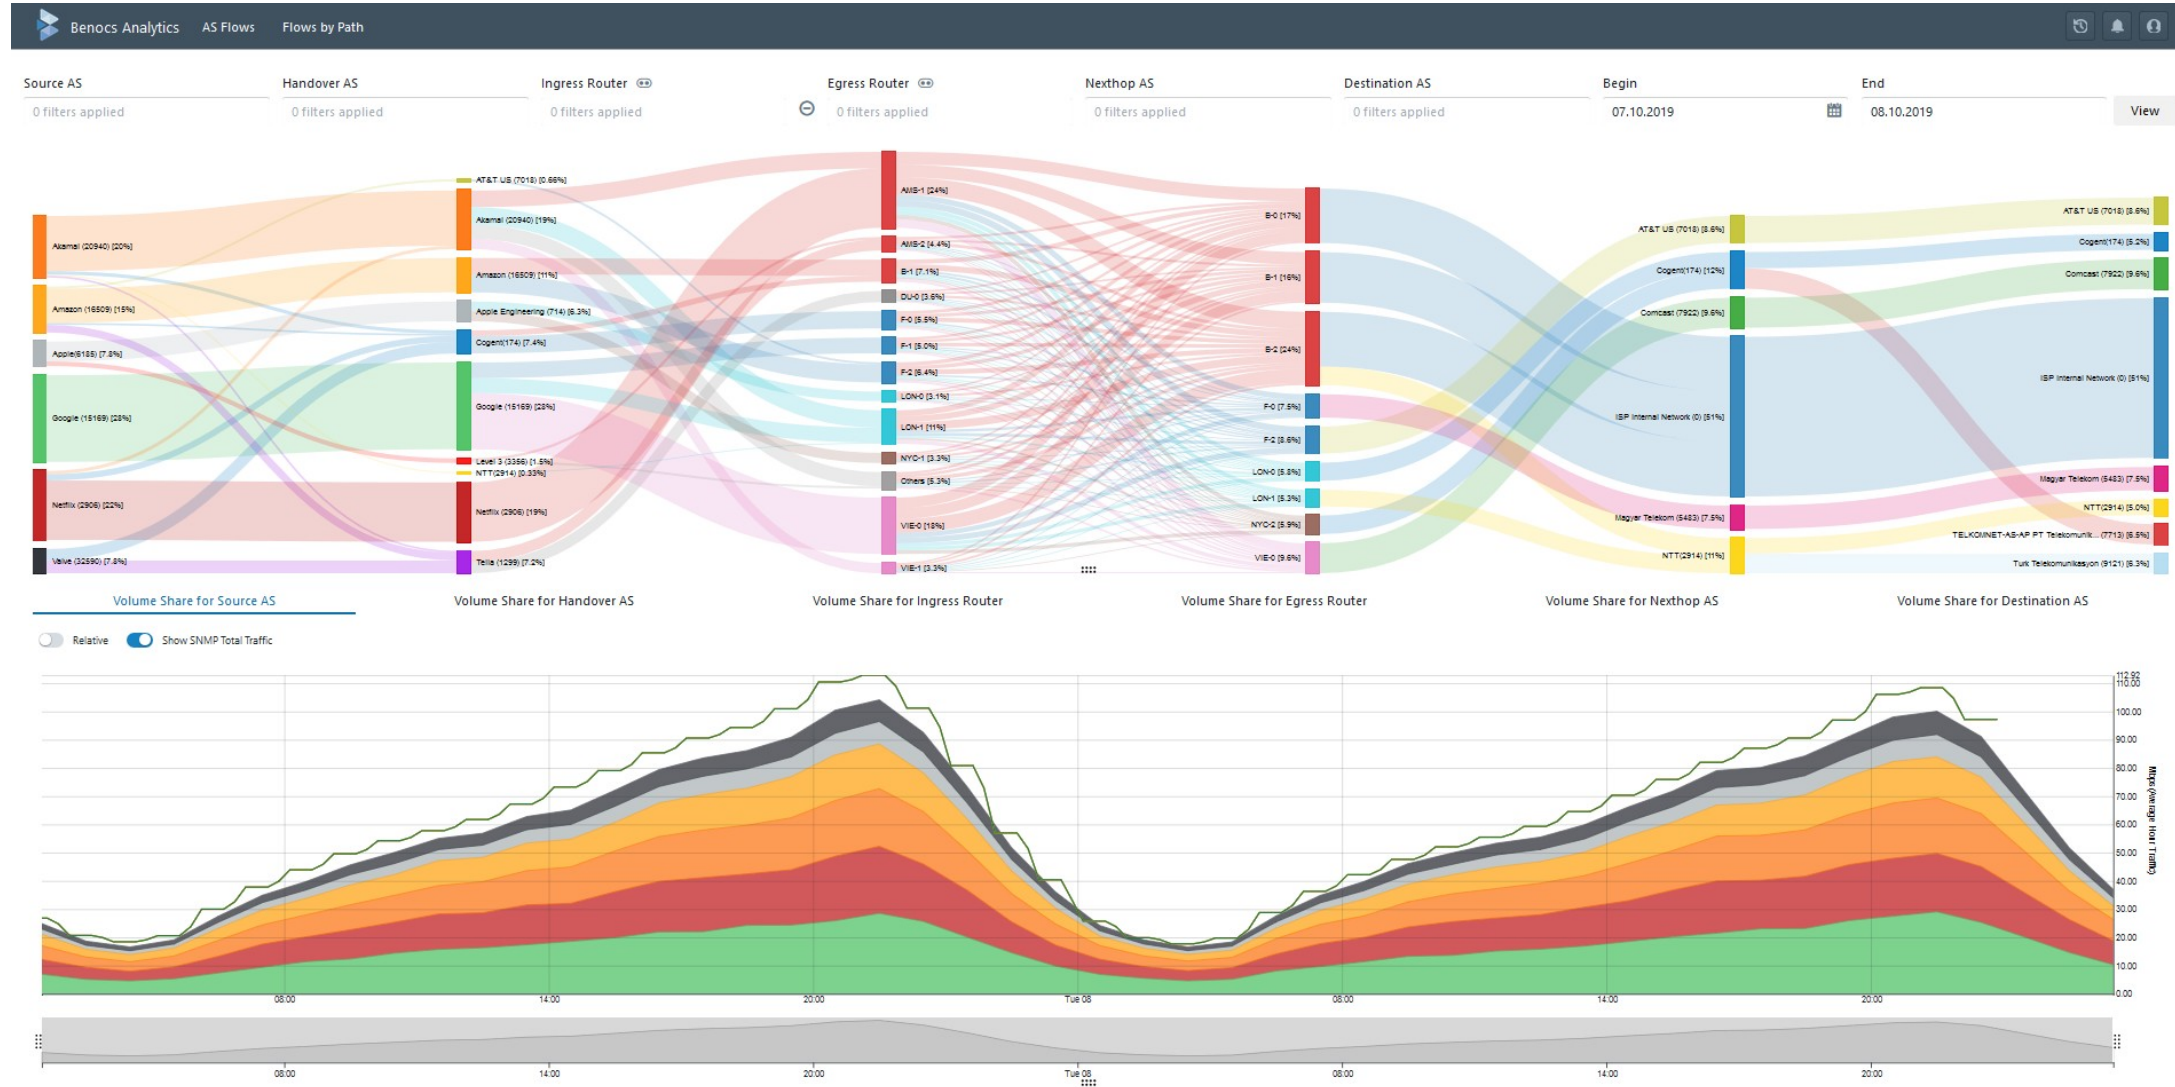
2018

Fully developed BENOCS Analytics
enters the market

2020

Customer base expanded globally

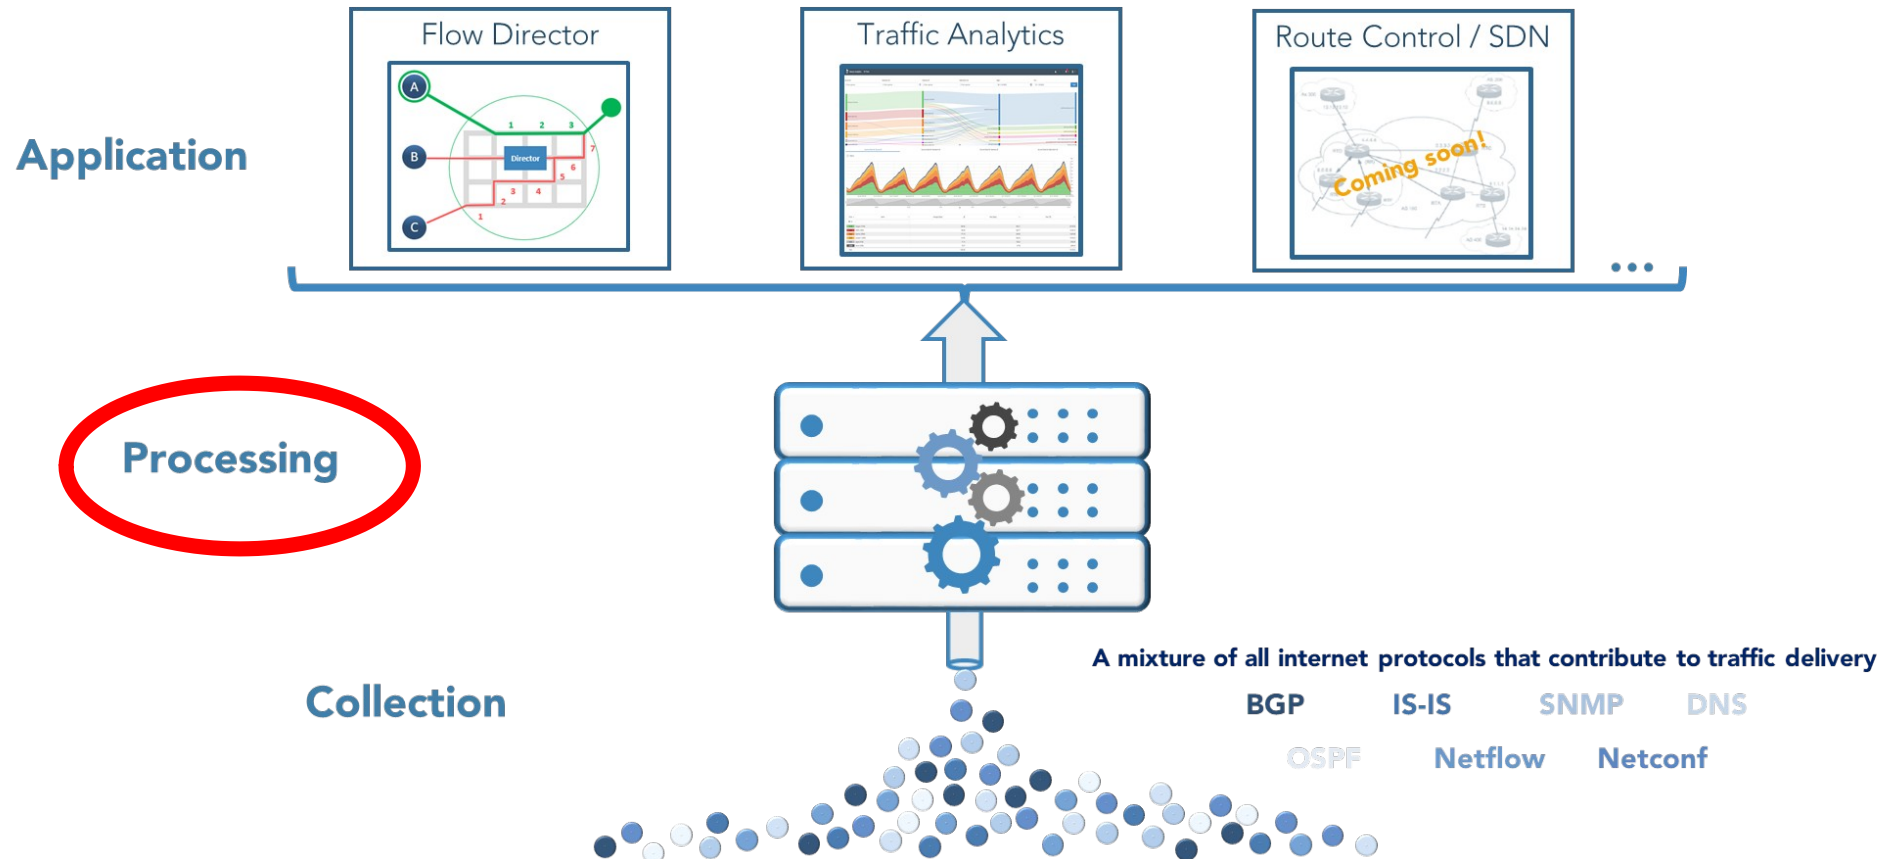
Visualizes end-to-end traffic at one glance



This is the BENOCs solution

Products:

Our products are made from our Core Engine meant to target efficient internet-traffic delivery.



The challenge in (our) data processing

No control over data sanity:

Expect data to be partially tainted, useless, broken, wrong or (maybe) malicious.

No Control over delivery timing:

Data will arrive whenever (Bursts, constant delays, intermittent failures),

most of the time

without any guarantees what is happening when

Multiple data sources:

Each source at each customer may have its own delivery, timing and failure/recovery model

All analysis is interconnected, single sources are (almost) useless:

Data cannot be processed until multiple/all data sources are ready – or when there is evidence it will not come

Installations at customer premises (SaaS):

We manage everything, but run in a highly distributed, fully separated, partially internet-less environment.

Our Goals

Clear dependency graph

Define a graph of **Targets** where each table has a set of requirements (predecessors) and a set of dependencies (successors).

Easy adaption of intermediate processing steps

Each **Target** needs to be fully interchangeable without touching/modifying any of the other Targets.

Statelessness

All state is kept in clickhouse. The Targets have no state and re-synchronize with the DB every time they run.

Easy re-processing of data or *Stuff* fails - all the time !

A **Target** will always process all data that is eligible for processing. No matter if new or old.

This allows for re-processing as well as normal time progression



SyncClickhouse (working title)

Definition: Target

A **Target** is a specific processing step in the graph – usually bound to a single table. A **Target**

- ▷ can have none to multiple requirements (i.e. target that have to finish beforehand)
- ▷ defines exactly one **Name** that other **Targets** can depend on
- ▷ can be bound to one table in clickhouse. This table will be maintained through the **Target** (create, update, alter...)
- ▷ Defines two stages for processing
 - 1) Tokens
 - 1) Tokens uniquely identify non-overlapping parts of a table (usually time ranges/timestamps) that are eligible for processing. This is done through a single SQL statement.
 - 2) The SQL statement can consider data from an external source (i.e. files from a disk)
 - 2) Processing
 - 1) Each Token is processed independently and in non-deterministic order (i.e. partially parallel) by spawning a templated query (usually an *INSERT* statement).
- ▷ There is a lot more Targets can do (multi-processing, external Commands, Query breakdown, etc)

Target Example:

```
[Target]
  require = (init)
  provide = Productive.version

[Target.Productive.version]
  tableName = Productive.version
  template = %General.templatePath%/Productive/version.sql

[external]
  command = grep %Deploy.srcDatabaseSchemaDB% %General.versionFile% | grep -v "\." | awk '{print $2}'
  structure = version UInt32

[Token]
  external = STDIN
  query = <<EOT
SELECT version FROM
(
  SELECT DISTINCT version
  FROM STDIN
)
ALL LEFT JOIN
(
  SELECT DISTINCT
    version,
    1 as exists
  FROM %Target.Productive.version.tableName%
) using version
where exists = 0
EOT

# all settings related to inserting the missing timestamps
[Process]
  query=insert into %Target.Productive.version.tableName% VALUES(0, __RAW__)
  optimize = 1
```

Table names are independent of the Target

External Command becomes a table

All queries are templated and can be change at run time

Each Row returned from Tokens will become a substitute in the Process section.


Each row will spawn it's own, independent statement.

Unintended work-around

Problem:

We have a set of customer specific RegEx and a table those RegEx are to be applied to.

If Desc

| timestamp | Router | Interface | IfDesc | | Result |
|---------------------|--------|-----------|-------------------|---|--------|
| 2022-12-05 14:53:46 | R1 | 12 | Customer C1: 1299 |  | 1299 |
| 2022-12-05 14:53:46 | R1 | 14 | BB Links to R2 | | remove |
| 2022-12-05 14:53:46 | R2 | 452 | BB Link to R1 | | remove |
| 2022-12-05 14:53:46 | R2 | 12 | Upstream: 3320 | | 3320 |
| 2022-12-05 14:53:46 | R2 | 25 | Google GGC | | 15169 |

Rules

- Customer\s*.*?\s*:\s*(\d+) → use match group as ASN
- Upstream:\s*(\d+) → use match group as ASN
- Google\s+(GGC) → map GGC to 15169

Token Query Example:

```

SELECT DISTINCT
    Timestamp, regEx, index
FROM
(
    SELECT DISTINCT
        Timestamp, regEx, index
    FROM
    (
        SELECT DISTINCT
            timestamp,
            %Target.Mangle.linkVRFOverwrite.regEx% as originalRegExArr,
            arrayPushFront(originalRegExArr, '') as regExArr,
            arrayPushFront(arrayEnumerate(originalRegExArr),0 ) as indexArr
        FROM %Target.Data.DMIPKPI.tableName%
        WHERE timestamp >= '___MIN_TS__'
    )
    ALL INNER JOIN
    (
        SELECT DISTINCT
            timestamp
        FROM %Target.Data.DMIPGeneric.tableName%
        WHERE timestamp >= '___MIN_TS__'
    )
    USING timestamp
    ARRAY JOIN
        regExArr as regEx,
        indexArr as index
    UNION ALL
    SELECT DISTINCT
        Timestamp, '' as regEx, 0 as index
    FROM %Target.Data.DMIPKPI.tableName%
    WHERE timestamp >= '___MIN_TS__' AND KPI = 'VRF'
)
ALL LEFT JOIN
(
    SELECT DISTINCT
        timestamp,
        1 as exists
    FROM %Target.Mangle.linkVRFOverwrite.tableName%
) USING timestamp
where exists = 0 AND
((length(%Target.Mangle.linkVRFOverwrite.regEx%) > 0) or ((select count() from %Target.Summary.VRFASMapping.tableName%) > 0))
ORDER BY timestamp asc

```

Summary

Problem:

Work with unreliable, incomplete data from diverse data sources at distributed sites with limited internet connectivity.

Our Solution:

Build a pipeline around clickhouse that manages the data and uses the DBMS as its processing engine.

- ▷ Automatically takes care of dependencies, scheduling, DB consistency, (re)-processing, etc.
- ▷ Completely stateless – all state is stored implicitly in clickhouse
- ▷ Data Driven – process data when all sources have reported in. No external synchronization.

But... Is there a better way to do this ?

Thank you !

Backup

Unintended work-around

Problem:

We have a set of customer specific RegEx and a table those RegEx are to applied to.

Match an array of RegEx against a column with

```
multiMatchAny(haystack, [pattern1, pattern2, ..., pattern])
```

Find which RegEx matched

```
multiMatchAnyIndex(haystack, [pattern1, pattern2, ..., pattern])
```

```
multiMatchAllIndex(haystack, [pattern1, pattern2, ..., pattern])
```

But – there is no function to extract the match groups from an dynamic/array of patterns

extract()* functions need constant patterns !

We can work around this by spawning one **Token** per timestamp/regEx and passing this via the config as a constant into clickhouse.

This spawns (potentially) a lot of queries ($O(nm)$, $n=timestamps$, $m=RegExs$).